



香港中文大學(深圳)
The Chinese University of Hong Kong, Shenzhen



Ack: Prof. Jignesh Patel @ CMU
Prof. Andy Pavlo @CMU

CSC3170

4: Database Storage

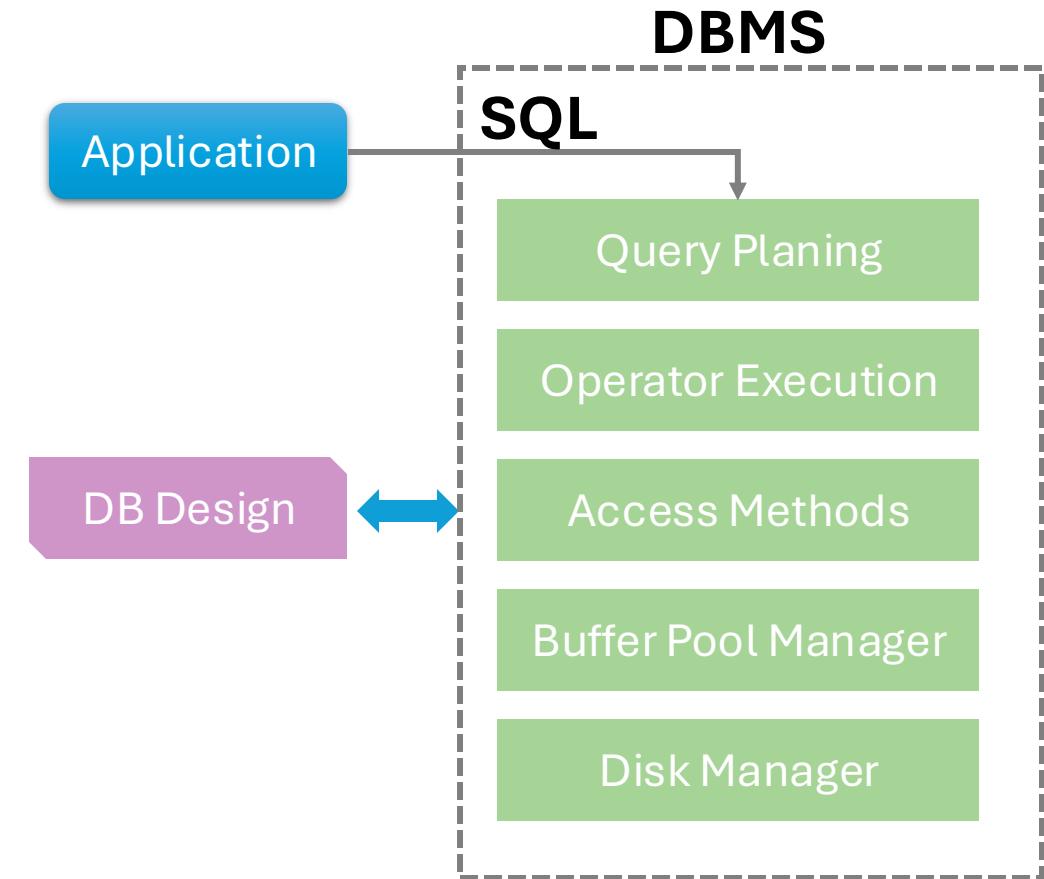
Chenhao Ma

School of Data Science

The Chinese University of Hong Kong, Shenzhen

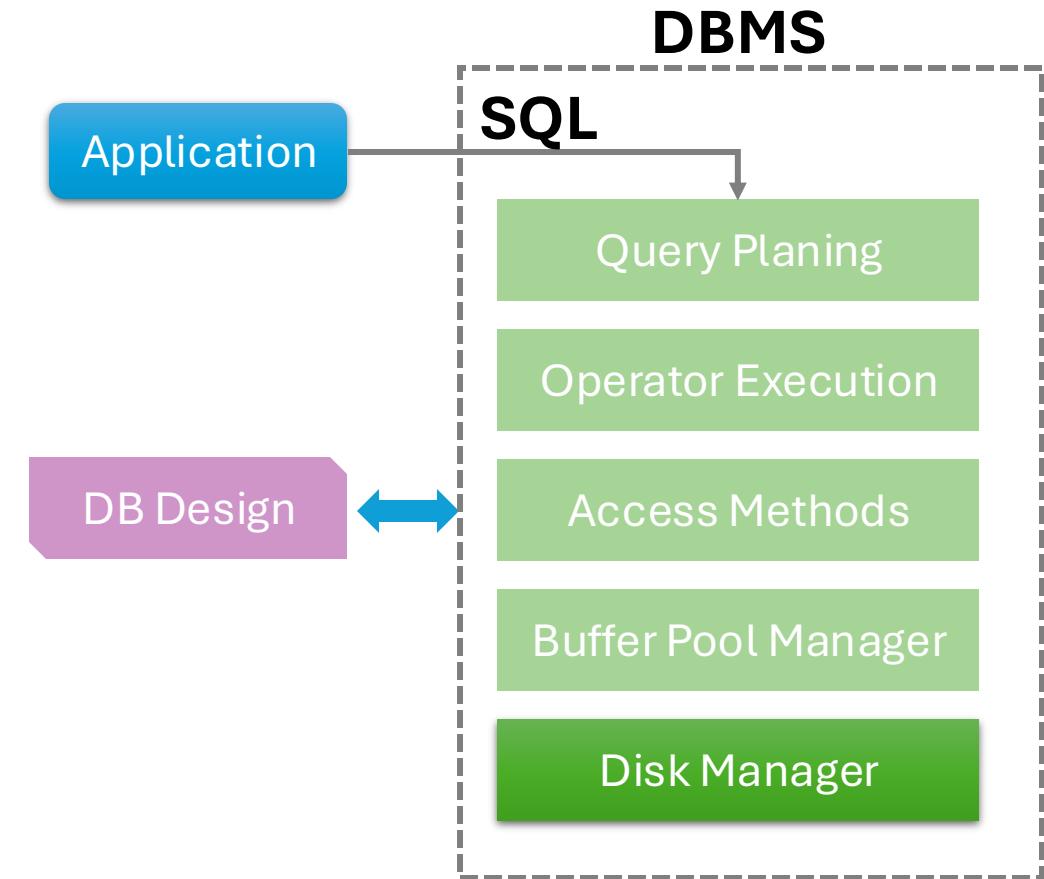
Previous Lectures

- We now understand what a database looks like at a logical level and how to write queries to read/write data (e.g., using SQL).
- We will next learn how to build software that manages a database (i.e., a DBMS).



Previous Lectures

- We now understand what a database looks like at a logical level and how to write queries to read/write data (e.g., using SQL).
- We will next learn how to build software that manages a database (i.e., a DBMS).

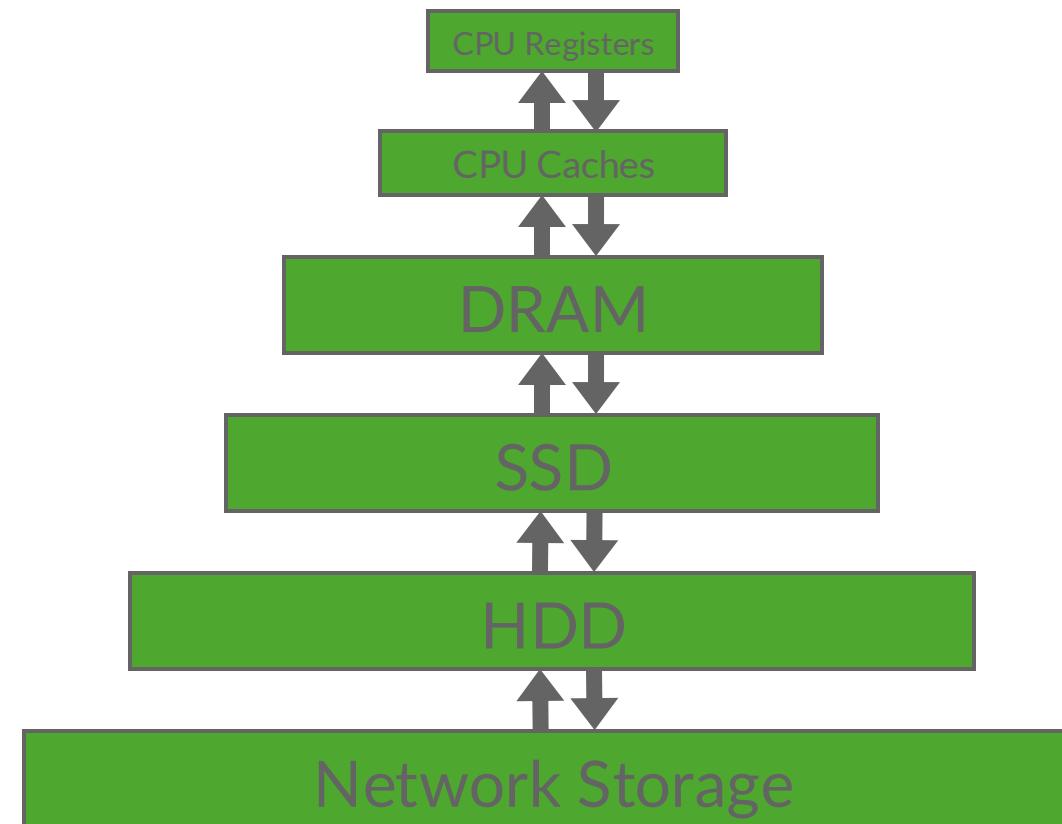


Disk-based Architecture

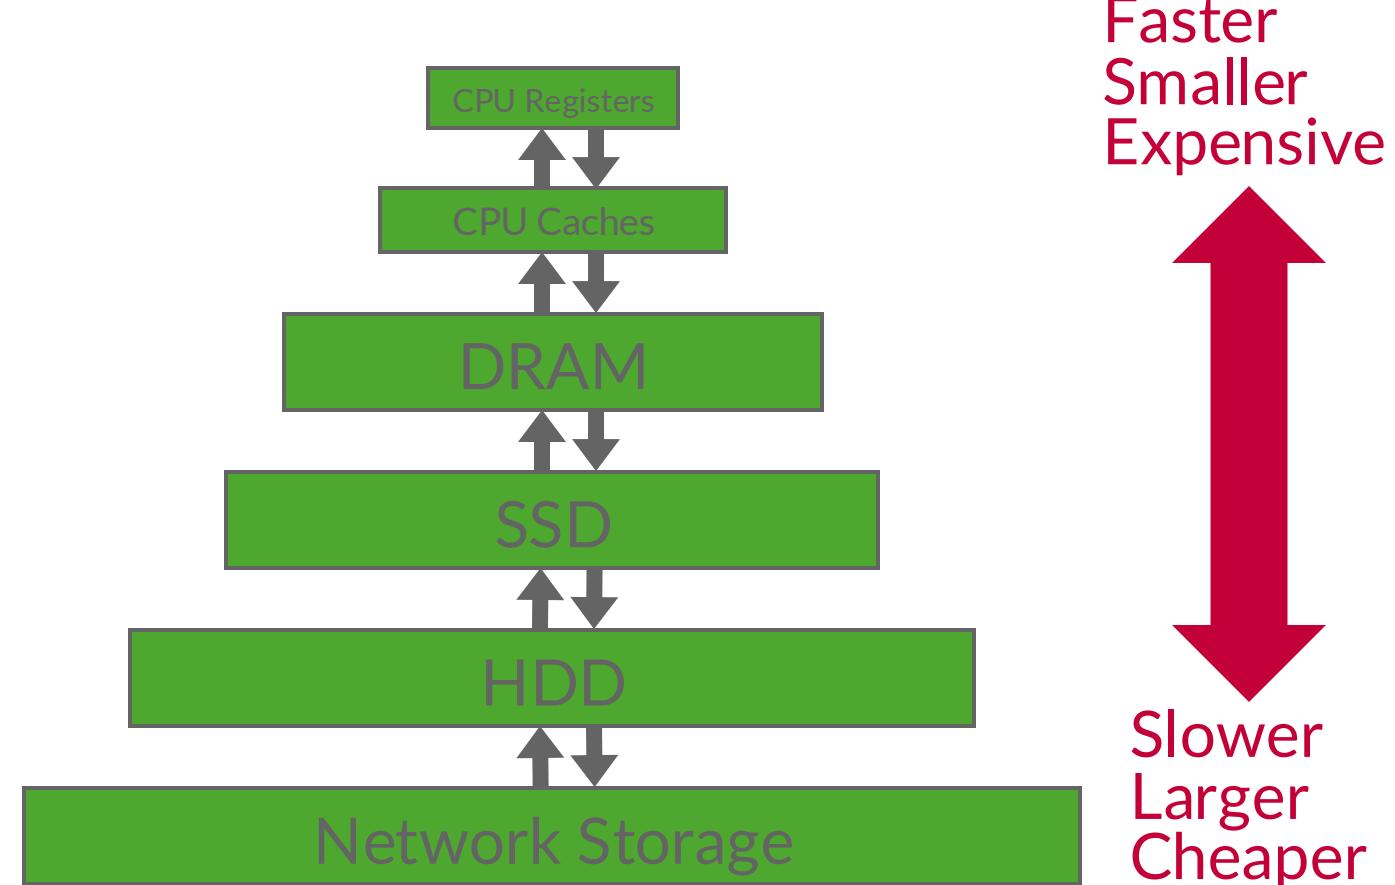
Disk-based Architecture

- The DBMS assumes that the primary storage location of the database is on non-volatile **disk**.
- The DBMS's components manage the movement of data between **non-volatile and volatile storage**.

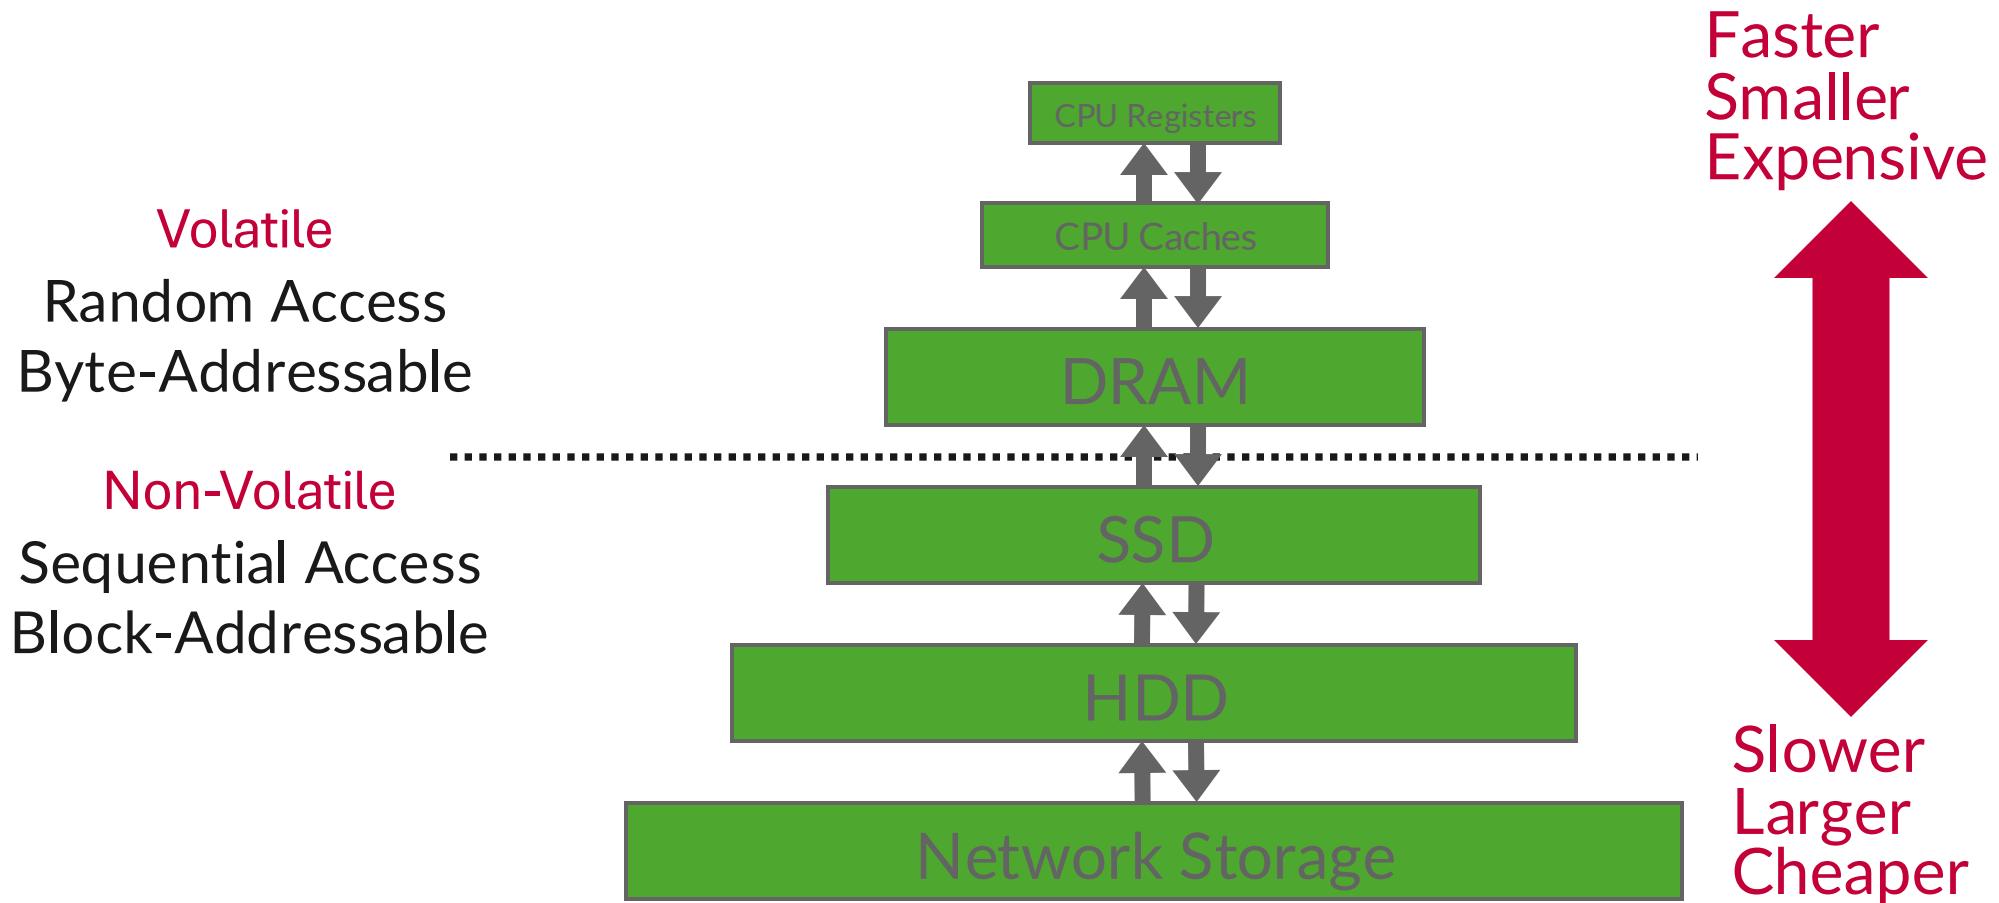
Storage Hierarchy



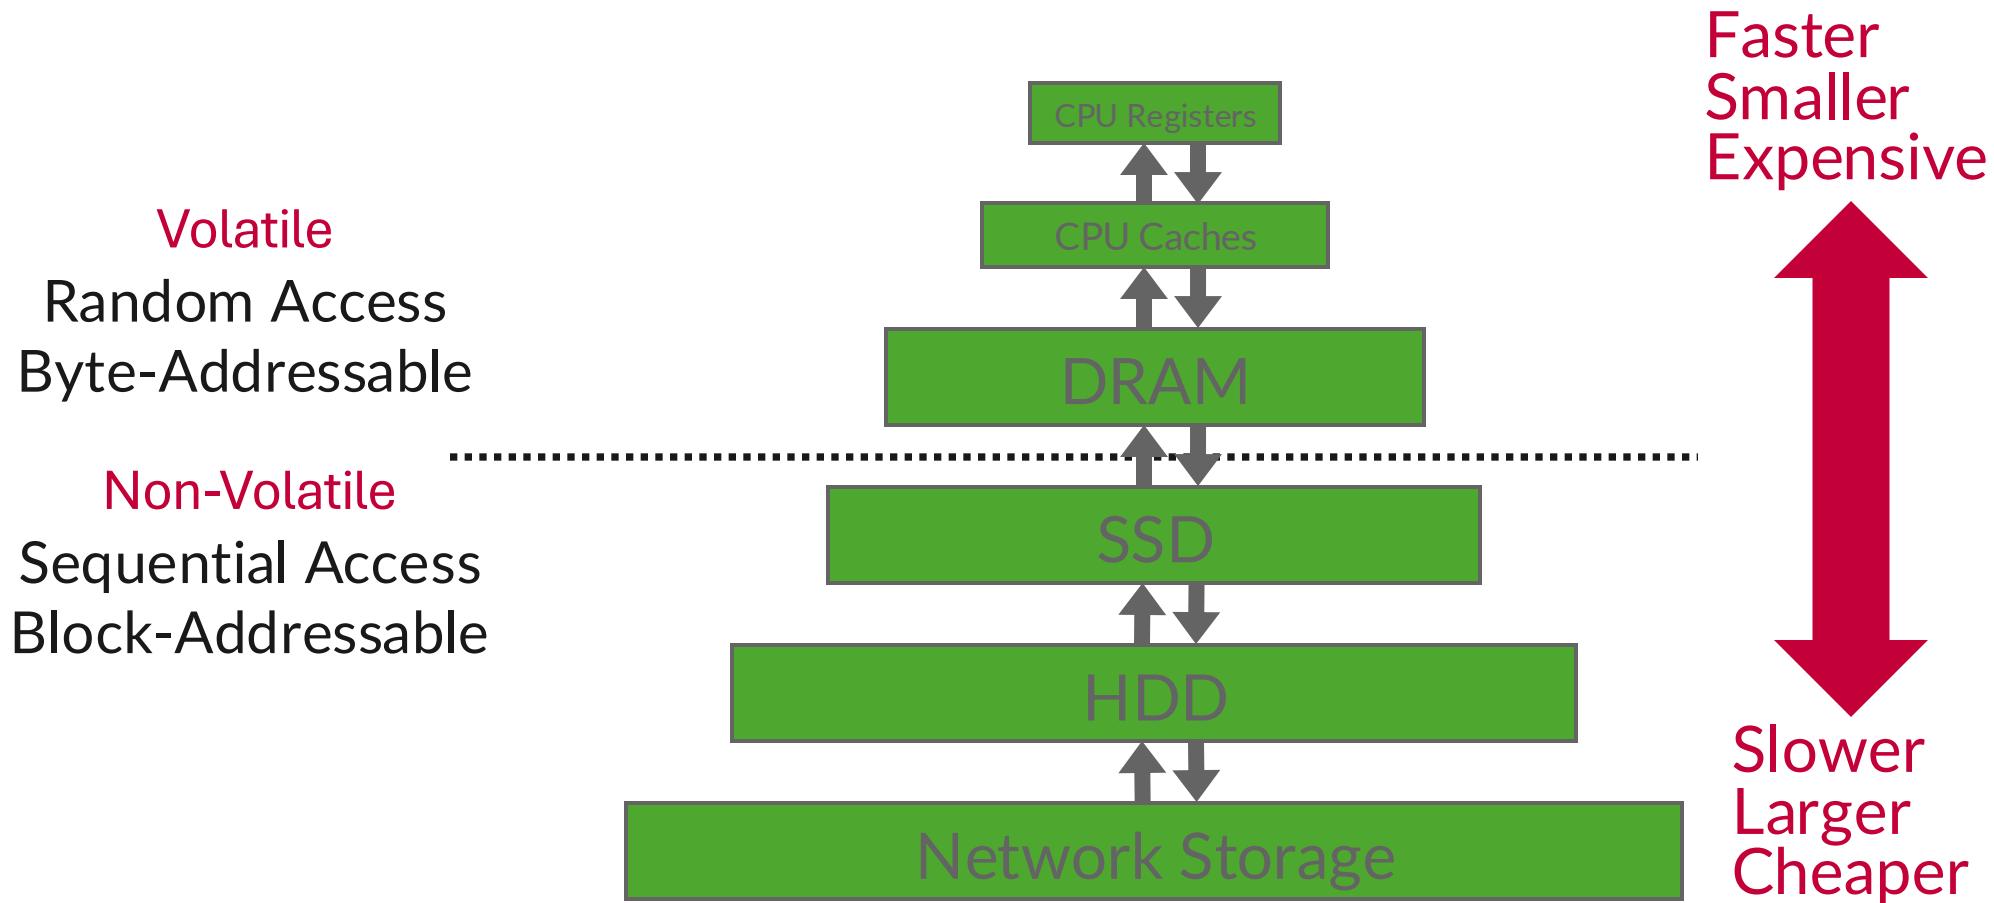
Storage Hierarchy



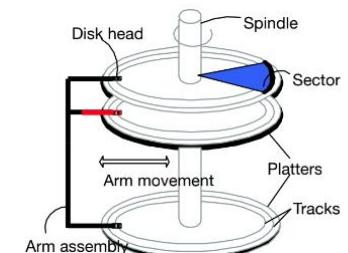
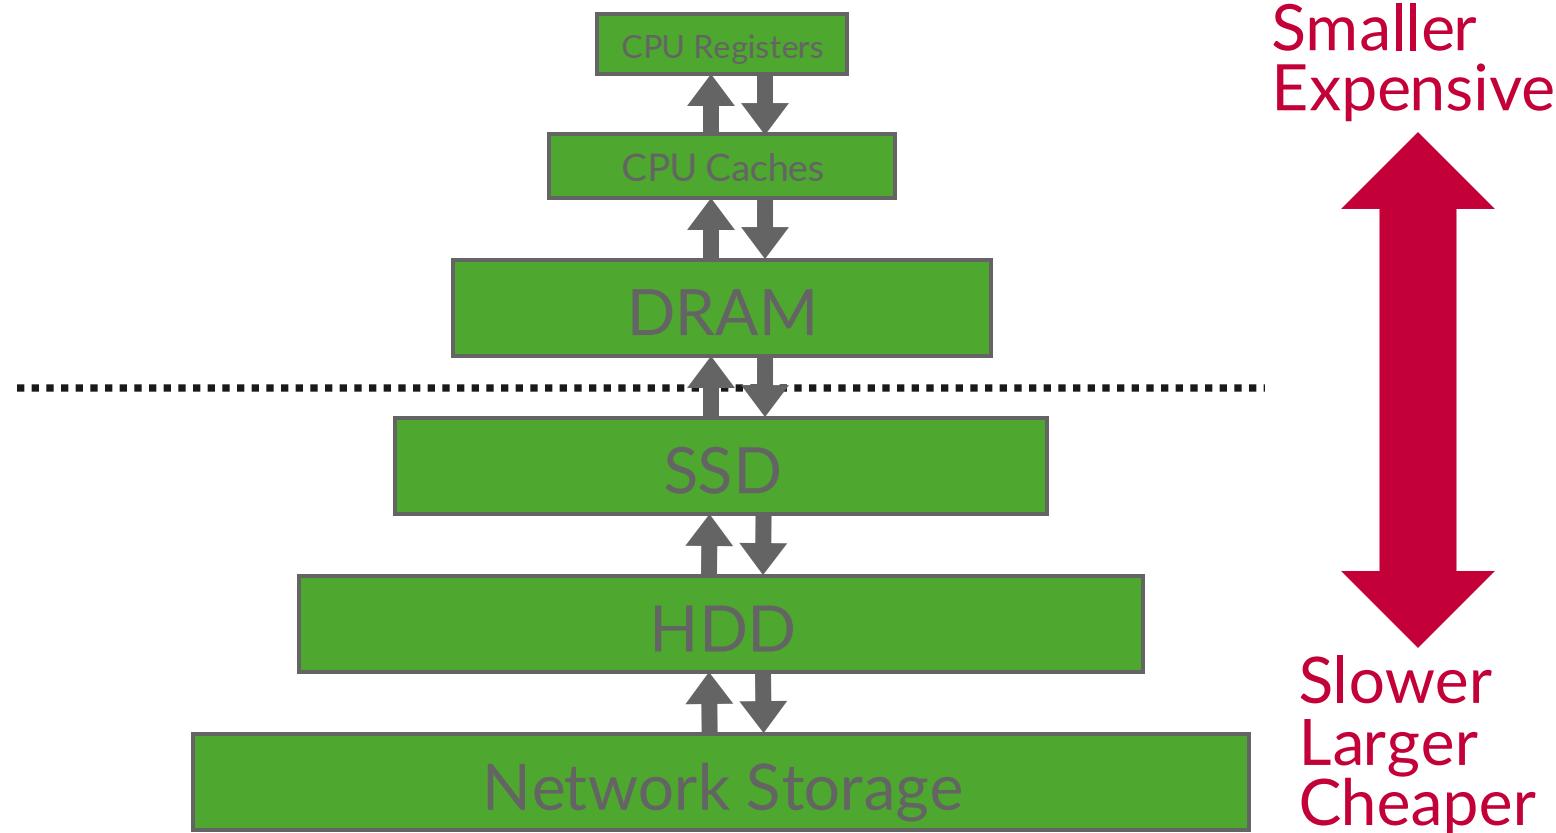
Storage Hierarchy



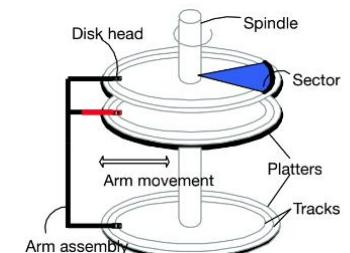
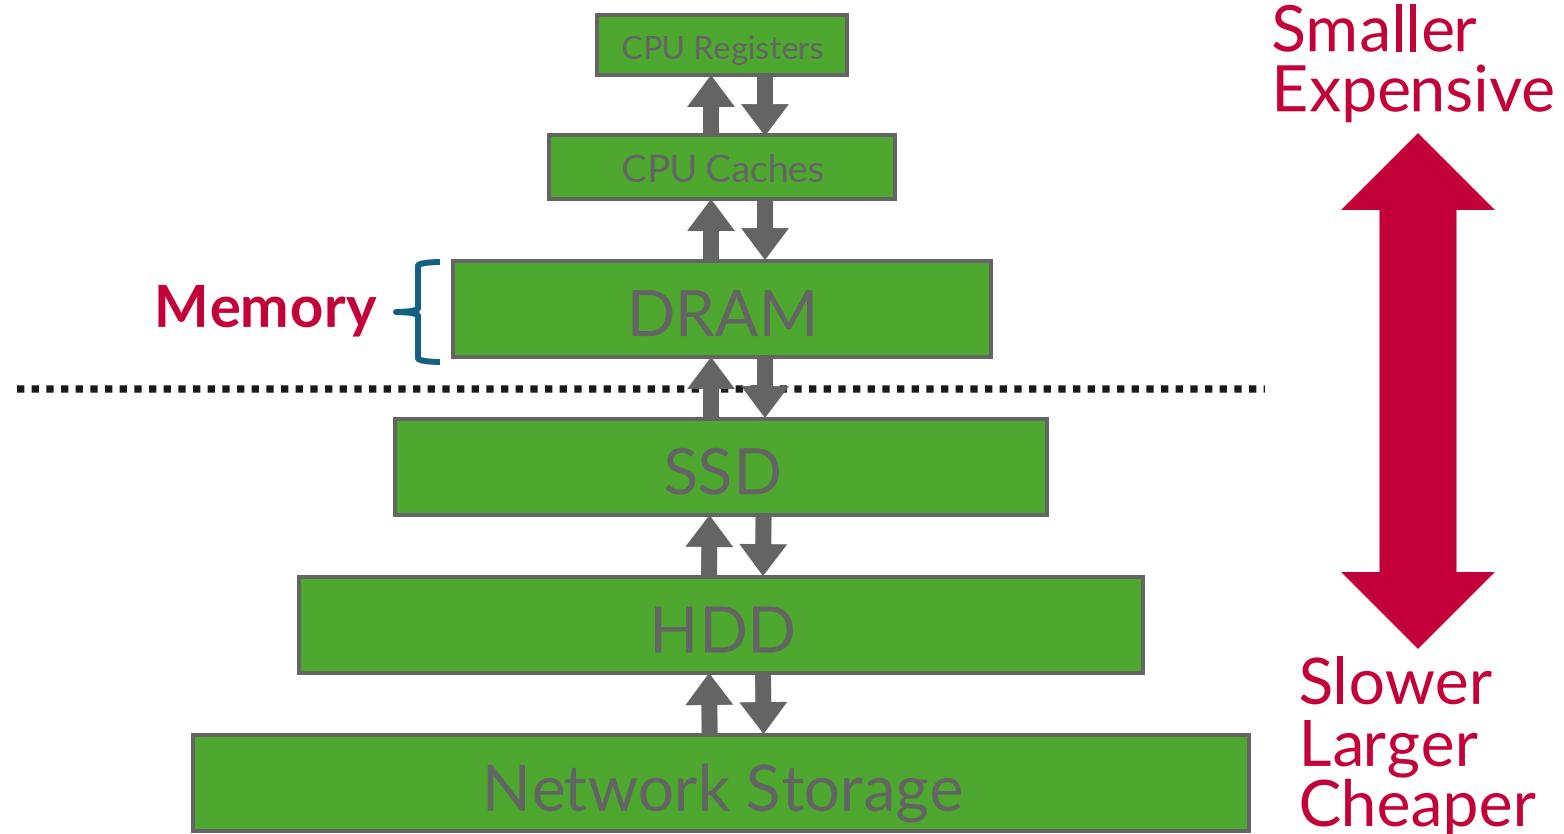
Storage Hierarchy



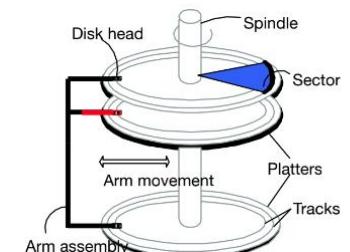
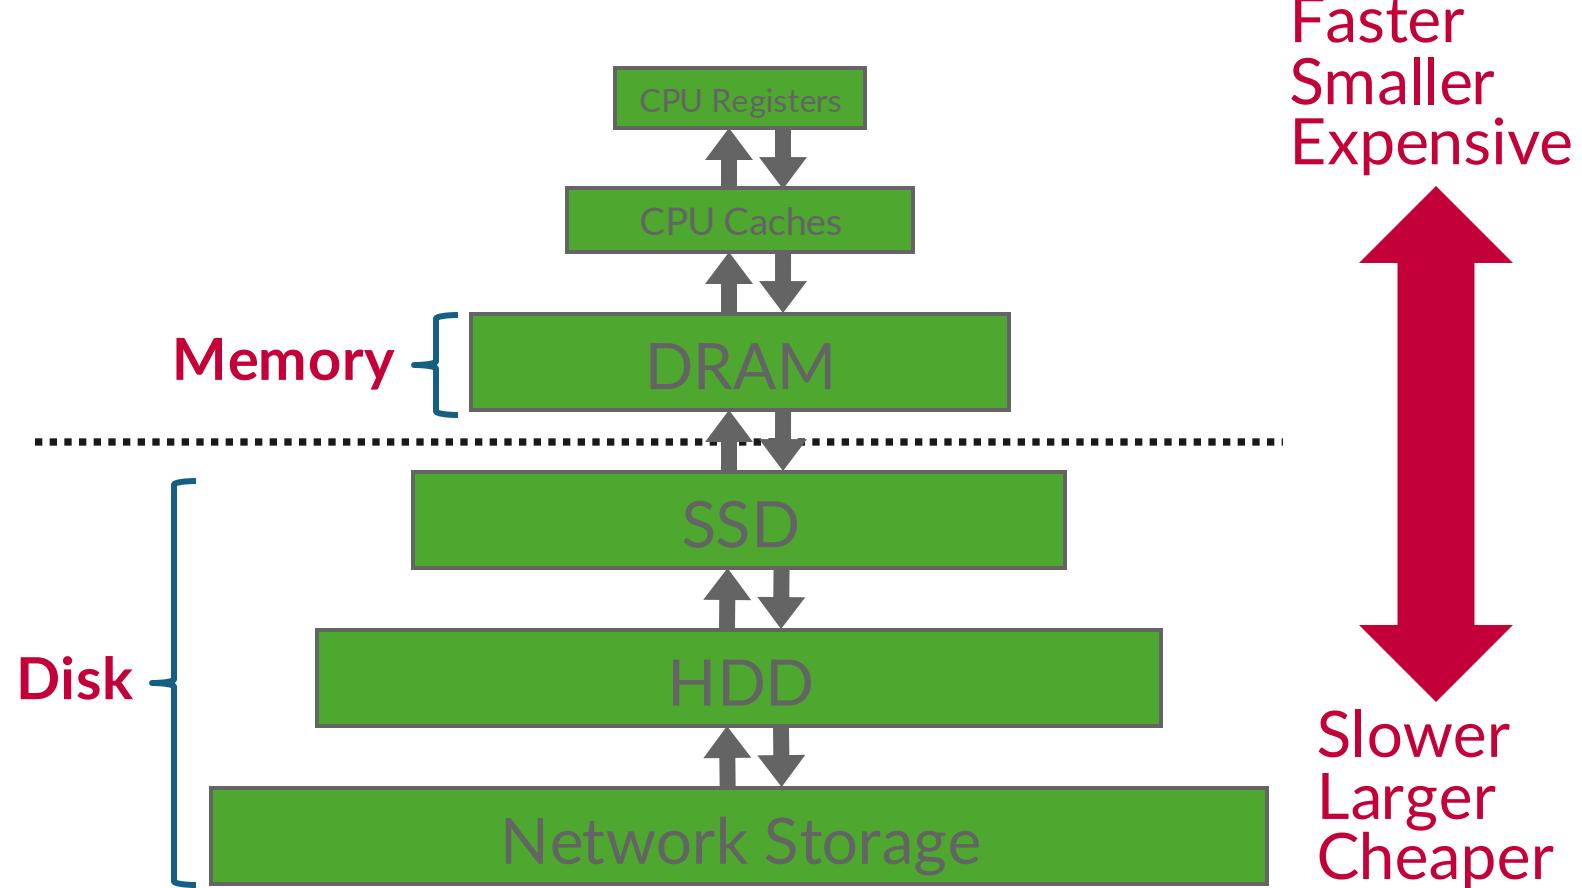
Storage Hierarchy



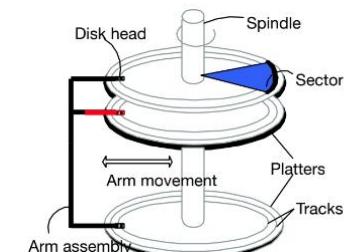
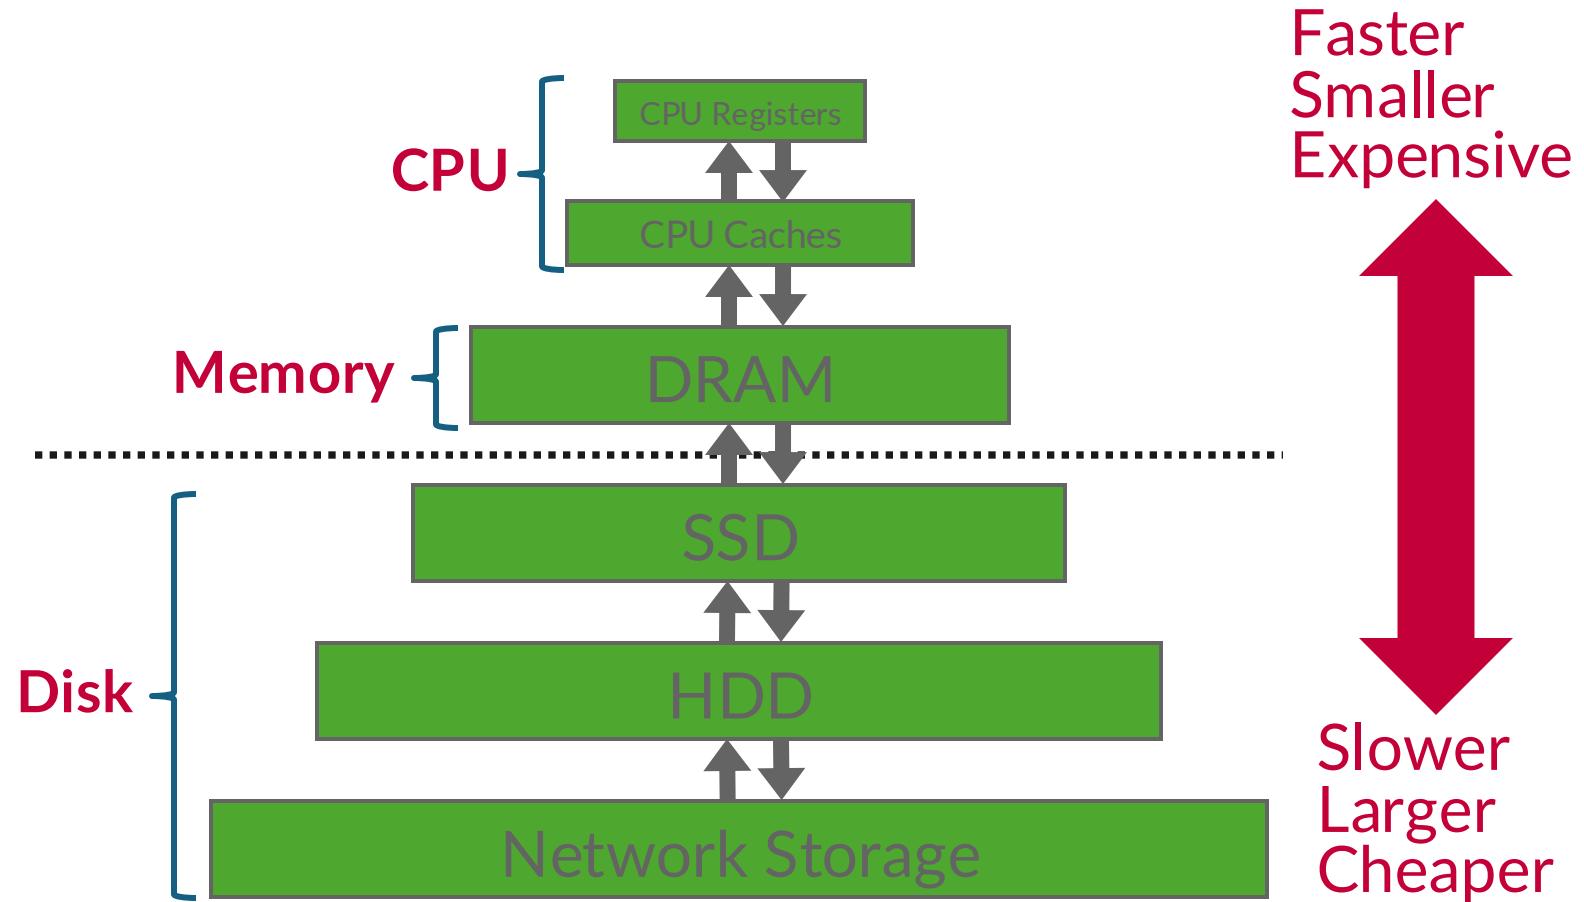
Storage Hierarchy



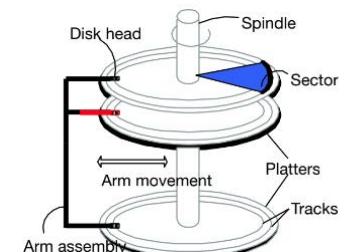
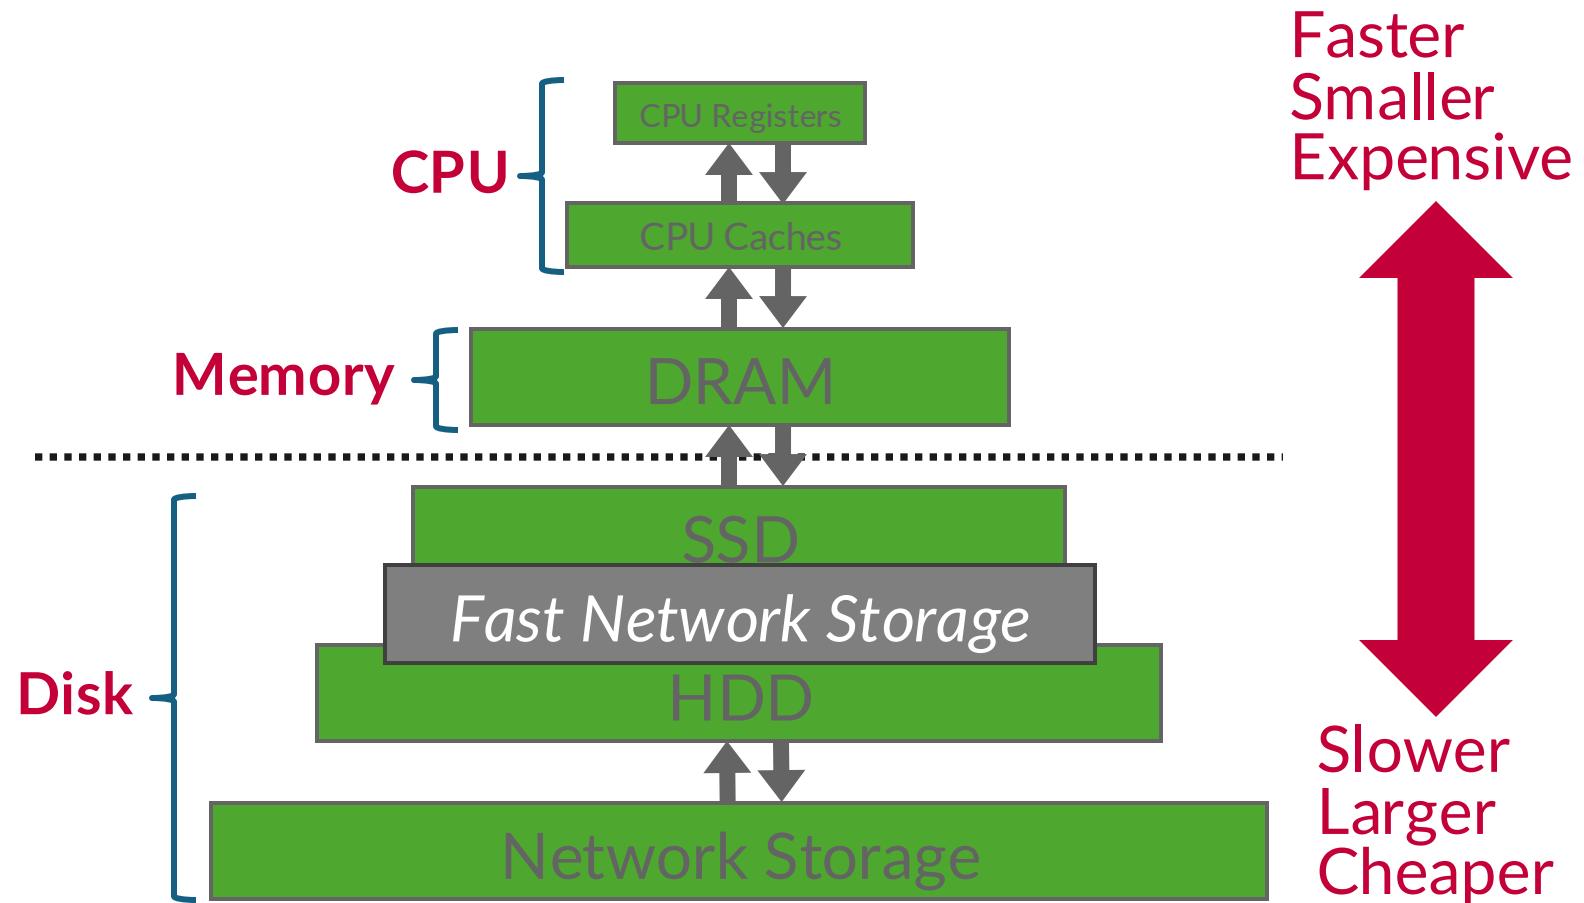
Storage Hierarchy



Storage Hierarchy

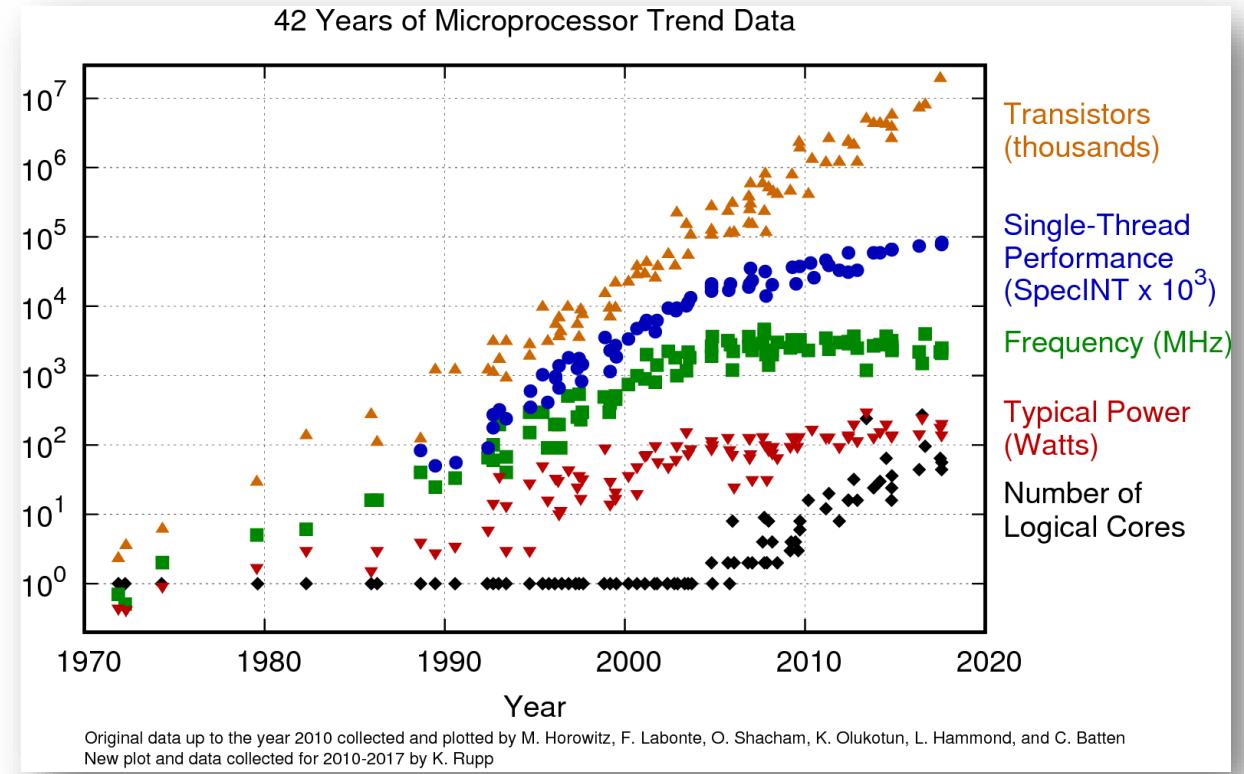


Storage Hierarchy



Hardware Trends

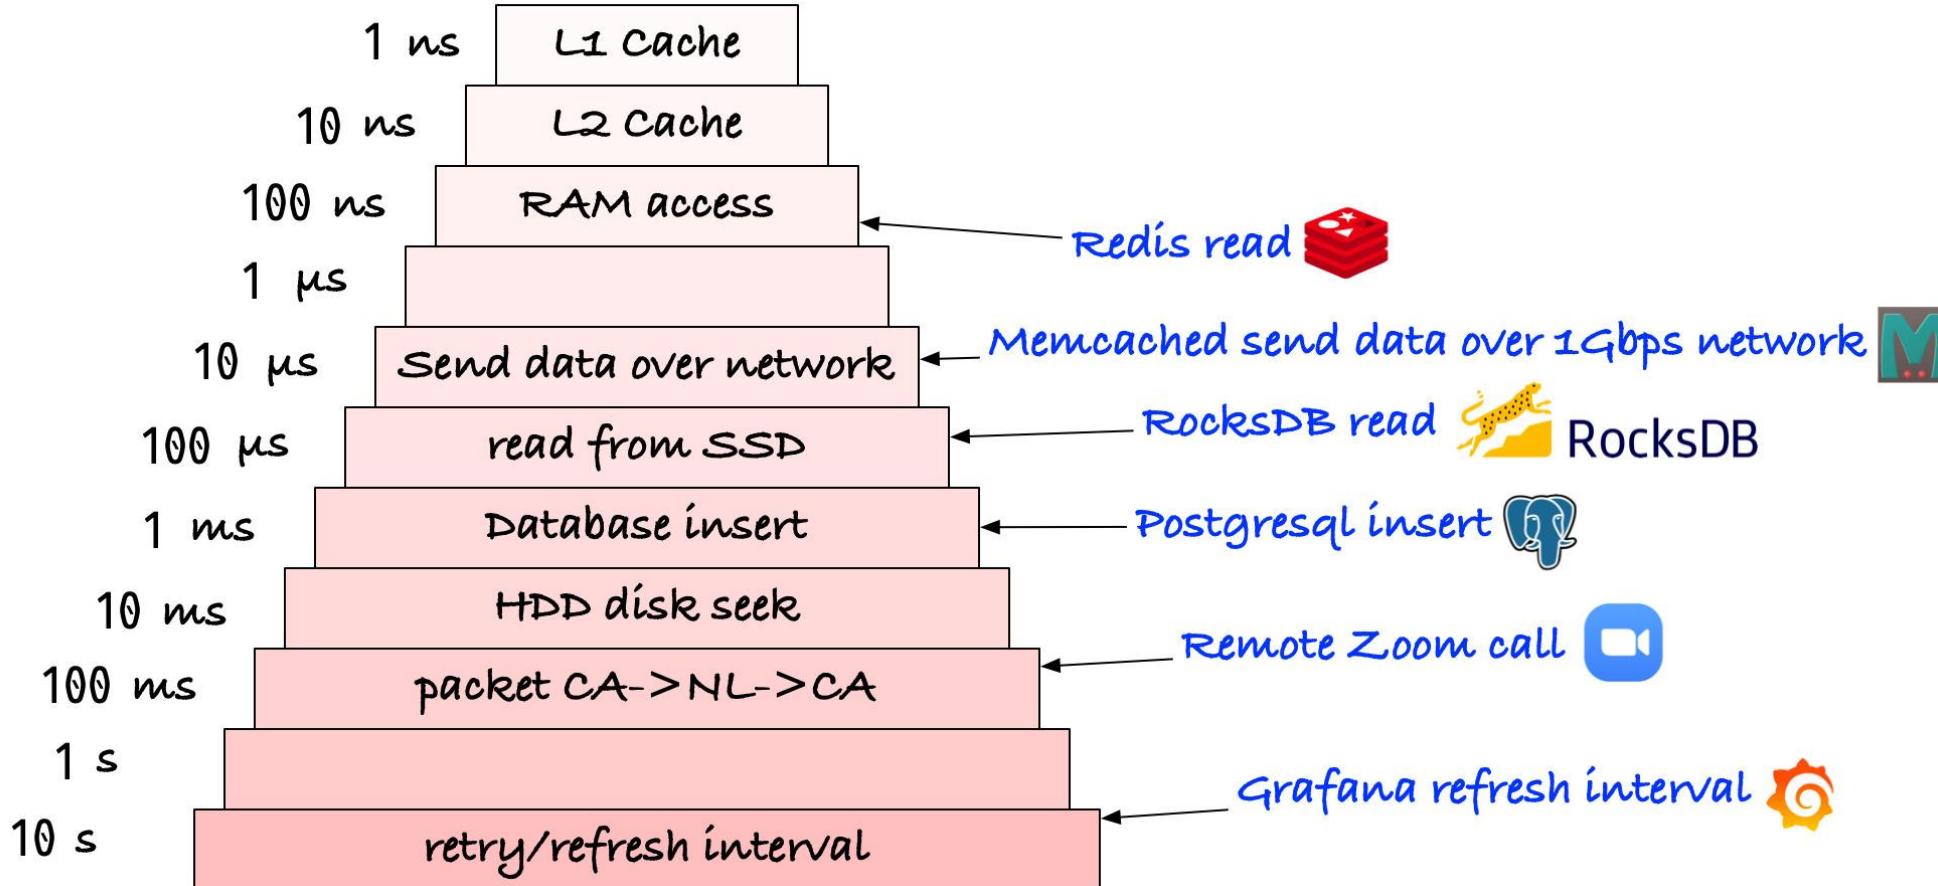
- Transistor growth continues.
- The question is how to use this hardware for higher application performance.
- Individual cores are not becoming faster, but there are more cores.
- Every processor is now a “parallel” data machine, and the degree of parallelism is increasing.



<https://www.karlrupp.net/2018/02/42-years-of-microprocessor-trend-data/>

Latency Numbers You Should Know

Every
programmer
must know
these numbers.



Jeff Dean

<https://blog.bytebytogo.com/p/ep22-latency-numbers-you-should-know>

Sequential v.s. Random Access

- Random access on non-volatile storage is almost always much slower than sequential access.
- DBMS will want to maximize sequential access.
 - Algorithms try to reduce number of writes to random pages so that data is stored in contiguous blocks.
 - Allocating multiple pages at the same time is called an extent.

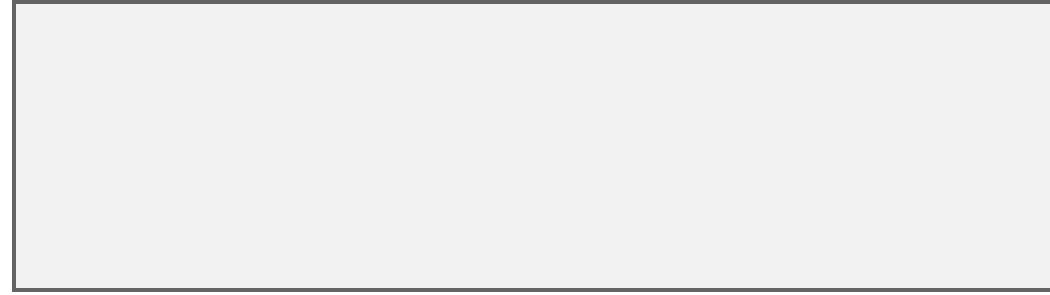
System Design Goals

- Allow the DBMS to manage databases that exceed the amount of memory available.
- Reading/writing to disk is expensive, so it must be managed carefully to avoid large stalls and performance degradation.
- Random access on disk is usually much slower than sequential access, so the DBMS will want to maximize sequential access.

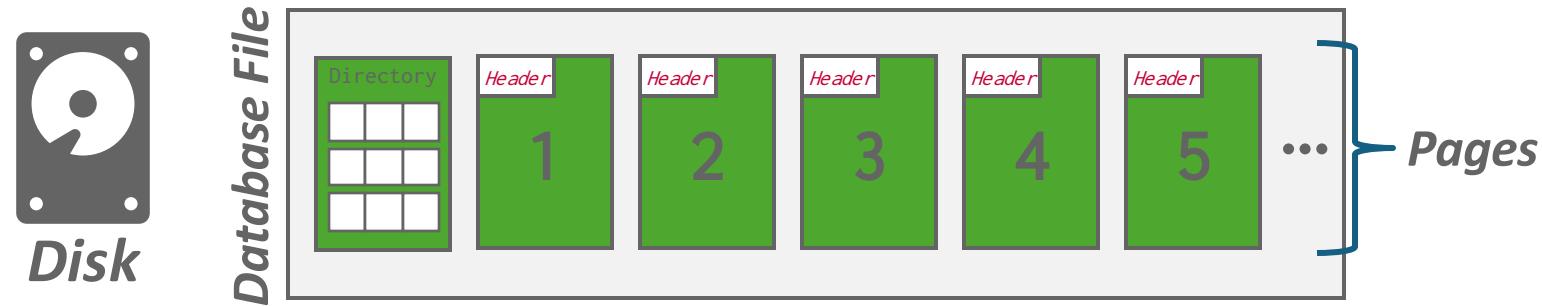
Disk-oriented DBMS



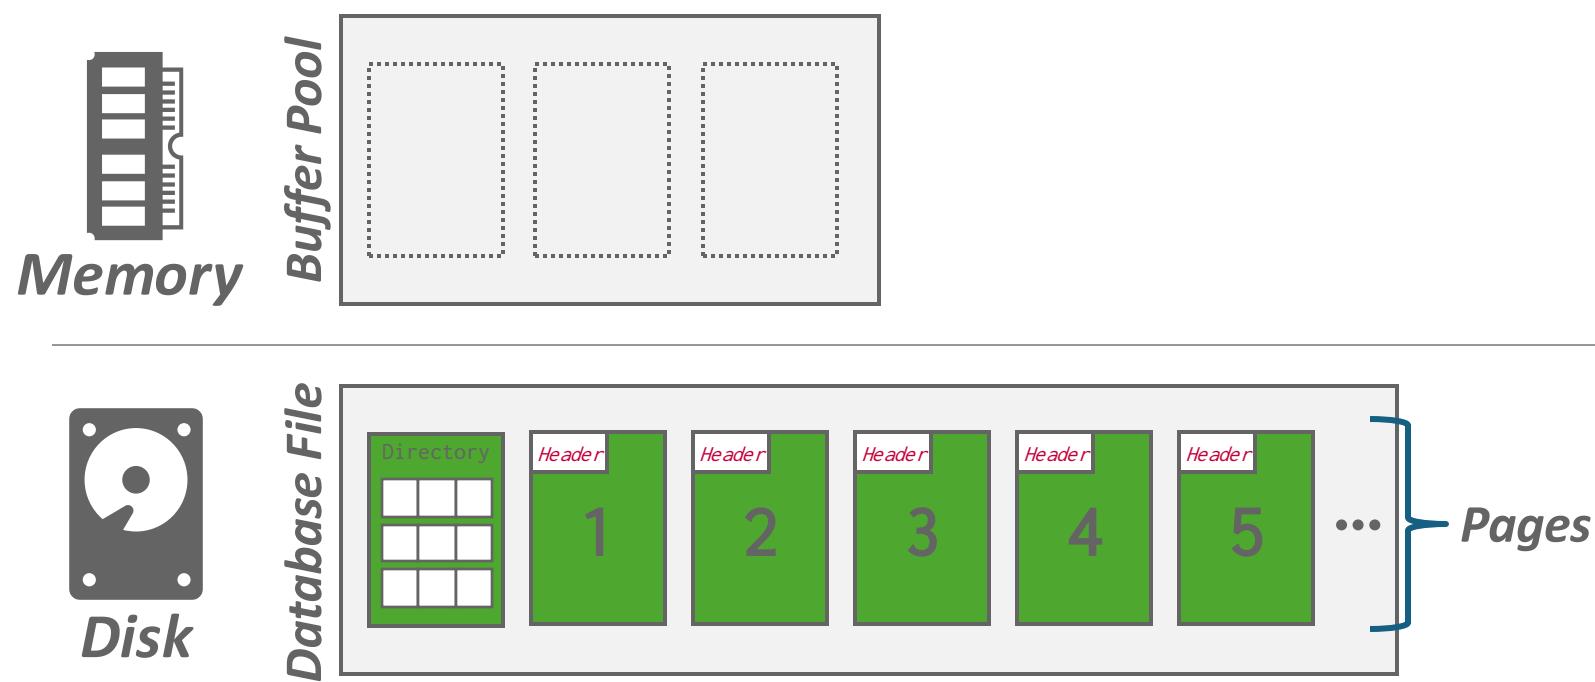
Database File



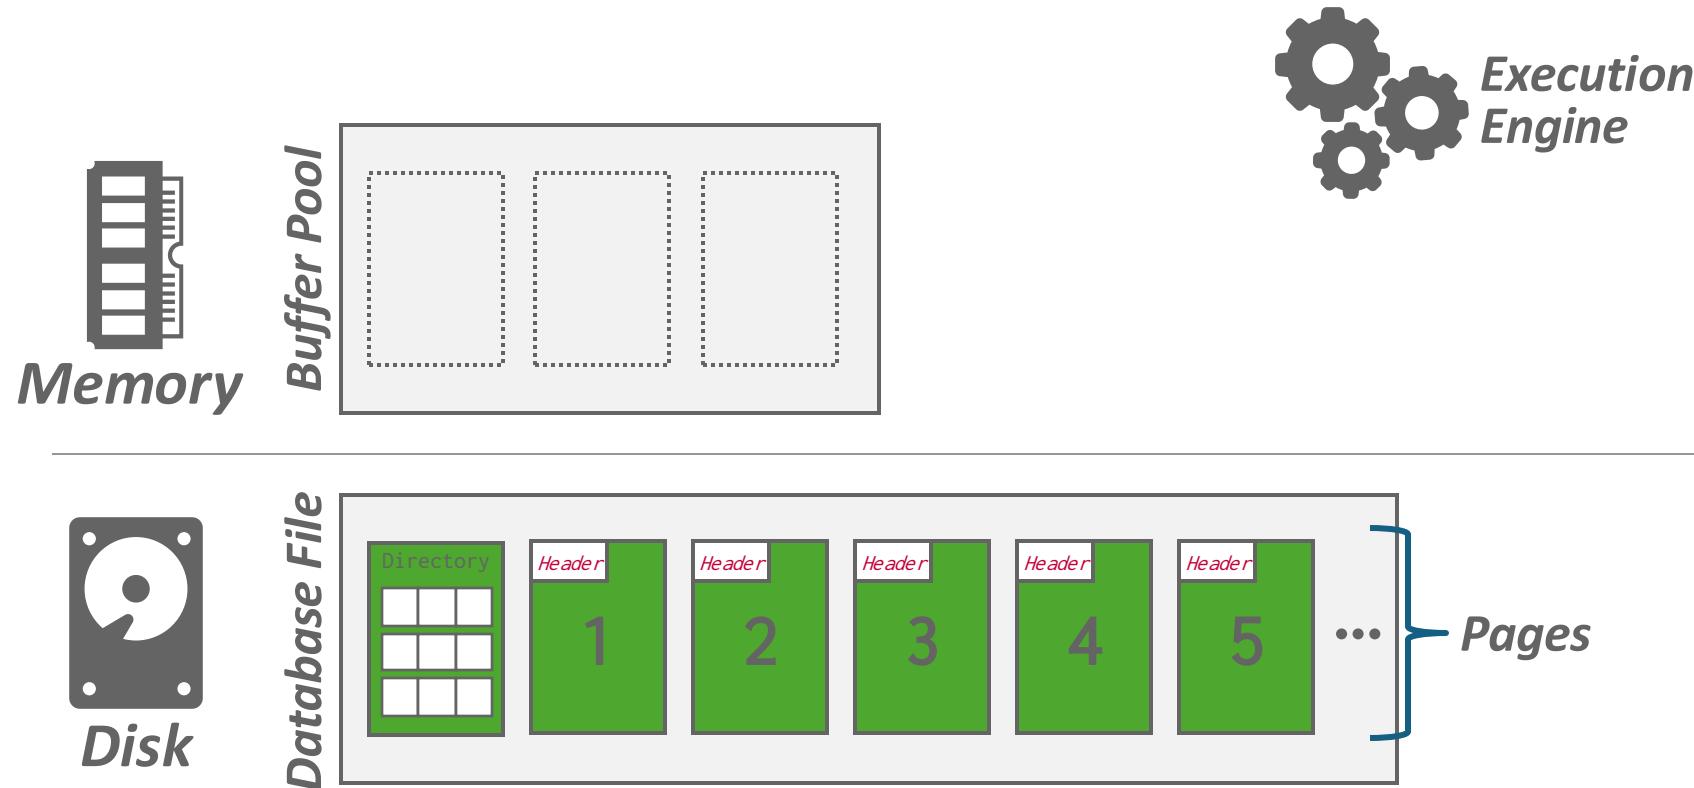
Disk-oriented DBMS



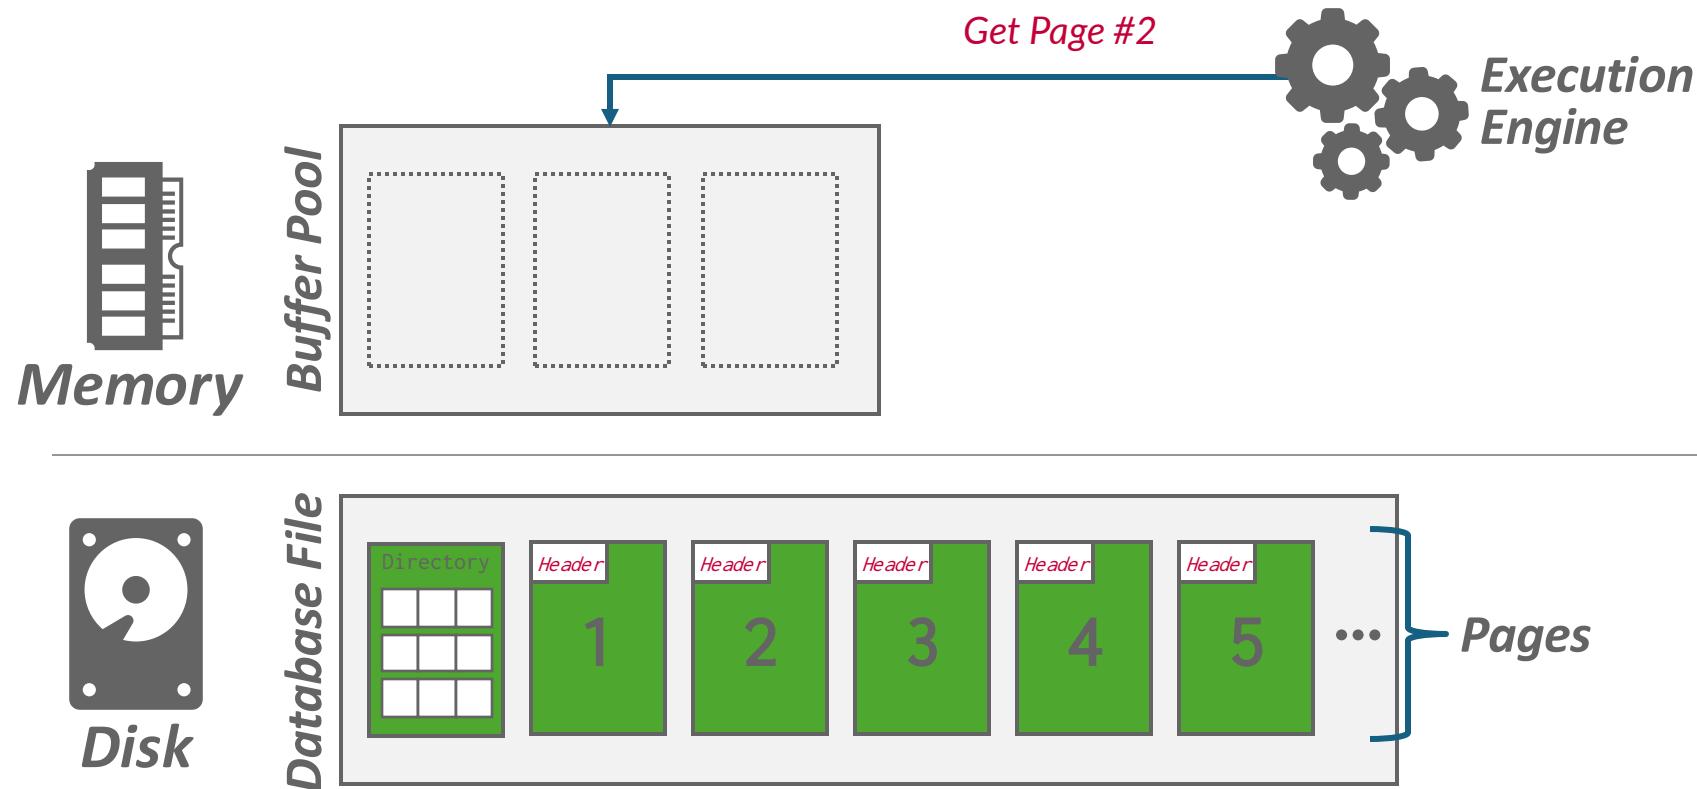
Disk-oriented DBMS



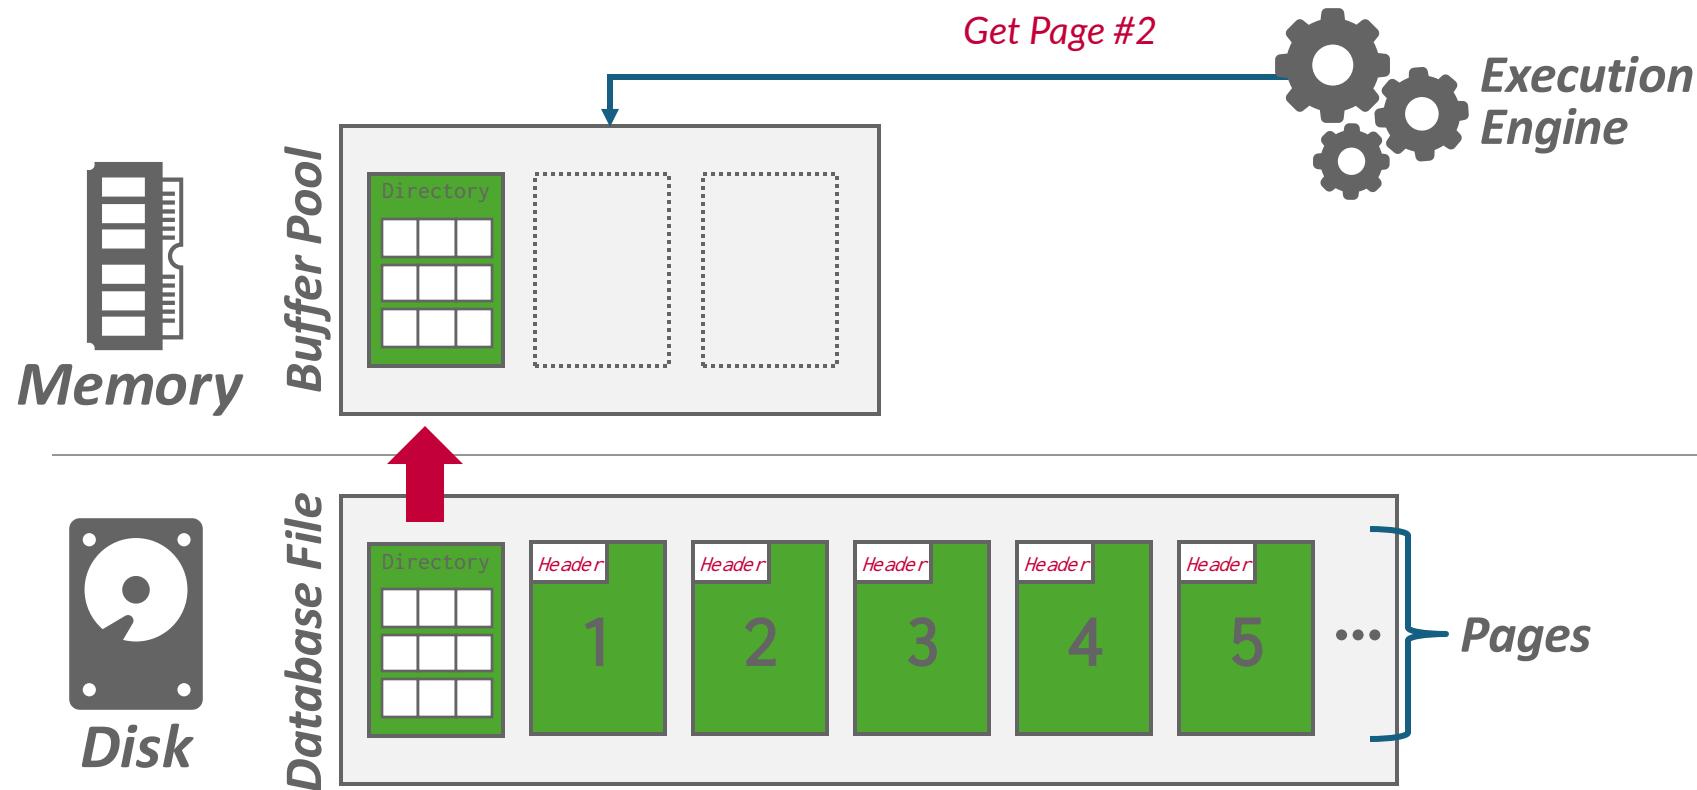
Disk-oriented DBMS



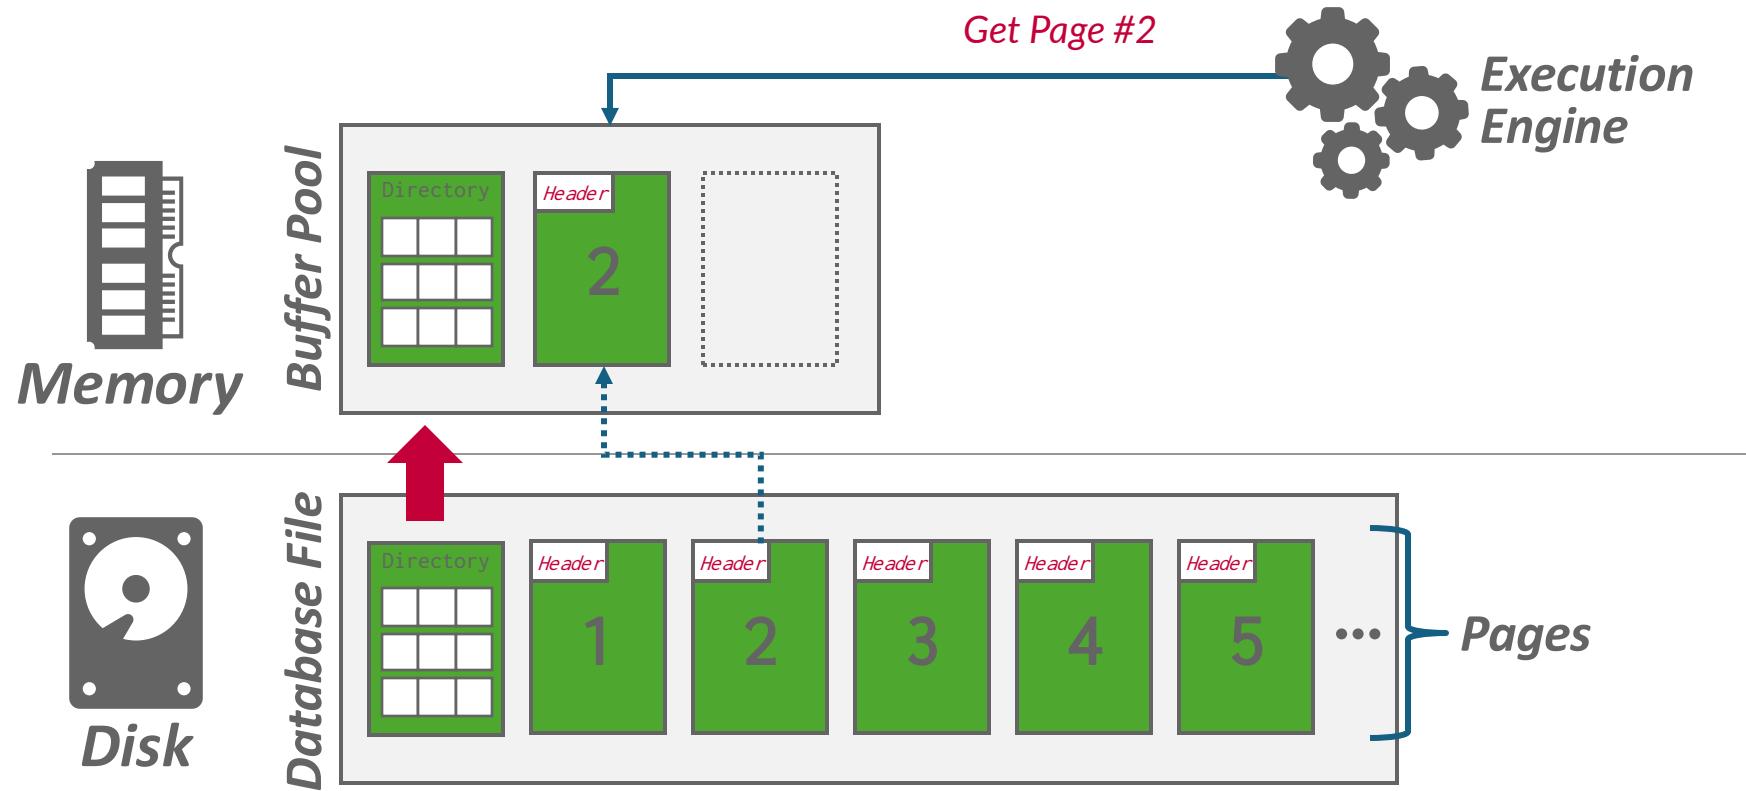
Disk-oriented DBMS



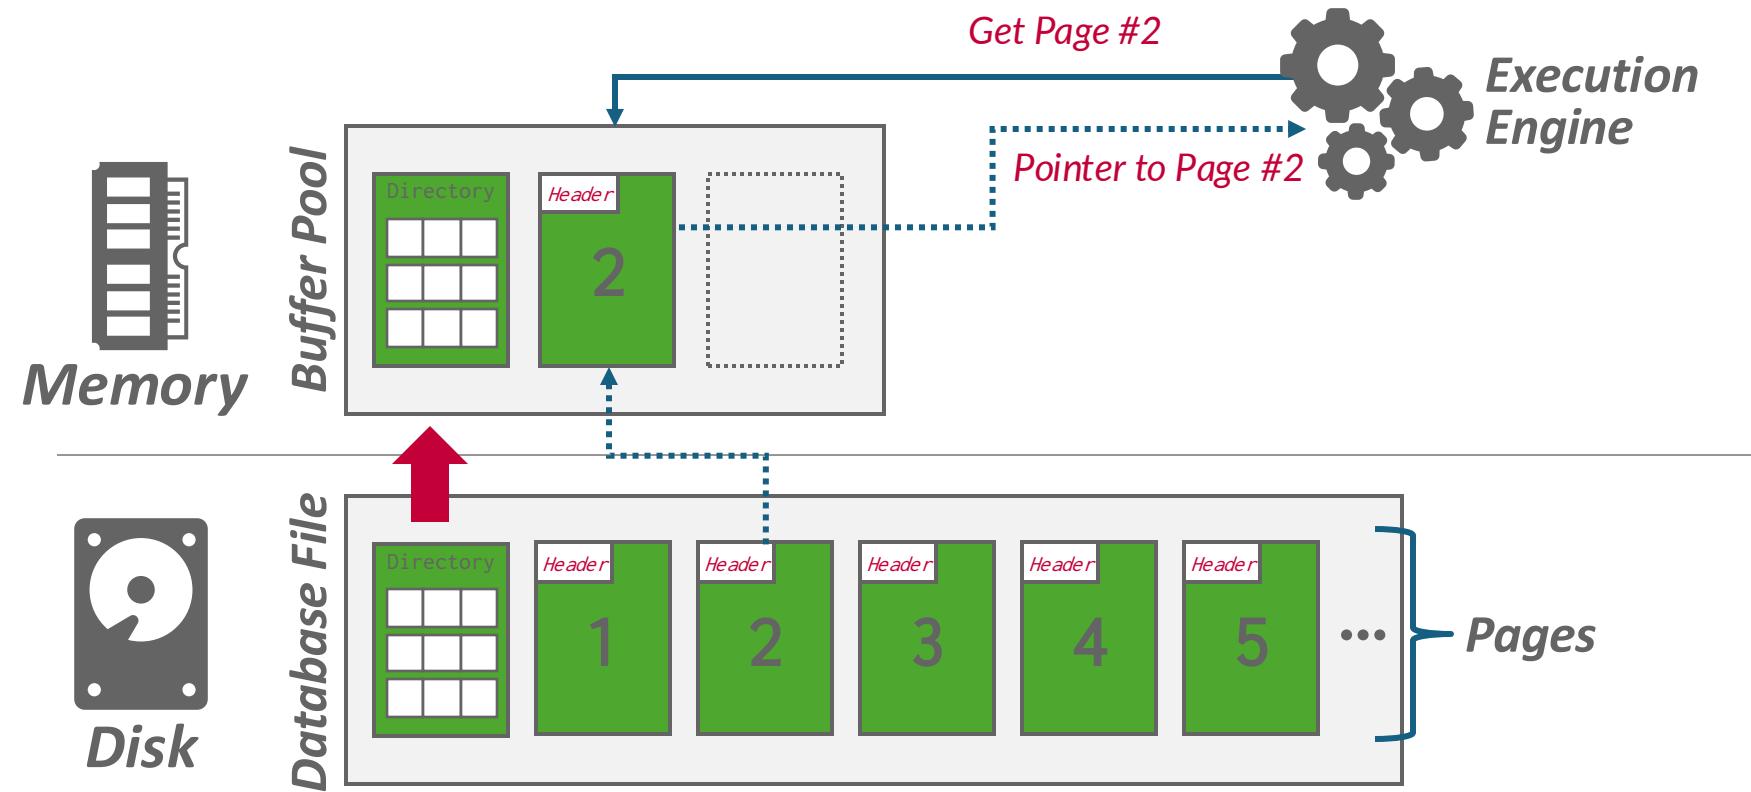
Disk-oriented DBMS



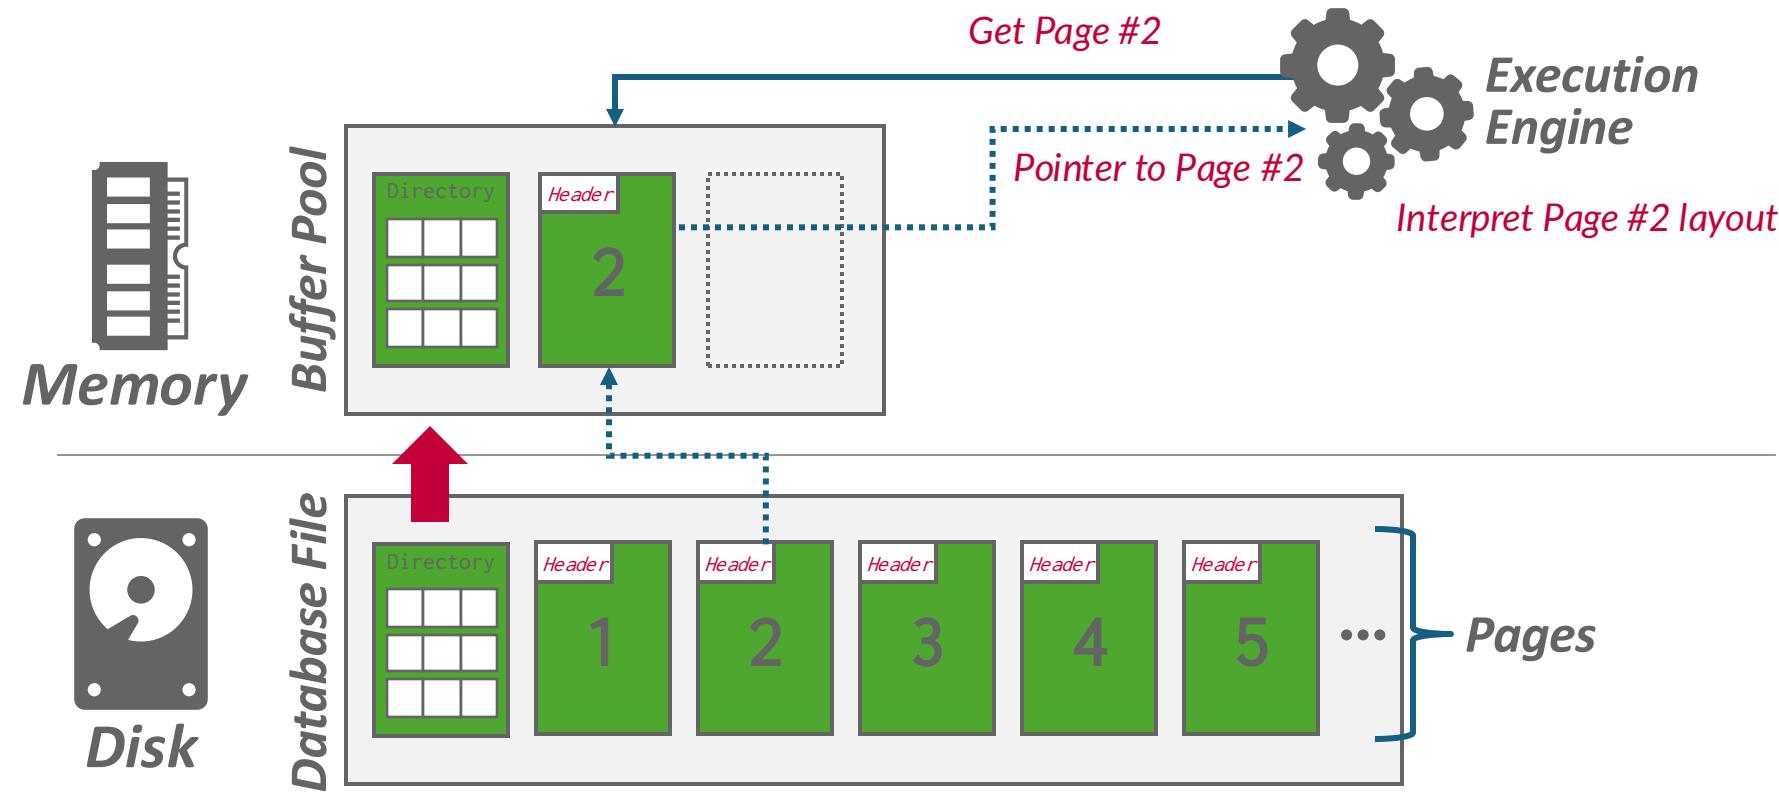
Disk-oriented DBMS



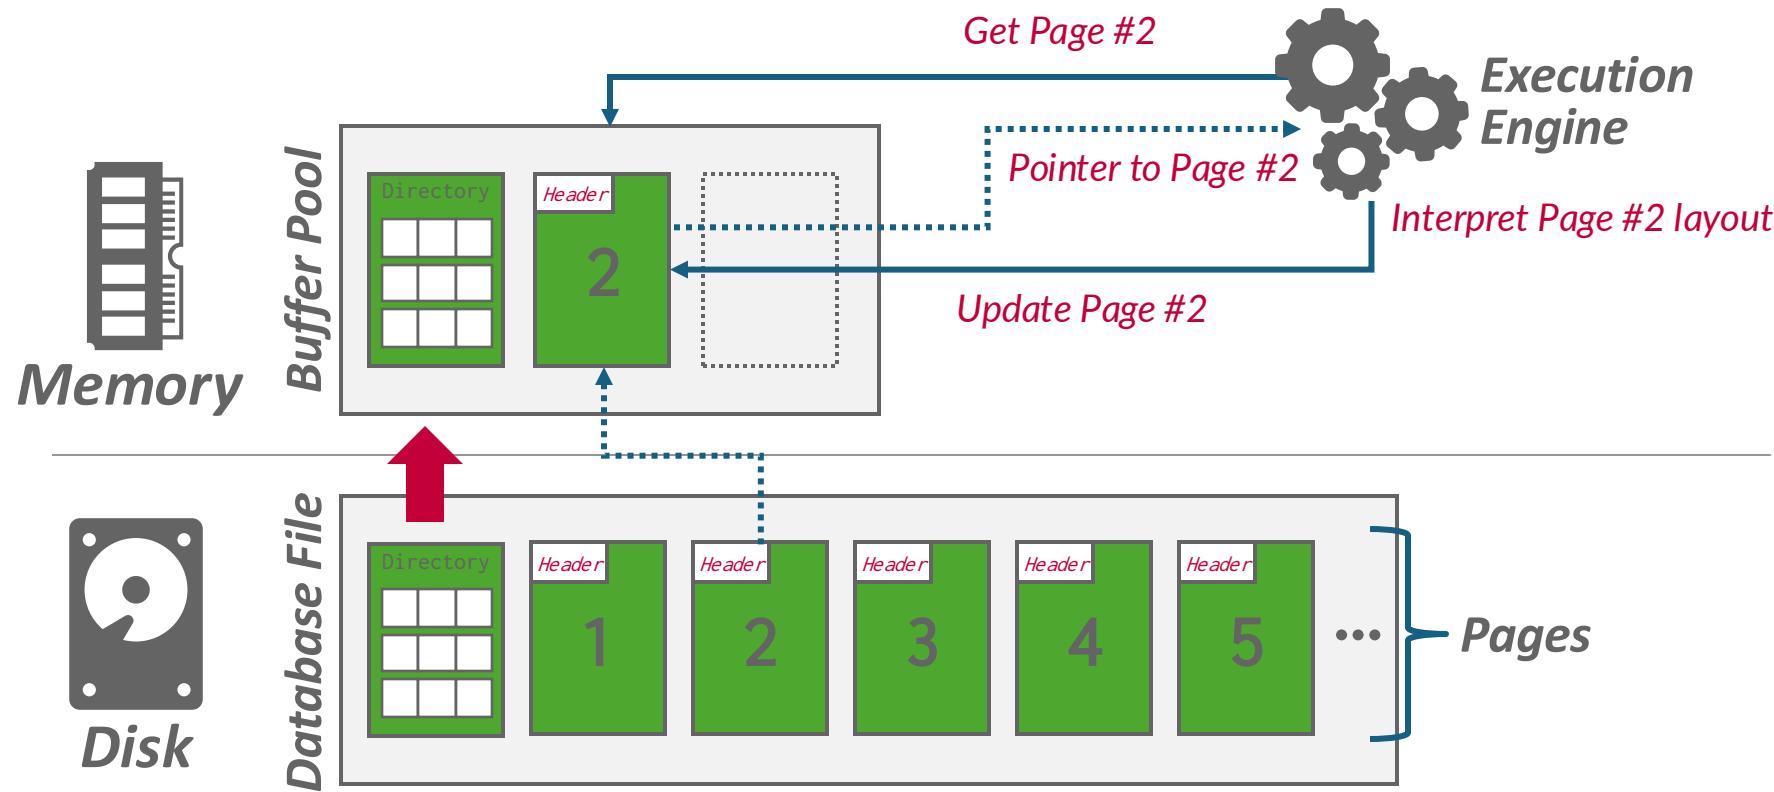
Disk-oriented DBMS



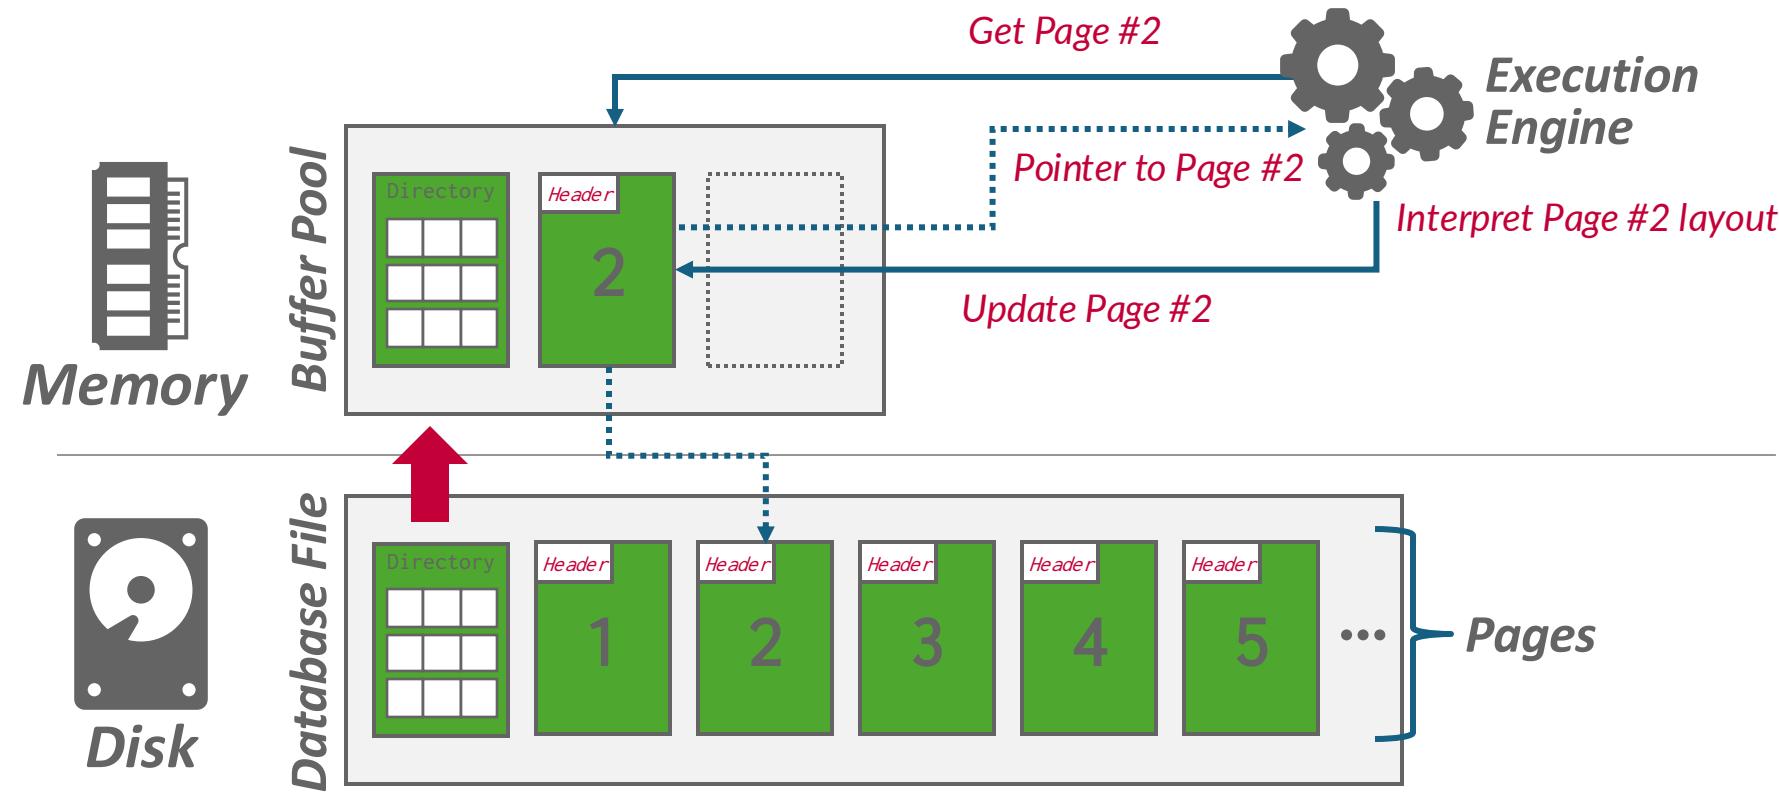
Disk-oriented DBMS



Disk-oriented DBMS



Disk-oriented DBMS



Why Not Use the OS?

- The DBMS can use memory mapping (**mmap**) to store the contents of a file into the address space of a program.
- OS is responsible for moving file pages in and out of memory, so the DBMS doesn't need to worry about it.

Why Not Use the OS?

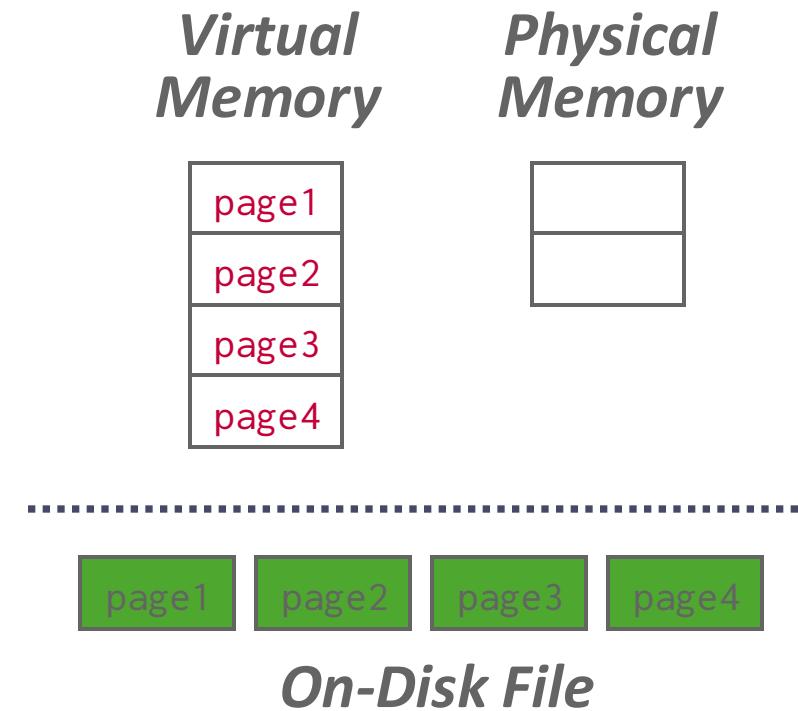
- The DBMS can use memory mapping (**mmap**) to store the contents of a file into the address space of a program.
- OS is responsible for moving file pages in and out of memory, so the DBMS doesn't need to worry about it.



On-Disk File

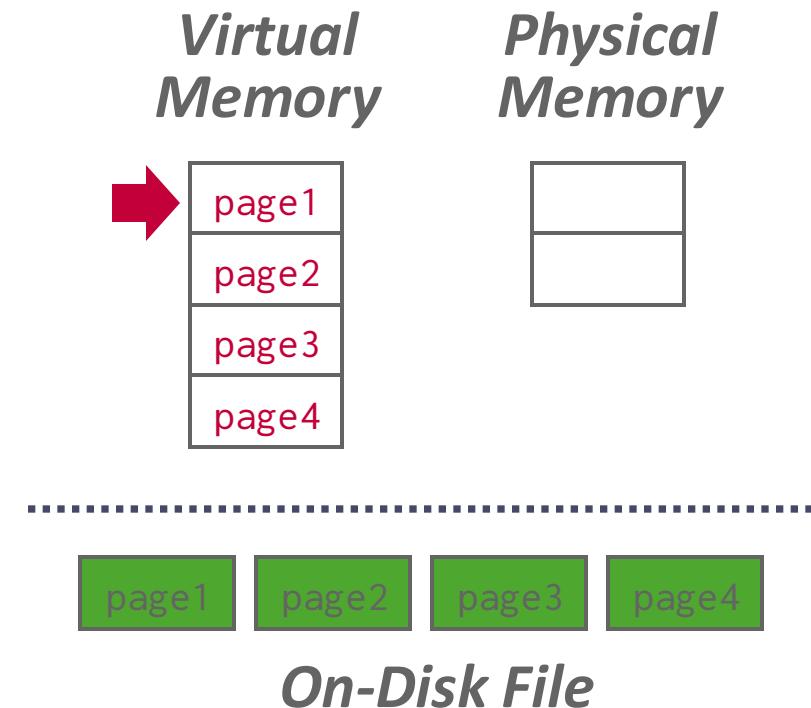
Why Not Use the OS?

- The DBMS can use memory mapping (`mmap`) to store the contents of a file into the address space of a program.
- OS is responsible for moving file pages in and out of memory, so the DBMS doesn't need to worry about it.



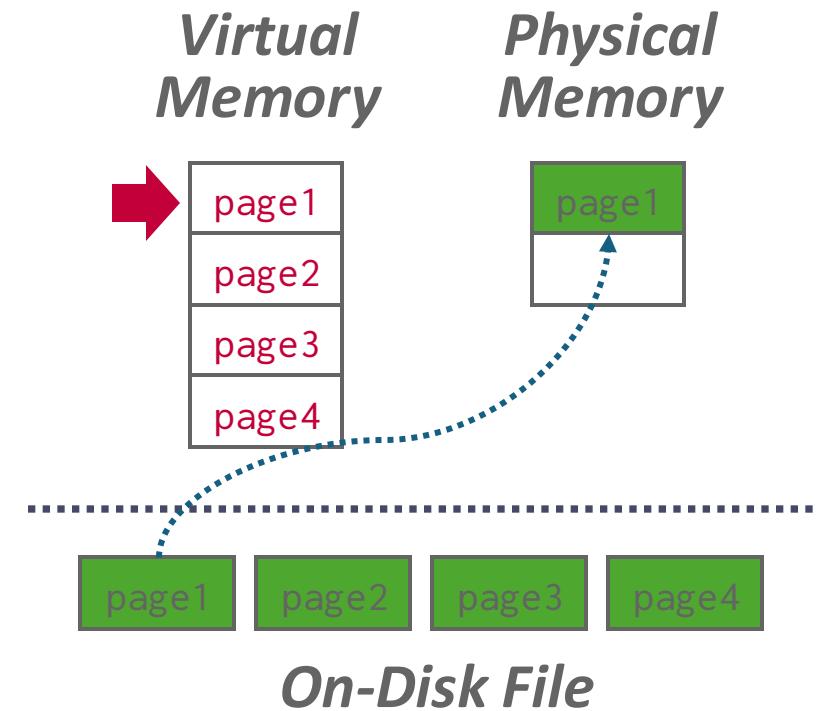
Why Not Use the OS?

- The DBMS can use memory mapping (`mmap`) to store the contents of a file into the address space of a program.
- OS is responsible for moving file pages in and out of memory, so the DBMS doesn't need to worry about it.



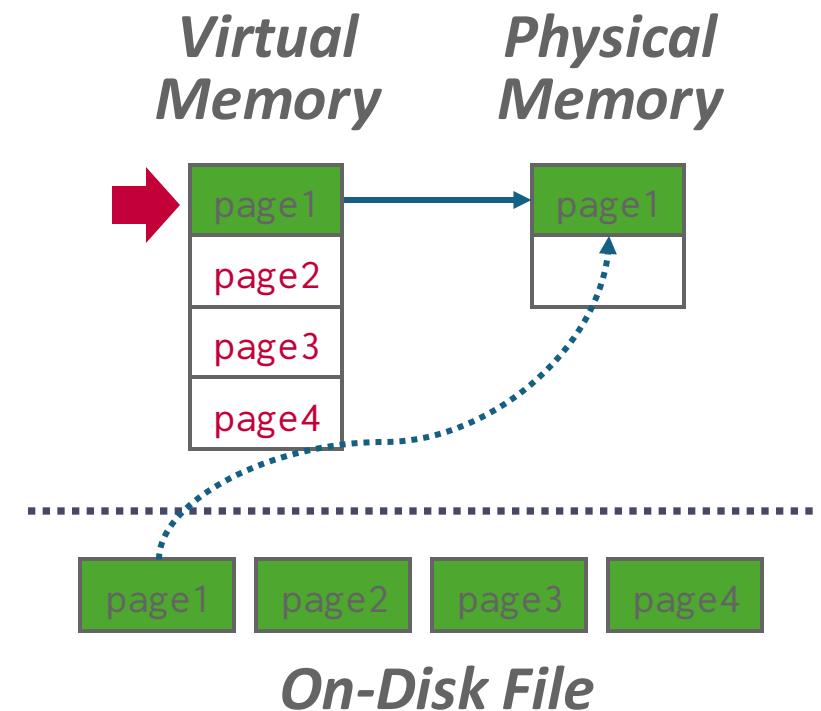
Why Not Use the OS?

- The DBMS can use memory mapping (`mmap`) to store the contents of a file into the address space of a program.
- OS is responsible for moving file pages in and out of memory, so the DBMS doesn't need to worry about it.



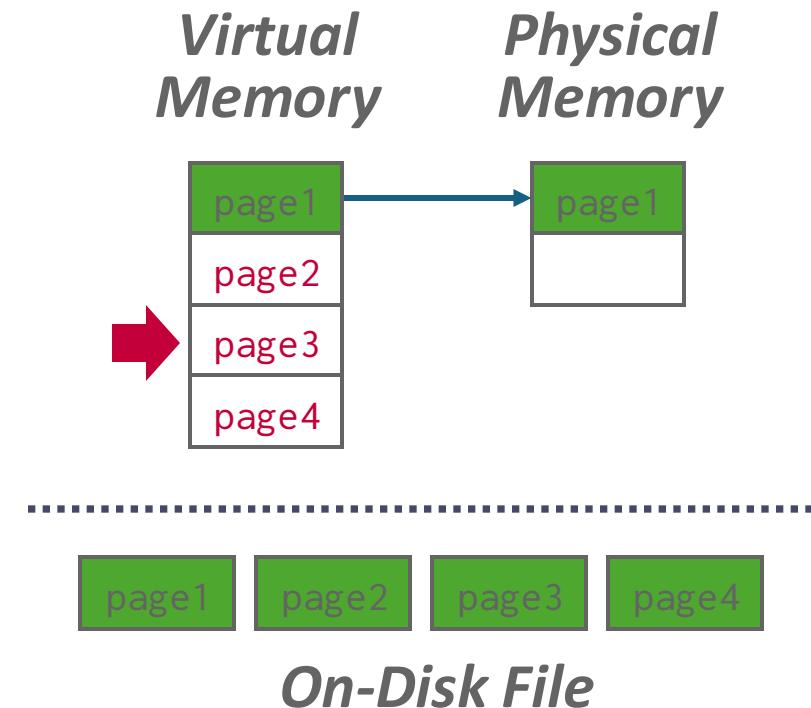
Why Not Use the OS?

- The DBMS can use memory mapping (`mmap`) to store the contents of a file into the address space of a program.
- OS is responsible for moving file pages in and out of memory, so the DBMS doesn't need to worry about it.



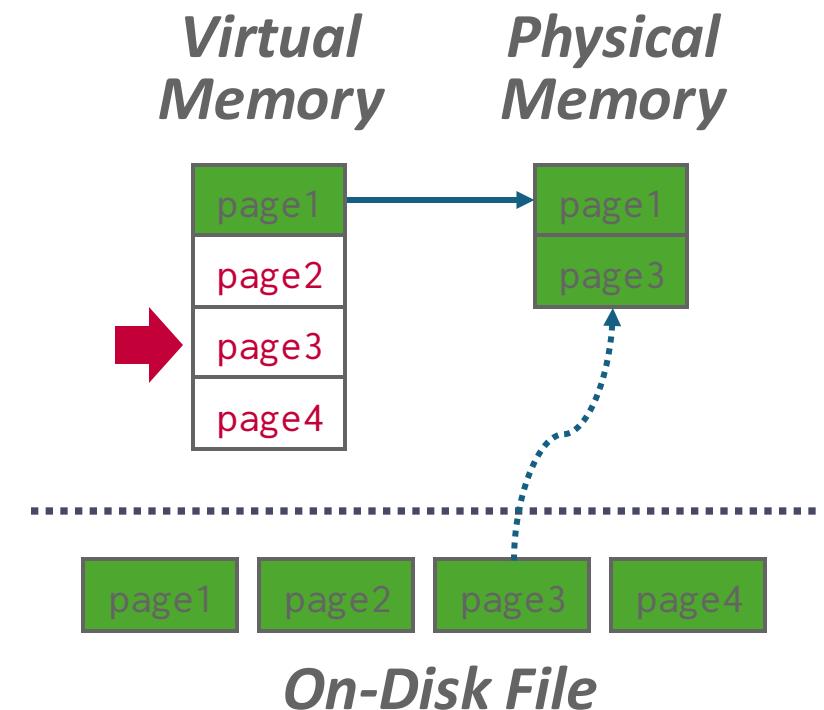
Why Not Use the OS?

- The DBMS can use memory mapping (`mmap`) to store the contents of a file into the address space of a program.
- OS is responsible for moving file pages in and out of memory, so the DBMS doesn't need to worry about it.



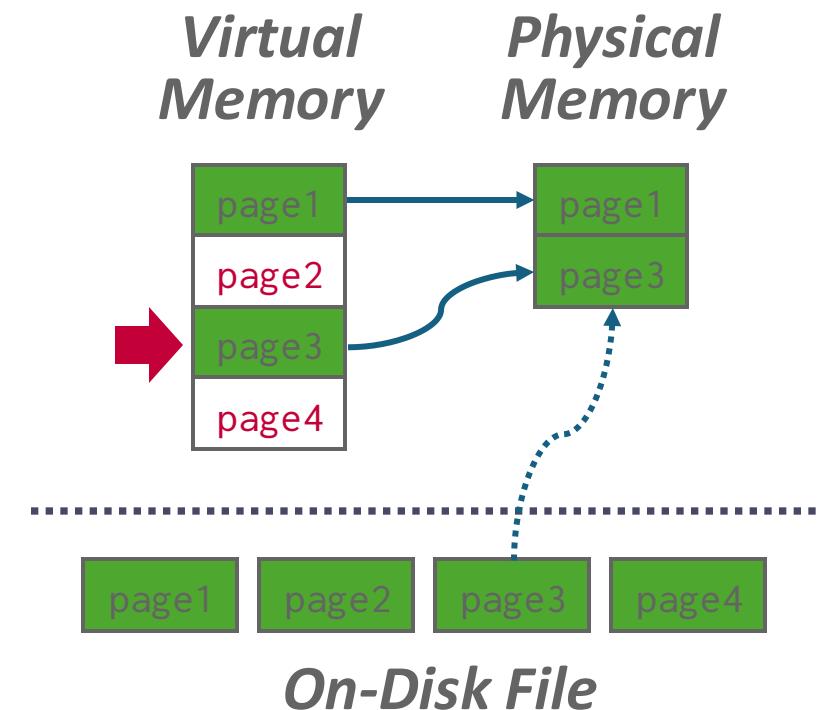
Why Not Use the OS?

- The DBMS can use memory mapping (`mmap`) to store the contents of a file into the address space of a program.
- OS is responsible for moving file pages in and out of memory, so the DBMS doesn't need to worry about it.



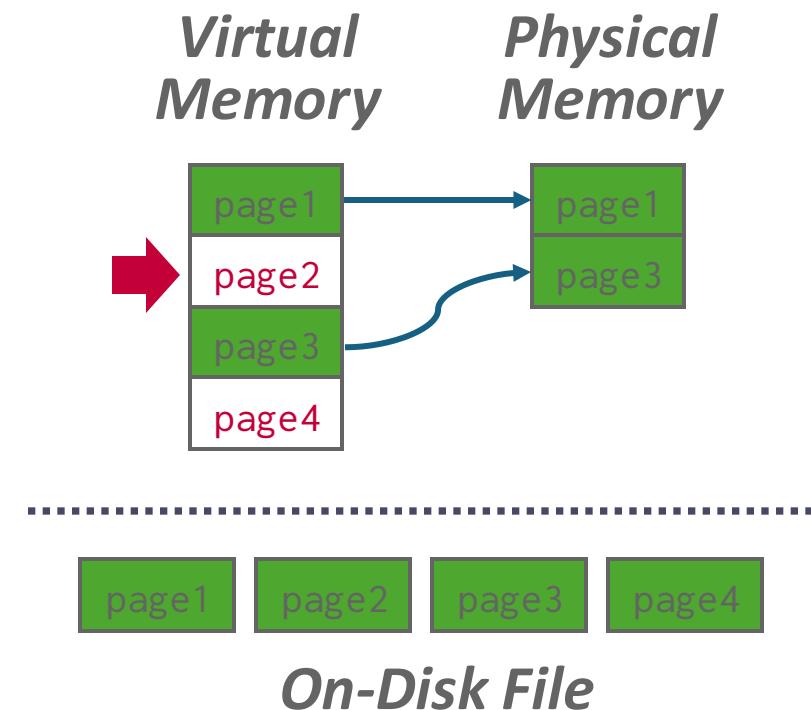
Why Not Use the OS?

- The DBMS can use memory mapping (`mmap`) to store the contents of a file into the address space of a program.
- OS is responsible for moving file pages in and out of memory, so the DBMS doesn't need to worry about it.



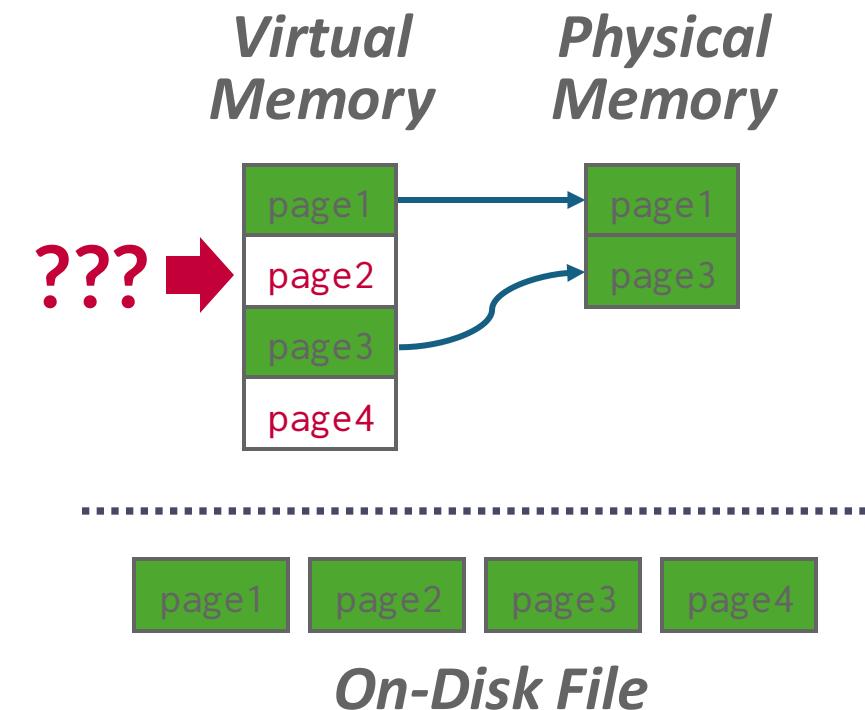
Why Not Use the OS?

- The DBMS can use memory mapping (`mmap`) to store the contents of a file into the address space of a program.
- OS is responsible for moving file pages in and out of memory, so the DBMS doesn't need to worry about it.



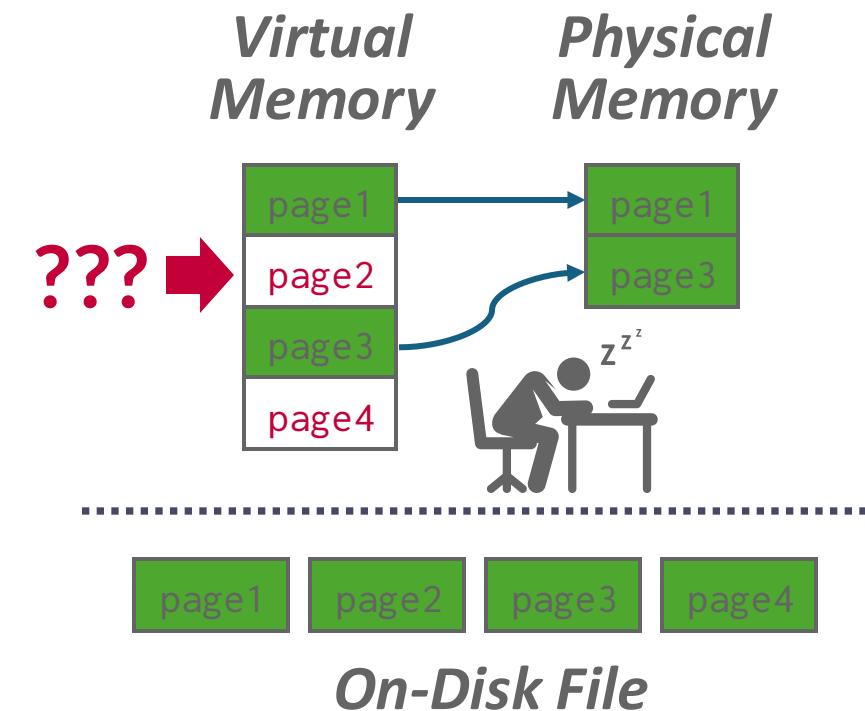
Why Not Use the OS?

- The DBMS can use memory mapping (`mmap`) to store the contents of a file into the address space of a program.
- OS is responsible for moving file pages in and out of memory, so the DBMS doesn't need to worry about it.



Why Not Use the OS?

- The DBMS can use memory mapping (**mmap**) to store the contents of a file into the address space of a program.
- OS is responsible for moving file pages in and out of memory, so the DBMS doesn't need to worry about it.



Why Not Use the OS?

- What if we allow multiple threads to access the `mmap` files to hide page fault stalls?
- This works reasonably well for read-only access.
It is complicated when there are multiple writers...

Memory Mapped I/O Problems

Problem #1: Transaction Safety

- OS can flush dirty pages at any time.

Problem #2: I/O Stalls

- DBMS doesn't know which pages are in memory. The OS will stall a thread on page fault.

Problem #3: Error Handling

- Difficult to validate pages. Any access can cause a **SIGBUS** that the DBMS must handle.

Problem #4: Performance Issues

- OS data structure contention. TLB shootdowns.

Why Not Use the OS?

There are some solutions to some of these problems:

- **madvise**: Tell the OS how you expect to read certain pages.
- **mlock**: Tell the OS that memory ranges cannot be paged out.
- **msync**: Tell the OS to flush memory ranges out to disk.

Using these syscalls to get the OS to behave correctly is just as onerous as managing memory yourself.

Why Not Use the OS?

There are some solutions to some of these problems:

- **madvise**: Tell the OS how you expect to read certain pages.
- **mlock**: Tell the OS that memory ranges cannot be paged out.
- **msync**: Tell the OS to flush memory ranges out to disk.

Using these syscalls to get the OS to behave correctly is just as onerous as managing memory yourself.

Full Usage



Partial Usage



Why Not Use the OS?

There are some solutions to some of these problems:

- **madvise**: Tell the OS how you expect to read certain pages.
- **mlock**: Tell the OS that memory ranges cannot be paged out.
- **msync**: Tell the OS to flush memory ranges out to disk.

Using these syscalls to get the OS to behave correctly is just as onerous as managing memory yourself.

Full Usage



Partial Usage



Why Not Use the OS?

- DBMS (almost) always wants to control things itself and can do a better job than the OS.
 - Flushing dirty pages to disk in the correct order.
 - Specialized prefetching.
 - Buffer replacement policy.
 - Thread/process scheduling.

The OS is not your friend.

Are You Sure You Want to Use MMAP in Your Database Management System?

Andrew Croft
Carnegie Mellon University
andrewc@cs.cmu.edu

Viktor Leis
University of Erlangen-Nuremberg
viktor.leis@fau.de

Andrew Pavlo
Carnegie Mellon University
pavlo@cs.cmu.edu

ABSTRACT

Memory-mapped (mmap) file I/O is an OS-provided feature that maps the contents of a file on secondary storage into a program's address space. The program then accesses pages via pointers as if the file resided entirely in memory. The OS transparently loads pages only when the program references them and automatically evicts pages if memory fills up.

mmap's perceived benefits have seduced database management system (DBMS) developers for decades as a viable alternative to implementing a buffer pool. There are, however, severe correctness and performance issues with mmap that are not immediately apparent. Such problems make it difficult, if not impossible, to use mmap correctly and efficiently in a modern DBMS. In fact, several popular DBMSs initially used mmap to support larger-than-memory databases but soon ended up with these hidden perks, forcing them to switch back to using file I/O themselves, introducing significant costs. In this way, mmap and DBMSs are like coffee and spicy food: an unfortunate combination that becomes obvious after the fact.

Since developers keep trying to use mmap in new DBMSs, we wrote this paper to provide a warning to others that mmap is not a suitable replacement for a traditional buffer pool. We discuss the main shortcomings of mmap in detail, and our experimental analysis demonstrates clear performance limitations. Based on these findings, we conclude with a prescription for when DBMSs *might* consider using mmap for file I/O.

1 INTRODUCTION

An important feature of disk-based DBMSs is their ability to support databases that are larger than the available physical memory. This functionality allows a user to query a database as if it resides entirely in memory, even if it does not fit at all once. DBMSs achieve this illusion by reading pages of data from secondary storage (e.g., HDD, SSD) on demand. If there is not enough memory for a new page, the DBMS will evict an existing page that is no longer needed in order to make room.

Traditionally, DBMSs implement the movement of pages between secondary storage and memory in a buffer pool, which interacts with secondary storage using system calls like `read` and `write`. These file I/O mechanisms copy data to and from a buffer in user space, with the DBMS maintaining complete control over how and when it transfers pages.

Alternatively, the DBMS relinquishes the responsibility of data movement to the OS, which maintains its own file mapping and

This paper is published under the Creative Commons Attribution 4.0 International (CC-BY) license. It is made available to download, store and use on the personal and corporate Web sites wth the appropriate attribution, provided that you attribute the original work to the authors and CIDR'22, 12th Annual Conference on Innovative Data Systems Research (CIDR '22), January 7–10, 2022, Champaign, IL, USA.

2 BACKGROUND

This section provides the relevant background on mmap. We begin with a high-level overview of memory-mapped file I/O and the POSIX mmap API. Then, we discuss real-world implementations of mmap-based systems.

Database Storage

- **Problem #1:** How the DBMS represents the database in files on disk.
- **Problem #2:** How the DBMS manages its memory and moves data back-and-forth from disk.

Database Storage

- **Problem #1:** How the DBMS represents the database in files on disk.
- **Problem #2:** How the DBMS manages its memory and moves data back-and-forth from disk.

← Today

This Lecture

- File Storage
- Page Layout
- Tuple Layout

File Storage

File Storage

Page Layout

Tuple Layout

File Storage

- The DBMS stores a database as **one or more files** on disk typically in a proprietary format.
 - The OS doesn't know anything about the contents of these files.
- Early systems in the 1980s used custom filesystems on raw block storage.
 - Some “enterprise” DBMSs still support this.
 - Most newer DBMSs do not do this.

Storage Manager

- The storage manager is responsible for maintaining a database's files.
 - Some do their own scheduling for reads and writes to improve spatial and temporal locality of pages.
- It organizes the files as a collection of pages.
 - Tracks data read/written to pages.
 - Tracks the available space.
- Assume that if there is replication (for fault tolerance), it happens outside the core storage manager function.

Database Pages

- A page is a fixed-size block of data.
 - It can contain tuples, meta-data, indexes, log records...
 - Most systems do not mix page types.
 - Some systems require a page to be self-contained.
- Each page is given a unique identifier.
 - The DBMS uses an indirection layer to map page IDs to physical locations.

Database Pages

- There are three different notions of “pages” in a DBMS:
 - Hardware Page (usually 4KB)
 - OS Page (usually 4KB, x64 2MB/1GB)
 - Database Page (512B-32KB)
- A hardware page is the largest block of data that the storage device can guarantee failsafe writes.

Database Pages

- There are three different notions of “pages” in a DBMS:
 - Hardware Page (usually 4KB)
 - OS Page (usually 4KB, x64 2MB/1GB)
 - Database Page (512B-32KB)
- A hardware page is the largest block of data that the storage device can guarantee failsafe writes.

Default DB Page Sizes

4KB	 SQLite	 ORACLE®
 IBM	 DB2	 RocksDB
<hr/>		
8KB	 Microsoft® SQL Server®	
	Postgre	SQL
<hr/>		
16KB	 MySQL™	

Database Pages

- There are three different notions of “pages” in a DBMS:
 - Hardware Page (usually 4KB)
 - OS Page (usually 4KB, x64 2MB/1GB)
 - Database Page (512B-32KB)
- A hardware page is the largest block of data that the storage device can guarantee failsafe writes.

Default DB Page Sizes

4KB	 SQLite	 ORACLE®
 IBM	 DB2	 RocksDB
<hr/>		
8KB	 Microsoft® SQL Server®	
 PostgreSQL		
<hr/>		
16KB	 MySQL™	

Page Storage Architecture

- Different DBMSs manage pages in files on disk in different ways.
 - Heap File Organization
 - Tree File Organization
 - Sequential / Sorted File Organization (ISAM)
 - Hashing File Organization
- At this point in the hierarchy we don't need to know anything about what is inside of the pages.

Page Storage Architecture

- Different DBMSs manage pages in files on disk in different ways.
 - Heap File Organization
 - Tree File Organization
 - Sequential / Sorted File Organization (ISAM)
 - Hashing File Organization
- At this point in the hierarchy we don't need to know anything about what is inside of the pages.

Heap File

- A heap file is an unordered collection of pages with tuples stored in random order.
 - Create / Get / Write / Delete Page
 - Must also support iterating over all pages.
- It is easy to find pages if there is only a single file.
- Need meta-data to track what pages exist in multiple files and which ones have free space.

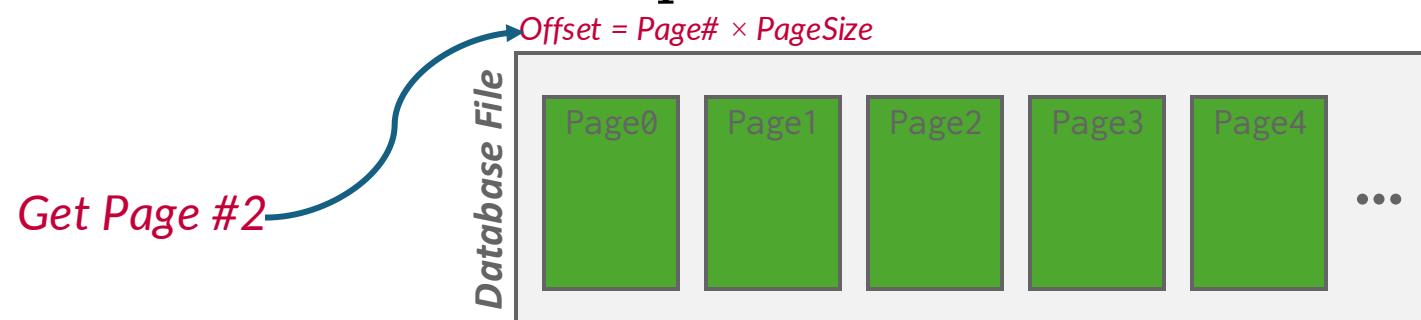
Heap File

- A heap file is an unordered collection of pages with tuples stored in random order.
 - Create / Get / Write / Delete Page
 - Must also support iterating over all pages.
- It is easy to find pages if there is only a single file.
- Need meta-data to track what pages exist in multiple files and which ones have free space.



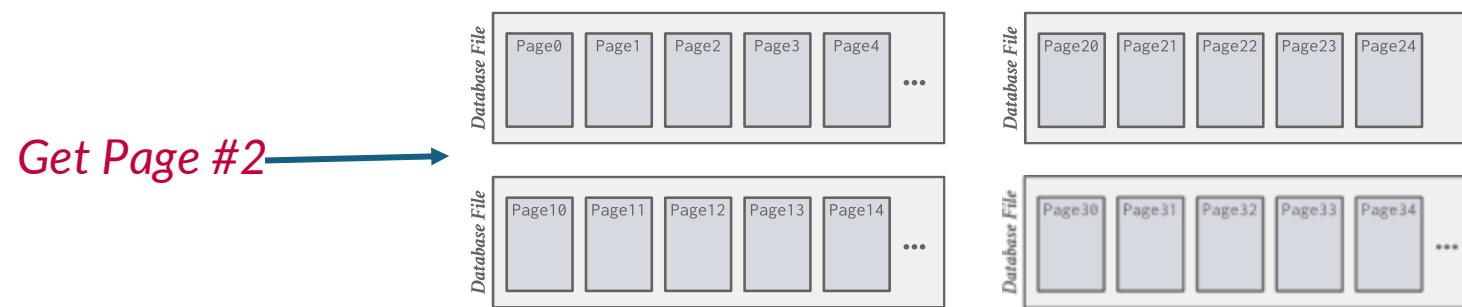
Heap File

- A heap file is an unordered collection of pages with tuples stored in random order.
 - Create / Get / Write / Delete Page
 - Must also support iterating over all pages.
- It is easy to find pages if there is only a single file.
- Need meta-data to track what pages exist in multiple files and which ones have free space.



Heap File

- A heap file is an unordered collection of pages with tuples stored in random order.
 - Create / Get / Write / Delete Page
 - Must also support iterating over all pages.
- It is easy to find pages if there is only a single file.
- Need meta-data to track what pages exist in multiple files and which ones have free space.

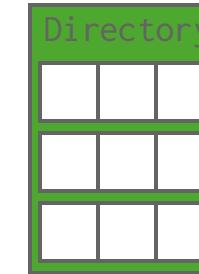


Heap File

- A heap file is an unordered collection of pages with tuples stored in random order.
 - Create / Get / Write / Delete Page
 - Must also support iterating over all pages.
- It is easy to find pages if there is only a single file.
- Need meta-data to track what pages exist in multiple files and which ones have free space.

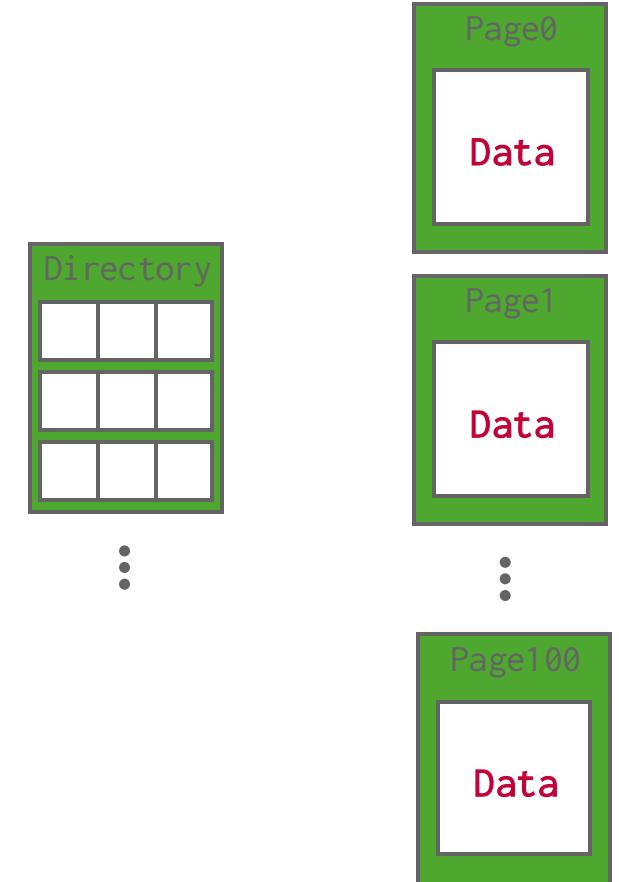
Heap File: Page Directory

- The DBMS maintains special pages that tracks the location of data pages in the database files.
 - Must make sure that the directory pages are in sync with the data pages.
- The directory also records meta-data about available space:
 - The number of free slots per page.
 - List of free / empty pages.



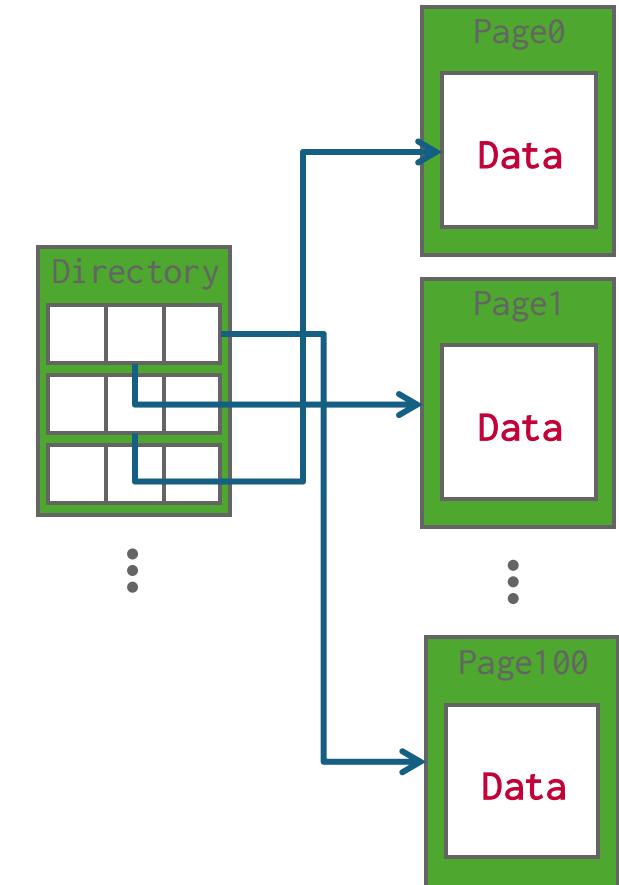
Heap File: Page Directory

- The DBMS maintains special pages that tracks the location of data pages in the database files.
 - Must make sure that the directory pages are in sync with the data pages.
- The directory also records meta-data about available space:
 - The number of free slots per page.
 - List of free / empty pages.



Heap File: Page Directory

- The DBMS maintains special pages that tracks the location of data pages in the database files.
 - Must make sure that the directory pages are in sync with the data pages.
- The directory also records meta-data about available space:
 - The number of free slots per page.
 - List of free / empty pages.



Page Layout

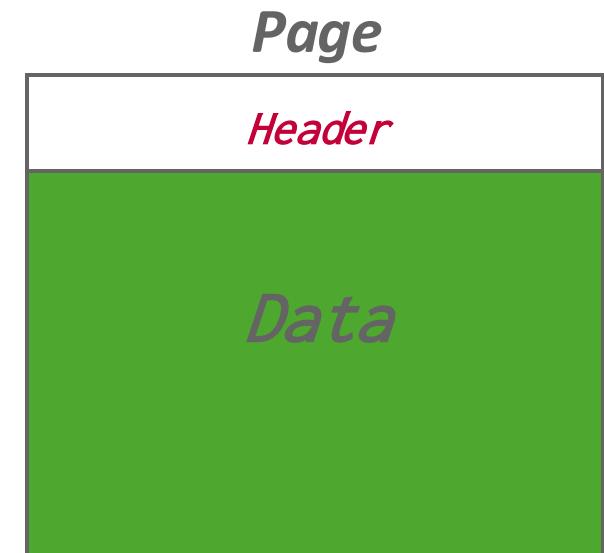
File Storage

Page Layout

Tuple Layout

Page Header

- Every page contains a header of meta-data about the page's contents.
 - Page Size
 - Checksum
 - DBMS Version
 - Transaction Visibility
 - Compression / Encoding Meta-data
 - Schema Information
 - Data Summary / Sketches
- Some systems require pages to be self-contained (e.g., Oracle).



Page Layout

- For any page storage architecture, we now need to decide how to organize the data inside of the page.
 - We are still assuming that we are only storing tuples in a **row-oriented storage model**.

Approach #1: Tuple-oriented Storage

Approach #2: Log-structured Storage

Approach #3: Index-organized Storage

Page Layout

- For any page storage architecture, we now need to decide how to organize the data inside of the page.
 - We are still assuming that we are only storing tuples in a **row-oriented storage model**.

Approach #1: Tuple-oriented Storage ← Today

Approach #2: Log-structured Storage

Approach #3: Index-organized Storage

Page Layout

- For any page storage architecture, we now need to decide how to organize the data inside of the page.
 - We are still assuming that we are only storing tuples in a **row-oriented storage model**.

Approach #1: Tuple-oriented Storage

Approach #2: Log-structured Storage

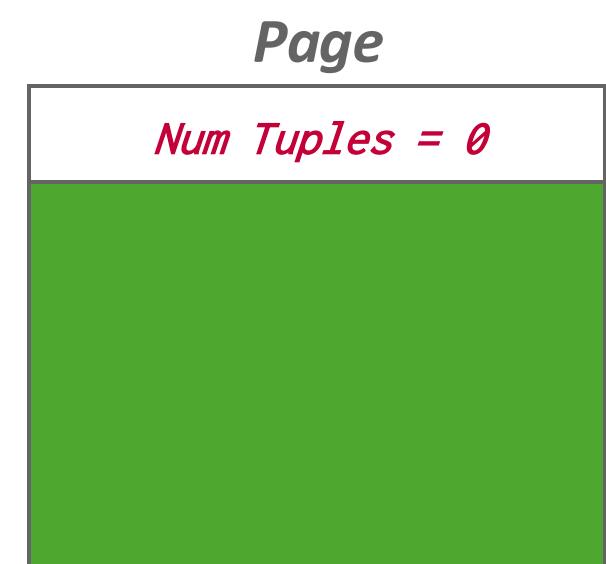
Approach #3: Index-organized Storage

Omitted

Tuple-oriented Storage

How to store tuples in a page?

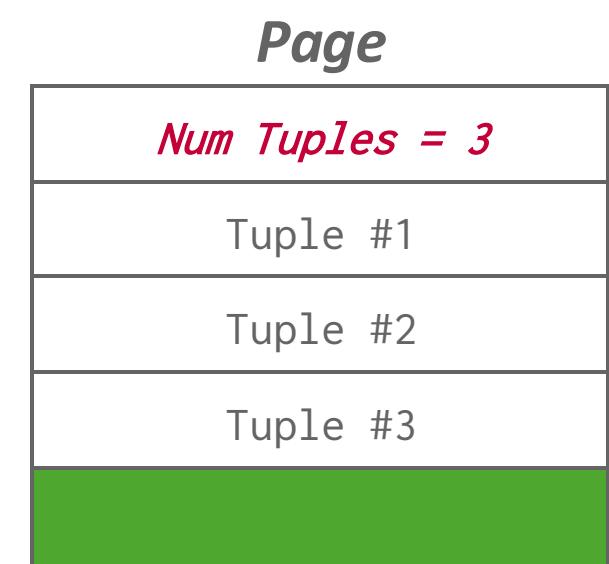
- **Strawman Idea:** Keep track of the number of tuples in a page and then just append a new tuple to the end.
 - What happens if we delete a tuple?
 - What happens if we have a variable-length attribute?



Tuple-oriented Storage

How to store tuples in a page?

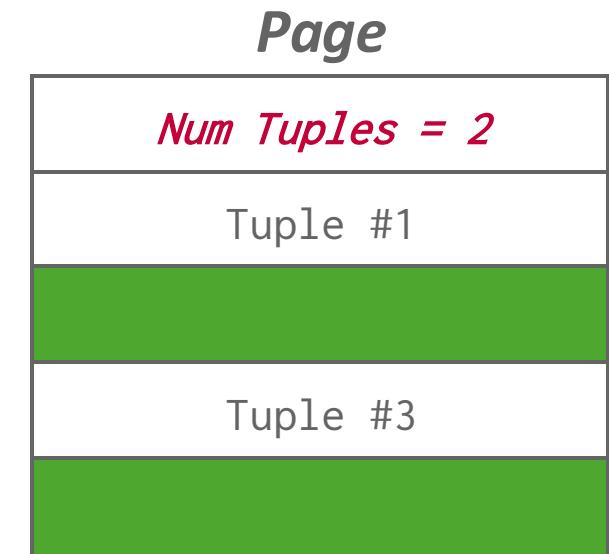
- **Strawman Idea:** Keep track of the number of tuples in a page and then just append a new tuple to the end.
 - What happens if we delete a tuple?
 - What happens if we have a variable-length attribute?



Tuple-oriented Storage

How to store tuples in a page?

- **Strawman Idea:** Keep track of the number of tuples in a page and then just append a new tuple to the end.
 - What happens if we delete a tuple?
 - What happens if we have a variable-length attribute?



Tuple-oriented Storage

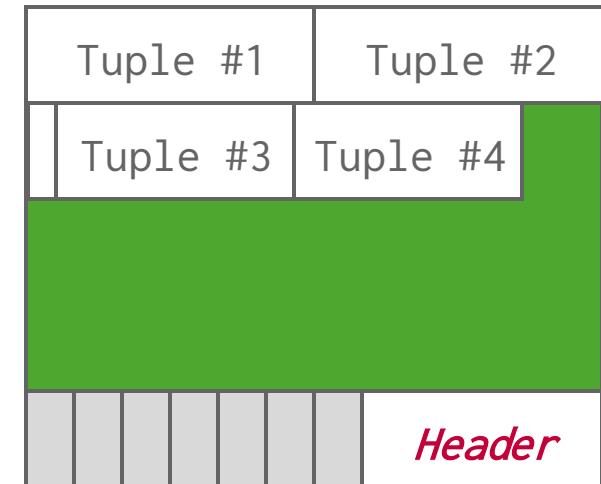
How to store tuples in a page?

- **Strawman Idea:** Keep track of the number of tuples in a page and then just append a new tuple to the end.
 - What happens if we delete a tuple?
 - What happens if we have a variable-length attribute?

<i>Page</i>
<i>Num Tuples = 3</i>
Tuple #1
Tuple #4
Tuple #3

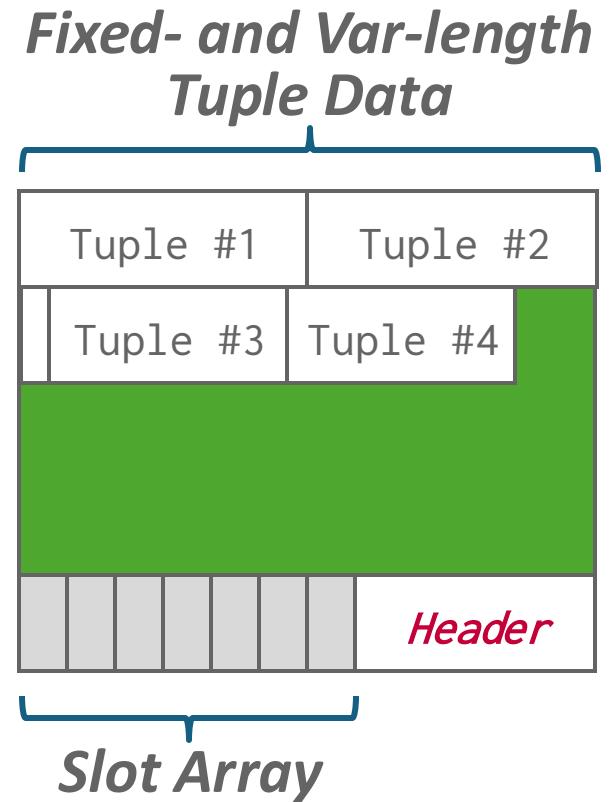
Slotted Pages

- The most common layout scheme is called slotted pages.
- The slot array maps “slots” to the tuples’ starting position offsets.
- The header keeps track of:
 - The # of used slots
 - The offset of the starting location of the last slot used.



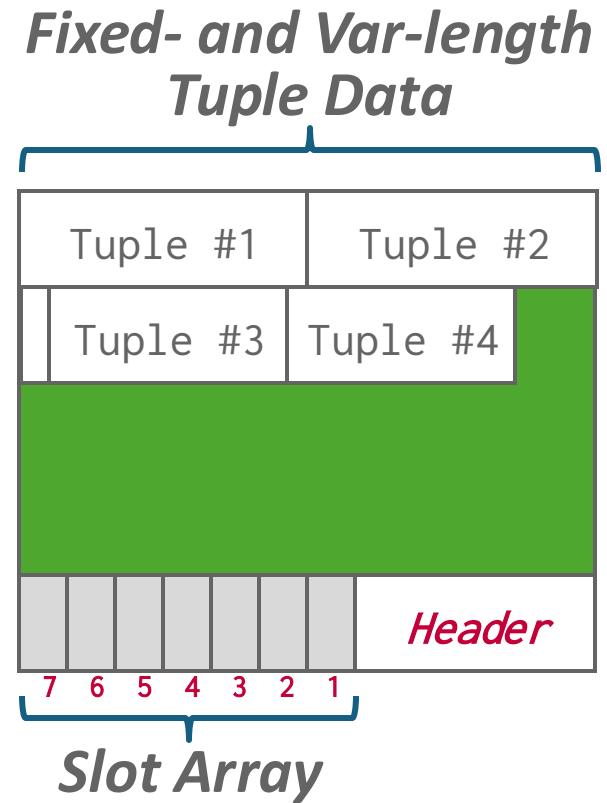
Slotted Pages

- The most common layout scheme is called slotted pages.
- The slot array maps “slots” to the tuples’ starting position offsets.
- The header keeps track of:
 - The # of used slots
 - The offset of the starting location of the last slot used.



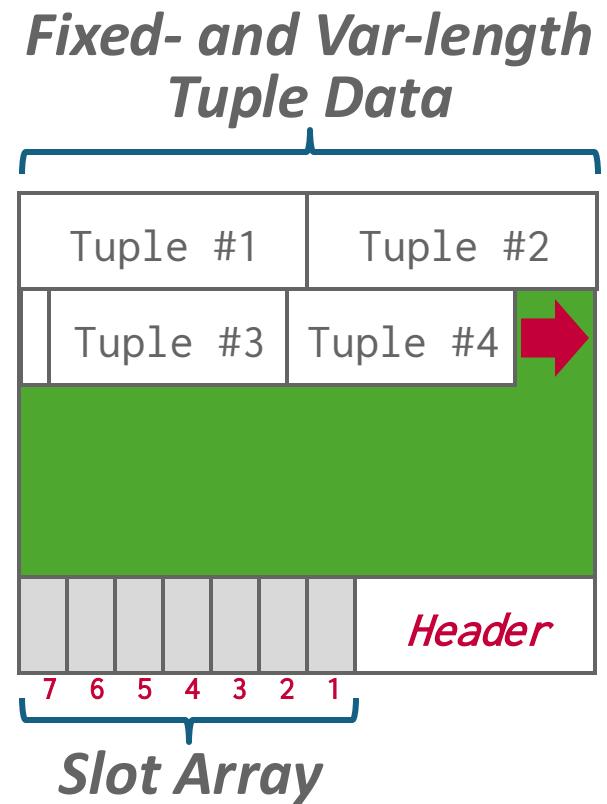
Slotted Pages

- The most common layout scheme is called slotted pages.
- The slot array maps “slots” to the tuples’ starting position offsets.
- The header keeps track of:
 - The # of used slots
 - The offset of the starting location of the last slot used.



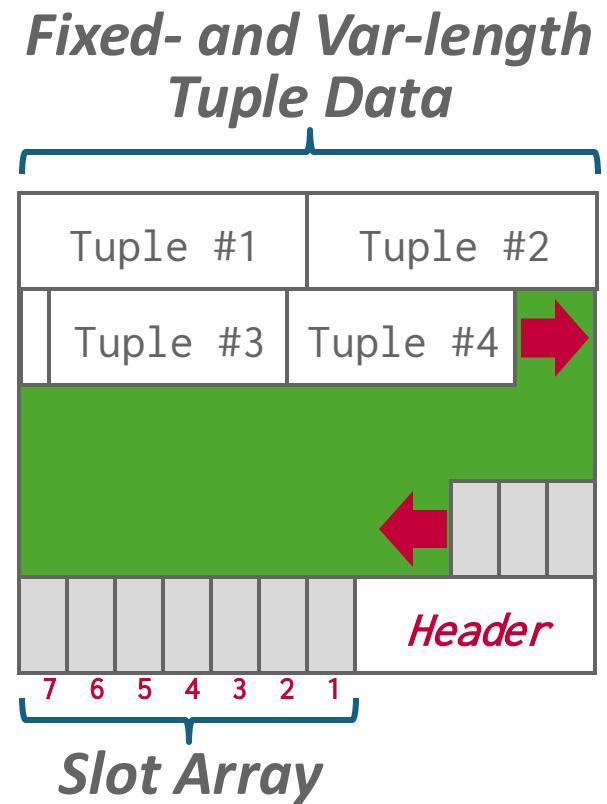
Slotted Pages

- The most common layout scheme is called slotted pages.
- The slot array maps “slots” to the tuples’ starting position offsets.
- The header keeps track of:
 - The # of used slots
 - The offset of the starting location of the last slot used.



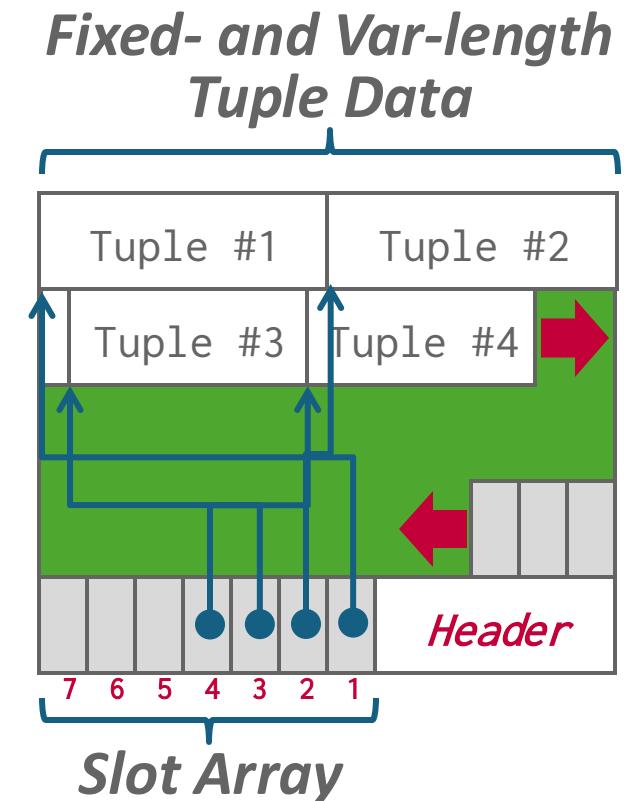
Slotted Pages

- The most common layout scheme is called slotted pages.
- The slot array maps “slots” to the tuples’ starting position offsets.
- The header keeps track of:
 - The # of used slots
 - The offset of the starting location of the last slot used.



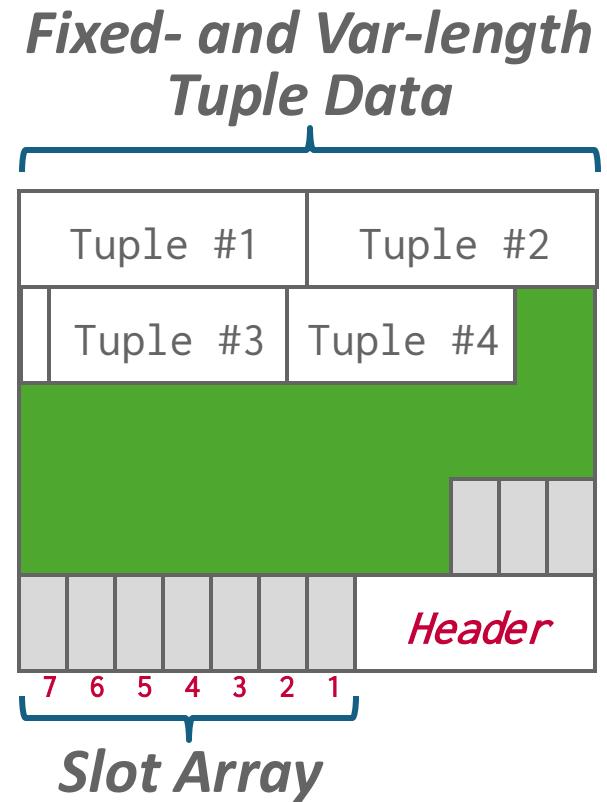
Slotted Pages

- The most common layout scheme is called slotted pages.
- The slot array maps “slots” to the tuples’ starting position offsets.
- The header keeps track of:
 - The # of used slots
 - The offset of the starting location of the last slot used.



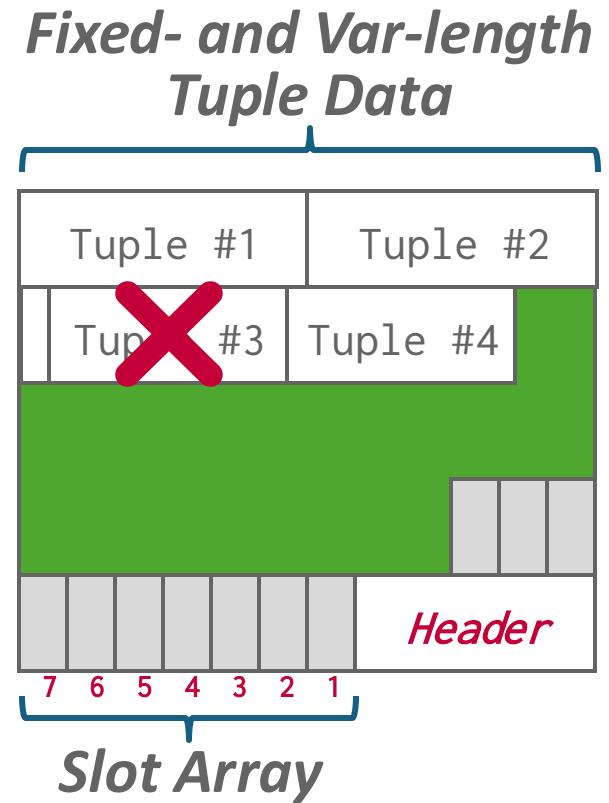
Slotted Pages

- The most common layout scheme is called slotted pages.
- The slot array maps “slots” to the tuples’ starting position offsets.
- The header keeps track of:
 - The # of used slots
 - The offset of the starting location of the last slot used.



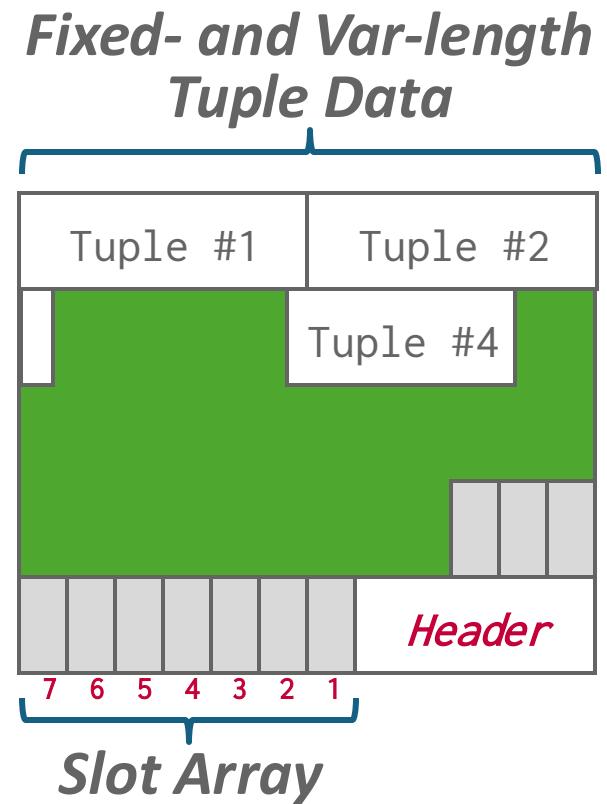
Slotted Pages

- The most common layout scheme is called slotted pages.
- The slot array maps “slots” to the tuples’ starting position offsets.
- The header keeps track of:
 - The # of used slots
 - The offset of the starting location of the last slot used.



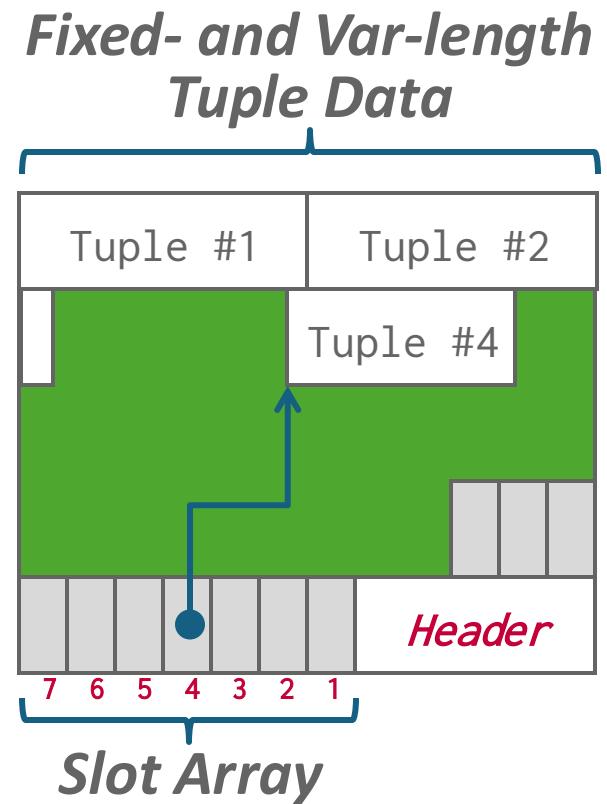
Slotted Pages

- The most common layout scheme is called slotted pages.
- The slot array maps “slots” to the tuples’ starting position offsets.
- The header keeps track of:
 - The # of used slots
 - The offset of the starting location of the last slot used.



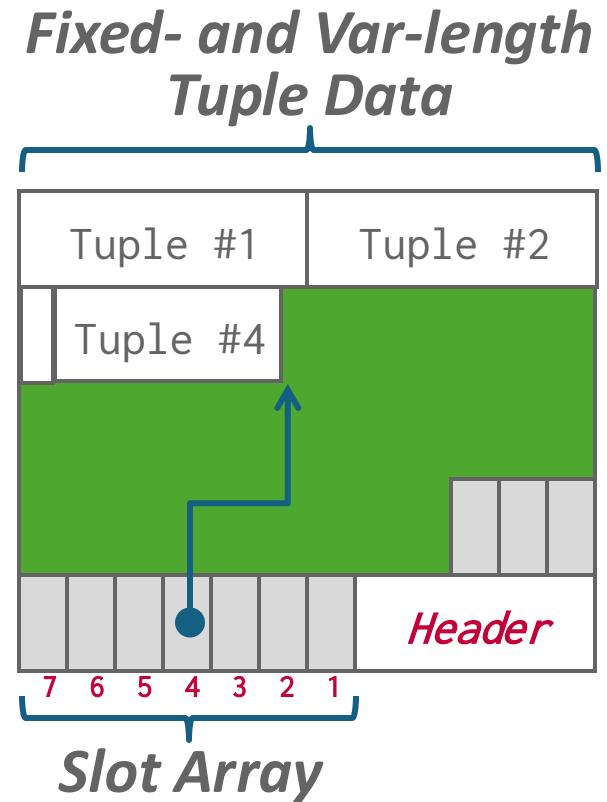
Slotted Pages

- The most common layout scheme is called slotted pages.
- The slot array maps “slots” to the tuples’ starting position offsets.
- The header keeps track of:
 - The # of used slots
 - The offset of the starting location of the last slot used.



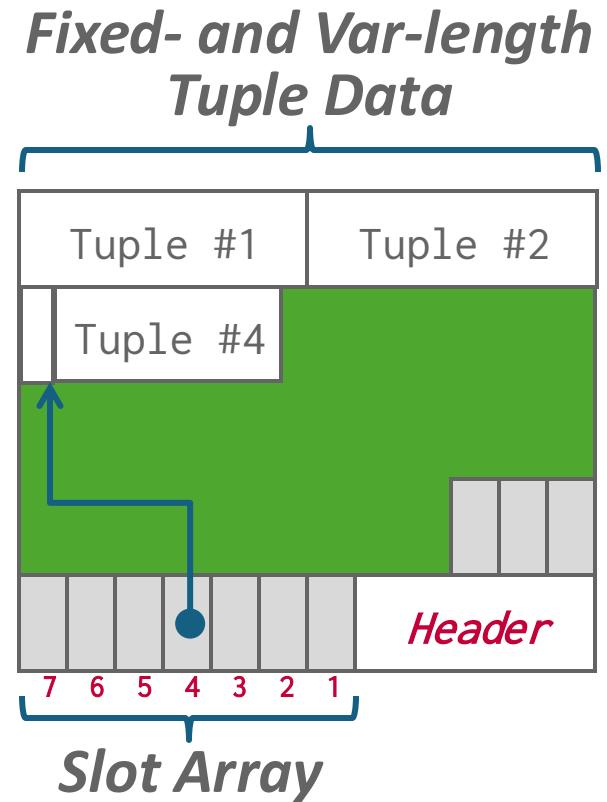
Slotted Pages

- The most common layout scheme is called slotted pages.
- The slot array maps “slots” to the tuples’ starting position offsets.
- The header keeps track of:
 - The # of used slots
 - The offset of the starting location of the last slot used.



Slotted Pages

- The most common layout scheme is called slotted pages.
- The slot array maps “slots” to the tuples’ starting position offsets.
- The header keeps track of:
 - The # of used slots
 - The offset of the starting location of the last slot used.

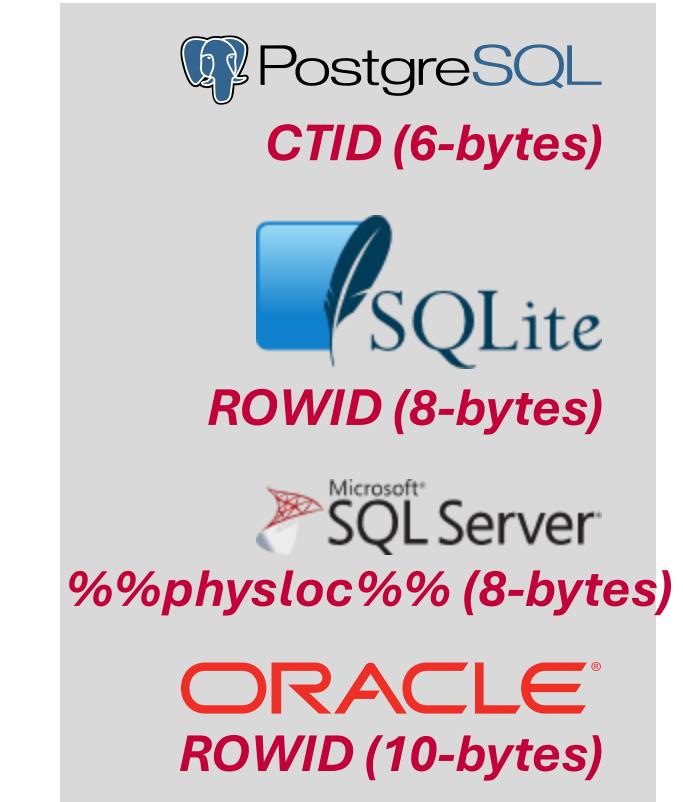


Record IDs

- The DBMS assigns each logical tuple a unique record identifier representing its physical location in the database.
 - File Id, Page Id, Slot #
 - Most DBMSs do not store IDs in tuples.
 - SQLite uses ROWID as the true primary key and stores it as a hidden attribute.
- Applications should never rely on these IDs to mean anything.

Record IDs

- The DBMS assigns each logical tuple a unique record identifier representing its physical location in the database.
 - File Id, Page Id, Slot #
 - Most DBMSs do not store IDs in tuples.
 - SQLite uses ROWID as the true primary key and stores it as a hidden attribute.
- Applications should never rely on these IDs to mean anything.



Tuple Layout

File Storage

Page Layout

Tuple Layout

Tuple Layout

- A tuple is essentially a sequence of bytes.
- It's the job of the DBMS to interpret those bytes into attribute types and values.

Tuple Header

- Each tuple is prefixed with a header that contains meta-data about it.
 - Visibility info (concurrency control)
 - Bit Map for **NULL** values.
- We do not need to store meta-data about the schema.



Tuple Data

- Attributes are typically stored in the order specified in the DDL used to create the table.
- This is done for software engineering reasons (i.e., simplicity).
- However, it might be more efficient to lay them out differently.

Tuple

<i>Header</i>	a	b	c	d	e

```
CREATE TABLE foo (
    a INT PRIMARY KEY,
    b INT NOT NULL,
    c INT,
    d DOUBLE,
    e FLOAT
);
```

Denormalized Tuple Data

- DBMS can physically *denormalize* (e.g., “pre-join”) related tuples and store them together in the same page.
 - Potentially reduces the amount of I/O for common workload patterns.
 - Can make updates more expensive.

```
CREATE TABLE foo (
    a INT PRIMARY KEY,
    b INT NOT NULL,
);
CREATE TABLE bar (
    c INT PRIMARY KEY,
    a INT
        ↗ REFERENCES foo (a),
);
;
```

Denormalized Tuple Data

- DBMS can physically *denormalize* (e.g., “pre-join”) related tuples and store them together in the same page.
 - Potentially reduces the amount of I/O for common workload patterns.
 - Can make updates more expensive.

```
CREATE TABLE foo (
    a INT PRIMARY KEY,
    b INT NOT NULL,
);
CREATE TABLE bar (
    c INT PRIMARY KEY,
    a INT
    ↗ REFERENCES foo (a),
);

```

Denormalized Tuple Data

- DBMS can physically *denormalize* (e.g., “pre-join”) related tuples and store them together in the same page.
 - Potentially reduces the amount of I/O for common workload patterns.
 - Can make updates more expensive.

foo		
Header	a	b
Header	c	a
Header	c	a
Header	c	a

bar		
Header	c	a
Header	c	a
Header	c	a
Header	c	a

Denormalized Tuple Data

- DBMS can physically *denormalize* (e.g., “pre-join”) related tuples and store them together in the same page.
 - Potentially reduces the amount of I/O for common workload patterns.
 - Can make updates more expensive.

```
SELECT * FROM foo JOIN bar
ON foo.a = bar.a;
```

foo

<i>Header</i>	a	b
<i>Header</i>	c	a
<i>Header</i>	c	a
<i>Header</i>	c	a

bar

Denormalized Tuple Data

foo

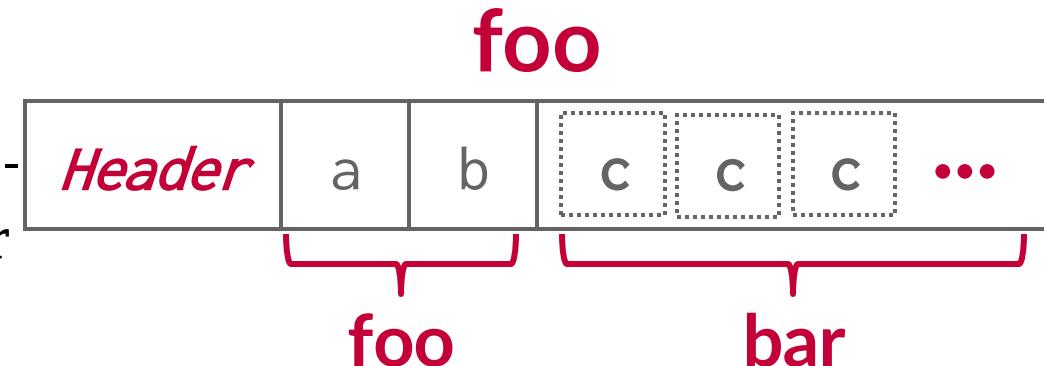
<i>Header</i>	a	b	c	c	c	...
---------------	---	---	---	---	---	-----

- DBMS can physically *denormalize* (e.g., “pre-join”) related tuples and store them together in the same page.
 - Potentially reduces the amount of I/O for common workload patterns.
 - Can make updates more expensive.

```
SELECT * FROM foo JOIN bar
  ON foo.a = bar.a;
```

Denormalized Tuple Data

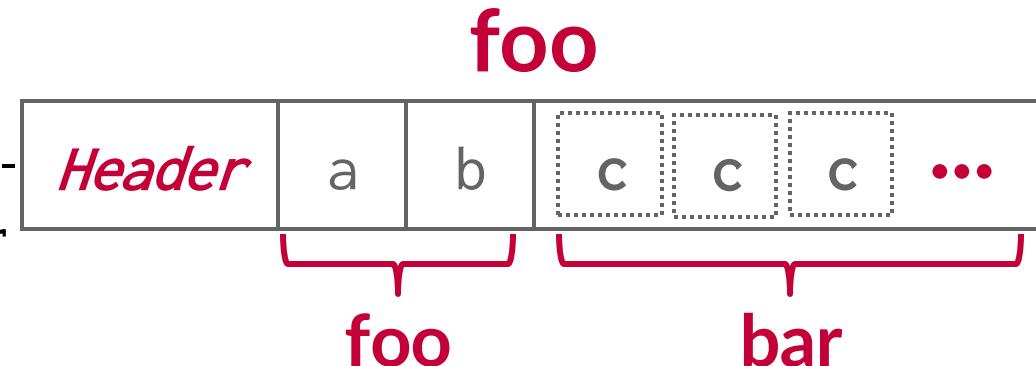
- DBMS can physically *denormalize* (e.g., “pre-join”) related tuples and store them together in the same page.
 - Potentially reduces the amount of I/O for common workload patterns.
 - Can make updates more expensive.



```
SELECT * FROM foo JOIN bar
  ON foo.a = bar.a;
```

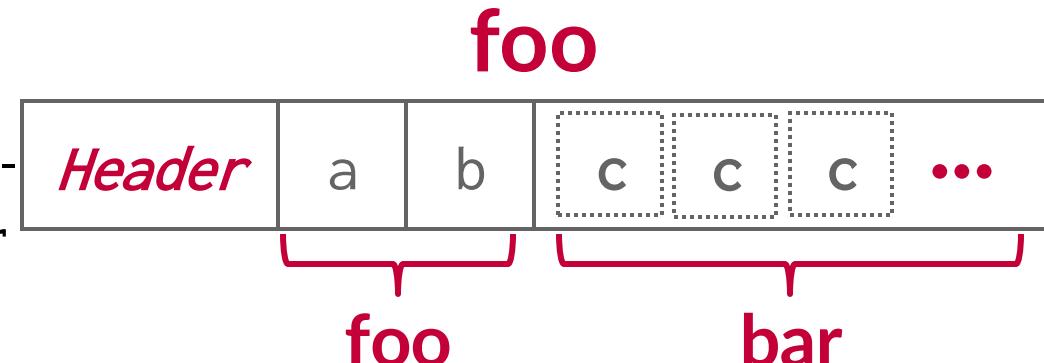
Denormalized Tuple Data

- DBMS can physically *denormalize* (e.g., “pre-join”) related tuples and store them together in the same page.
 - Potentially reduces the amount of I/O for common workload patterns.
 - Can make updates more expensive.
 - Not a new idea.
 - IBM System R did this in the 1970s.
 - Several NoSQL DBMSs do this without calling it physical denormalization.



Denormalized Tuple Data

- DBMS can physically *denormalize* (e.g., “pre-join”) related tuples and store them together in the same page.
 - Potentially reduces the amount of I/O for common workload patterns.
 - Can make updates more expensive.
- Not a new idea.
 - IBM System R did this in the 1970s.
 - Several NoSQL DBMSs do this without calling it physical denormalization.



Tuple Storage

- A tuple is essentially a sequence of bytes.
- It is the job of the DBMS to interpret those bytes into attribute types and values.
- The DBMS's catalogs contain the schema information about tables that the system uses to figure out the tuple's layout.

Word-aligned Tuples

- All attributes in a tuple must be word aligned to enable the CPU to access it without any unexpected behavior or additional work.

```
CREATE TABLE foo (
    id INT PRIMARY KEY,
    cdate TIMESTAMP,
    color CHAR(2),
    zipcode INT
);
```

char[]



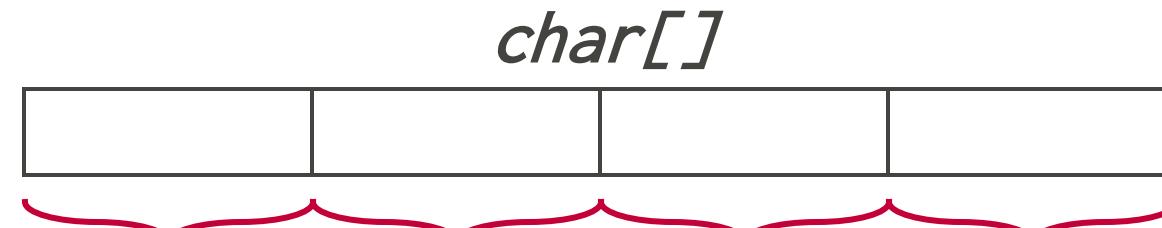
In the old days, the DBMS programmer had to worry about “unaligned word memory reference.”

Today: Processors handle it. It will read multiple words from memory, so it may have a performance impact.

Word-aligned Tuples

- All attributes in a tuple must be word aligned to enable the CPU to access it without any unexpected behavior or additional work.

```
CREATE TABLE foo (
    id INT PRIMARY KEY,
    cdate TIMESTAMP,
    color CHAR(2),
    zipcode INT
);
```



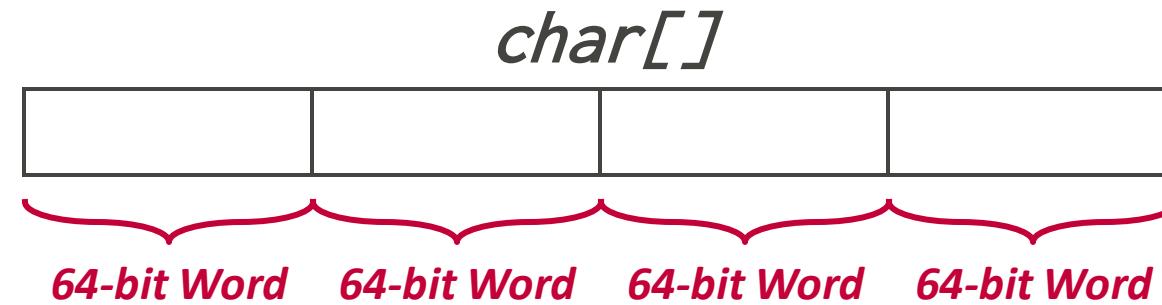
In the old days, the DBMS programmer had to worry about “unaligned word memory reference.”

Today: Processors handle it. It will read multiple words from memory, so it may have a performance impact.

Word-aligned Tuples

- All attributes in a tuple must be word aligned to enable the CPU to access it without any unexpected behavior or additional work.

```
CREATE TABLE foo (
32-bits id INT PRIMARY KEY,
    cdate TIMESTAMP,
    color CHAR(2),
    zipcode INT
);
```



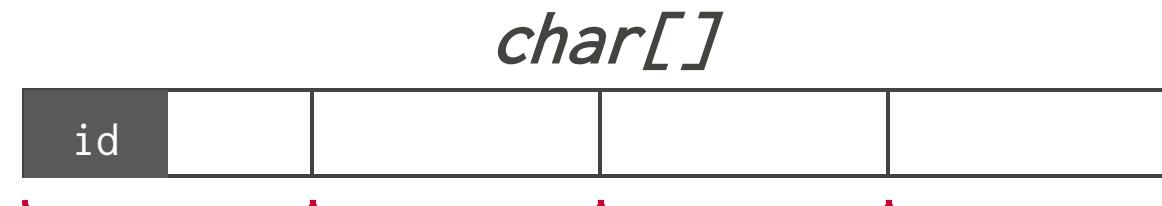
In the old days, the DBMS programmer had to worry about “unaligned word memory reference.”

Today: Processors handle it. It will read multiple words from memory, so it may have a performance impact.

Word-aligned Tuples

- All attributes in a tuple must be word aligned to enable the CPU to access it without any unexpected behavior or additional work.

```
CREATE TABLE foo (
32-bits id INT PRIMARY KEY,
    cdate TIMESTAMP,
    color CHAR(2),
    zipcode INT
);
```



64-bit Word 64-bit Word 64-bit Word 64-bit Word

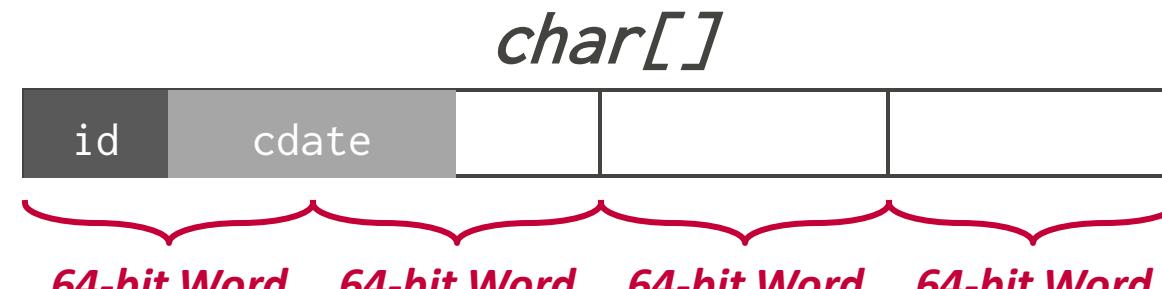
In the old days, the DBMS programmer had to worry about “unaligned word memory reference.”

Today: Processors handle it. It will read multiple words from memory, so it may have a performance impact.

Word-aligned Tuples

- All attributes in a tuple must be word aligned to enable the CPU to access it without any unexpected behavior or additional work.

```
CREATE TABLE foo (
  id INT PRIMARY KEY,
  cdate TIMESTAMP,
  color CHAR(2),
  zipcode INT
);
```



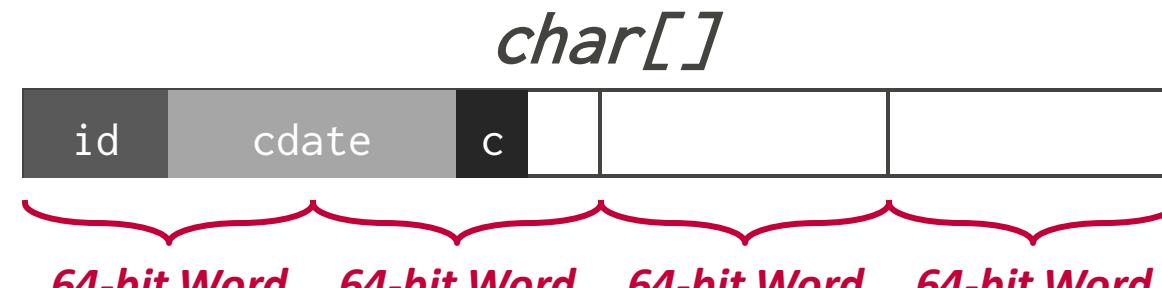
In the old days, the DBMS programmer had to worry about “unaligned word memory reference.”

Today: Processors handle it. It will read multiple words from memory, so it may have a performance impact.

Word-aligned Tuples

- All attributes in a tuple must be word aligned to enable the CPU to access it without any unexpected behavior or additional work.

```
CREATE TABLE foo (
  id INT PRIMARY KEY,
  cdate TIMESTAMP,
  color CHAR(2),
  zipcode INT
);
```



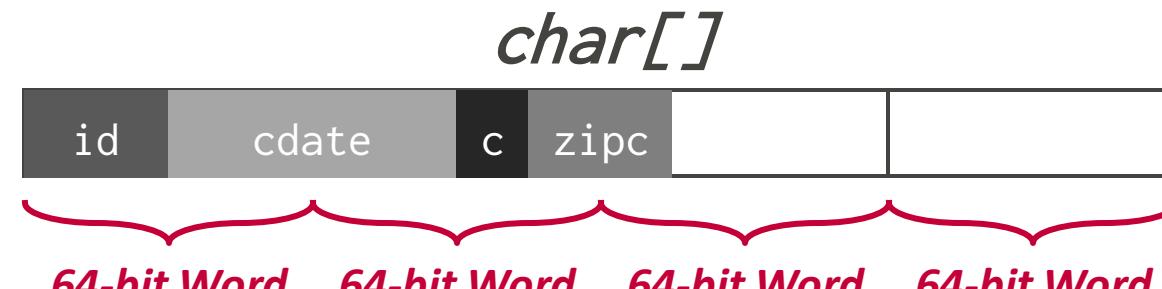
In the old days, the DBMS programmer had to worry about “unaligned word memory reference.”

Today: Processors handle it. It will read multiple words from memory, so it may have a performance impact.

Word-aligned Tuples

- All attributes in a tuple must be word aligned to enable the CPU to access it without any unexpected behavior or additional work.

```
CREATE TABLE foo (
  32-bits id INT PRIMARY KEY,
  64-bits cdate TIMESTAMP,
  16-bits color CHAR(2),
  32-bits zipcode INT
);
```



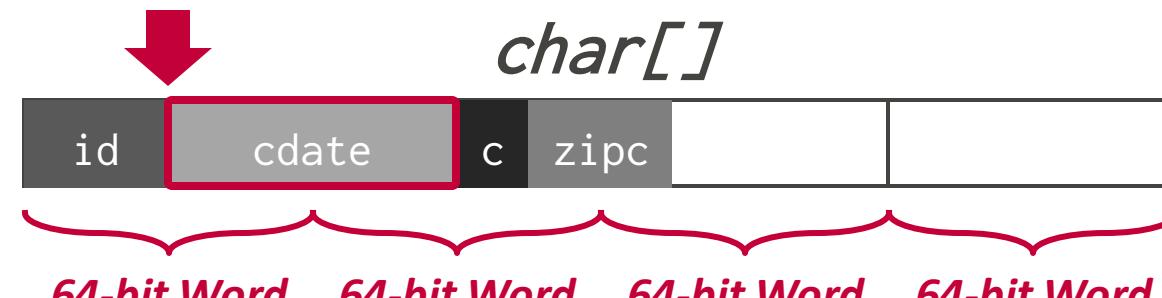
In the old days, the DBMS programmer had to worry about “unaligned word memory reference.”

Today: Processors handle it. It will read multiple words from memory, so it may have a performance impact.

Word-aligned Tuples

- All attributes in a tuple must be word aligned to enable the CPU to access it without any unexpected behavior or additional work.

```
CREATE TABLE foo (
  32-bits id INT PRIMARY KEY,
  64-bits cdate TIMESTAMP,
  16-bits color CHAR(2),
  32-bits zipcode INT
);
```



In the old days, the DBMS programmer had to worry about “unaligned word memory reference.”

Today: Processors handle it. It will read multiple words from memory, so it may have a performance impact.

Word-alignment: Padding

- Add empty bits after attributes to ensure that tuple is word aligned.

```
CREATE TABLE dj2pl (
32-bits id INT PRIMARY KEY,
64-bits cdate TIMESTAMP,
16-bits color CHAR(2),
32-bits zipcode INT
);
```



Word-alignment: Reordering

- Switch the order of attributes in the tuples' physical layout to make sure they are aligned.
 - May still have to use padding.

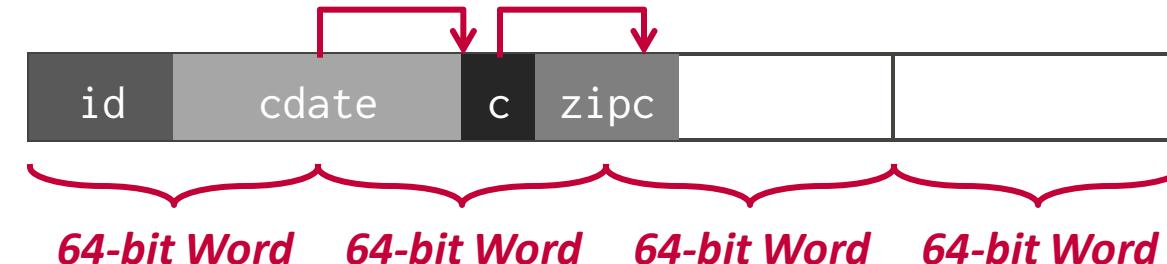
```
CREATE TABLE dj2pl (
32-bits id INT PRIMARY KEY,
64-bits cdate TIMESTAMP,
16-bits color CHAR(2),
32-bits zipcode INT
);
```



Word-alignment: Reordering

- Switch the order of attributes in the tuples' physical layout to make sure they are aligned.
 - May still have to use padding.

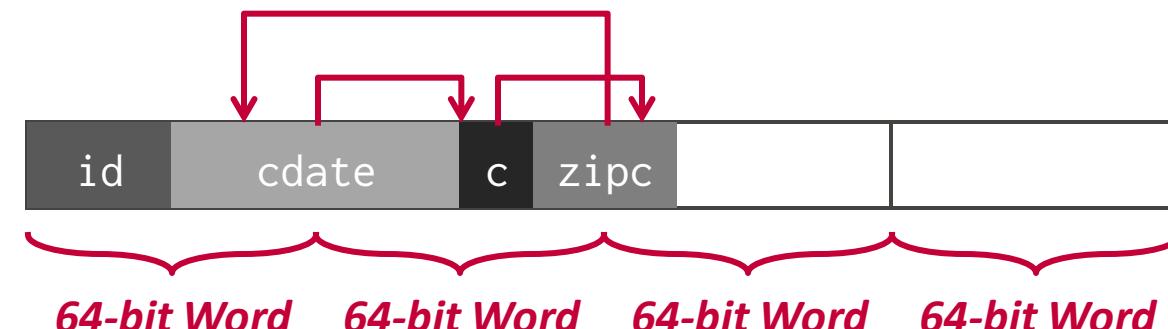
```
CREATE TABLE dj2pl (
32-bits id INT PRIMARY KEY,
64-bits cdate TIMESTAMP,
16-bits color CHAR(2),
32-bits zipcode INT
);
```



Word-alignment: Reordering

- Switch the order of attributes in the tuples' physical layout to make sure they are aligned.
 - May still have to use padding.

```
CREATE TABLE dj2pl (
32-bits id INT PRIMARY KEY,
64-bits cdate TIMESTAMP,
16-bits color CHAR(2),
32-bits zipcode INT
);
```



Word-alignment: Reordering

- Switch the order of attributes in the tuples' physical layout to make sure they are aligned.
 - May still have to use padding.

```
CREATE TABLE dj2pl (
32-bits id INT PRIMARY KEY,
64-bits cdate TIMESTAMP,
16-bits color CHAR(2),
32-bits zipcode INT
);
```



Word-alignment: Reordering

- Switch the order of attributes in the tuples' physical layout to make sure they are aligned.
 - May still have to use padding.

```
CREATE TABLE dj2pl (
32-bits id INT PRIMARY KEY,
64-bits cdate TIMESTAMP,
16-bits color CHAR(2),
32-bits zipcode INT
);
```



Data Representation

- **INTEGER/BIGINT/SMALLINT/TINYINT**
 - Same as in C/C++.
- **FLOAT/REAL** vs. **NUMERIC/DECIMAL**
 - IEEE-754 Standard / Fixed-point Decimals.
- **VARCHAR/VARBINARY/TEXT/BLOB**
 - Header with length, followed by data bytes OR pointer to another page/offset with data.
 - Need to worry about collations / sorting.
- **TIME/DATETIME/TIMESTAMP/INTERVAL**
 - 32/64-bit integer of (micro/milli)-seconds since Unix epoch (January 1st, 1970).

Data Representation

- **INTEGER/BIGINT/SMALLINT/TINYINT**
 - Same as in C/C++.
- **FLOAT/REAL vs. NUMERIC/DECIMAL**
 - IEEE-754 Standard / Fixed-point Decimals.
- **VARCHAR/VARBINARY/TEXT/BLOB**
 - Header with length, followed by data bytes OR pointer to another page/offset with data.
 - Need to worry about collations / sorting.
- **TIME/DATETIME/TIMESTAMP/INTERVAL**
 - 32/64-bit integer of (micro/milli)-seconds since Unix epoch (January 1st, 1970).

Variable Precision Numbers

- Inexact, variable-precision numeric type that uses the “native” C/C++ types.
- Store directly as specified by [IEEE-754](#).
 - Example: **FLOAT, REAL/DOUBLE**
- These types are typically faster than fixed precision numbers because CPU ISA’s (Xeon, Arm) have instructions / registers to support them.
- But they do not guarantee exact values...

Variable Precision Numbers

Rounding Example

```
#include <stdio.h>

int main(int argc, char* argv[]) {
    float x = 0.1;
    float y = 0.2;
    printf("x+y = %f\n", x+y);
    printf("0.3 = %f\n", 0.3);
}
```

Variable Precision Numbers

Rounding Example

```
#include <stdio.h>

int main(int argc, char* argv[]) {
    float x = 0.1;
    float y = 0.2;
    printf("x+y = %f\n", x+y);
    printf("0.3 = %f\n", 0.3);
}
```

Output

```
x+y = 0.300000
0.3 = 0.300000
```

Variable Precision Numbers

Rounding Example

```
#include <stdio.h>  
  
int main(int argc, char* argv[]) {  
    float x = 0.1;  
    float y = 0.2;  
    printf("x+y = %.20f\n", x+y);  
    printf("0.3 = %.20f\n", 0.3);  
}
```

Output

```
x+y = 0.300000  
0.3 = 0.300000
```

Variable Precision Numbers

Rounding Example

```
#include <stdio.h>

in #include <stdio.h>

    int main(int argc, char* argv[]) {
        float x = 0.1;
        float y = 0.2;
        printf("x+y = %.20f\n", x+y);
        printf("0.3 = %.20f\n", 0.3);
    }
}
```

Output

```
x+y = 0.300000
0.3 = 0.300000
```

```
x+y = 0.3000001192092895508
0.3 = 0.2999999999999998890
```

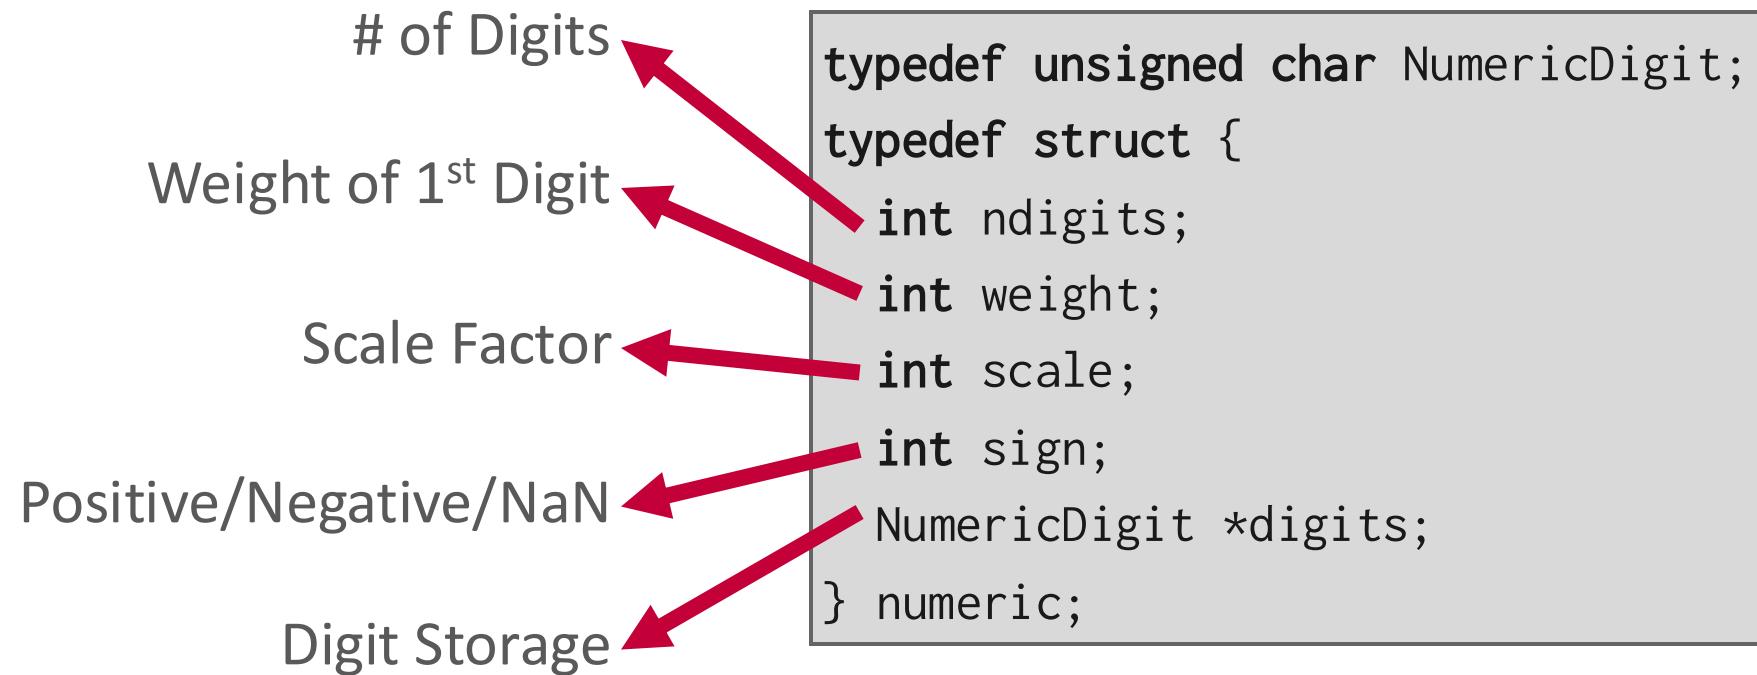
Fixed Precision Numbers

- Numeric data types with (potentially) arbitrary precision and scale. Used when rounding errors are unacceptable.
 - Example: **NUMERIC, DECIMAL**
- Many different implementations.
 - Example: Store in an exact, variable-length binary representation with additional meta-data.
 - Can be less expensive if the DBMS does not provide arbitrary precision (e.g., decimal point can be in a different position per value).

Postgres: NUMERIC

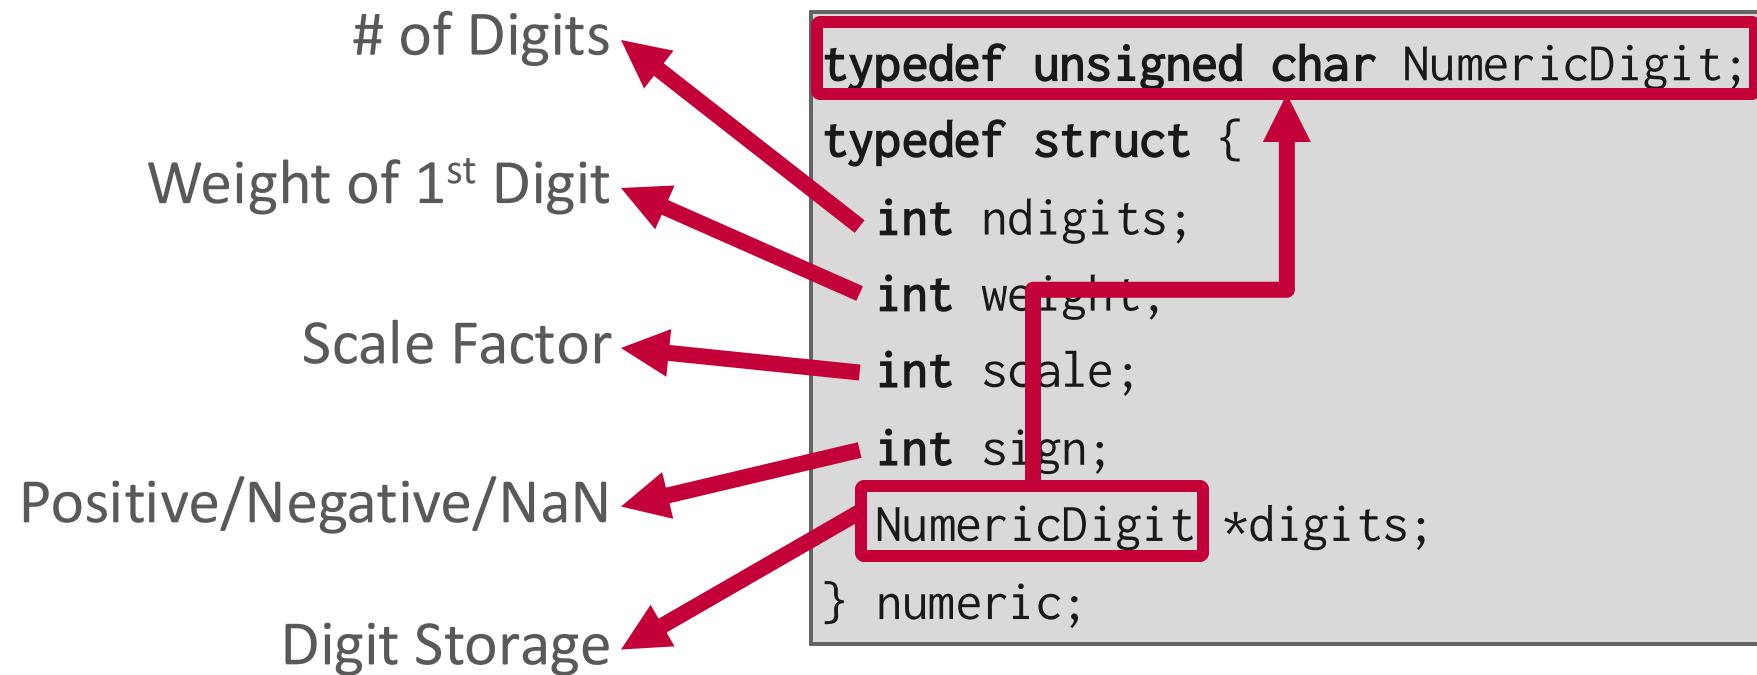
```
typedef unsigned char NumericDigit;  
typedef struct {  
    int ndigits;  
    int weight;  
    int scale;  
    int sign;  
    NumericDigit *digits;  
} numeric;
```

Postgres: NUMERIC



```
typedef unsigned char NumericDigit;
typedef struct {
    int ndigits;
    int weight;
    int scale;
    int sign;
    NumericDigit *digits;
} numeric;
```

Postgres: NUMERIC



```
typedef unsigned char NumericDigit;
typedef struct {
    int ndigits;
    int weight,
    int scale;
    int sign;
    NumericDigit *digits;
} numeric;
```

NULL Data Types

- **Choice #1: Null Column Bitmap Header**

- Store a bitmap in a centralized header that specifies what attributes are null.
 - This is the most common approach.

- **Choice #2: Special Values**

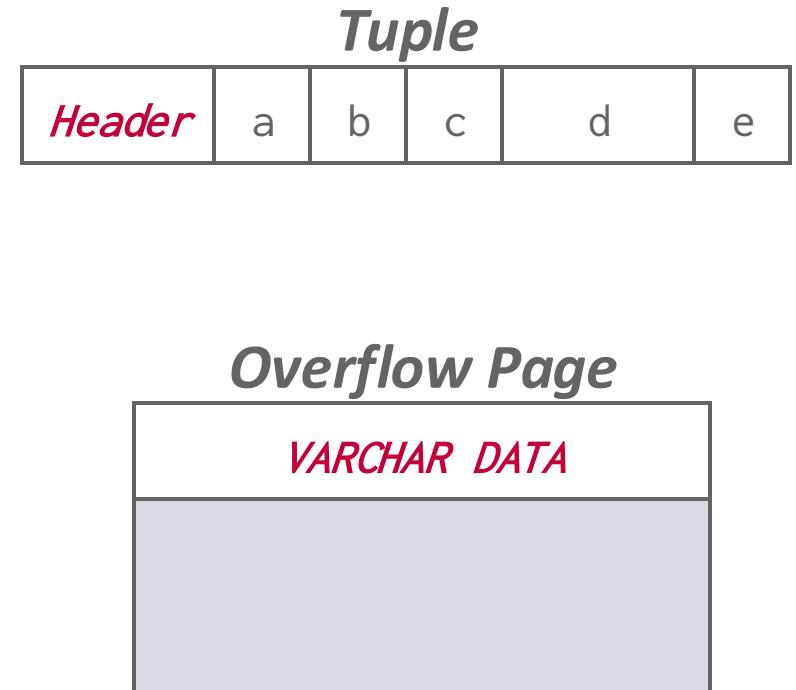
- Designate a value to represent **NULL** for a data type (e.g., `INT32_MIN`).

- **Choice #3: Per Attribute Null Flag**

- Store a flag that marks that a value is null.
- Must use more space than just a single bit because this messes up with word alignment.

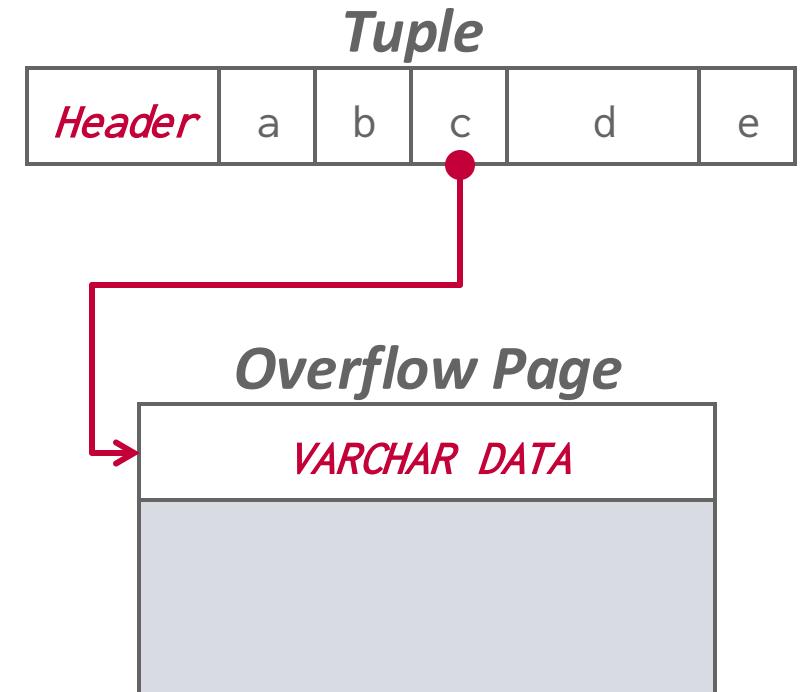
Large Values

- Most DBMSs don't allow a tuple to exceed the size of a single page.
- To store values that are larger than a page, the DBMS uses separate **overflow** storage pages.
 - Postgres: TOAST (>2KB)
 - MySQL: Overflow (> $\frac{1}{2}$ size of page)
 - SQL Server: Overflow (>size of page)



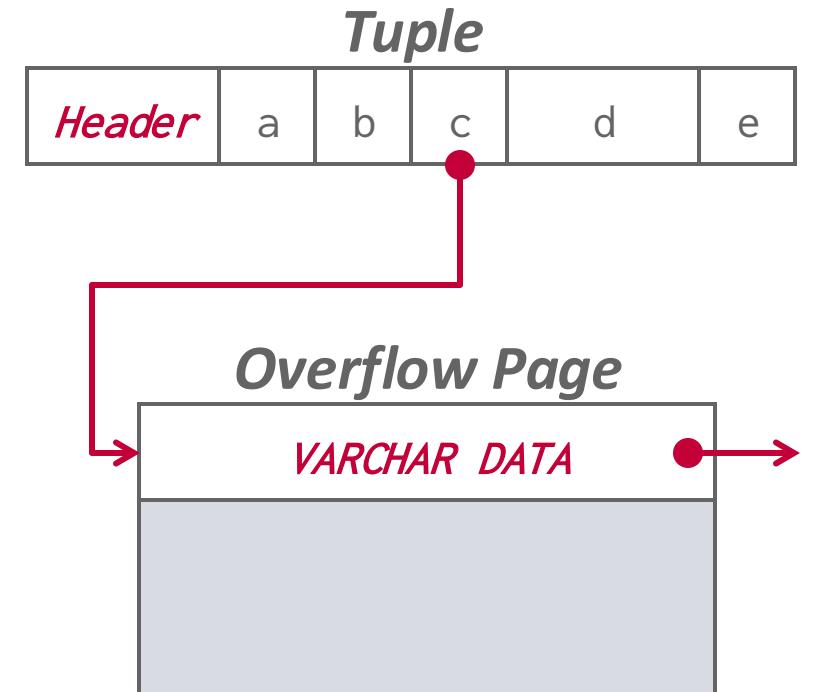
Large Values

- Most DBMSs don't allow a tuple to exceed the size of a single page.
- To store values that are larger than a page, the DBMS uses separate **overflow** storage pages.
 - Postgres: TOAST (>2KB)
 - MySQL: Overflow (> $\frac{1}{2}$ size of page)
 - SQL Server: Overflow (>size of page)



Large Values

- Most DBMSs don't allow a tuple to exceed the size of a single page.
- To store values that are larger than a page, the DBMS uses separate **overflow** storage pages.
 - Postgres: TOAST (>2KB)
 - MySQL: Overflow (> $\frac{1}{2}$ size of page)
 - SQL Server: Overflow (>size of page)

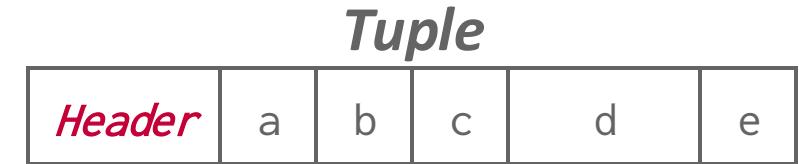


External Value Storage

- Some systems allow you to store a large value in an external file.
Treated as a **BLOB** type.

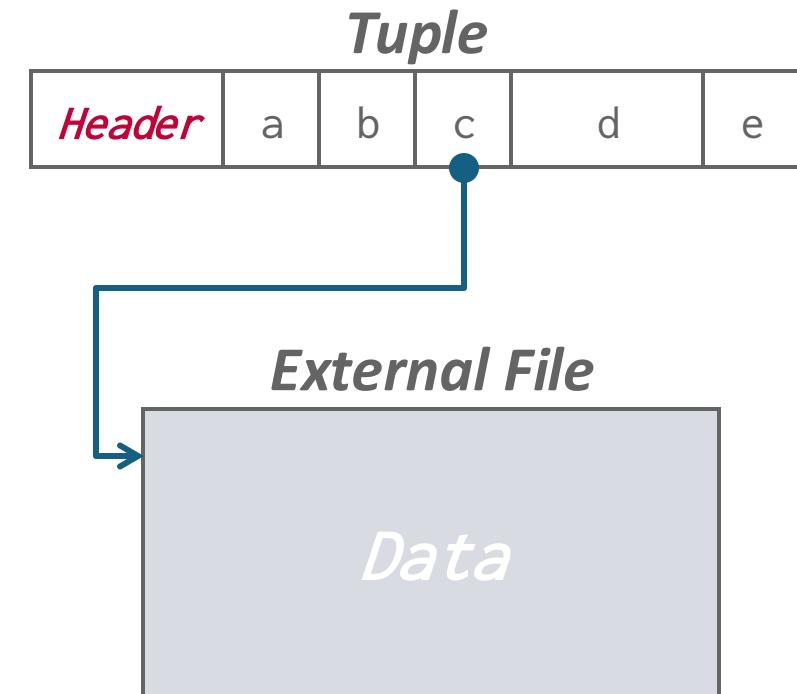
→ Oracle: **BFILE** data type
→ Microsoft: **FILESTREAM** data type

- The DBMS **cannot** manipulate the contents of an external file.
 - No durability protections.
 - No transaction protections.



External Value Storage

- Some systems allow you to store a large value in an external file.
Treated as a **BLOB** type.
 - Oracle: **BFILE** data type
 - Microsoft: **FILESTREAM** data type
- The DBMS cannot manipulate the contents of an external file.
 - No durability protections.
 - No transaction protections.



External Value Storage

- Some systems allow you to store a large value in an external file.
Treated as a **BLOB** type.
 - Oracle: **BFILE** data type
 - Microsoft: **FILESTREAM** data type
- The DBMS cannot manipulate the contents of an external file.
 - No durability protections.
 - No transaction protections.

**To BLOB or Not To BLOB:
Large Object Storage in a Database or a Filesystem?**

Russell Sears¹, Catharina van Ingen¹, Jim Gray¹
¹: Microsoft Research, 2: University of California at Berkeley
 sears@cs.berkeley.edu, vaningen@microsoft.com, gray@microsoft.com
 MSR-TR-2006-45
 April 2006 Revised June 2006

Abstract

Application designers must decide whether to store large objects (BLOBs) in a filesystem or in a database. Generally, this decision is based on factors such as application simplicity or maintainability. Often, system performance affects these factors.

Folklore tells us that databases efficiently handle large numbers of small objects, while filesystems are more efficient for large objects. What is the break-even point? When is accessing a BLOB stored as a file cheaper than accessing a BLOB stored as a database record?

Of course, this depends on the particular filesystem, database system, and workload in question. This study shows that when comparing the NTFS filesystem and SQL Server 2005 database system on a *create*, *(read, replace)** *delete* workload, BLOBs smaller than 256KB are more efficiently handled by SQL Server, while NTFS is most efficient BLOBS larger than 1MB. Of course, this break-even point will vary among different database systems, filesystems, and workloads.

By measuring the performance of a storage server workload typical of web applications which use get/put protocols such as WebDAV [WebDAV], we found that the break-even point depends on many factors. However, our experiments suggest that *storage age*, the ratio of bytes in deleted to replaced objects to bytes in live objects, is dominant. As storage age increases, fragmentation tends to increase. The filesystem we study has better fragmentation control than the database we used, suggesting the database system would benefit from incorporating ideas from filesystem architecture. Conversely, filesystem performance may be improved by using database techniques to handle small files.

Surprisingly, for these studies, when average object size is held constant, the distribution of object sizes did not significantly affect performance. We also found that, in addition to low percentage free space, a low ratio of free space to average object size leads to fragmentation and performance degradation.

2

Catalogs

System Catalogs

- A DBMS stores meta-data about databases in its internal catalogs.
 - Tables, columns, indexes, views
 - Users, permissions
 - Internal statistics
- Almost every DBMS stores the database's catalog inside itself (i.e., as tables).
 - Wrap object abstraction around tuples.
 - Specialized code for “bootstrapping” catalog tables.

System Catalogs

- You can query the DBMS's internal **INFORMATION_SCHEMA** catalog to get info about the database.
 - ANSI standard set of read-only views that provide info about all the tables, views, columns, and procedures in a database
- DBMSs also have non-standard shortcuts to retrieve this information.

Accessing Table Schema

- *List all the tables in the current database:*

```
SELECT *                                SQL-92
      FROM INFORMATION_SCHEMA.TABLES
     WHERE table_catalog = '<db name>';
```

```
\d;                                     Postgres
```

```
SHOW TABLES;                            MySQL
```

```
.tables                                SQLite
```

Accessing Table Schema

- *List the schema of the student table:*

```
SELECT *                                SQL-92
      FROM INFORMATION_SCHEMA.TABLES
     WHERE table_name = 'student'
```

```
\d student;                            Postgres
```

```
DESCRIBE student;                      MySQL
```

```
.schema student                          SQLite
```

Conclusion

- Database is organized in pages.
 - Different ways to track pages.
 - Different ways to store pages.
 - Different ways to store tuples.
-
- The storage manager is not entirely independent from the rest of the DBMS.