



香港中文大學(深圳)
The Chinese University of Hong Kong, Shenzhen



Ack: Prof. Jignesh Patel @ CMU
Prof. Andy Pavlo @CMU

CSC3170

6: Memory & Disk I/O Management


Chenhao Ma

School of Data Science

The Chinese University of Hong Kong, Shenzhen

Last Lecture


- **Problem #1:** How the DBMS represents the database in files on disk.
- **Problem #2:** How the DBMS manages its memory and move data back-and-forth from disk.




Database Storage

- **Spatial Control:**
 - Where to write pages on disk.
 - The goal is to keep pages that are used together often as physically close together as possible on disk.
- **Temporal Control:**
 - When to read pages into memory, and when to write them to disk.
 - The goal is to minimize the number of stalls from having to read data from disk.


Disk-oriented DBMS




Disk-oriented DBMS




Disk-oriented DBMS




Disk-oriented DBMS




Disk-oriented DBMS




Disk-oriented DBMS




Disk-oriented DBMS




Disk-oriented DBMS




Disk-oriented DBMS



Disk-oriented DBMS



Disk-oriented DBMS




This Lecture

- Buffer Pool Manager
- Disk I/O Scheduling
- Replacement Policies
- Other Memory Pools

Buffer Pool Manager


Buffer Pool Organization

- Memory region organized as an array of fixed-size pages.
An array entry is called a **frame**.
- When the DBMS requests a page, an exact copy is placed into one of these frames.
- Dirty pages are buffered and not written to disk immediately
 - Write-Back Cache




Buffer Pool Organization

- Memory region organized as an array of fixed-size pages.
An array entry is called a **frame**.
- When the DBMS requests a page, an exact copy is placed into one of these frames.
- Dirty pages are buffered and not written to disk immediately
 - Write-Back Cache




Buffer Pool Organization

- Memory region organized as an array of fixed-size pages.
An array entry is called a **frame**.
- When the DBMS requests a page, an exact copy is placed into one of these frames.
- Dirty pages are buffered and not written to disk immediately
 - Write-Back Cache




Buffer Pool Meta-data

- The **page table** keeps track of pages that are currently in memory.
 - Usually a fixed-size hash table protected with latches to ensure thread-safe access.
- Additional meta-data per page:
 - Dirty Flag
 - Pin/Reference Counter
 - Access Tracking Information




Buffer Pool Meta-data

- The **page table** keeps track of pages that are currently in memory.
 - Usually a fixed-size hash table protected with latches to ensure thread-safe access.
- Additional meta-data per page:
 - Dirty Flag
 - Pin/Reference Counter
 - Access Tracking Information




Buffer Pool Meta-data

- The **page table** keeps track of pages that are currently in memory.
 - Usually a fixed-size hash table protected with latches to ensure thread-safe access.
- Additional meta-data per page:
 - Dirty Flag
 - Pin/Reference Counter
 - Access Tracking Information




Buffer Pool Meta-data

- The **page table** keeps track of pages that are currently in memory.
 - Usually a fixed-size hash table protected with latches to ensure thread-safe access.
- Additional meta-data per page:
 - Dirty Flag
 - Pin/Reference Counter
 - Access Tracking Information




Buffer Pool Meta-data

- The **page table** keeps track of pages that are currently in memory.
 - Usually a fixed-size hash table protected with latches to ensure thread-safe access.
- Additional meta-data per page:
 - Dirty Flag
 - Pin/Reference Counter
 - Access Tracking Information




Buffer Pool Meta-data

- The **page table** keeps track of pages that are currently in memory.
 - Usually a fixed-size hash table protected with latches to ensure thread-safe access.
- Additional meta-data per page:
 - Dirty Flag
 - Pin/Reference Counter
 - Access Tracking Information



Buffer Pool Meta-data

- The **page table** keeps track of pages that are currently in memory.
 - Usually a fixed-size hash table protected with latches to ensure thread-safe access.
- Additional meta-data per page:
 - Dirty Flag
 - Pin/Reference Counter
 - Access Tracking Information



Locks vs. Latches

- **Locks:**

- Protects the database's logical contents from other transactions.
- Held for transaction duration.
- Need to be able to rollback changes.

- **Latches:**

- Protects the critical sections of the DBMS's internal data structure from other threads.
- Held for operation duration.
- Do not need to be able to rollback changes.

Locks vs. Latches

- **Locks:**

- Protects the database's logical contents from other transactions.
- Held for transaction duration.
- Need to be able to rollback changes.

- **Latches:**

- Protects the critical sections of the DBMS's internal data structure from other threads.
- Held for operation duration.
- Do not need to be able to rollback changes.

← Mutex

Page Tables vs. Page Directory

- The **page directory** is the mapping from page ids to page locations in the database files.
 - All changes must be recorded on disk to allow the DBMS to find on restart.
- The **page table** is the mapping from page ids to a copy of the page in buffer pool frames.
 - This is an in-memory data structure that does not need to be stored on disk.

Allocation Policies

- **Global Policies:**
 - Make decisions for all active queries.
- **Local Policies:**
 - Allocate frames to a specific queries without considering the behavior of concurrent queries.
 - Still need to support sharing pages.

Buffer Pool Optimizations

Buffer Pool Optimizations

- Multiple Buffer Pools
- Pre-Fetching
- Scan Sharing
- Buffer Pool Bypass



Multiple Buffer Pools

- The DBMS does not always have a single buffer pool for the entire system.
 - Multiple buffer pool instances
 - Per-database buffer pool
 - Per-page type buffer pool
- Partitioning memory across multiple pools helps reduce latch contention and improve locality.



Multiple Buffer Pools


- **Approach #1: Object Id**
 - Embed an object identifier in record ids and then maintain a mapping from objects to specific buffer pools.
- **Approach #2: Hashing**
 - Hash the page id to select which buffer pool to access.



Multiple Buffer Pools

- **Approach #1: Object Id**
 - Embed an object identifier in record ids and then maintain a mapping from objects to specific buffer pools.
- **Approach #2: Hashing**
 - Hash the page id to select which buffer pool to access.

Q1 GET RECORD #123




Multiple Buffer Pools

- **Approach #1: Object Id**
 - Embed an object identifier in record ids and then maintain a mapping from objects to specific buffer pools.
- **Approach #2: Hashing**
 - Hash the page id to select which buffer pool to access.

Q1 GET RECORD #123


<ObjectId, PageId, SlotNum>



Multiple Buffer Pools


- **Approach #1: Object Id**
 - Embed an object identifier in record ids and then maintain a mapping from objects to specific buffer pools.
- **Approach #2: Hashing**
 - Hash the page id to select which buffer pool to access.

Q1 GET RECORD #123
`<ObjectId, PageId, SlotNum>`



Multiple Buffer Pools


- **Approach #1: Object Id**
 - Embed an object identifier in record ids and then maintain a mapping from objects to specific buffer pools.
- **Approach #2: Hashing**
 - Hash the page id to select which buffer pool to access.



Multiple Buffer Pools

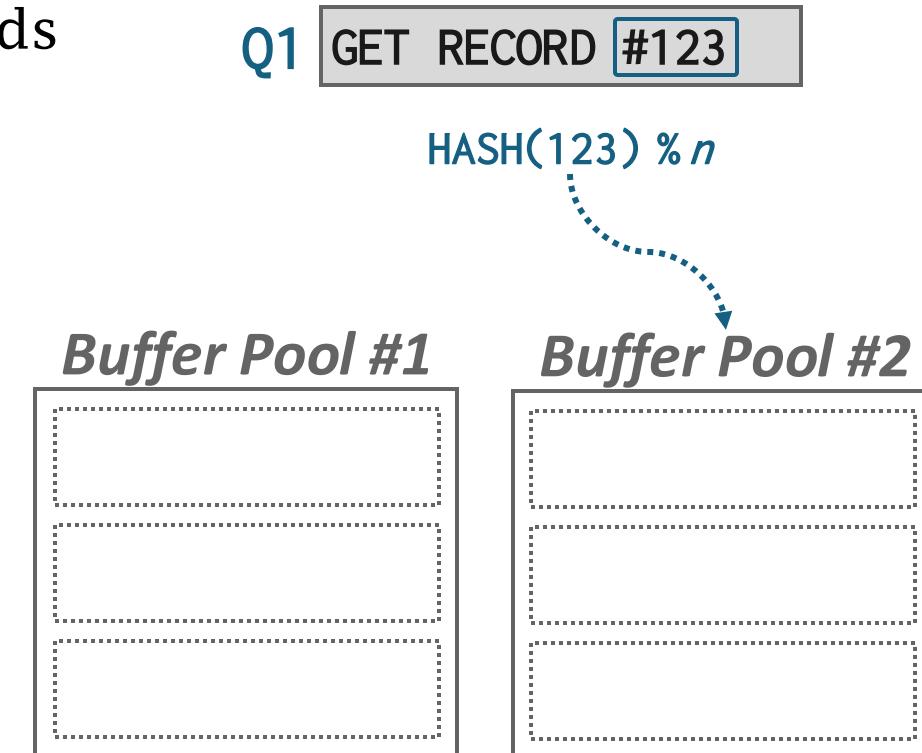
- **Approach #1: Object Id**
 - Embed an object identifier in record ids and then maintain a mapping from objects to specific buffer pools.
- **Approach #2: Hashing**
 - Hash the page id to select which buffer pool to access.

Q1 GET RECORD #123




Multiple Buffer Pools

- **Approach #1: Object Id**
 - Embed an object identifier in record ids and then maintain a mapping from objects to specific buffer pools.
- **Approach #2: Hashing**
 - Hash the page id to select which buffer pool to access.




Pre-fetching

- The DBMS can also prefetch pages based on a query plan.
 - Sequential Scans
 - Index Scans




Pre-fetching

- The DBMS can also prefetch pages based on a query plan.
 - Sequential Scans
 - Index Scans




Pre-fetching

- The DBMS can also prefetch pages based on a query plan.
 - Sequential Scans
 - Index Scans




Pre-fetching

- The DBMS can also prefetch pages based on a query plan.
 - Sequential Scans
 - Index Scans




Pre-fetching

- The DBMS can also prefetch pages based on a query plan.
 - Sequential Scans
 - Index Scans




Pre-fetching

- The DBMS can also prefetch pages based on a query plan.
 - Sequential Scans
 - Index Scans




Pre-fetching

- The DBMS can also prefetch pages based on a query plan.
 - Sequential Scans
 - Index Scans




Pre-fetching

- The DBMS can also prefetch pages based on a query plan.
 - Sequential Scans
 - Index Scans




Pre-fetching

- The DBMS can also prefetch pages based on a query plan.
 - Sequential Scans
 - Index Scans



Pre-fetching



Disk Pages

index-page0
index-page1
index-page2
index-page3
index-page4
index-page5


Buffer Pool



Q1

```
SELECT * FROM A  
WHERE val BETWEEN 100 AND 250
```

Pre-fetching



Disk Pages




Buffer Pool



Q1

```
SELECT * FROM A
WHERE val BETWEEN 100 AND 250
```

Pre-fetching




Buffer Pool





Q1


```
SELECT * FROM A
WHERE val BETWEEN 100 AND 250
```




Pre-fetching


Q1

```
SELECT * FROM A
WHERE val BETWEEN 100 AND 250
```





Pre-fetching



Q1



```
SELECT * FROM A
WHERE val BETWEEN 100 AND 250
```



Pre-fetching

Q1



```
SELECT * FROM A
WHERE val BETWEEN 100 AND 250
```



Pre-fetching

Q1

```
SELECT * FROM A
WHERE val BETWEEN 100 AND 250
```



Scan Sharing

- Queries can reuse data retrieved from storage or operator computations.
 - Also called *synchronized scans*.
 - This is different from result caching.
- Allow multiple queries to attach to a single cursor that scans a table.
 - Queries do not have to be the same.
 - Can also share intermediate results.

Scan Sharing

- If a query wants to scan a table and another query is already doing this, then the DBMS will attach the second query's cursor to the existing cursor.
- Examples:
 - Fully supported in DB2, MSSQL, Teradata, and Postgres.
 - Oracle only supports cursor sharing for identical queries.




ORACLE



PostgreSQL

Scan Sharing

- If a query wants to scan a table and another query is already doing this, then the DBMS will attach the second query's cursor to the existing cursor.
- Examples:
 - Fully supported in DB2, MSSQL, Teradata, and Postgres.
 - Oracle only supports cursor sharing for identical queries.



The screenshot shows a section of Oracle Database documentation. It includes a note about cursor sharing requirements, three examples of SQL statements, a 'Copy' button, and the Oracle logo.

For a textual match to occur, the text of the SQL statements or PL/SQL blocks must be character-for-character identical, including spaces, case, and comments. For example, the following statements cannot use the same shared SQL area:

```
SELECT * FROM employees;
SELECT * FROM Employees;
SELECT *  FROM employees;
```

 Copy


TER  reSQL 

Scan Sharing


Q1

```
SELECT SUM(val) FROM A
```

Buffer Pool



Disk Pages





Scan Sharing

Q1

```
SELECT SUM(val) FROM A
```

Buffer Pool




Scan Sharing

Q1

```
SELECT SUM(val) FROM A
```

Buffer Pool




Scan Sharing


Q1

```
SELECT SUM(val) FROM A
```

Buffer Pool



Disk Pages




Scan Sharing


Q1

```
SELECT SUM(val) FROM A
```

Buffer Pool



Disk Pages




Scan Sharing


Q1

```
SELECT SUM(val) FROM A
```

Buffer Pool



Disk Pages




Scan Sharing


Q1 `SELECT SUM(val) FROM A`

Q2 `SELECT AVG(val) FROM A`

Buffer Pool



Disk Pages





Scan Sharing

Q1 `SELECT SUM(val) FROM A`

Q2 `SELECT AVG(val) FROM A`

Buffer Pool





Scan Sharing

Q1 `SELECT SUM(val) FROM A`

Q2 `SELECT AVG(val) FROM A`

Buffer Pool




Scan Sharing


Q1 `SELECT SUM(val) FROM A`

Q2 `SELECT AVG(val) FROM A`

Buffer Pool



Disk Pages





Scan Sharing

Q1 `SELECT SUM(val) FROM A`

Q2 `SELECT AVG(val) FROM A`

Buffer Pool




Scan Sharing


Q1 `SELECT SUM(val) FROM A`

Q2 `SELECT AVG(val) FROM A`

Buffer Pool



Disk Pages




Scan Sharing

Q1


```
SELECT SUM(val) FROM A
```

Buffer Pool



Q2

Disk Pages




Scan Sharing


Q1

```
SELECT SUM(val) FROM A
```

Buffer Pool




Disk Pages



Continuous Scan Sharing

- Instead of trying to be clever, the DBMS continuously scans the database files repeatedly.
 - One continuous cursor per table.
 - Queries “hop” on board the cursor while it is running and then disconnect once they have enough data.
- Not viable if you pay per IOP.
Only done in academic prototypes.




Continuous Scan Sharing

- Instead of trying to be clever, the DBMS continuously scans the database files repeatedly.
 - One continuous cursor per table.
 - Queries “hop” on board the cursor while it is running and then disconnect once they have enough data.
- Not viable if you pay per IOP.
Only done in academic prototypes.



Continuous Scan Sharing

- Instead of trying to be clever, the DBMS continuously scans the database files repeatedly.
 - One continuous cursor per table.
 - Queries “hop” on board the cursor while it is running and then disconnect once they have enough data.
- Not viable if you pay per IOP.
Only done in academic prototypes.




Continuous Scan Sharing

- Instead of trying to be clever, the DBMS continuously scans the database files repeatedly.
 - One continuous cursor per table.
 - Queries “hop” on board the cursor while it is running and then disconnect once they have enough data.
- Not viable if you pay per IOP.
Only done in academic prototypes.



Continuous Scan Sharing

- Instead of trying to be clever, the DBMS continuously scans the database files repeatedly.
 - One continuous cursor per table.
 - Queries “hop” on board the cursor while it is running and then disconnect once they have enough data.
- Not viable if you pay per IOP.
Only done in academic prototypes.



Buffer Pool Bypass

- The sequential scan operator will not store fetched pages in the buffer pool to avoid overhead.
 - Memory is local to running query.
 - Works well if operator needs to read a large sequence of pages that are contiguous on disk.
 - Can also be used for temporary data (sorting, joins).
- Called “Light Scans” in Informix.





Buffer Replacement Policies

Buffer Replacement Policies

- When the DBMS needs to free up a frame to make room for a new page, it must decide which page to evict from the buffer pool.
- Goals:
 - Correctness
 - Accuracy
 - Speed
 - Metadata overhead



Least-Recently Used

- Maintain a single timestamp of when each page was last accessed. When the DBMS needs to evict a page, select the one with the oldest timestamp.
 - Keep the pages in sorted order to reduce the search time on eviction.





Least-Recently Used

- Maintain a single timestamp of when each page was last accessed. When the DBMS needs to evict a page, select the one with the oldest timestamp.
 - Keep the pages in sorted order to reduce the search time on eviction.





Least-Recently Used

- Maintain a single timestamp of when each page was last accessed. When the DBMS needs to evict a page, select the one with the oldest timestamp.
 - Keep the pages in sorted order to reduce the search time on eviction.




Least-Recently Used

- Maintain a single timestamp of when each page was last accessed. When the DBMS needs to evict a page, select the one with the oldest timestamp.
 - Keep the pages in sorted order to reduce the search time on eviction.



Least-Recently Used

- Maintain a single timestamp of when each page was last accessed. When the DBMS needs to evict a page, select the one with the oldest timestamp.
 - Keep the pages in sorted order to reduce the search time on eviction.



Clock

- Approximation of LRU that does not need a separate timestamp per page.
 - Each page has a reference bit.
 - When a page is accessed, set to 1.
- Organize the pages in a circular buffer with a “clock hand”:
 - Upon sweeping, check if a page’s bit is set to 1.
 - If yes, set to zero. If no, then evict.

Clock

- Approximation of LRU that does not need a separate timestamp per page.
 - Each page has a reference bit.
 - When a page is accessed, set to 1.
- Organize the pages in a circular buffer with a “clock hand”:
 - Upon sweeping, check if a page’s bit is set to 1.
 - If yes, set to zero. If no, then evict.

page1


page4

page2

page3


Clock

- Approximation of LRU that does not need a separate timestamp per page.
 - Each page has a reference bit.
 - When a page is accessed, set to 1.
- Organize the pages in a circular buffer with a “clock hand”:
 - Upon sweeping, check if a page’s bit is set to 1.
 - If yes, set to zero. If no, then evict.




Clock

- Approximation of LRU that does not need a separate timestamp per page.
 - Each page has a reference bit.
 - When a page is accessed, set to 1.
- Organize the pages in a circular buffer with a “clock hand”:
 - Upon sweeping, check if a page’s bit is set to 1.
 - If yes, set to zero. If no, then evict.




Clock

- Approximation of LRU that does not need a separate timestamp per page.
 - Each page has a reference bit.
 - When a page is accessed, set to 1.
- Organize the pages in a circular buffer with a “clock hand”:
 - Upon sweeping, check if a page’s bit is set to 1.
 - If yes, set to zero. If no, then evict.




Clock

- Approximation of LRU that does not need a separate timestamp per page.
 - Each page has a reference bit.
 - When a page is accessed, set to 1.
- Organize the pages in a circular buffer with a “clock hand”:
 - Upon sweeping, check if a page’s bit is set to 1.
 - If yes, set to zero. If no, then evict.




Clock

- Approximation of LRU that does not need a separate timestamp per page.
 - Each page has a reference bit.
 - When a page is accessed, set to 1.
- Organize the pages in a circular buffer with a “clock hand”:
 - Upon sweeping, check if a page’s bit is set to 1.
 - If yes, set to zero. If no, then evict.




Clock

- Approximation of LRU that does not need a separate timestamp per page.
 - Each page has a reference bit.
 - When a page is accessed, set to 1.
- Organize the pages in a circular buffer with a “clock hand”:
 - Upon sweeping, check if a page’s bit is set to 1.
 - If yes, set to zero. If no, then evict.




Clock

- Approximation of LRU that does not need a separate timestamp per page.
 - Each page has a reference bit.
 - When a page is accessed, set to 1.
- Organize the pages in a circular buffer with a “clock hand”:
 - Upon sweeping, check if a page’s bit is set to 1.
 - If yes, set to zero. If no, then evict.




Clock

- Approximation of LRU that does not need a separate timestamp per page.
 - Each page has a reference bit.
 - When a page is accessed, set to 1.
- Organize the pages in a circular buffer with a “clock hand”:
 - Upon sweeping, check if a page’s bit is set to 1.
 - If yes, set to zero. If no, then evict.




Clock

- Approximation of LRU that does not need a separate timestamp per page.
 - Each page has a reference bit.
 - When a page is accessed, set to 1.
- Organize the pages in a circular buffer with a “clock hand”:
 - Upon sweeping, check if a page’s bit is set to 1.
 - If yes, set to zero. If no, then evict.




Clock

- Approximation of LRU that does not need a separate timestamp per page.
 - Each page has a reference bit.
 - When a page is accessed, set to 1.
- Organize the pages in a circular buffer with a “clock hand”:
 - Upon sweeping, check if a page’s bit is set to 1.
 - If yes, set to zero. If no, then evict.




Clock

- Approximation of LRU that does not need a separate timestamp per page.
 - Each page has a reference bit.
 - When a page is accessed, set to 1.
- Organize the pages in a circular buffer with a “clock hand”:
 - Upon sweeping, check if a page’s bit is set to 1.
 - If yes, set to zero. If no, then evict.



Clock

- Approximation of LRU that does not need a separate timestamp per page.
 - Each page has a reference bit.
 - When a page is accessed, set to 1.
- Organize the pages in a circular buffer with a “clock hand”:
 - Upon sweeping, check if a page’s bit is set to 1.
 - If yes, set to zero. If no, then evict.



Observation

- LRU + CLOCK replacement policies are susceptible to sequential flooding.
 - A query performs a sequential scan that reads every page.
 - This pollutes the buffer pool with pages that are read once and then never again.
 - In OLAP workloads, the ***most recently used*** page is often the best page to evict.
- LRU + CLOCK only tracks when a page was last accessed, but not how often a page is accessed.

Sequential Flooding


Q1

```
SELECT * FROM A WHERE primKey = 1
```

Buffer Pool



Disk Pages




Sequential Flooding

Q1

```
SELECT * FROM A WHERE primKey = 1
```

Buffer Pool



Sequential Flooding

Q1


```
SELECT * FROM A WHERE primKey = 1
```

Buffer Pool



Q1

Disk Pages



Sequential Flooding


Q1 `SELECT * FROM A WHERE primKey = 1`

Q2 `SELECT AVG(val) FROM A`

Buffer Pool




Disk Pages



Sequential Flooding

Q1 `SELECT * FROM A WHERE primKey = 1`


Q2 `SELECT AVG(val) FROM A`



Sequential Flooding

Q1 `SELECT * FROM A WHERE primKey = 1`


Q2 `SELECT AVG(val) FROM A`



Sequential Flooding

Q1 `SELECT * FROM A WHERE primKey = 1`

Q2 `SELECT AVG(val) FROM A`





Sequential Flooding

Q1 `SELECT * FROM A WHERE primKey = 1`

Q2 `SELECT AVG(val) FROM A`

Q3 `SELECT * FROM A WHERE primKey = 1`





Sequential Flooding

Q1 `SELECT * FROM A WHERE primKey = 1`

Q2 `SELECT AVG(val) FROM A`

Q3 `SELECT * FROM A WHERE primKey = 1`





Sequential Flooding

Q1 `SELECT * FROM A WHERE primKey = 1`

Q2 `SELECT AVG(val) FROM A`

Q3 `SELECT * FROM A WHERE primKey = 1`

Sequential flooding can occur when a table is scanned multiple times within one query, such as with a Nested-Blocks Join.



Better Policies: LRU- k

- Track the last k references to each page and compute the interval between subsequent accesses.
 - Can get fancy with distinguishing between reference types.
- The DBMS then uses this history to estimate the next time that page is going to be accessed.
 - Replace the page with the oldest “ k -th” access.
 - A balance between recency and frequency of access.
 - Maintain an ephemeral in-memory cache for recently evicted pages to prevent them from always being evicted.

The LRU-K Page Replacement Algorithm For Database Disk Buffering

Elizabeth J. O’Neil¹, Patrick E. O’Neil¹, Gerhard Weikum²
¹ Department of Mathematics and Computer Science
 University of Massachusetts at Boston
 Harbor Mtn. Campus
 Boston, MA 02125-3393
 E-mail: conell@cs.umb.edu, pouch@cs.umb.edu, weikum@inf.ethz.ch

² Department of Computer Science
 ETH Zurich
 CH-8092 Zurich
 Switzerland

ABSTRACT

This paper introduces a new approach to database disk buffering, called the LRU-K method. The basic idea of LRU-K is to keep track of the times of the last K references to popular database pages, using this information to statistically predict when pages will be accessed again, on a page basis. Although the LRU-K approach performs optimal statistical inference under relatively standard assumptions, it is far simpler than existing approaches. As we demonstrate with simulation experiments, the LRU-K algorithm surpasses conventional buffering algorithms in discriminating between frequently and infrequently referenced pages. In fact, LRU-K exhibits the same behavior of buffering algorithms in which page sets with known access frequencies are manually analyzed. Unlike customized buffering algorithms however, the LRU-K method is self-tuning, and does not rely on external hints about workload characteristics. Furthermore, the LRU-K algorithm adapts in real time to changing patterns of access.

1. Introduction

1.1 Problem Statement
 All database systems retain disk pages in memory buffers for a period of time after they have been read in from disk and accessed by a particular application. The purpose is to keep popular pages memory resident and reduce disk I/O. In “Five Minutes to Cache and Parallel Page Buffering”, we argue that more memory buffers up to a certain point, in order to reduce the cost of disk arms for a system (GCRAYPUT, see also (CKS)), is needed for a page to be read in from disk, and all current buffers are in use. What current pages should be replaced when a new page is read in depends on the replacement policy, and the different buffering algorithms take their names from the type of replacement policy they impose (see, for example, (BMR)).

SIGMOD ’93 (Washington, DC, USA).
 © 1993 ACM 0-89791-392-5/93/0005/0297...\$1.50.
 297

The algorithm utilized by almost all commercial systems is known as LRU, for Least Recently Used. When a new buffer is needed, the LRU policy drops the page from buffer that has not been accessed for the longest time. LRU has been used mainly for memory management and instruction logic (for example, [DENENING], [COFFDENN], and does not always fit well into the database context, and [CHOUDEW]). In fact, the LRU buffering algorithm has a problem which is addressed by the current paper: that it defines what page to drop from buffer based on too little information, namely the number of times a page is accessed. Specifically, LRU is unable to differentiate between pages that have relatively frequent references and pages that have very infrequent references until the system has wasted a lot of resources keeping infrequently referenced pages in buffer for an extended period.

Example 1.1. Consider a multi-user database application,



through a clustered B-tree indexed key, CUST-ID, to retrieve a customer record (T1). It is statistically known that 20,000 customers exist, that a customer record is 2000 bytes in length, and that space needed for the B-tree index is negligible. We assume that each record is 100 bytes and each key entry. Then if disk pages contain 4000 bytes of usable space, to be packed full, we require 100 pages to hold the leaf level of the B-tree (there is one single B-tree root node), and 10,000 pages to hold the records.

The pattern of reference to these pages (ignoring the B-tree root node) is clearly: 1, 1, 1, 1, 1, 1, ..., alternately referencing a leaf level record and a root record. If we can only afford to buffer 101 pages in memory for this application, the B-tree root node is eliminated; we should buffer all leaf level records. We can easily determine, with a probability of .005 (once in each 20,000 general page references), while it is clearly wasteful to displaced one of the leaf level records, that the leaf pages have only .00005 probability of reference (once in each 20,000 general page references). Using the LRU algorithm, however, we will evict the 101st page, which is one of the most recently referenced ones. To a first approximation, this means 50 B-tree leaf pages and 50 record pages. Given that a page gets no credit for being referenced twice in the recent past and that it is more likely to happen with B-tree leaf pages, there will even be slightly more data





MySQL Approximate LRU- k

- Single LRU linked list but with two entry points (“old” vs “young”).
- New pages are always inserted to the head of the old list.
- If pages in the old list is accessed again, then insert into the head of the young list.





MySQL Approximate LRU- k

- Single LRU linked list but with two entry points (“old” vs “young”).
- New pages are always inserted to the head of the old list.
- If pages in the old list is accessed again, then insert into the head of the young list.





MySQL Approximate LRU- k

- Single LRU linked list but with two entry points (“old” vs “young”).
- New pages are always inserted to the head of the old list.
- If pages in the old list is accessed again, then insert into the head of the young list.





MySQL Approximate LRU- k

- Single LRU linked list but with two entry points (“old” vs “young”).
- New pages are always inserted to the head of the old list.
- If pages in the old list is accessed again, then insert into the head of the young list.





MySQL Approximate LRU- k

- Single LRU linked list but with two entry points (“old” vs “young”).
- New pages are always inserted to the head of the old list.
- If pages in the old list is accessed again, then insert into the head of the young list.





MySQL Approximate LRU- k

- Single LRU linked list but with two entry points (“old” vs “young”).
- New pages are always inserted to the head of the old list.
- If pages in the old list is accessed again, then insert into the head of the young list.





MySQL Approximate LRU- k

- Single LRU linked list but with two entry points (“old” vs “young”).
- New pages are always inserted to the head of the old list.
- If pages in the old list is accessed again, then insert into the head of the young list.





MySQL Approximate LRU- k

- Single LRU linked list but with two entry points (“old” vs “young”).
- New pages are always inserted to the head of the old list.
- If pages in the old list is accessed again, then insert into the head of the young list.



MySQL Approximate LRU- k

- Single LRU linked list but with two entry points (“old” vs “young”).
- New pages are always inserted to the head of the old list.
- If pages in the old list is accessed again, then insert into the head of the young list.



Better Policies: Localization


- The DBMS chooses which pages to evict on a per query basis. This minimizes the pollution of the buffer pool from each query.
 - Keep track of the pages that a query has accessed.
- Example: Postgres maintains a small ring buffer that is private to the query.

Better Policies: Priority Hints

- The DBMS knows about the context of each page during query execution.
- It can provide hints to the buffer pool on whether a page is important or not.

Better Policies: Priority Hints


- The DBMS knows about the context of each page during query execution.
- It can provide hints to the buffer pool on whether a page is important or not.



Better Policies: Priority Hints

- The DBMS knows about the context of each page during query execution.
- It can provide hints to the buffer pool on whether a page is important or not.


Q1 `INSERT INTO A VALUES (id++)`



Better Policies: Priority Hints

- The DBMS knows about the context of each page during query execution.
- It can provide hints to the buffer pool on whether a page is important or not.


Q1 **INSERT INTO A VALUES (*id++*)**



Better Policies: Priority Hints

- The DBMS knows about the context of each page during query execution.
- It can provide hints to the buffer pool on whether a page is important or not.


Q1 `INSERT INTO A VALUES (id++)`



Better Policies: Priority Hints

- The DBMS knows about the context of each page during query execution.
- It can provide hints to the buffer pool on whether a page is important or not.

Q1 **INSERT INTO A VALUES (*id++*)**




Better Policies: Priority Hints

- The DBMS knows about the context of each page during query execution.
- It can provide hints to the buffer pool on whether a page is important or not.

Q1 `INSERT INTO A VALUES (id++)`

Q2 `SELECT * FROM A WHERE id = ?`




Better Policies: Priority Hints

- The DBMS knows about the context of each page during query execution.
- It can provide hints to the buffer pool on whether a page is important or not.

Q1 `INSERT INTO A VALUES (id++)`

Q2 `SELECT * FROM A WHERE id = ?`



Dirty Pages

- **Fast Path:** If a page in the buffer pool is not dirty, then the DBMS can simply “drop” it.
- **Slow Path:** If a page is dirty, then the DBMS must write back to disk to ensure that its changes are persisted.
- Trade-off between fast evictions versus dirty writing pages that will not be read again in the future.

Background Writing

- The DBMS can periodically walk through the page table and write dirty pages to disk.
- When a dirty page is safely written, the DBMS can either evict the page or just unset the dirty flag.
- Need to be careful that the system doesn't write dirty pages before their log records are written...

Disk I/O Scheduling

Observation


- OS/hardware tries to maximize disk bandwidth by reordering and batching I/O requests.
- But they do not know which I/O requests are more important than others.
- Many DBMSs tell you to switch Linux to use the deadline or noop (FIFO) scheduler.
 - Example: [Oracle](#), [Vertica](#), [MySQL](#)

Disk I/O Scheduling

- The DBMS maintain internal queue(s) to track page read/write requests from the entire system.
- Compute priorities based on several factors:
 - Sequential vs. Random I/O
 - Critical Path Task vs. Background Task
 - Table vs. Index vs. Log vs. Ephemeral Data
 - Transaction Information
 - User-based SLAs
- The OS doesn't know these things and is going to get into the way...


OS Page Cache

- Most disk operations go through the OS API. Unless the DBMS tells it not to, the OS maintains its own filesystem cache (aka page cache, buffer cache).
- Most DBMSs use direct I/O (`O_DIRECT`) to bypass the OS's cache.
 - Redundant copies of pages.
 - Different eviction policies.
 - Loss of control over file I/O.




OS Page Cache

- Most disk operations go through the OS API. Unless the DBMS tells it not to, the OS maintains its own filesystem cache (aka page cache, buffer cache).
- Most DBMSs use direct I/O (`O_DIRECT`) to bypass the OS's cache.
 - Redundant copies of pages.
 - Different eviction policies.
 - Loss of control over file I/O.




OS Page Cache

- Most disk operations go through the OS API. Unless the DBMS tells it not to, the OS maintains its own filesystem cache (aka page cache, buffer cache).
- Most DBMSs use direct I/O (`O_DIRECT`) to bypass the OS's cache.
 - Redundant copies of pages.
 - Different eviction policies.
 - Loss of control over file I/O.



OS Page Cache

- Most disk operations go through the OS API. Unless the DBMS tells it not to, the OS maintains its own filesystem cache (aka page cache, buffer cache).
- Most DBMSs use direct I/O (`O_DIRECT`) to bypass the OS's cache.
 - Redundant copies of pages.
 - Different eviction policies.
 - Loss of control over file I/O.



FSYNC Problems

- If the DBMS calls `fwrite`, what happens?
- If the DBMS calls `fsync`, what happens?
- If `fsync` fails (EIO), what happens?
 - Linux marks the dirty pages as clean.
 - If the DBMS calls `fsync` again, then Linux tells you that the flush was successful.

FSYNC Problems

- If the DBMS calls `fwrite`, what happens?
- If the DBMS calls `fsync`, what happens?
- If `fsync` fails (EIO), what happens?
 - Linux marks the dirty pages as clean.
 - If the DBMS calls `fsync` again, then Linux tells you that the flush was successful.

[page](#) [discussion](#) [view source](#) [history](#) [log in](#)

Fsync Errors

This article covers the current status, history, and OS and OS version differences relating to the circa 2018 `fsync()` reliability issues discussed on the PostgreSQL mailing list and elsewhere. It has sometimes been referred to as "fsyncgate 2018".

Contents [hide]	
1	Current status
2	Articles and news
3	Research notes and OS differences
3.1	Open source kernels
3.2	Closed source kernels
3.3	Special cases
3.4	History and notes

Current status

As of this PostgreSQL 12 commit, PostgreSQL will now PANIC on `fsync()` failure. It was backpatched to PostgreSQL 11, 10, 9.6, 9.5 and 9.4. Thanks to Thomas Munro, Andres Freund, Robert Haas, and Craig Ringer. Linux kernel 4.13 improved `fsync()` error handling and the man page for `fsync()` is somewhat improved as well. See:

- Kernelnewbies for 4.13
- Particularly significant 4.13 commits include:
 - "fs: new infrastructure for writeback error handling and reporting"
 - "ext4: use erseq_t based error handling for reporting data writeback errors"
 - "Documentation: flesh out the section in vfs.txt on storing and reporting writeback errors"
 - "mm: set both AS_EIO/AS_ENOSPC and erseq_t in mapping_set_error"

Many thanks to Jeff Layton for work done in this area. Similar changes were made in InnoDB/MySQL, WiredTiger/MongoDB and no doubt other software as a result of the PR around this.

A proposed follow-up change to PostgreSQL was discussed in the thread Refactoring the checkpointer's sync request queue. The patch that was committed did not incorporate the file-descriptor passing changes proposed. There is still discussion open or some additional safeguards that may use file system error counters and/or filesystem-wide flushing.

Articles and news

- The "fsyncgate 2018" mailing list thread
- LVN.net article "PostgreSQL's fsync() surprise"
- LVN.net article "Improved block-layer error handling"

Other Memory Pools

- The DBMS needs memory for things other than just tuples and indexes.
- These other memory pools may not always be backed by disk.
Depends on implementation.
 - Sorting + Join Buffers
 - Query Caches
 - Maintenance Buffers
 - Log Buffers
 - Dictionary Caches

Conclusion

- The DBMS can almost always manage memory better than the OS.
- Leverage the semantics about the query plan to make better decisions:
 - Evictions
 - Allocations
 - Pre-fetching
- Next: Hash Table