香港中文大學(深圳)
The Chinese University of Hong Kong, Shenzhen

SCHOOL OF
DATA SCIENCE
數據科學學院

Ack: Prof. Jignesh Patel @ CMU
Prof. Andy Pavlo @CMU

CSC3170
7: Hash Tables

Chenhao Ma
School of Data Science
The Chinese University of Hong Kong, Shenzhen

Course Status

- We are now going to talk about how to support the DBMS's execution engine to read/write data from pages.

- Two types of data structures:
  - Hash Tables (Unordered)
  - Trees (Ordered)

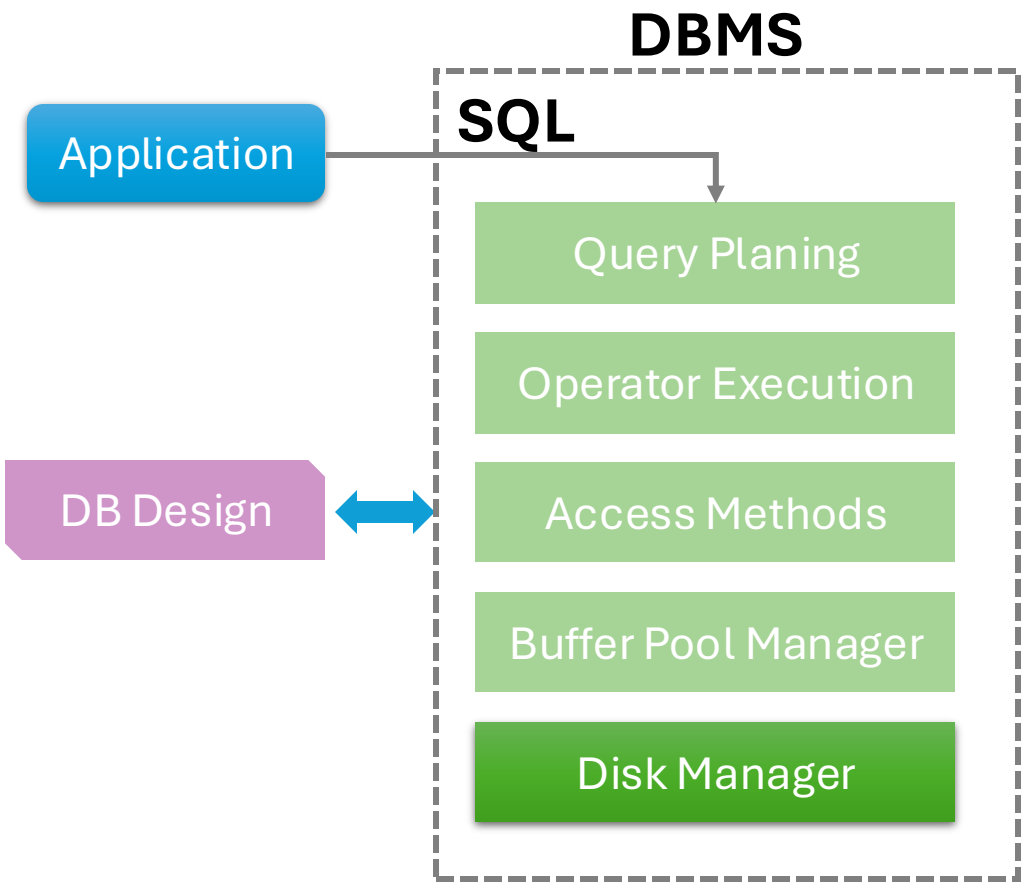

Course Status

- We are now going to talk about how to support the DBMS's execution engine to read/write data from pages.

- Two types of data structures:
  - Hash Tables (Unordered)
  - Trees (Ordered)

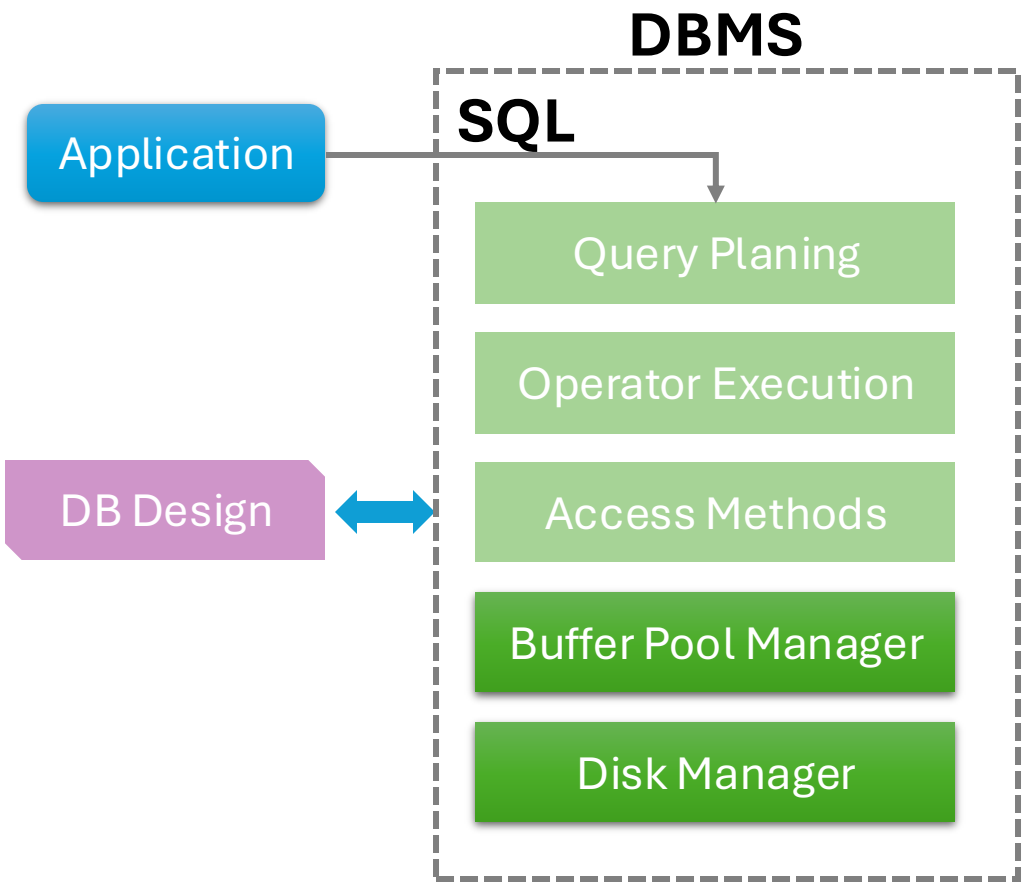

Data Structures

- Internal Meta-data
- Core Data Storage
- Temporary Data Structures
- Table Indexes

Design Decisions

- **Data Organization**
  - How we layout data structure in memory/pages and what information to store to support efficient access.
- **Concurrency**
  - How to enable multiple threads to access the data structure at the same time without causing problems.

Hash Tables

- A **hash table** implements an unordered associative array that maps keys to values.
- It uses a **hash function** to compute an offset into this array for a given key, from which the desired value can be found.
- Space Complexity: $O(n)$
  Time Complexity:
  - Average: $O(1)$
  - Worst: $O(n)$

Hash Tables

- A **hash table** implements an unordered associative array that maps keys to values.
- It uses a **hash function** to compute an offset into this array for a given key, from which the desired value can be found.
- Space Complexity: $O(n)$
  Time Complexity:
  - Average: $O(1)$
  - Worst: $O(n)$

⬅ *Databases care about constants!*

Static Hash Table

- Allocate a giant array that has one slot for every element you need to store.
- To find an entry, mod the key by the number of elements to find the offset in the array.

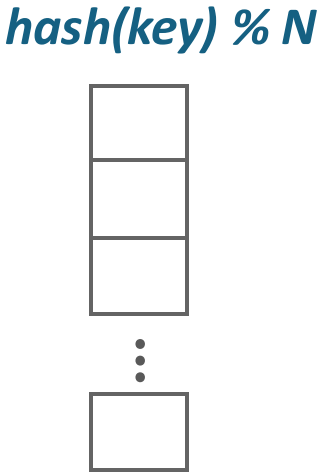

Static Hash Table

- Allocate a giant array that has one slot for every element you need to store.
- To find an entry, mod the key by the number of elements to find the offset in the array.

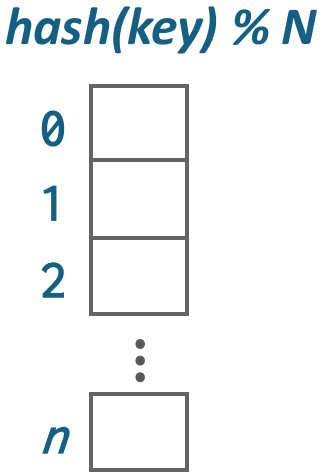

Static Hash Table

- Allocate a giant array that has one slot for every element you need to store.
- To find an entry, mod the key by the number of elements to find the offset in the array.

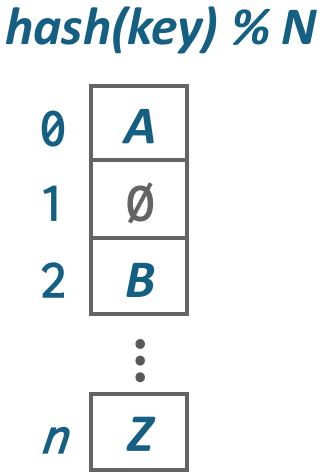

Static Hash Table

- Allocate a giant array that has one slot for every element you need to store.
- To find an entry, mod the key by the number of elements to find the offset in the array.

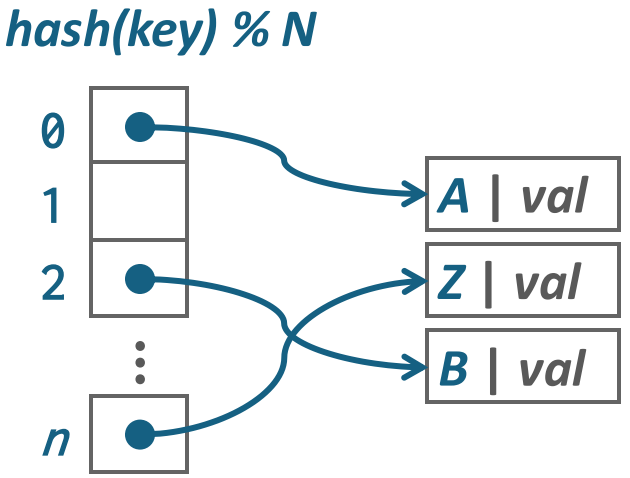

Unrealistic Assumptions

- **Assumption #1:** Number of elements is known ahead of time and fixed.
- **Assumption #2:** Each key is unique.
- **Assumption #3:** Perfect hash function guarantees no collisions.
  - If key1≠key2, then hash(key1)≠hash(key2)

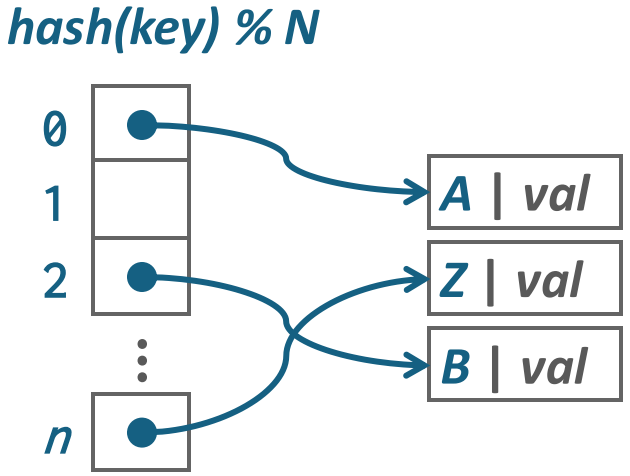

Hash Table

- **Design Decision #1: Hash Function**
  - How to map a large key space into a smaller domain.
  - Trade-off between being fast vs. collision rate.

- **Design Decision #2: Hashing Scheme**
  - How to handle key collisions after hashing.
  - Trade-off between allocating a large hash table vs. additional instructions to get/put keys.

This Lecture

- Hash Functions
- Static Hashing Schemes
- Dynamic Hashing Schemes

Hash Functions

Hash Functions

- For any input key, return an integer representation of that key.
- We do not want to use a cryptographic hash function for DBMS hash tables (e.g., SHA-2).
- We want something that is fast and has a low collision rate.

Hash Functions

- **CRC-64 (1975)**
  - Used in networking for error detection.
- **MurmurHash (2008)**
  - Designed as a fast, general-purpose hash function.
- **Google CityHash (2011)**
  - Designed to be faster for short keys (<64 bytes).
- **Facebook XXHash (2012)**
  - From the creator of zstd compression.
- **Google FarmHash (2014)**
  - Newer version of CityHash with better collision rates.

Hash Functions

- **CRC-64 (1975)**
  - Used in networking for error detection.
- **MurmurHash (2008)**
  - Designed as a fast, general-purpose hash function.
- **Google CityHash (2011)**
  - Designed to be faster for short keys (<64 bytes).
- **Facebook XXHash (2012)**
  - From the creator of zstd compression.

← State-of-the-art

- **Google FarmHash (2014)**
  - Newer version of CityHash with better collision rates.

Hash Functions

- **CRC-64 (1975)**
  - Used in networking for error detection.
- **MurmurHash (2008)**
  - Designed as a fast, general-purpose hash function.
- **Google CityHash (2011)**
  - Designed to be faster for short keys (<64 bytes).
- **Facebook XXHash (2012)**
  - From the creator of zstd compression.

← State-of-the-art

- **Google FarmHash (2014)**
  - Newer version of CityHash with better collision rates.

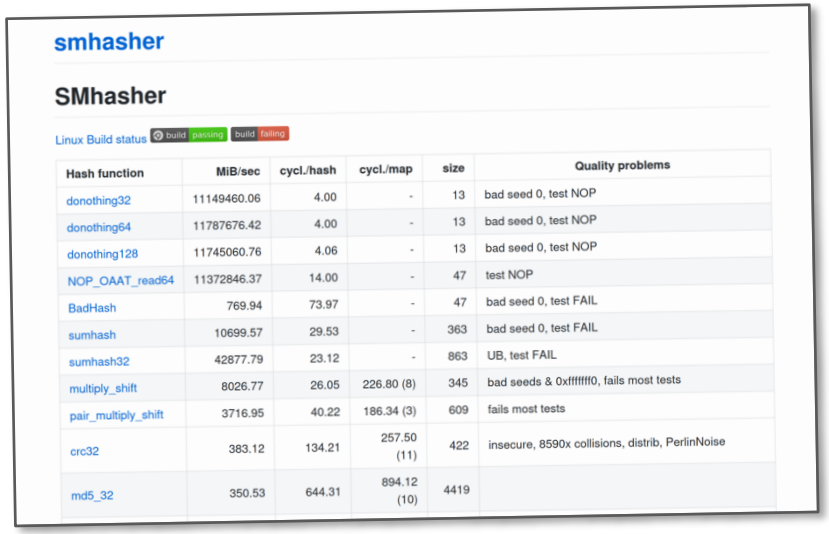

smhasher

SMhasher

Linux Build status build passing build failing

Hash function	MiB/sec	cycl./hash	cycl./map	size	Quality problems
donothing32	11149460.06	4.00	-	13	bad seed 0, test NOP
donothing64	11787676.42	4.00	-	13	bad seed 0, test NOP
donothing128	11745060.76	4.06	-	13	bad seed 0, test NOP
NOP_OAAT_read64	11372846.37	14.00	-	47	test NOP
BadHash	769.94	73.97	-	47	bad seed 0, test FAIL
sumhash	10699.57	29.53	-	363	bad seed 0, test FAIL
sumhash32	42877.79	23.12	-	863	UB, test FAIL
multiply_shift	8026.77	26.05	226.80 (8)	345	bad seeds & 0xffffffff0, fails most tests
pair_multiply_shift	3716.95	40.22	186.34 (3)	609	fails most tests
crc32	383.12	134.21	257.50 (11)	422	insecure, 8590x collisions, distrib, PerlinNoise
md5_32	350.53	644.31	894.12 (10)	4419	

Hash Functions

- **CRC-64 (1975)**
  - Used in networking for error detection.
- **MurmurHash (2008)**
  - Designed as a fast, general-purpose hash function.
- **Google CityHash (2011)**
  - Designed to be faster for short keys (<64 bytes).
- **Facebook XXHash (2012)**
  - From the creator of zstd compression.
- **Google FarmHash (2014)**
  - Newer version of CityHash with better collision rates.

← State-of-th

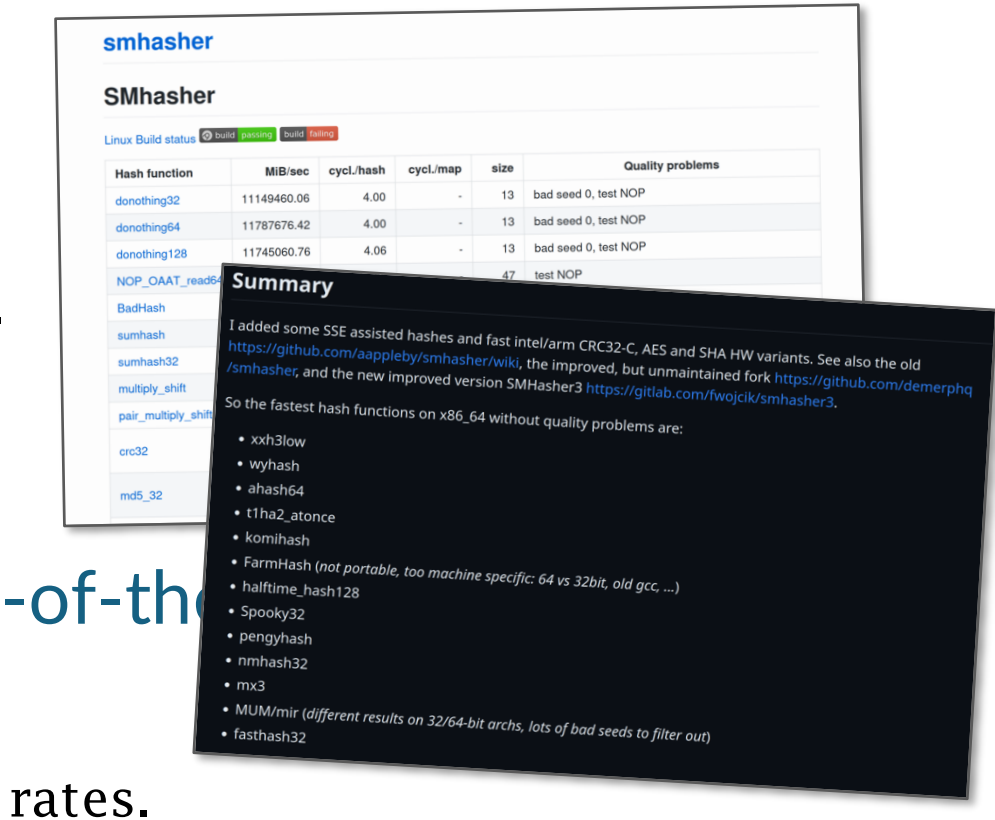

Static Hashing Schemes

- **Approach #1: Linear Probe Hashing**
- **Approach #2: Cuckoo Hashing**
- There are several other schemes (not covered in this course):
  - Robin Hood Hashing
  - Hopscotch Hashing
  - Swiss Tables

Static Hashing Schemes

- **Approach #1: Linear Probe Hashing**
- **Approach #2: Cuckoo Hashing**

← Both are members of a broader class called "**Open Addressing**".

- There are several other schemes (not covered in this course):
  - Robin Hood Hashing
  - Hopscotch Hashing
  - Swiss Tables

Linear Probe Hashing

Static Hashing Schemes

Linear Probe Hashing

- Single giant table of slots.
- Resolve collisions by linearly searching for the next free slot in the table.
  - To determine whether an element is present, hash to a location in the index and scan for it.
  - Must store the key in the index to know when to stop scanning.
  - Insertions and deletions are generalizations of lookups.

- **Example**: Google's `absl::flat_hash_map`

Linear Probe Hashing

hash(key) % N

A
B
C
D
E
F

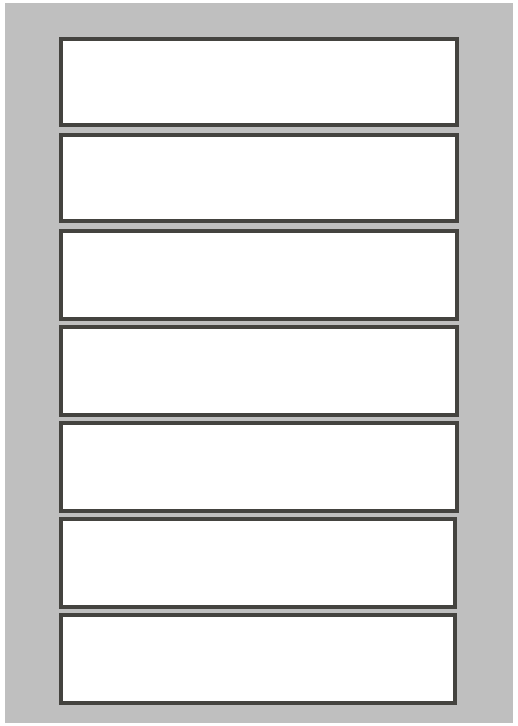

Linear Probe Hashing

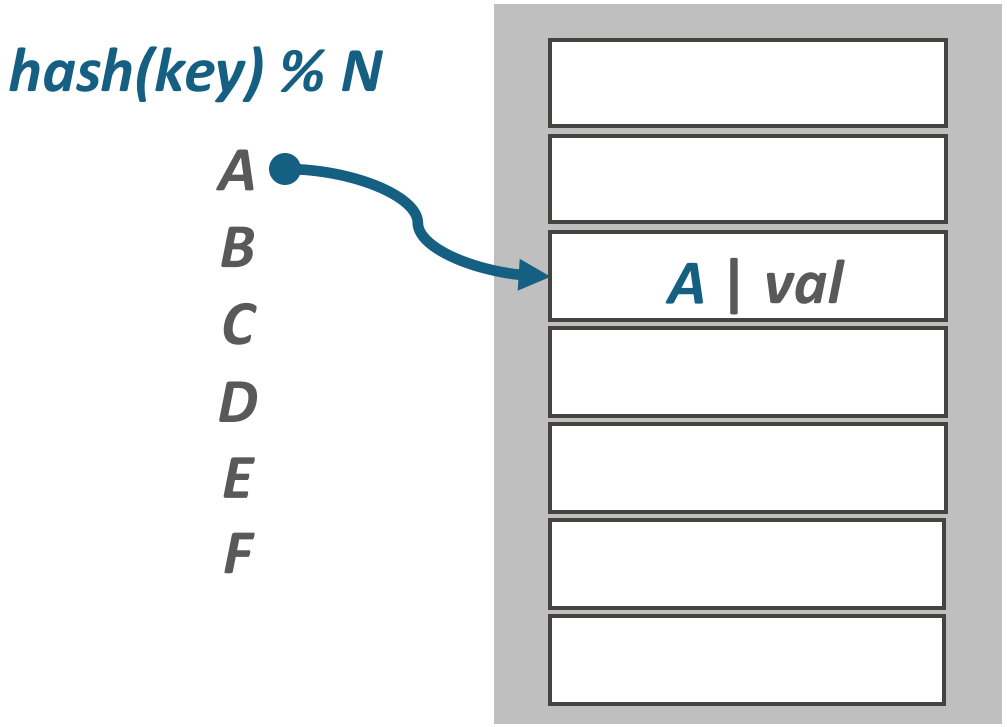

Linear Probe Hashing

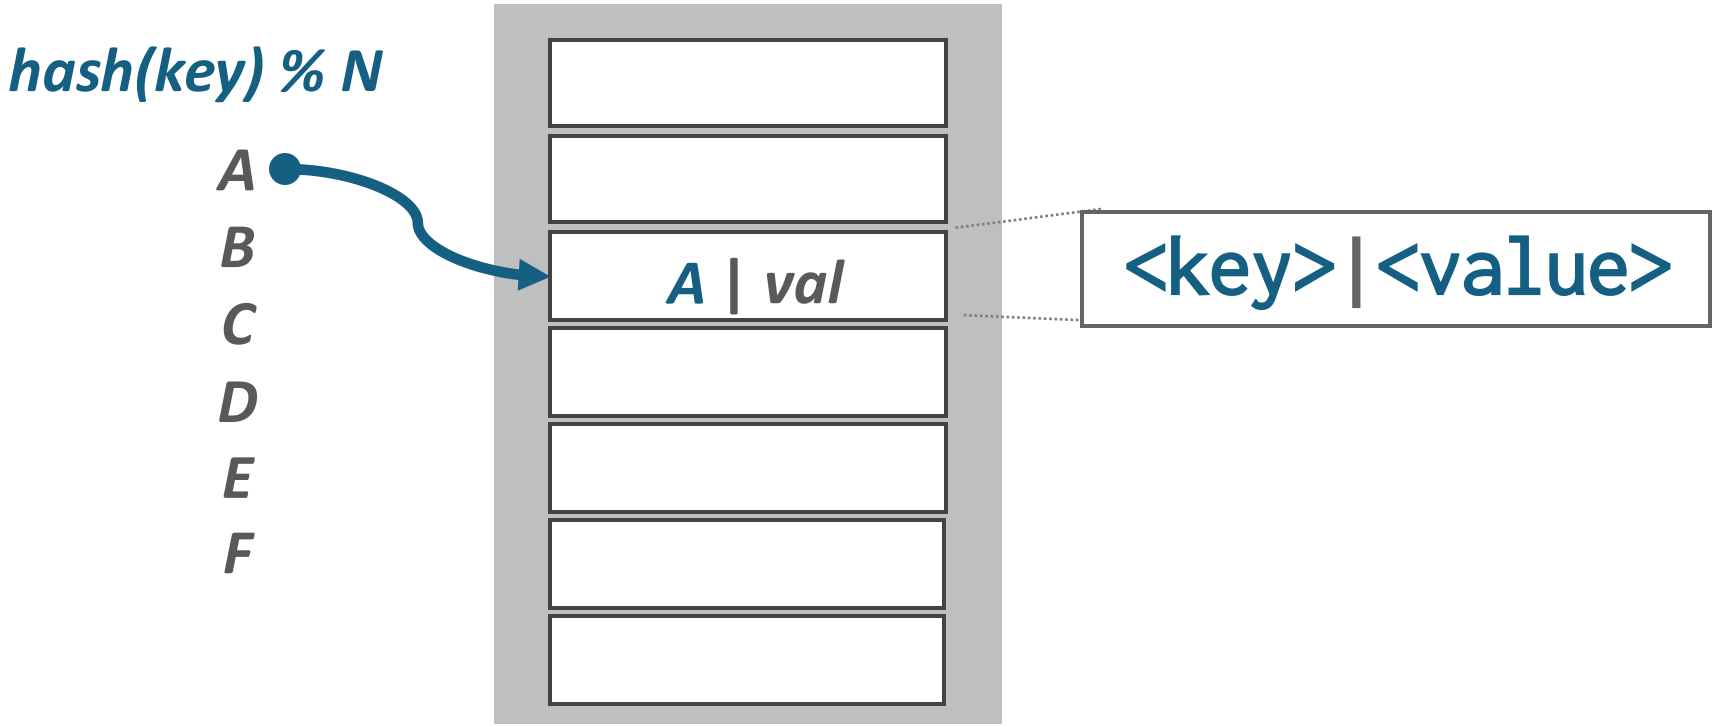

Linear Probe Hashing

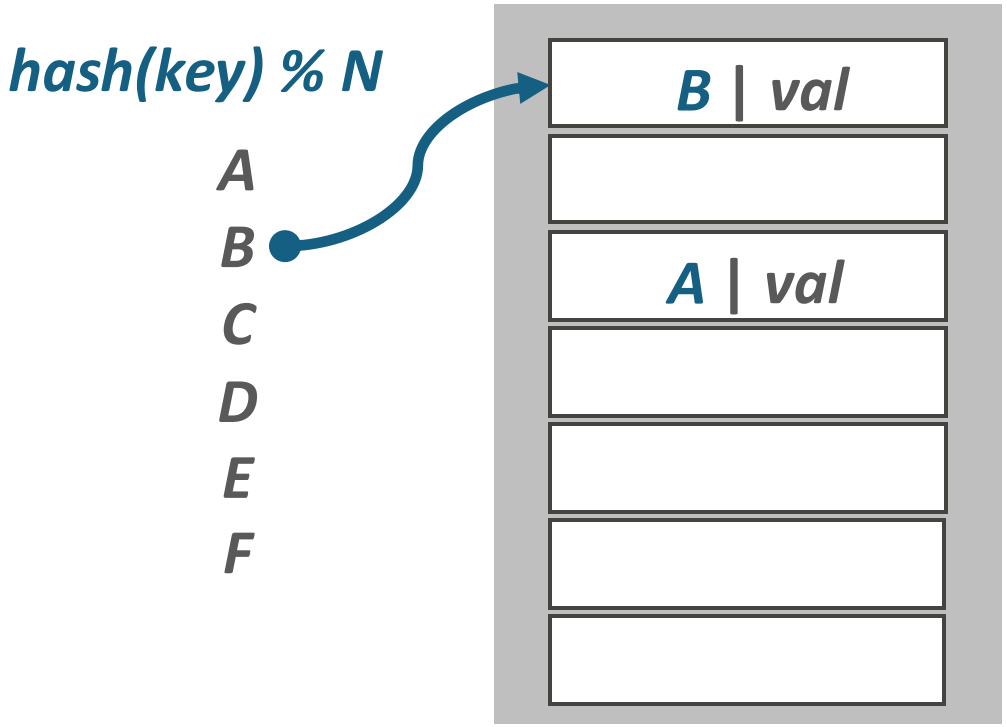

Linear Probe Hashing

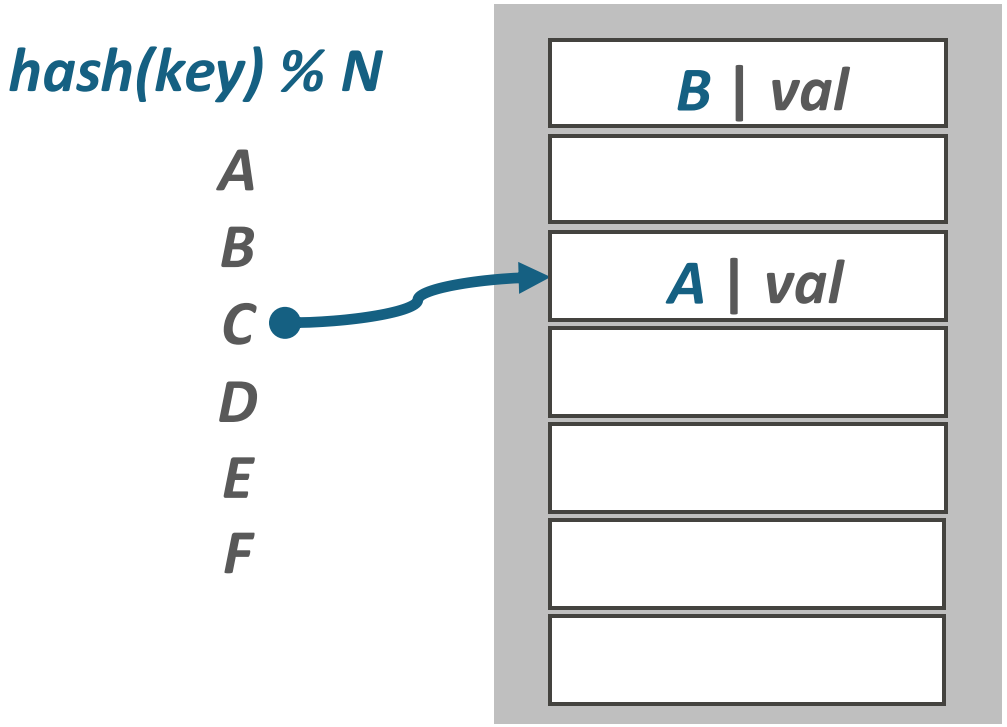

Linear Probe Hashing

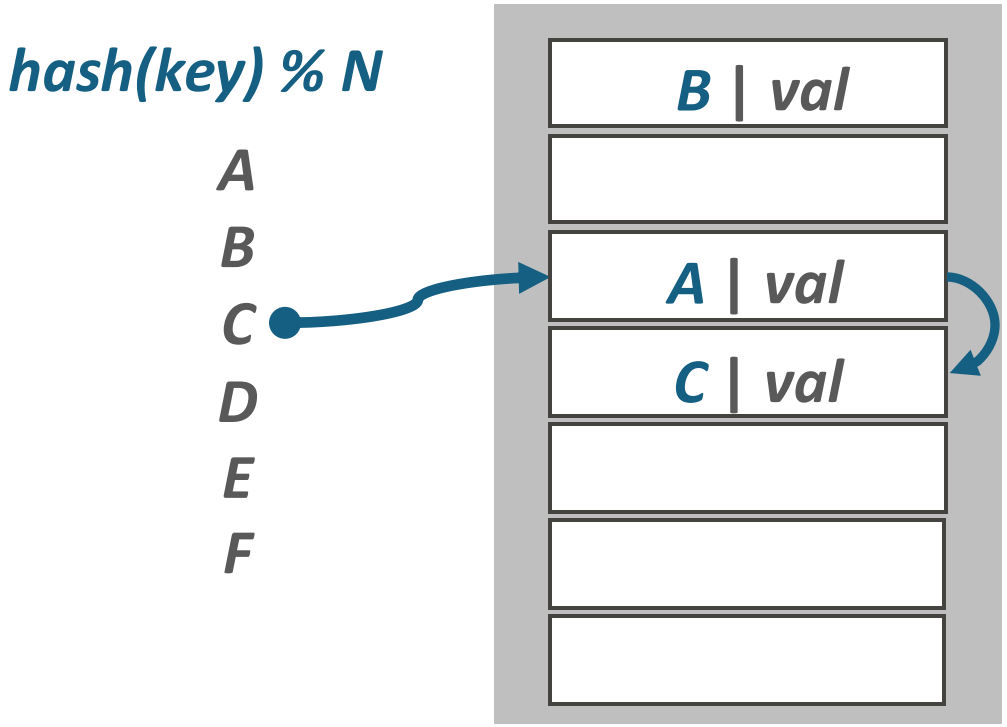

Linear Probe Hashing

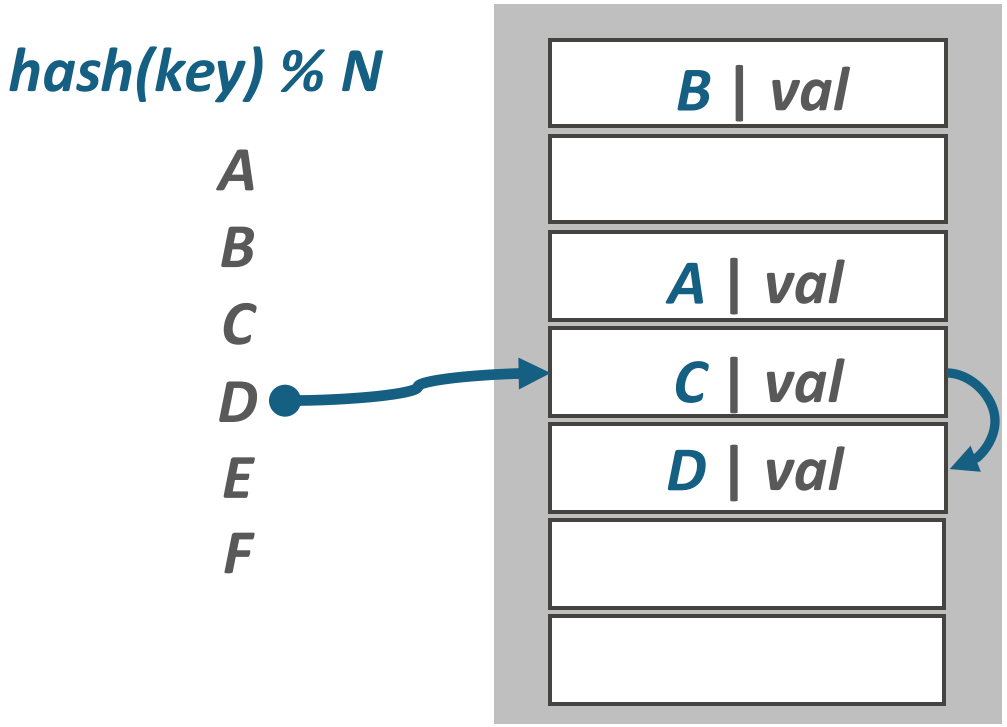

Linear Probe Hashing

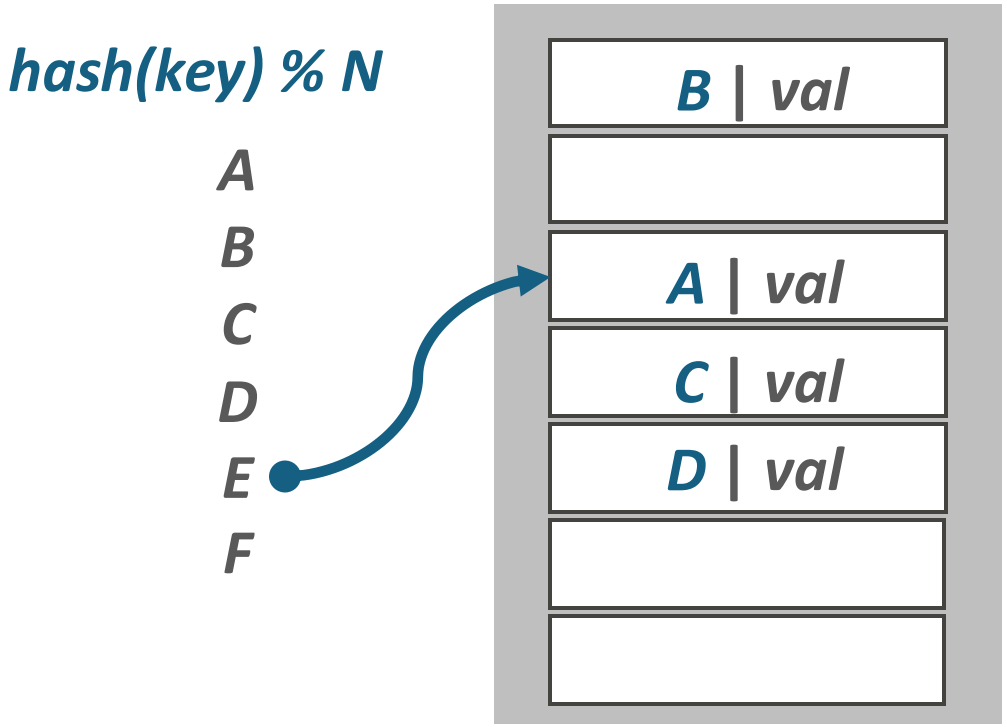

Linear Probe Hashing

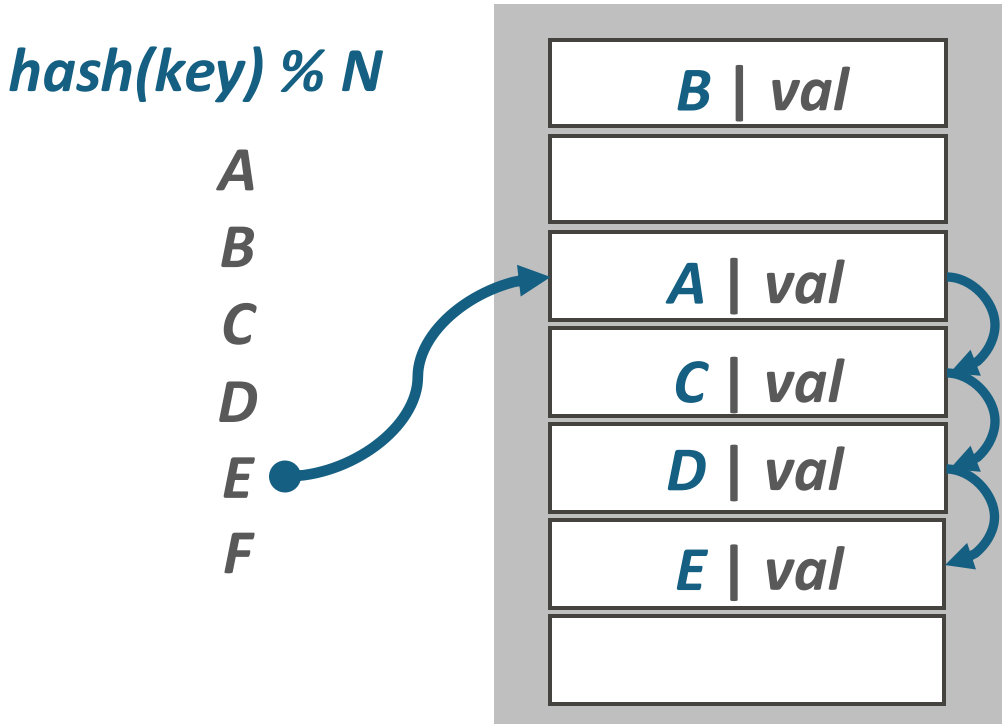

Linear Probe Hashing

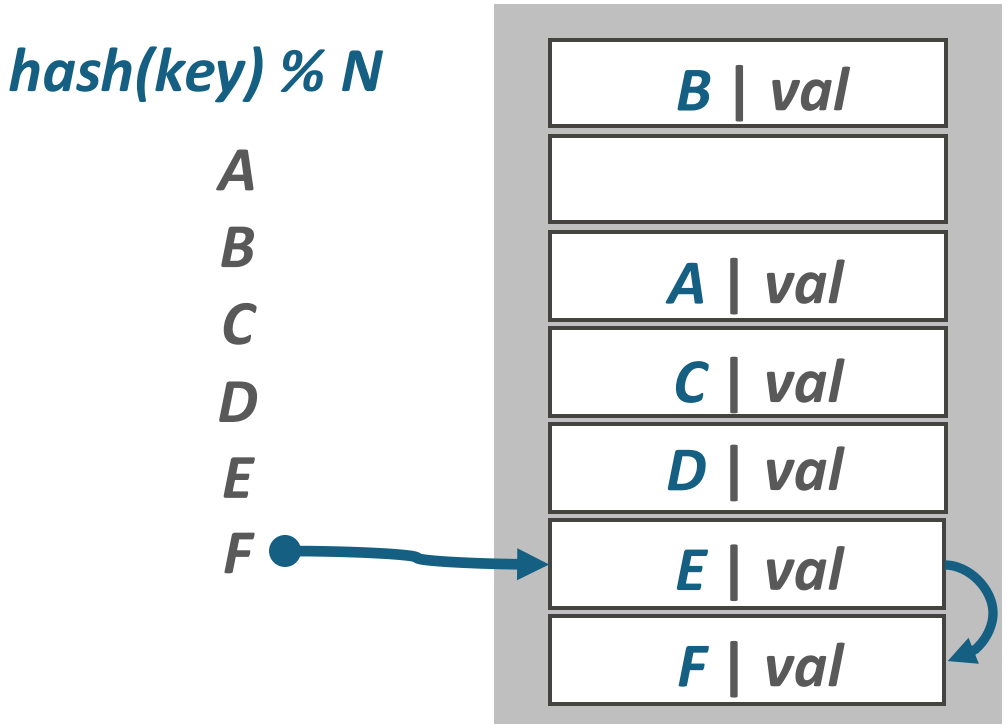

SCHOOL OF DATA SCIENCE

Linear Probe Hashing - Deletes

hash(key) % N	
A	*B | val*
B	
C	*A | val*
D	*C | val*
E	*D | val*
F	*E | val*
	F | val

SCHOOL OF DATA SCIENCE

Linear Probe Hashing - Deletes

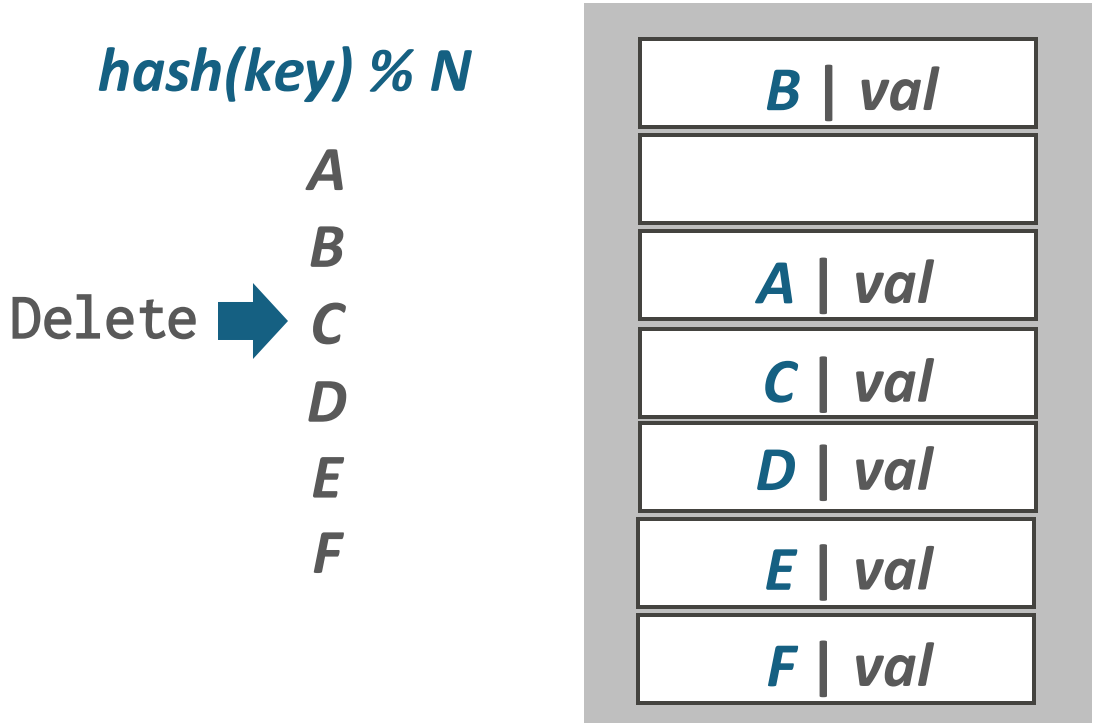

Linear Probe Hashing - Deletes

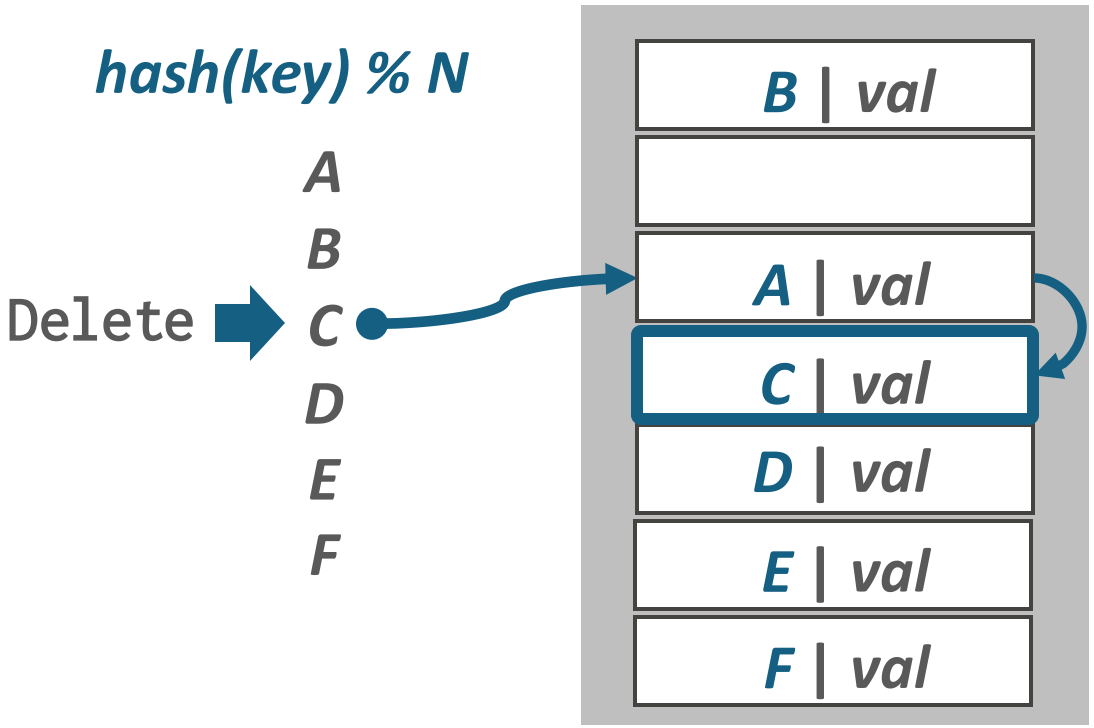

Linear Probe Hashing - Deletes

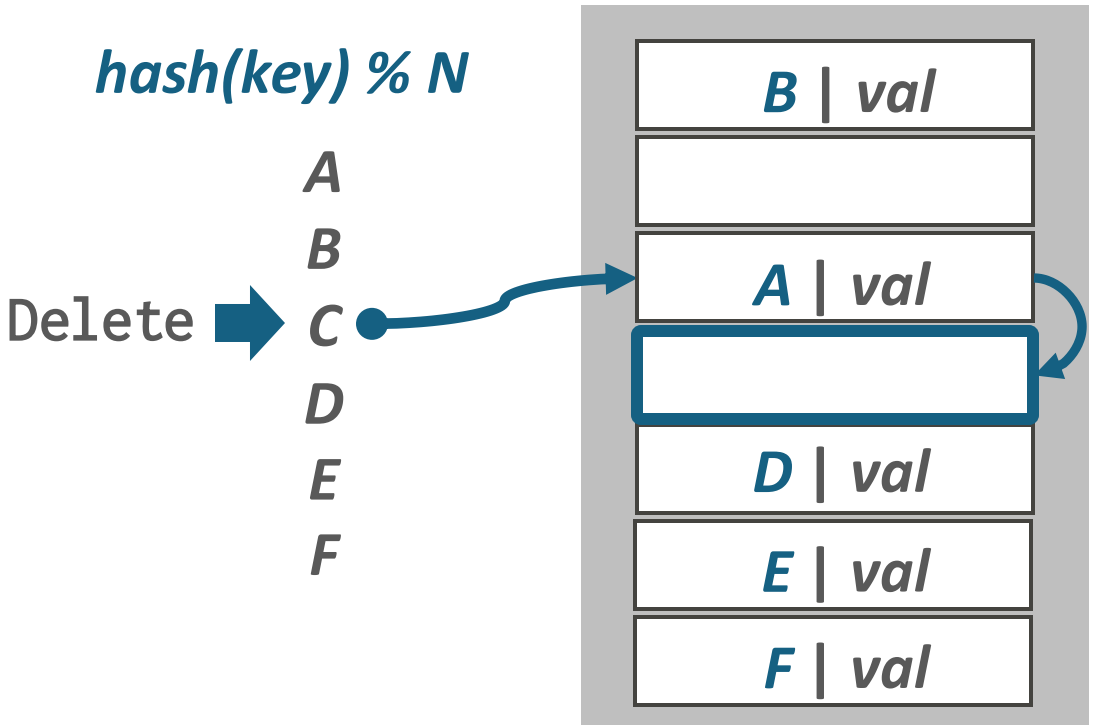

Linear Probe Hashing - Deletes

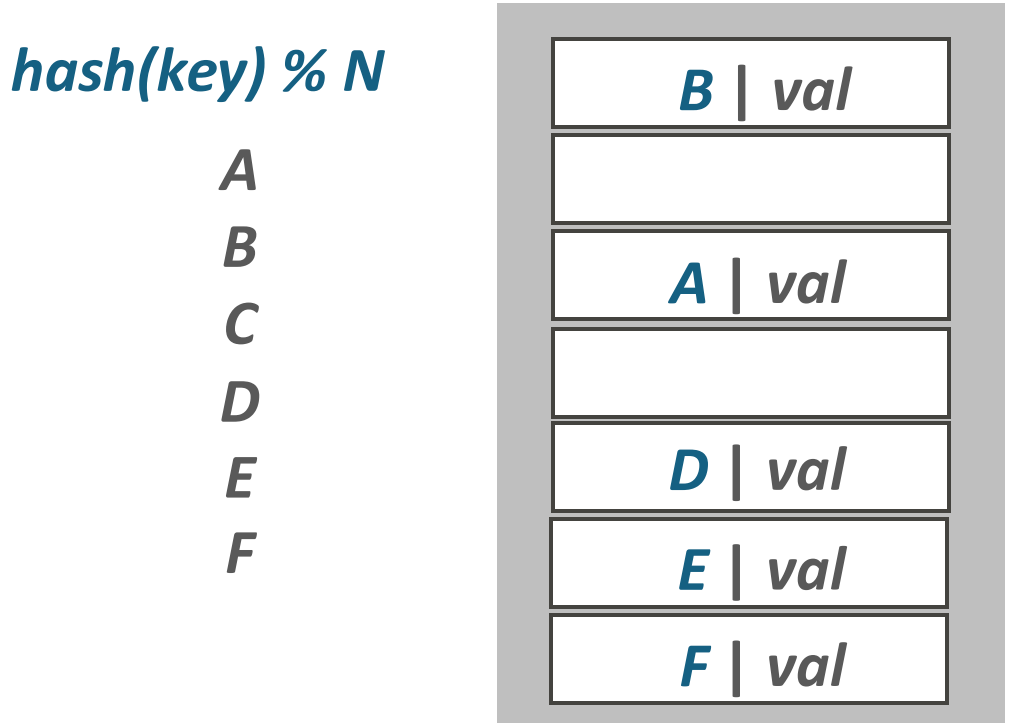

Linear Probe Hashing - Deletes

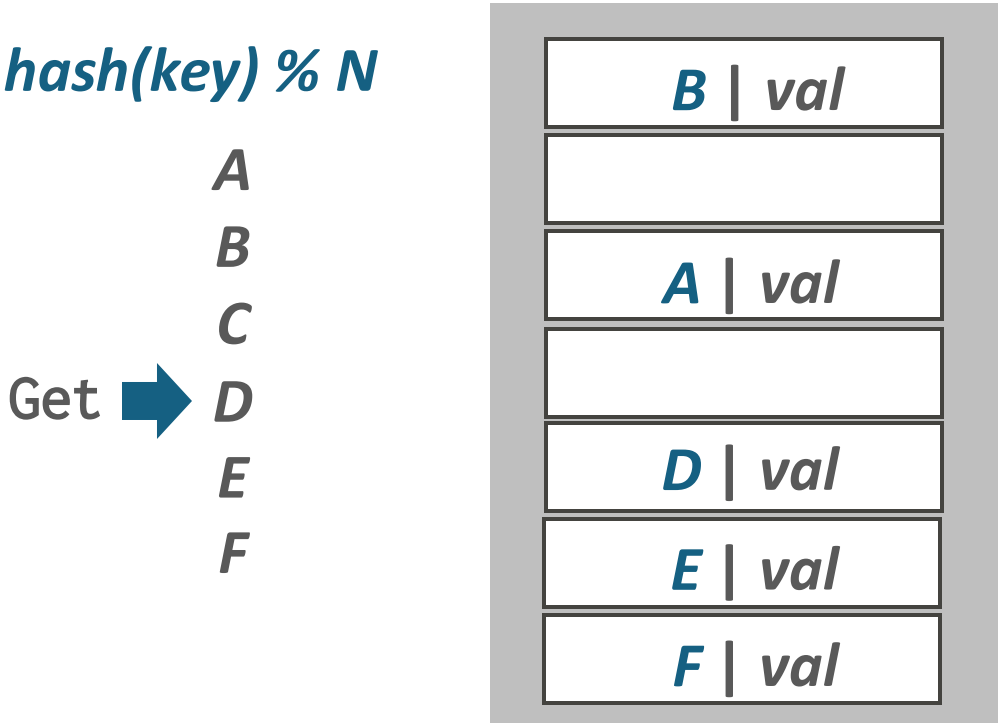

Linear Probe Hashing - Deletes

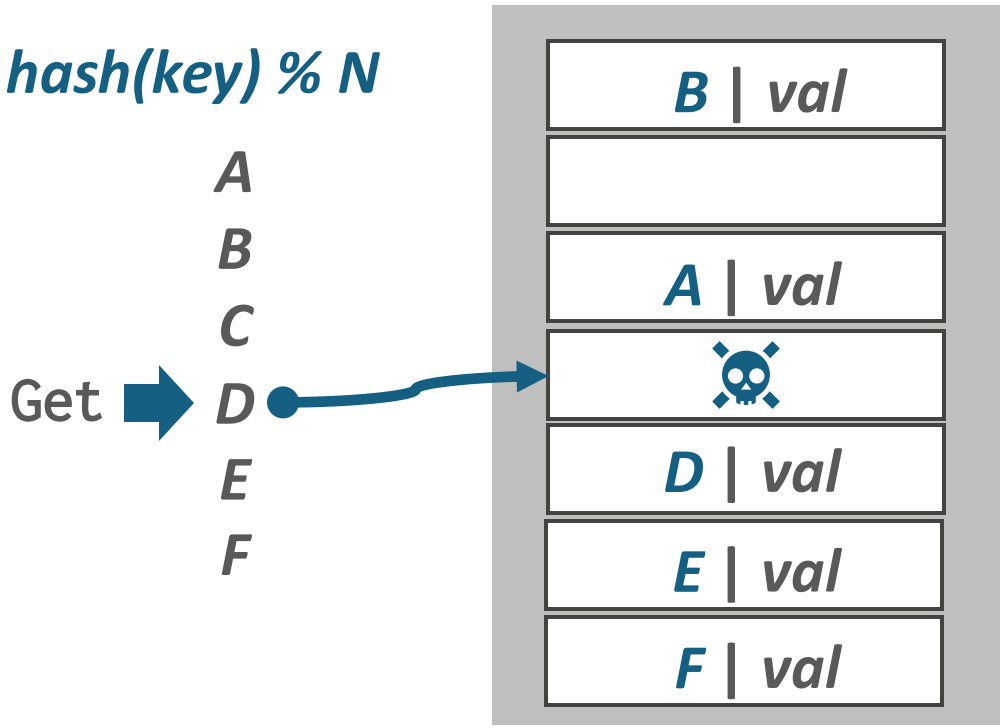

Linear Probe Hashing - Deletes

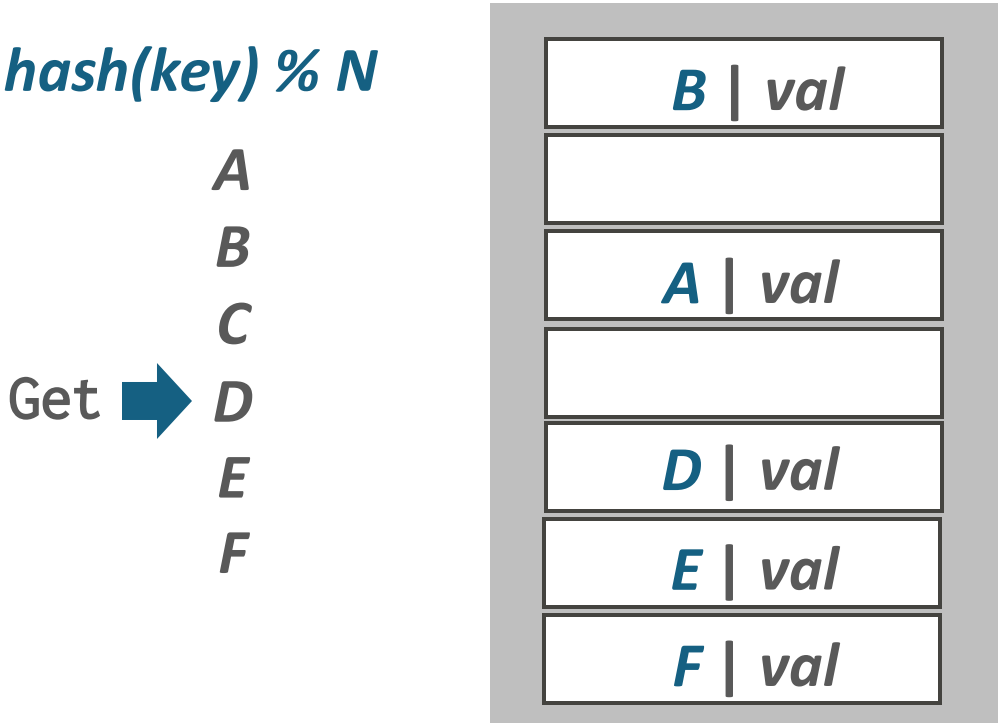

Linear Probe Hashing - Deletes

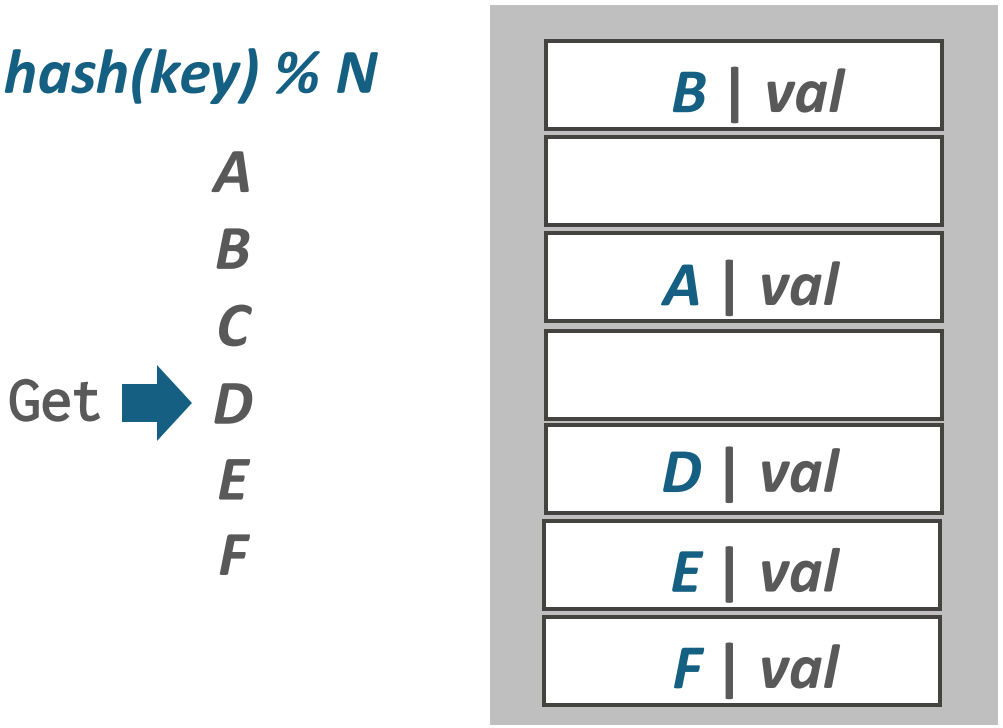

- **Approach #1: Movement**
  - Rehash keys until you find the first empty slot.

Linear Probe Hashing - Deletes

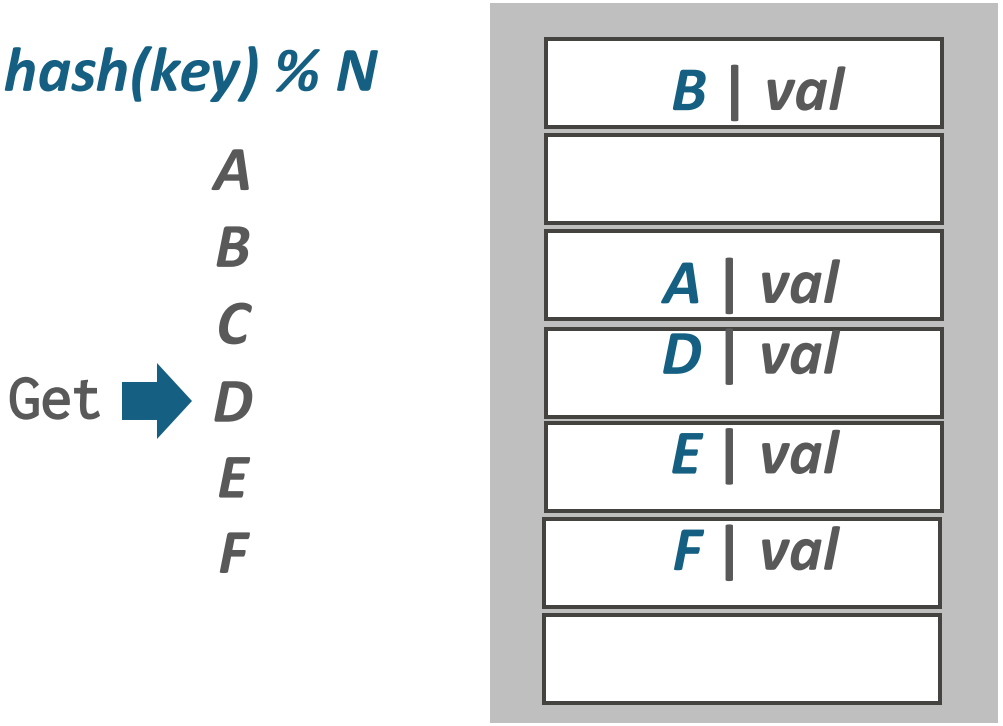

- **Approach #1: Movement**
  - Rehash keys until you find the first empty slot.

Linear Probe Hashing - Deletes

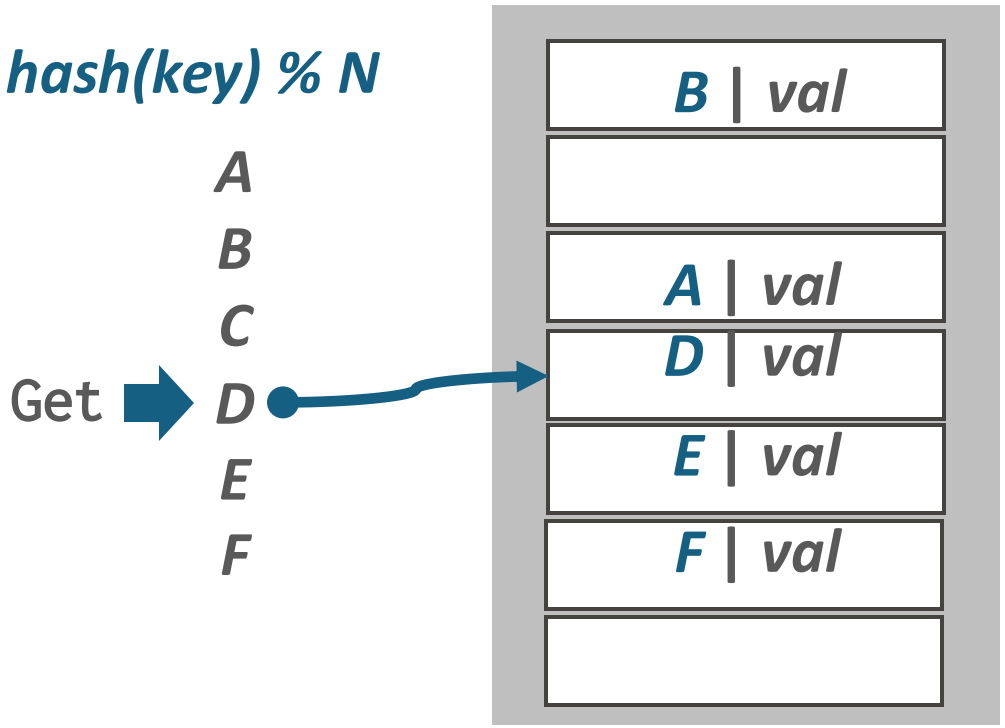

- **Approach #1: Movement**
  - Rehash keys until you find the first empty slot.

Linear Probe Hashing - Deletes

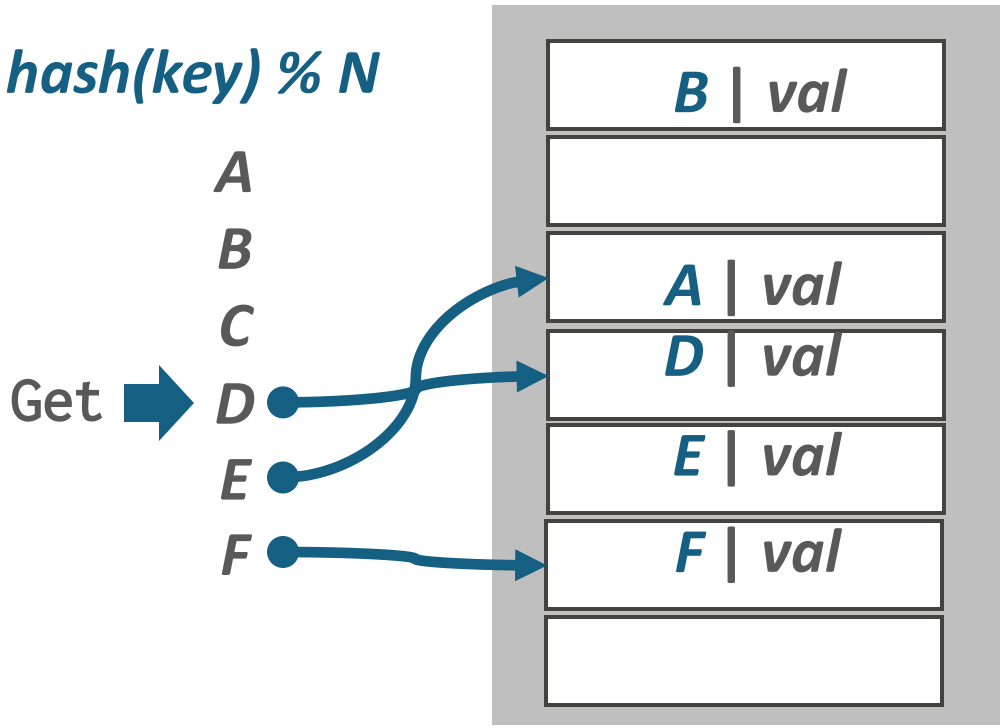

- **Approach #1: Movement**
  - Rehash keys until you find the first empty slot.

Linear Probe Hashing - Deletes

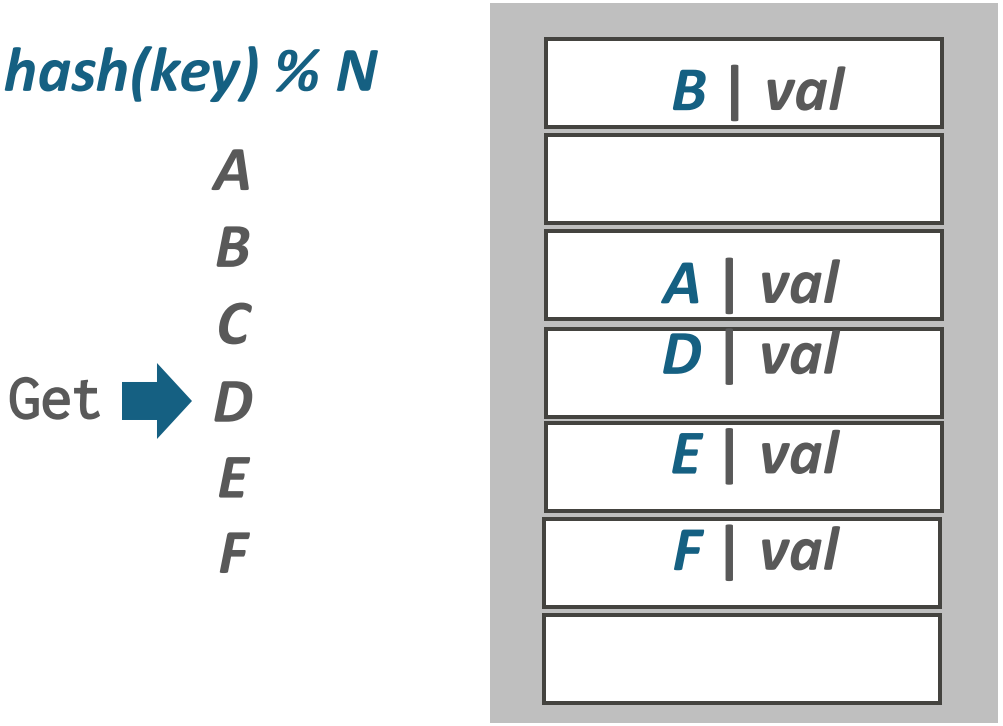

- **Approach #1: Movement**
  - Rehash keys until you find the first empty slot.

Linear Probe Hashing - Deletes

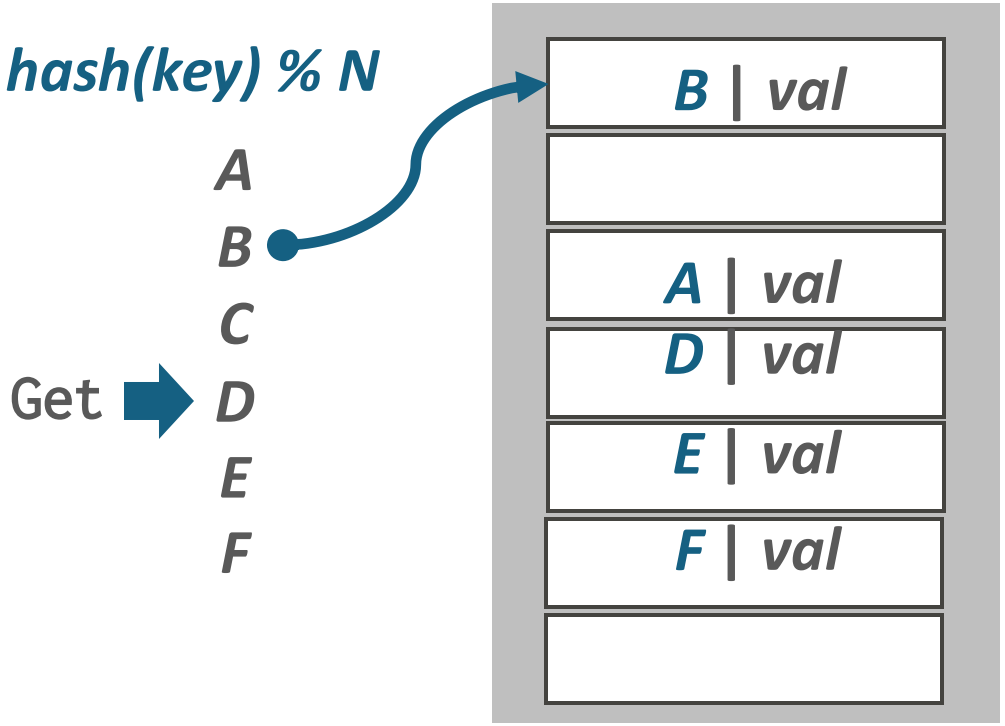

- **Approach #1: Movement**
  - Rehash keys until you find the first empty slot.

Linear Probe Hashing - Deletes

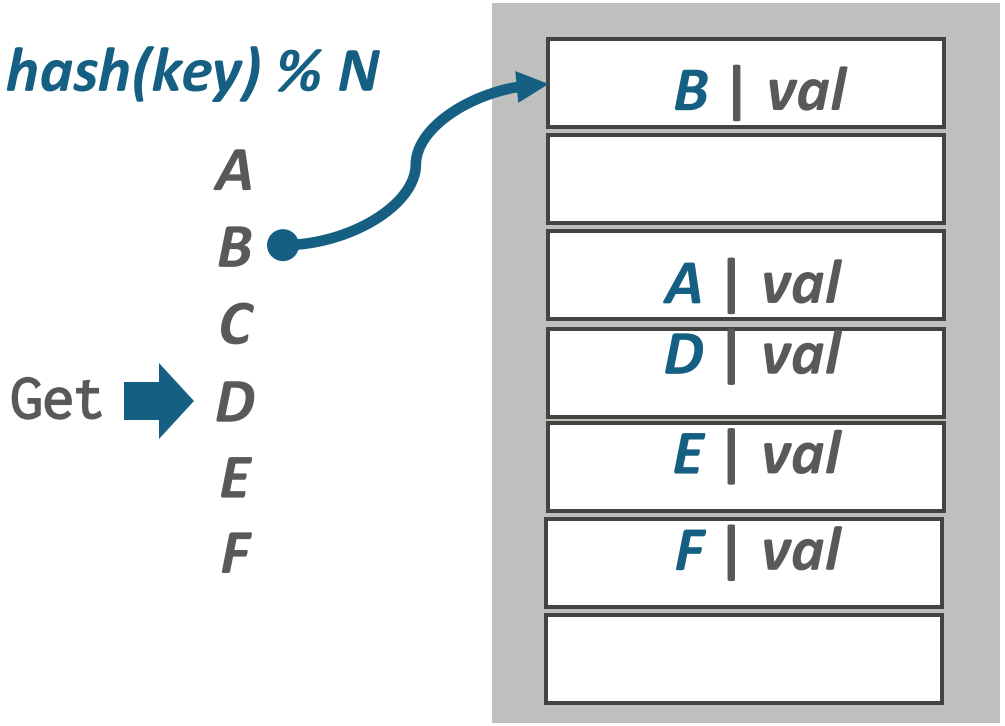

- **Approach #1: Movement**
  - Rehash keys until you find the first empty slot.
  - Expensive! May need to reorganize the entire table.
  - No DBMS does this.

Linear Probe Hashing - Deletes

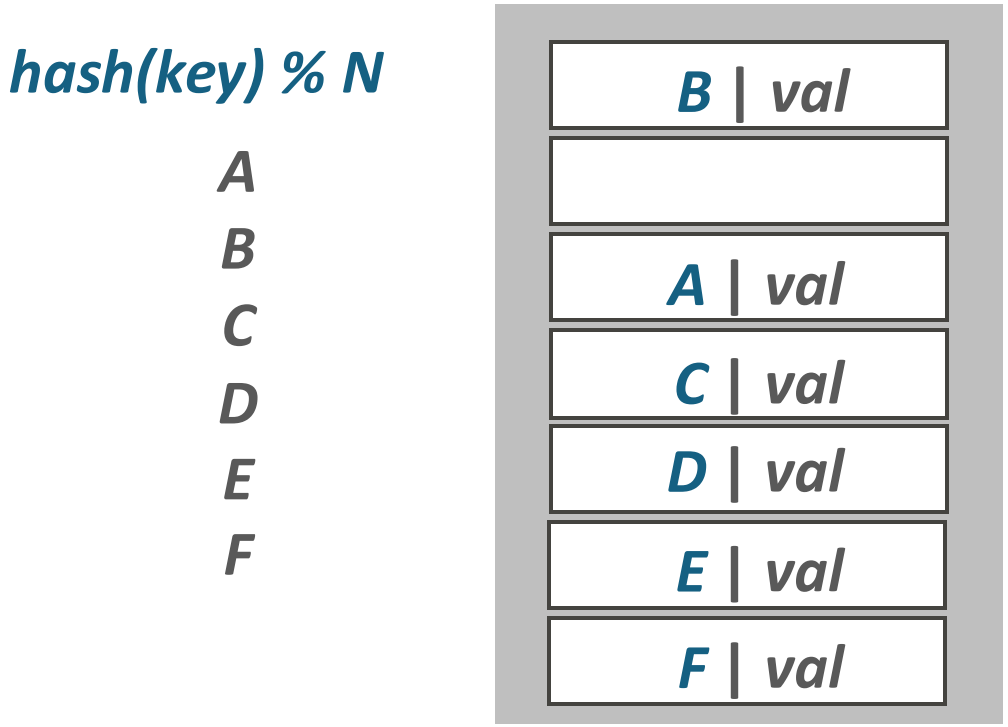

- **Approach #2: Tombstone**
  - Set a marker to indicate that the entry in the slot is logically deleted.
  - Reuse the slot for new keys.
  - May need periodic garbage collection.

Linear Probe Hashing - Deletes

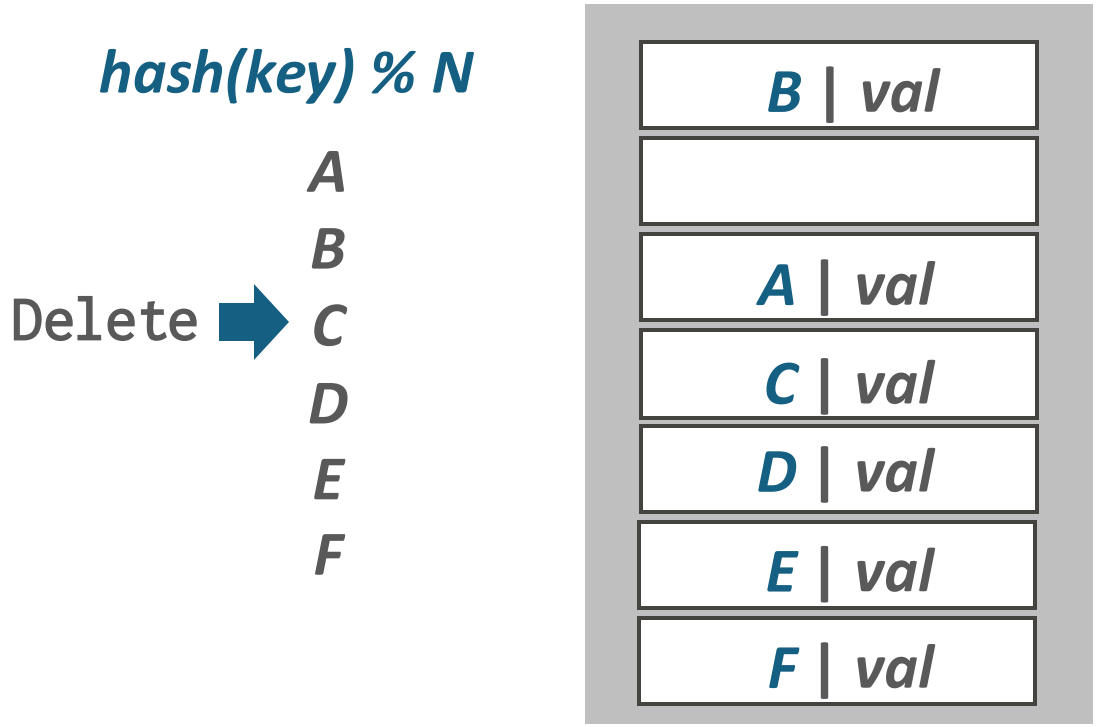

- **Approach #2: Tombstone**
  - Set a marker to indicate that the entry in the slot is logically deleted.
  - Reuse the slot for new keys.
  - May need periodic garbage collection.

Linear Probe Hashing - Deletes

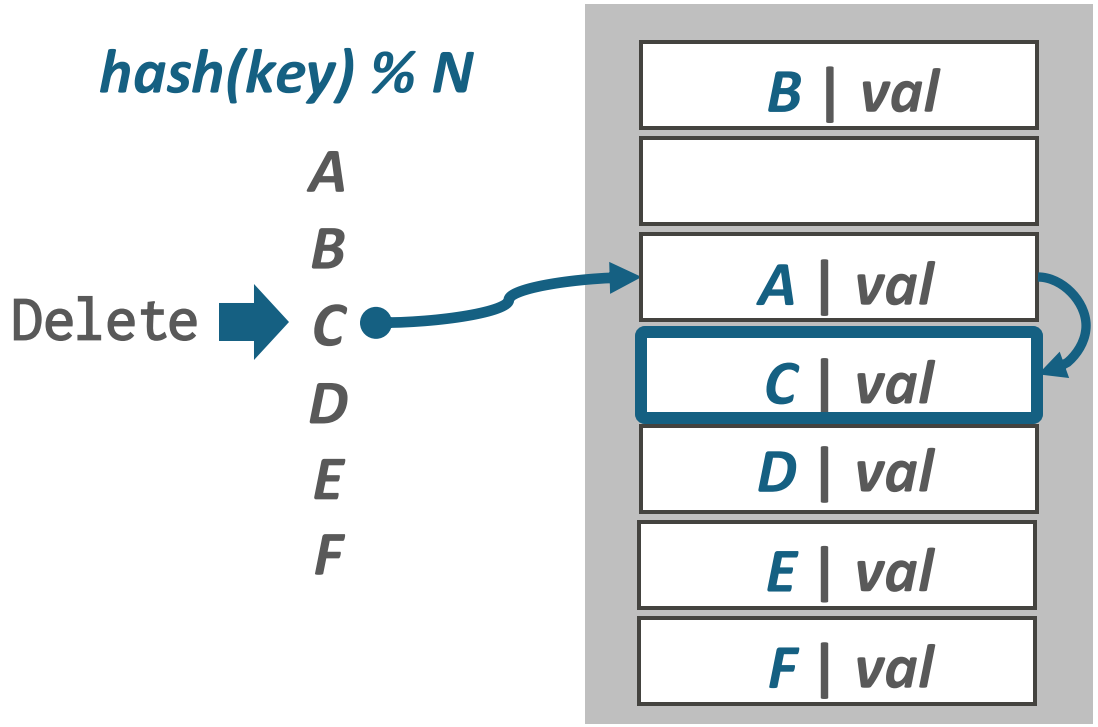

- **Approach #2: Tombstone**
  - Set a marker to indicate that the entry in the slot is logically deleted.
  - Reuse the slot for new keys.
  - May need periodic garbage collection.

Linear Probe Hashing - Deletes

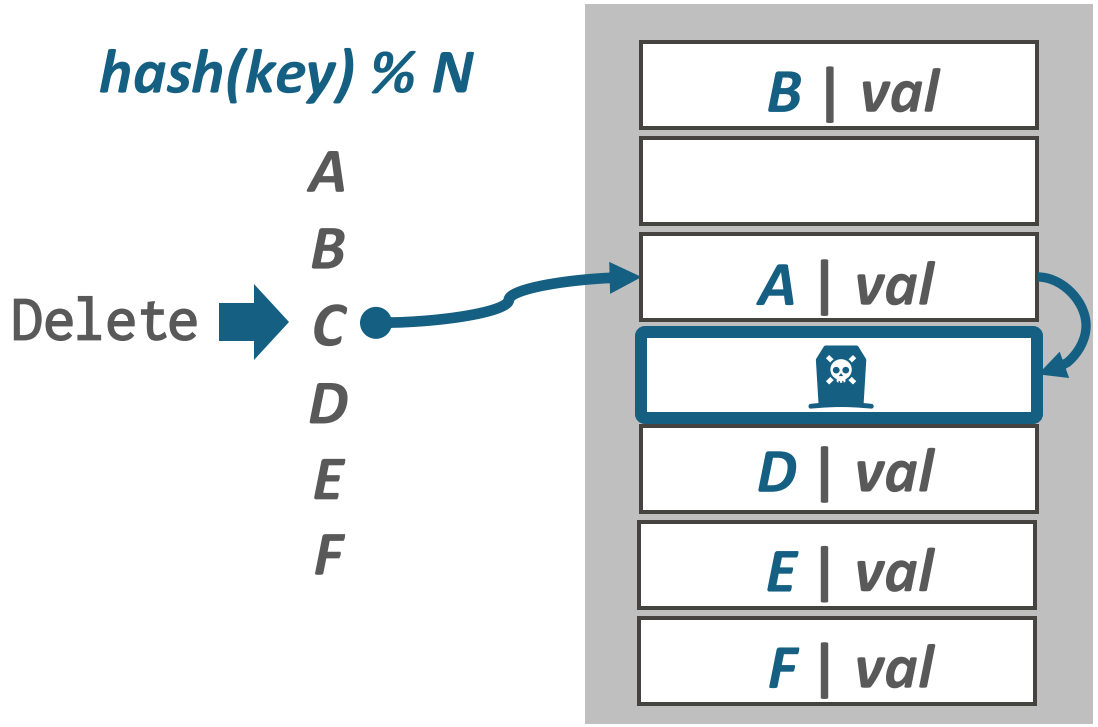

- **Approach #2: Tombstone**
  - Set a marker to indicate that the entry in the slot is logically deleted.
  - Reuse the slot for new keys.
  - May need periodic garbage collection.

Linear Probe Hashing - Deletes

hash(key) % N

A
B
C
D
E
F

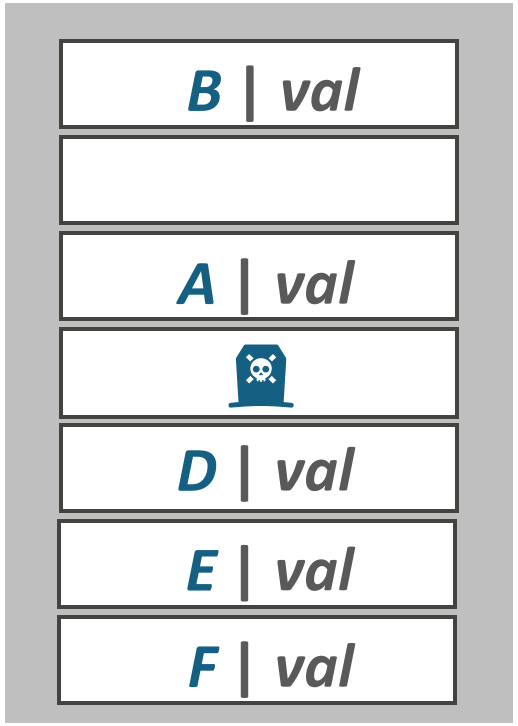

- **Approach #2: Tombstone**
  - Set a marker to indicate that the entry in the slot is logically deleted.
  - Reuse the slot for new keys.
  - May need periodic garbage collection.

Linear Probe Hashing - Deletes

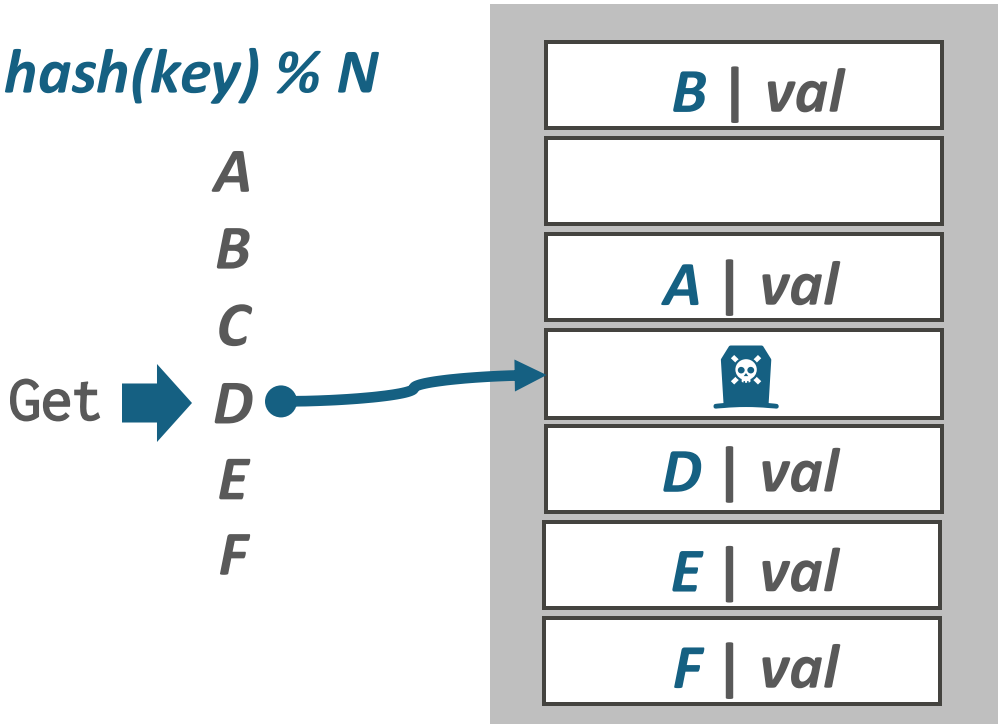

- **Approach #2: Tombstone**
  - Set a marker to indicate that the entry in the slot is logically deleted.
  - Reuse the slot for new keys.
  - May need periodic garbage collection.

Linear Probe Hashing - Deletes

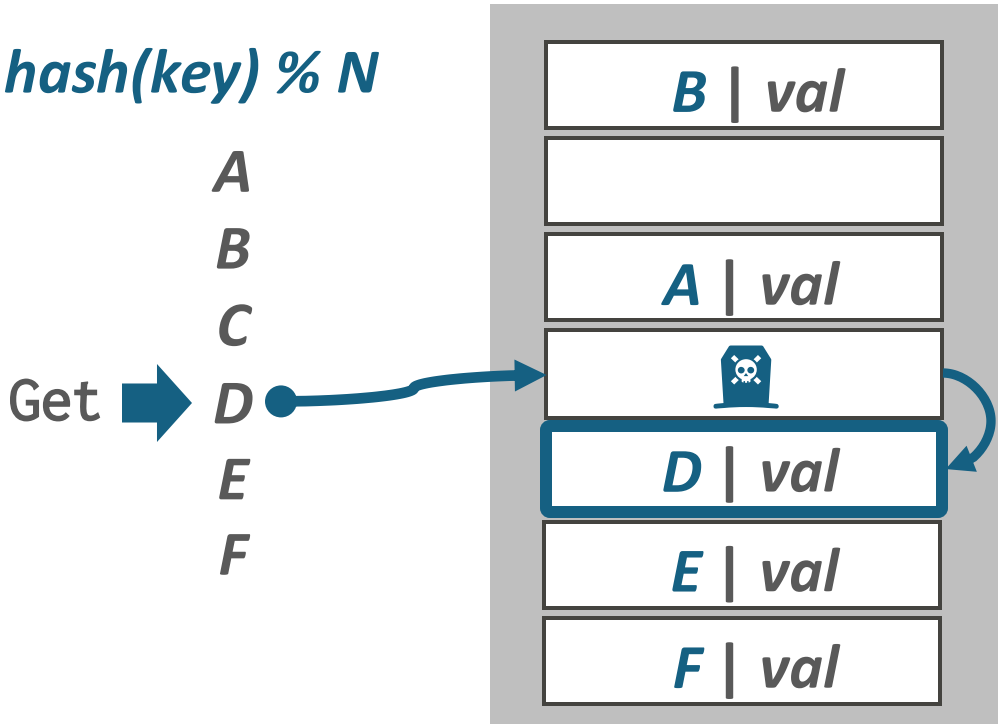

- **Approach #2: Tombstone**
  - Set a marker to indicate that the entry in the slot is logically deleted.
  - Reuse the slot for new keys.
  - May need periodic garbage collection.

Linear Probe Hashing - Deletes

hash(key) % N

A
B
C
D
E
F

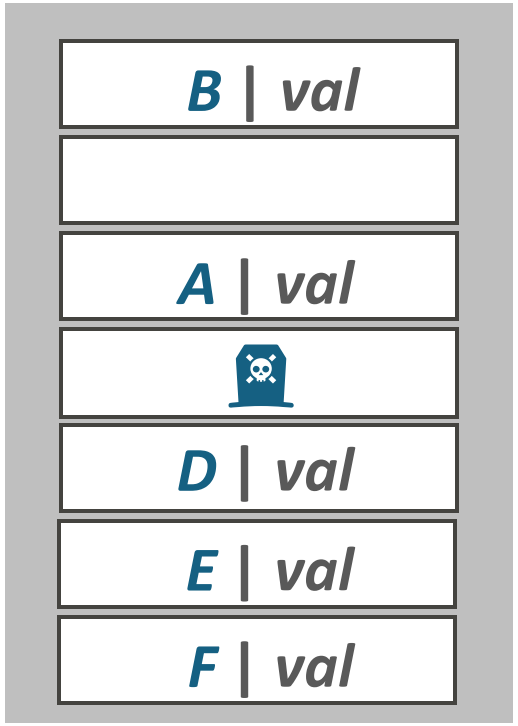

- **Approach #2: Tombstone**
  - Set a marker to indicate that the entry in the slot is logically deleted.
  - Reuse the slot for new keys.
  - May need periodic garbage collection.

Linear Probe Hashing - Deletes

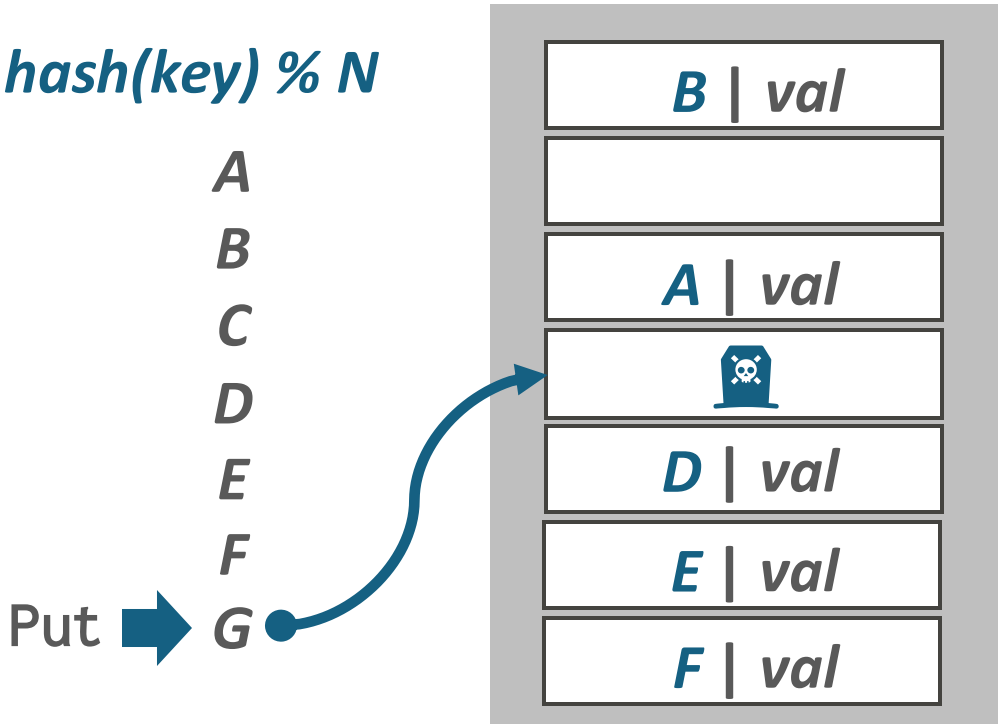

- **Approach #2: Tombstone**
  - Set a marker to indicate that the entry in the slot is logically deleted.
  - Reuse the slot for new keys.
  - May need periodic garbage collection.

Linear Probe Hashing - Deletes

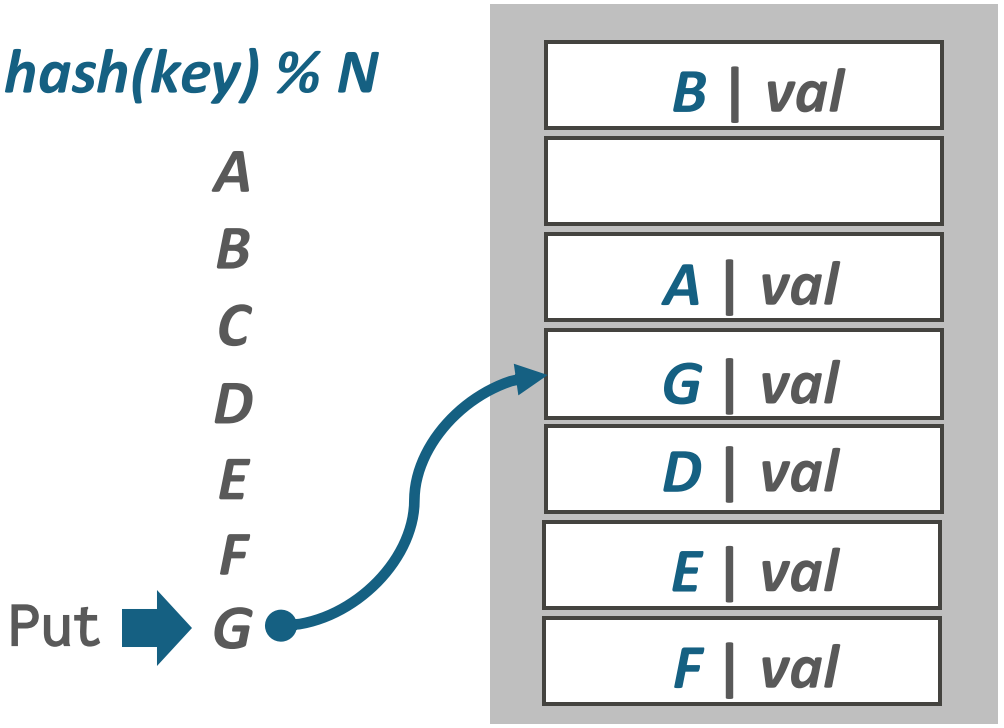

- **Approach #2: Tombstone**
  - Set a marker to indicate that the entry in the slot is logically deleted.
  - Reuse the slot for new keys.
  - May need periodic garbage collection.

Non-Unique Keys

- **Choice #1: Separate Linked List**
  - Store values in separate storage area for each key.
  - Value lists can overflow to multiple pages if the number of duplicates is large.

Non-Unique Keys

- **Choice #1: Separate Linked List**
  - Store values in separate storage area for each key.
  - Value lists can overflow to multiple pages if the number of duplicates is large.

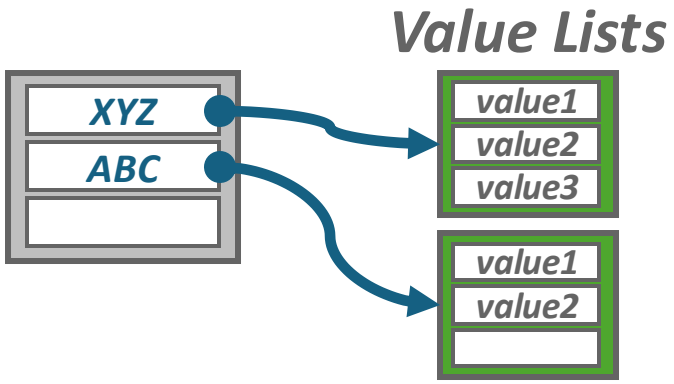

Non-Unique Keys

- **Choice #1: Separate Linked List**
  - Store values in separate storage area for each key.
  - Value lists can overflow to multiple pages if the number of duplicates is large.
- **Choice #2: Redundant Keys**
  - Store duplicate keys entries together in the hash table.
  - This is what most systems do.

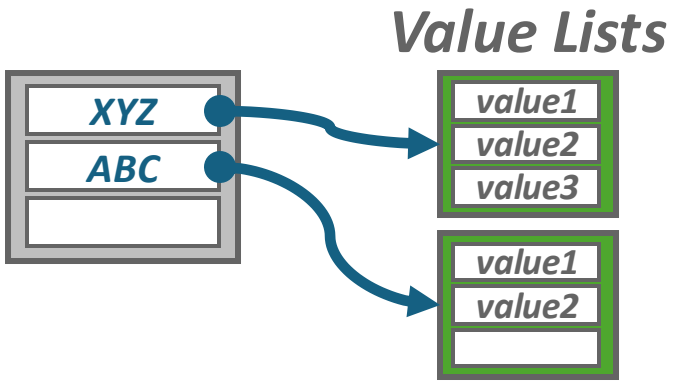

Non-Unique Keys

- **Choice #1: Separate Linked List**
  - Store values in separate storage area for each key.
  - Value lists can overflow to multiple pages if the number of duplicates is large.
- **Choice #2: Redundant Keys**
  - Store duplicate keys entries together in the hash table.
  - This is what most systems do.

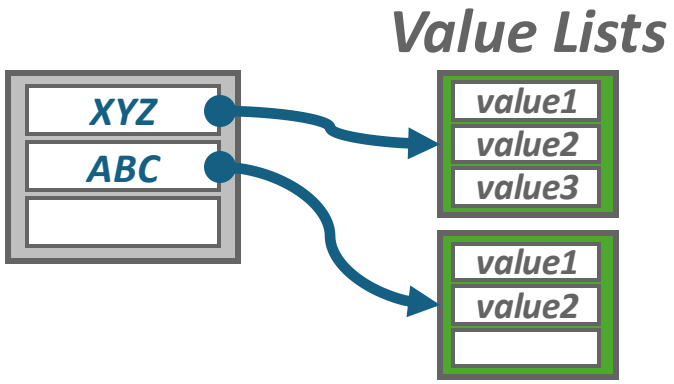

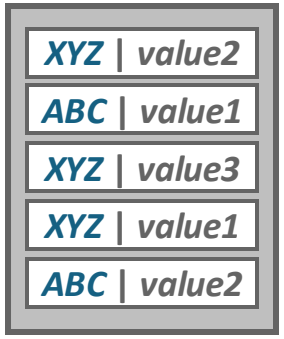

Optimizations

- Specialized hash table implementations based on key type(s) and sizes.
  - Example: Maintain multiple hash tables for different string sizes for a set of keys.
- Store metadata separate in a separate array.
  - Packed bitmap tracks whether a slot is empty/tombstone.
- Use table + slot versioning metadata to quickly invalidate all entries in the hash table.
  - Example: If table version does not match slot version, then treat the slot as empty.

Source: Maksim Kita

Cuckoo Hashing

Static Hashing Schemes

Cuckoo Hashing

- Use multiple hash functions to find multiple locations in the hash table to insert records.
  - On insert, check multiple locations and pick the one that is empty.
  - If no location is available, evict the element from one of them and then re-hash it find a new location.
- Look-ups and deletions are always $\boldsymbol{O(1)}$ because only one location per hash table is checked.
- Best open-source implementation is from CMU.

Cuckoo Hashing

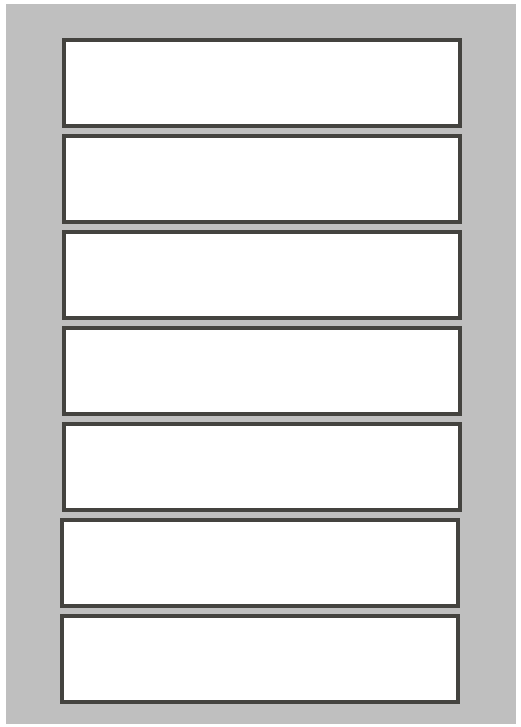

Cuckoo Hashing

Put A: $hash_1(A)$
$hash_2(A)$

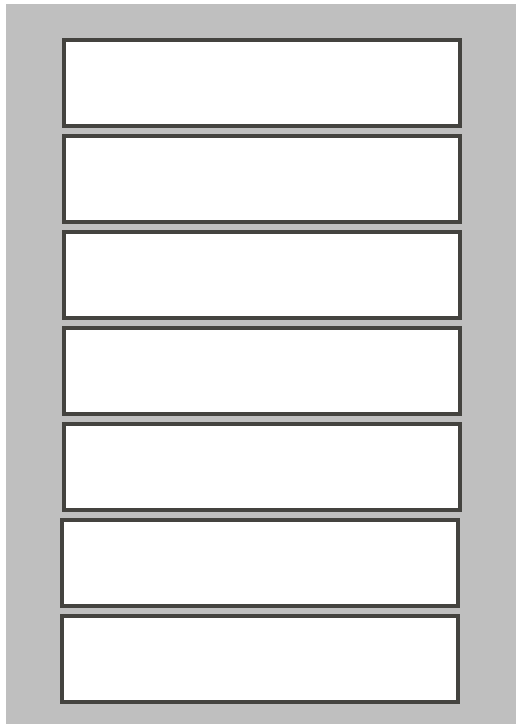

Cuckoo Hashing

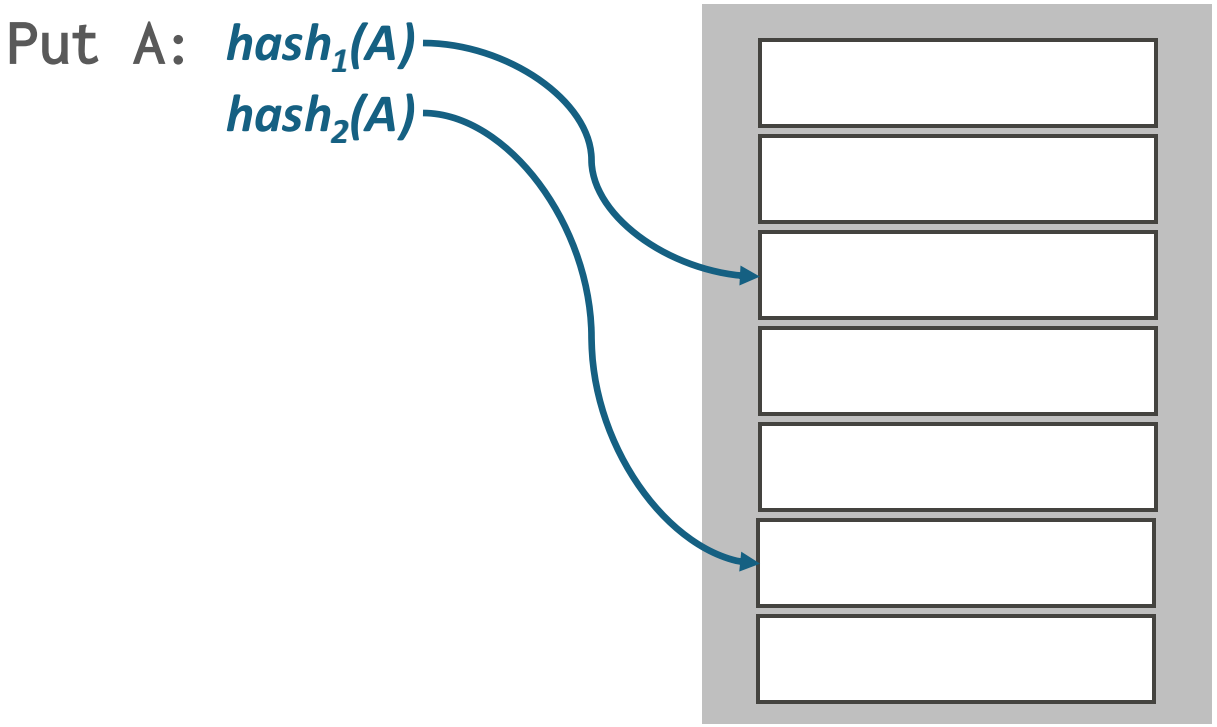

Cuckoo Hashing

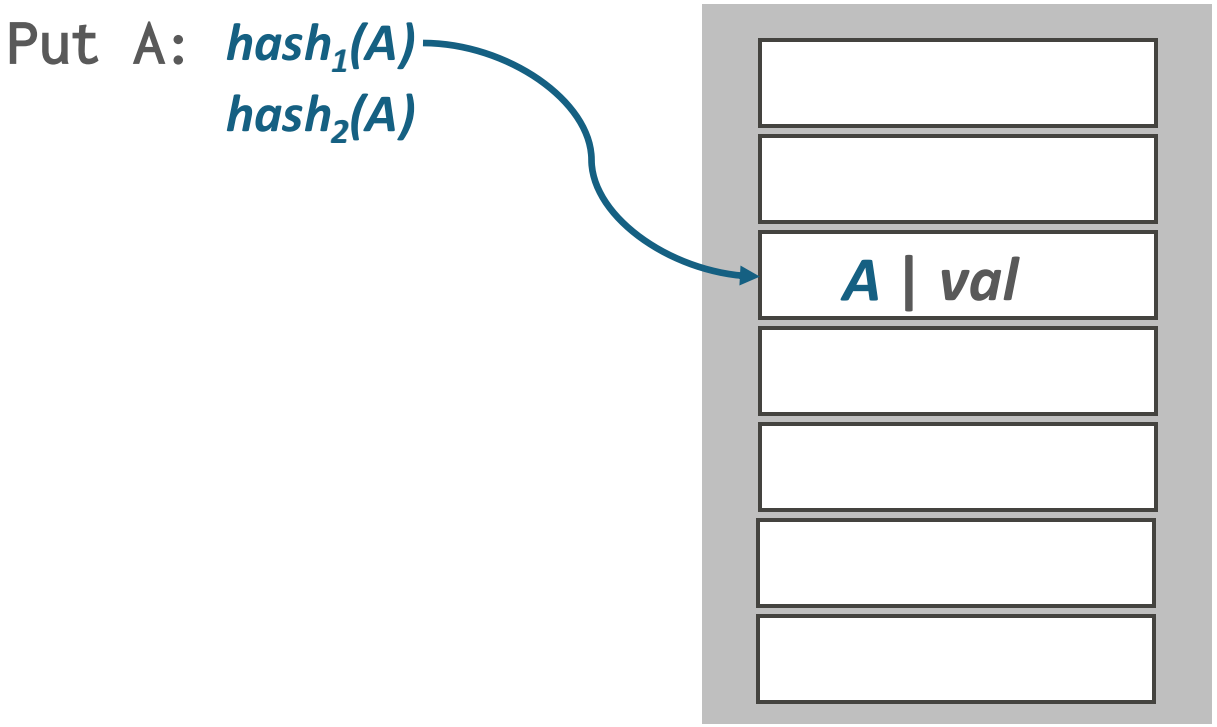

Cuckoo Hashing

Put A: $hash_1(A)$
$hash_2(A)$

Put B: $hash_1(B)$
$hash_2(B)$

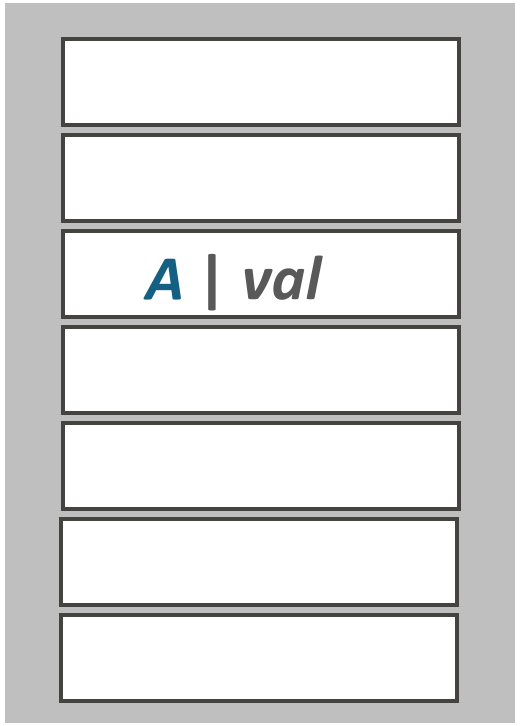

Cuckoo Hashing

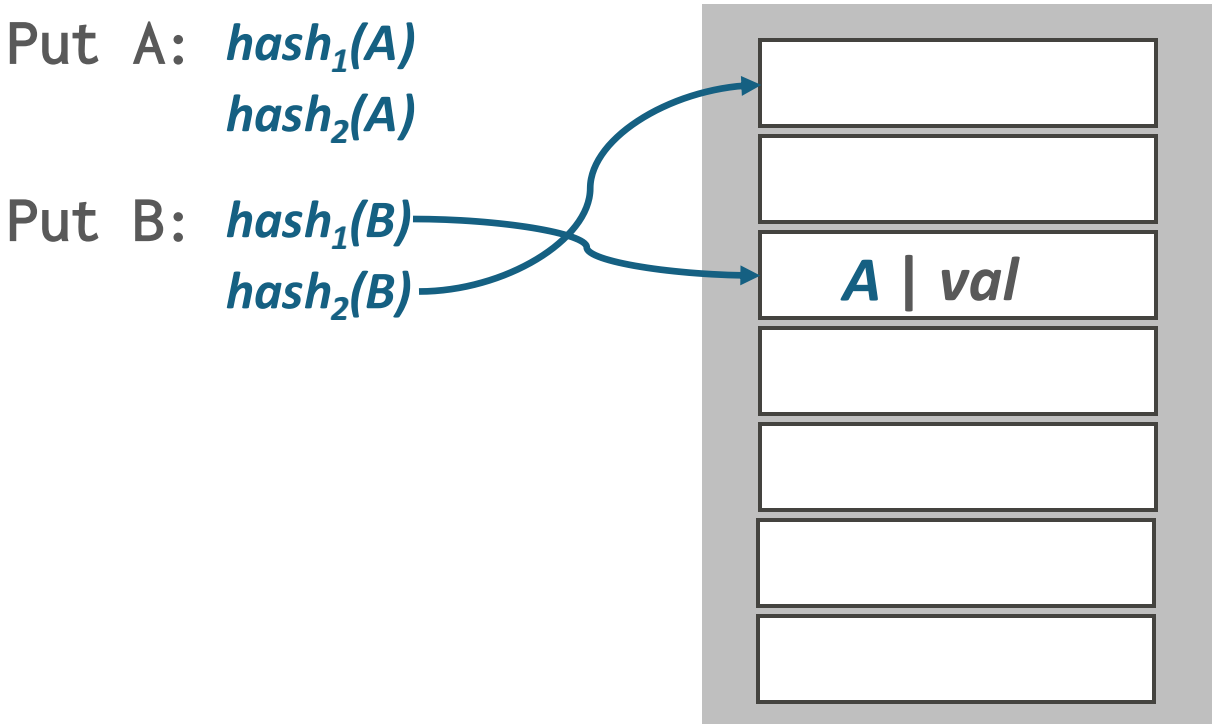

Cuckoo Hashing

Put A: $hash_1(A)$
$hash_2(A)$

Put B: $hash_1(B)$
$hash_2(B)$

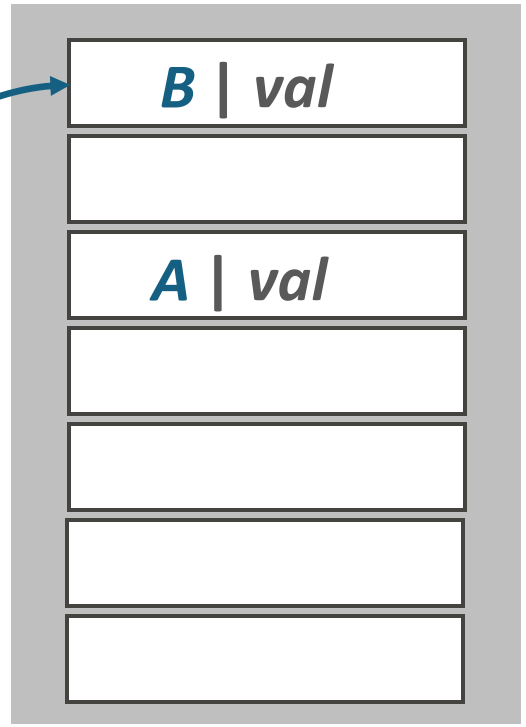

Cuckoo Hashing

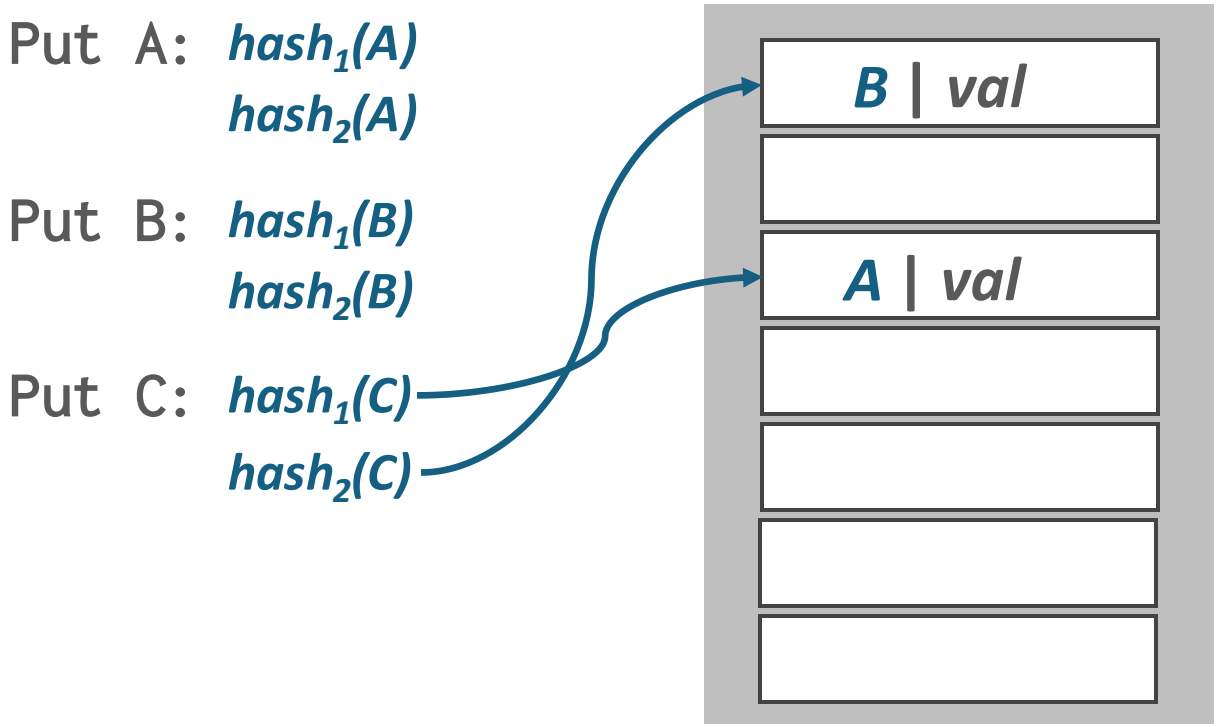

Cuckoo Hashing

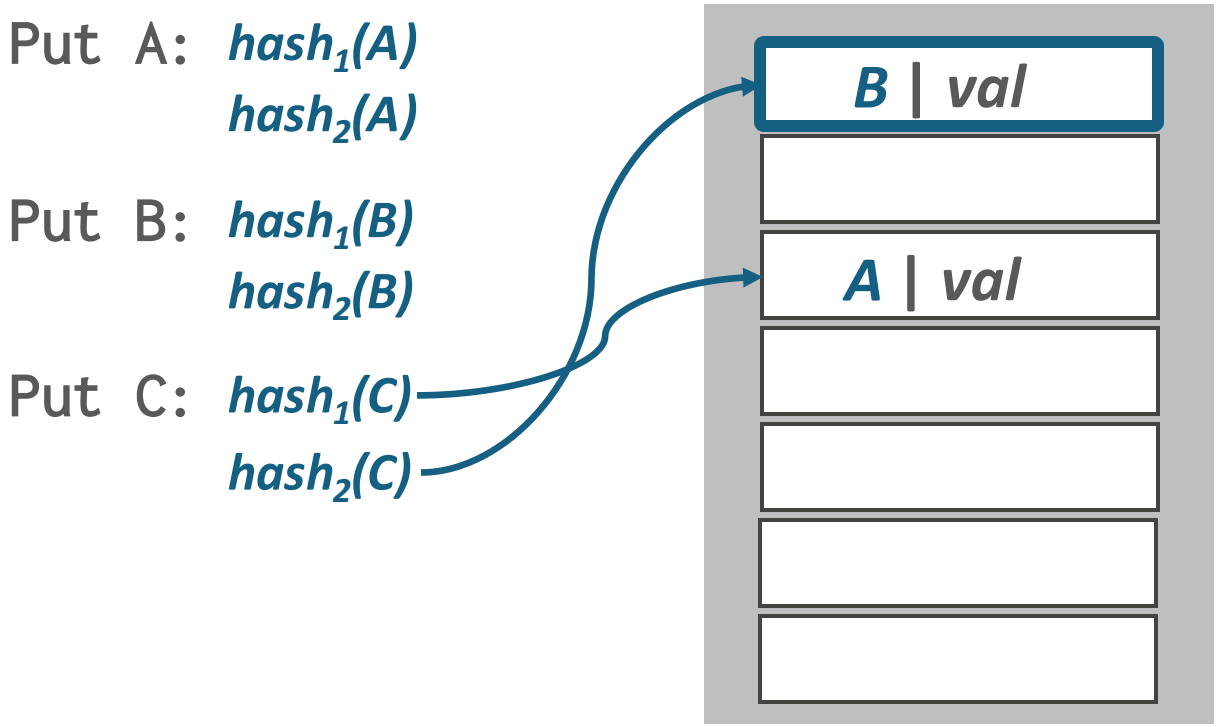

Cuckoo Hashing

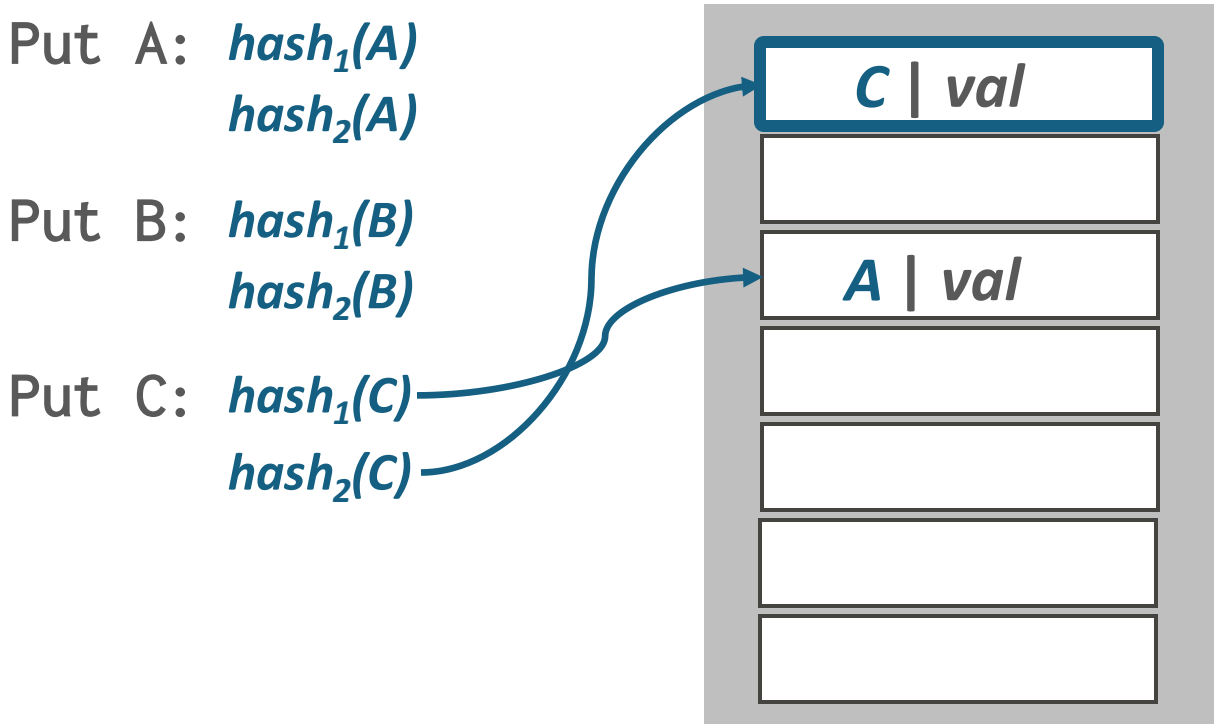

Cuckoo Hashing

Put A: $hash_1(A)$
$hash_2(A)$

Put B: $hash_1(B)$
$hash_2(B)$

Put C: $hash_1(C)$
$hash_2(C)$
$hash_1(B)$

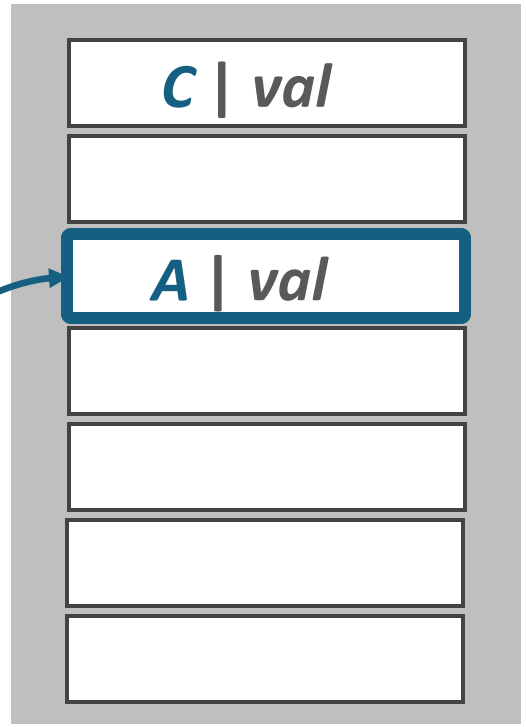

Cuckoo Hashing

Put A: $hash_1(A)$
$hash_2(A)$

Put B: $hash_1(B)$
$hash_2(B)$

Put C: $hash_1(C)$
$hash_2(C)$
$hash_1(B)$

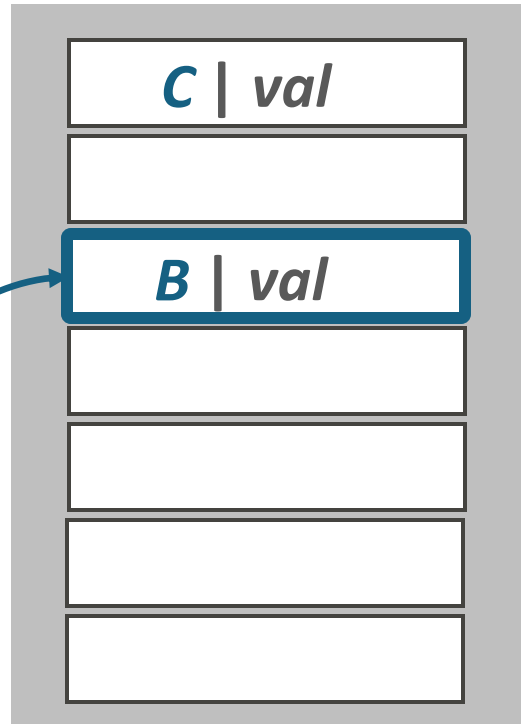

Cuckoo Hashing

Put A: $hash_1(A)$
$hash_2(A)$

Put B: $hash_1(B)$
$hash_2(B)$

Put C: $hash_1(C)$
$hash_2(C)$
$hash_1(B)$
$hash_2(A)$

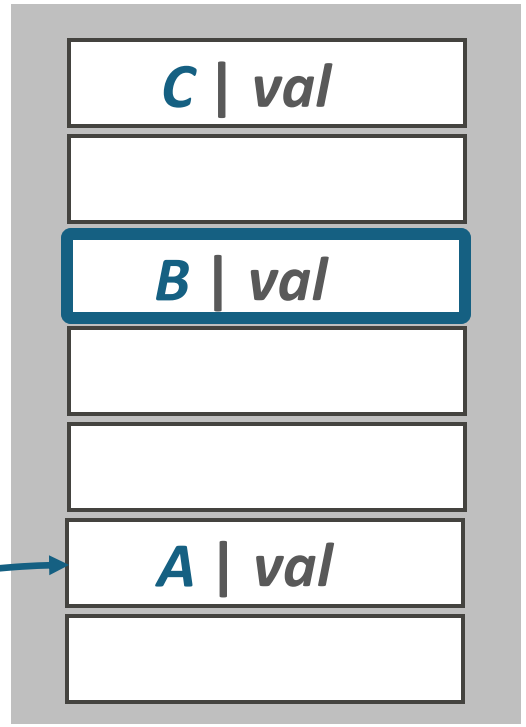

Cuckoo Hashing

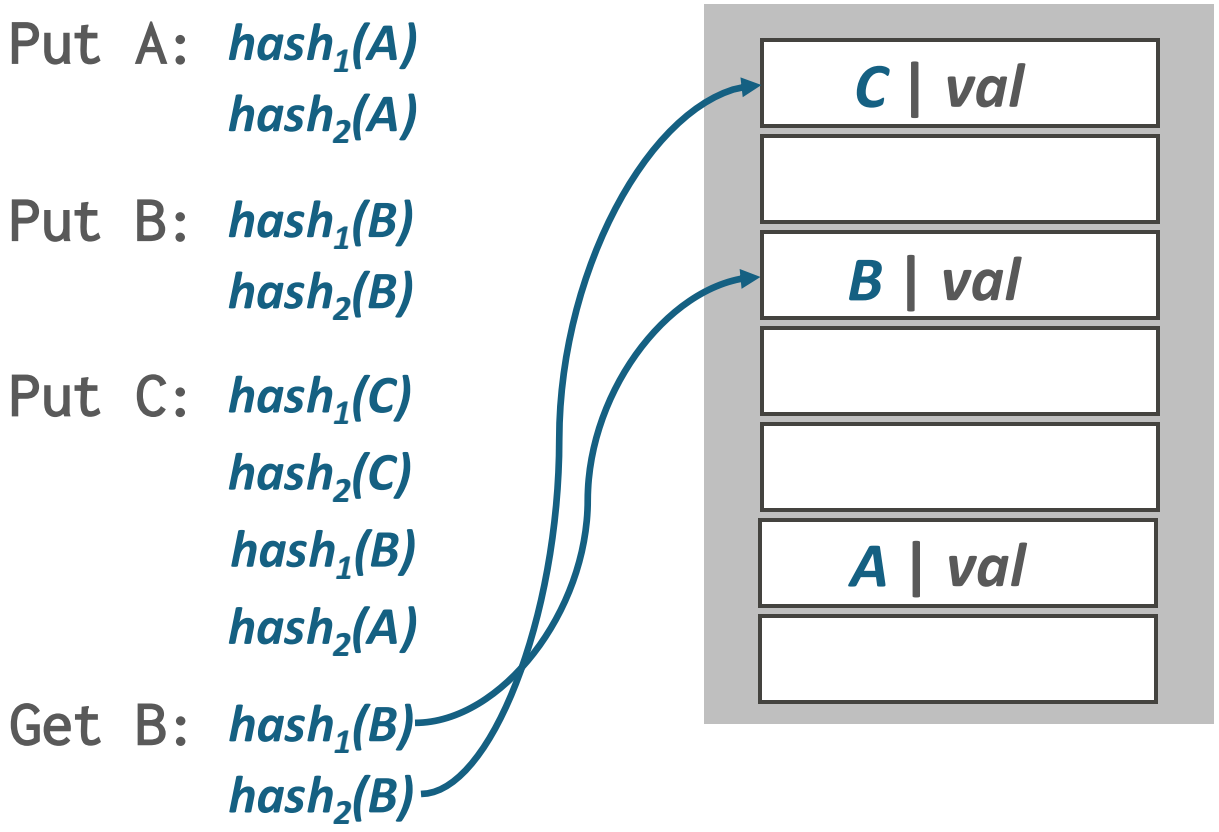

Chained Hash Table

Dynamic Hashing Schemes

Observation

- The previous hash tables require the DBMS to know the number of elements it wants to store.
  - Otherwise, it must rebuild the table if it needs to grow/shrink in size.

- Dynamic hash tables incrementally resize themselves as needed.
  - Chained Hashing
  - Extendible Hashing
  - Linear Hashing

Chained Hashing

- Maintain a linked list of buckets for each slot in the hash table.
- Resolve collisions by placing all elements with the same hash key into the same bucket.
  - To determine whether an element is present, hash to its bucket and scan for it.
  - Insertions and deletions are generalizations of lookups.

Chained Hashing

hash(key) % N

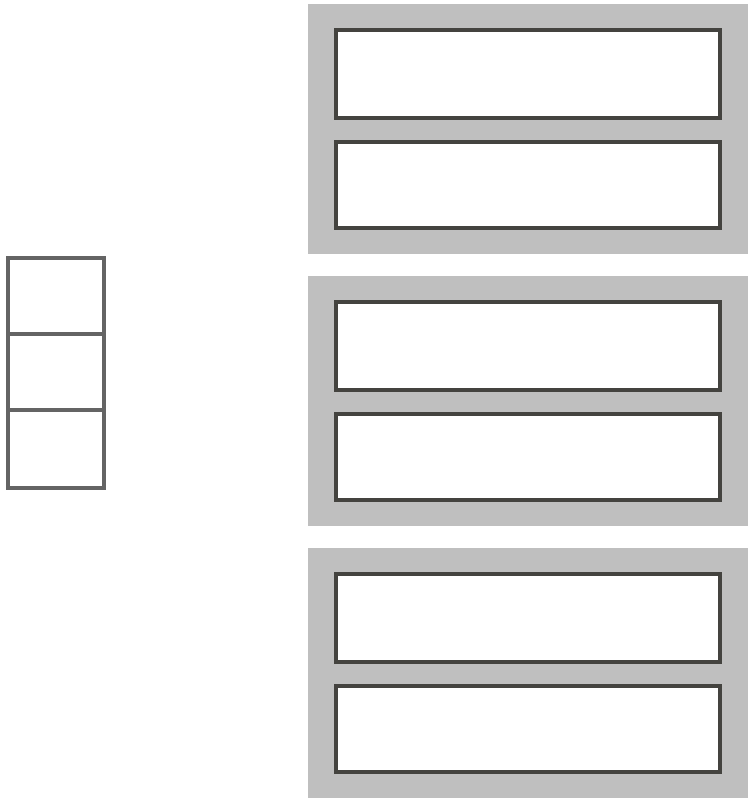

Chained Hashing

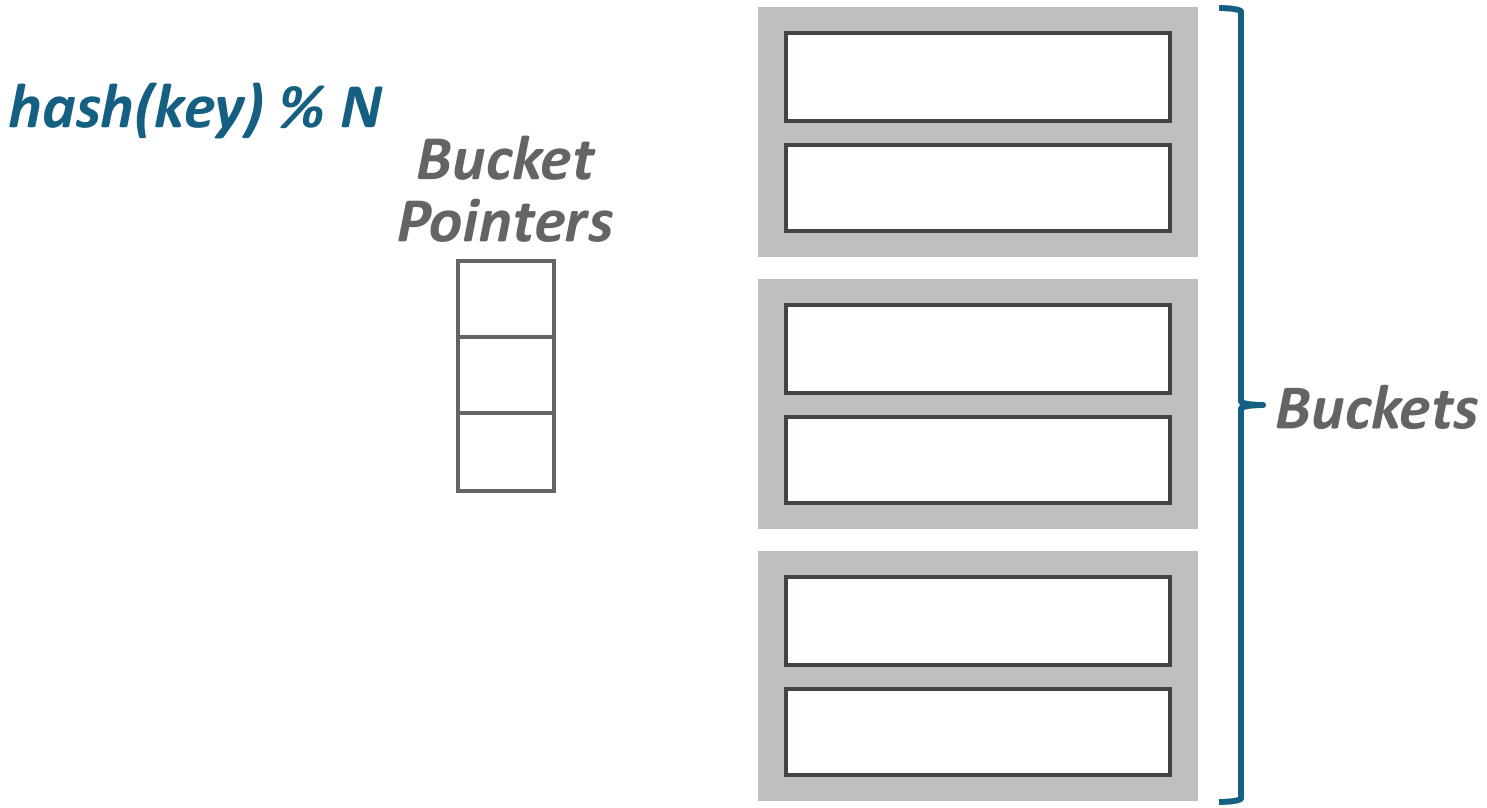

Chained Hashing

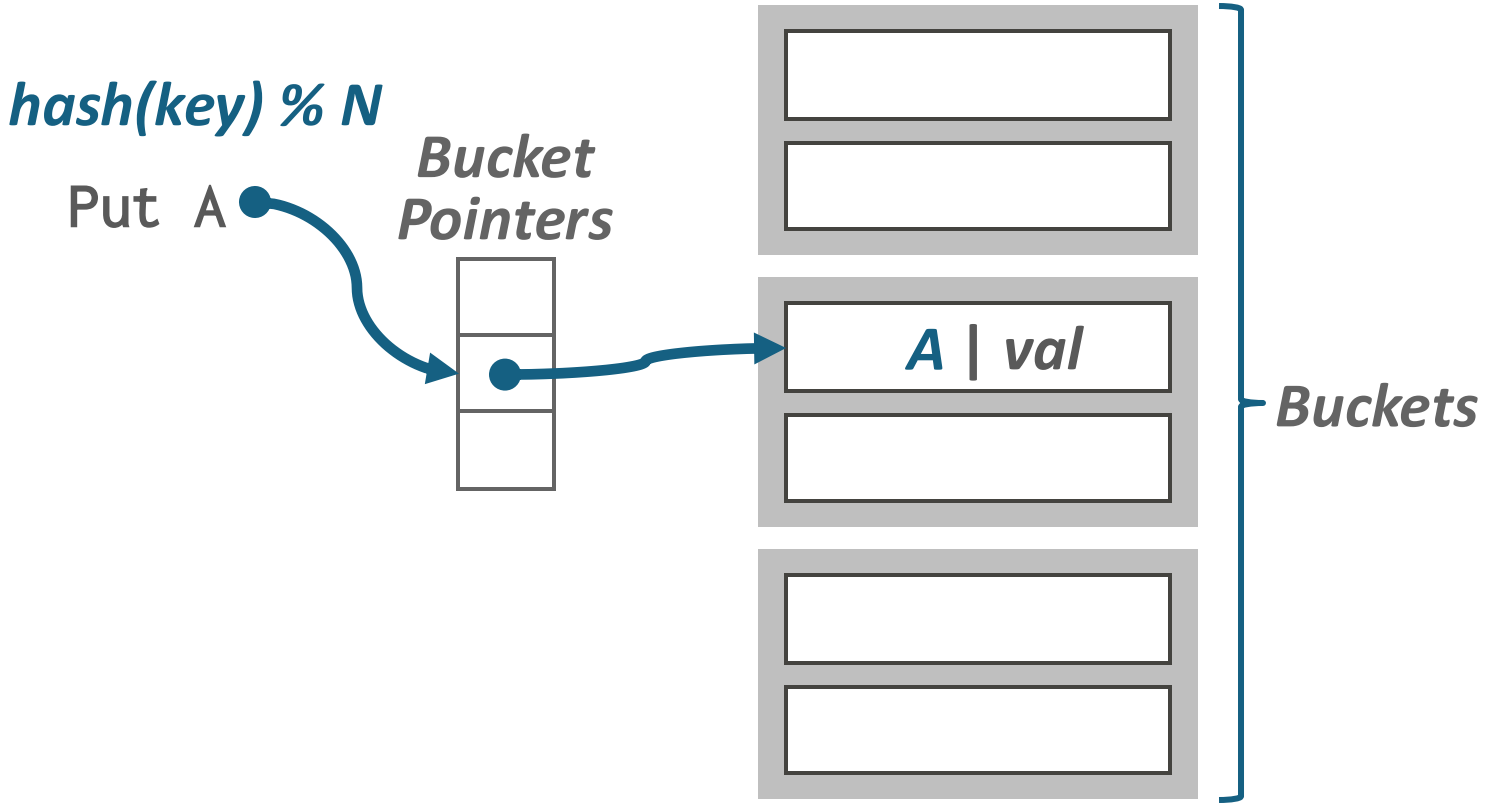

Chained Hashing

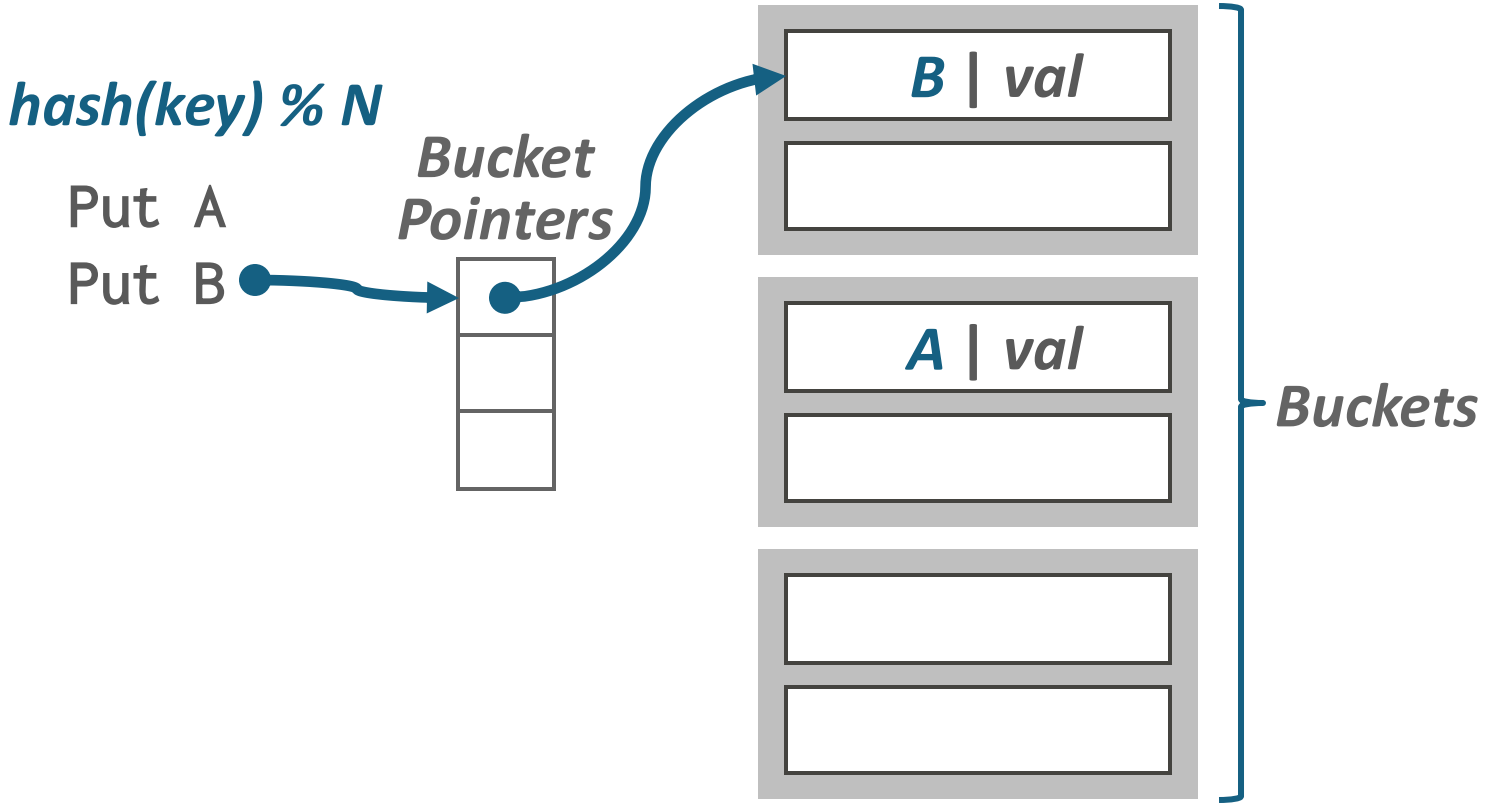

Chained Hashing

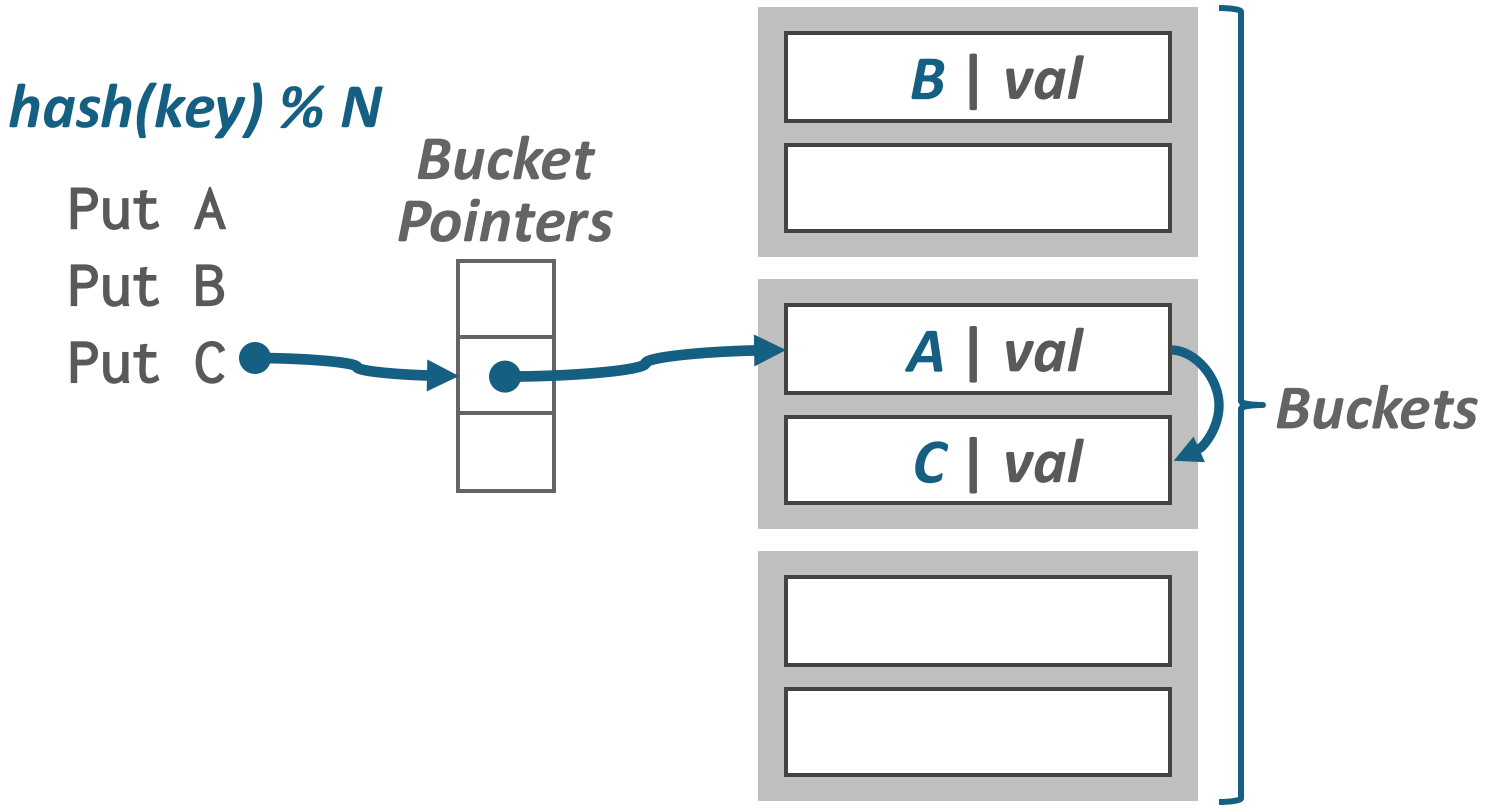

Chained Hashing

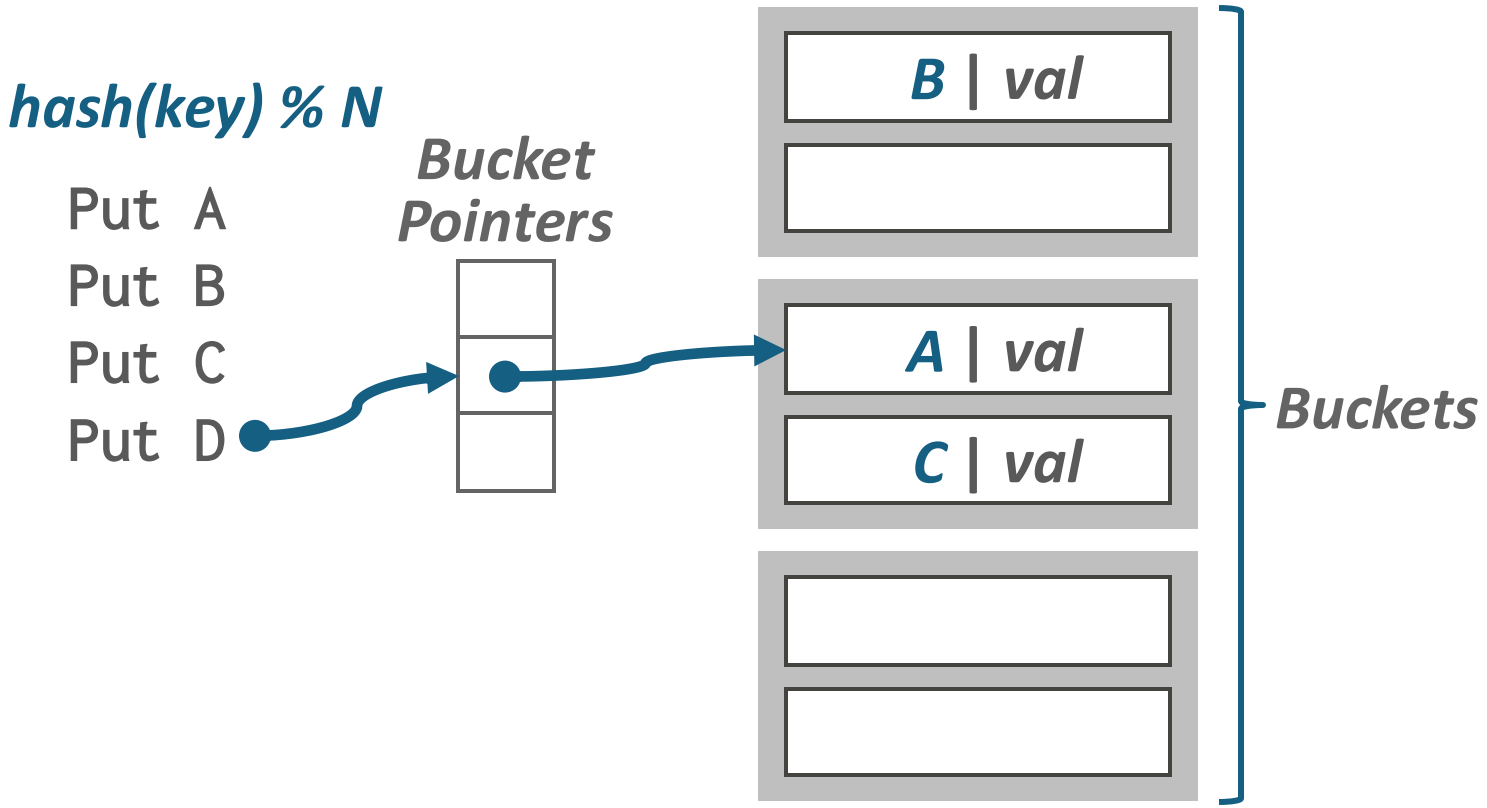

Chained Hashing

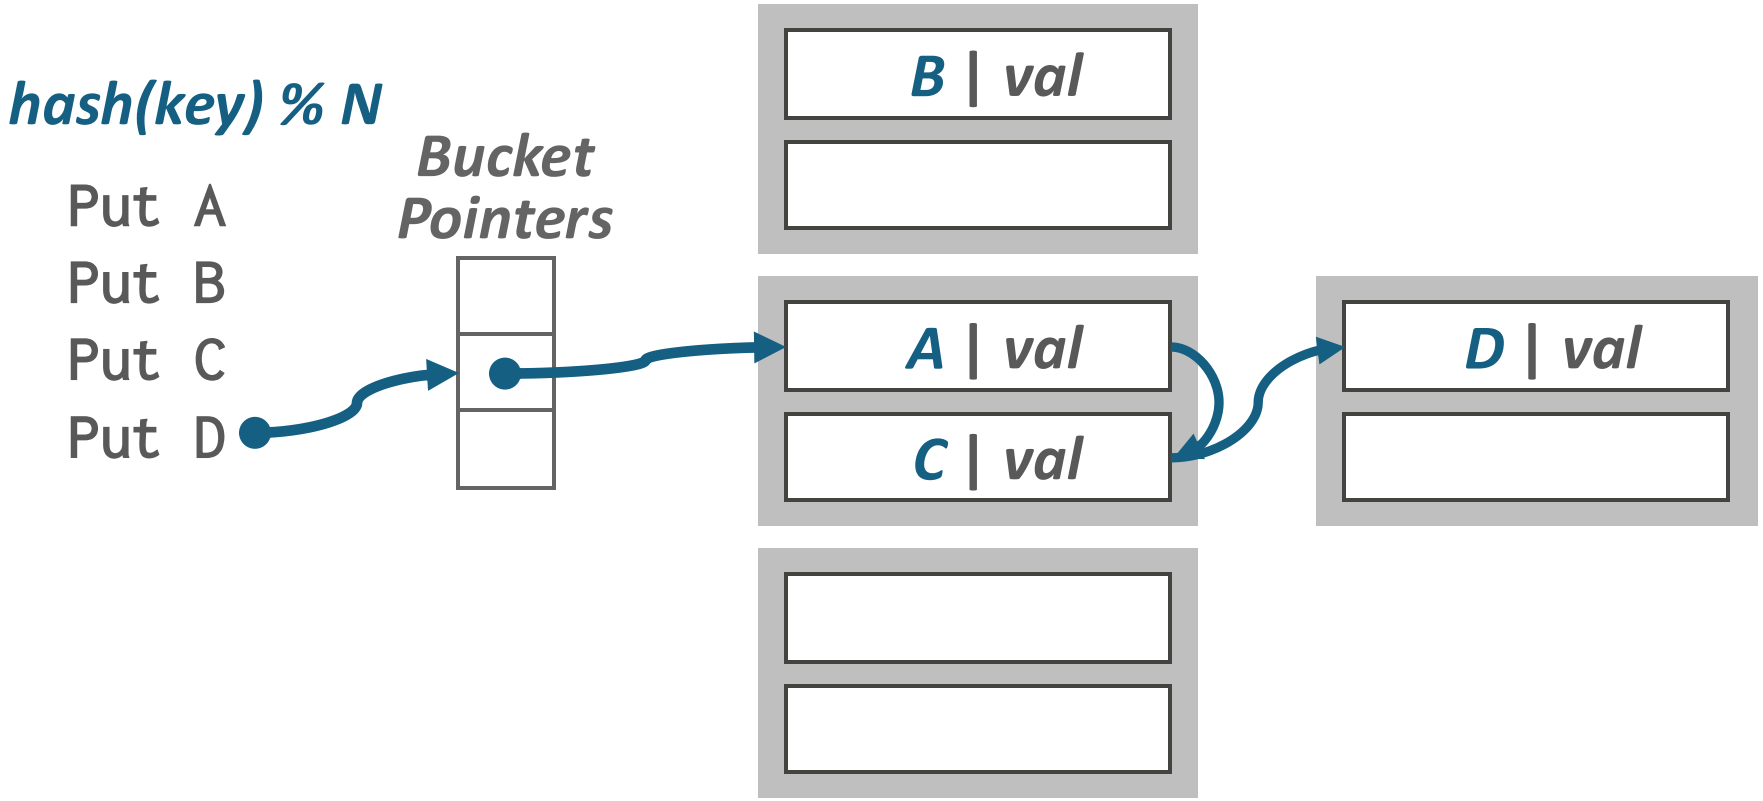

Chained Hashing

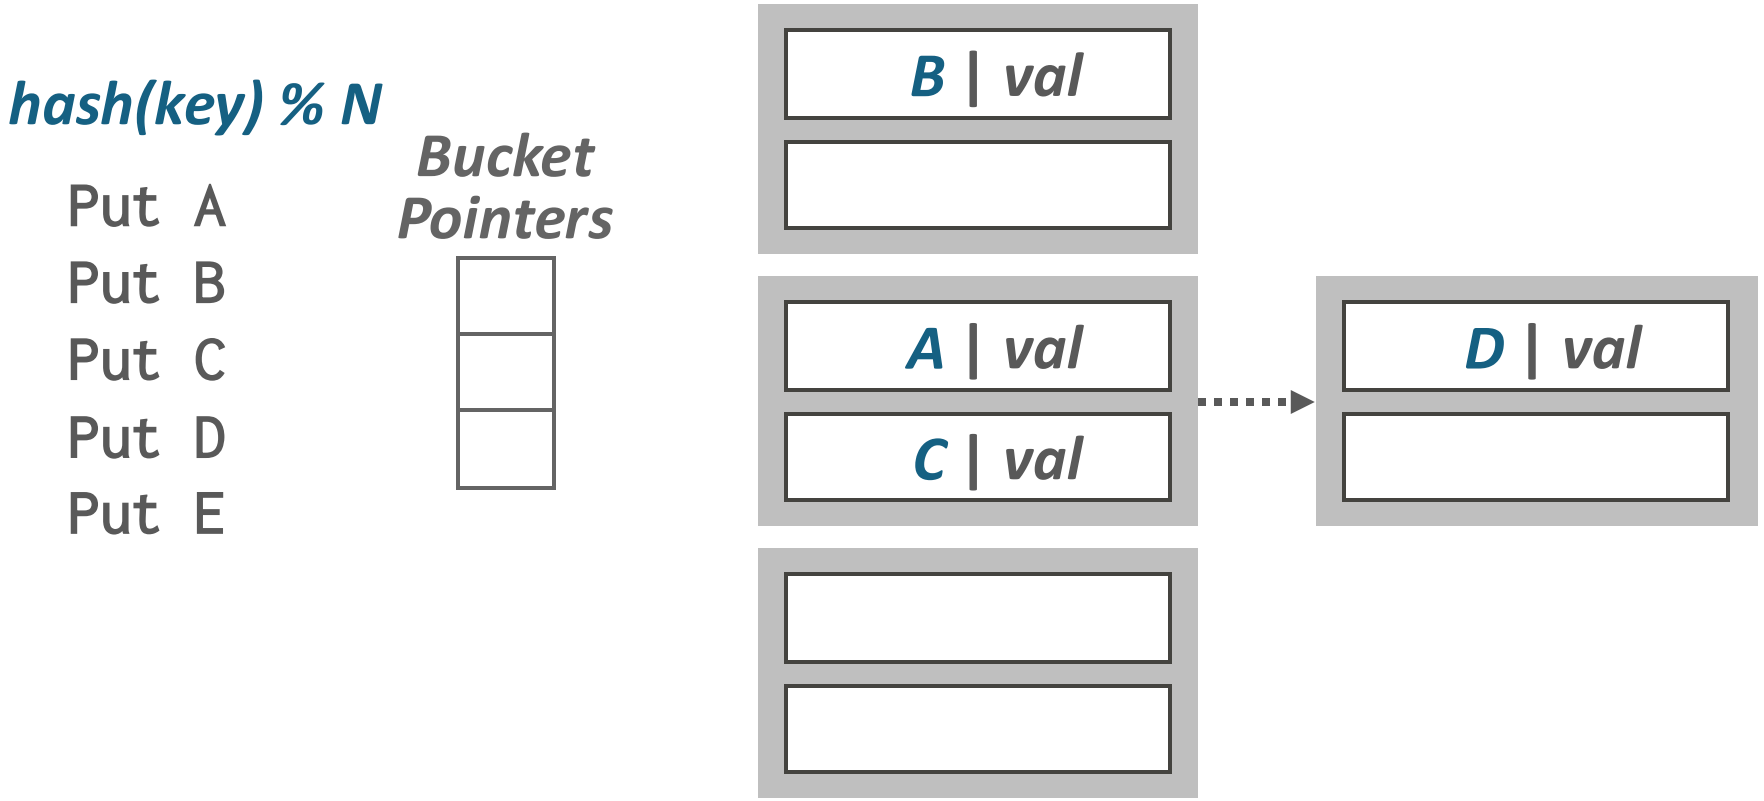

Chained Hashing

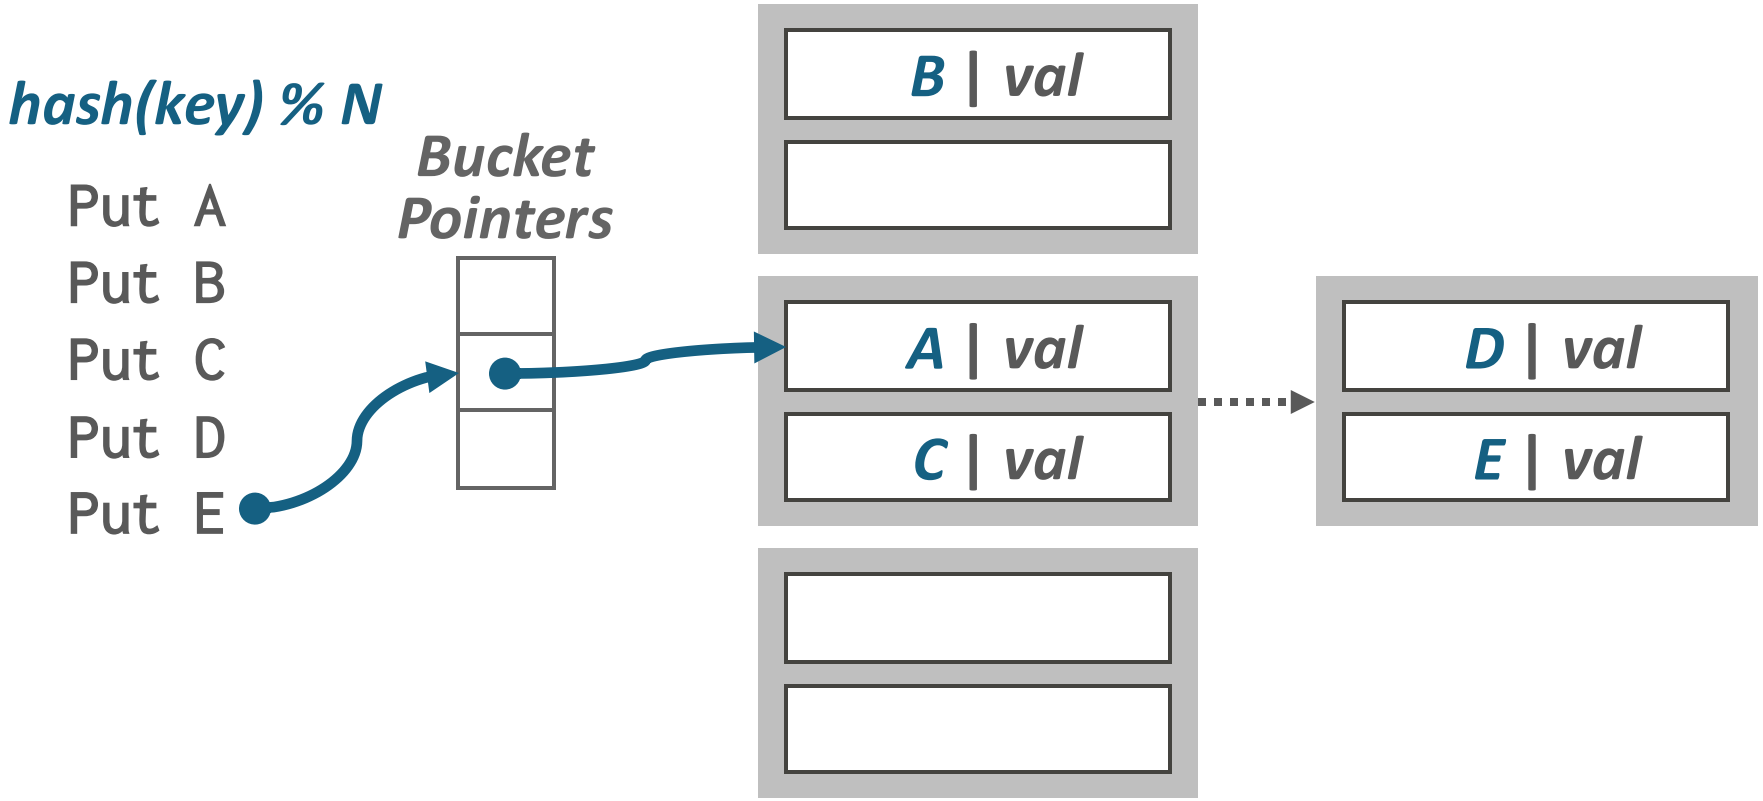

Chained Hashing

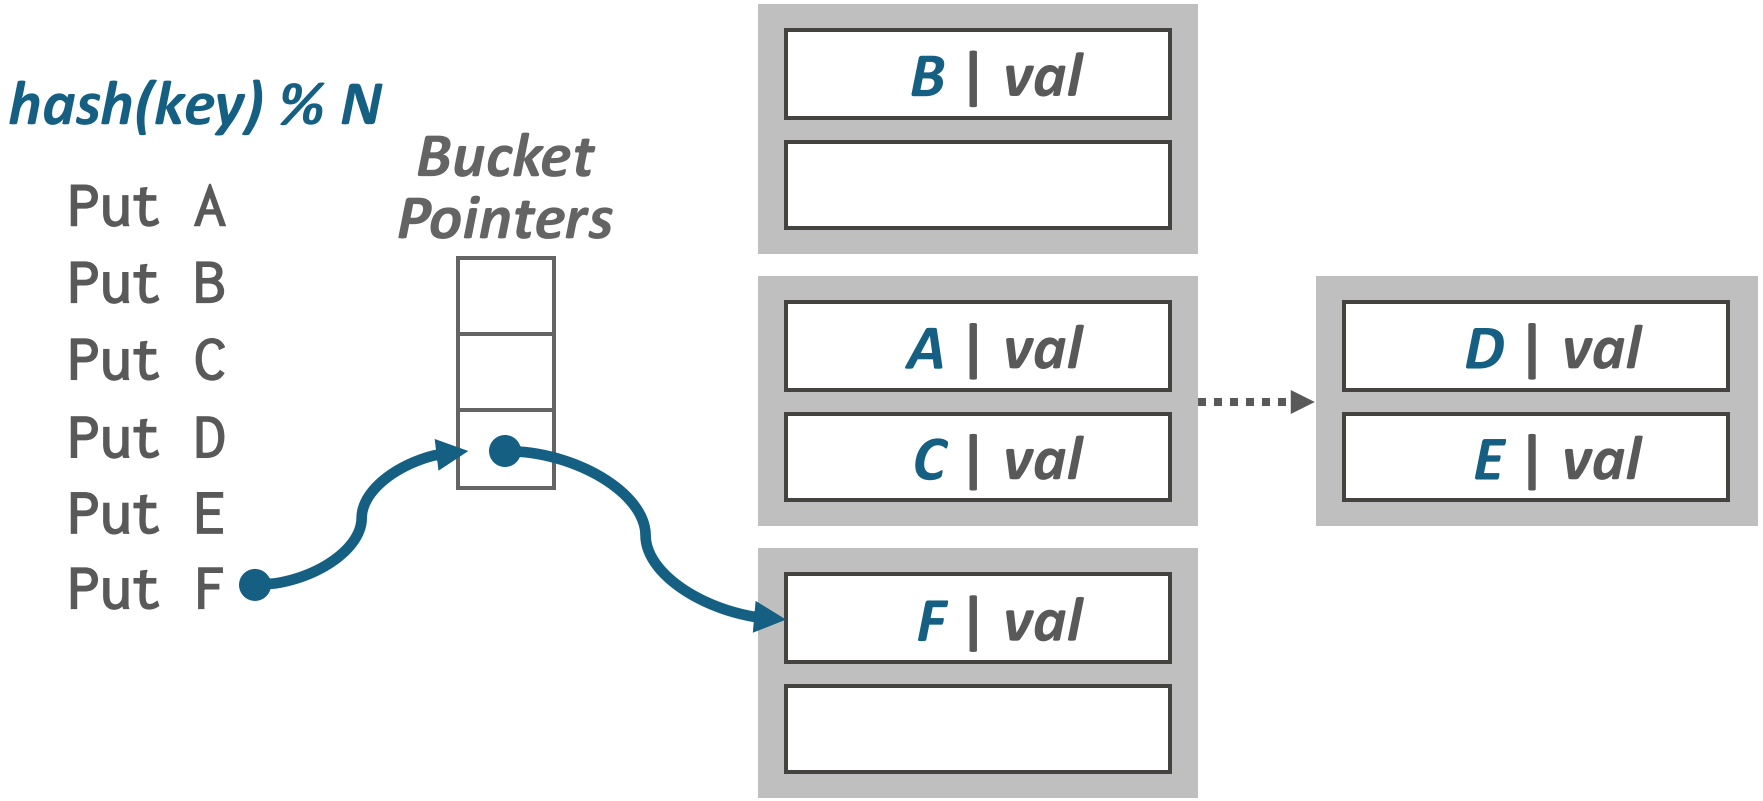

Chained Hashing

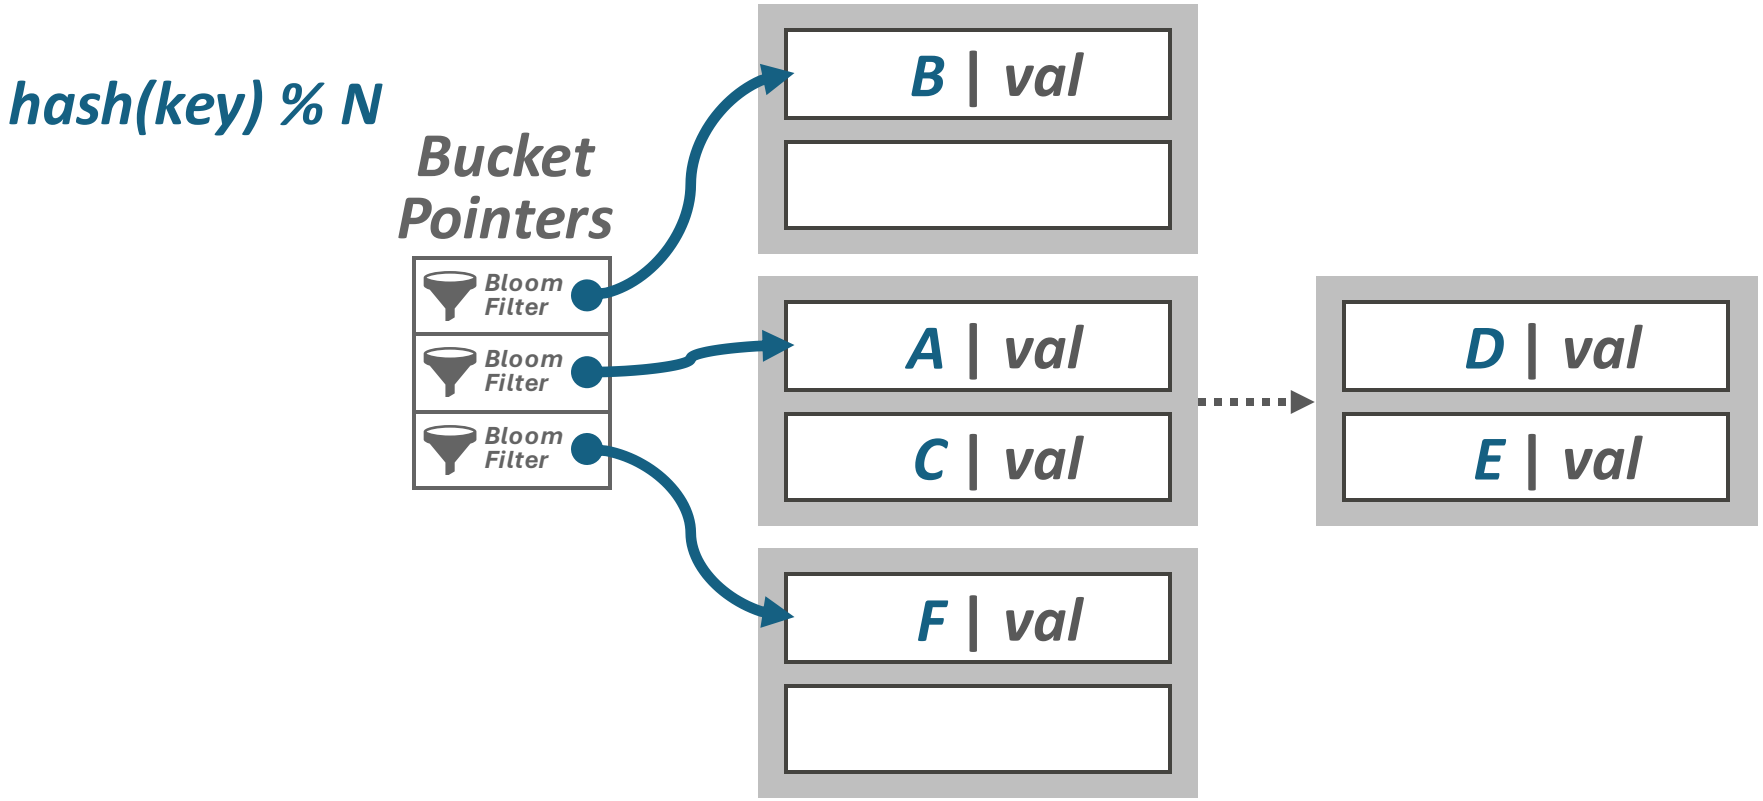

Chained Hashing

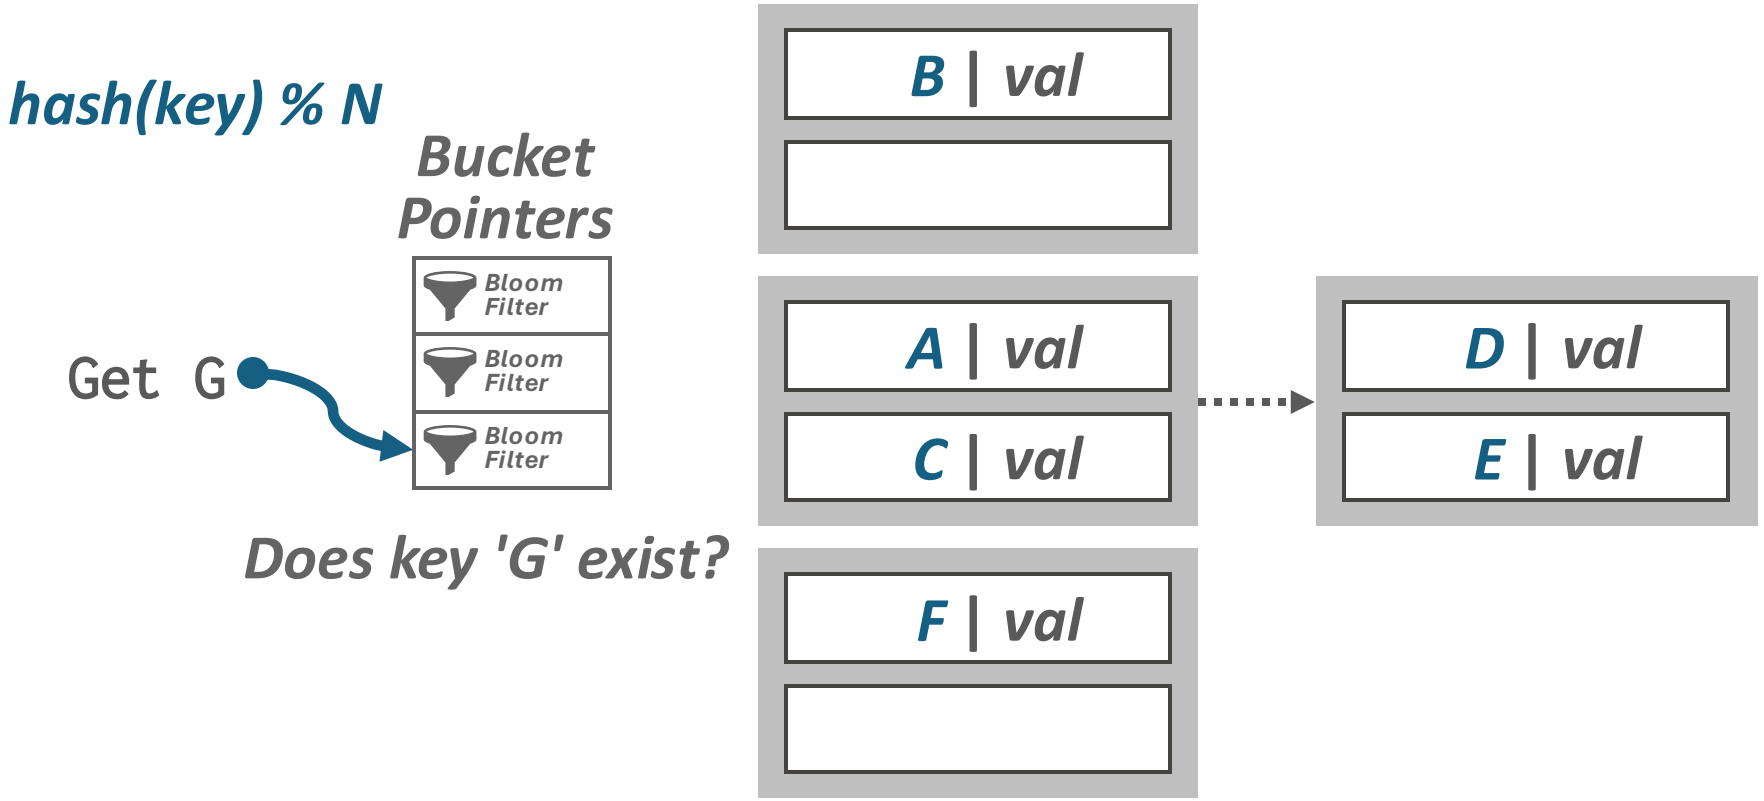

Chained Hashing

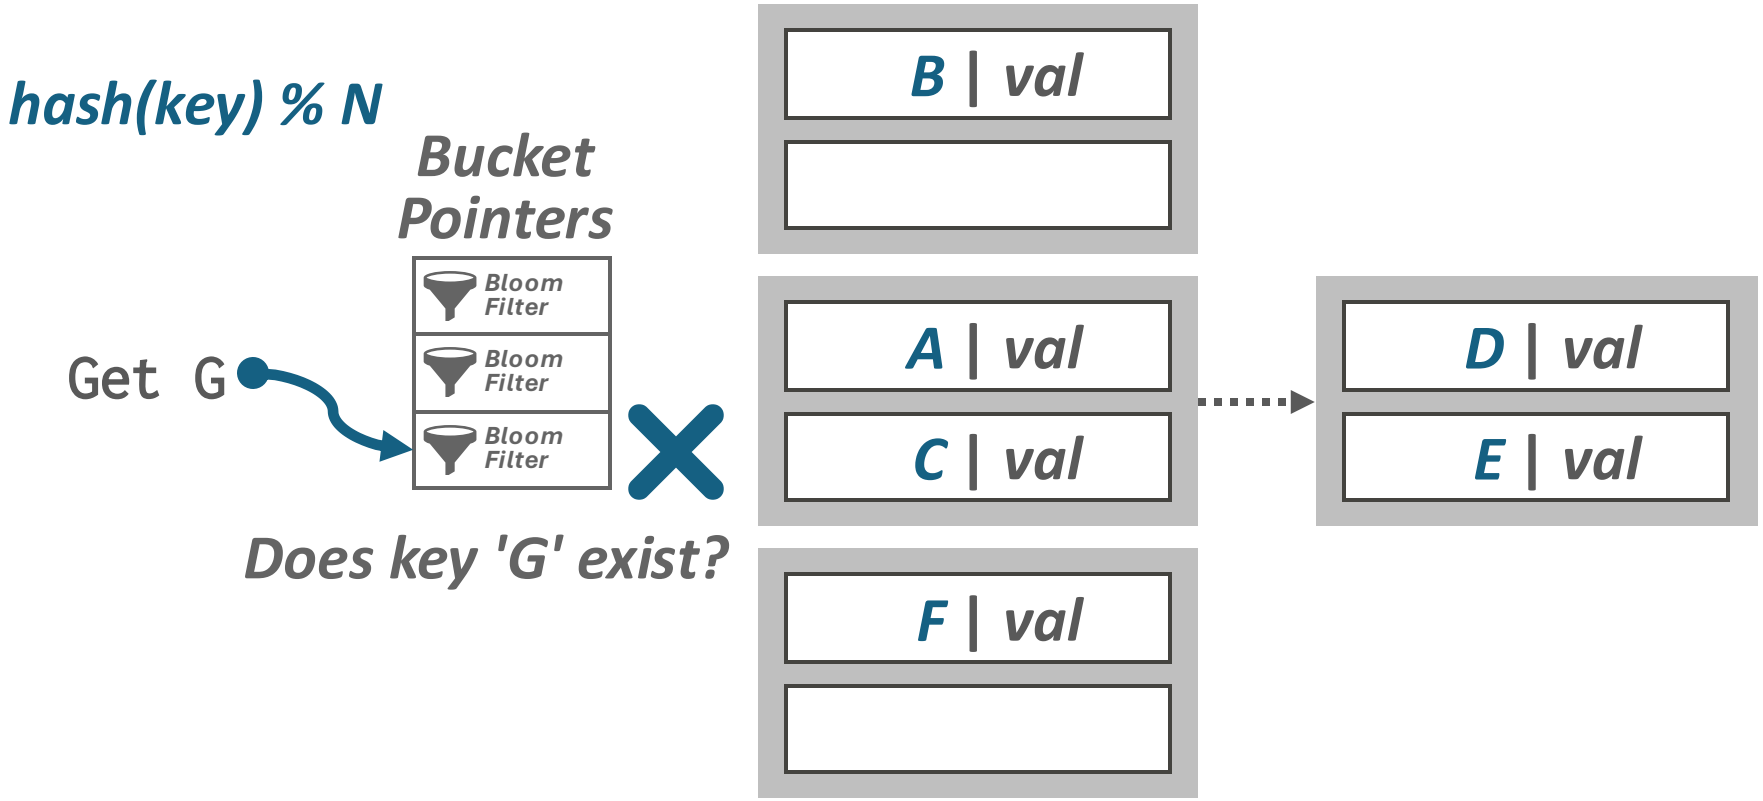

Bloom Filters

- Probabilistic data structure (bitmap) that answers set membership queries.
  - **False negatives** will **never occur**.
  - **False positives** can **sometimes occur**.
  - See Bloom Filter Calculator.
- **Insert(x):**
  - Use *k* hash functions to set bits in the filter to 1.
- **Lookup(x):**
  - Check whether the bits are 1 for each hash function.

Bloom Filters

Bloom Filter

0	1	2	3	4	5	6	7
0	0	0	0	0	0	0	0

Bloom Filters

Bloom Filter

0	1	2	3	4	5	6	7
0	0	0	0	0	0	0	0

- Insert('RZA')

Bloom Filters

Bloom Filter

0	1	2	3	4	5	6	7
0	0	0	0	0	0	0	0

$hash_1('RZA') = 2222 \% 8 = 6$

$hash_2('RZA') = 4444 \% 8 = 4$

- Insert('RZA')

Bloom Filters

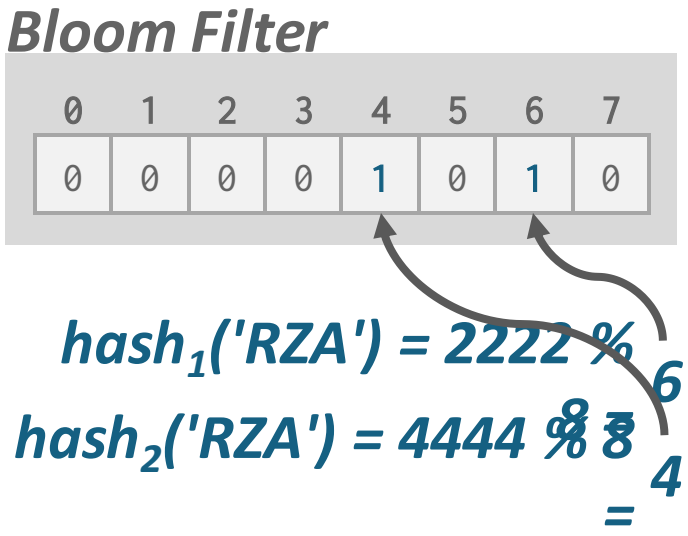

- Insert('RZA')

Bloom Filters

Bloom Filter

$hash_1('GZA') = 5555 \% 8 = 3$

$hash_2('GZA') = 7777 \% 8 = 1$

- Insert('RZA')
- Insert('GZA')

Bloom Filters

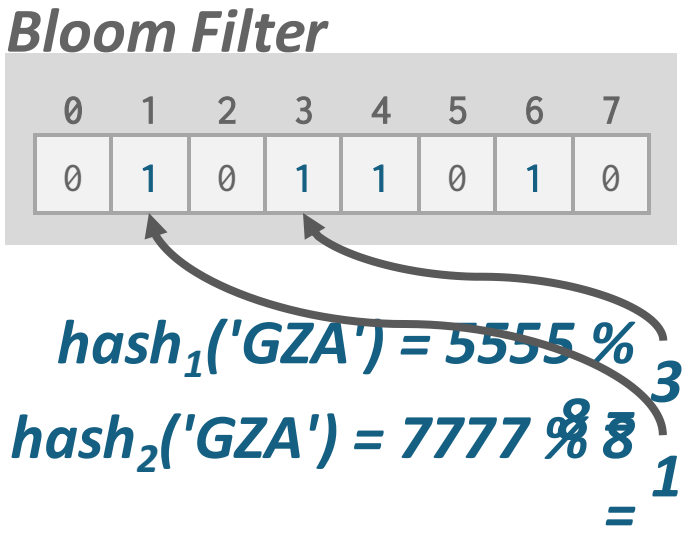

- Insert('RZA')
- Insert('GZA')

Bloom Filters

Bloom Filter

0	1	2	3	4	5	6	7
0	1	0	1	1	0	1	0

- Insert('RZA')
- Insert('GZA')
- Lookup('RZA')

Bloom Filters

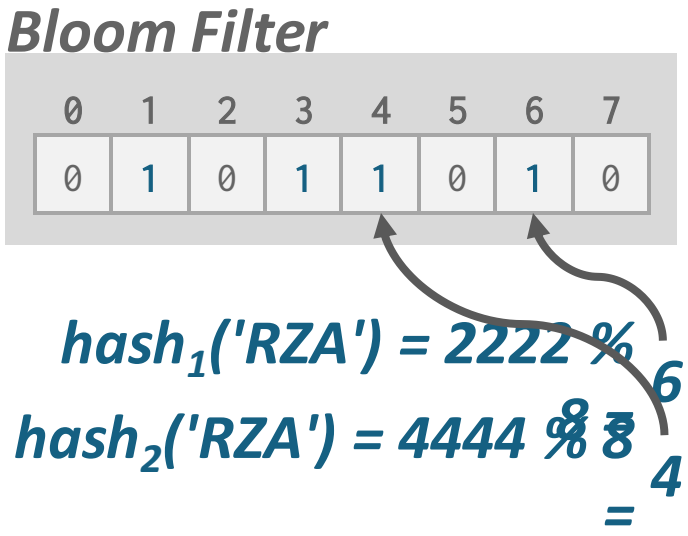

- Insert('RZA')
- Insert('GZA')
- Lookup('RZA')

Bloom Filters

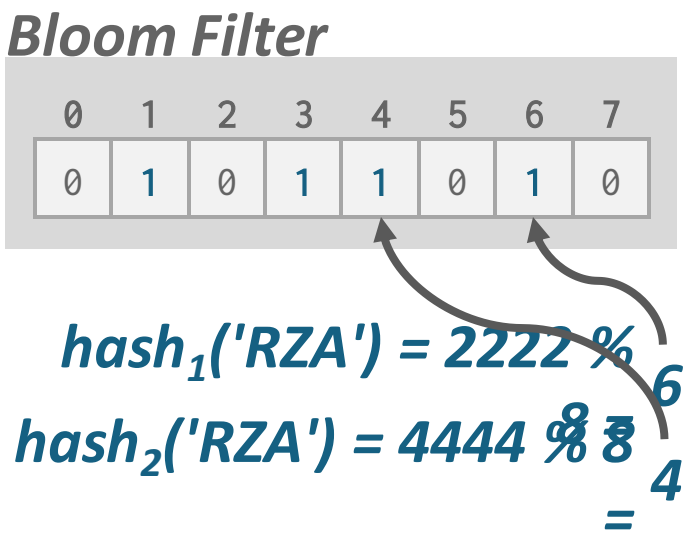

- Insert('RZA')
- Insert('GZA')
- Lookup('RZA') → *TRUE*

Bloom Filters

Bloom Filter

- Insert('RZA')
- Insert('GZA')
- Lookup('RZA') → ***TRUE***
- Lookup('Raekwon')

Bloom Filters

Bloom Filter

0	1	2	3	4	5	6	7
0	1	0	1	1	0	1	0

$hash_1$('Raekwon') = 3333 % 8 = 5

$hash_2$('Raekwon') = 8899 % 8 = 3

- Insert('RZA')
- Insert('GZA')
- Lookup('RZA') → *TRUE*
- Lookup('Raekwon')

Bloom Filters

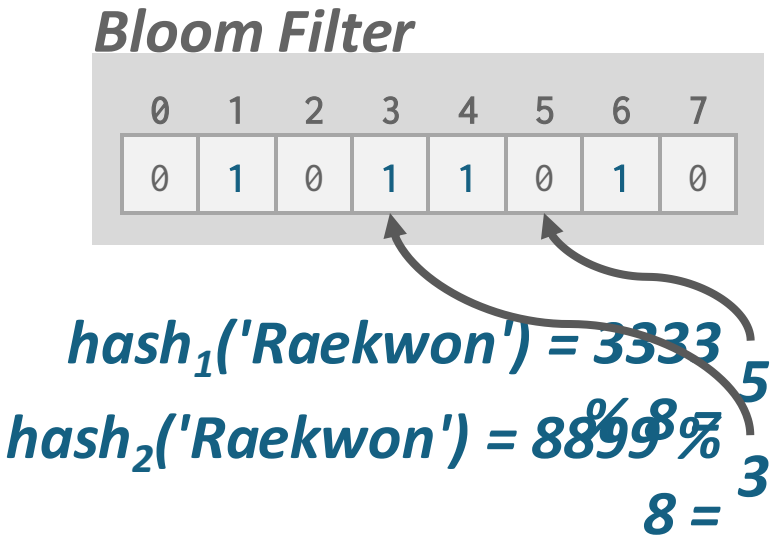

- Insert('RZA')
- Insert('GZA')
- Lookup('RZA') → *TRUE*
- Lookup('Raekwon')

Bloom Filters

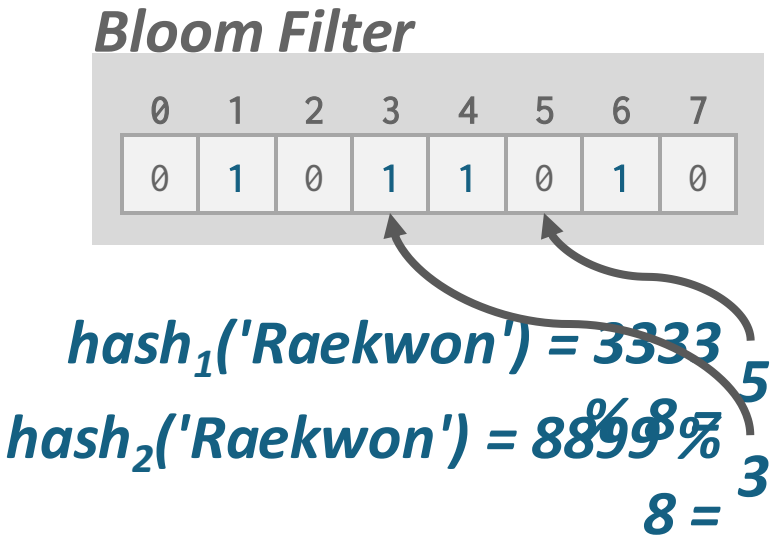

- Insert('RZA')
- Insert('GZA')
- Lookup('RZA') → *TRUE*
- Lookup('Raekwon') → *FALSE*

Bloom Filters

Bloom Filter

- Insert('RZA')
- Insert('GZA')
- Lookup('RZA') → ***TRUE***
- Lookup('Raekwon') → ***FALSE***

Bloom Filters

Bloom Filter

$hash_1$('ODB') = 6699 % 8 = 3

$hash_2$('ODB') = 9966 % 8 = 6

- Insert('RZA')
- Insert('GZA')
- Lookup('RZA') → *TRUE*
- Lookup('Raekwon') → *FALSE*
- Lookup('ODB')

Bloom Filters

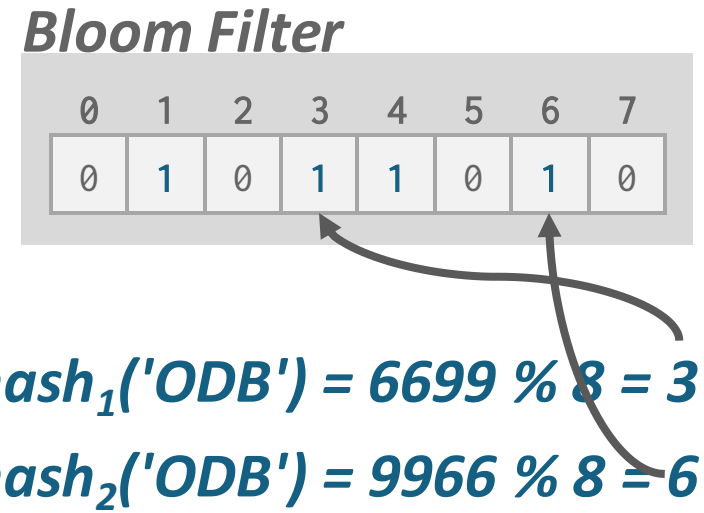

$hash_1('ODB') = 6699 \% 8 = 3$

$hash_2('ODB') = 9966 \% 8 = 6$

- Insert('RZA')
- Insert('GZA')
- Lookup('RZA') → *TRUE*
- Lookup('Raekwon') → *FALSE*
- Lookup('ODB')

Bloom Filters

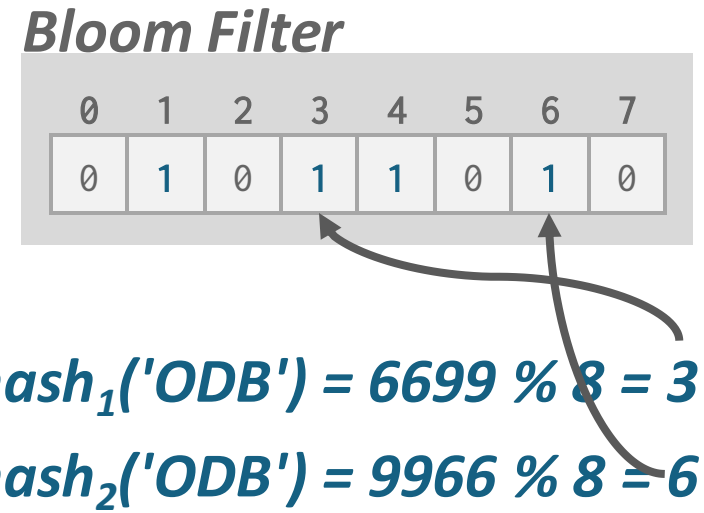

$hash_1('ODB') = 6699 \% 8 = 3$

$hash_2('ODB') = 9966 \% 8 = 6$

- Insert('RZA')
- Insert('GZA')
- Lookup('RZA') → *TRUE*
- Lookup('Raekwon') → *FALSE*
- Lookup('ODB') → *TRUE (false positive)*

Extendible Hashing

Dynamic Hashing Schemes

Extendible Hashing

- Extendible-hashing approach that splits buckets incrementally instead of letting the linked list grow forever.
- Multiple slot locations can point to the same bucket chain.
- Reshuffle bucket entries on split and increase the number of bits to examine.
  - Data movement is localized to just the split chain.

Extendible Hashing

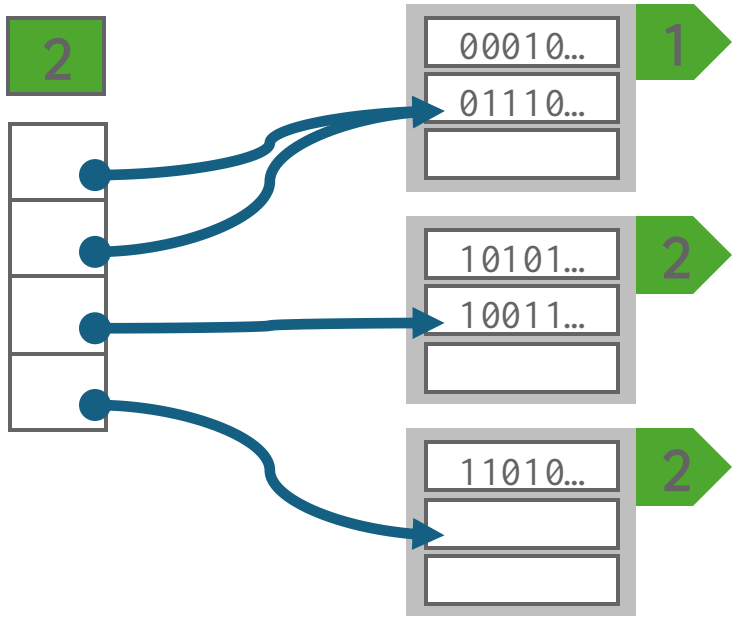

Extendible Hashing

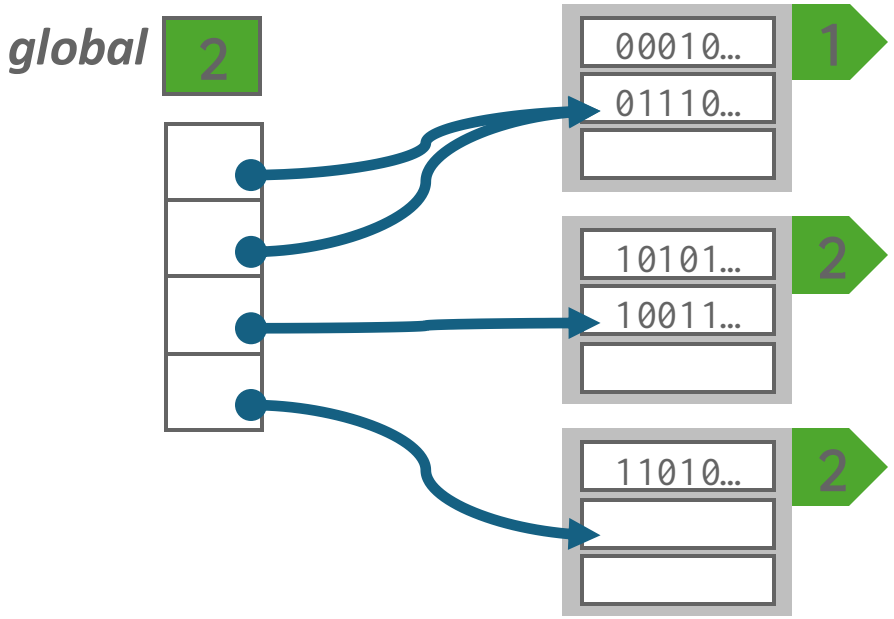

Extendible Hashing

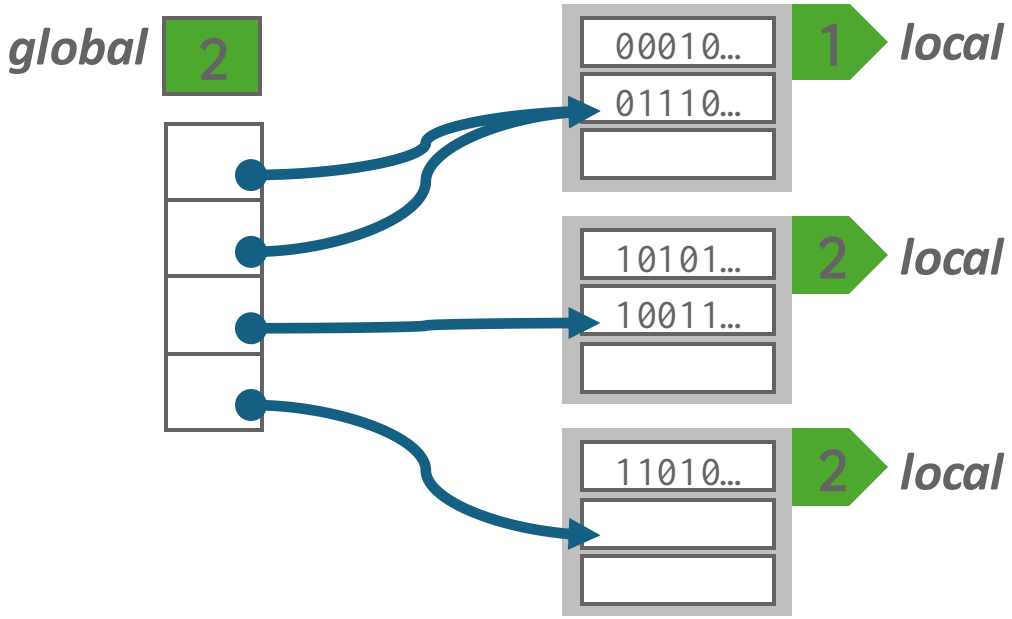

Extendible Hashing

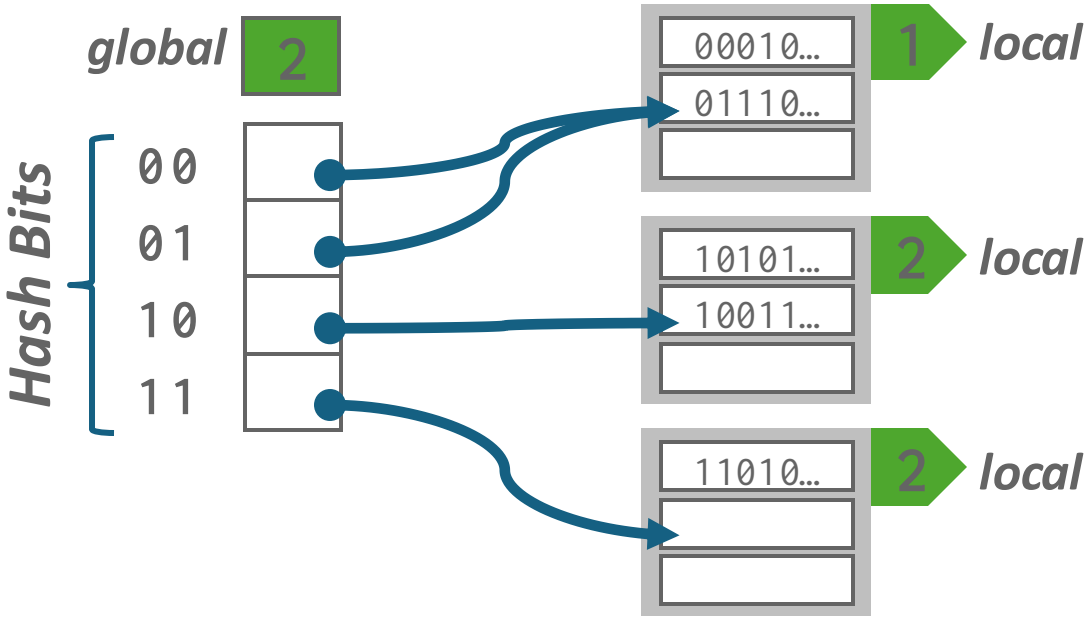

Extendible Hashing

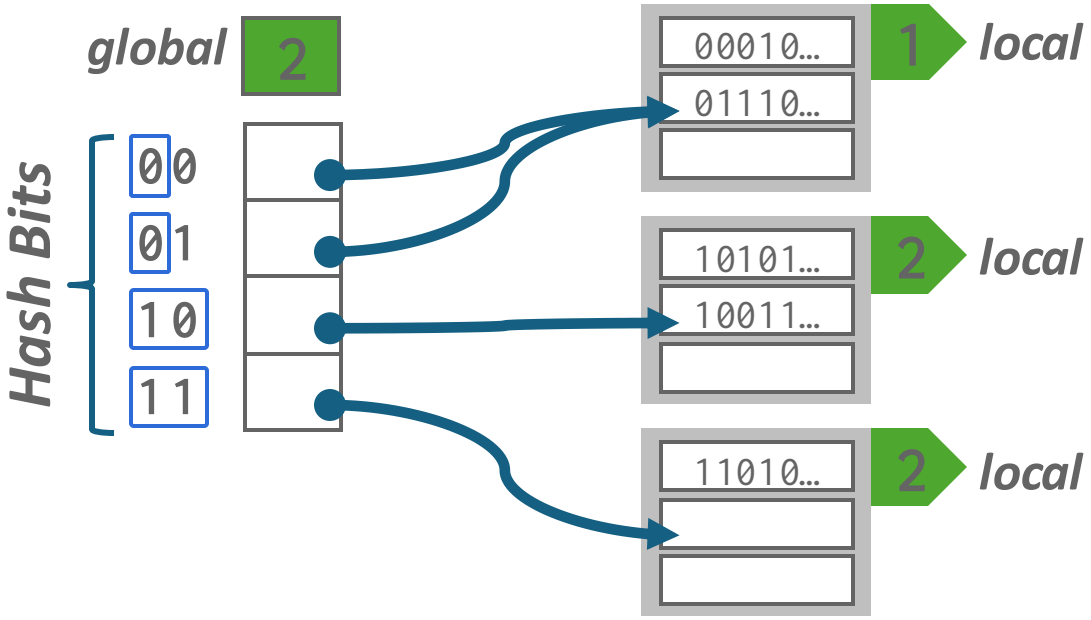

Extendible Hashing

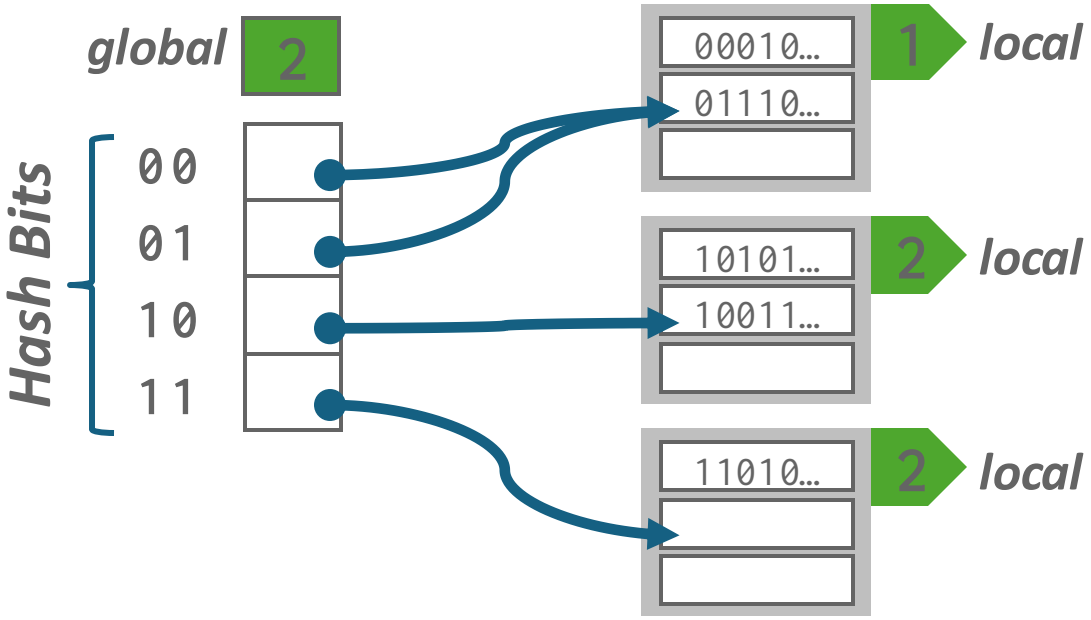

Get A
hash(A) = 01110...

Extendible Hashing

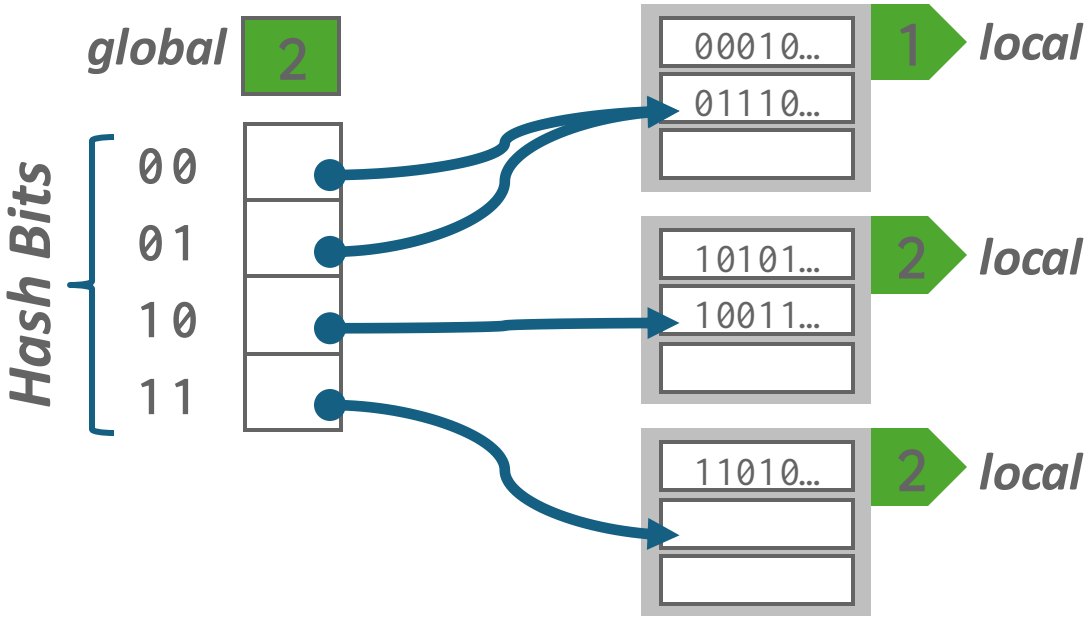

Get A
hash(A) = 01110...

Extendible Hashing

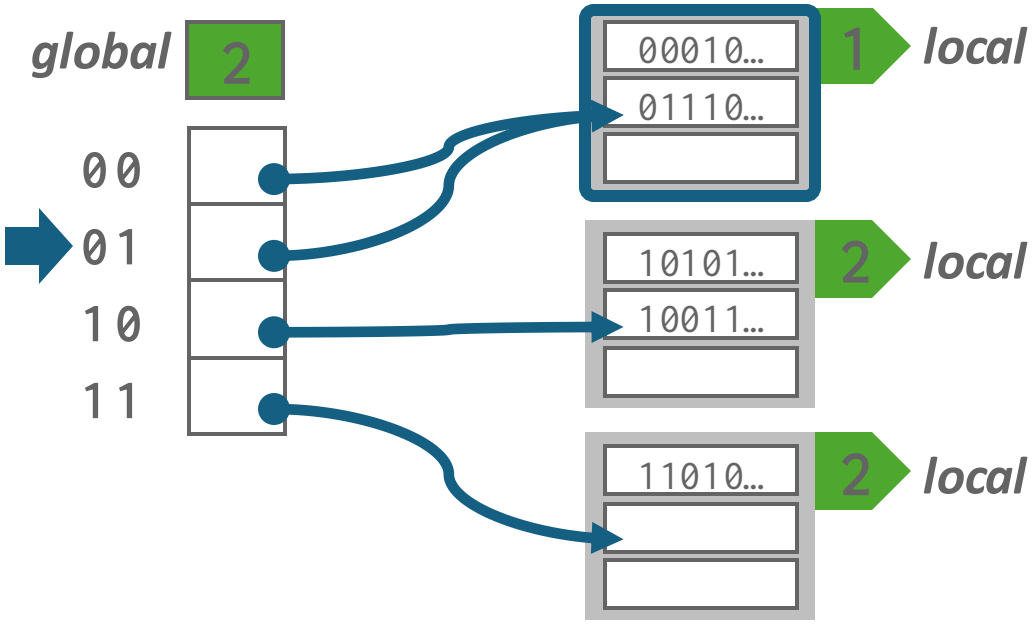

Get A
hash(A) = 01110...

Extendible Hashing

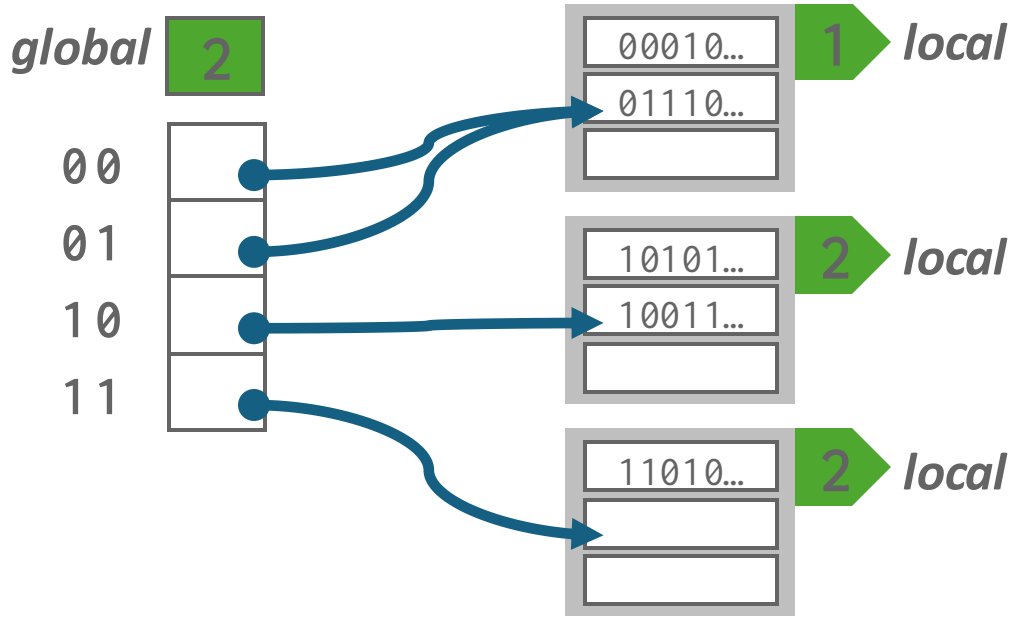

Get A
hash(A) = 01110...

Put B
hash(B) = 10111...

Extendible Hashing

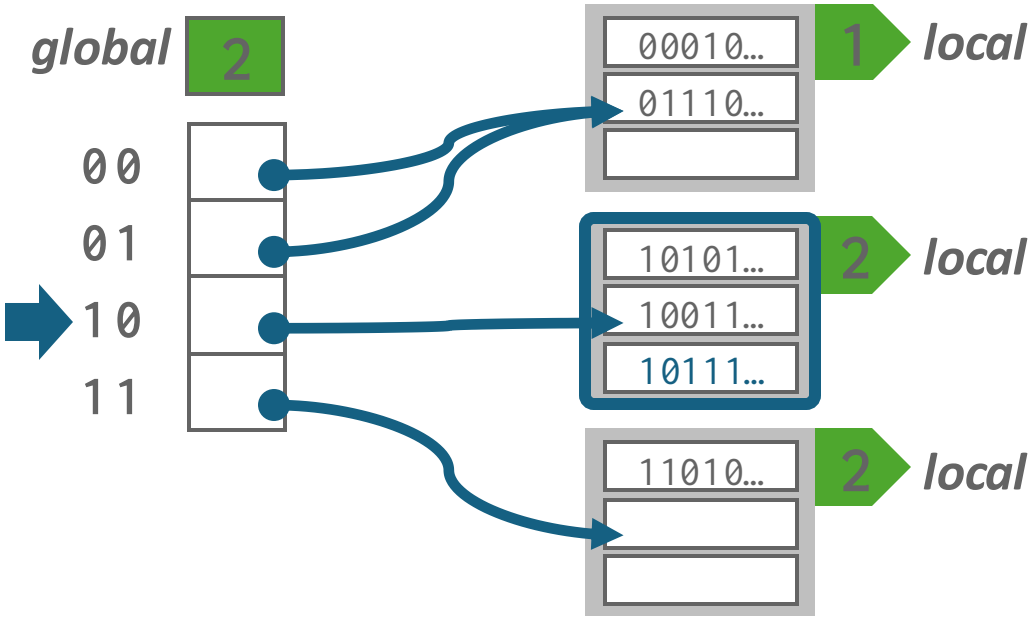

Get A
hash(A) = 01110...

Put B
hash(B) = 10111...

Extendible Hashing

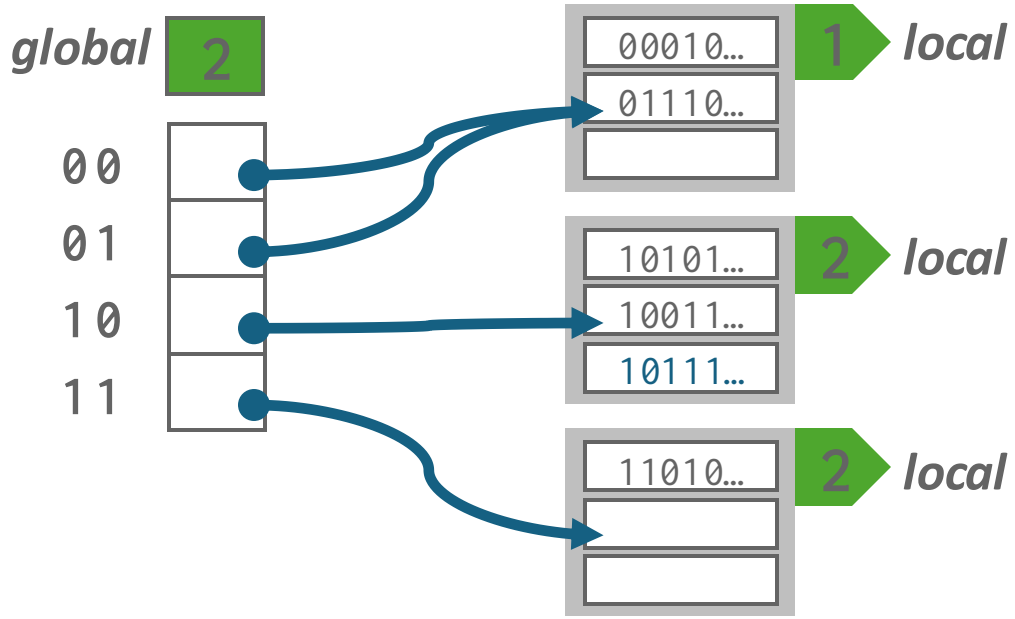

Get A
hash(A) = 01110...

Put B
hash(B) = 10111...

Put C
hash(C) = 10100...

Extendible Hashing

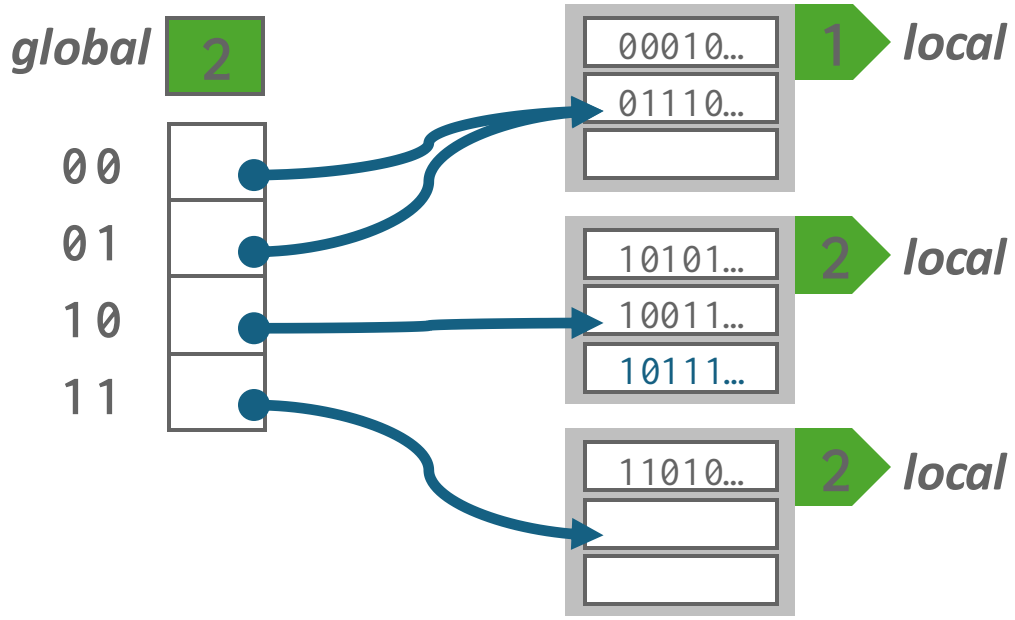

Get A
hash(A) = 01110...

Put B
hash(B) = 10111...

Put C
hash(C) = 10100...

Extendible Hashing

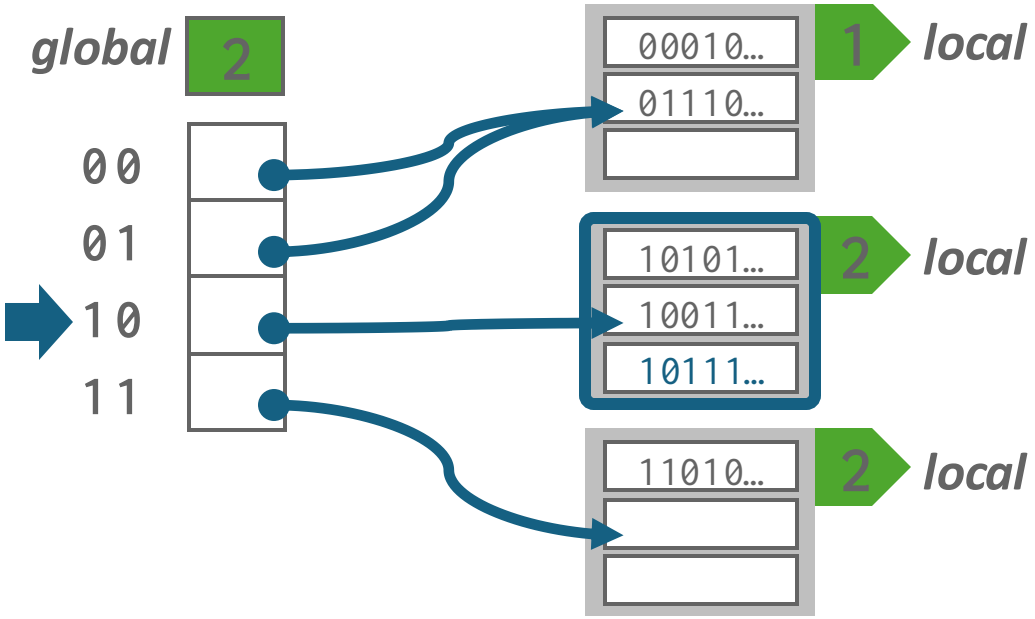

Get A
hash(A) = 01110...

Put B
hash(B) = 10111...

Put C
hash(C) = 10100...

Extendible Hashing

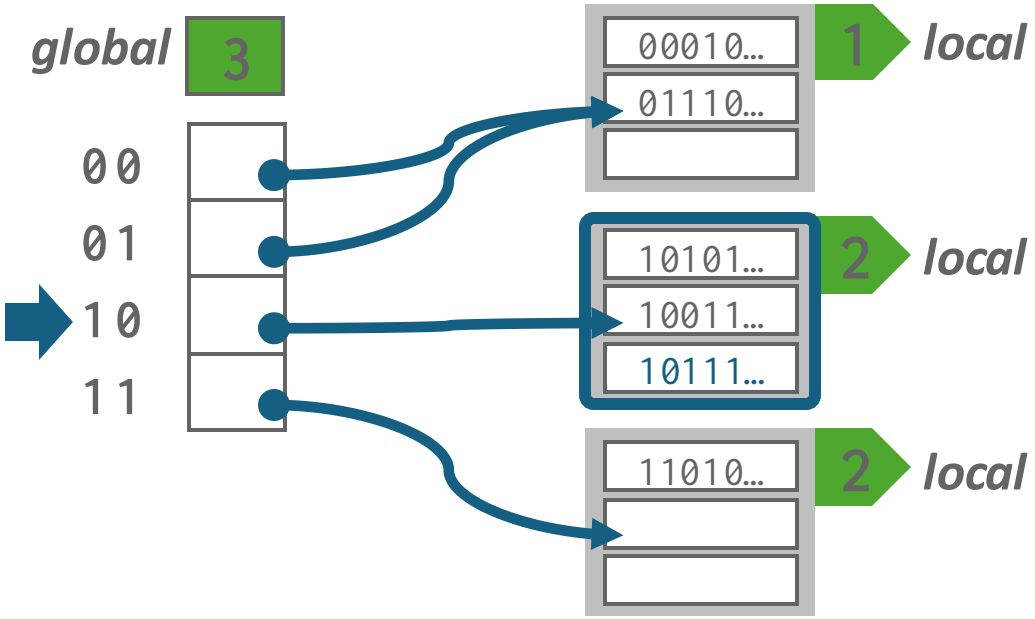

Get A
hash(A) = 01110...

Put B
hash(B) = 10111...

Put C
hash(C) = 10100...

Extendible Hashing

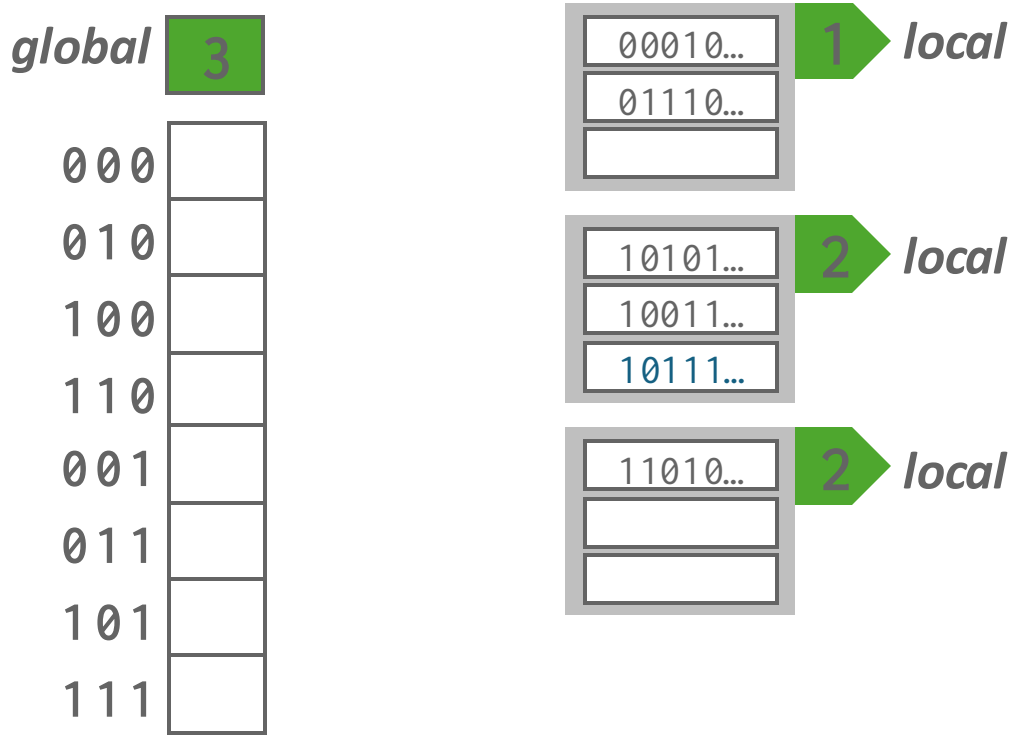

Get A
hash(A) = 01110...

Put B
hash(B) = 10111...

Put C
hash(C) = 10100...

Extendible Hashing

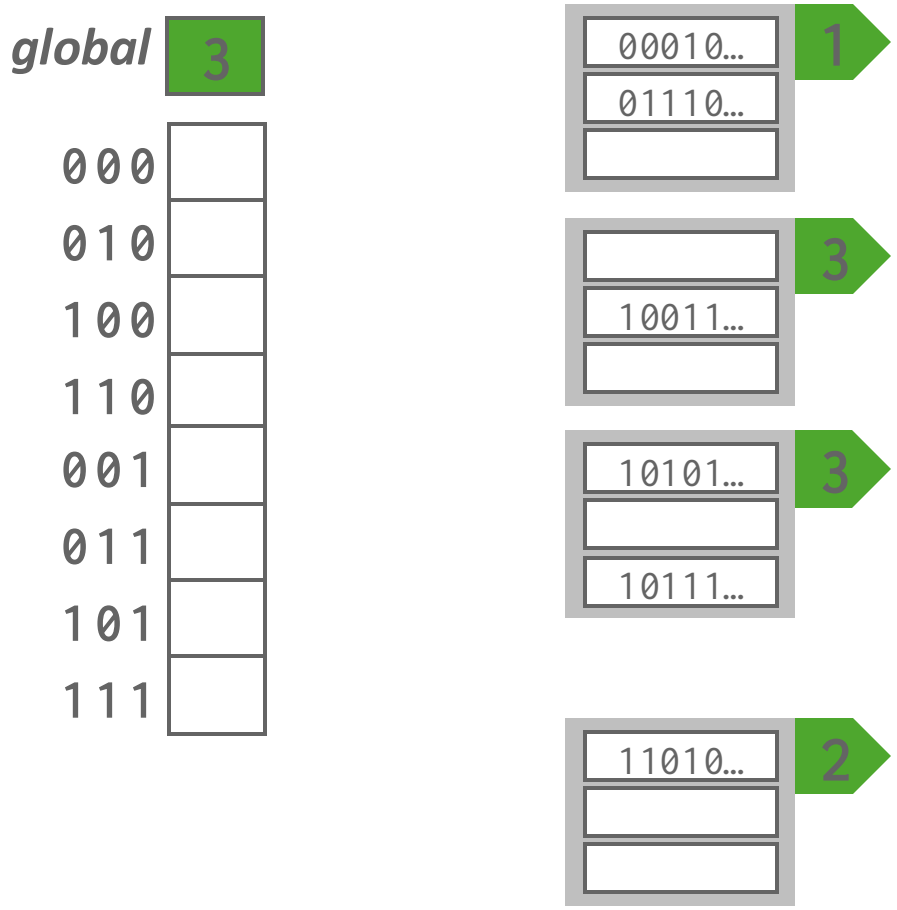

Get A
hash(A) = 01110...

Put B
hash(B) = 10111...

Put C
hash(C) = 10100...

Extendible Hashing

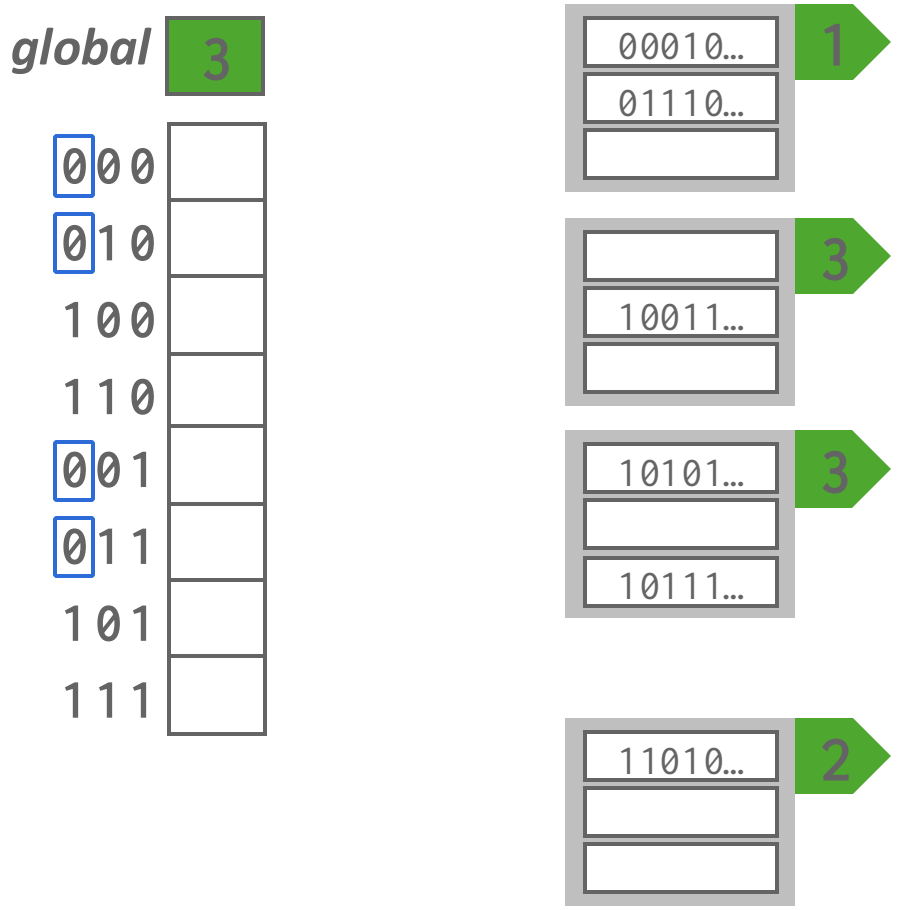

Get A
hash(A) = 01110...

Put B
hash(B) = 10111...

Put C
hash(C) = 10100...

Extendible Hashing

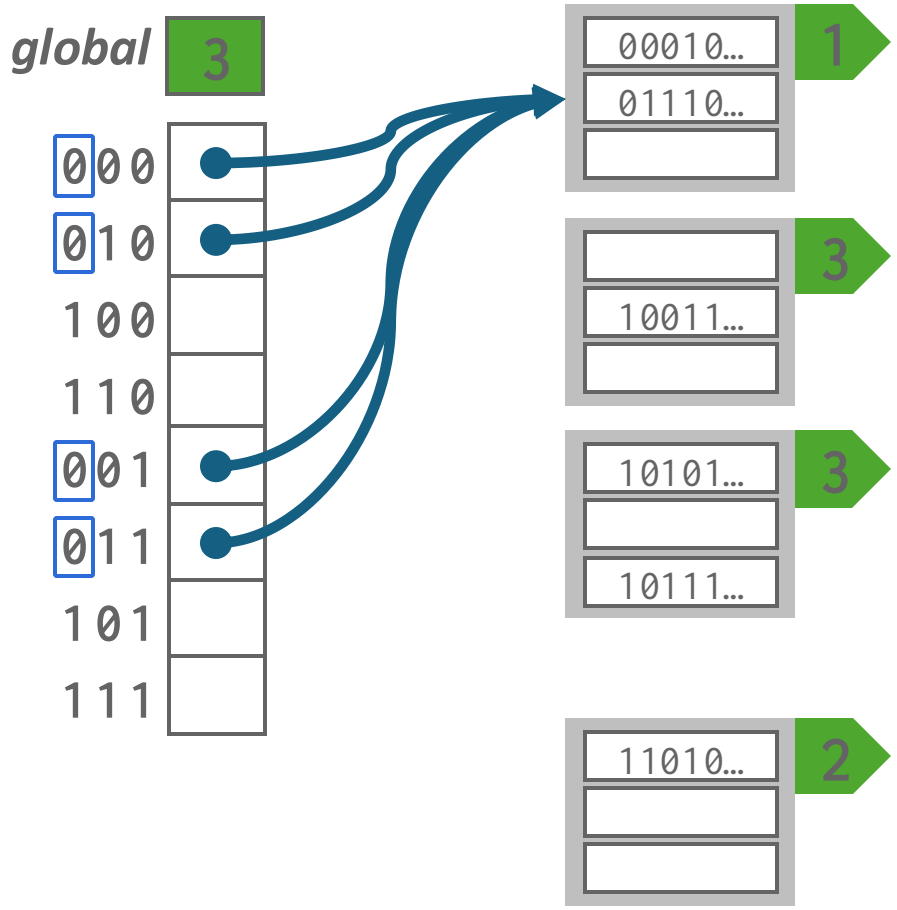

Get A
hash(A) = 01110...

Put B
hash(B) = 10111...

Put C
hash(C) = 10100...

Extendible Hashing

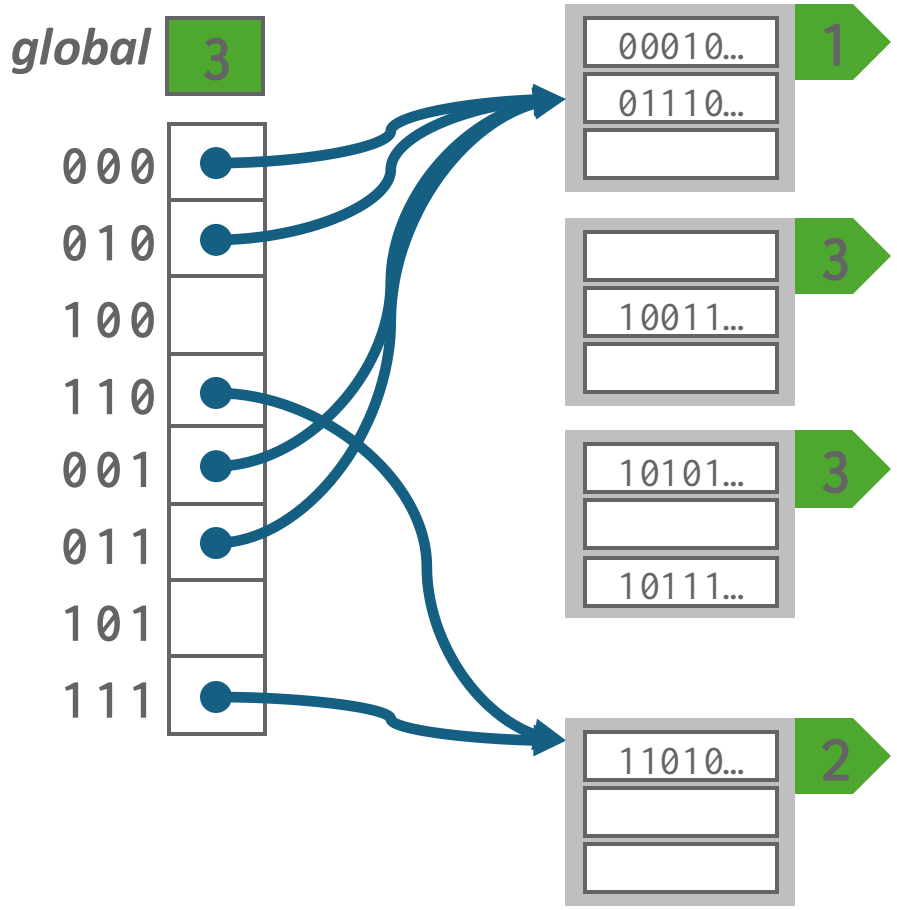

Get A
hash(A) = 01110...

Put B
hash(B) = 10111...

Put C
hash(C) = 10100...

Extendible Hashing

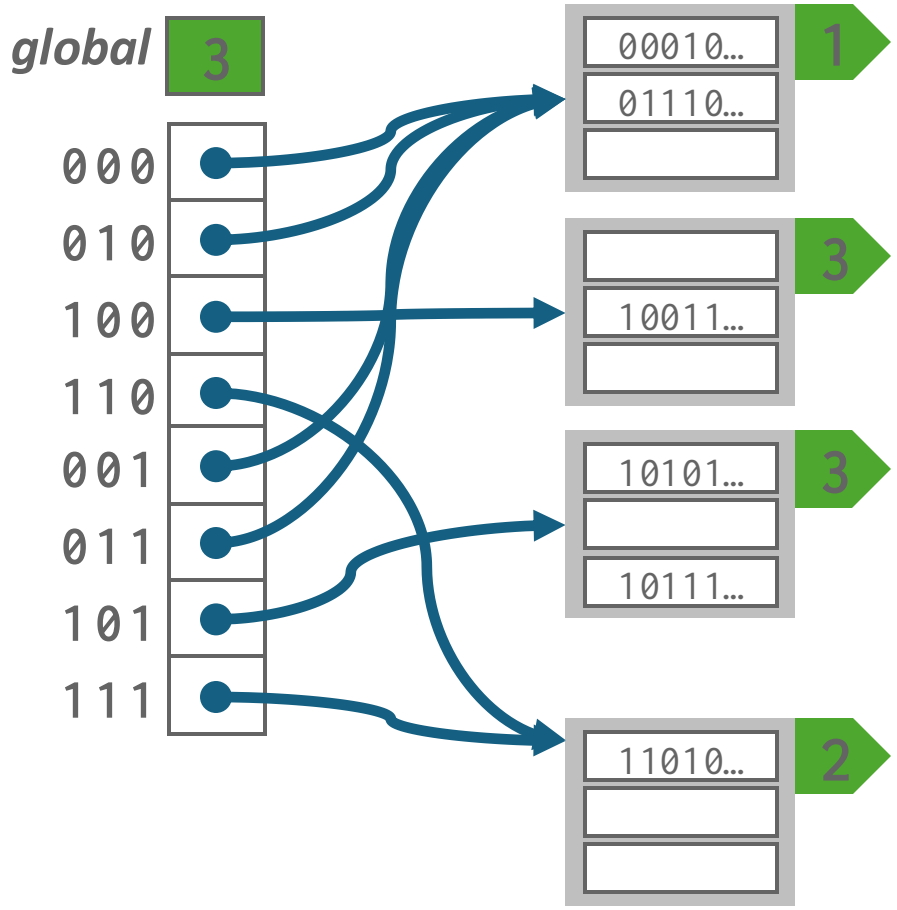

Get A
hash(A) = 01110...

Put B
hash(B) = 10111...

Put C
hash(C) = 10100...

Extendible Hashing

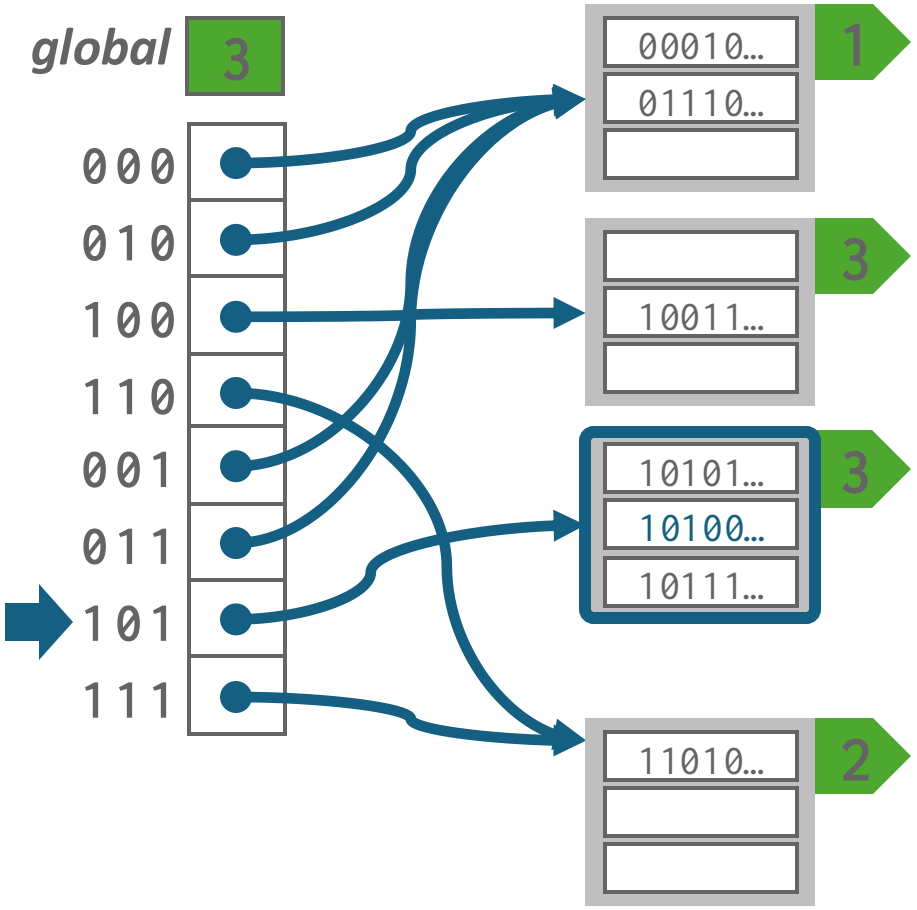

Get A
hash(A) = 01110...

Put B
hash(B) = 10111...

Put C
hash(C) = 10100...

Linear Hashing

Dynamic Hashing Schemes

Linear Hashing

- The hash table maintains a pointer that tracks the next bucket to split.
  - When any bucket overflows, split the bucket at the pointer location.
- Use multiple hashes to find the right bucket for a given key.
- Can use different overflow criterion:
  - Space Utilization
  - Average Length of Overflow Chains

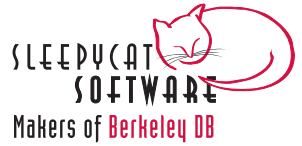

Linear Hashing

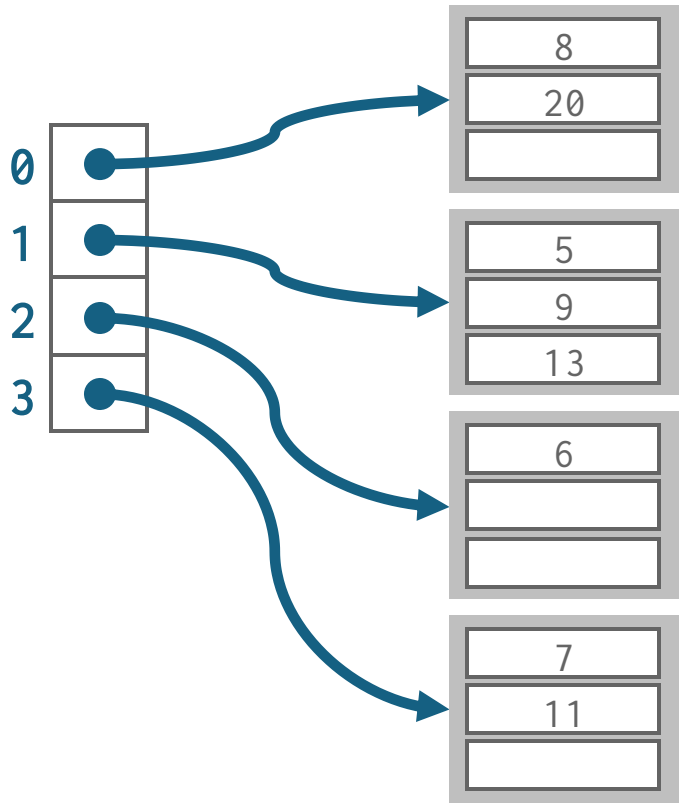

Linear Hashing

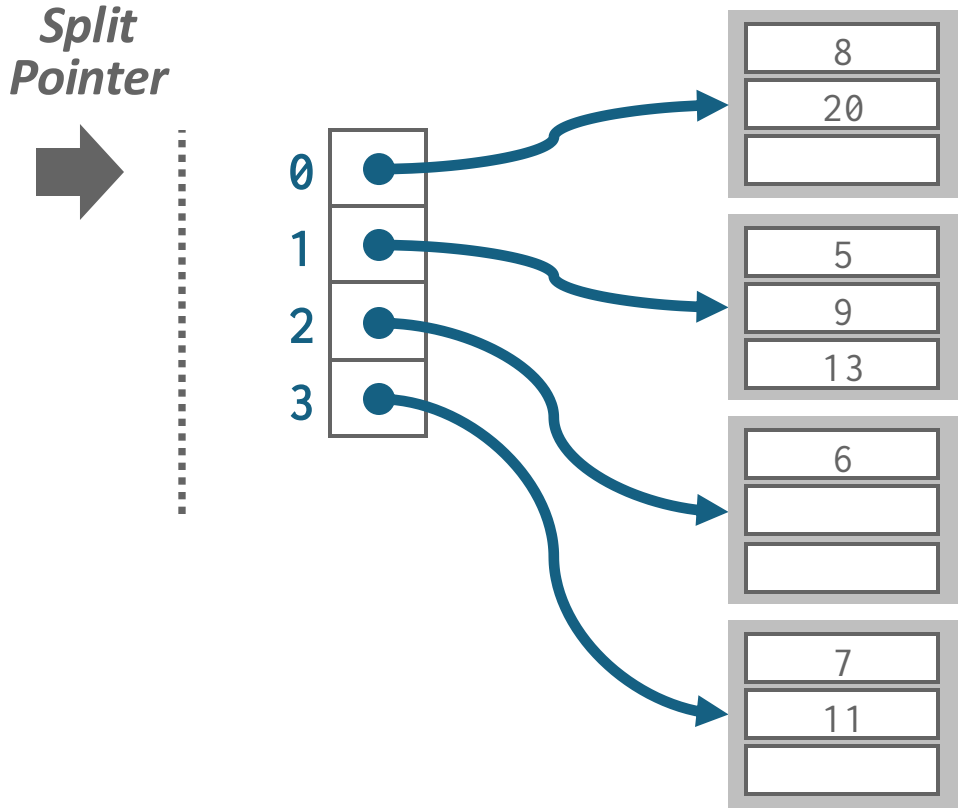

Linear Hashing

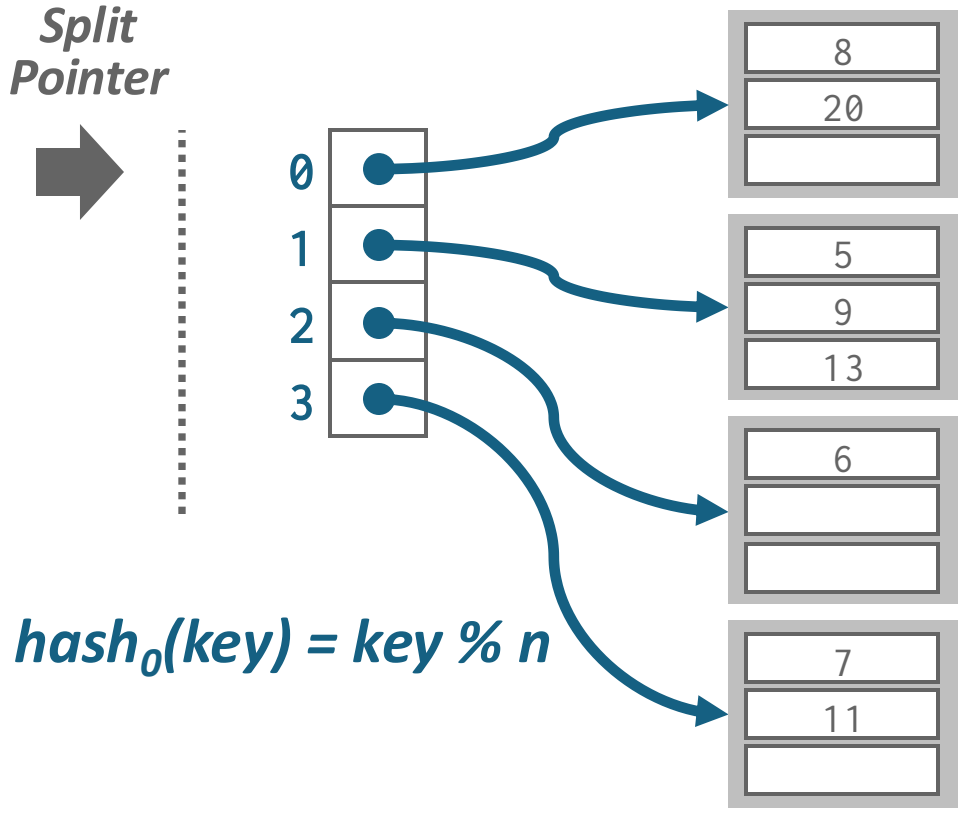

Linear Hashing

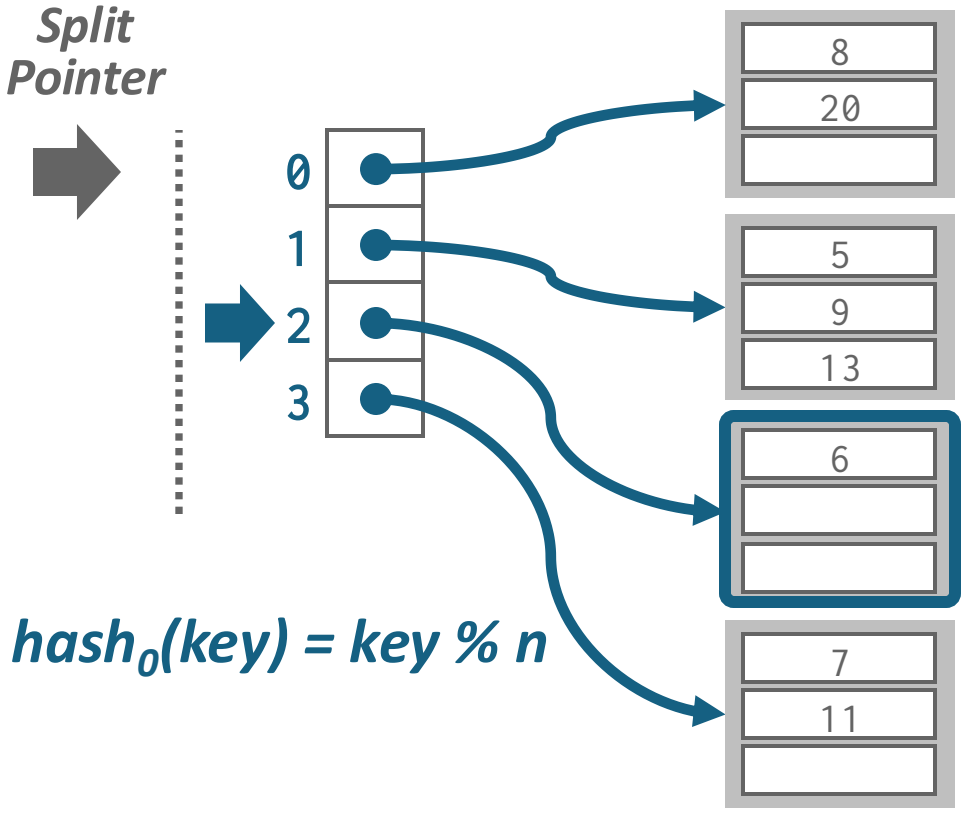

Get 6

$hash_0(6) = 6 \% 4 = 2$

Linear Hashing

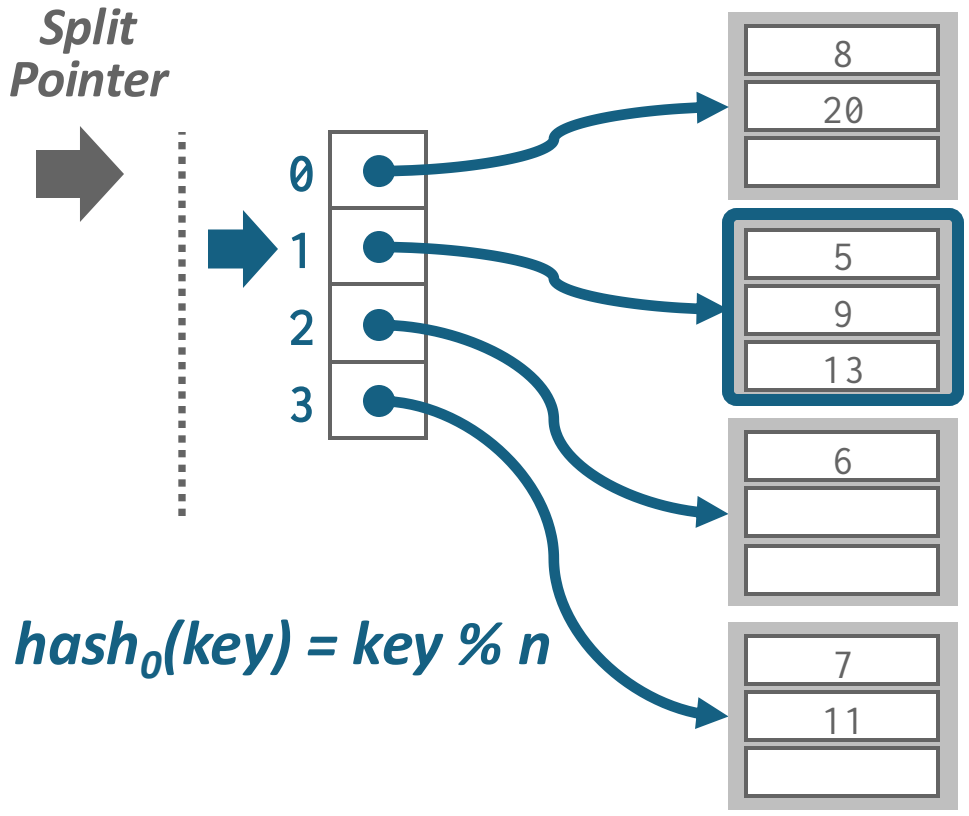

Get 6

$hash_0(6) = 6 \% 4 = 2$

Put 17

$hash_0(17) = 17 \% 4 = 1$

Linear Hashing

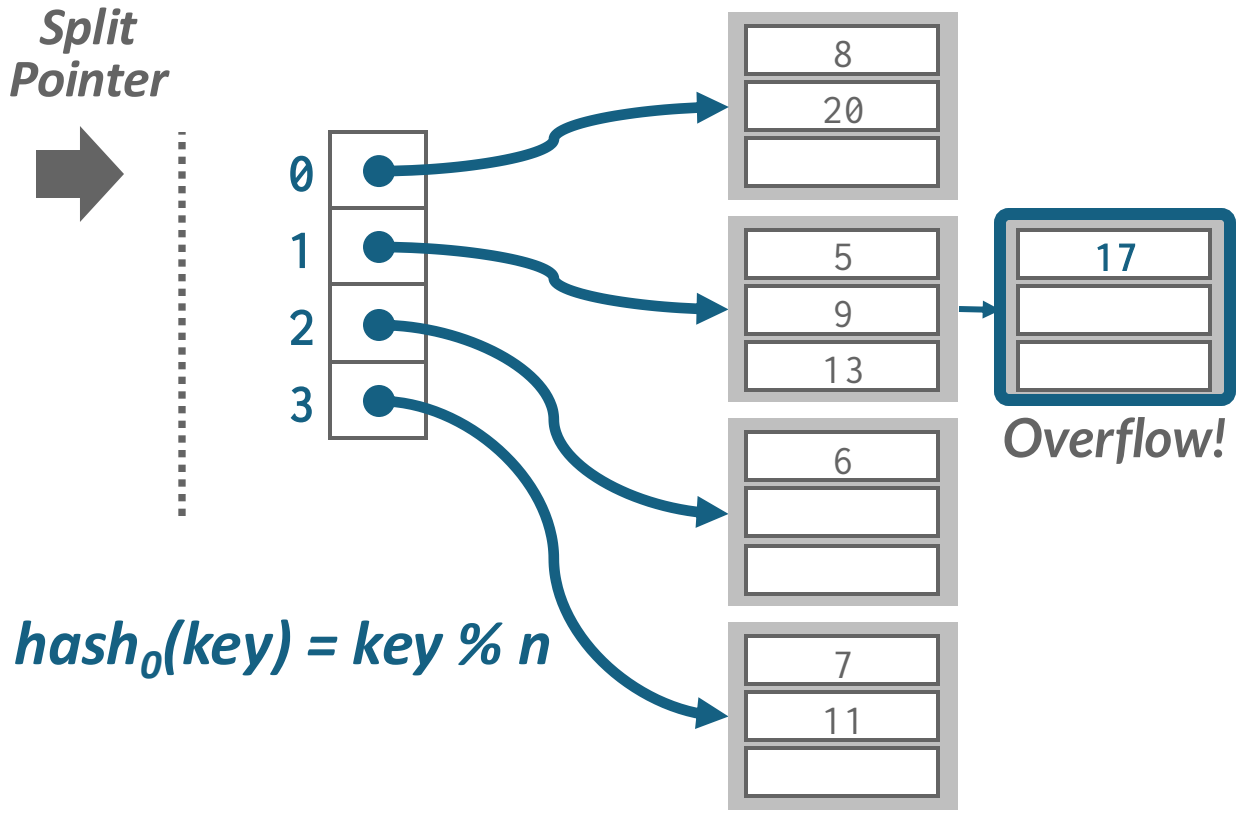

Get 6

$hash_0(6) = 6 \% 4 = 2$

Put 17

$hash_0(17) = 17 \% 4 = 1$

Linear Hashing

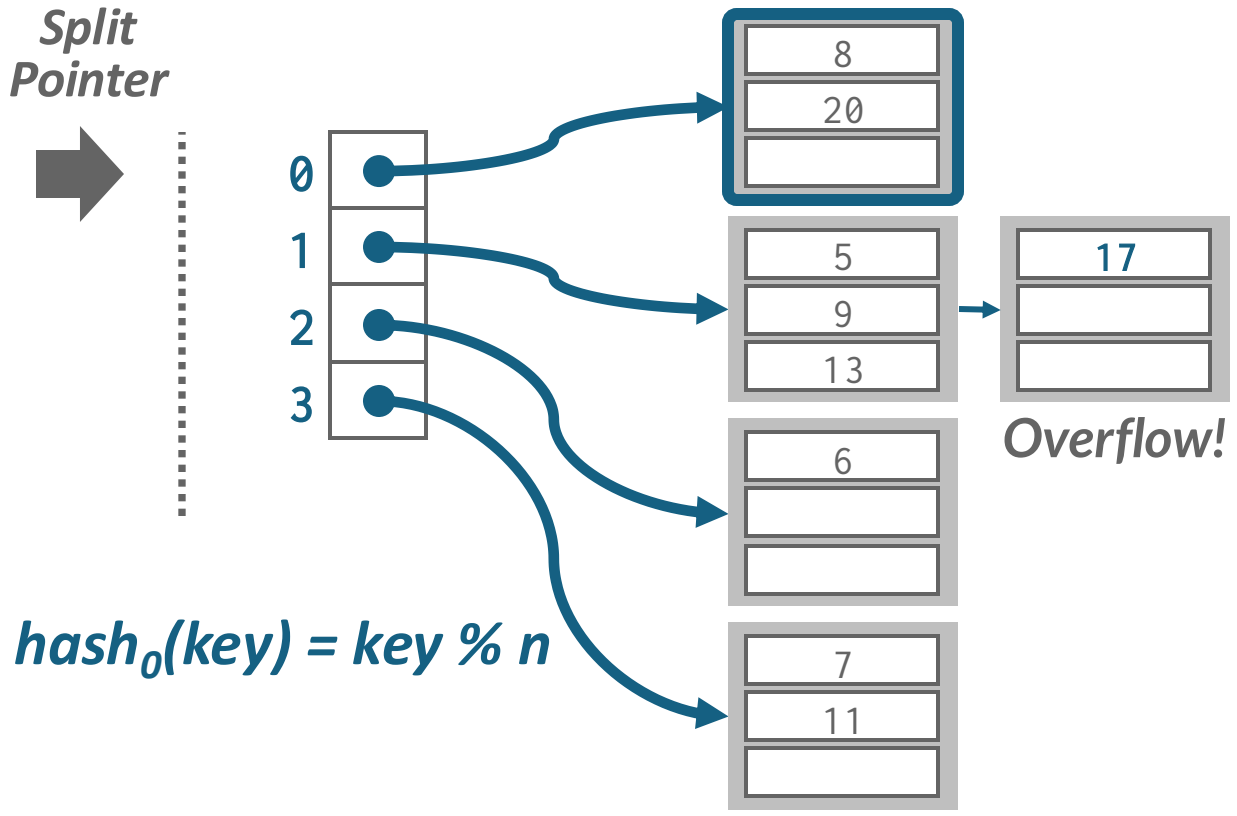

Get 6

$hash_0(6) = 6 \% 4 = 2$

Put 17

$hash_0(17) = 17 \% 4 = 1$

Linear Hashing

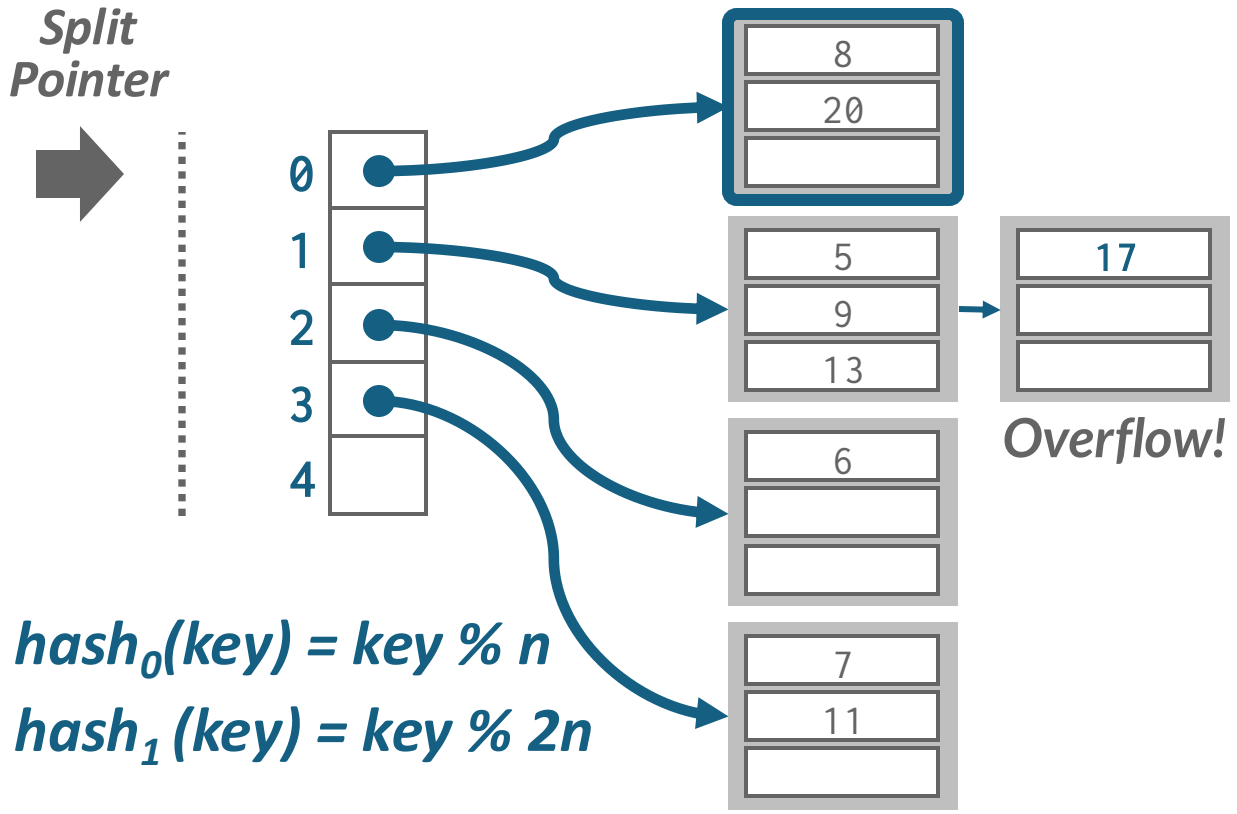

Get 6

$hash_0(6) = 6 \% 4 = 2$

Put 17

$hash_0(17) = 17 \% 4 = 1$

Linear Hashing

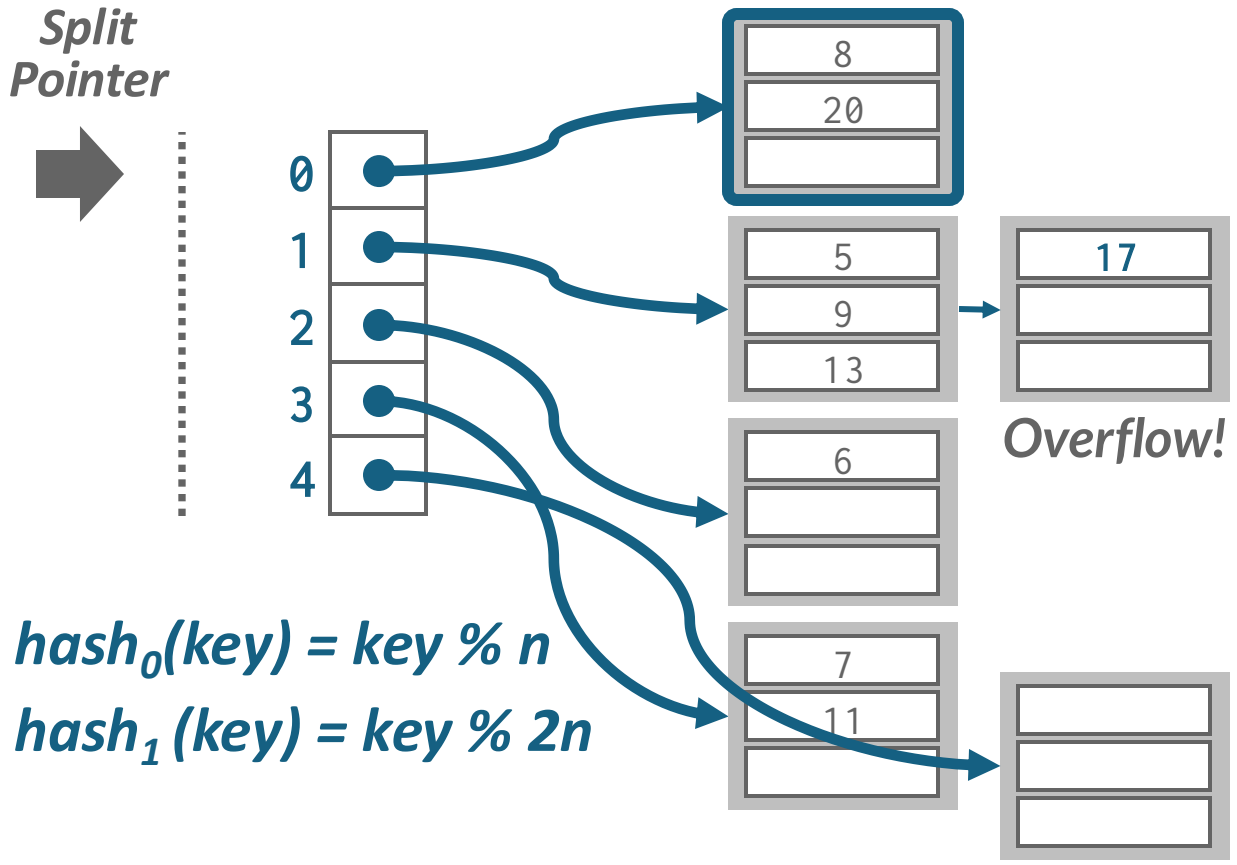

Get 6

$hash_0(6) = 6 \% 4 = 2$

Put 17

$hash_0(17) = 17 \% 4 = 1$

Linear Hashing

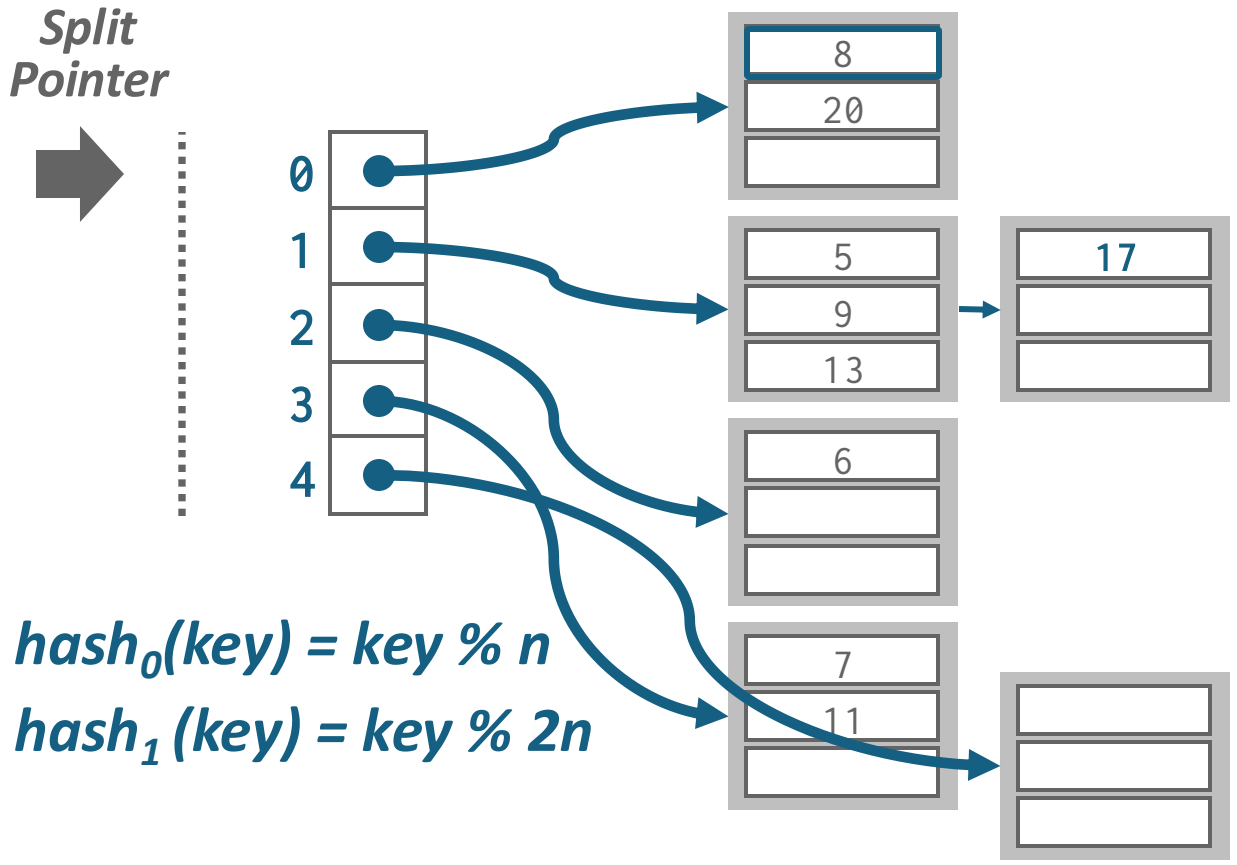

Get 6

$hash_0(6) = 6 \% 4 = 2$

Put 17

$hash_0(17) = 17 \% 4 = 1$

$hash_1(8) = 8 \% 8 = 0$

Linear Hashing

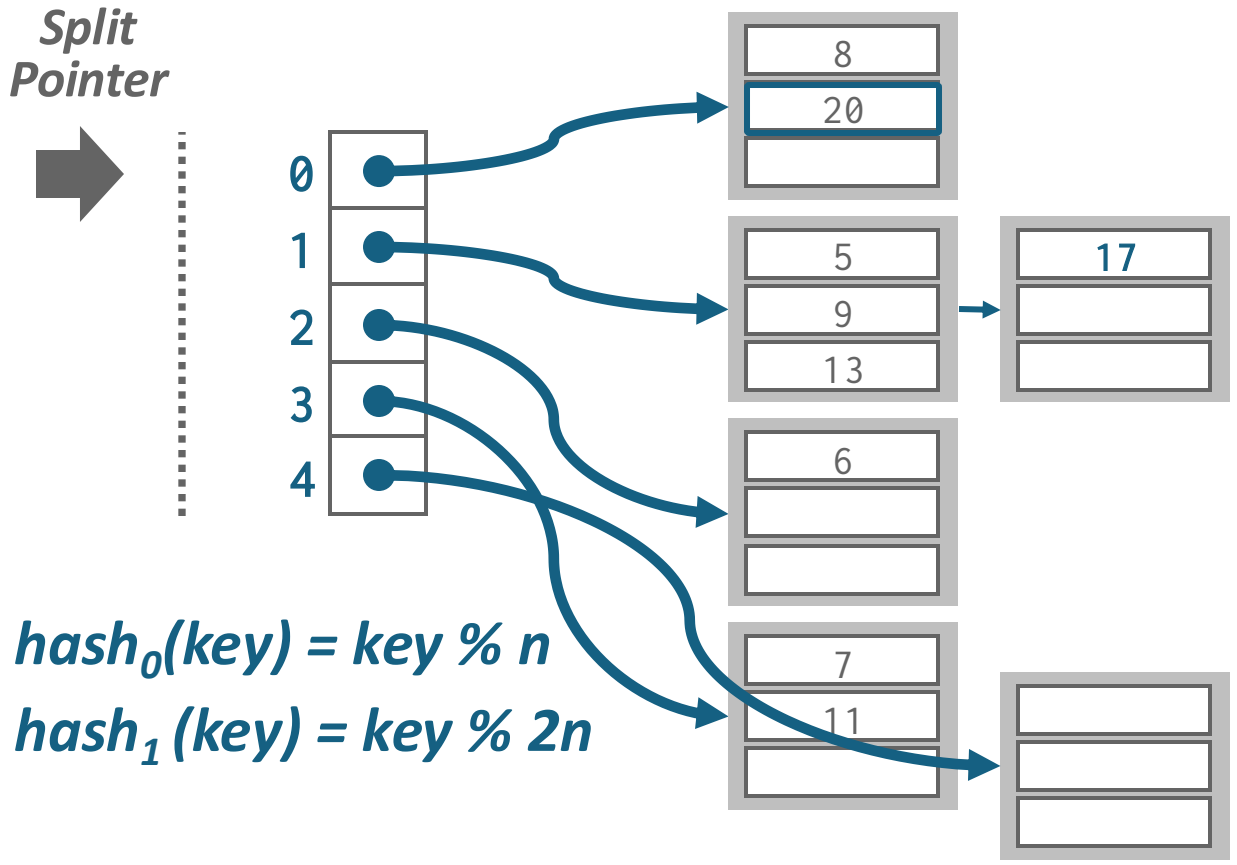

Get 6

$hash_0(6) = 6 \% 4 = 2$

Put 17

$hash_0(17) = 17 \% 4 = 1$

$hash_1(8) = 8 \% 8 = 0$

$hash_1(20) = 20 \% 8 = 4$

Linear Hashing

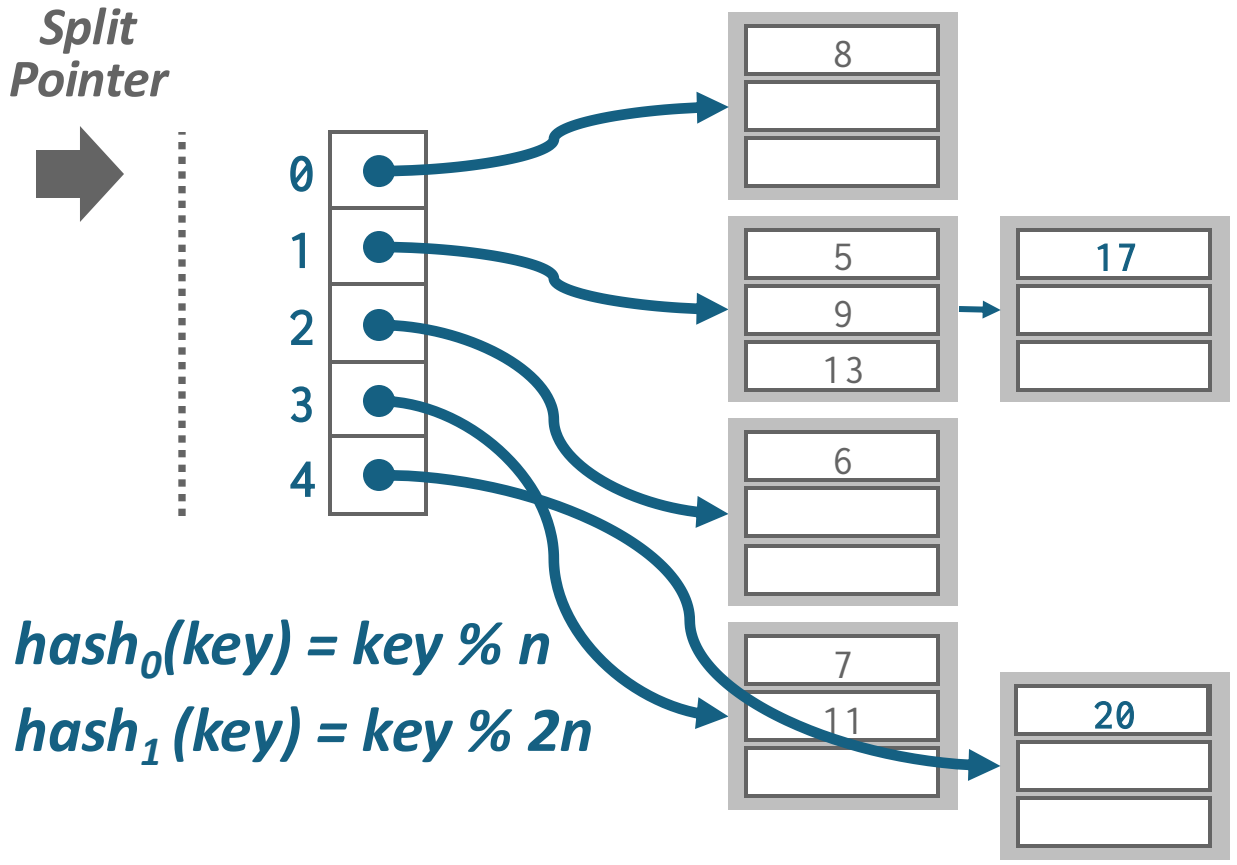

Get 6

$hash_0(6) = 6 \% 4 = 2$

Put 17

$hash_0(17) = 17 \% 4 = 1$

$hash_1(8) = 8 \% 8 = 0$

$hash_1(20) = 20 \% 8 = 4$

Linear Hashing

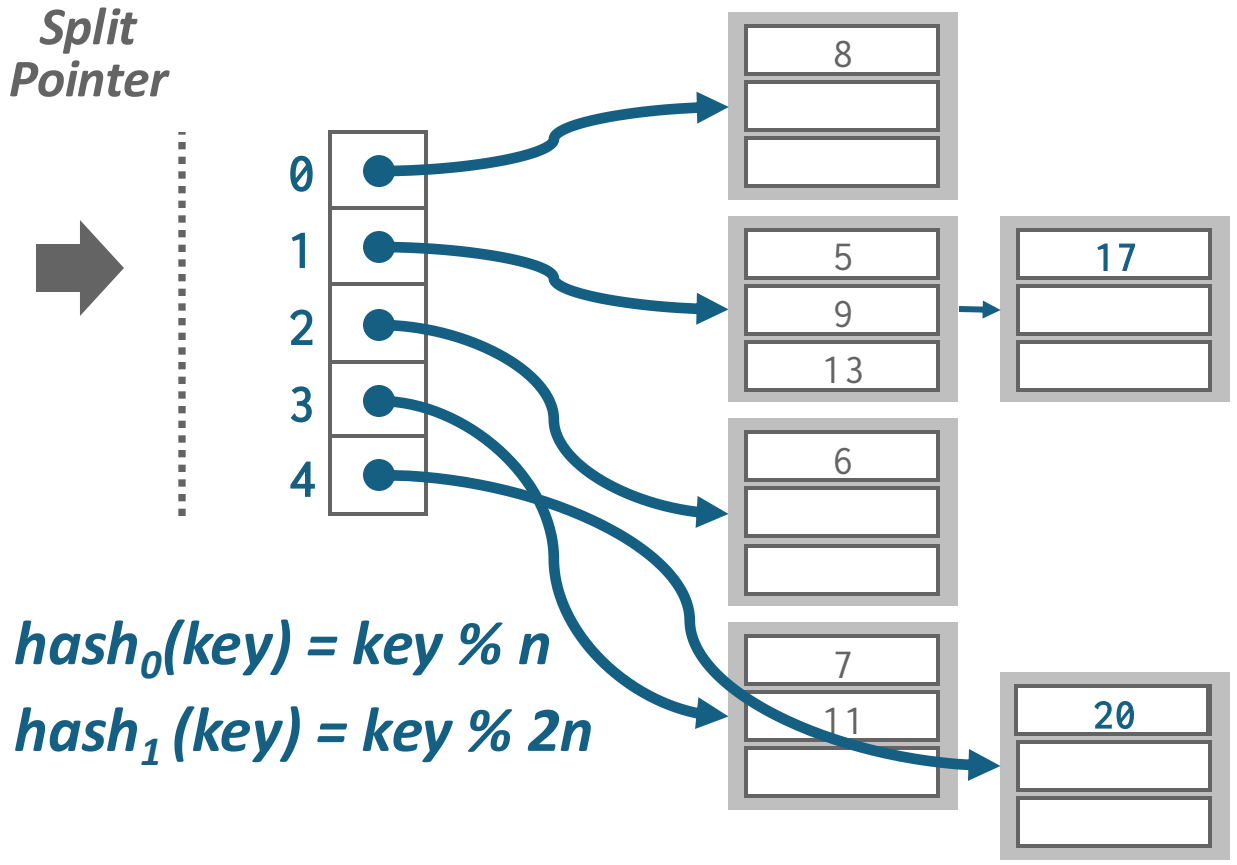

Get 6

$hash_0(6) = 6 \% 4 = 2$

Put 17

$hash_0(17) = 17 \% 4 = 1$

$hash_1(8) = 8 \% 8 = 0$

$hash_1(20) = 20 \% 8 = 4$

Linear Hashing

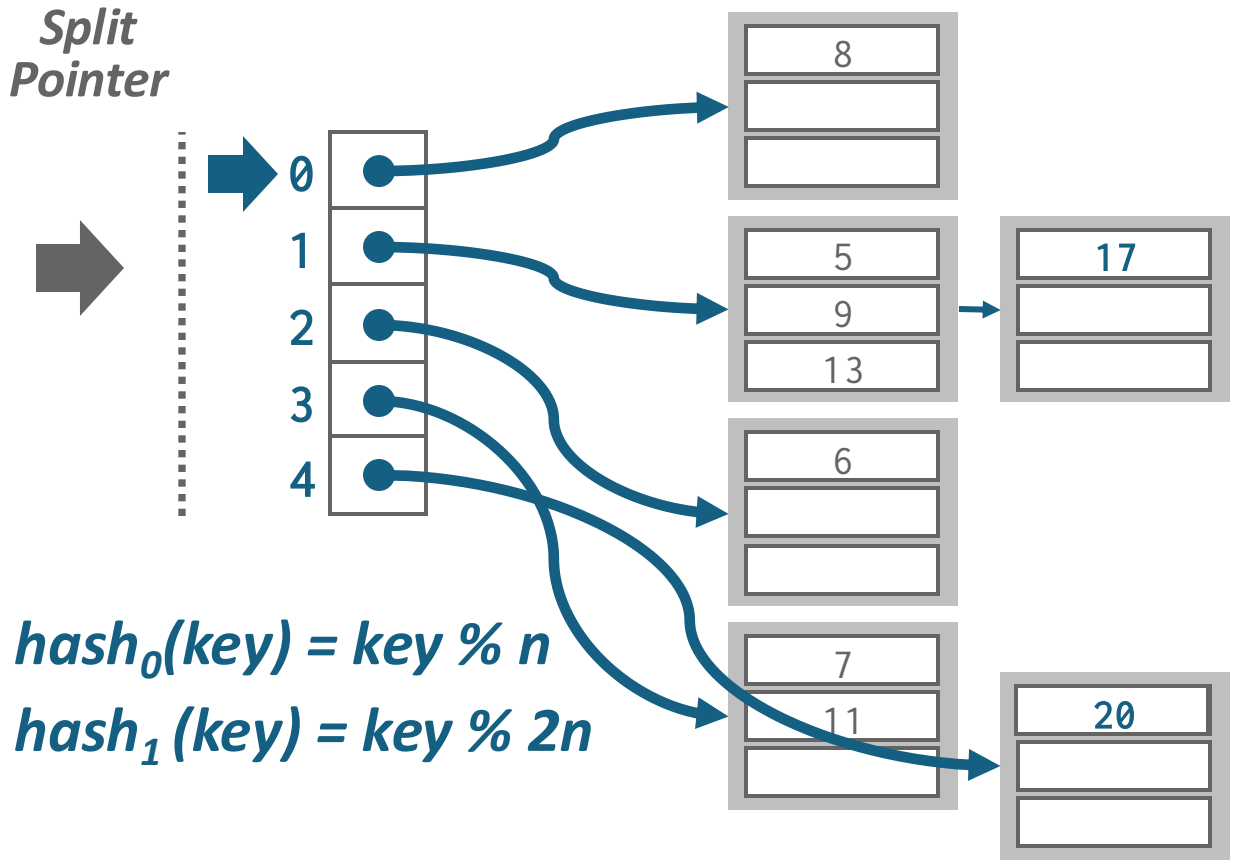

Get 6

$hash_0(6) = 6 \% 4 = 2$

Put 17

$hash_0(17) = 17 \% 4 = 1$

$hash_1(8) = 8 \% 8 = 0$

$hash_1(20) = 20 \% 8 = 4$

Get 20

$hash_0(20) = 20 \% 4 = 0$

Linear Hashing

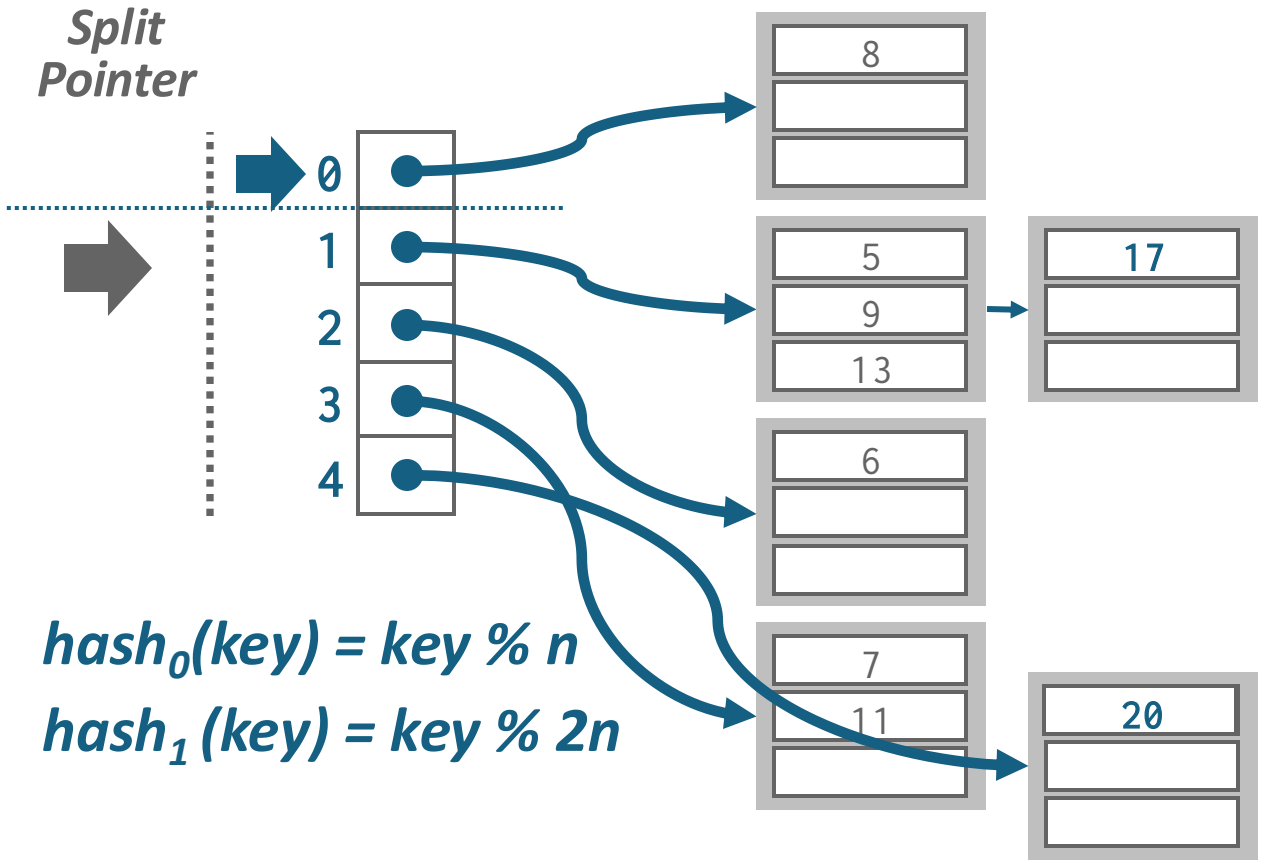

Get 6

$hash_0(6) = 6 \% 4 = 2$

Put 17

$hash_0(17) = 17 \% 4 = 1$

$hash_1(8) = 8 \% 8 = 0$

$hash_1(20) = 20 \% 8 = 4$

Get 20

$hash_0(20) = 20 \% 4 = 0$

Linear Hashing

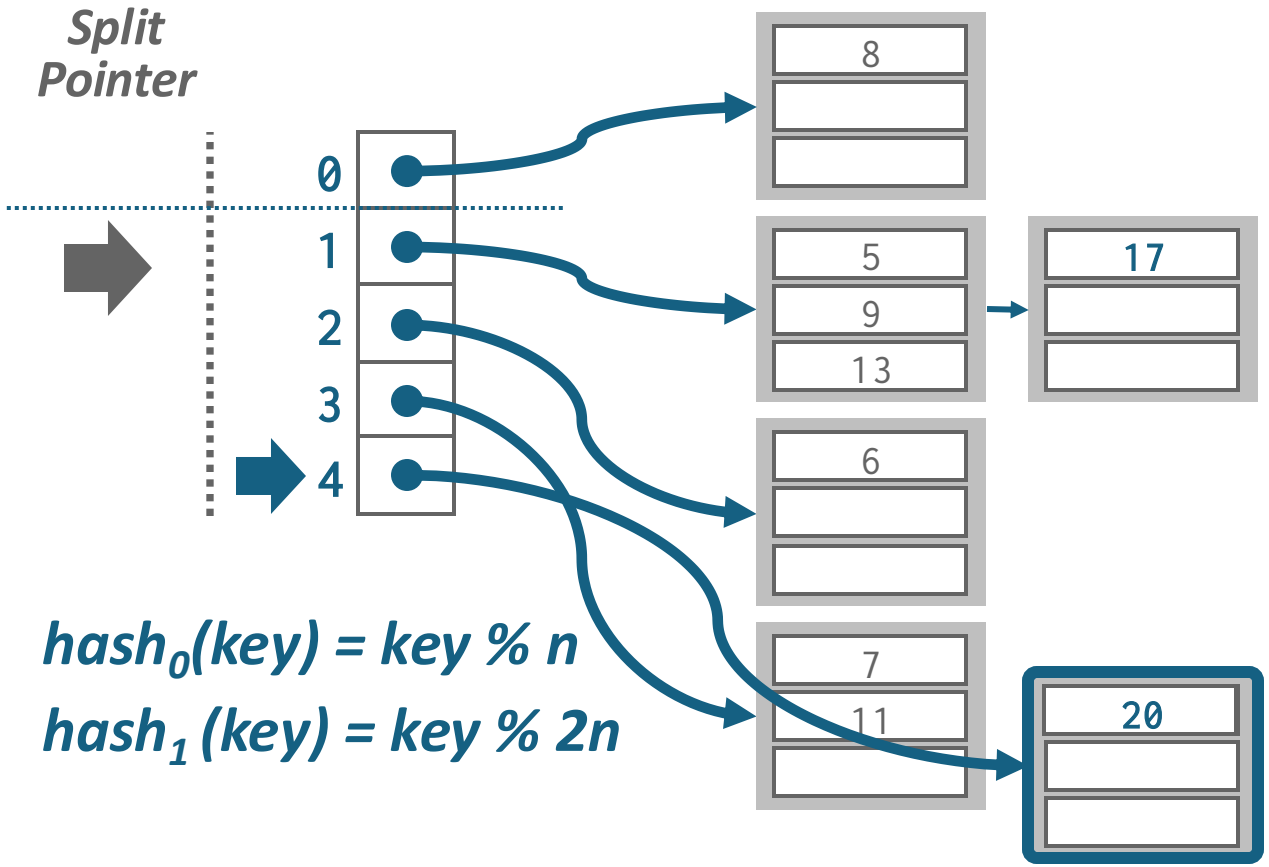

Get 6

$hash_0(6) = 6 \% 4 = 2$

Put 17

$hash_0(17) = 17 \% 4 = 1$

$hash_1(8) = 8 \% 8 = 0$

$hash_1(20) = 20 \% 8 = 4$

Get 20

$hash_0(20) = 20 \% 4 = 0$

$hash_1(20) = 20 \% 8 = 4$

Linear Hashing

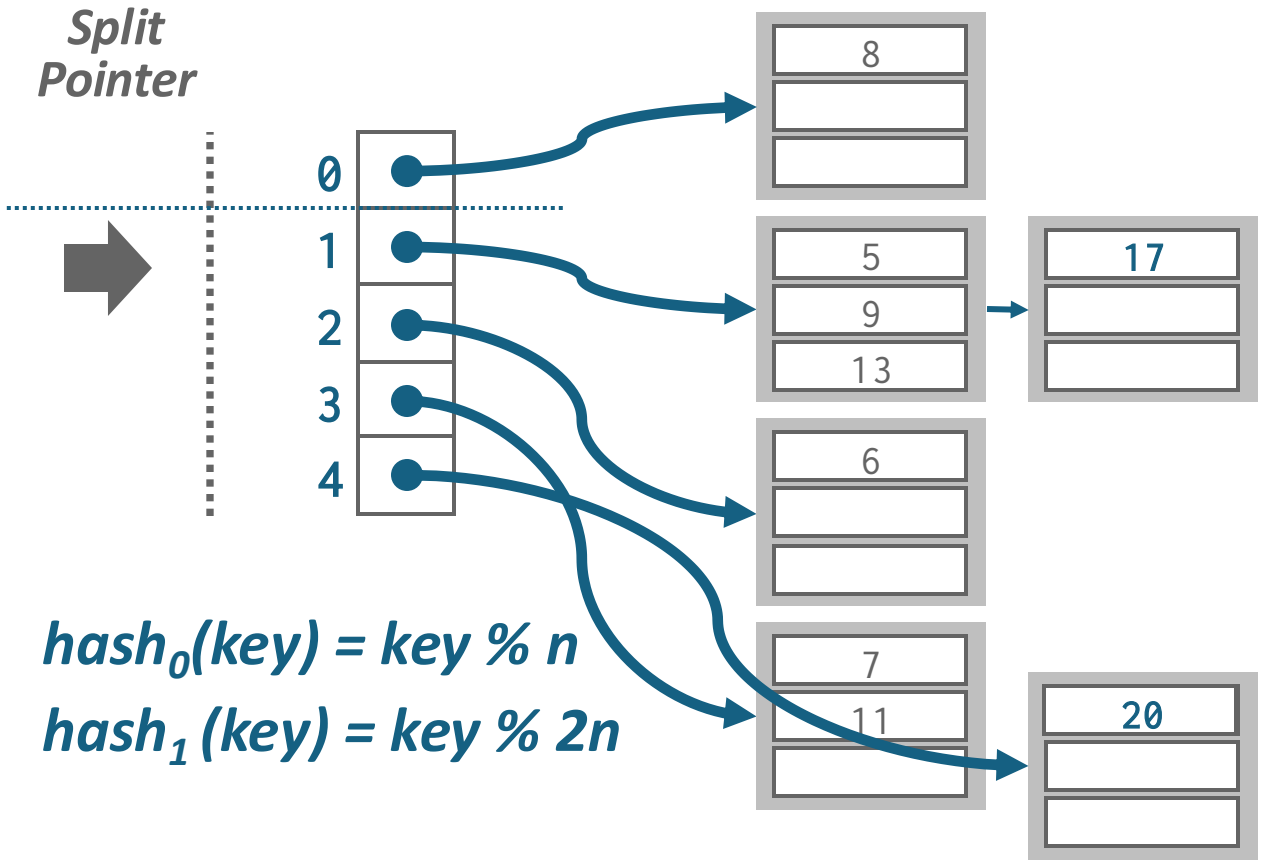

Get 6

$hash_0(6) = 6 \% 4 = 2$

Put 17

$hash_0(17) = 17 \% 4 = 1$

$hash_1(8) = 8 \% 8 = 0$

$hash_1(20) = 20 \% 8 = 4$

Get 20

$hash_0(20) = 20 \% 4 = 0$

$hash_1(20) = 20 \% 8 = 4$

Get 9

$hash_0(9) = 9 \% 4 = 1$

Linear Hashing

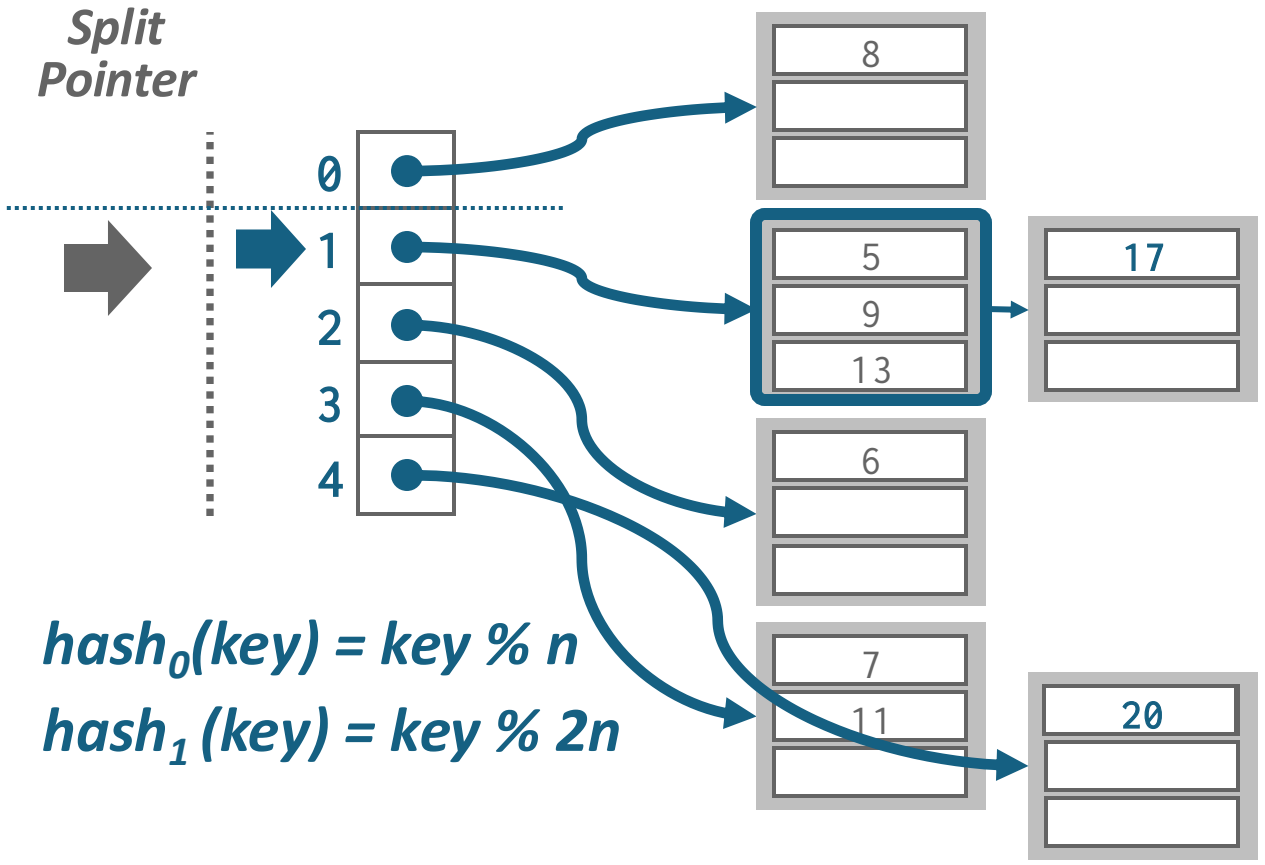

Get 6
$hash_0(6) = 6 \% 4 = 2$

Put 17
$hash_0(17) = 17 \% 4 = 1$
$hash_1(8) = 8 \% 8 = 0$
$hash_1(20) = 20 \% 8 = 4$

Get 20
$hash_0(20) = 20 \% 4 = 0$
$hash_1(20) = 20 \% 8 = 4$

Get 9
$hash_0(9) = 9 \% 4 = 1$

Linear Hashing - Resizing

- Splitting buckets based on the split pointer will eventually get to all overflowed buckets.
  - When the pointer reaches the last slot, remove the first hash function and move pointer back to beginning.

- If the “highest” bucket below the split pointer is empty, the hash table could remove it and move the splinter pointer in reverse direction.

Linear Hashing - Deletes

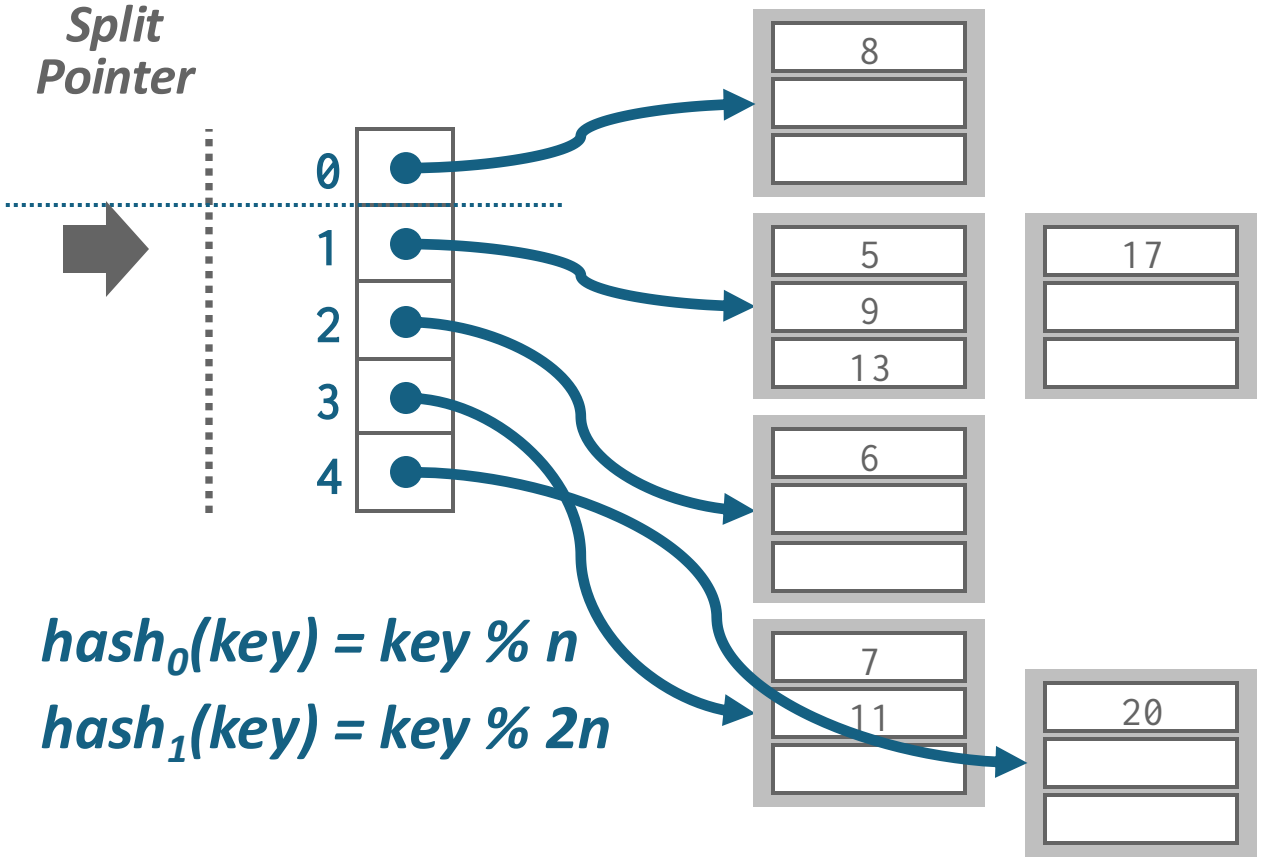

Linear Hashing - Deletes

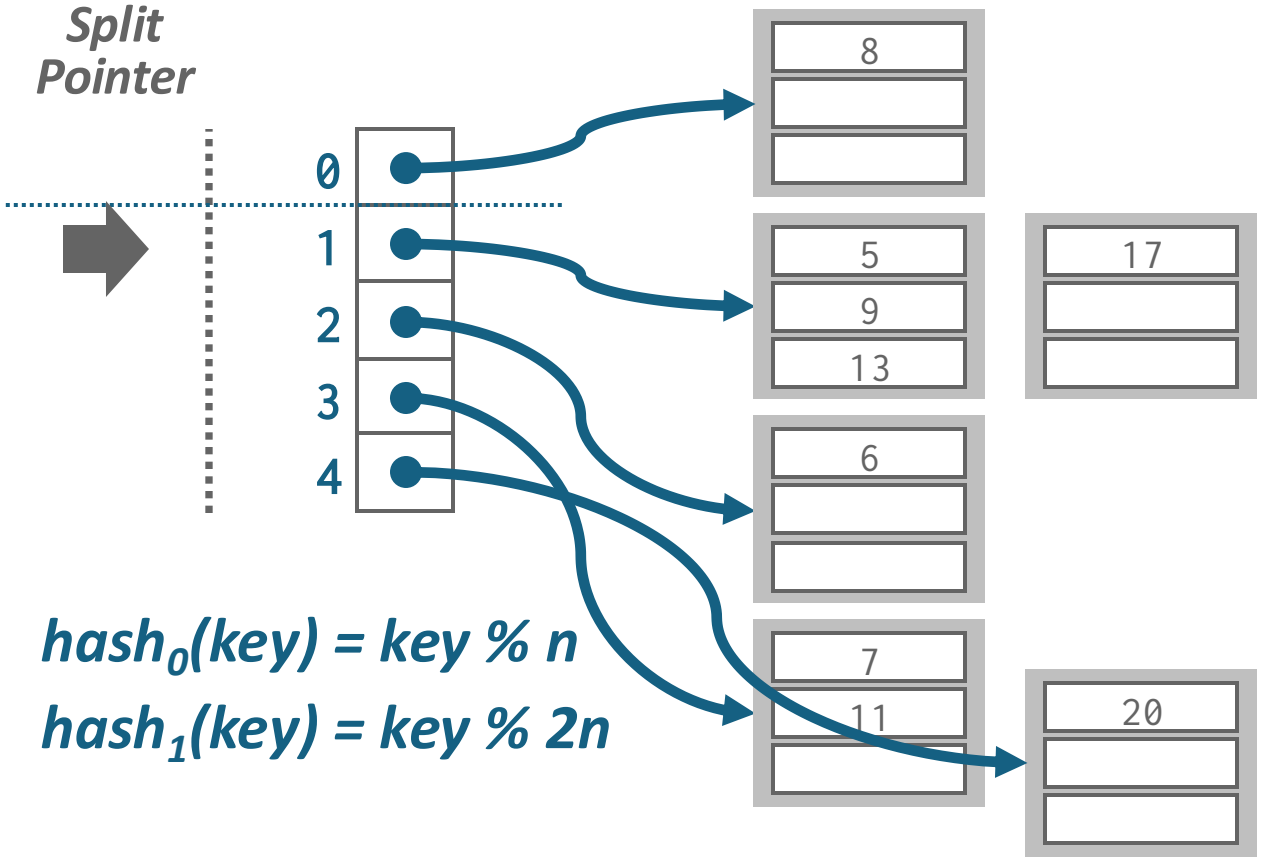

Delete 20

$hash_0(20) = 20 \% 4 = 0$

Linear Hashing - Deletes

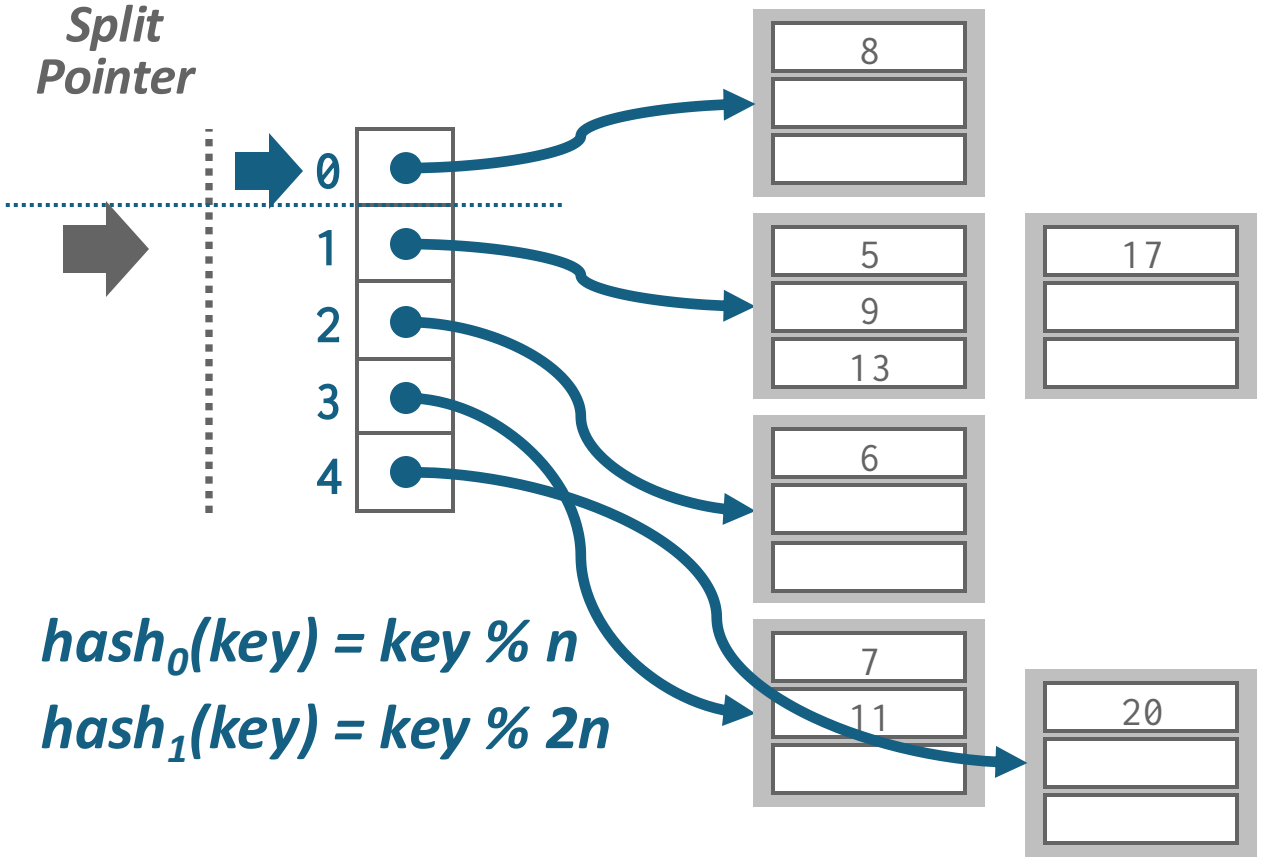

Delete 20

$hash_0(20) = 20 \% 4 = 0$

Linear Hashing - Deletes

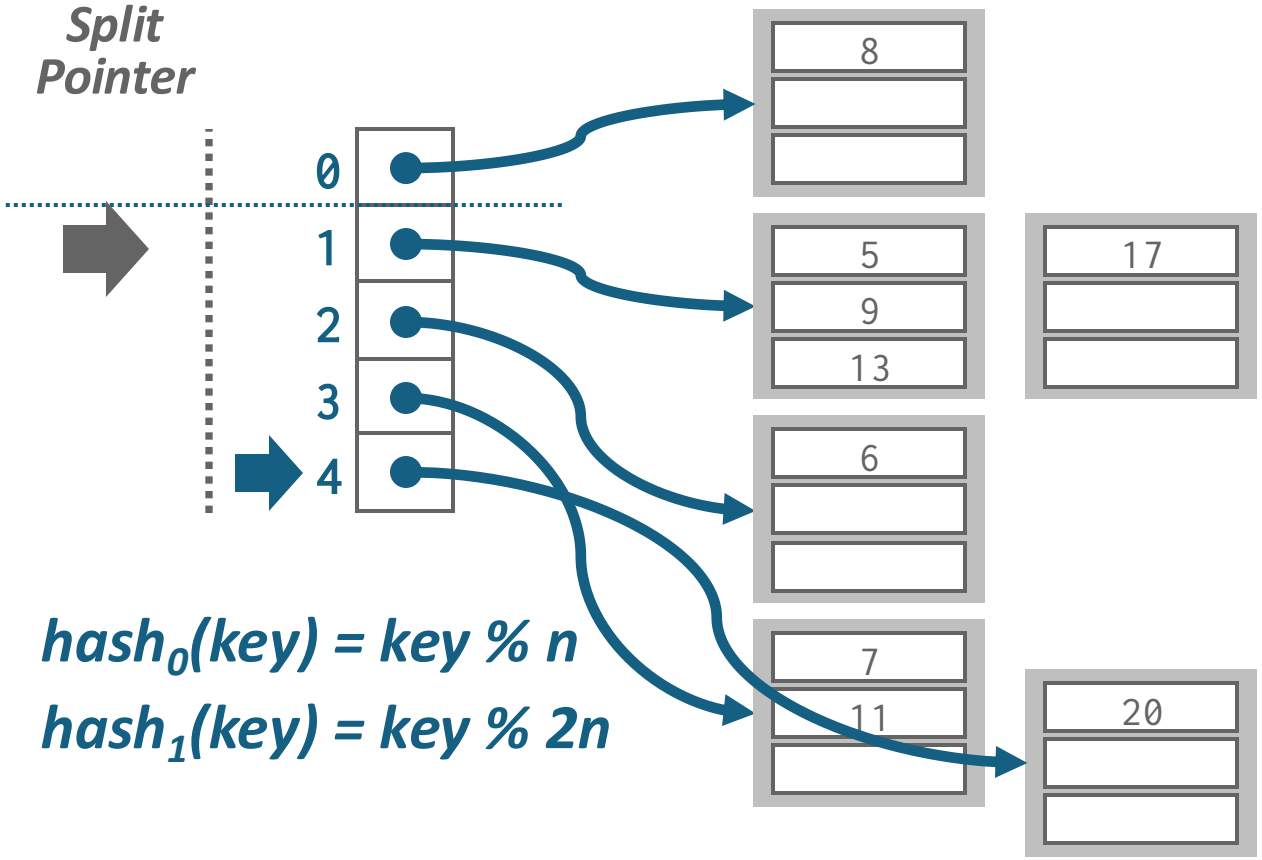

Delete 20

$hash_0(20) = 20 \% 4 = 0$

$hash_1(20) = 20 \% 8 = 4$

Linear Hashing - Deletes

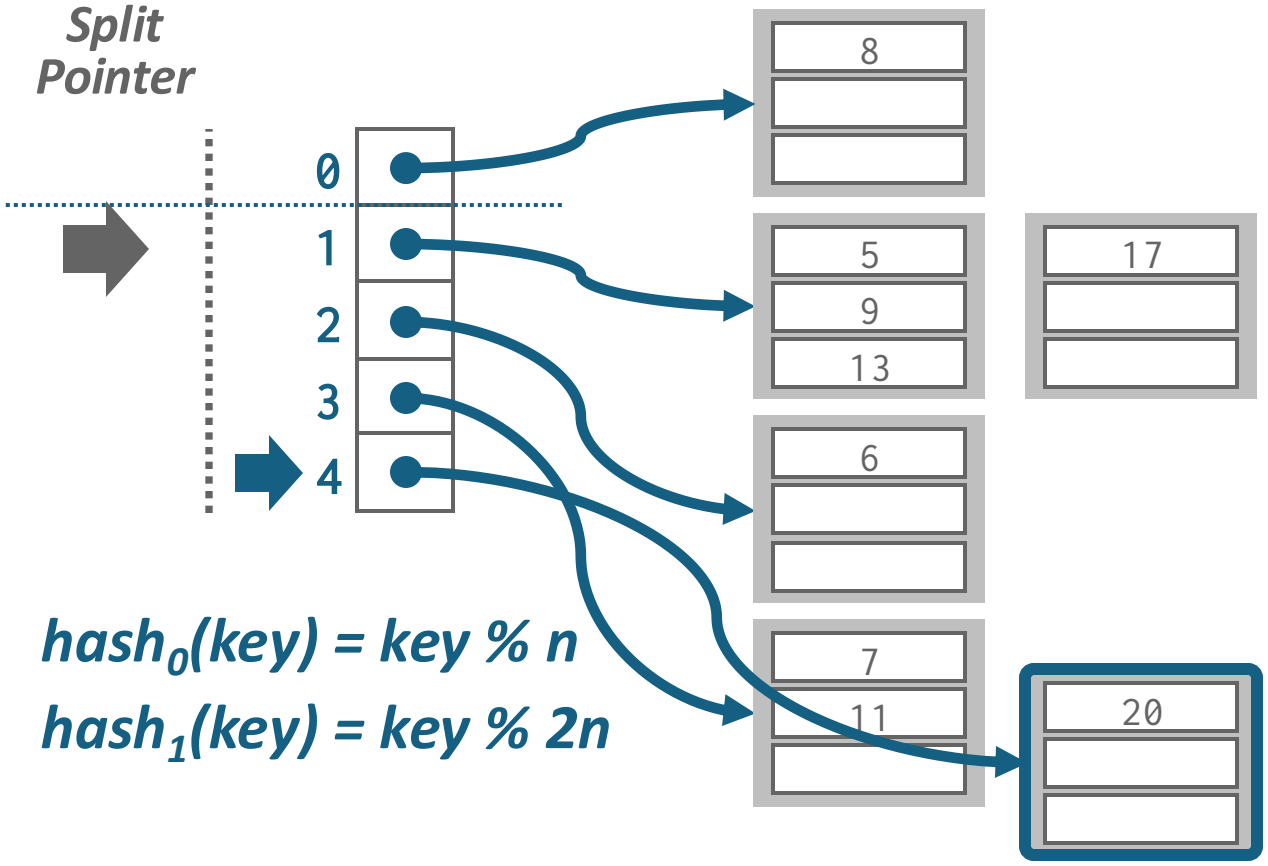

Delete 20

$hash_0(20) = 20 \% 4 = 0$

$hash_1(20) = 20 \% 8 = 4$

Linear Hashing - Deletes

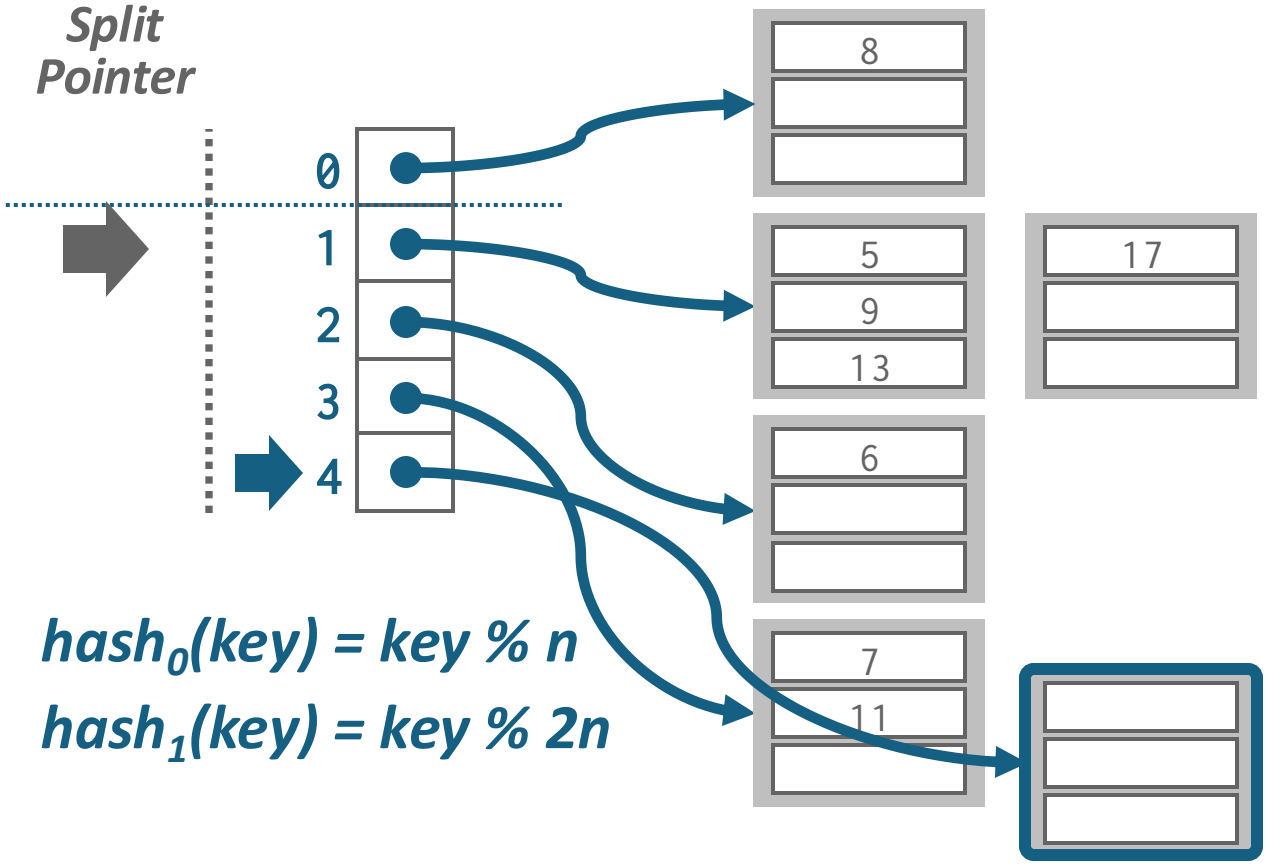

Delete 20

$hash_0(20) = 20 \% 4 = 0$

$hash_1(20) = 20 \% 8 = 4$

Linear Hashing - Deletes

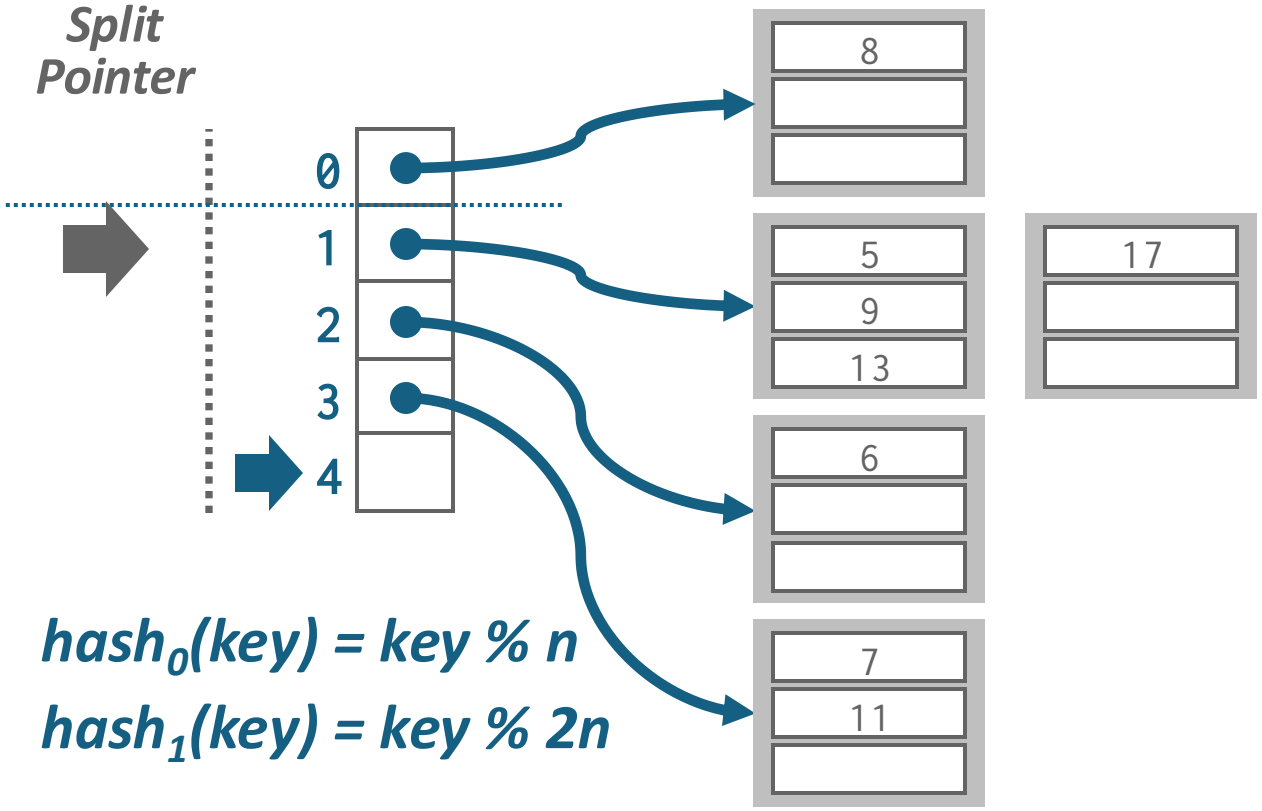

Delete 20

$hash_0(20) = 20 \% 4 = 0$

$hash_1(20) = 20 \% 8 = 4$

Linear Hashing - Deletes

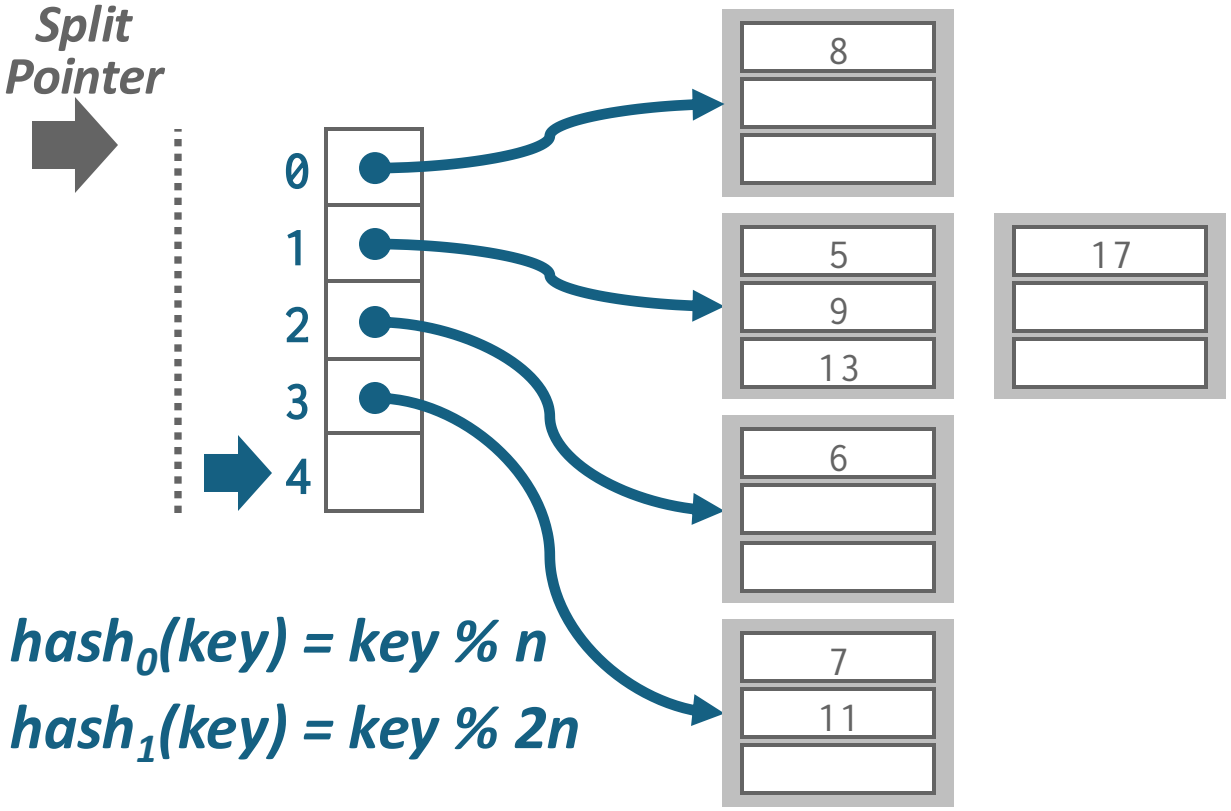

Delete 20

$hash_0(20) = 20 \% 4 = 0$

$hash_1(20) = 20 \% 8 = 4$

Linear Hashing - Deletes

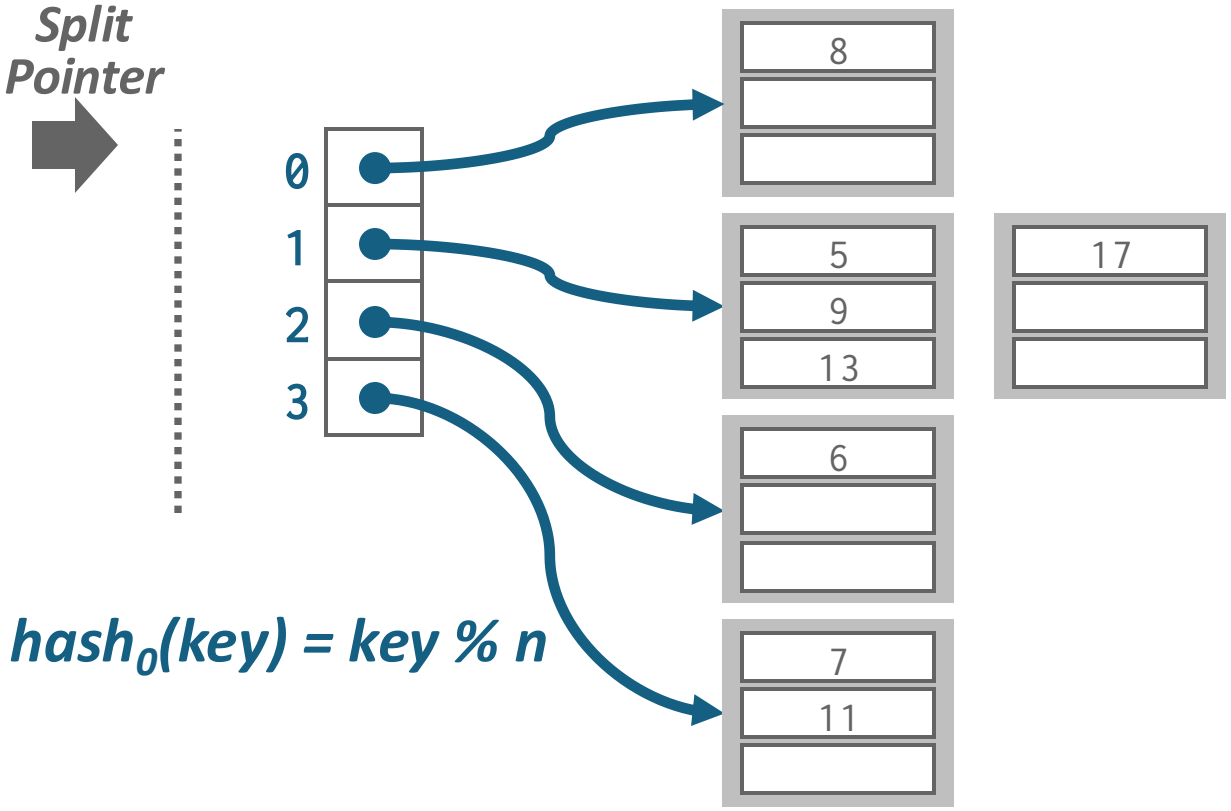

Delete 20

$hash_0(20) = 20 \% 4 = 0$

$hash_1(20) = 20 \% 8 = 4$

Linear Hashing - Deletes

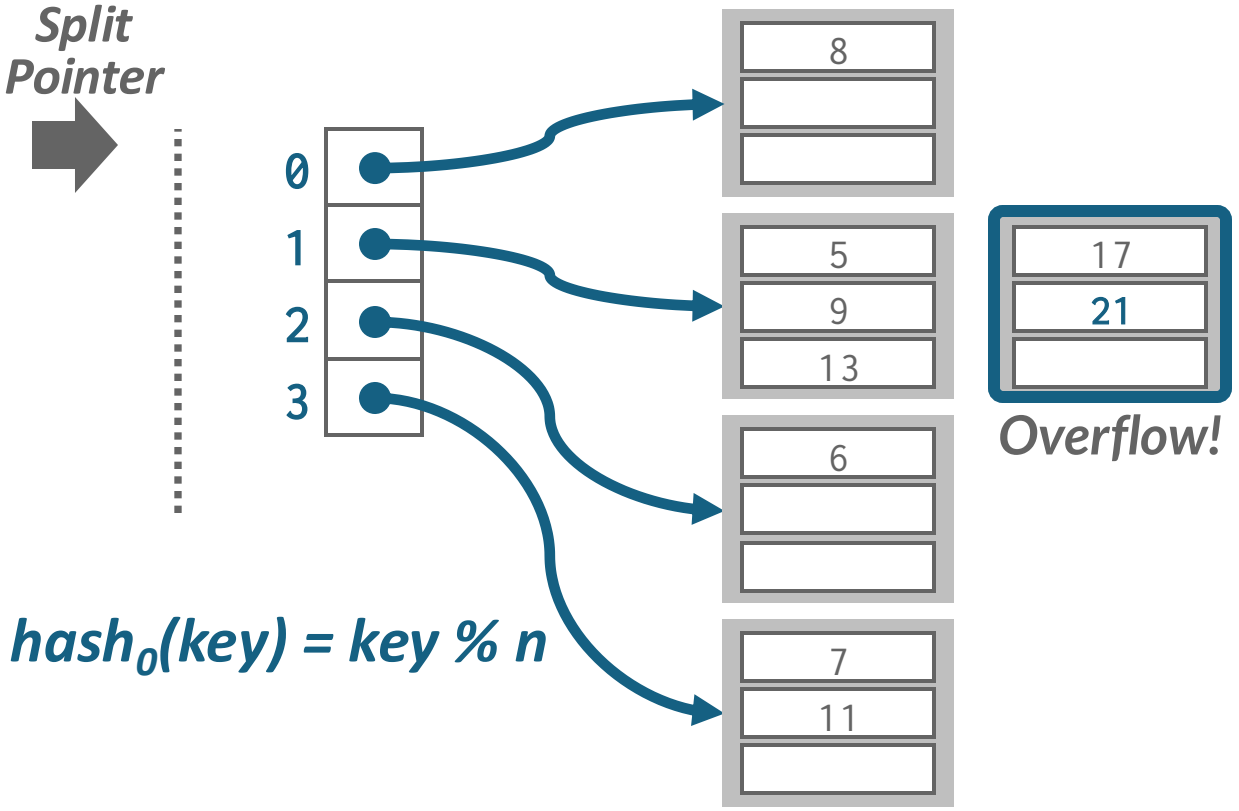

Delete 20

$hash_0(20) = 20 \% 4 = 0$

$hash_1(20) = 20 \% 8 = 4$

Put 21

$hash_0(21) = 21 \% 4 = 1$

Conclusion

- Fast data structures that support $O(1)$ look-ups that are used all throughout DBMS internals.
  - Trade-off between speed and flexibility.

- Hashing schemes get used for both in-memory and on-disk.
  - Linear Probing, Cuckoo Hashing, Chained Hashing: Generally used for in-memory hash tables; e.g., in hash join and aggregate operators.
  - Extendible Hashing, Linear Hashing: for disk-based hash indexing.

Next Lecture

- **B+Trees**
  - aka “The Greatest Data Structure of All Time”