



香港中文大學(深圳)  
The Chinese University of Hong Kong, Shenzhen



Ack: Prof. Jignesh Patel @ CMU  
Prof. Andy Pavlo @CMU

# CSC3170

# 9: Index Concurrency Control

Chenhao Ma

School of Data Science

The Chinese University of Hong Kong, Shenzhen

# Observation

- We (mostly) assumed all the data structures that we have discussed so far are single-threaded.
- But a DBMS needs to allow **multiple threads** to safely access data structures to take advantage of additional CPU cores and hide disk I/O stalls.

# Observation

- We (mostly) assumed all the data structures that we have discussed so far are single-threaded.
- But a DBMS needs to allow **multiple threads** to safely access data structures to take advantage of additional CPU cores and hide disk I/O stalls.



# Concurrency Control

- A **concurrency control** protocol is the method that the DBMS uses to ensure “correct” results for concurrent operations on a shared object.
- A protocol’s correctness criteria can vary:
  - **Logical Correctness:** Can a thread see the data that it is supposed to see?
  - **Physical Correctness:** Is the internal representation of the object sound?

# Concurrency Control

- A **concurrency control** protocol is the method that the DBMS uses to ensure “correct” results for concurrent operations on a shared object.
- A protocol’s correctness criteria can vary:
  - **Logical Correctness:** Can a thread see the data that it is supposed to see?
  - **Physical Correctness:** Is the internal representation of the object sound?

# This Lecture

- Latches Overview
- Hash Table Latching
- B+Tree Latching
- Leaf Node Scans

# Latches

# Locks vs. Latches

- **Locks (Transactions)**
  - Protect the database's logical contents from other transactions.
  - Held for transaction's duration.
  - Need to be able to rollback changes.
- **Latches (Workers)**
  - Protect the critical sections of the DBMS's internal data structure from other workers (e.g., threads).
  - Held for operation duration.
  - Do not need to be able to rollback changes.

# Locks vs. Latches

	<b><i>Locks</i></b>	<b><i>Latches</i></b>
<b>Separate...</b>	Transactions	Workers (threads, processes)
<b>Protect...</b>	Database Contents	In-Memory Data Structures
<b>During...</b>	Entire Transactions	Critical Sections
<b>Modes...</b>	Shared, Exclusive, Update, Intention	Read, Write
<b>Deadlock</b>	Detection & Resolution	Avoidance
<b>...by...</b>	Waits-for, Timeout, Aborts	Coding Discipline
<b>Kept in...</b>	Lock Manager	Protected Data Structure

Source: [Goetz Graefe](#)

# Locks vs. Latches

<b><i>Locks</i></b>		<b><i>Latches</i></b>
<b>Separate...</b>	Transactions	Workers (threads, processes)
<b>Protect...</b>	Database Contents	In-Memory Data Structures
<b>During...</b>	Entire Transactions	Critical Sections
<b>Modes...</b>	Shared, Exclusive, Update, Intention	Read, Write
<b>Deadlock</b>	Detection & Resolution	Avoidance
<b>...by...</b>	Waits-for, Timeout, Aborts	Coding Discipline
<b>Kept in...</b>	Lock Manager	Protected Data Structure

Source: [Goetz Graefe](#)

# Locks vs. Latches

Future  
Lecture



## Locks

- Separate...** Transactions
- Protect...** Database Contents
- During...** Entire Transactions
- Modes...** Shared, Exclusive, Update, Intention
- Deadlock** Detection & Resolution
  - ...by...** Waits-for, Timeout, Aborts
- Kept in...** Lock Manager

## Latches

- Workers (threads, processes)
- In-Memory Data Structures
- Critical Sections
- Read, Write
- Avoidance
- Coding Discipline
- Protected Data Structure

Source: [Goetz Graefe](#)

# Latch Modes

- **Read Mode**

- Multiple threads can read the same object at the same time.
- A thread can acquire the read latch if another thread has it in read mode.

- **Write Mode**

- Only one thread can access the object.
- A thread cannot acquire a write latch if another thread has it in any mode.

# Latch Modes

- **Read Mode**

- Multiple threads can read the same object at the same time.
- A thread can acquire the read latch if another thread has it in read mode.

- **Write Mode**

- Only one thread can access the object.
- A thread cannot acquire a write latch if another thread has it in any mode.

		Compatibility Matrix	
		Read	Write
Read	Read	✓	✗
	Write	✗	✗

# Latch Implementation Goals

- Small memory footprint.
- Fast execution path when no contention.
- Deschedule thread when it has been waiting for too long to avoid burning cycles.
- Each latch should not have to implement their own queue to track waiting threads.

Source: [Filip Pizlo](#)

# Latch Implementations

- Test-and-Set Spinlock
- Blocking OS Mutex
- Reader-Writer Locks
- Advanced approaches:
  - Adaptive Spinlock ([Apple ParkingLot](#))
  - Queue-based Spinlock ([MCS Locks](#))

# Latch Implementations

- **Approach #1: Test-and-Set Spin Latch (TAS)**
  - Very efficient (single instruction to latch/unlatch)
  - Non-scalable, not cache friendly, not OS friendly.
  - Example: `std::atomic<T>`

# Latch Implementations

- **Approach #1: Test-and-Set Spin Latch (TAS)**
  - Very efficient (single instruction to latch/unlatch)
  - Non-scalable, not cache friendly, not OS friendly.
  - Example: `std::atomic<T>`

```
std::atomic_flag latch;  
:  
while (latch.test_and_set(...)) {  
    // Retry? Yield? Abort?  
}
```

# Latch Implementations

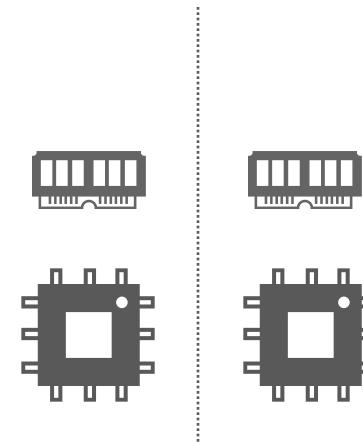
- **Approach #1: Test-and-Set Spin Latch (TAS)**
  - Very efficient (single instruction to latch/unlatch)
  - Non-scalable, not cache friendly, not OS friendly.
  - Example: `std::atomic<T>`

```
std::atomic<bool>
{
    std::atomic_flag latch;
    :
    while (latch.test_and_set(...)) {
        // Retry? Yield? Abort?
    }
}
```

# Latch Implementations

- **Approach #1: Test-and-Set Spin Latch (TAS)**
  - Very efficient (single instruction to latch/unlatch)
  - Non-scalable, not cache friendly, not OS friendly.
  - Example: `std::atomic<T>`

```
std::atomic<bool>
{
    std::atomic_flag latch;
    :
    while (latch.test_and_set(...)) {
        // Retry? Yield? Abort?
    }
}
```

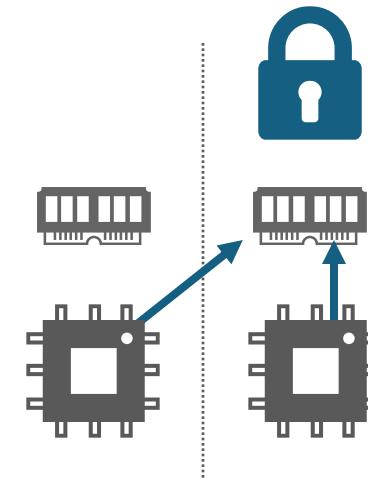


# Latch Implementations

- **Approach #1: Test-and-Set Spin Latch (TAS)**

- Very efficient (single instruction to latch/unlatch)
- Non-scalable, not cache friendly, not OS friendly.
- Example: `std::atomic<T>`

```
std::atomic<bool>
{
    std::atomic_flag latch;
    :
    while (latch.test_and_set(...)) {
        // Retry? Yield? Abort?
    }
}
```



# Latch Implementations

- **Approach #2: Blocking OS Mutex**
  - Simple to use
  - Non-scalable (about 25ns per lock/unlock invocation)
  - Example: `std::mutex`

# Latch Implementations

- **Approach #2: Blocking OS Mutex**
  - Simple to use
  - Non-scalable (about 25ns per lock/unlock invocation)
  - Example: `std::mutex`

```
std::mutex m;  
:  
m.lock();  
// Do something special...  
m.unlock();
```

# Latch Implementations

- **Approach #2: Blocking OS Mutex**
  - Simple to use
  - Non-scalable (about 25ns per lock/unlock invocation)
  - Example: `std::mutex` → `pthread_mutex`

```
std::mutex m;  
:  
m.lock();  
// Do something special...  
m.unlock();
```

# Latch Implementations

- **Approach #2: Blocking OS Mutex**
  - Simple to use
  - Non-scalable (about 25ns per lock/unlock invocation)
  - Example: `std::mutex` → `pthread_mutex` → `futex`

```
std::mutex m;  
:  
m.lock();  
// Do something special...  
m.unlock();
```

# Latch Implementations

- **Approach #2: Blocking OS Mutex**
  - Simple to use
  - Non-scalable (about 25ns per lock/unlock invocation)
  - Example: `std::mutex` → `pthread_mutex` → `futex`

```
std::mutex m;  
:  
m.lock();  
// Do something special...  
m.unlock();
```

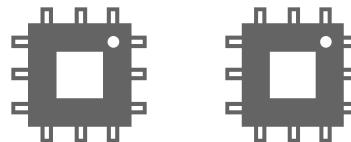


# Latch Implementations

- **Approach #2: Blocking OS Mutex**
  - Simple to use
  - Non-scalable (about 25ns per lock/unlock invocation)
  - Example: `std::mutex` → `pthread_mutex` → `futex`

```
std::mutex m;  
:  
m.lock();  
// Do something special...  
m.unlock();
```

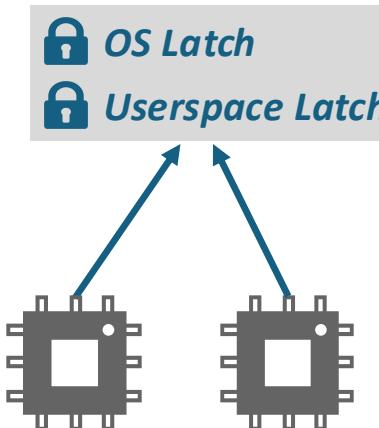
 *OS Latch*  
 *Userspace Latch*



# Latch Implementations

- **Approach #2: Blocking OS Mutex**
  - Simple to use
  - Non-scalable (about 25ns per lock/unlock invocation)
  - Example: `std::mutex` → `pthread_mutex` → `futex`

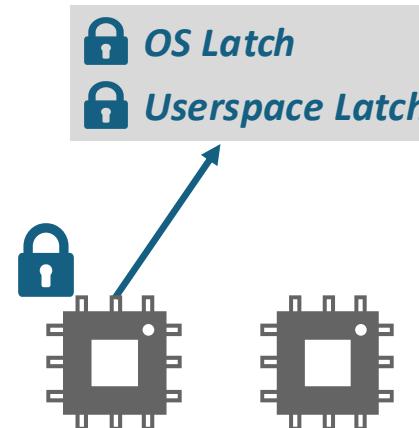
```
std::mutex m;  
:  
m.lock();  
// Do something special...  
m.unlock();
```



# Latch Implementations

- **Approach #2: Blocking OS Mutex**
  - Simple to use
  - Non-scalable (about 25ns per lock/unlock invocation)
  - Example: `std::mutex` → `pthread_mutex` → `futex`

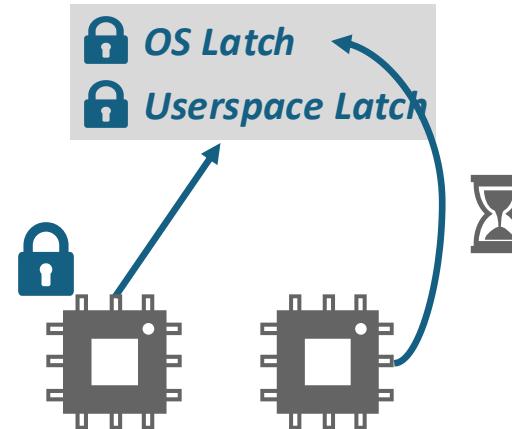
```
std::mutex m;  
:  
m.lock();  
// Do something special...  
m.unlock();
```



# Latch Implementations

- **Approach #2: Blocking OS Mutex**
  - Simple to use
  - Non-scalable (about 25ns per lock/unlock invocation)
  - Example: `std::mutex` → `pthread_mutex` → `futex`

```
std::mutex m;  
:  
m.lock();  
// Do something special...  
m.unlock();
```



# Latch Implementations

- **Approach #3: Reader-Writer Latches**
  - Allows for concurrent readers. Must manage read/write queues to avoid starvation.
  - Can be implemented on top of spinlocks.
  - Example: `std::shared_mutex`

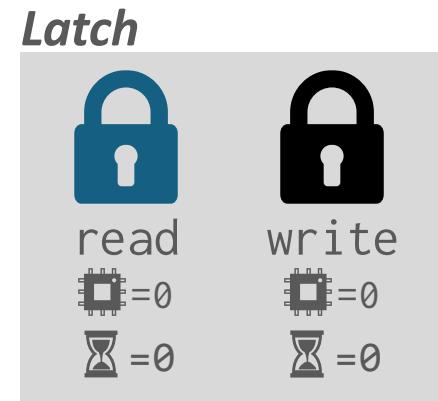
# Latch Implementations

- **Approach #3: Reader-Writer Latches**
  - Allows for concurrent readers. Must manage read/write queues to avoid starvation.
  - Can be implemented on top of spinlocks.
  - Example: `std::shared_mutex` → `pthread_rwlock`

# Latch Implementations

- **Approach #3: Reader-Writer Latches**

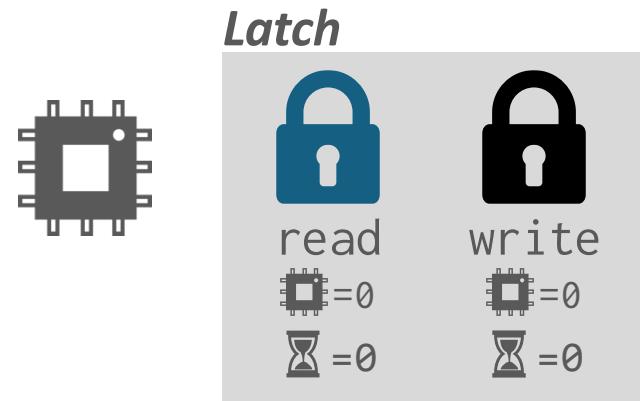
- Allows for concurrent readers. Must manage read/write queues to avoid starvation.
- Can be implemented on top of spinlocks.
- Example: `std::shared_mutex` → `pthread_rwlock`



# Latch Implementations

- **Approach #3: Reader-Writer Latches**

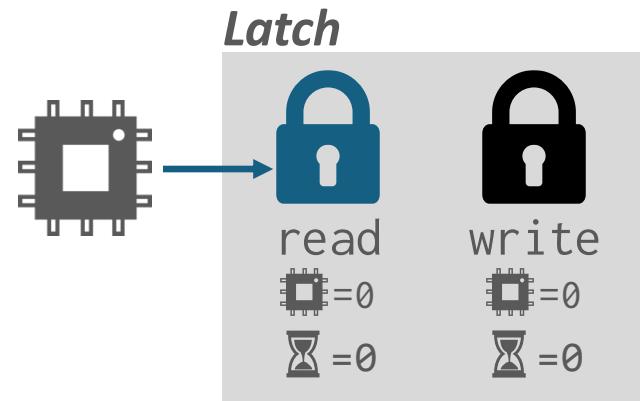
- Allows for concurrent readers. Must manage read/write queues to avoid starvation.
- Can be implemented on top of spinlocks.
- Example: `std::shared_mutex` → `pthread_rwlock`



# Latch Implementations

- **Approach #3: Reader-Writer Latches**

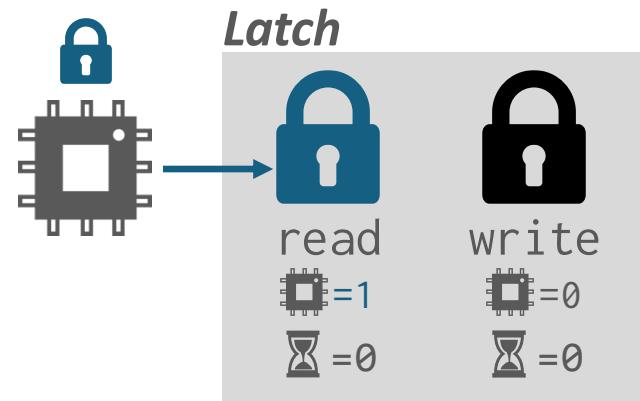
- Allows for concurrent readers. Must manage read/write queues to avoid starvation.
- Can be implemented on top of spinlocks.
- Example: `std::shared_mutex` → `pthread_rwlock`



# Latch Implementations

- **Approach #3: Reader-Writer Latches**

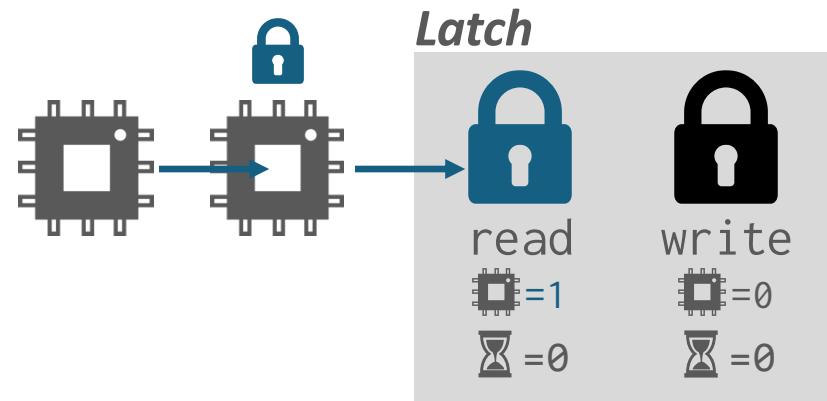
- Allows for concurrent readers. Must manage read/write queues to avoid starvation.
- Can be implemented on top of spinlocks.
- Example: `std::shared_mutex` → `pthread_rwlock`



# Latch Implementations

- **Approach #3: Reader-Writer Latches**

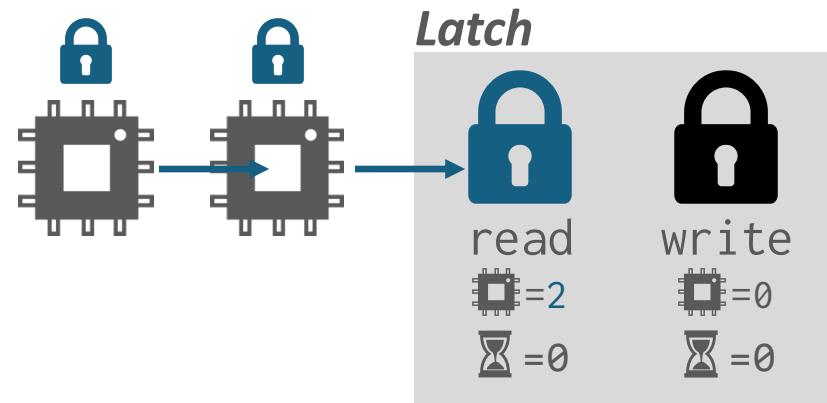
- Allows for concurrent readers. Must manage read/write queues to avoid starvation.
- Can be implemented on top of spinlocks.
- Example: `std::shared_mutex` → `pthread_rwlock`



# Latch Implementations

- **Approach #3: Reader-Writer Latches**

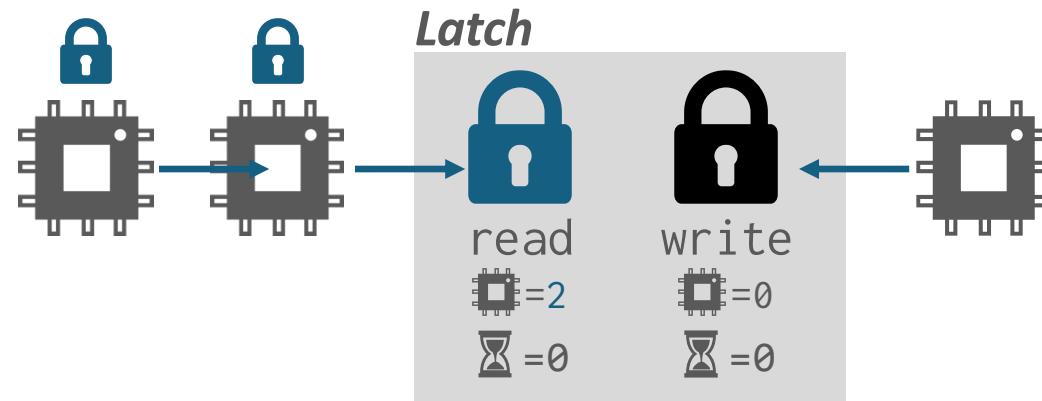
- Allows for concurrent readers. Must manage read/write queues to avoid starvation.
- Can be implemented on top of spinlocks.
- Example: `std::shared_mutex` → `pthread_rwlock`



# Latch Implementations

- **Approach #3: Reader-Writer Latches**

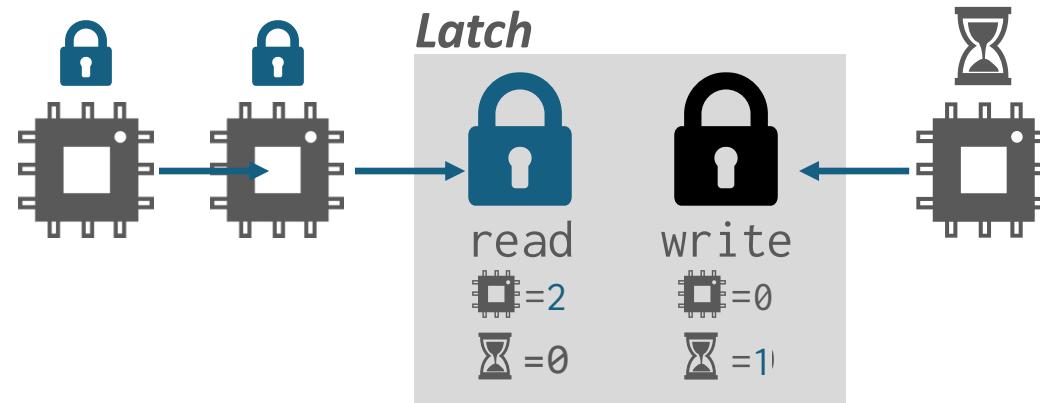
- Allows for concurrent readers. Must manage read/write queues to avoid starvation.
- Can be implemented on top of spinlocks.
- Example: `std::shared_mutex` → `pthread_rwlock`



# Latch Implementations

- **Approach #3: Reader-Writer Latches**

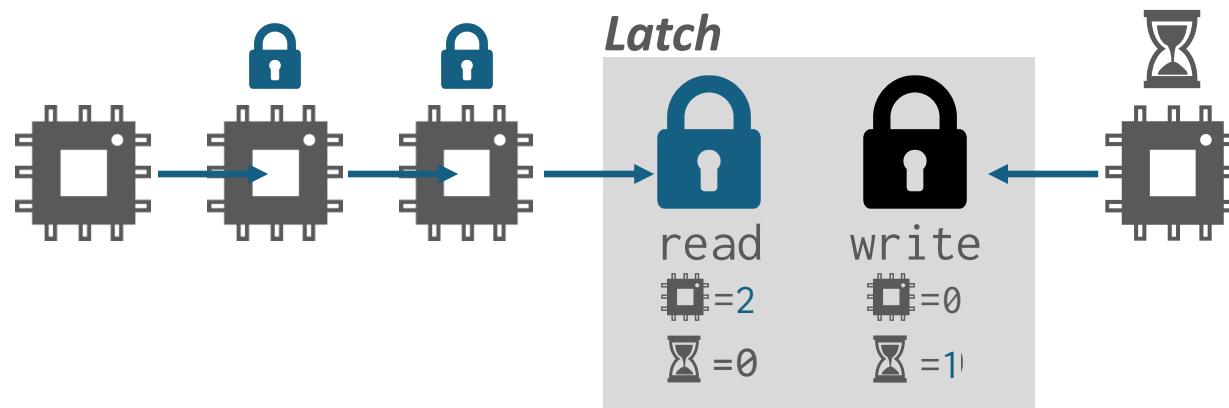
- Allows for concurrent readers. Must manage read/write queues to avoid starvation.
- Can be implemented on top of spinlocks.
- Example: `std::shared_mutex` → `pthread_rwlock`



# Latch Implementations

- **Approach #3: Reader-Writer Latches**

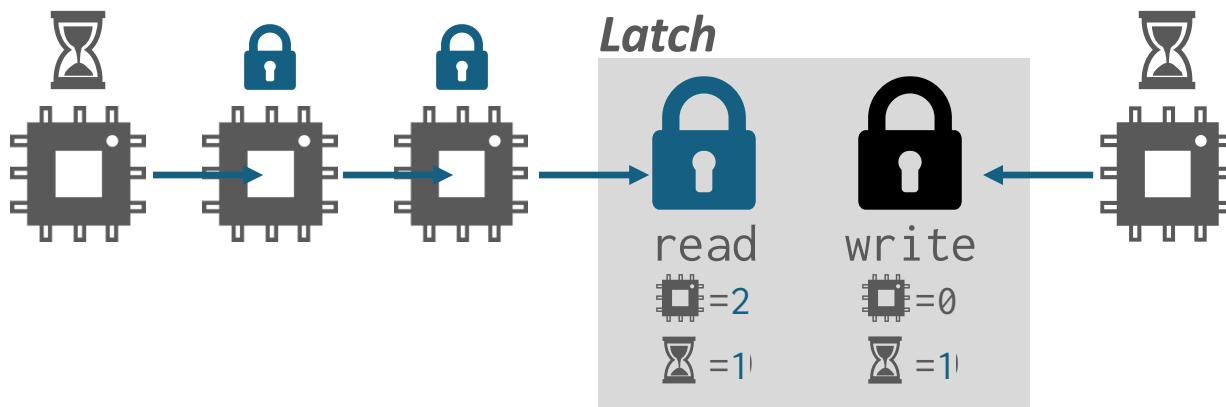
- Allows for concurrent readers. Must manage read/write queues to avoid starvation.
- Can be implemented on top of spinlocks.
- Example: `std::shared_mutex` → `pthread_rwlock`



# Latch Implementations

- **Approach #3: Reader-Writer Latches**

- Allows for concurrent readers. Must manage read/write queues to avoid starvation.
- Can be implemented on top of spinlocks.
- Example: `std::shared_mutex` → `pthread_rwlock`



# Compare-and-Swap

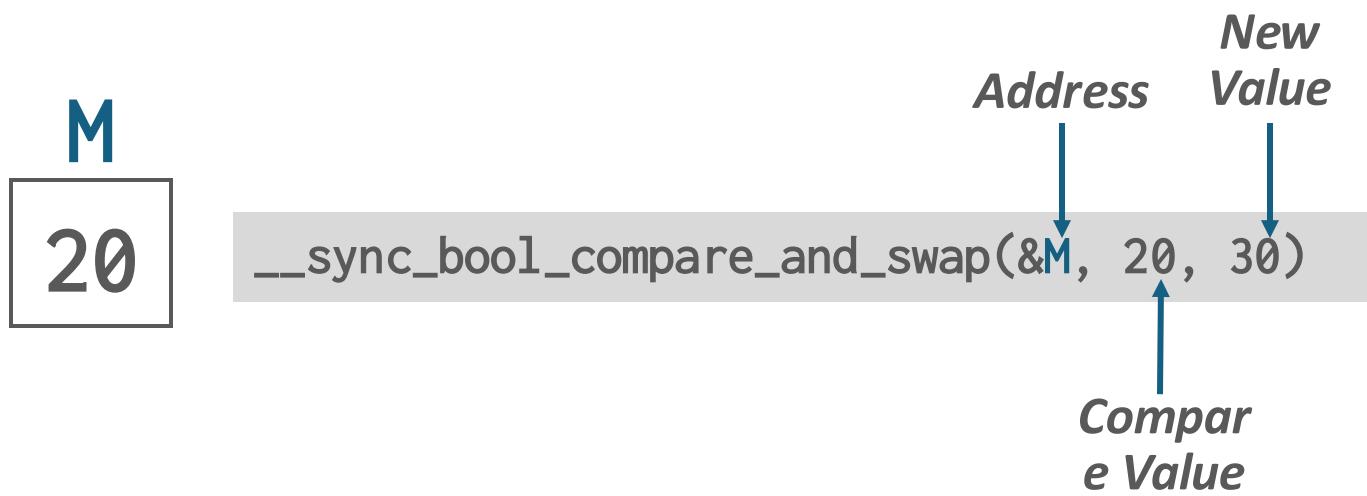
- Atomic instruction that compares contents of a memory location **M** to a given value **V**
  - If values are equal, installs new given value **V'** in **M**
  - Otherwise, operation fails

**M**  
**20**

```
__sync_bool_compare_and_swap(&M, 20, 30)
```

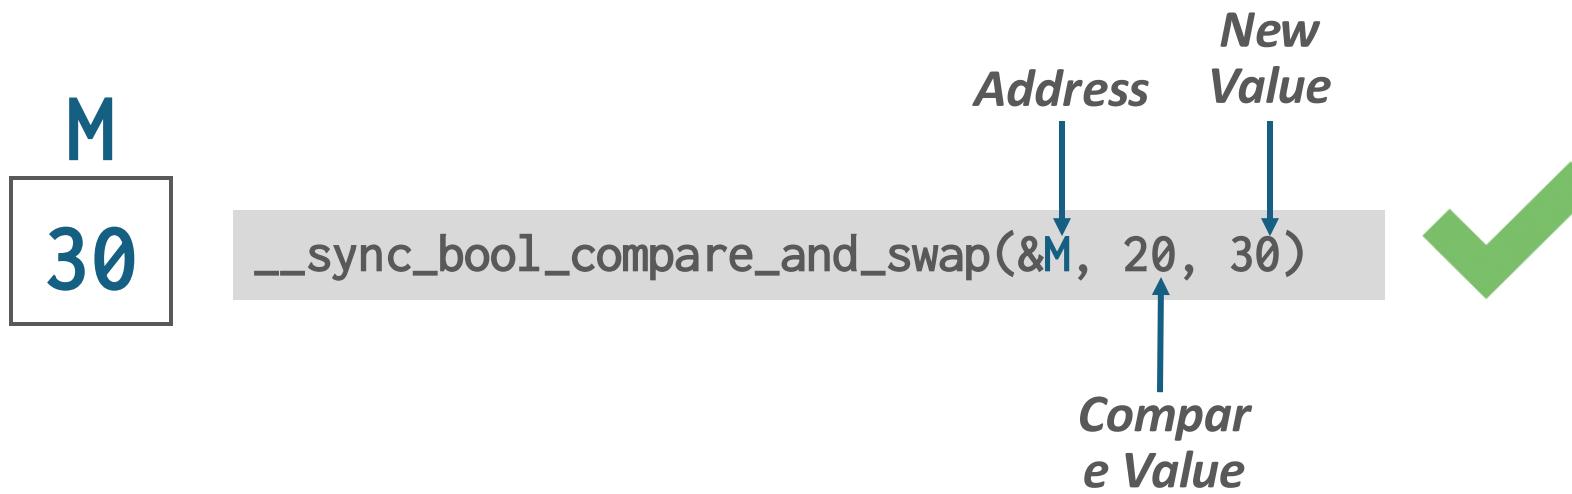
# Compare-and-Swap

- Atomic instruction that compares contents of a memory location **M** to a given value **V**
  - If values are equal, installs new given value **V'** in **M**
  - Otherwise, operation fails



# Compare-and-Swap

- Atomic instruction that compares contents of a memory location **M** to a given value **V**
  - If values are equal, installs new given value **V'** in **M**
  - Otherwise, operation fails



# Hash Table Latching

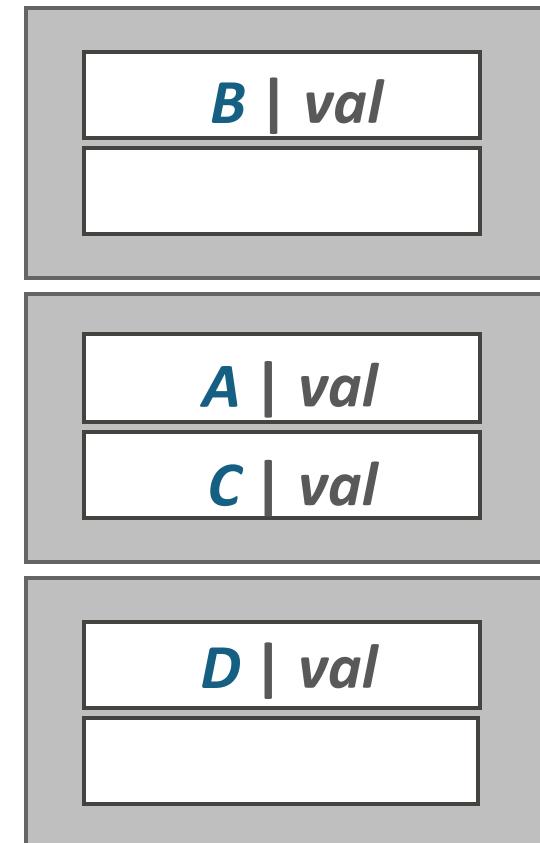
# Hash Table Latching

- Easy to support concurrent access due to the limited ways threads access the data structure.
  - All threads move in the same direction and only access a single page/slot at a time.
  - Deadlocks are not possible.
- To resize the table, take a global write latch on the entire table (e.g., in the header page).

# Hash Table Latching

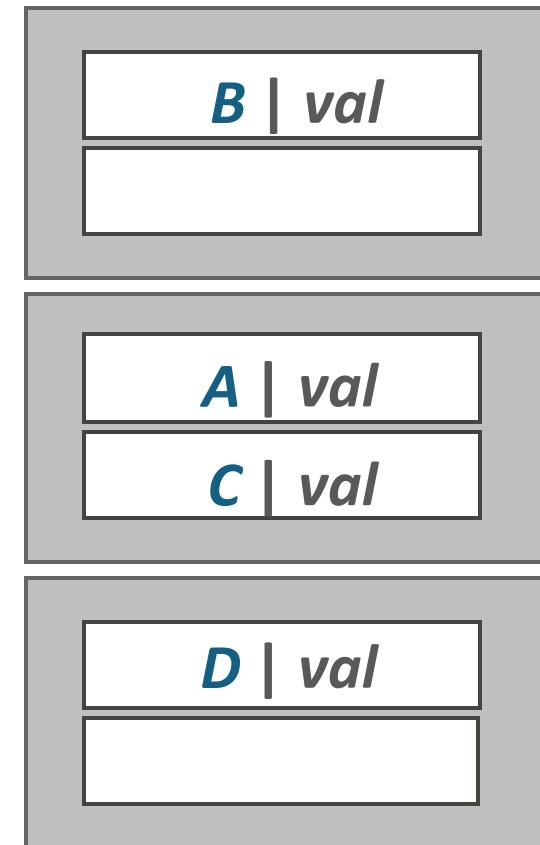
- **Approach #1: Page/Block-level Latches**
  - Each page/block has its own reader-writer latch that protects its entire contents.
  - Threads acquire either a read or write latch before they access a page/block.
- **Approach #2: Slot Latches**
  - Each slot has its own latch.
  - Can use a single-mode latch to reduce meta-data and computational overhead.

# Hash Table - Page/Block Latches

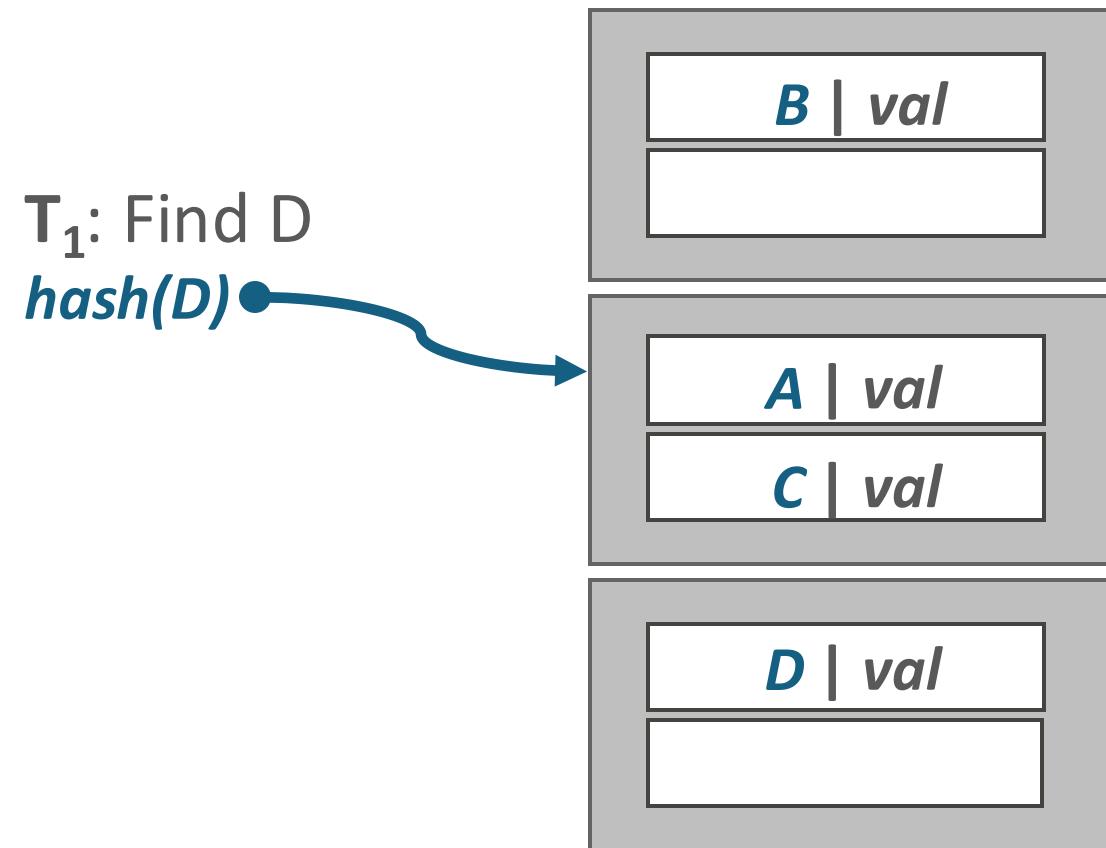


# Hash Table - Page/Block Latches

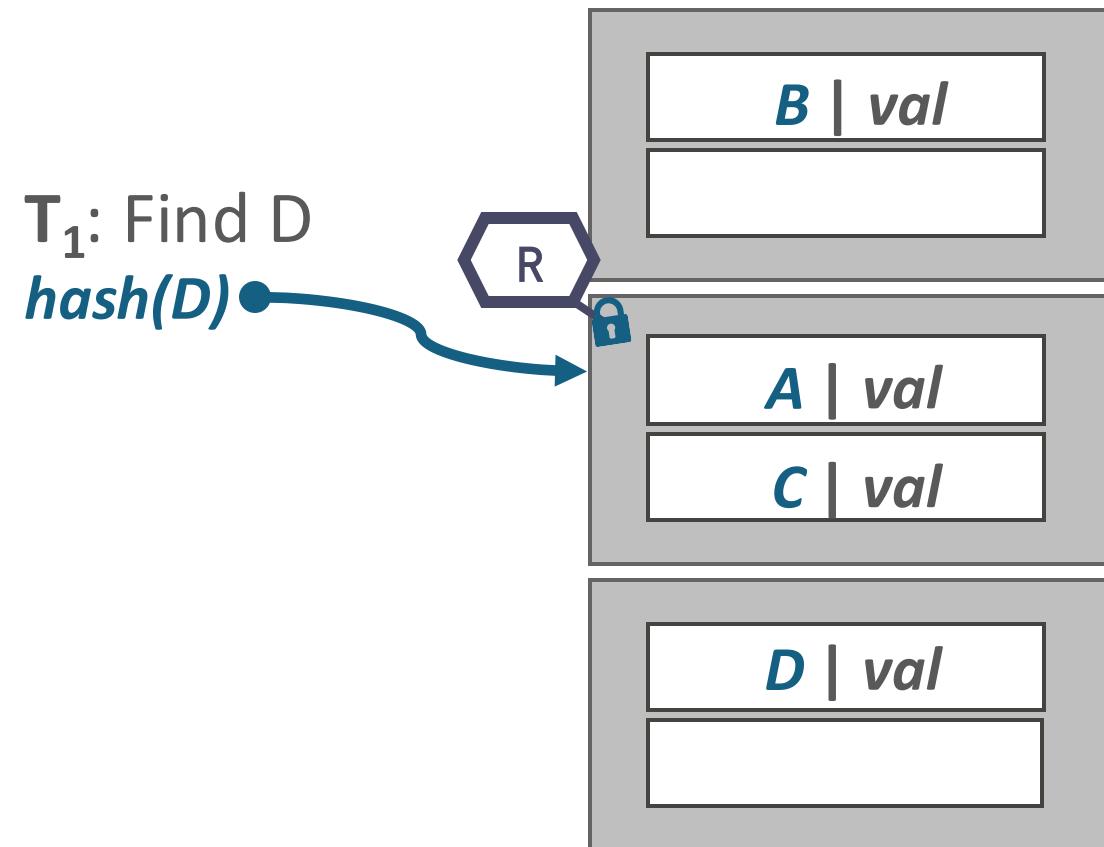
$T_1$ : Find D  
 $hash(D)$



# Hash Table - Page/Block Latches

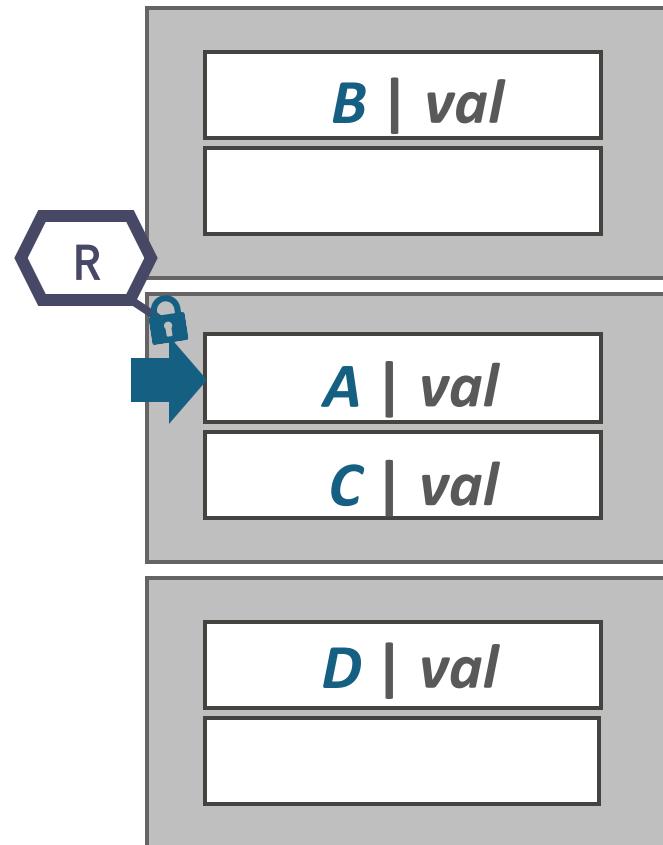


# Hash Table - Page/Block Latches



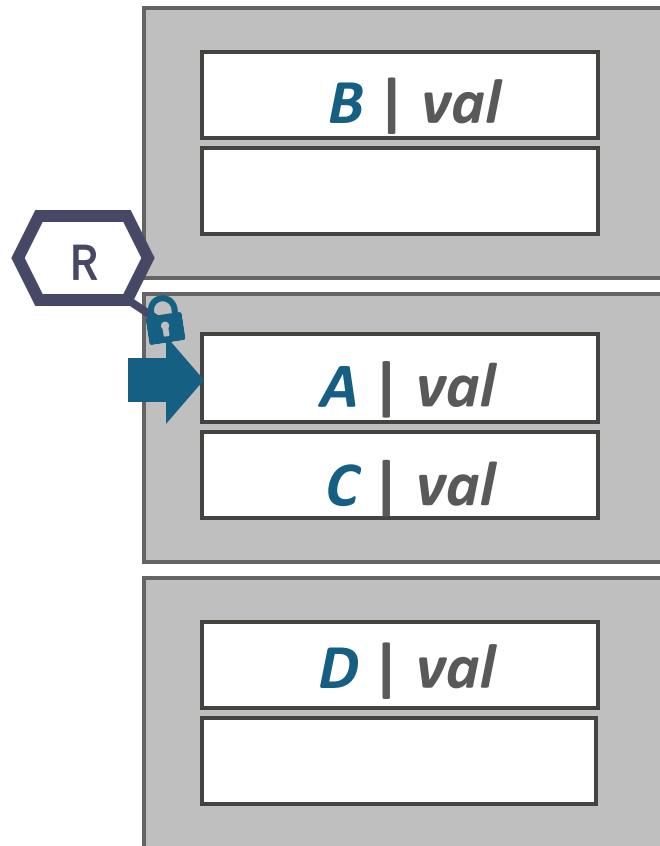
# Hash Table - Page/Block Latches

$T_1$ : Find D  
 $hash(D)$



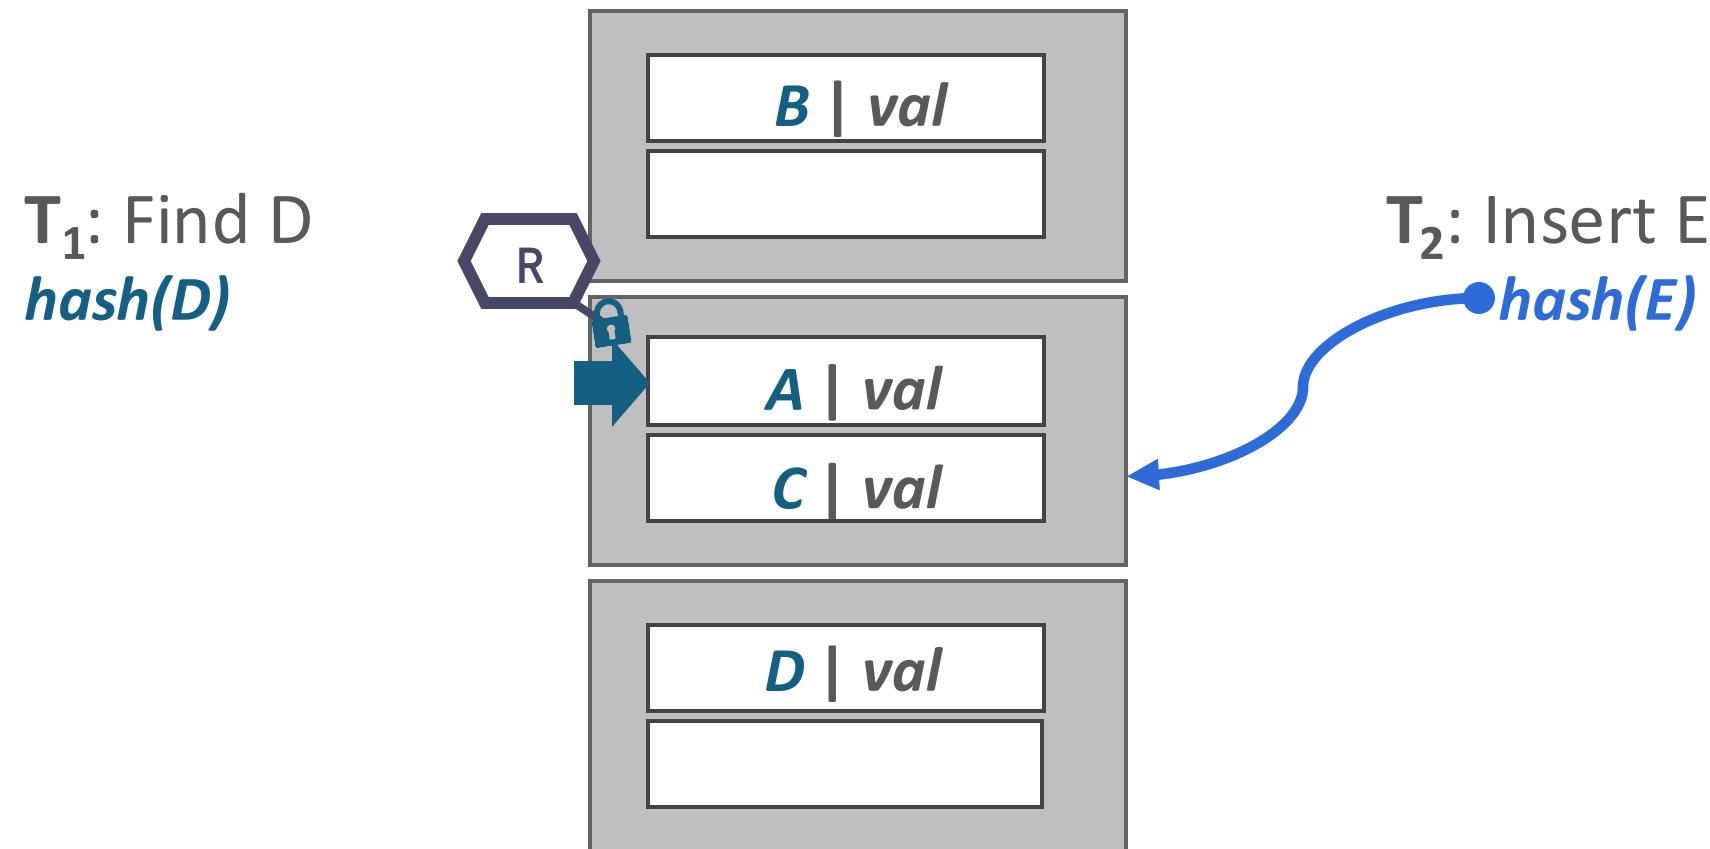
# Hash Table - Page/Block Latches

$T_1$ : Find D  
 $hash(D)$

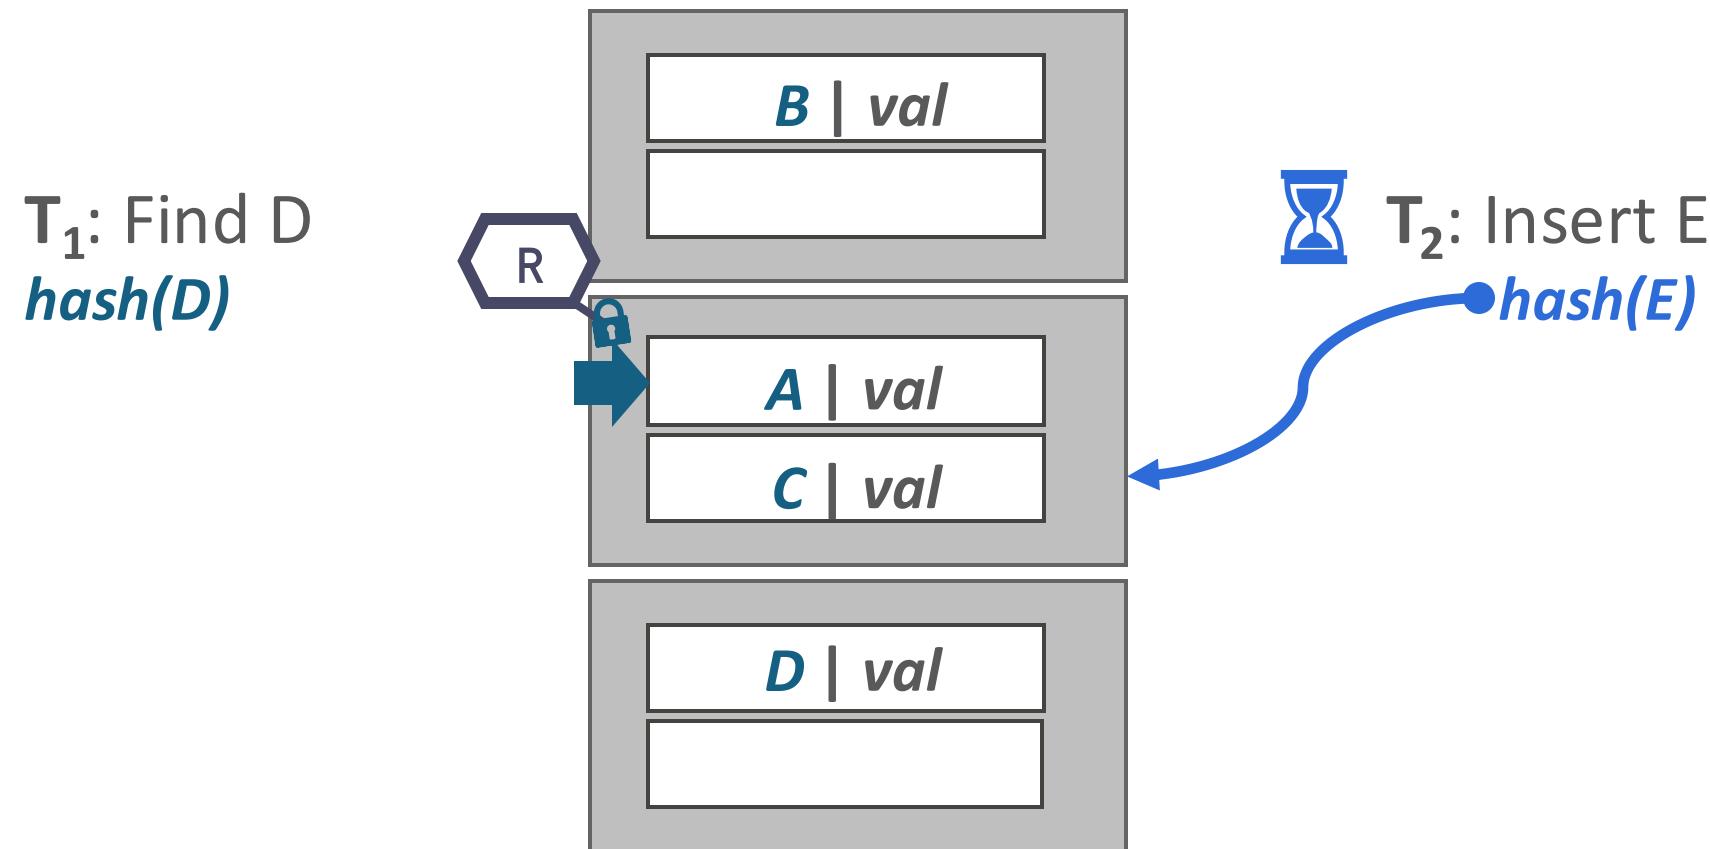


$T_2$ : Insert E  
 $hash(E)$

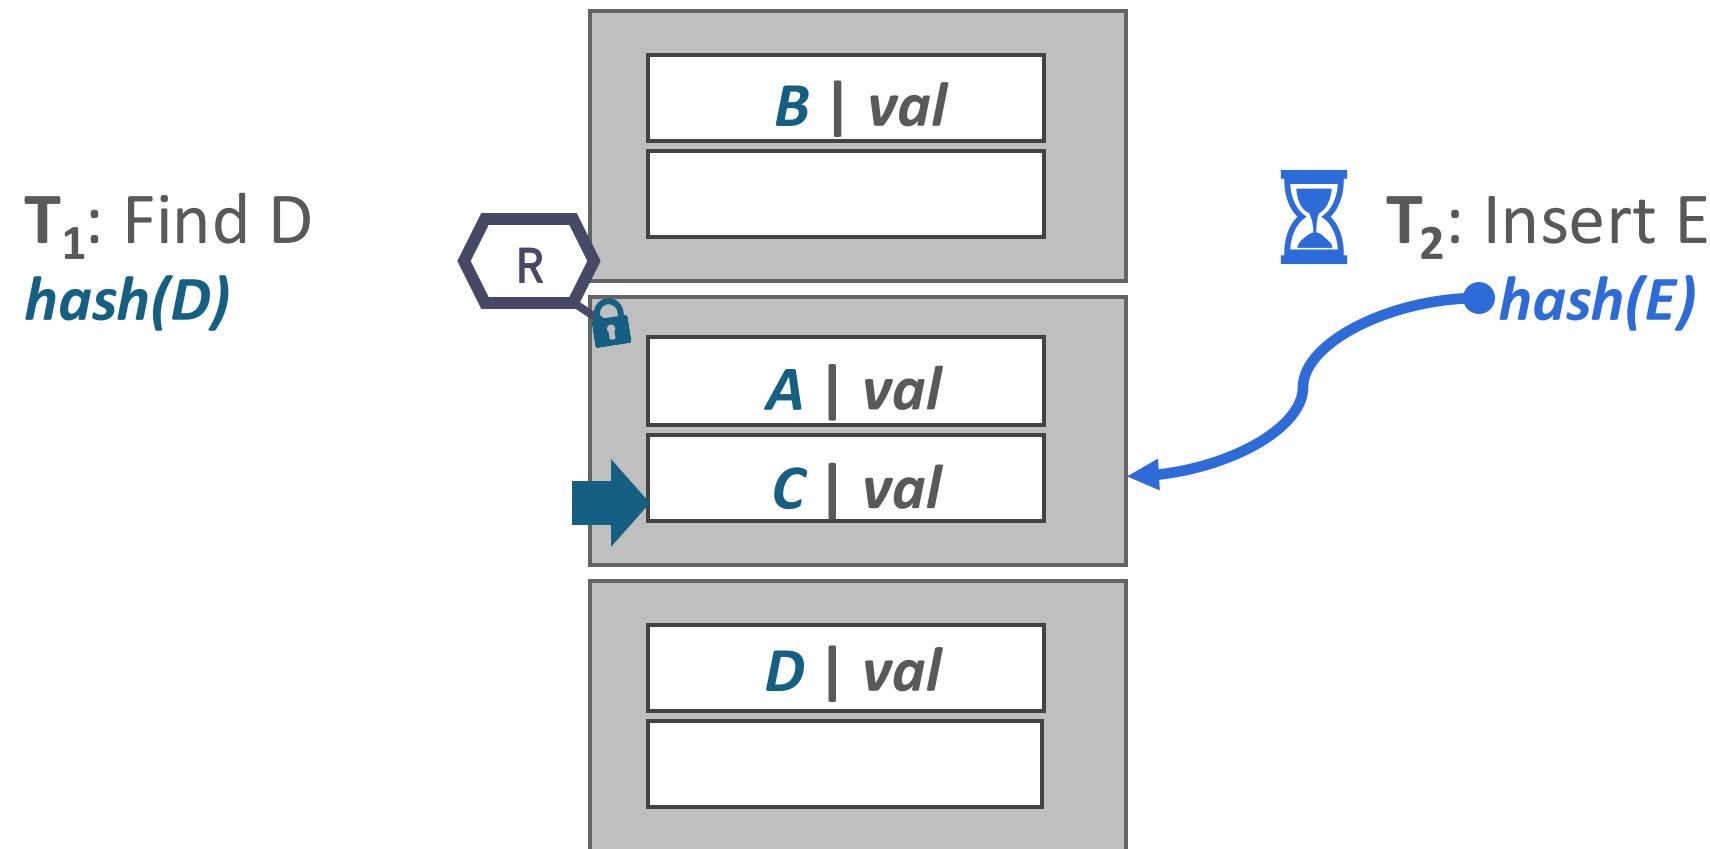
# Hash Table - Page/Block Latches



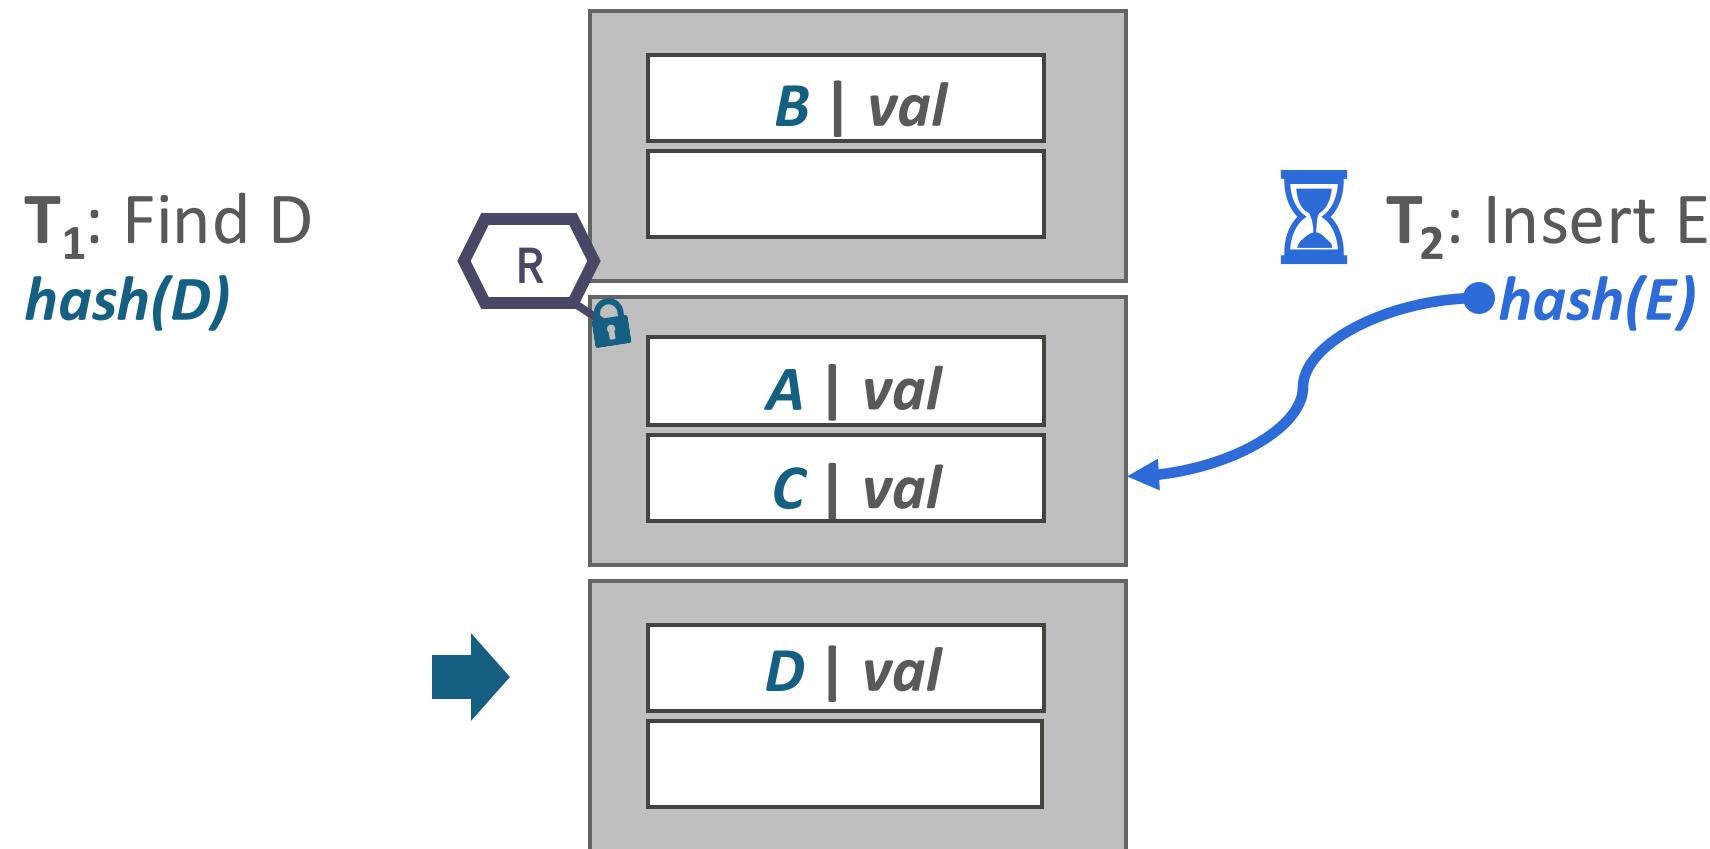
# Hash Table - Page/Block Latches



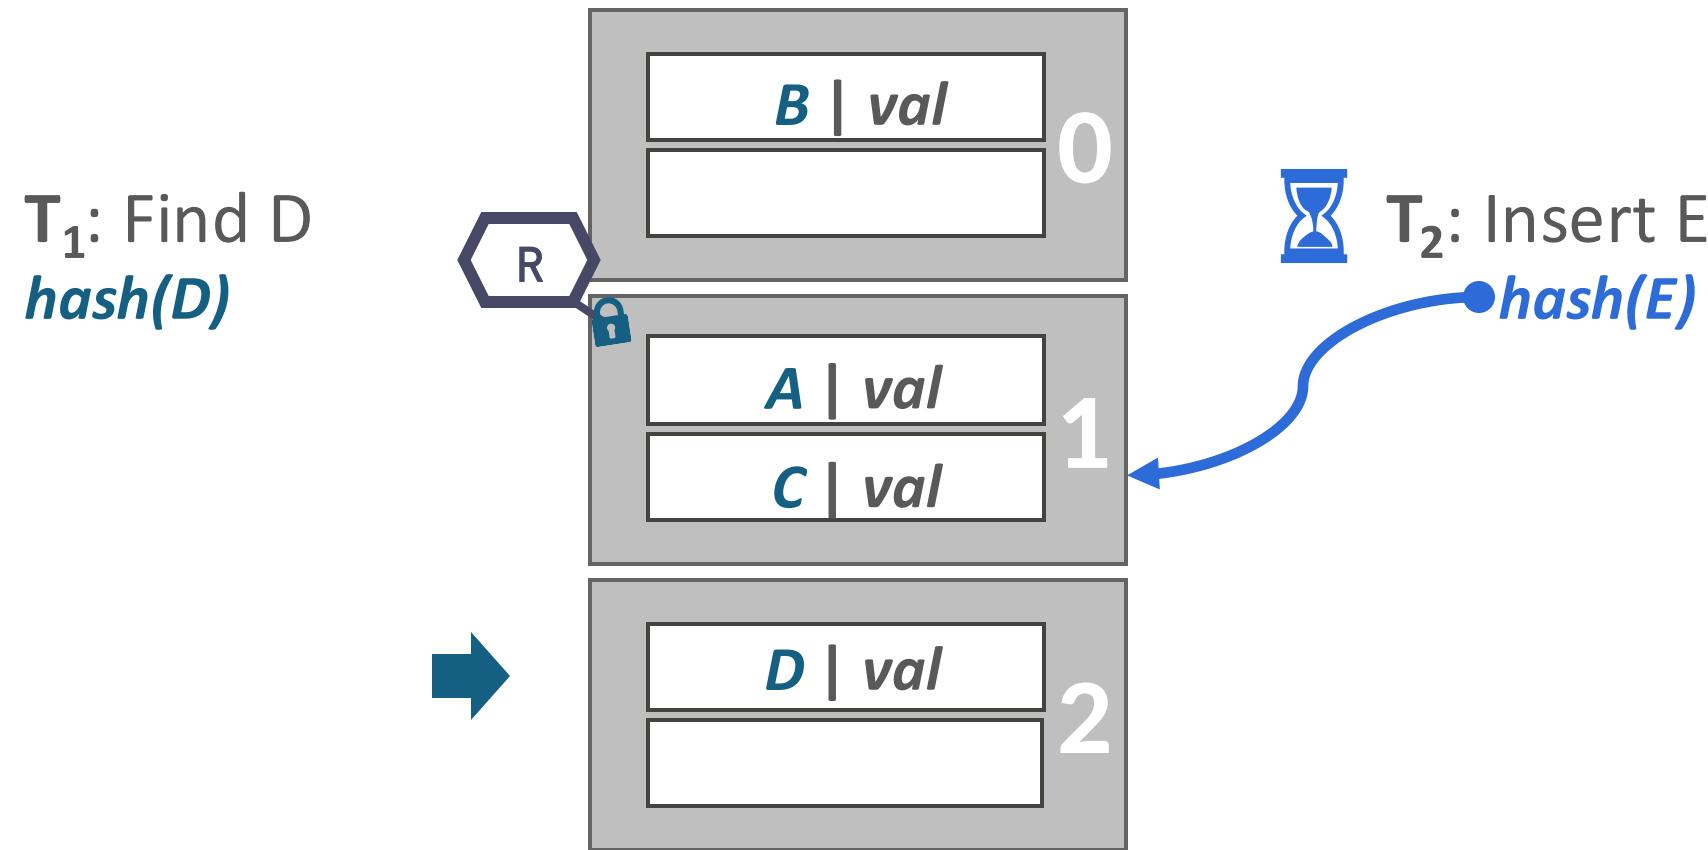
# Hash Table - Page/Block Latches



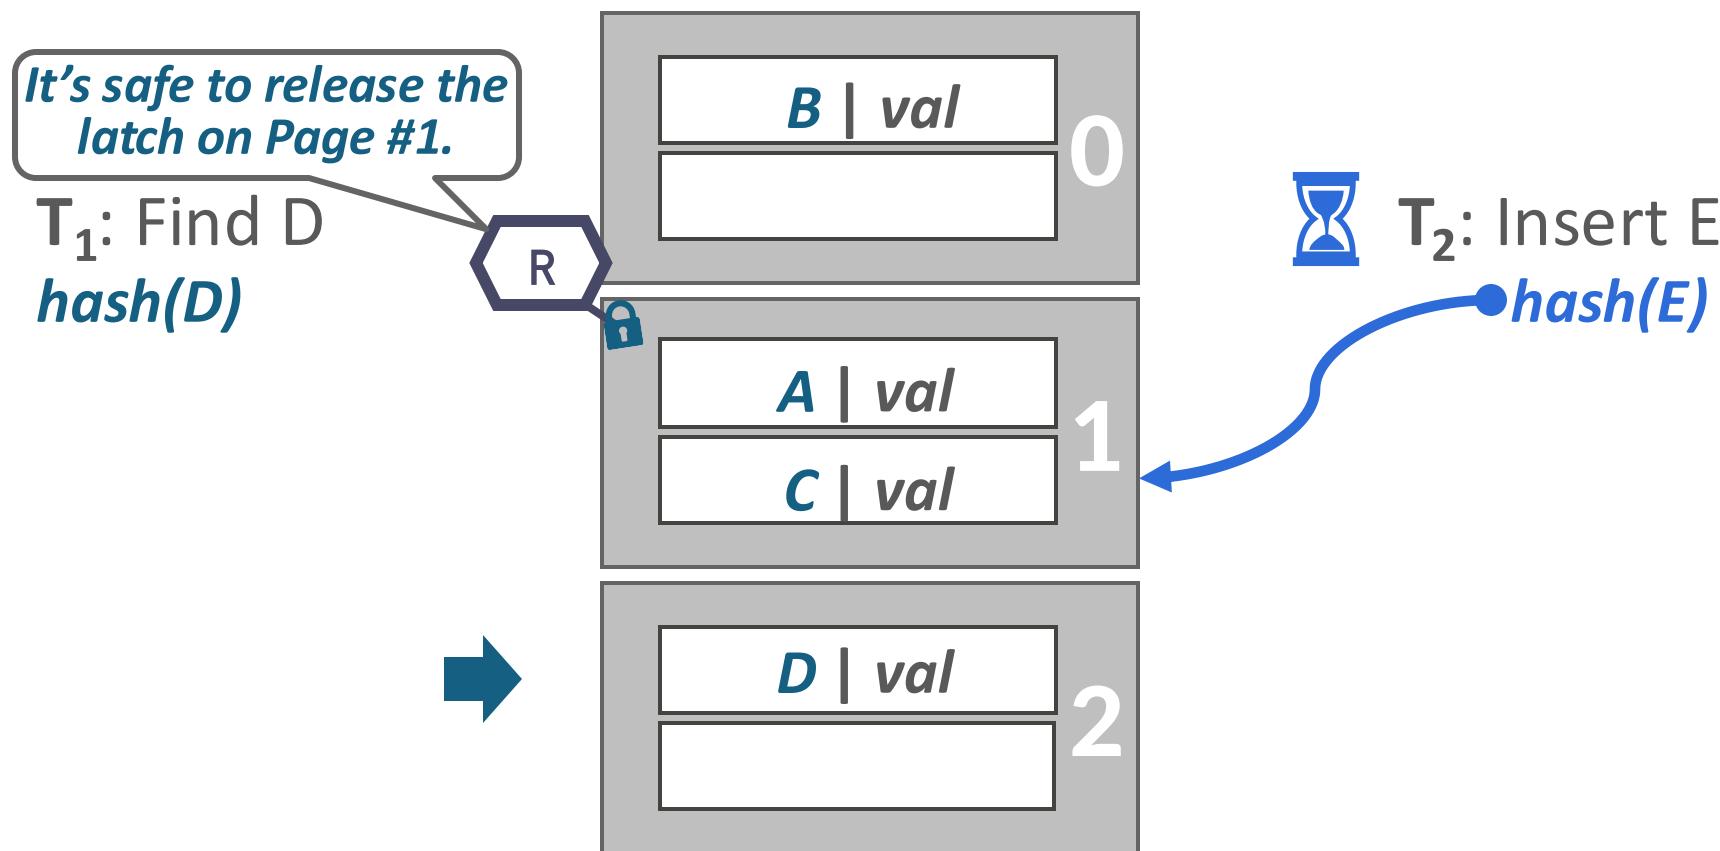
# Hash Table - Page/Block Latches



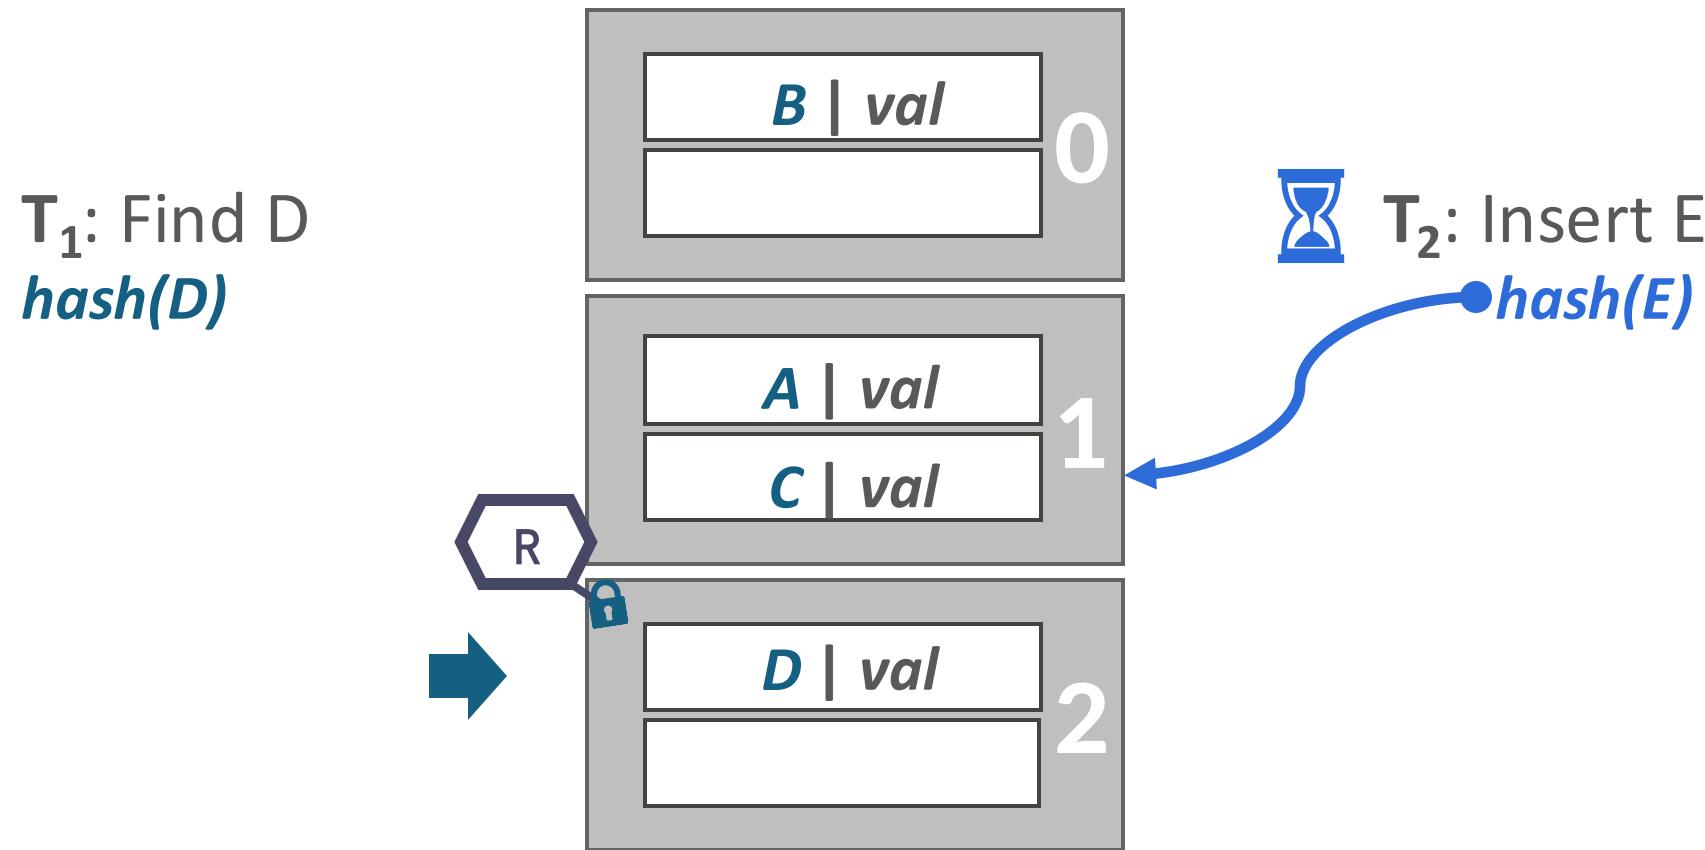
# Hash Table - Page/Block Latches



# Hash Table - Page/Block Latches

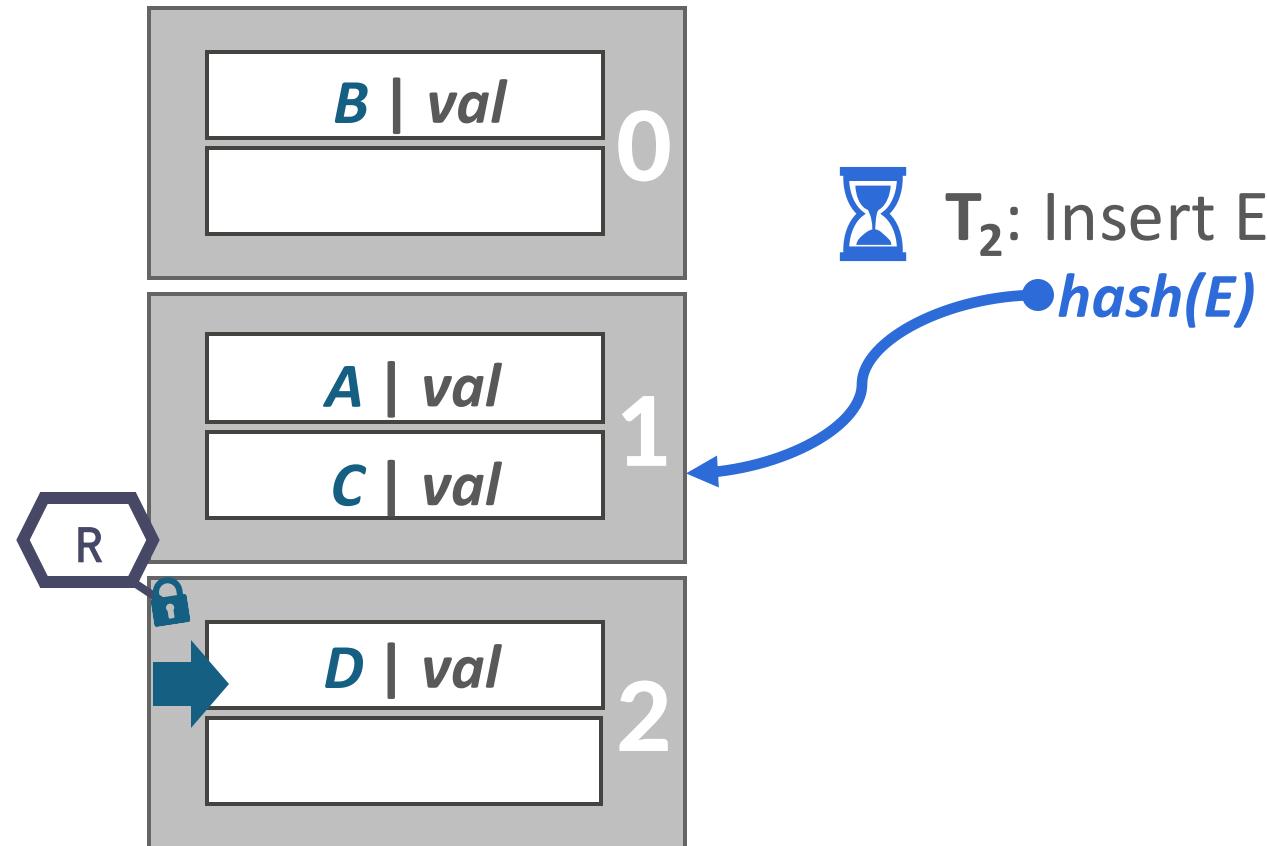


# Hash Table - Page/Block Latches



# Hash Table - Page/Block Latches

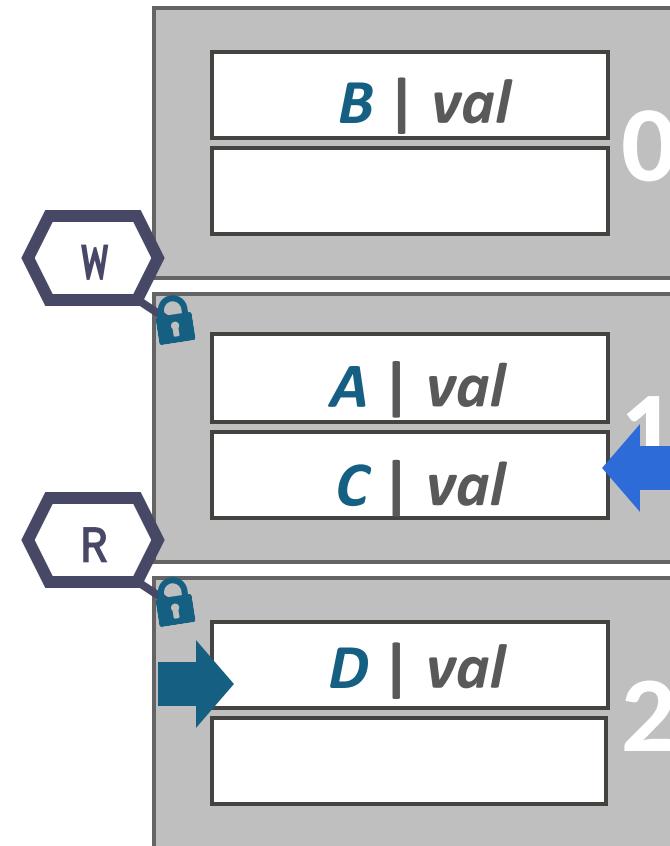
$T_1$ : Find D  
 $hash(D)$



# Hash Table - Page/Block Latches

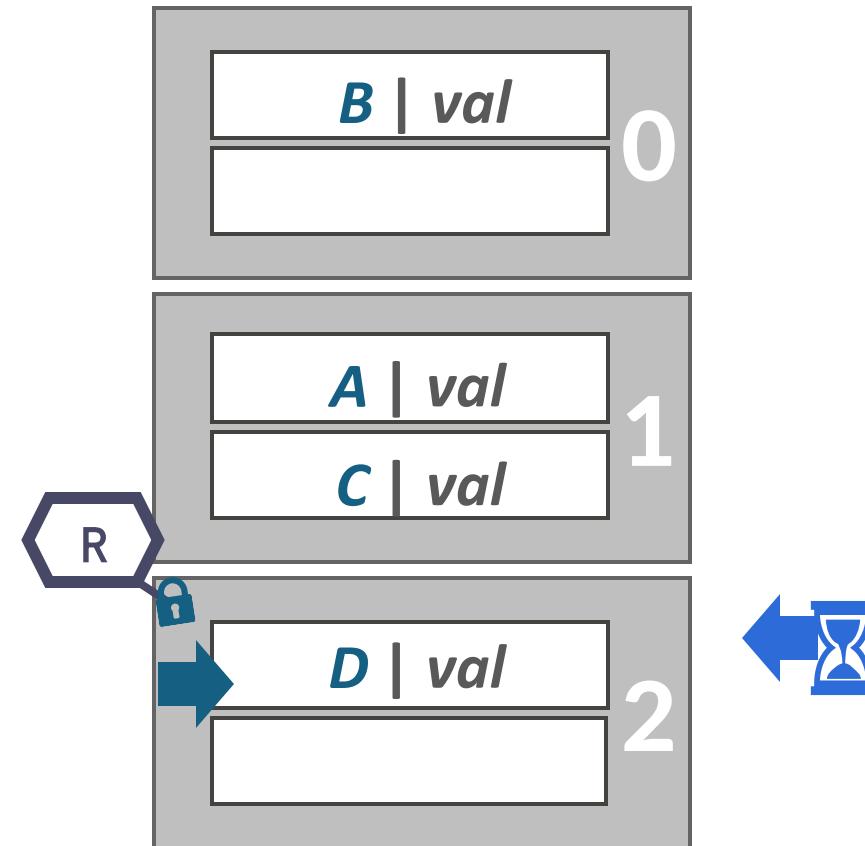
$T_1$ : Find D  
 $hash(D)$

$T_2$ : Insert E  
 $hash(E)$



# Hash Table - Page/Block Latches

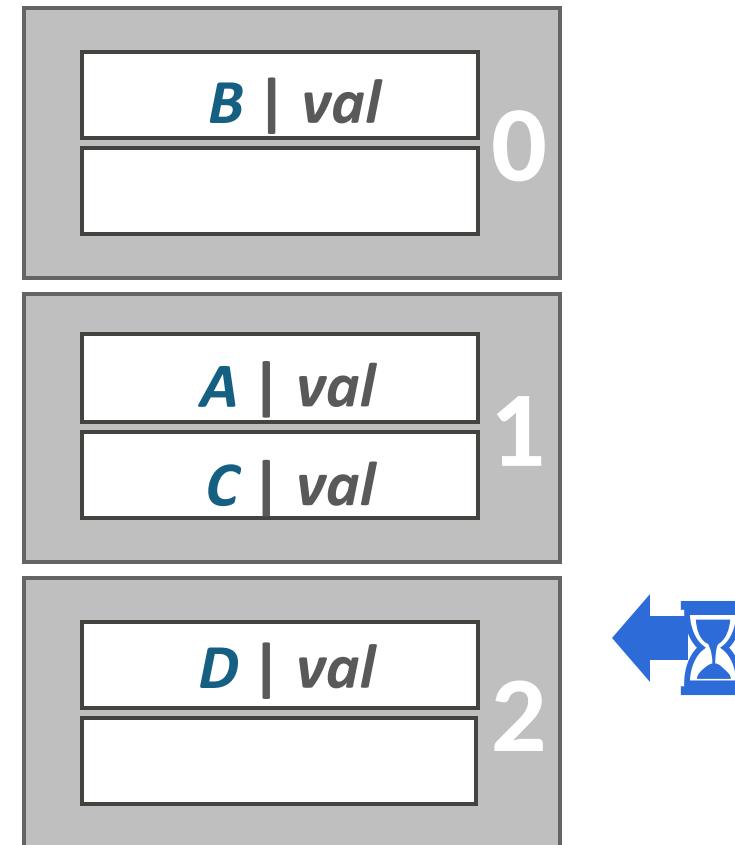
$T_1$ : Find D  
 $hash(D)$



$T_2$ : Insert E  
 $hash(E)$

# Hash Table - Page/Block Latches

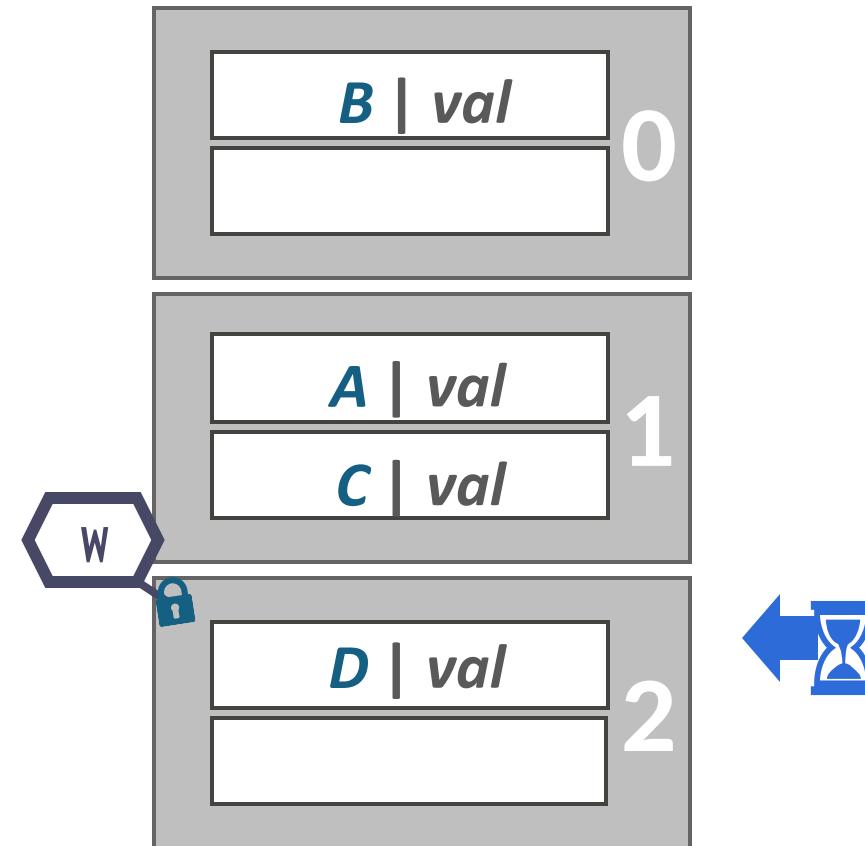
$T_1$ : Find D  
 $hash(D)$



$T_2$ : Insert E  
 $hash(E)$

# Hash Table - Page/Block Latches

$T_1$ : Find D  
 $hash(D)$



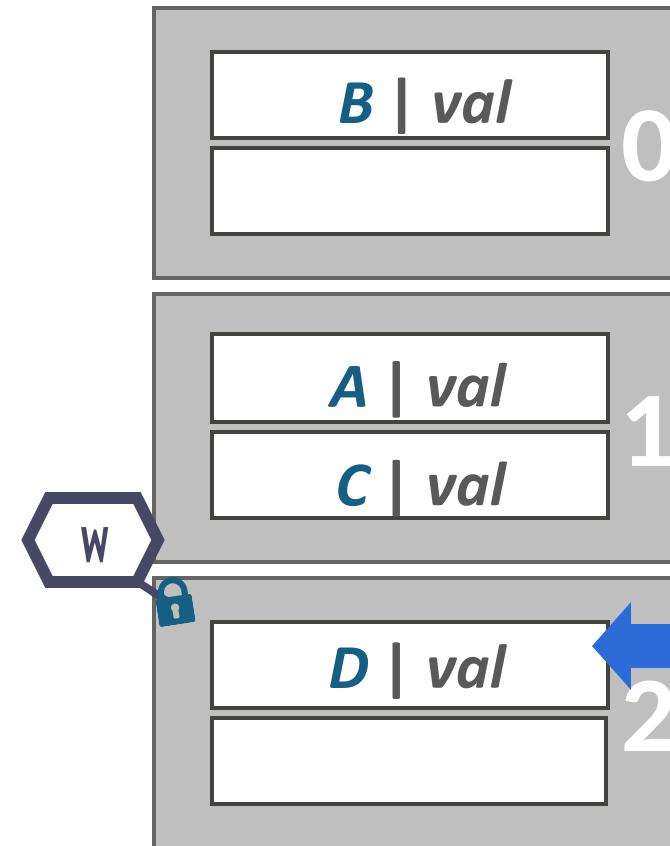
$T_2$ : Insert E  
 $hash(E)$



# Hash Table - Page/Block Latches

$T_1$ : Find D  
 $hash(D)$

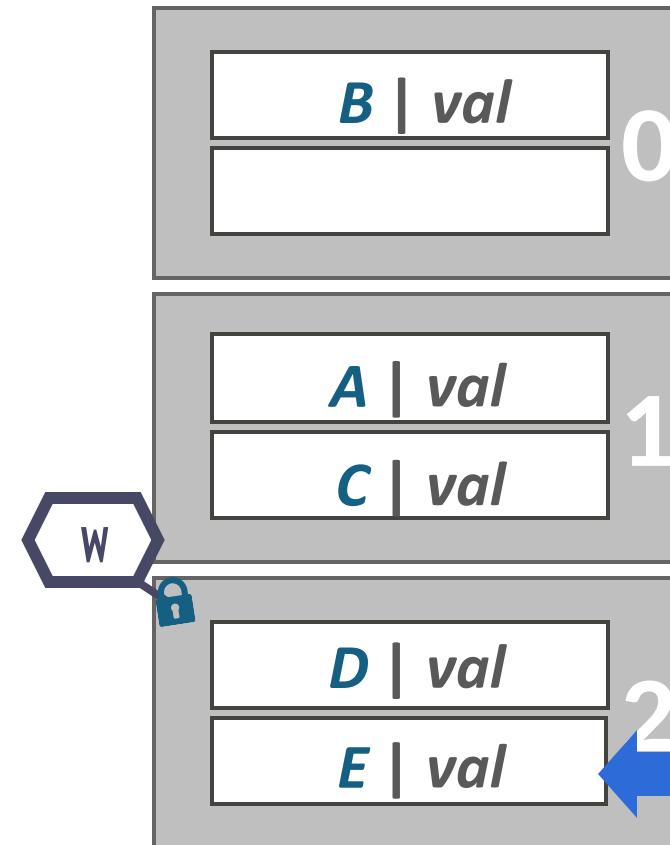
$T_2$ : Insert E  
 $hash(E)$



# Hash Table - Page/Block Latches

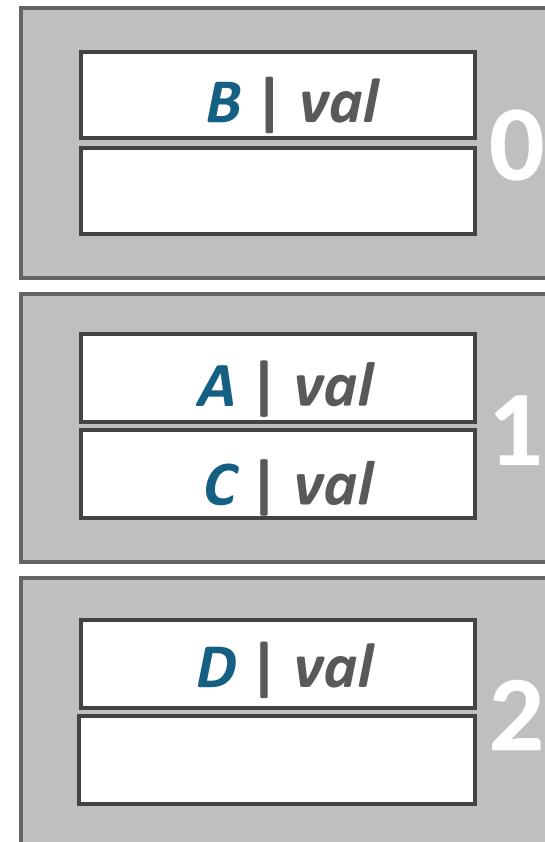
$T_1$ : Find D  
 $hash(D)$

$T_2$ : Insert E  
 $hash(E)$



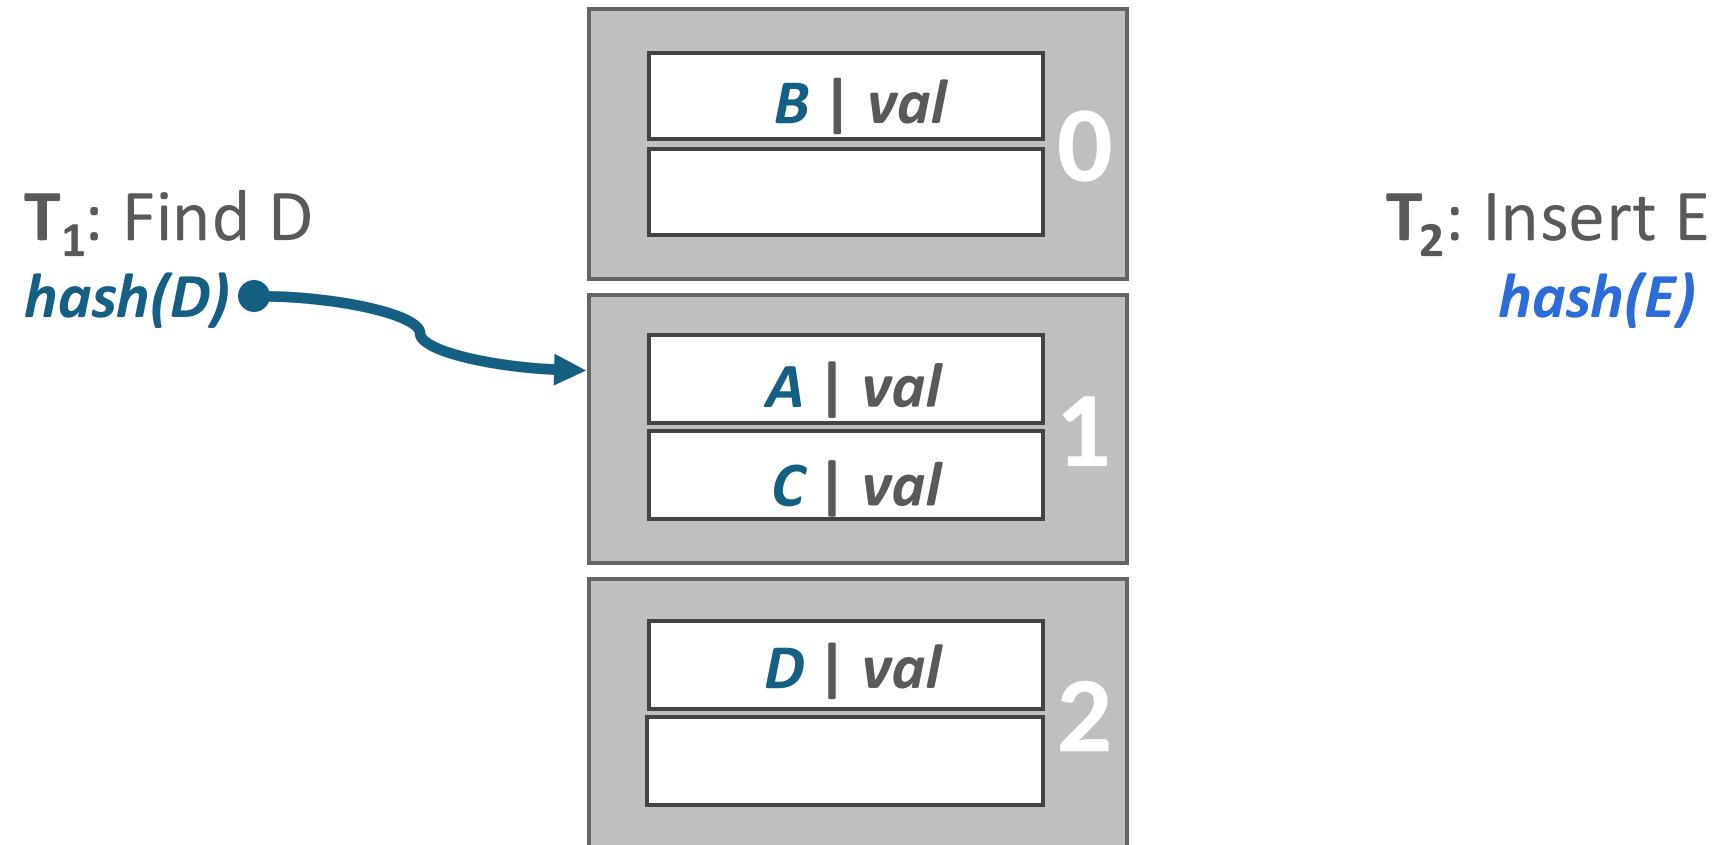
# Hash Table – Slot Latches

$T_1$ : Find D  
 $hash(D)$

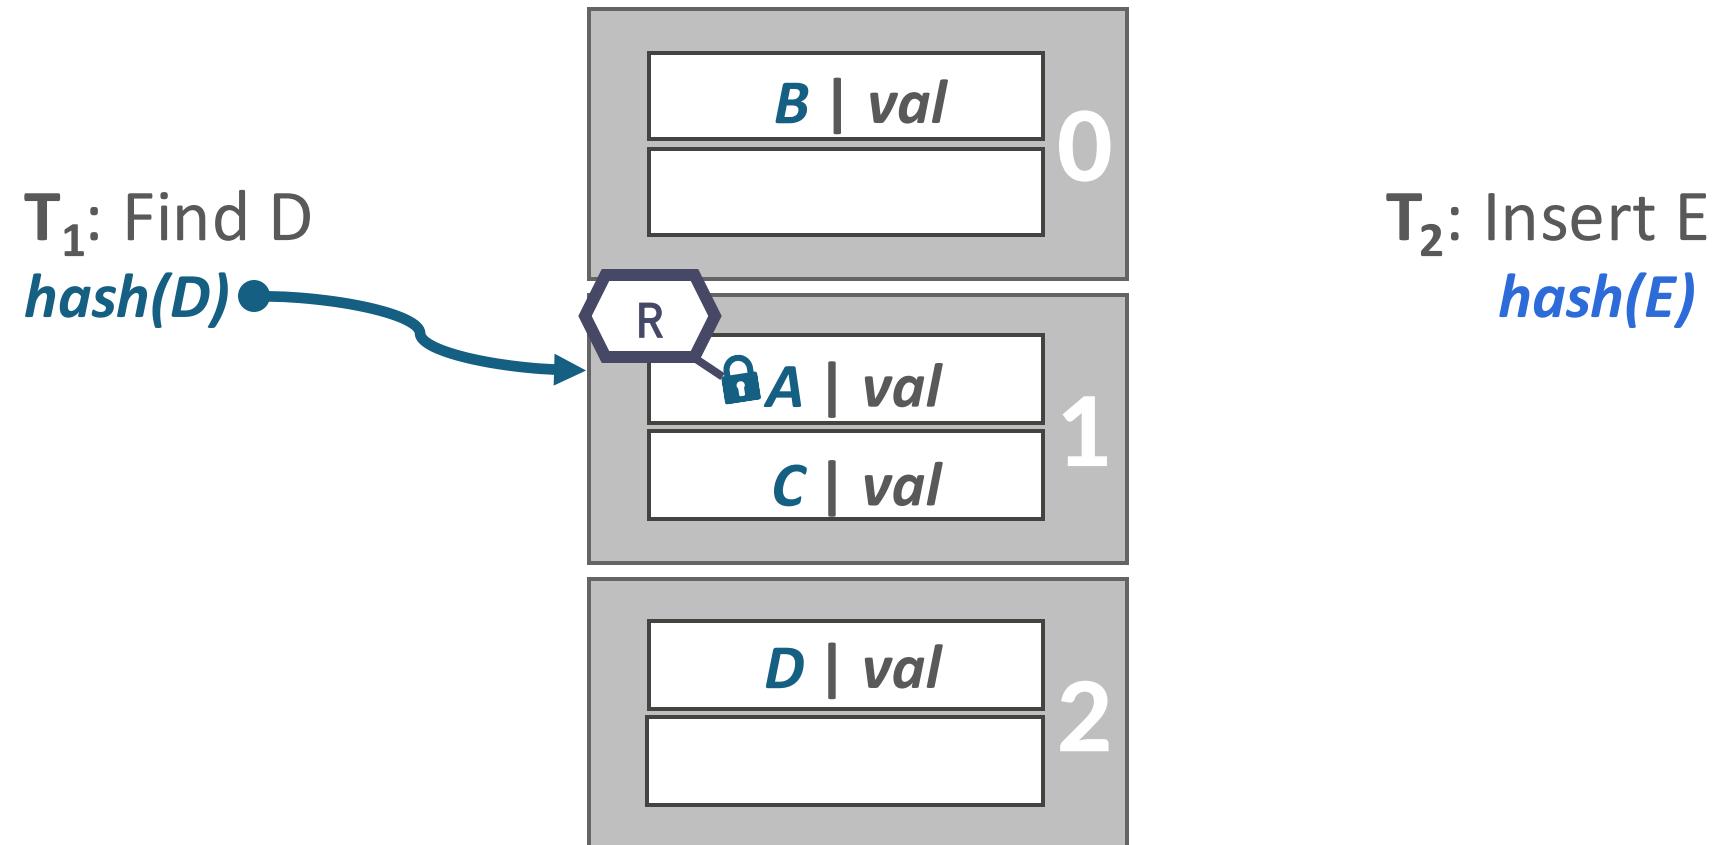


$T_2$ : Insert E  
 $hash(E)$

# Hash Table – Slot Latches

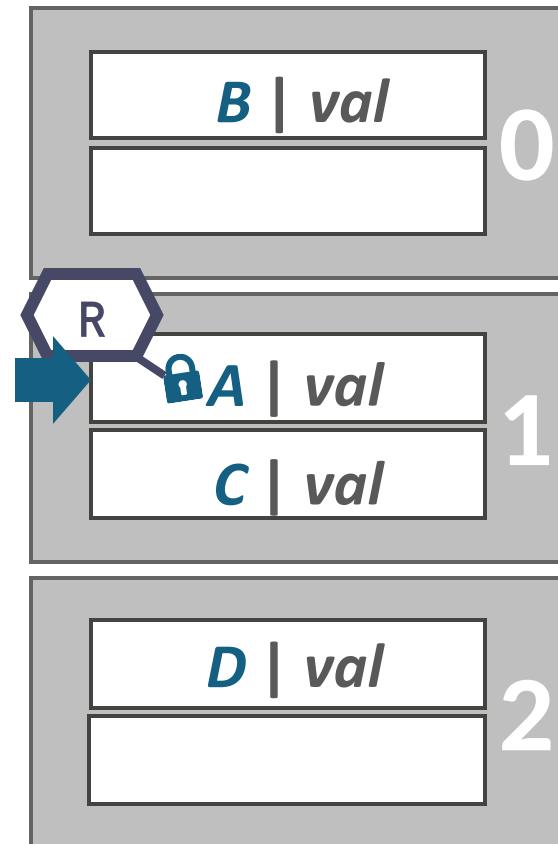


# Hash Table – Slot Latches



# Hash Table – Slot Latches

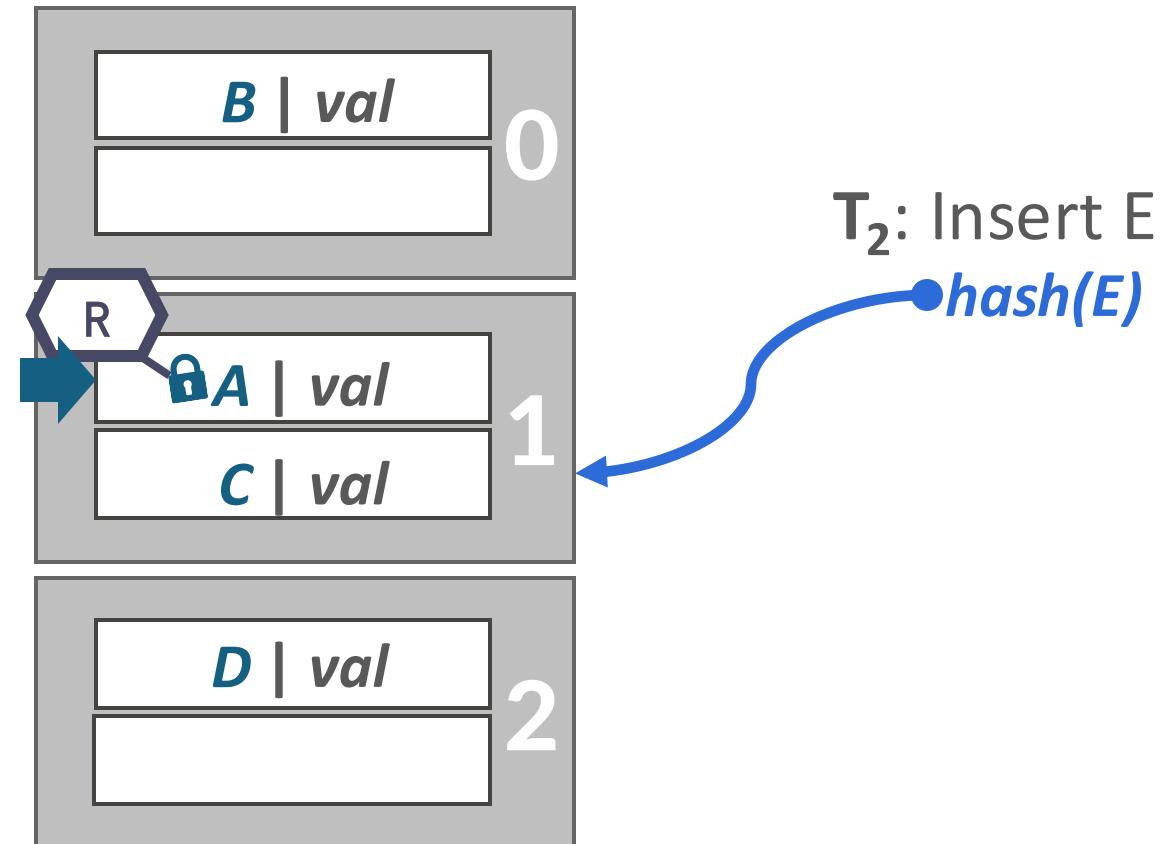
$T_1$ : Find D  
 $hash(D)$



$T_2$ : Insert E  
 $hash(E)$

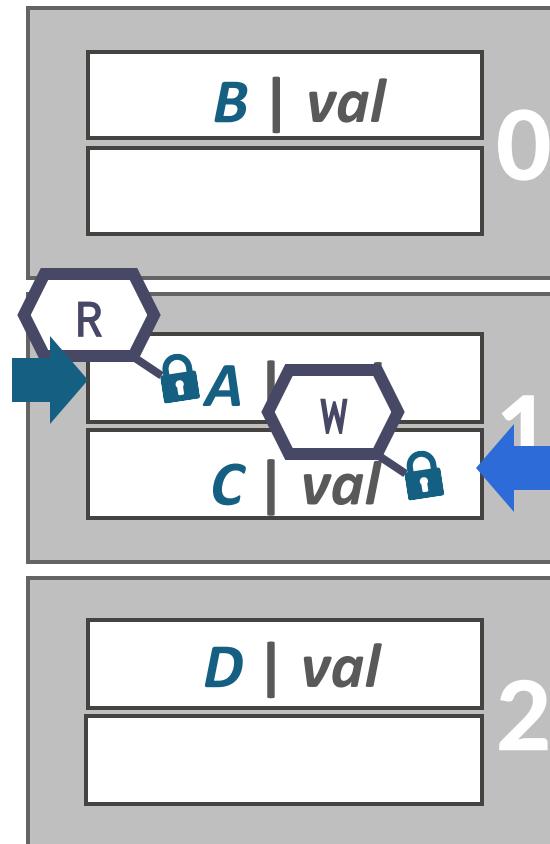
# Hash Table – Slot Latches

$T_1$ : Find D  
 $hash(D)$



# Hash Table – Slot Latches

$T_1$ : Find D  
 $hash(D)$

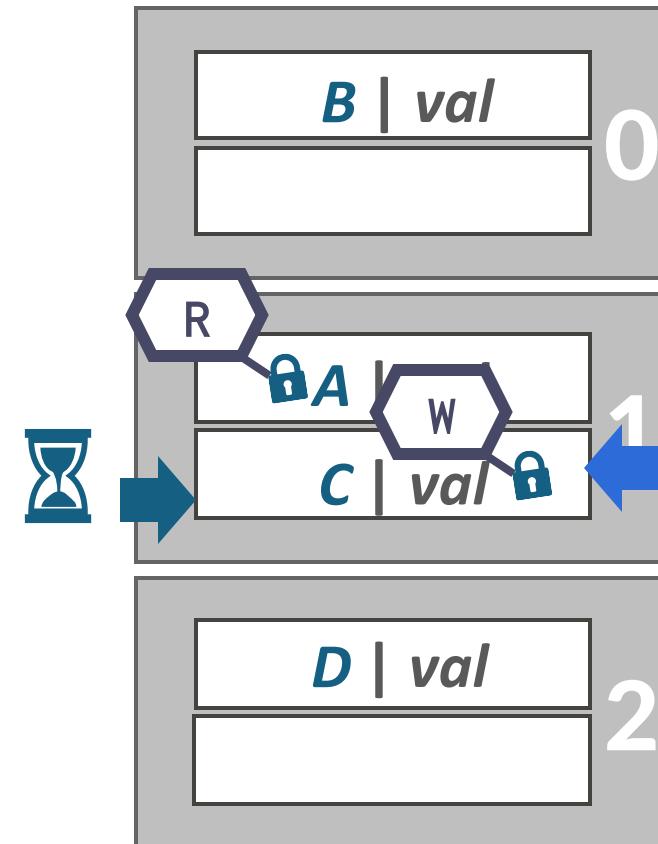


$T_2$ : Insert E  
 $hash(E)$

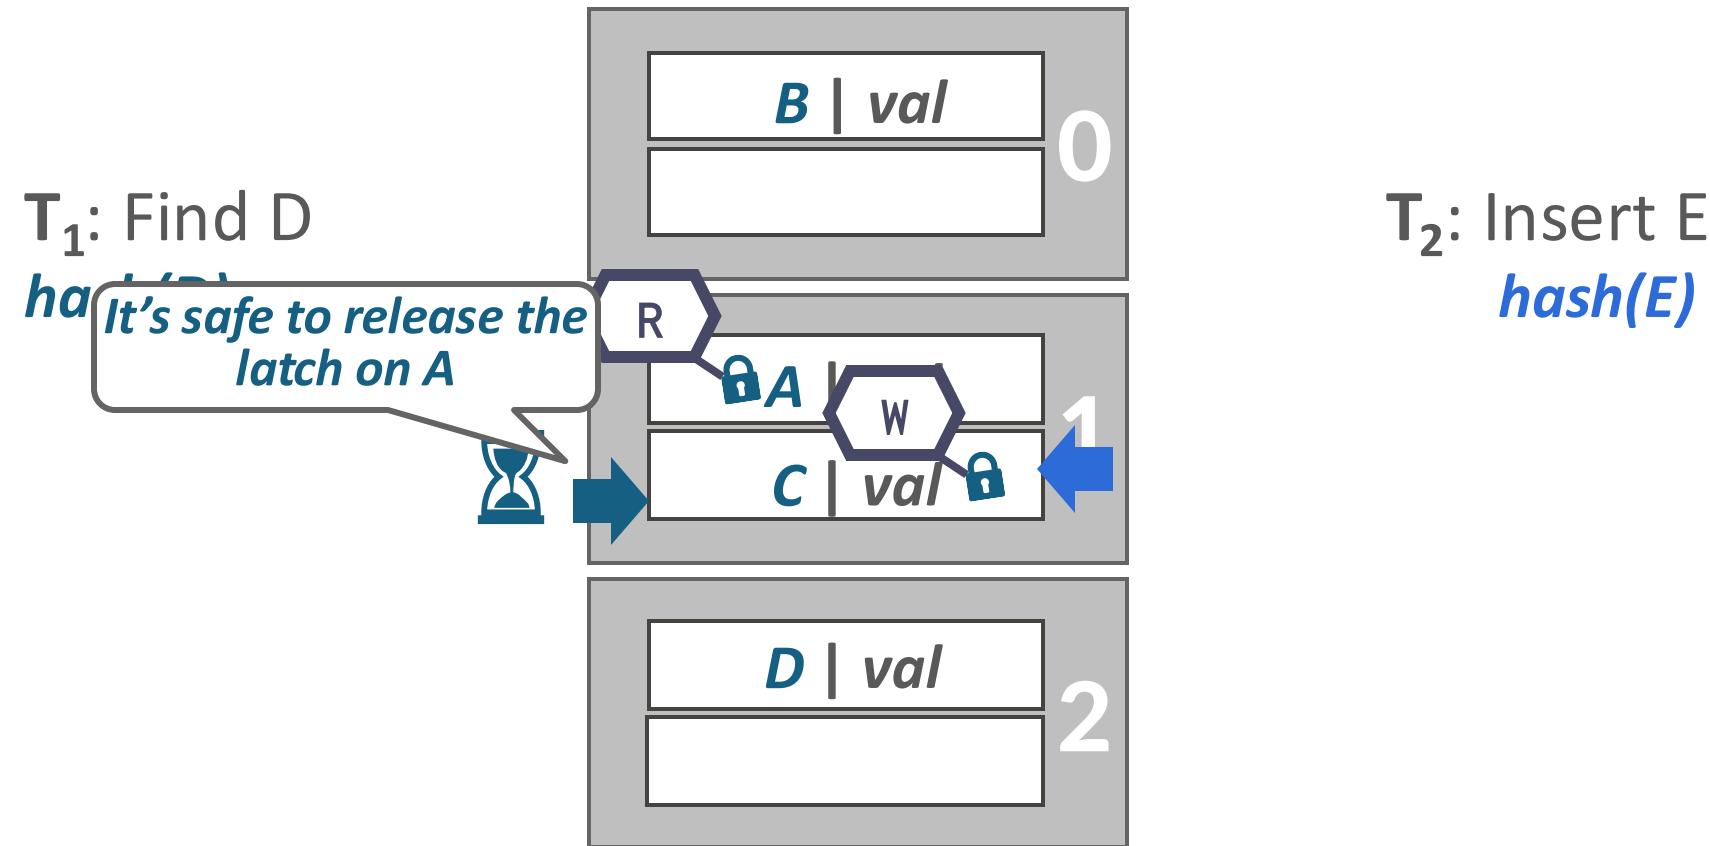
# Hash Table – Slot Latches

$T_1$ : Find D  
 $hash(D)$

$T_2$ : Insert E  
 $hash(E)$



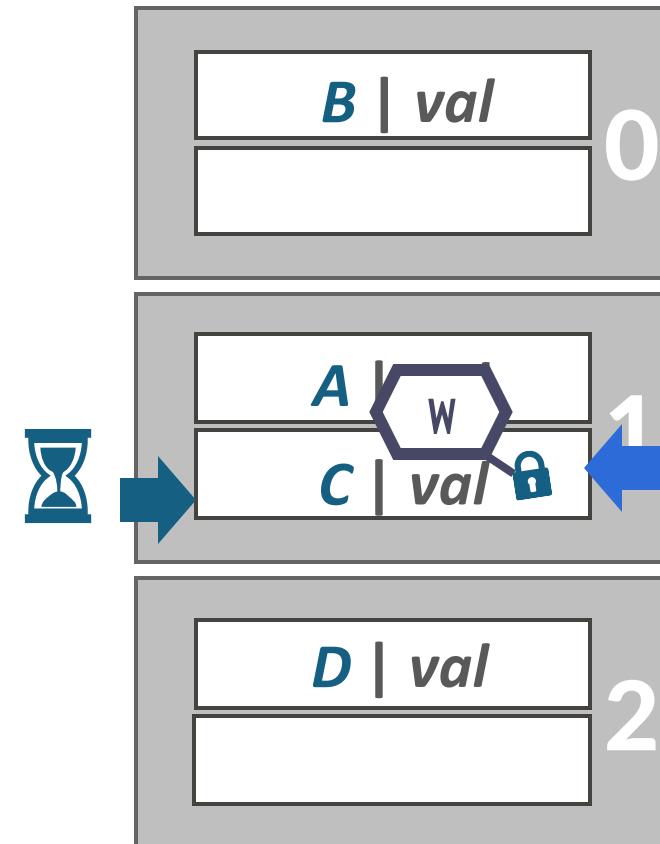
# Hash Table – Slot Latches



# Hash Table – Slot Latches

$T_1$ : Find D  
 $hash(D)$

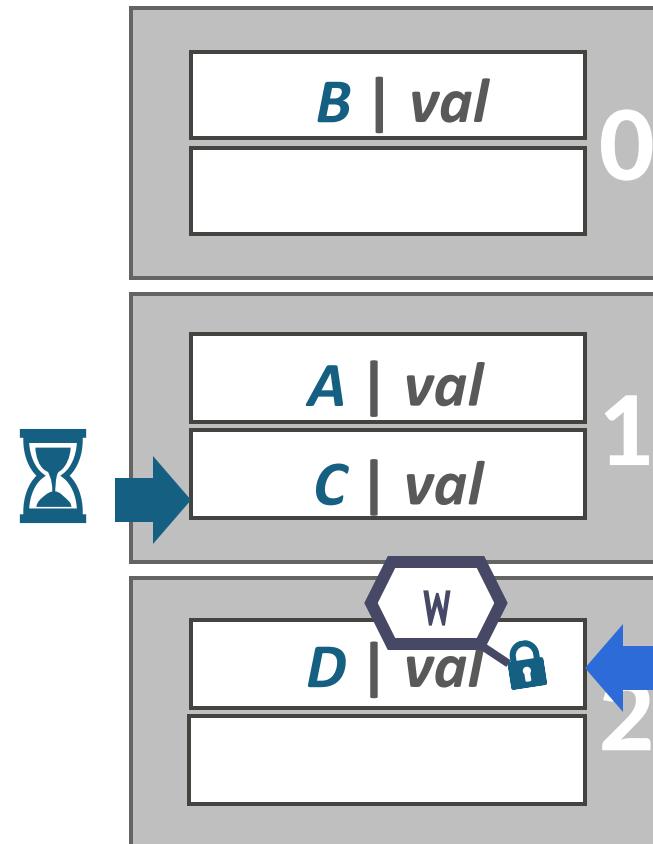
$T_2$ : Insert E  
 $hash(E)$



# Hash Table – Slot Latches

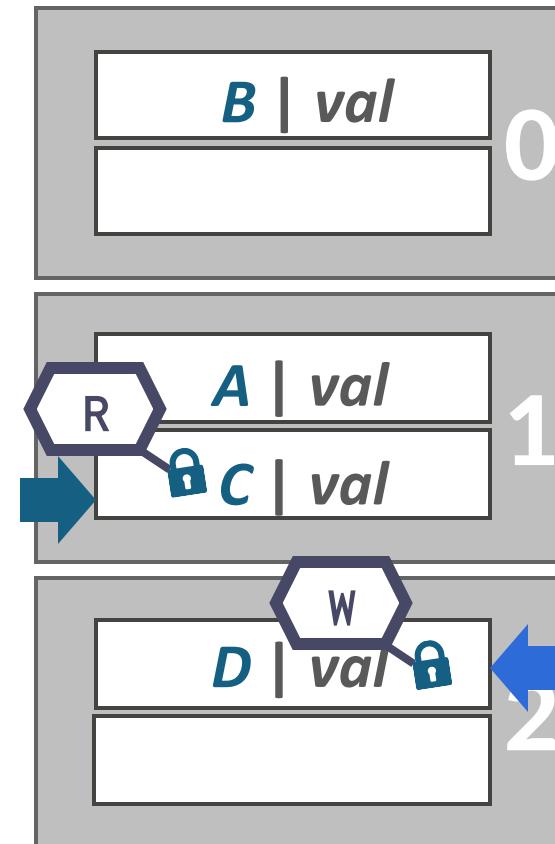
$T_1$ : Find D  
 $hash(D)$

$T_2$ : Insert E  
 $hash(E)$



# Hash Table – Slot Latches

$T_1$ : Find D  
 $hash(D)$

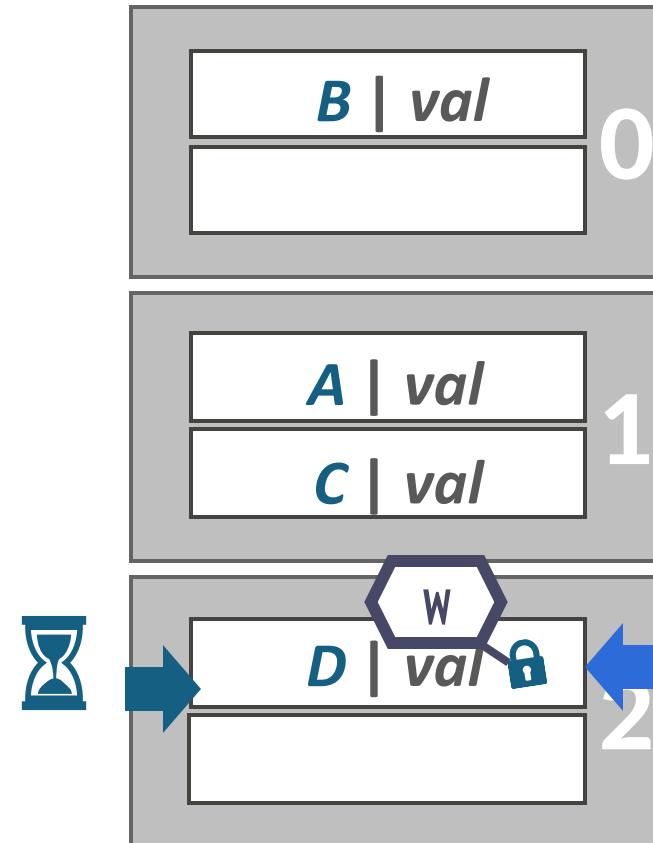


$T_2$ : Insert E  
 $hash(E)$

# Hash Table – Slot Latches

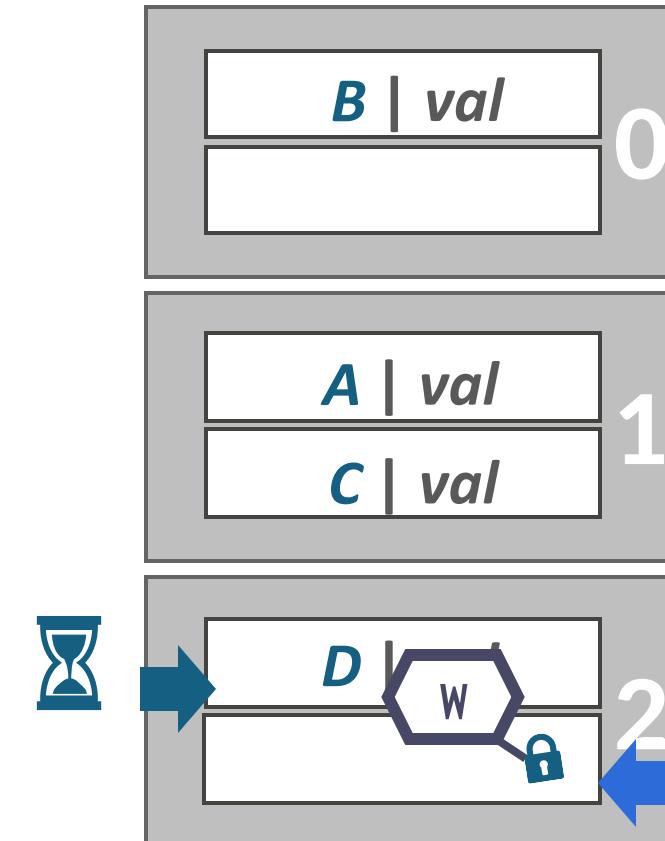
$T_1$ : Find D  
 $hash(D)$

$T_2$ : Insert E  
 $hash(E)$



# Hash Table – Slot Latches

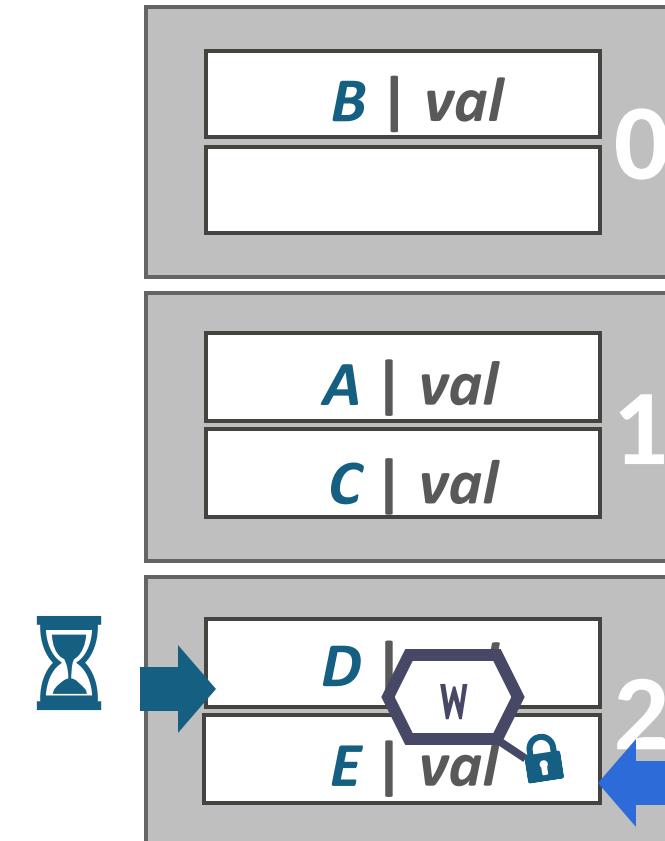
$T_1$ : Find D  
 $\text{hash}(D)$



$T_2$ : Insert E  
 $\text{hash}(E)$

# Hash Table – Slot Latches

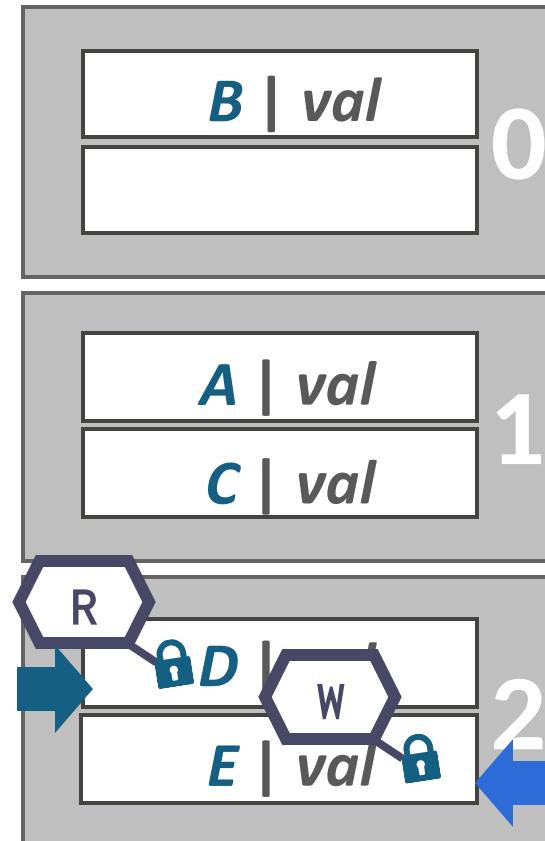
$T_1$ : Find D  
 $\text{hash}(D)$



$T_2$ : Insert E  
 $\text{hash}(E)$

# Hash Table – Slot Latches

$T_1$ : Find D  
 $\text{hash}(D)$



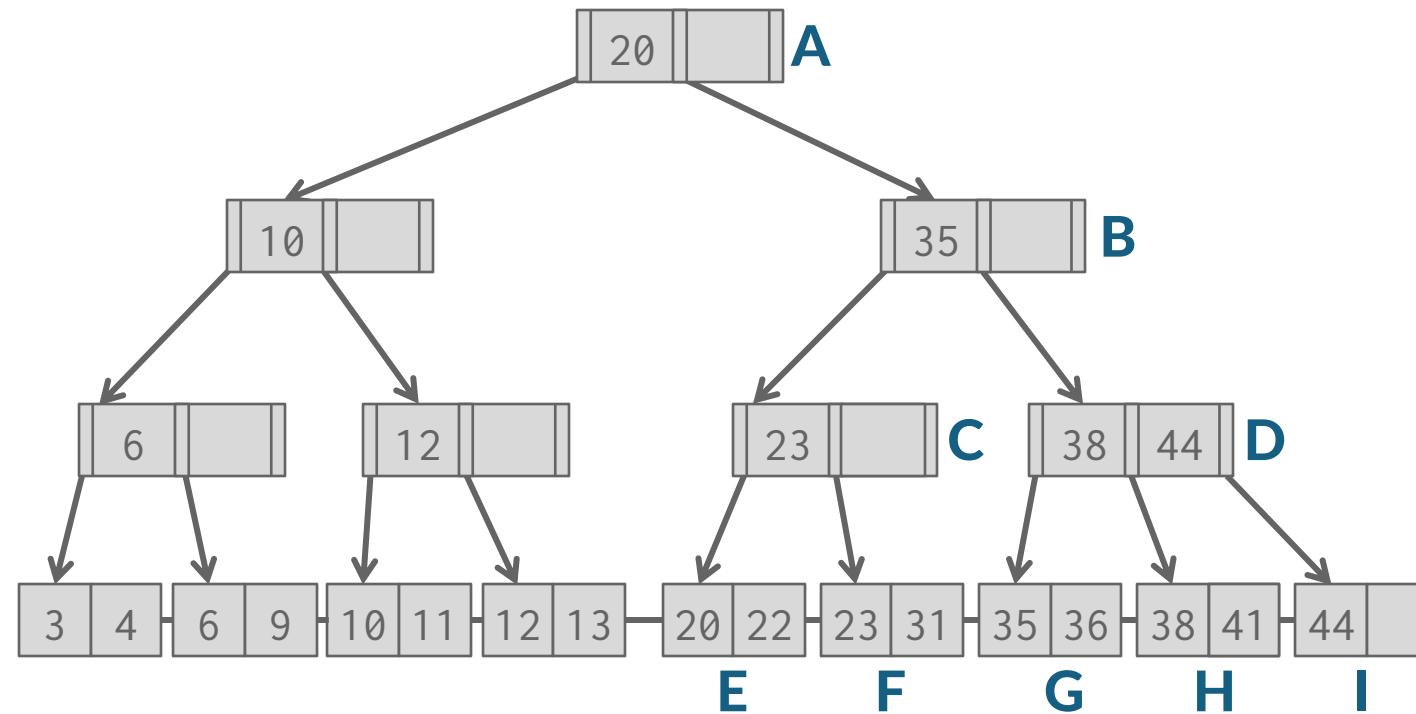
$T_2$ : Insert E  
 $\text{hash}(E)$

# B+Tree Concurrency Control

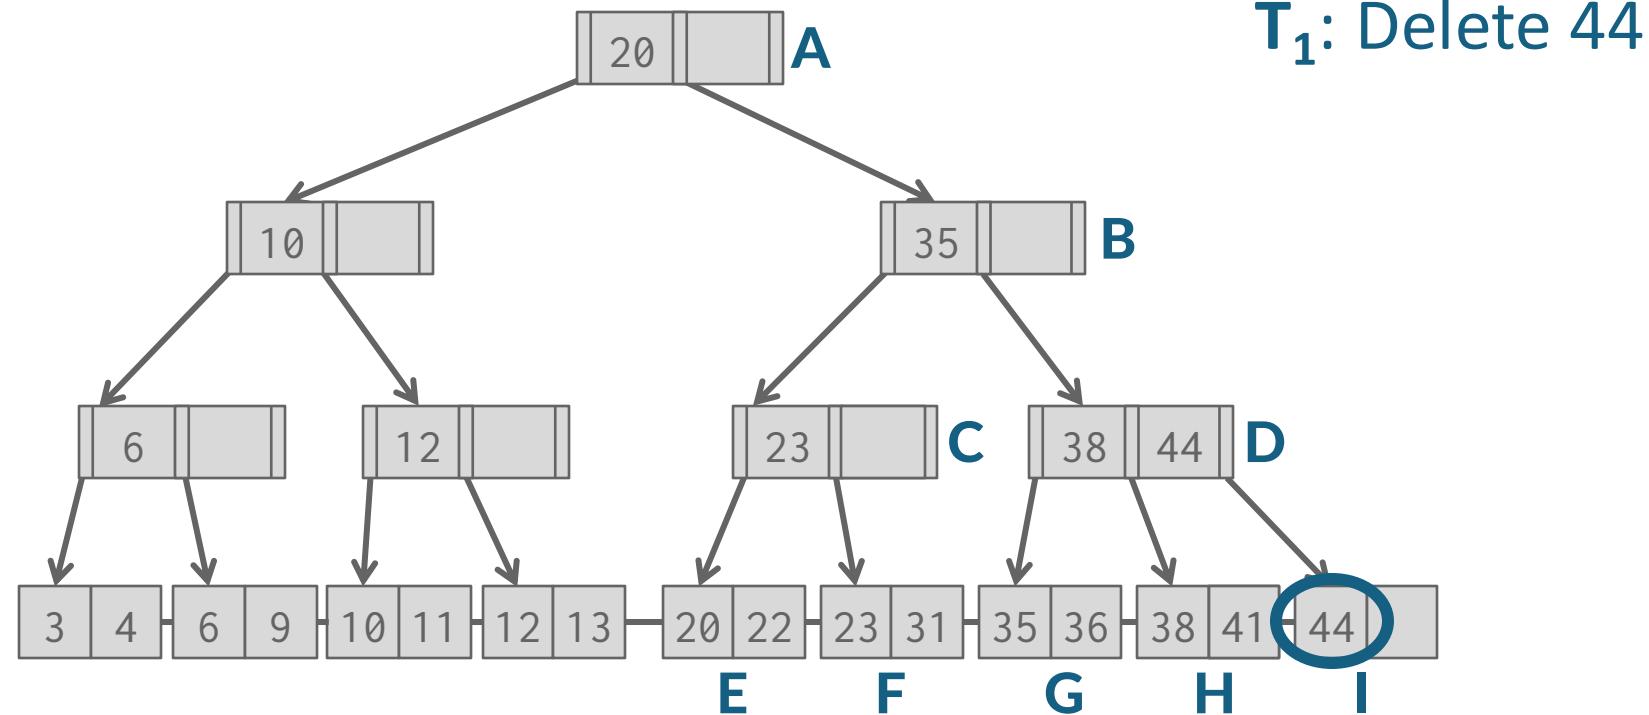
# B+Tree Concurrency Control

- We want to allow multiple threads to read and update a B+Tree at the same time.
- We need to protect against two types of problems:
  - Threads trying to modify the contents of a node at the same time.
  - One thread traversing the tree while another thread splits/merges nodes.

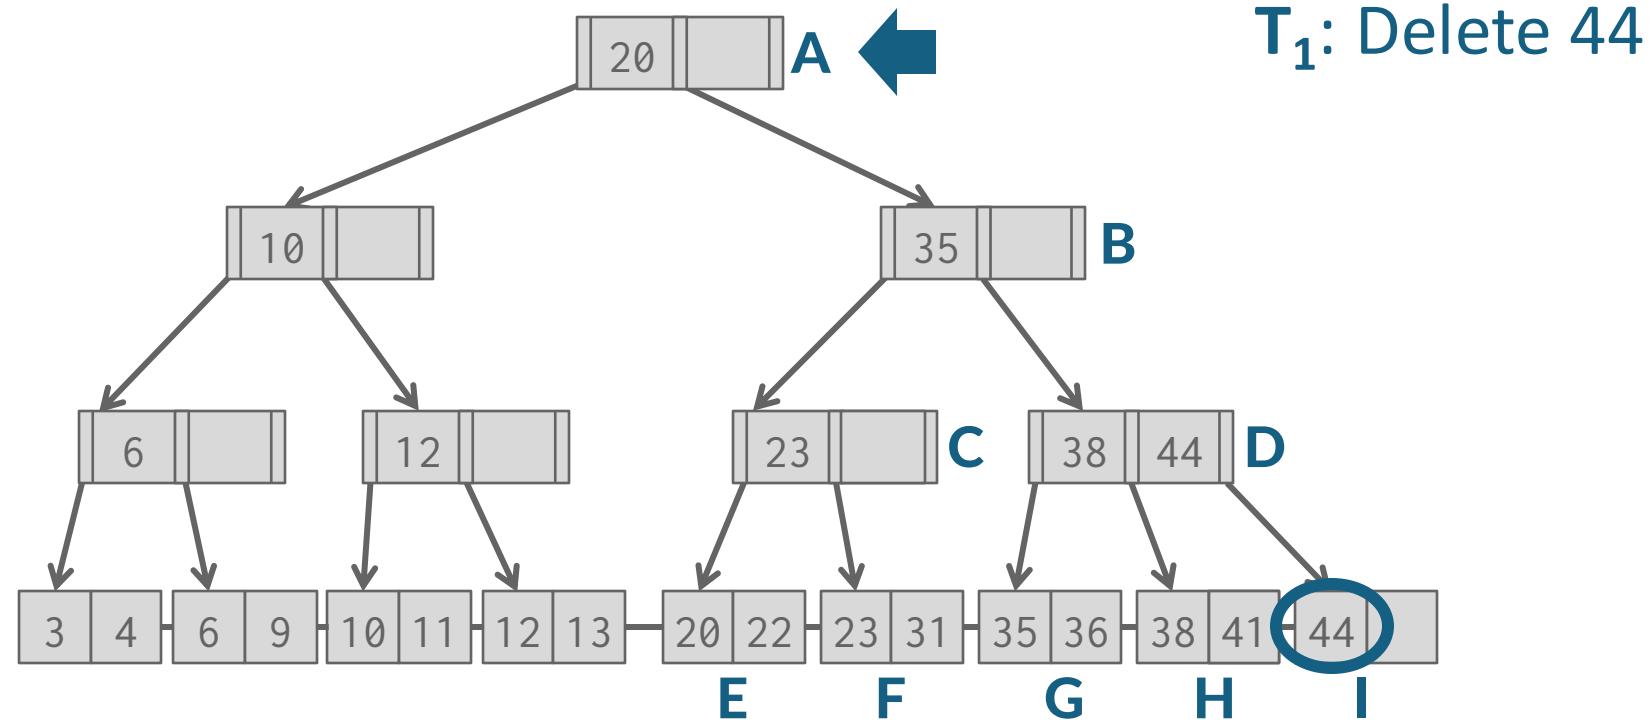
# B+Tree Multi-Threaded Example



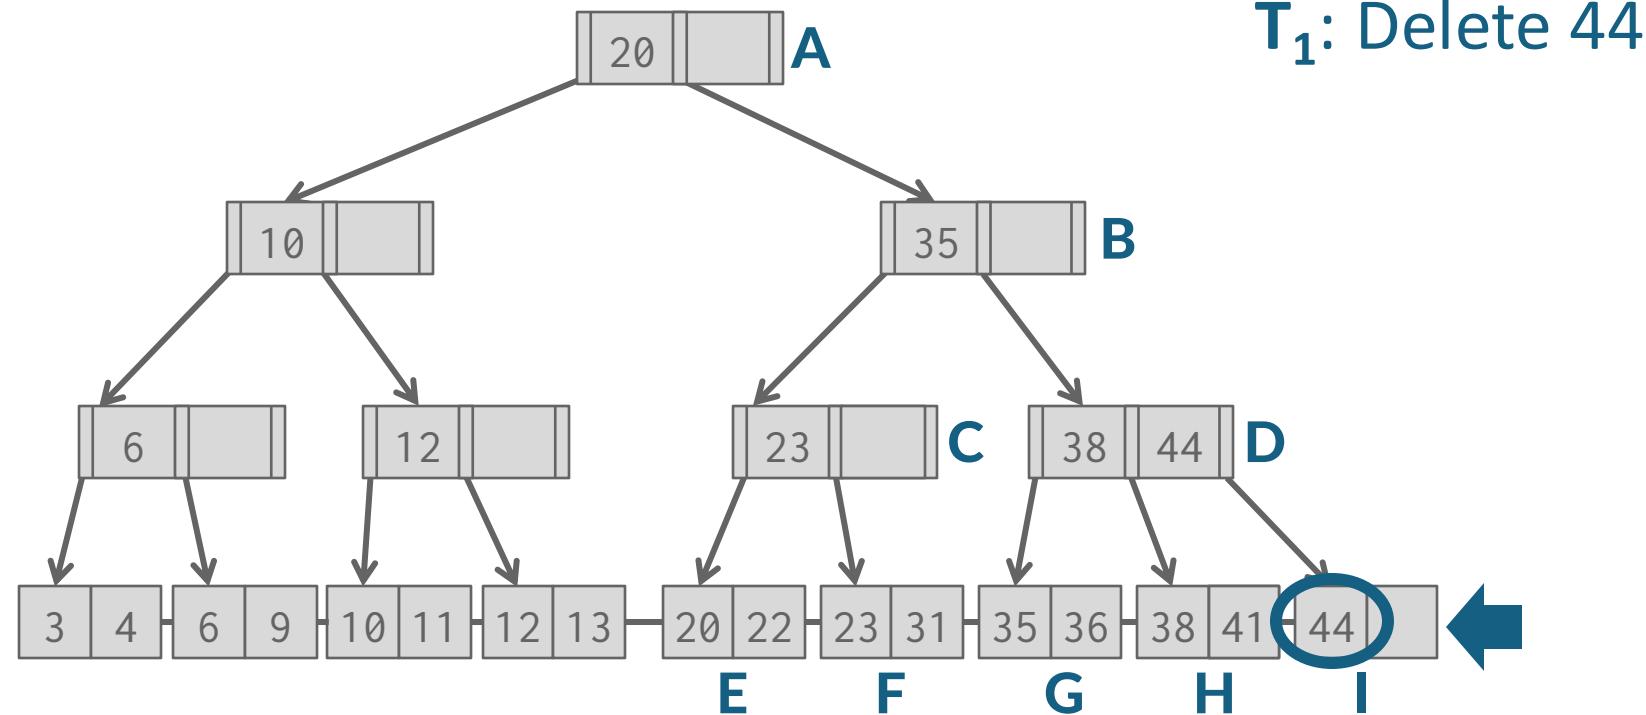
# B+Tree Multi-Threaded Example



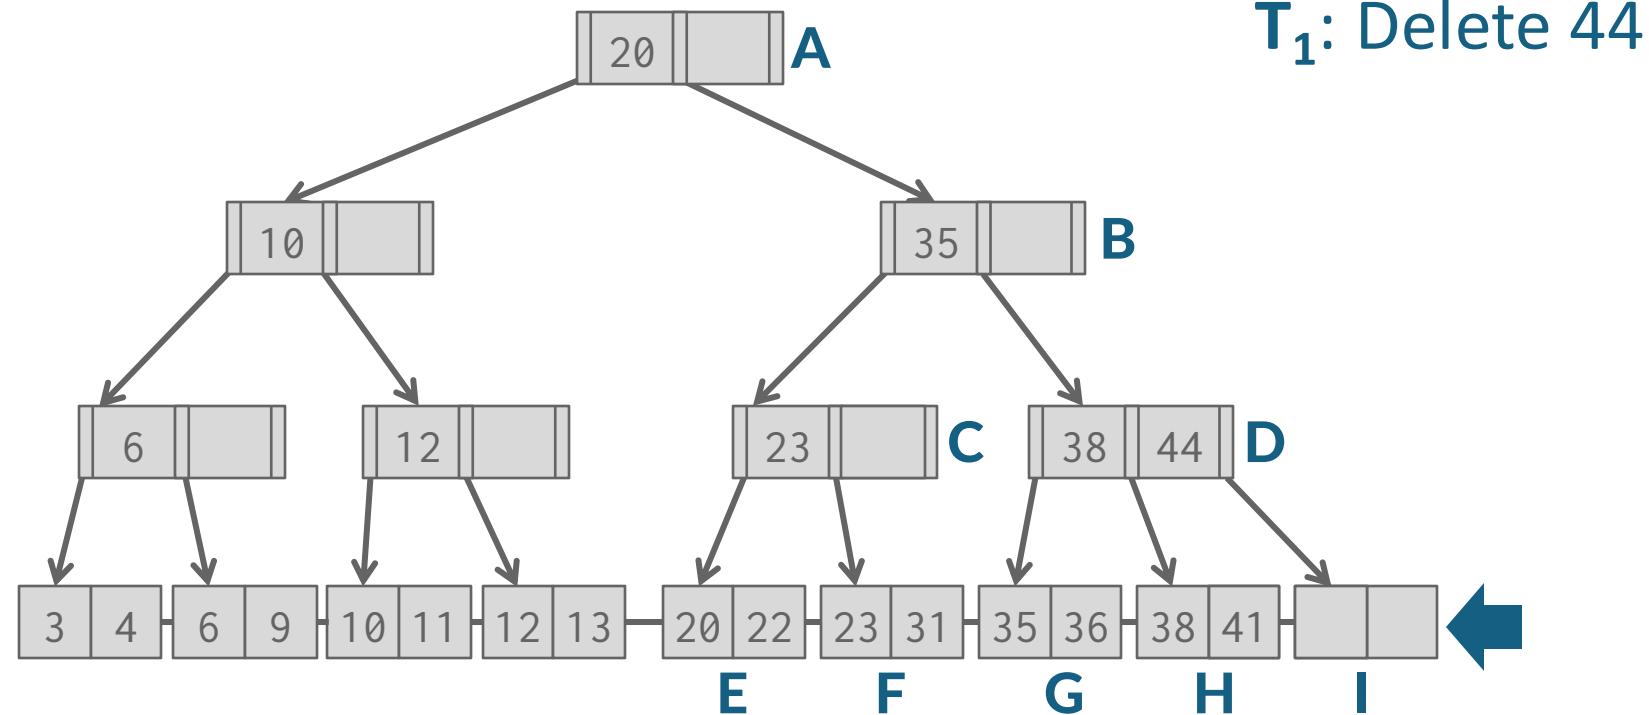
# B+Tree Multi-Threaded Example



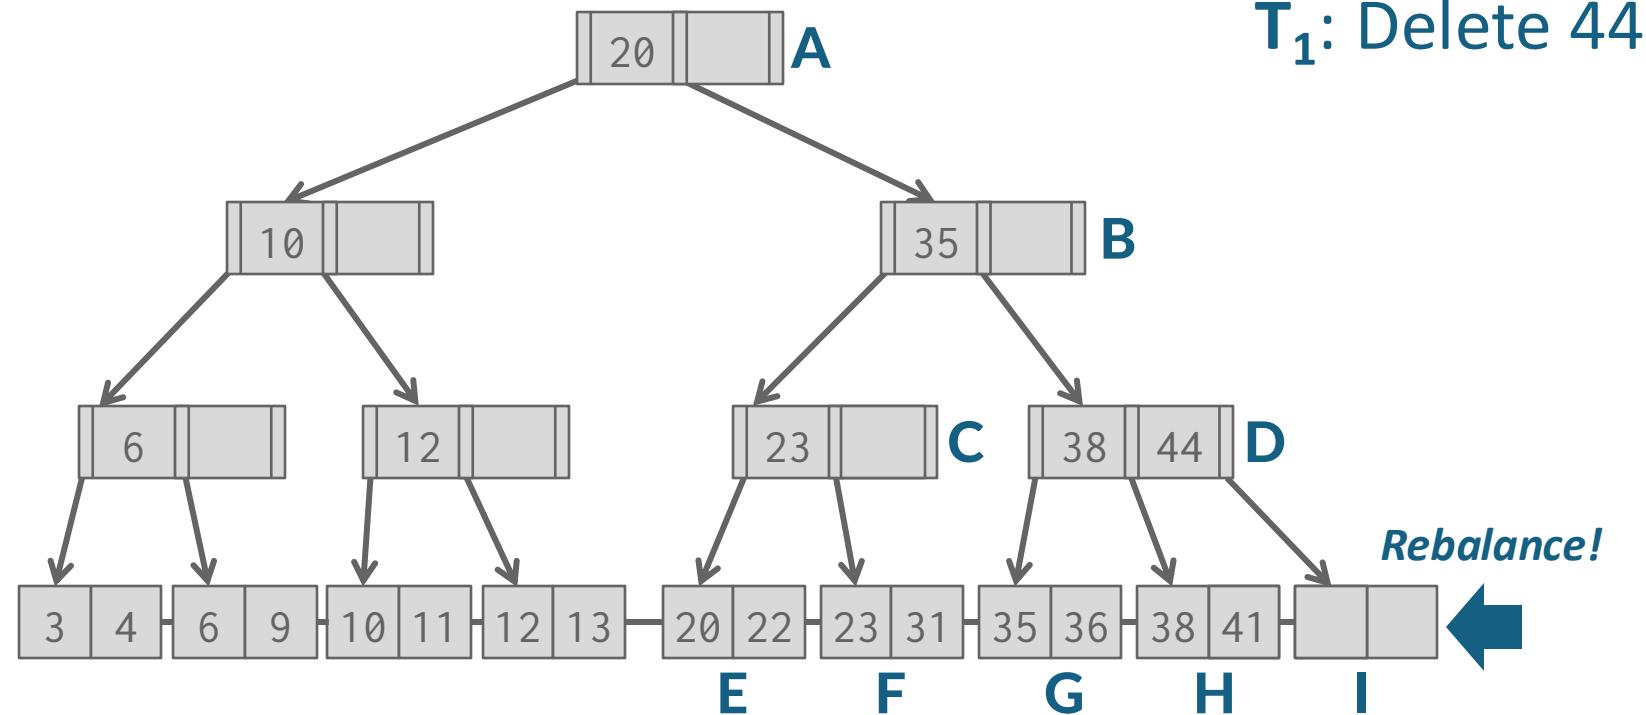
# B+Tree Multi-Threaded Example



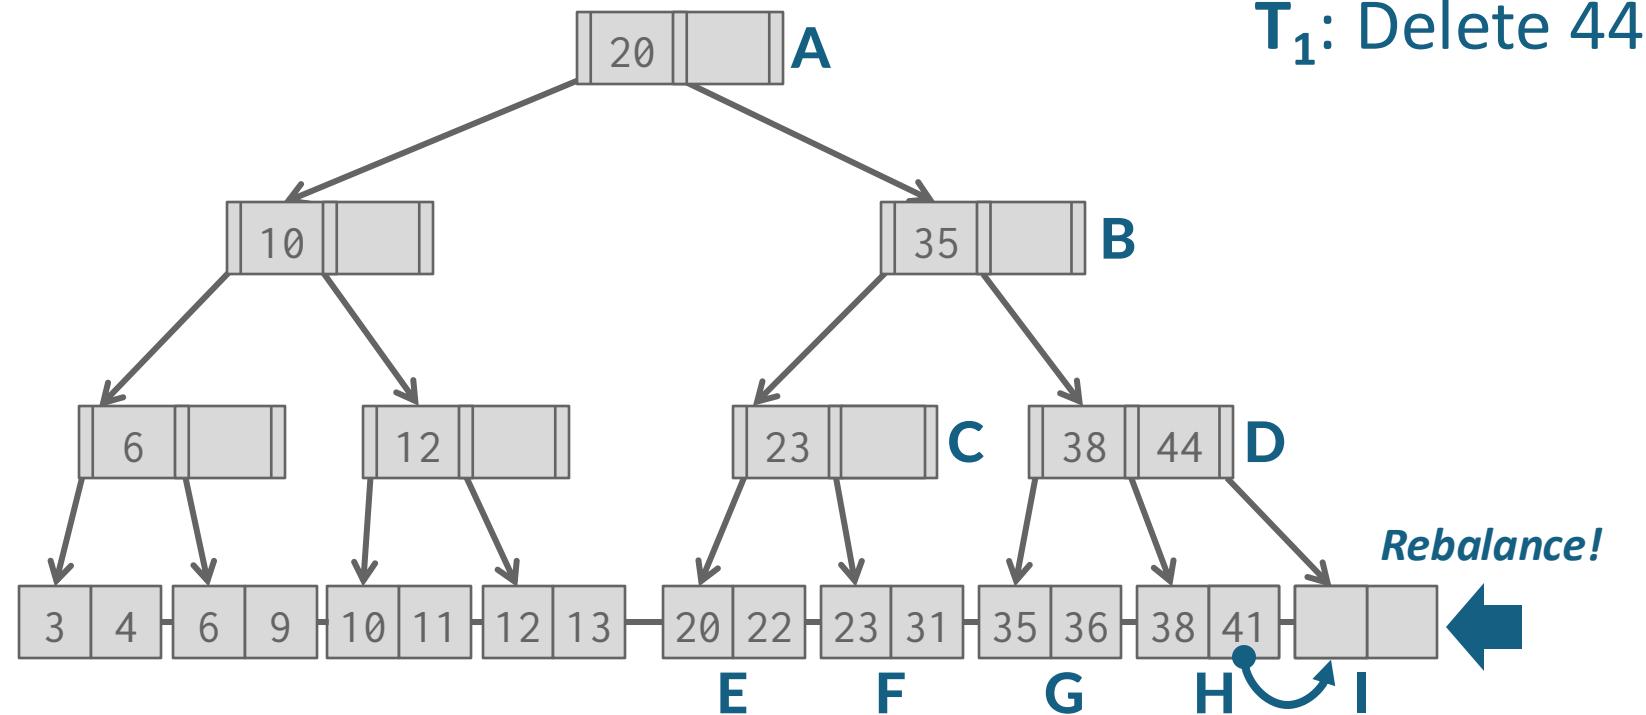
# B+Tree Multi-Threaded Example



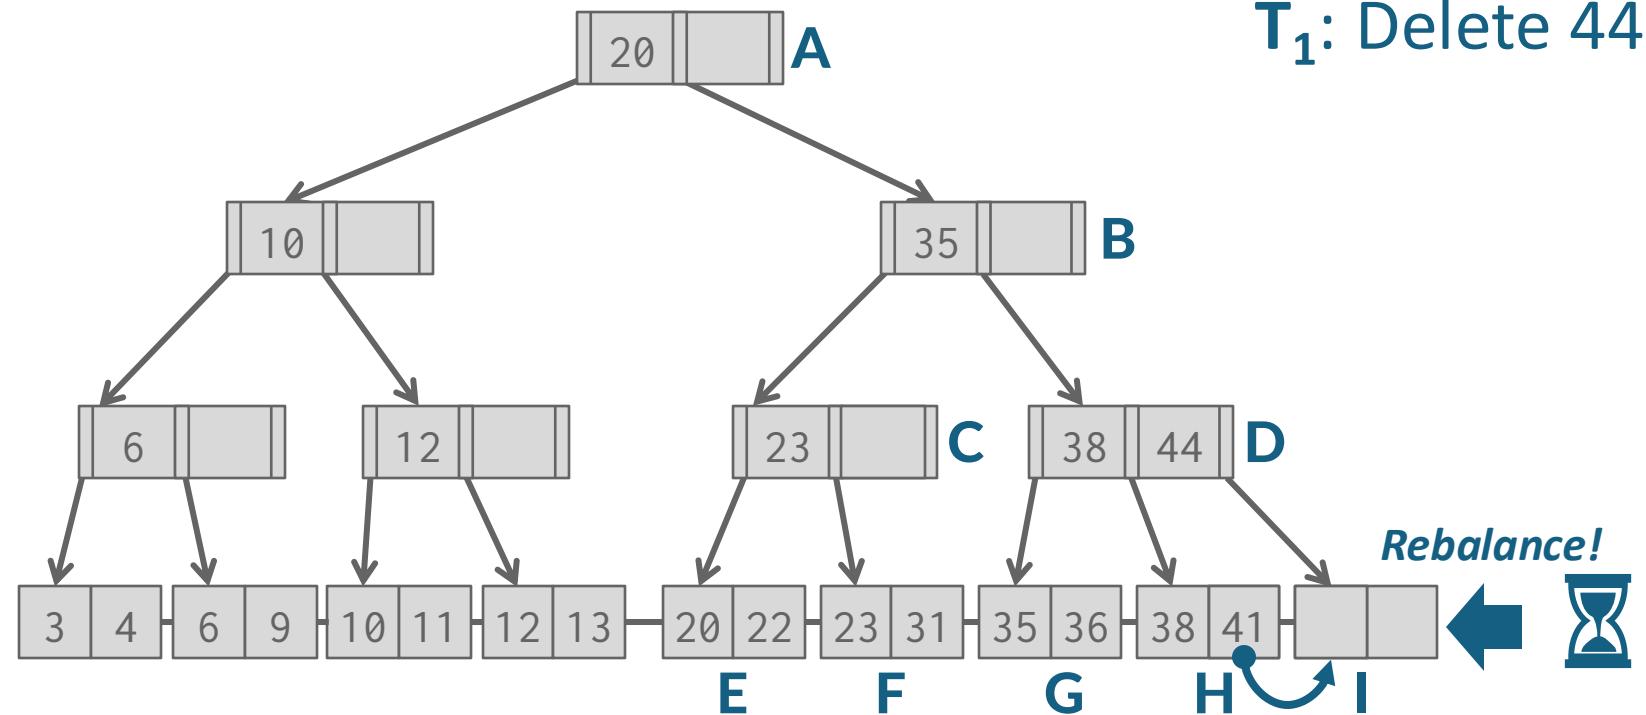
# B+Tree Multi-Threaded Example



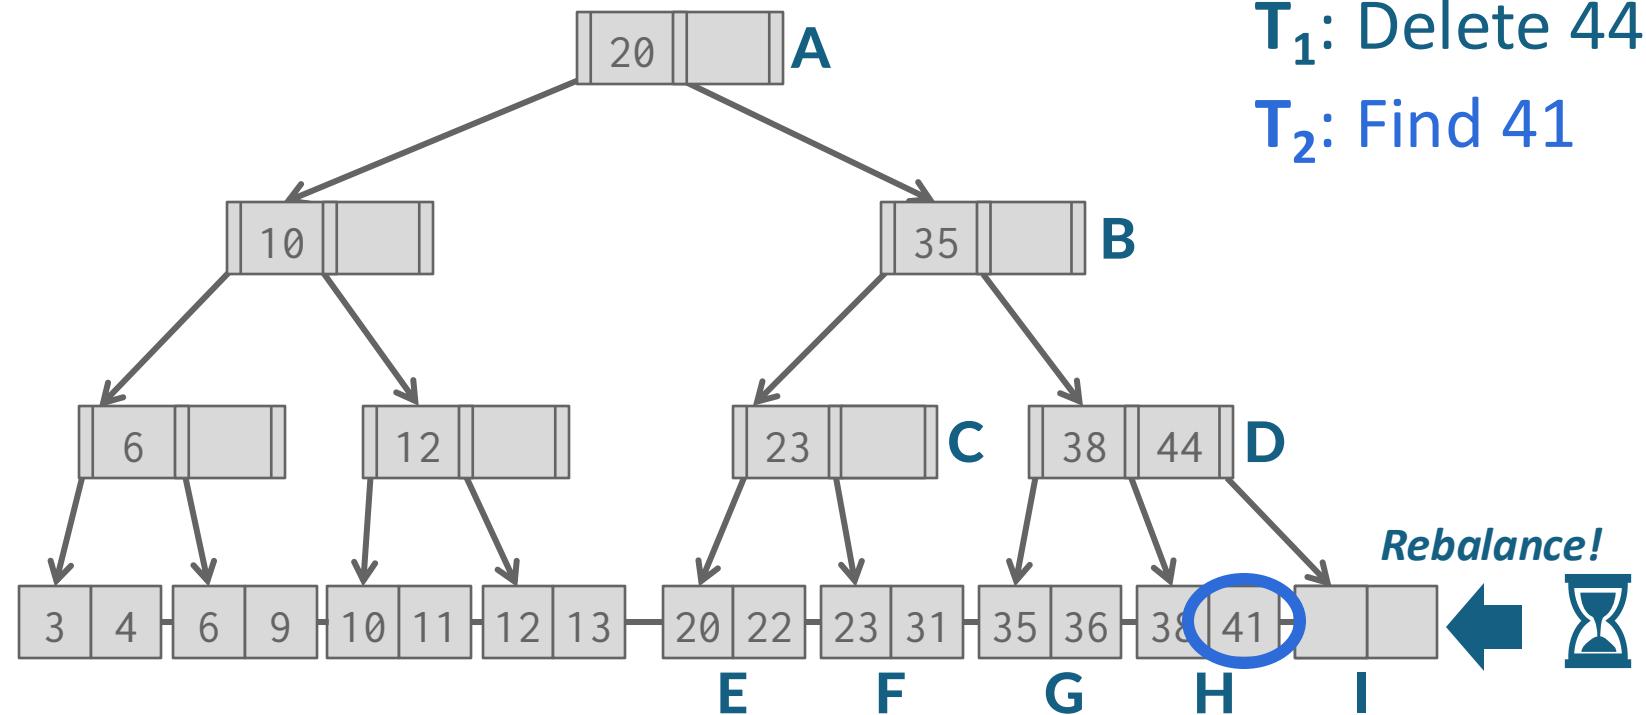
# B+Tree Multi-Threaded Example



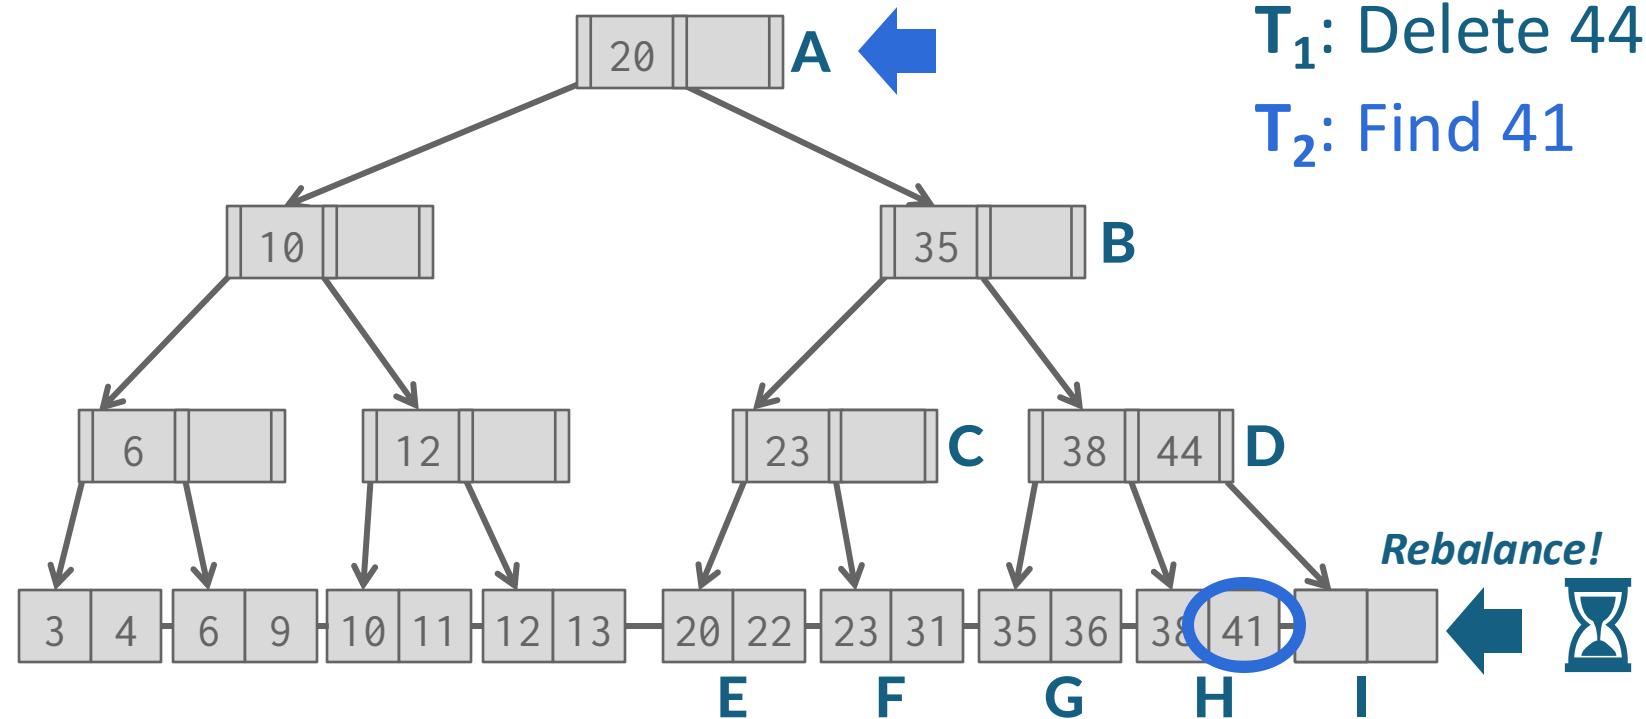
# B+Tree Multi-Threaded Example



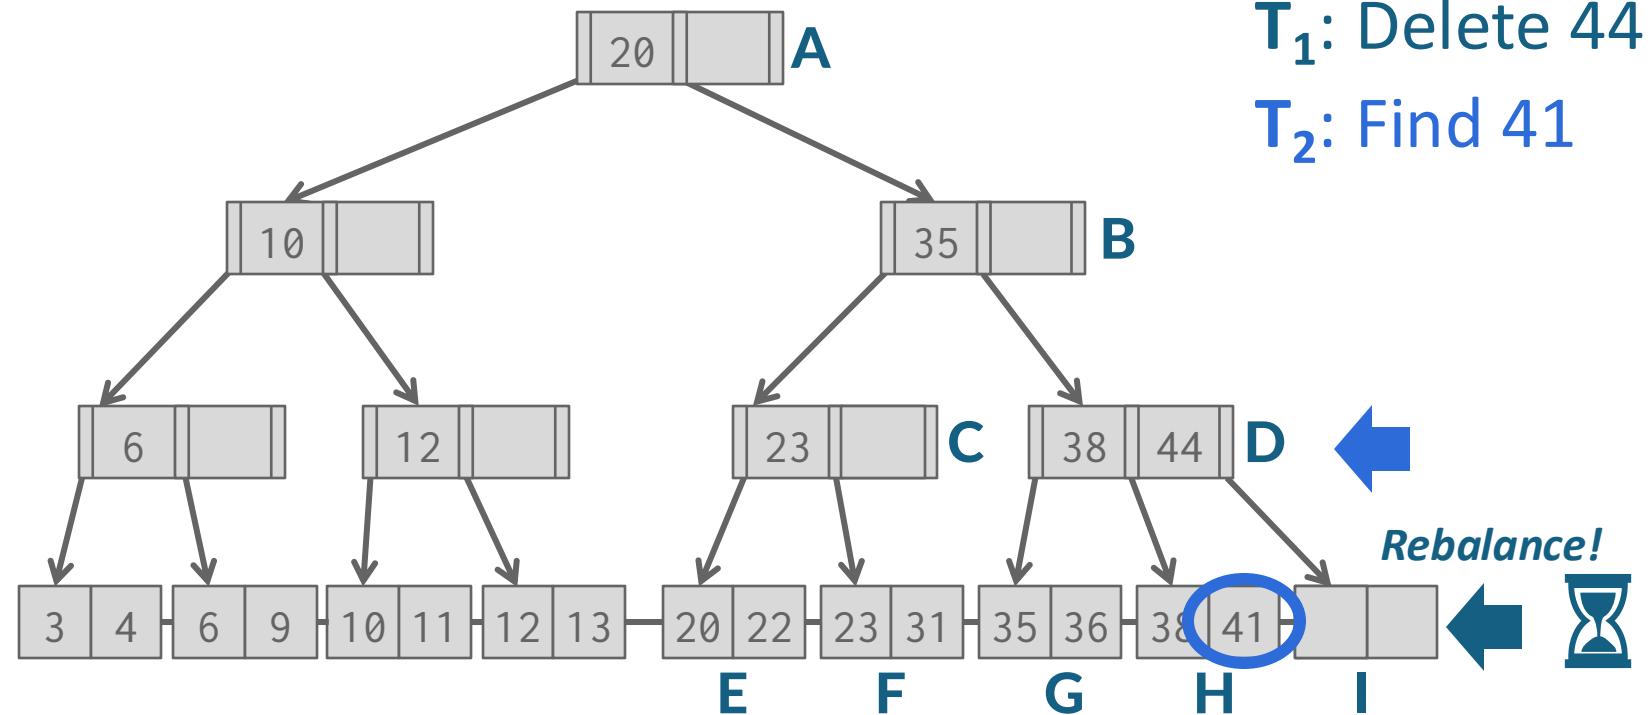
# B+Tree Multi-Threaded Example



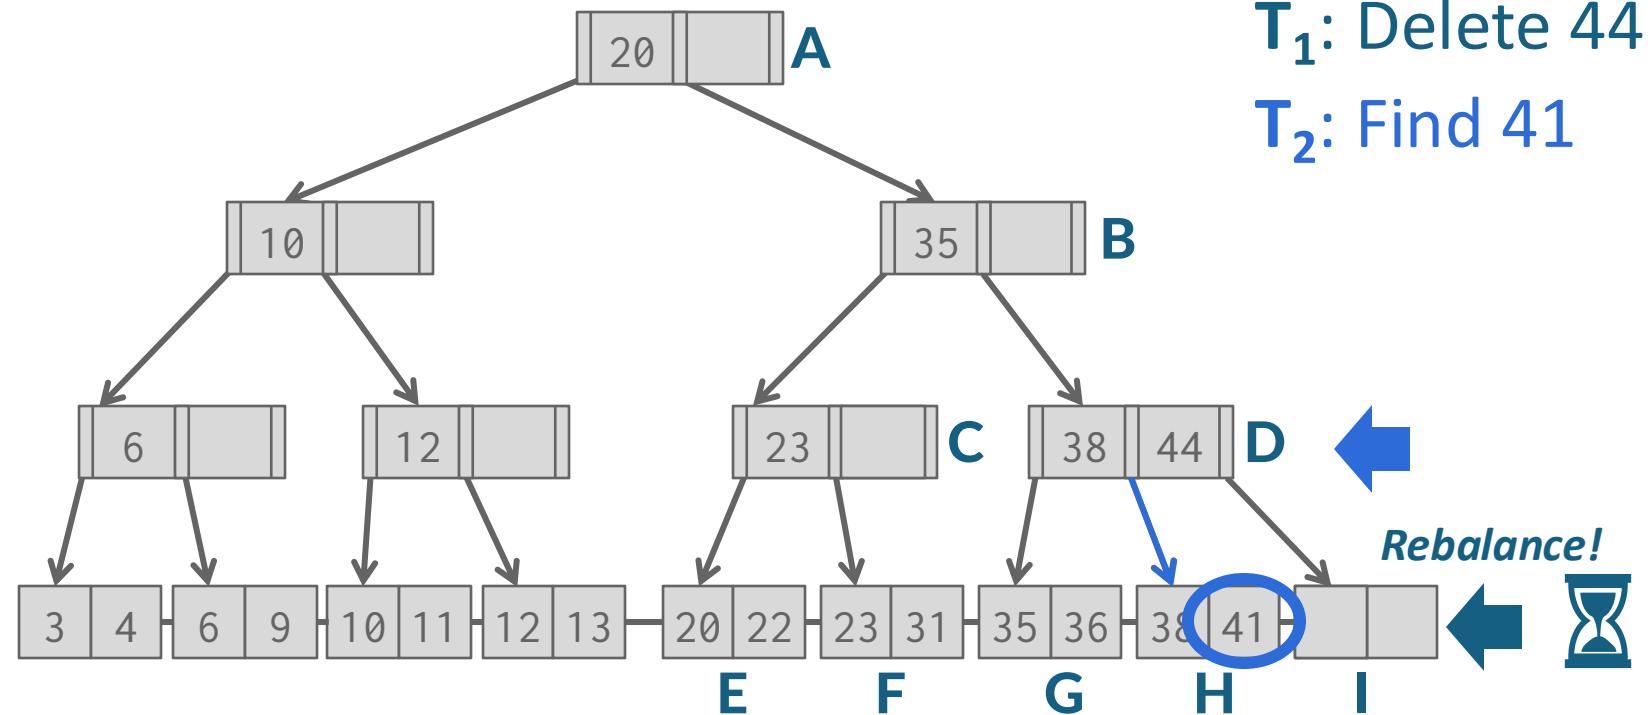
# B+Tree Multi-Threaded Example



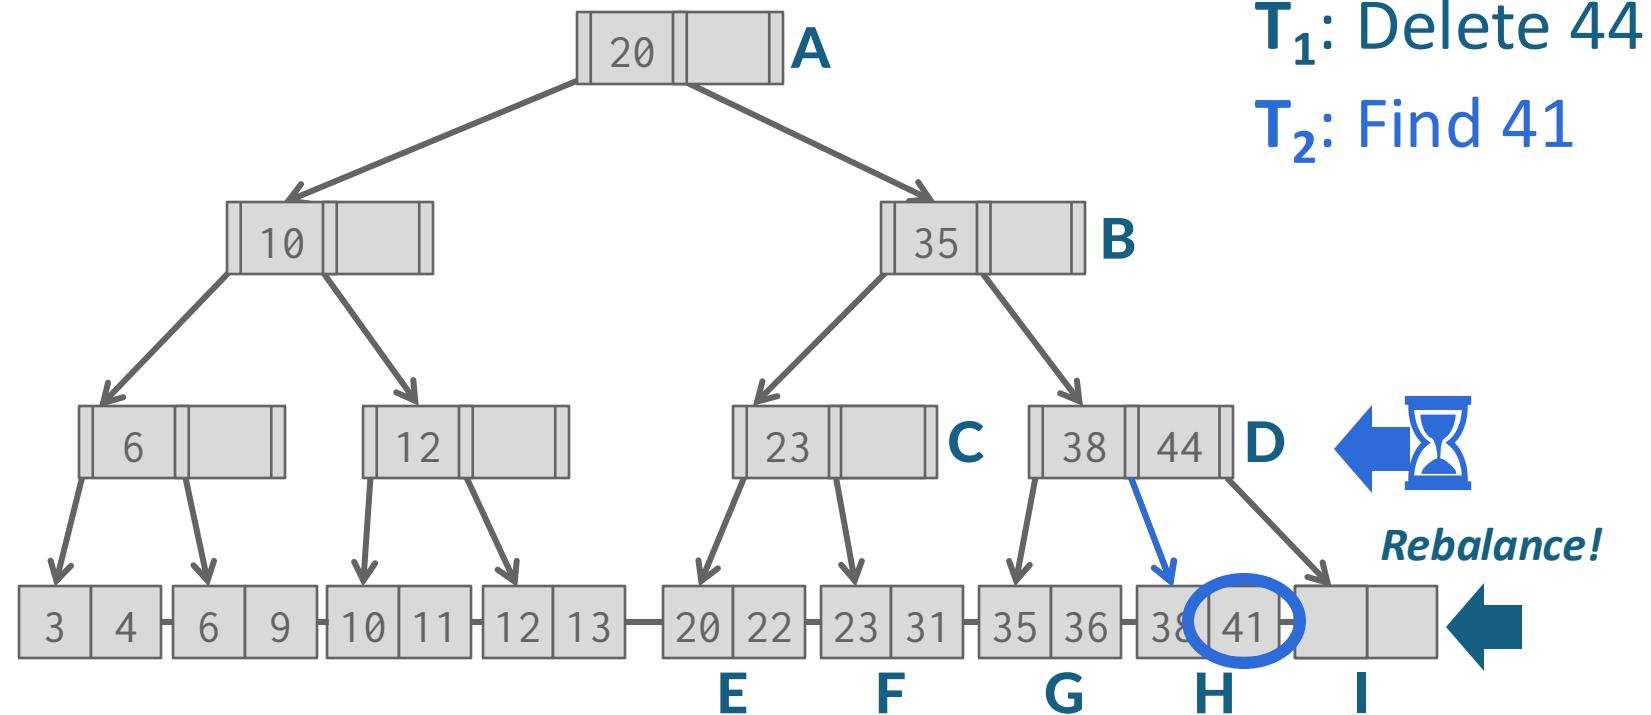
# B+Tree Multi-Threaded Example



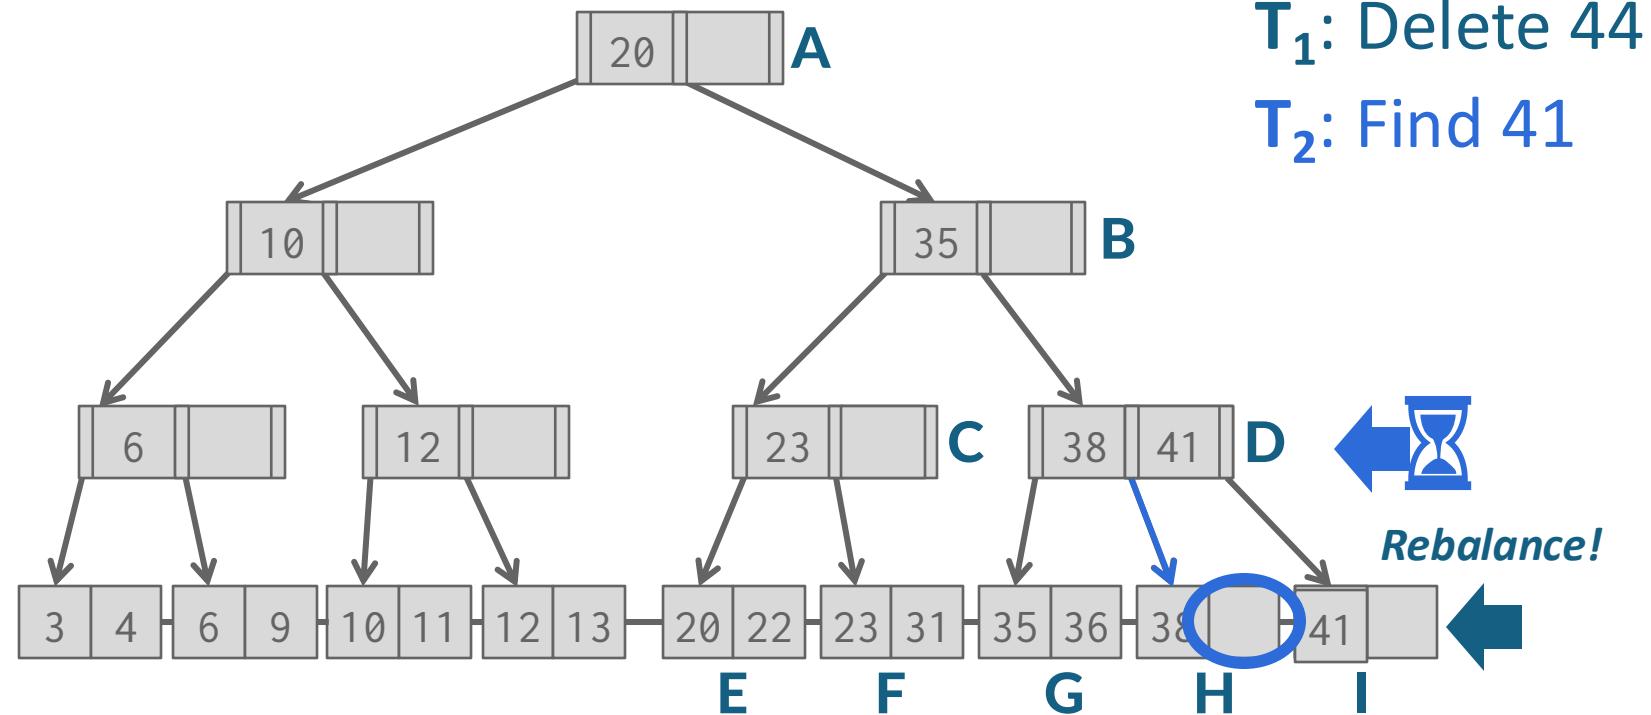
# B+Tree Multi-Threaded Example



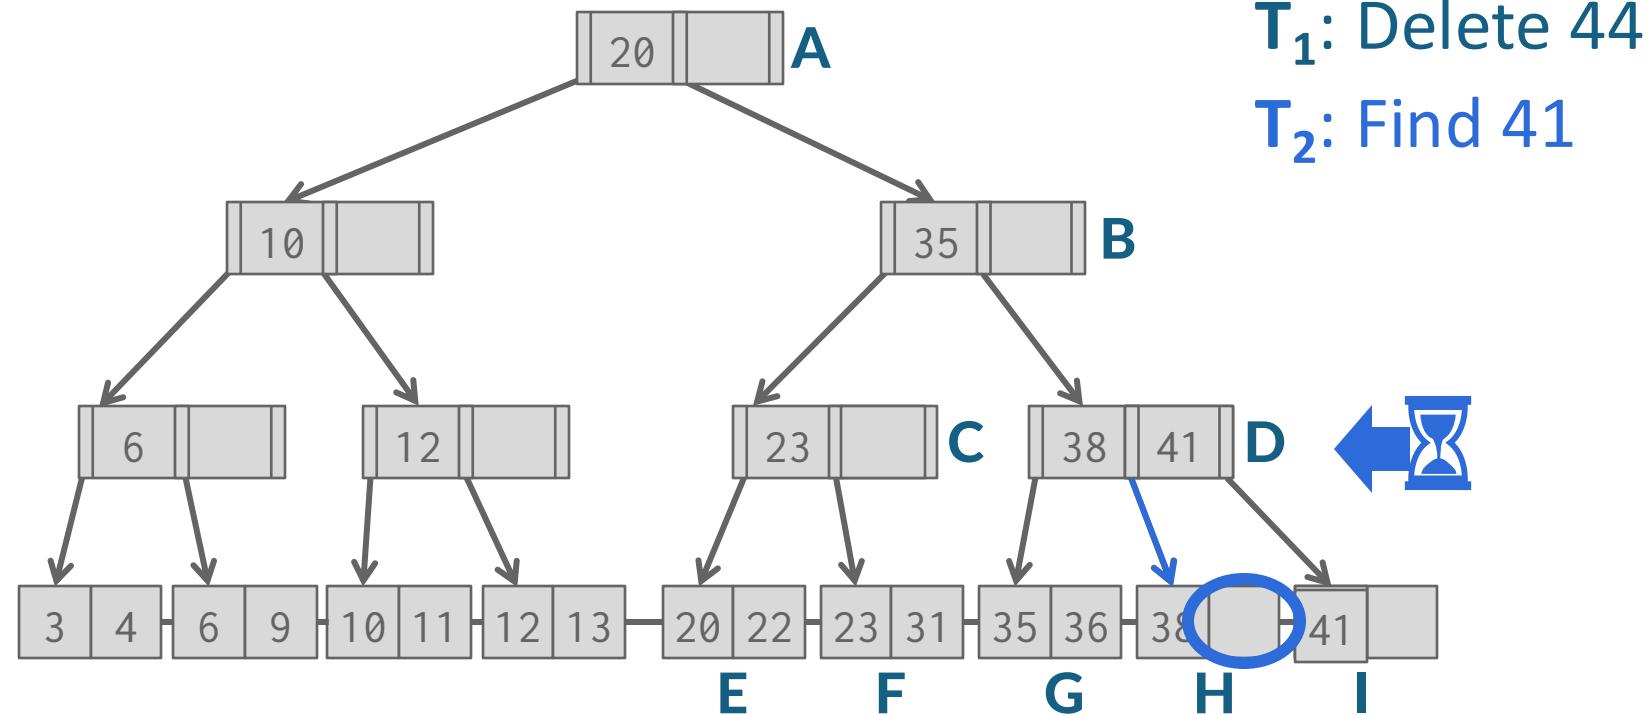
# B+Tree Multi-Threaded Example



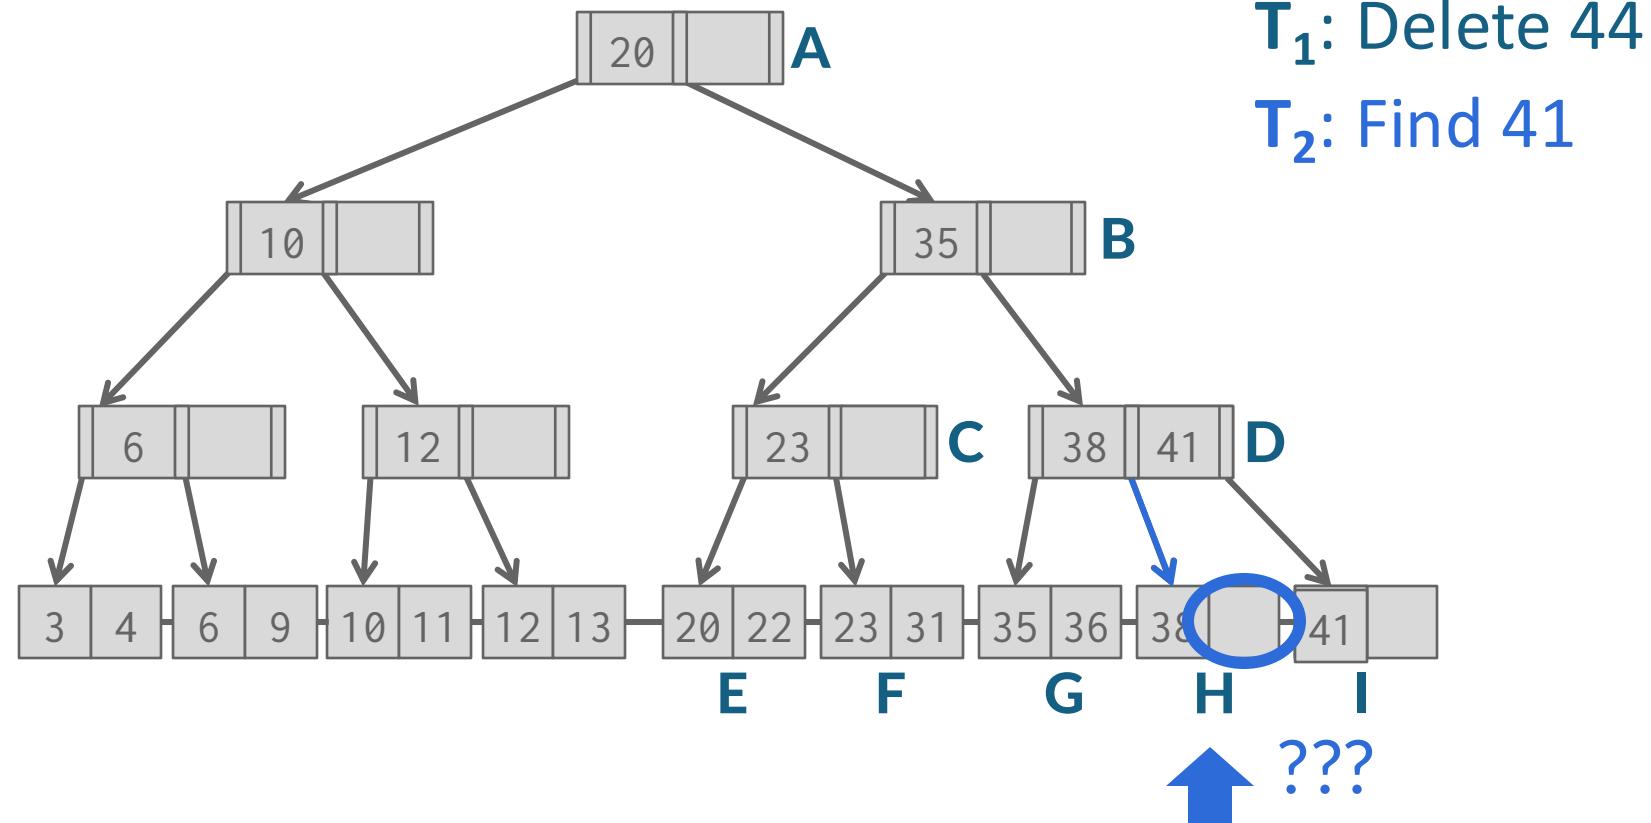
# B+Tree Multi-Threaded Example



# B+Tree Multi-Threaded Example



# B+Tree Multi-Threaded Example



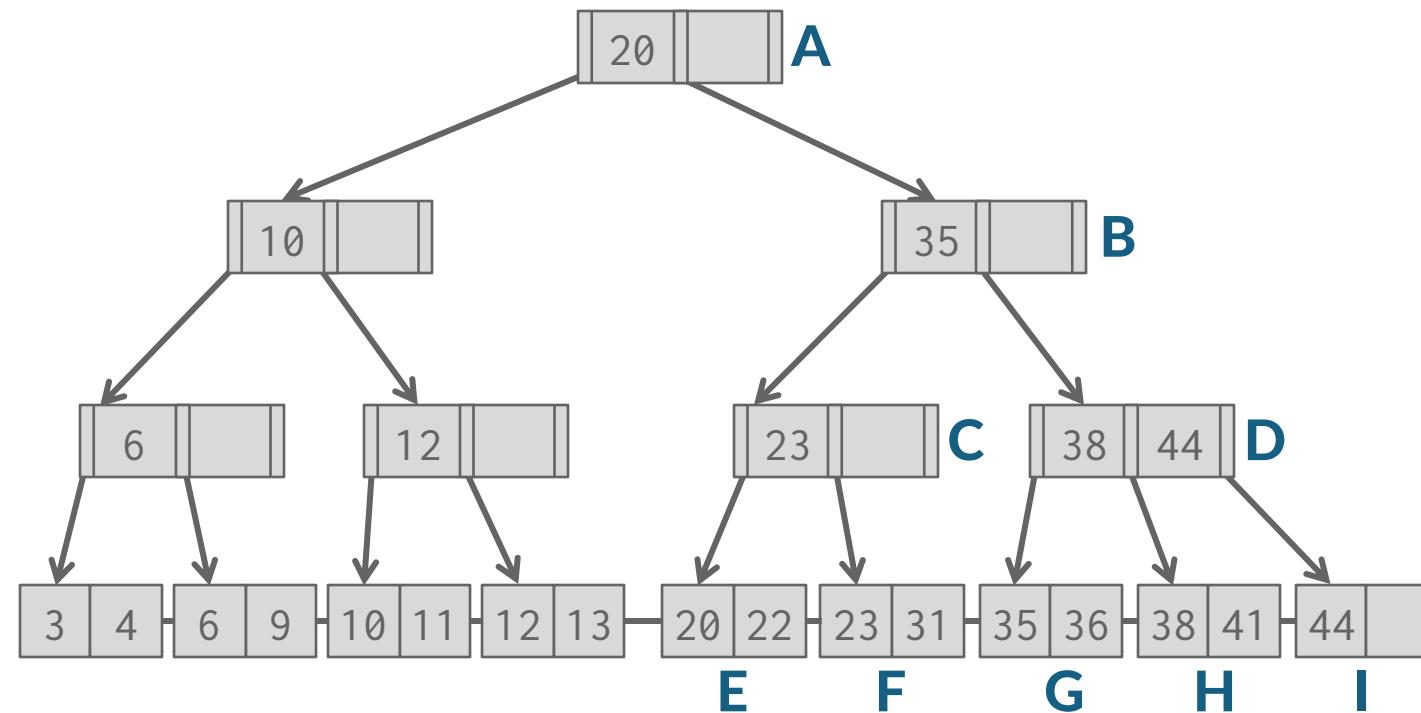
# Latch Crabbling/Coupling

- Protocol to allow multiple threads to access/modify B+Tree at the same time.
  - Get latch for parent
  - Get latch for child
  - Release latch for parent if “safe”
- A **safe node** is one that will not split or merge when updated.
  - Not full (on insertion)
  - More than half-full (on deletion)

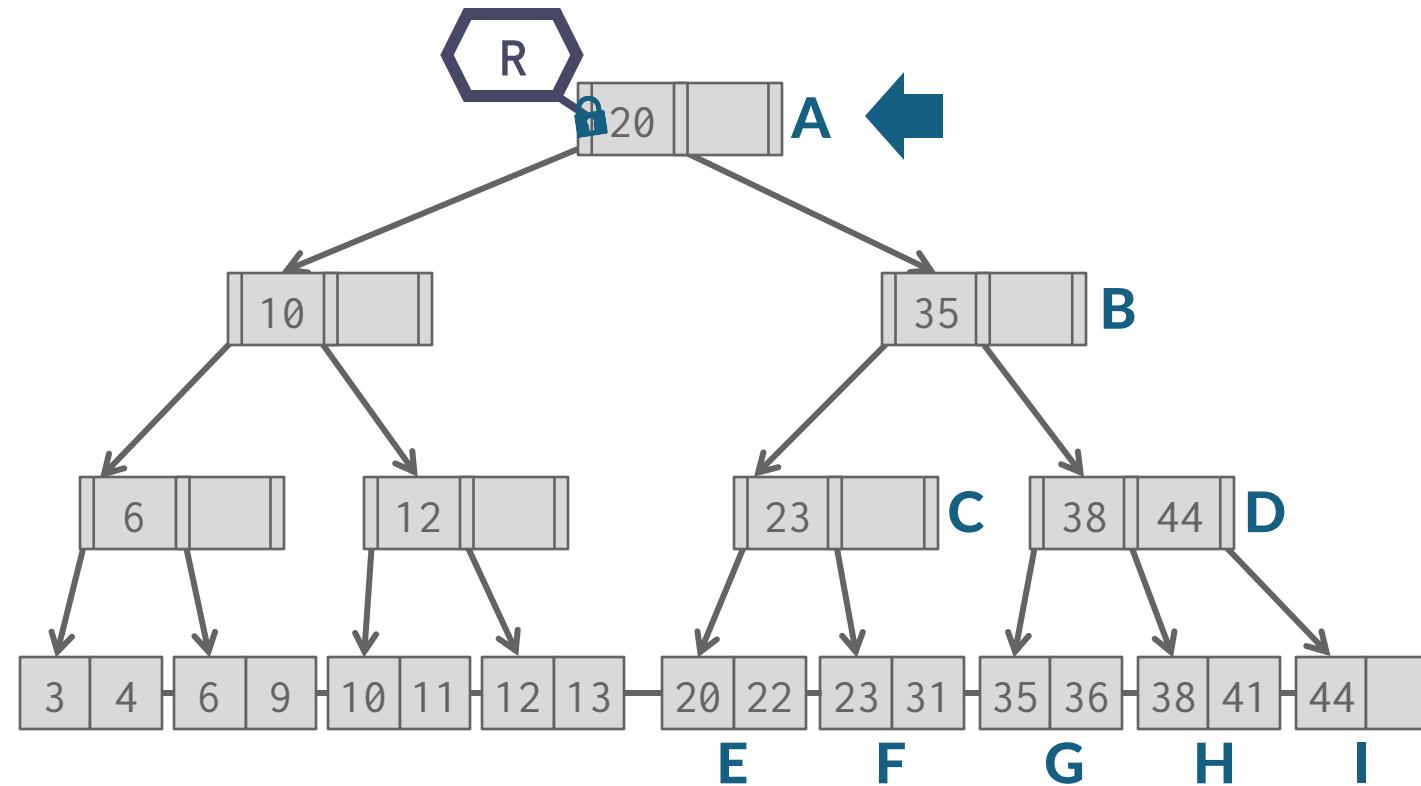
# Latch Crabbling/Coupling

- **Find:** Start at root and traverse down the tree:
  - Acquire **R** latch on child,
  - Then unlatch parent.
  - Repeat until we reach the leaf node.
- **Insert/Delete:** Start at root and go down, obtaining **W** latches as needed. Once child is latched, check if it is safe:
  - If child is safe, release all latches on ancestors

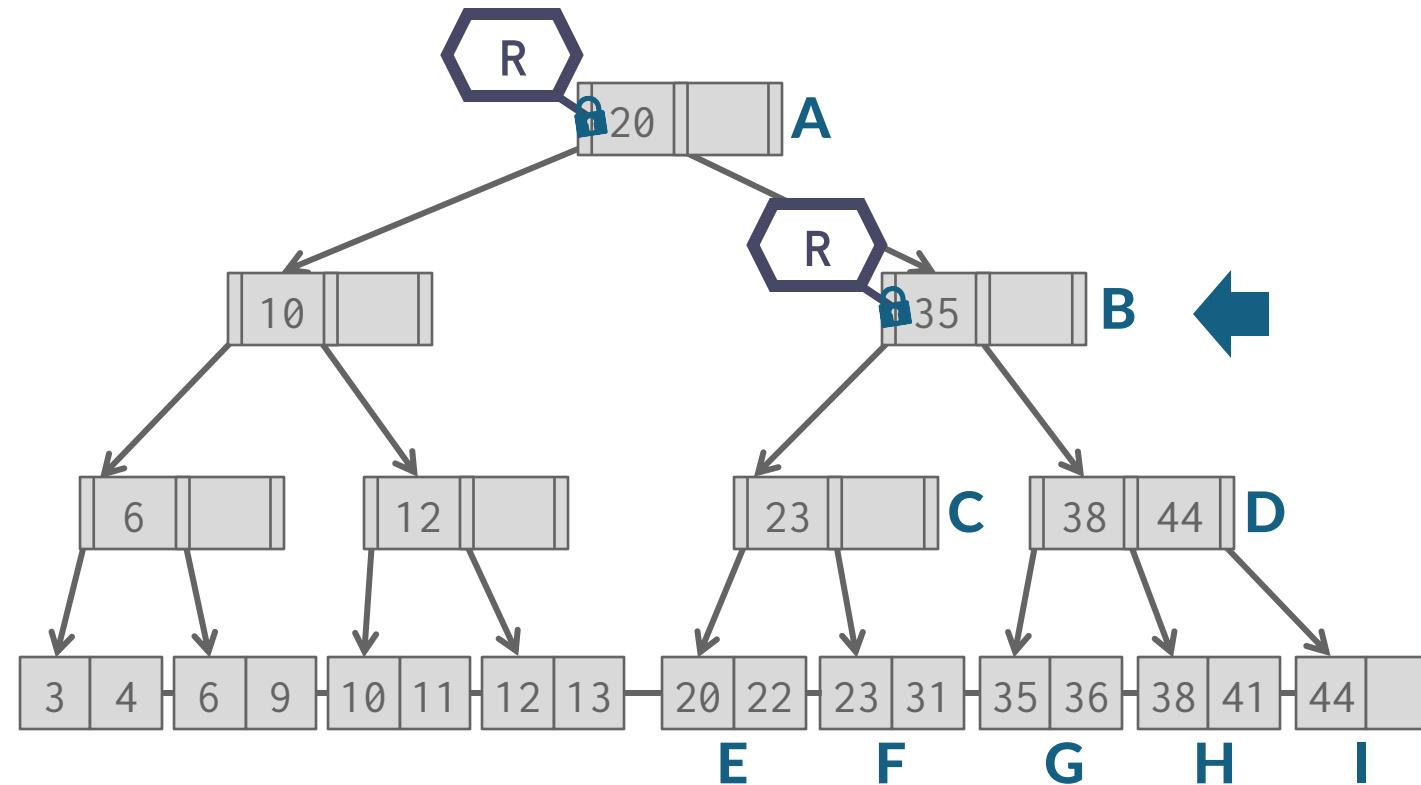
# Example #1 – Find 38



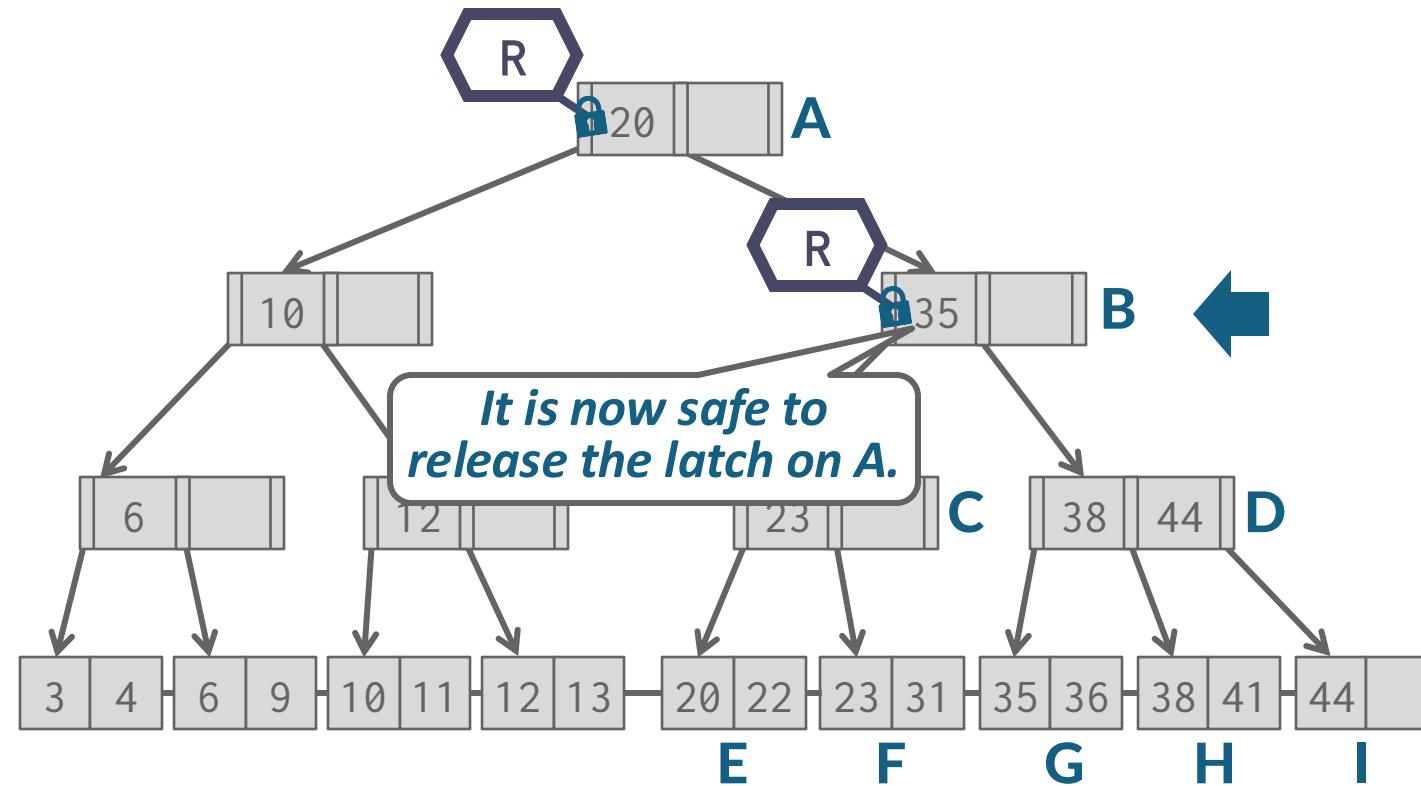
# Example #1 – Find 38



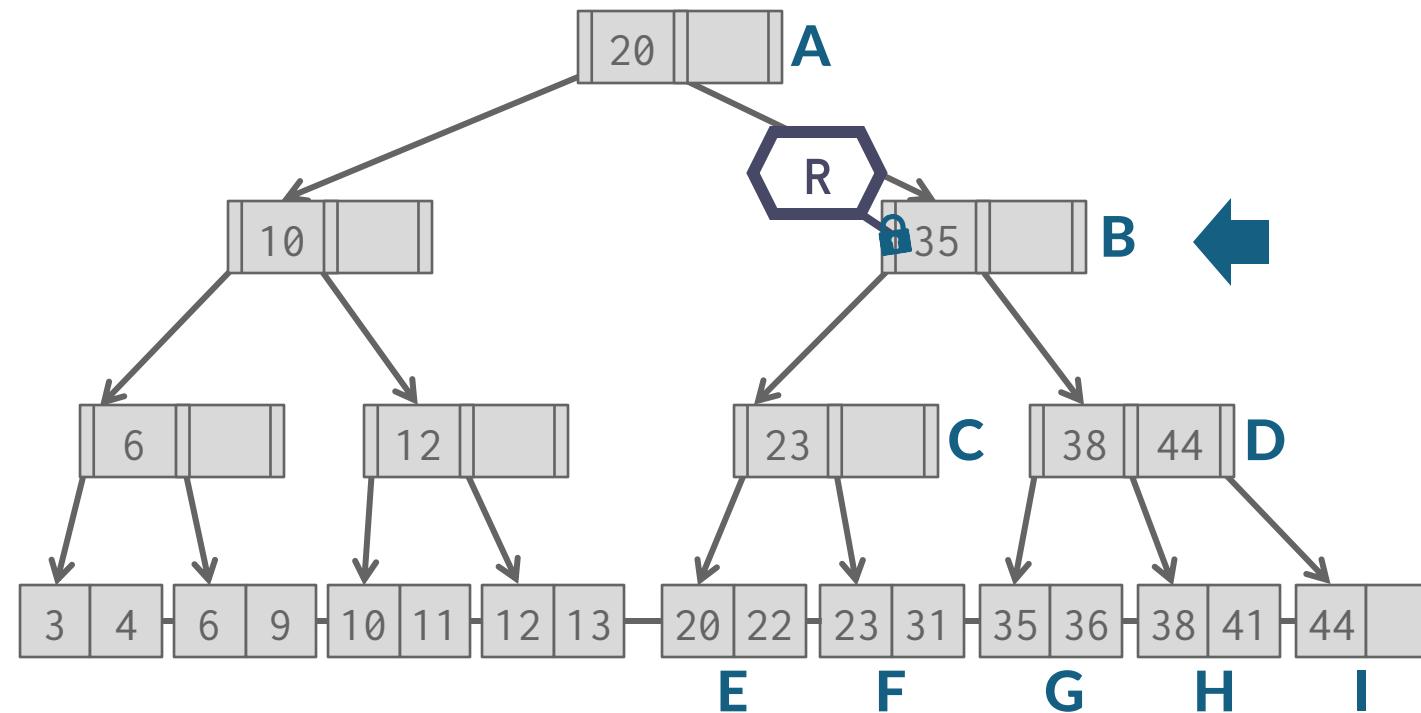
# Example #1 – Find 38



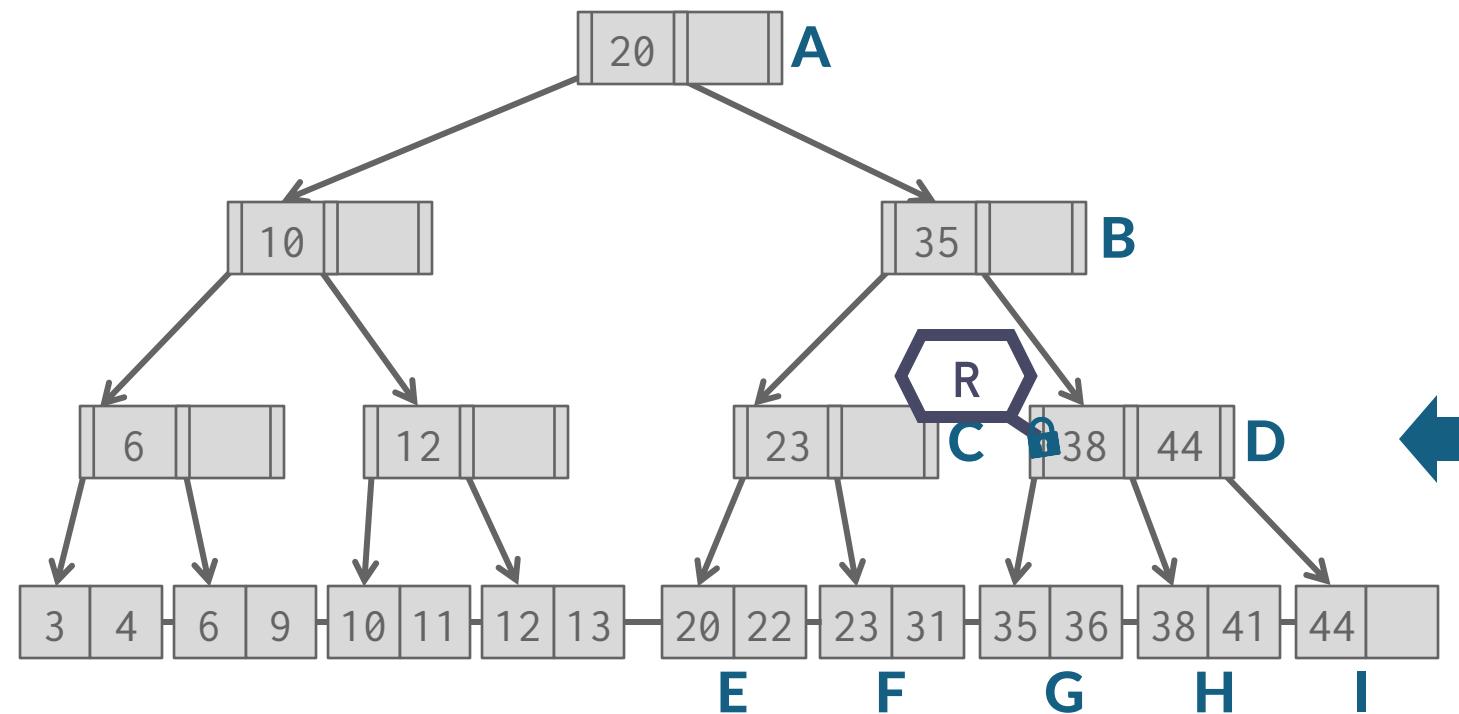
# Example #1 – Find 38



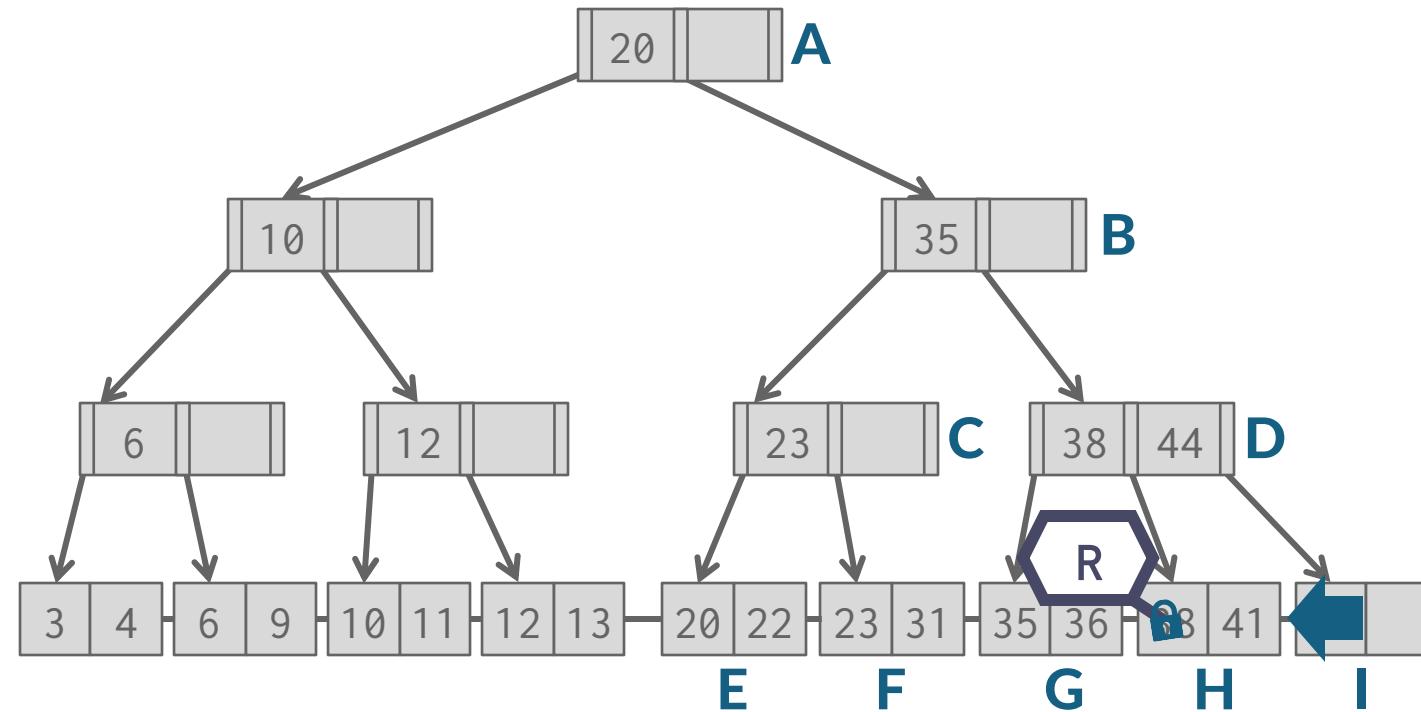
# Example #1 – Find 38



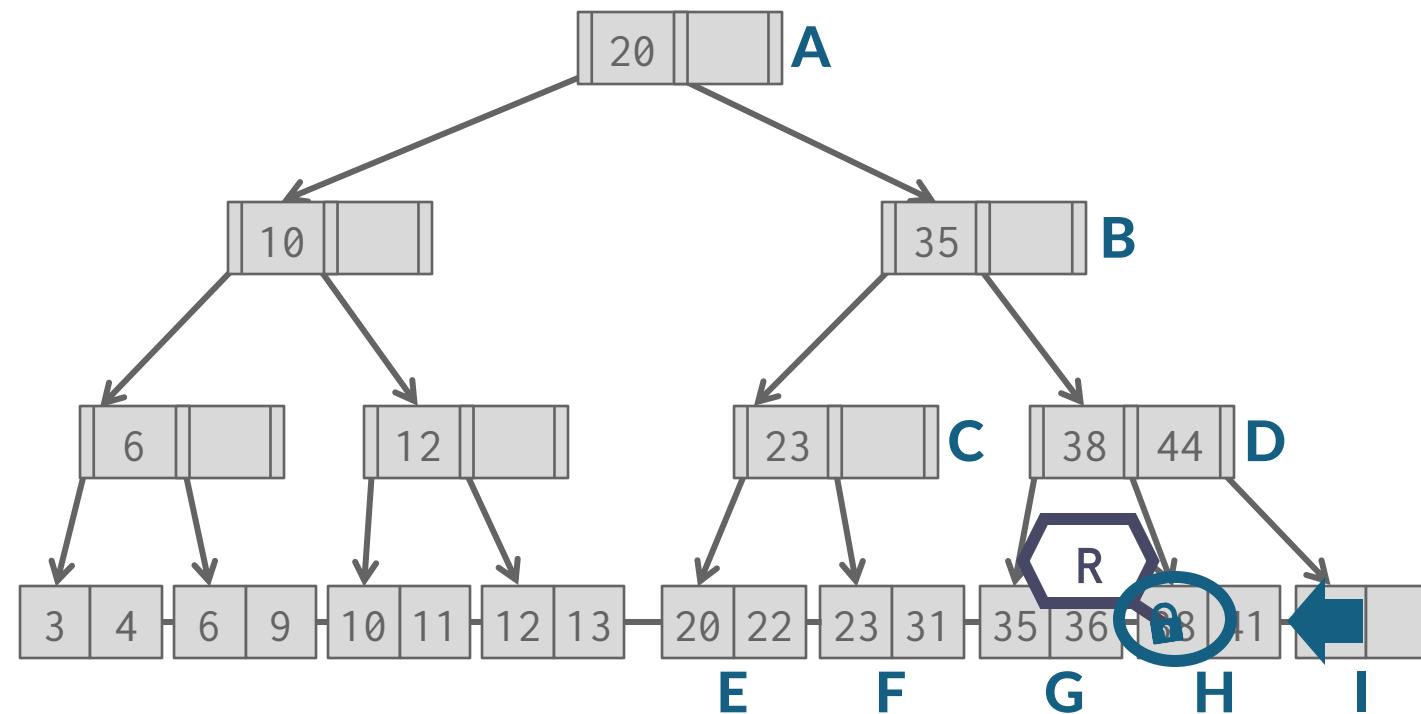
# Example #1 – Find 38



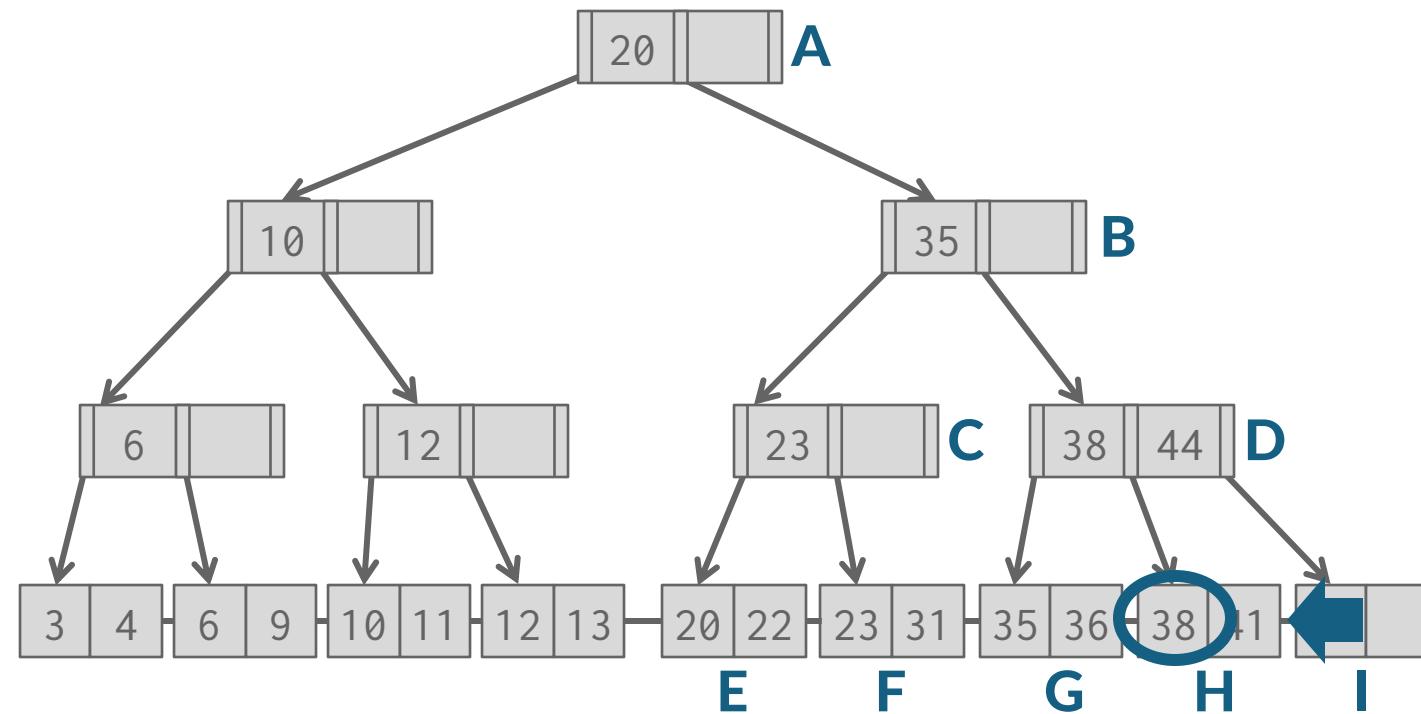
# Example #1 – Find 38



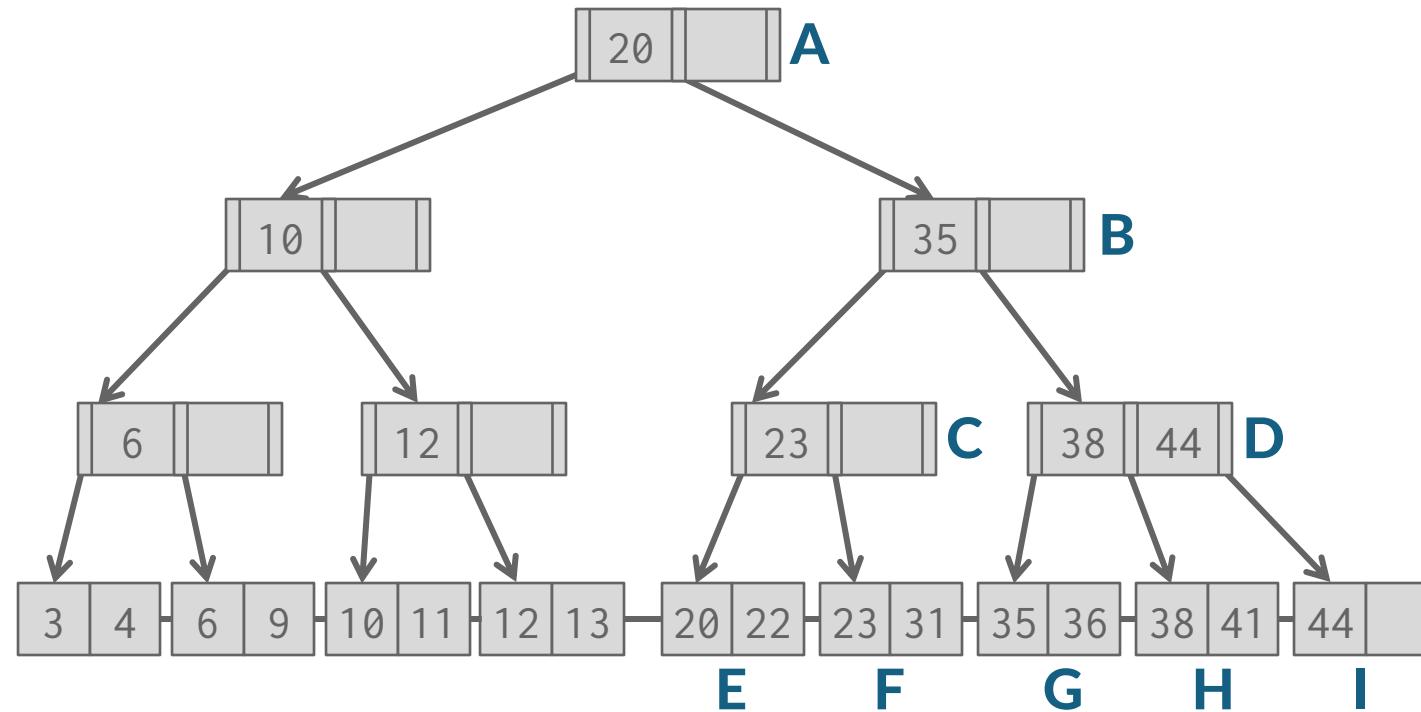
# Example #1 – Find 38



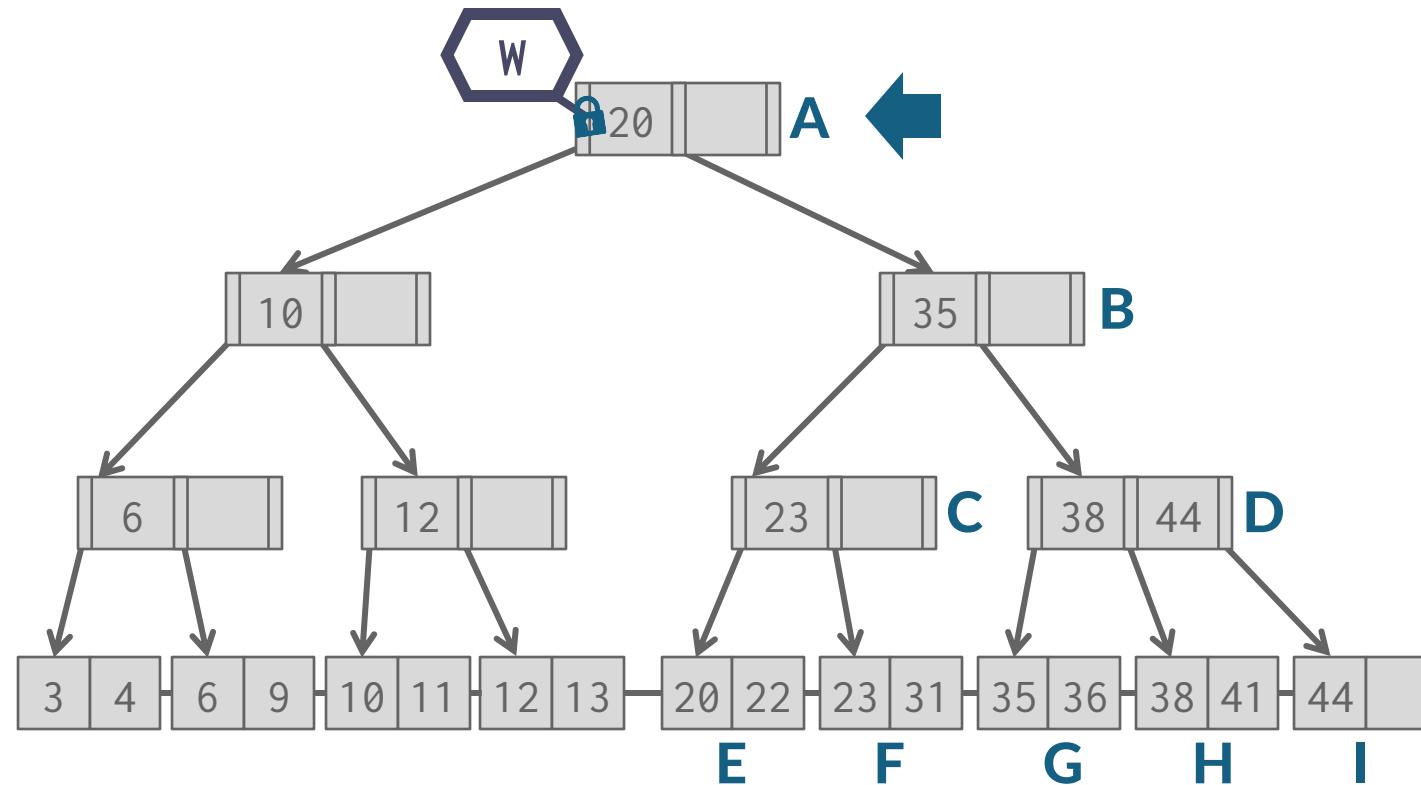
# Example #1 – Find 38



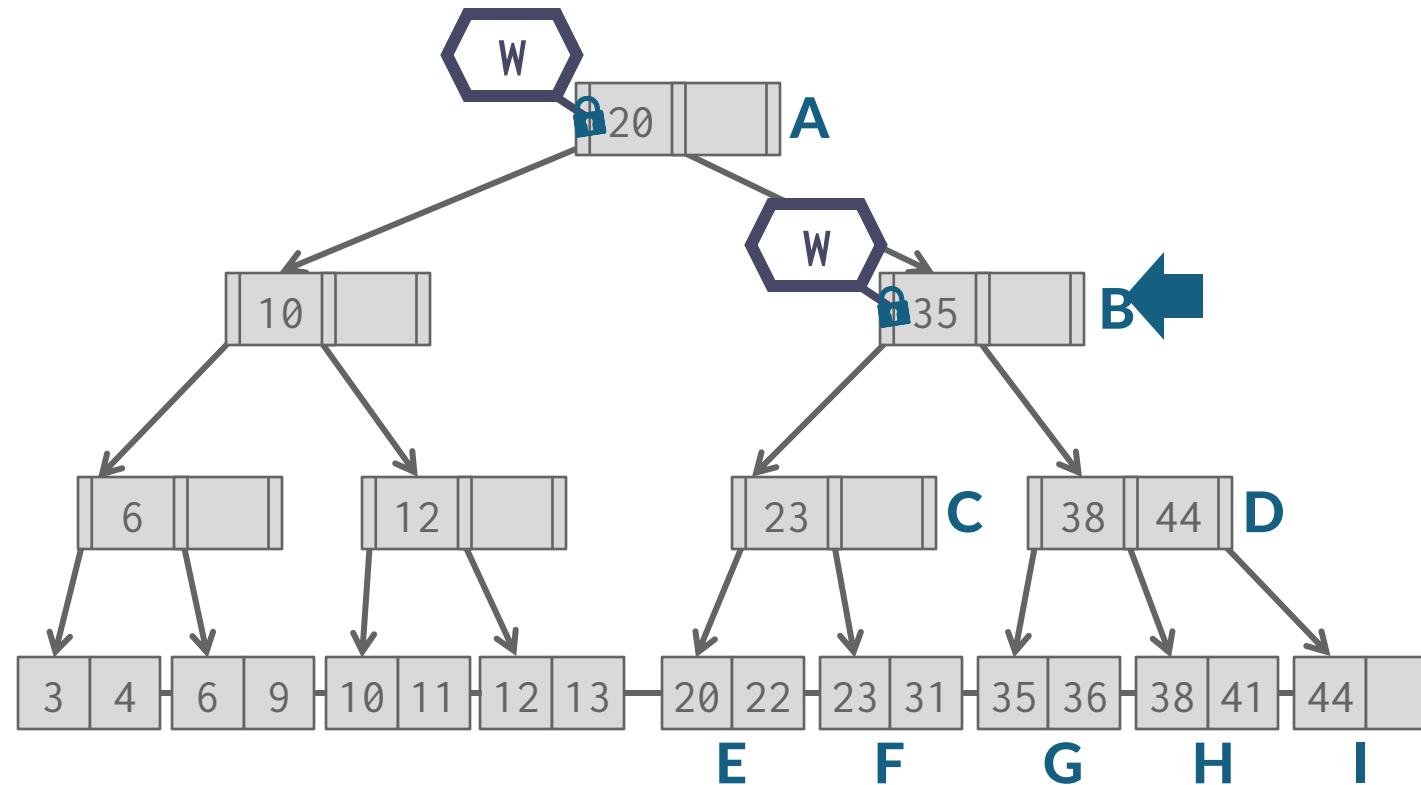
# Example #1 – Delete 38



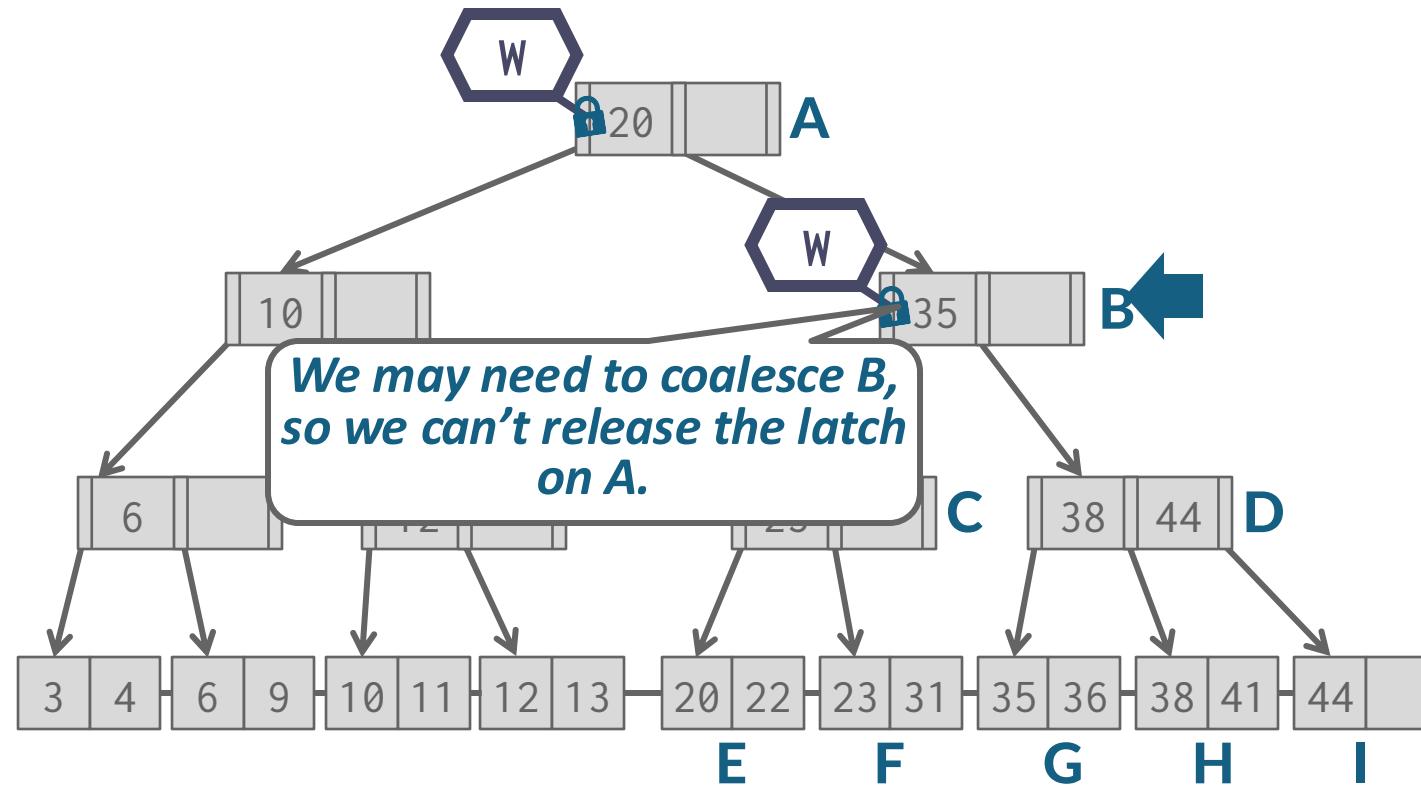
# Example #1 – Delete 38



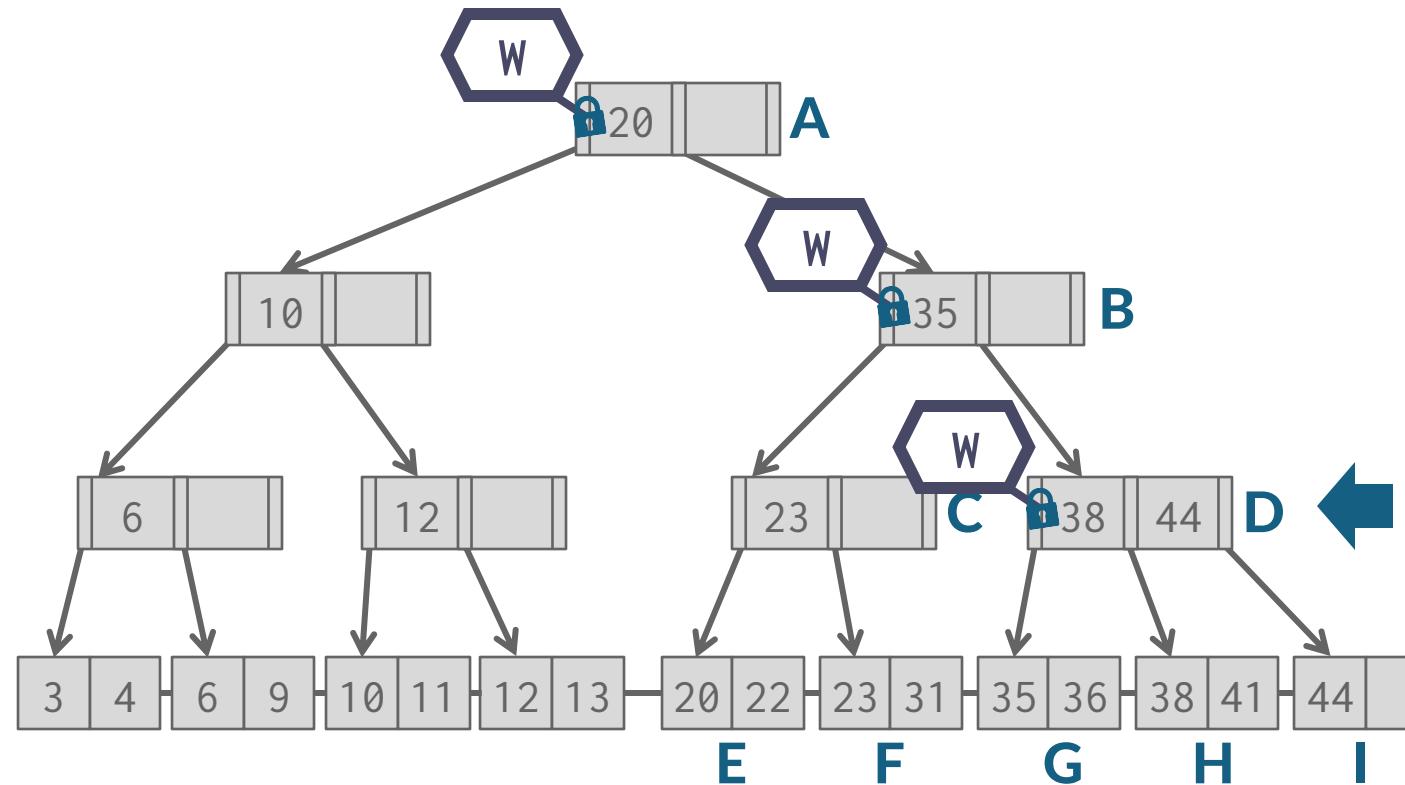
# Example #1 – Delete 38



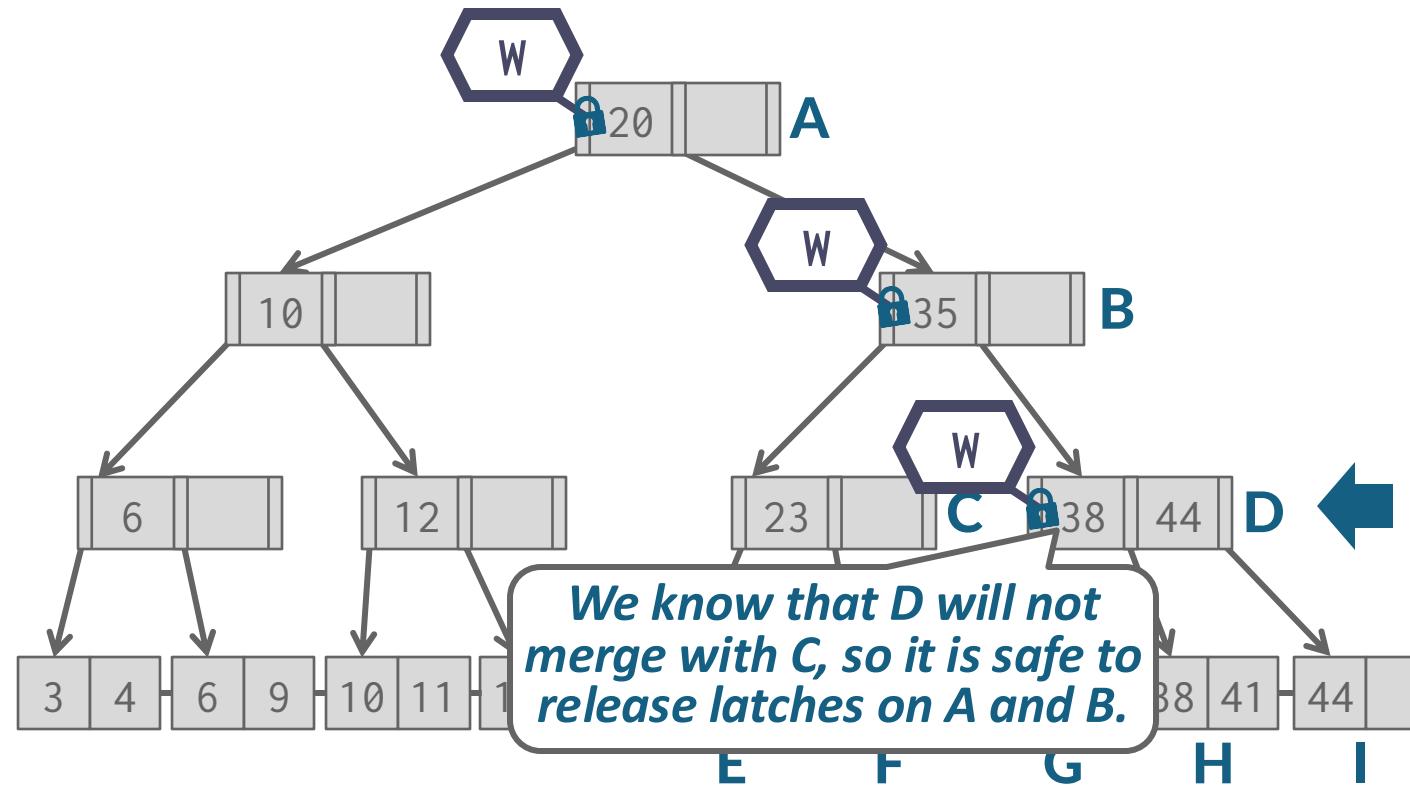
# Example #1 – Delete 38



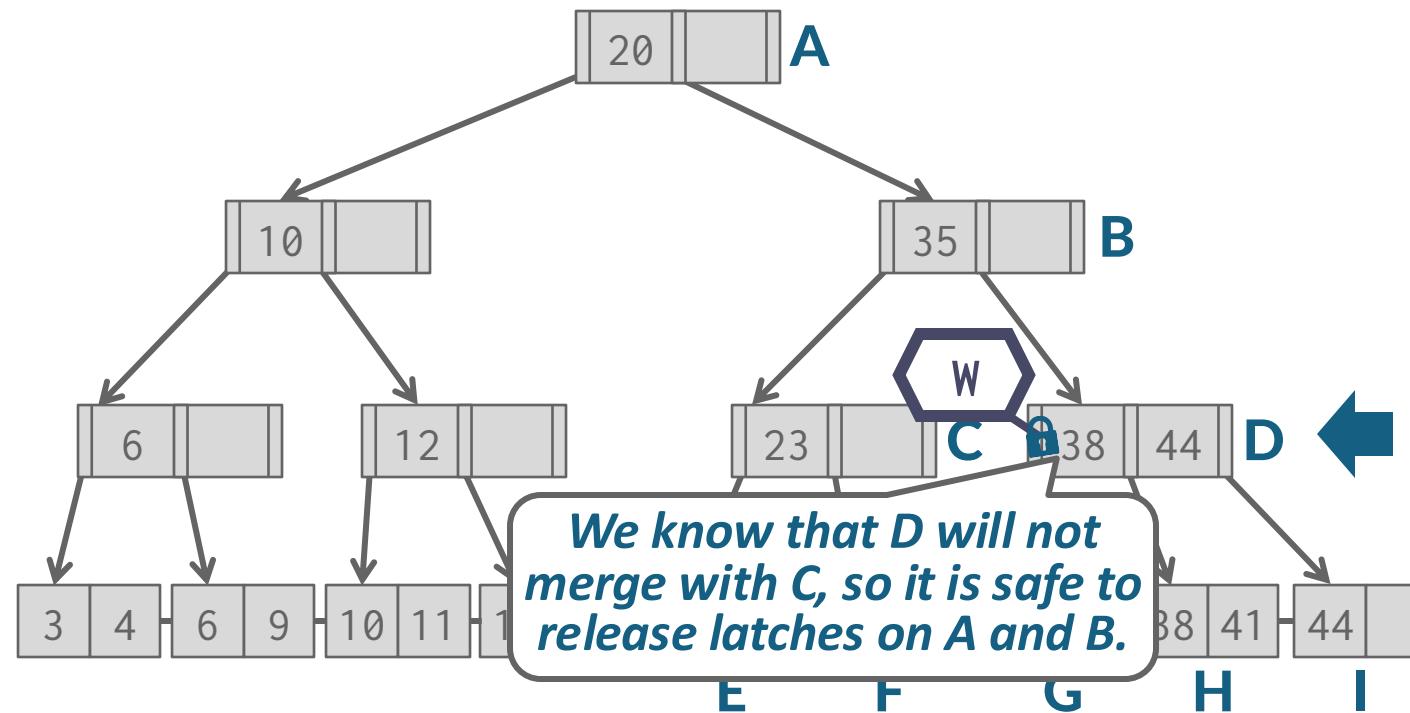
# Example #1 – Delete 38



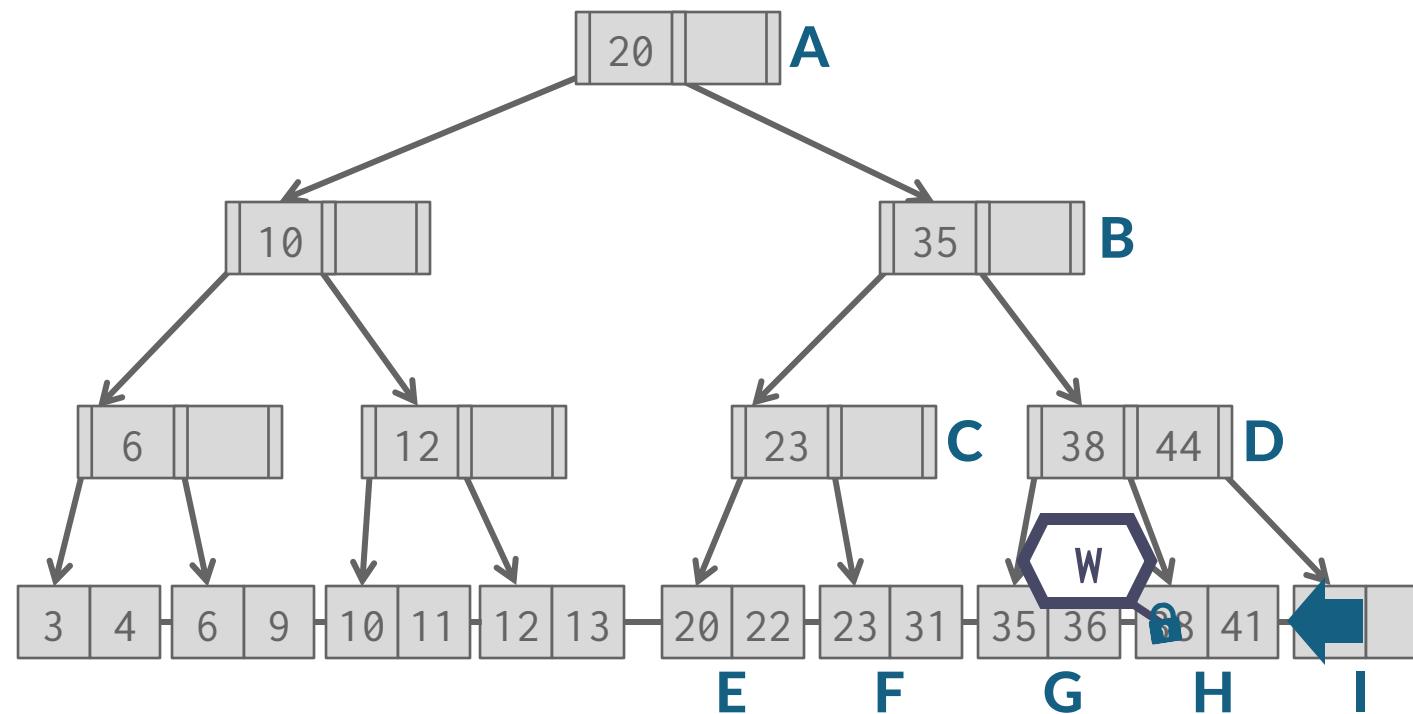
# Example #1 – Delete 38



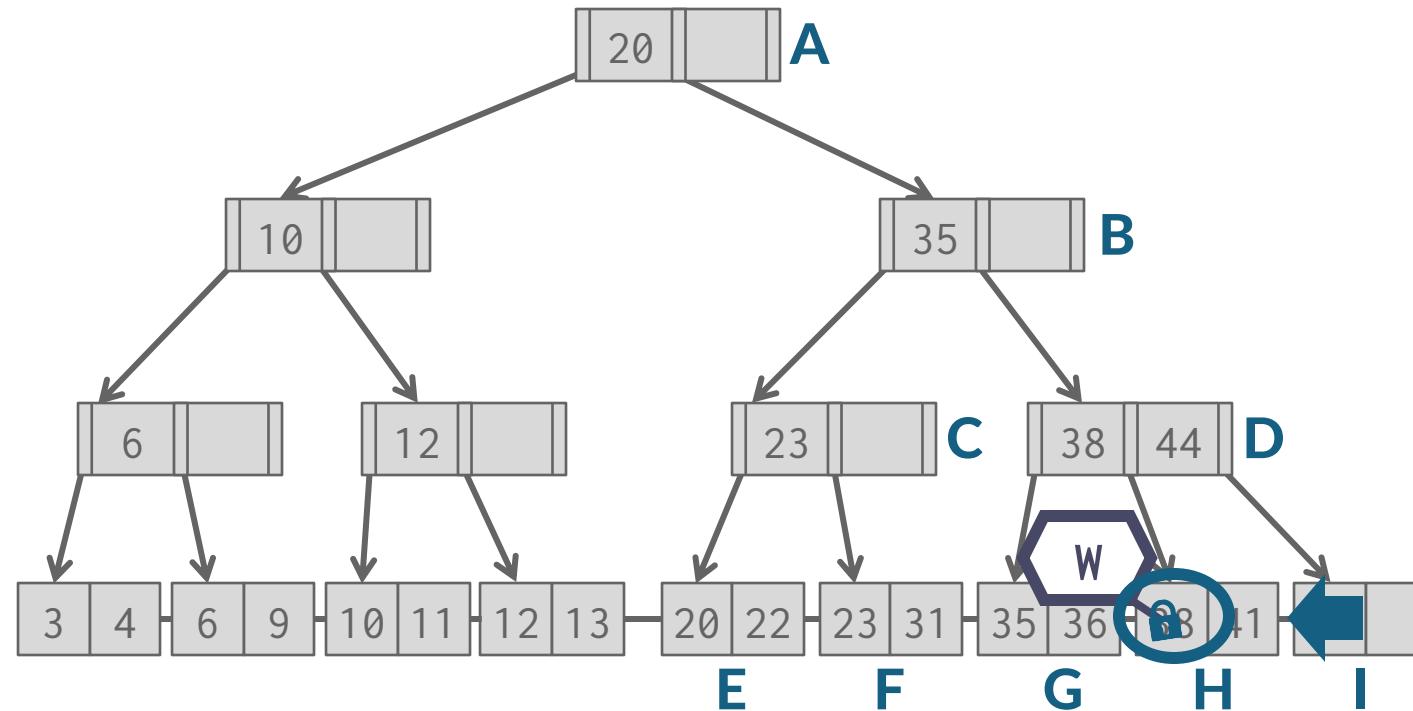
# Example #1 – Delete 38



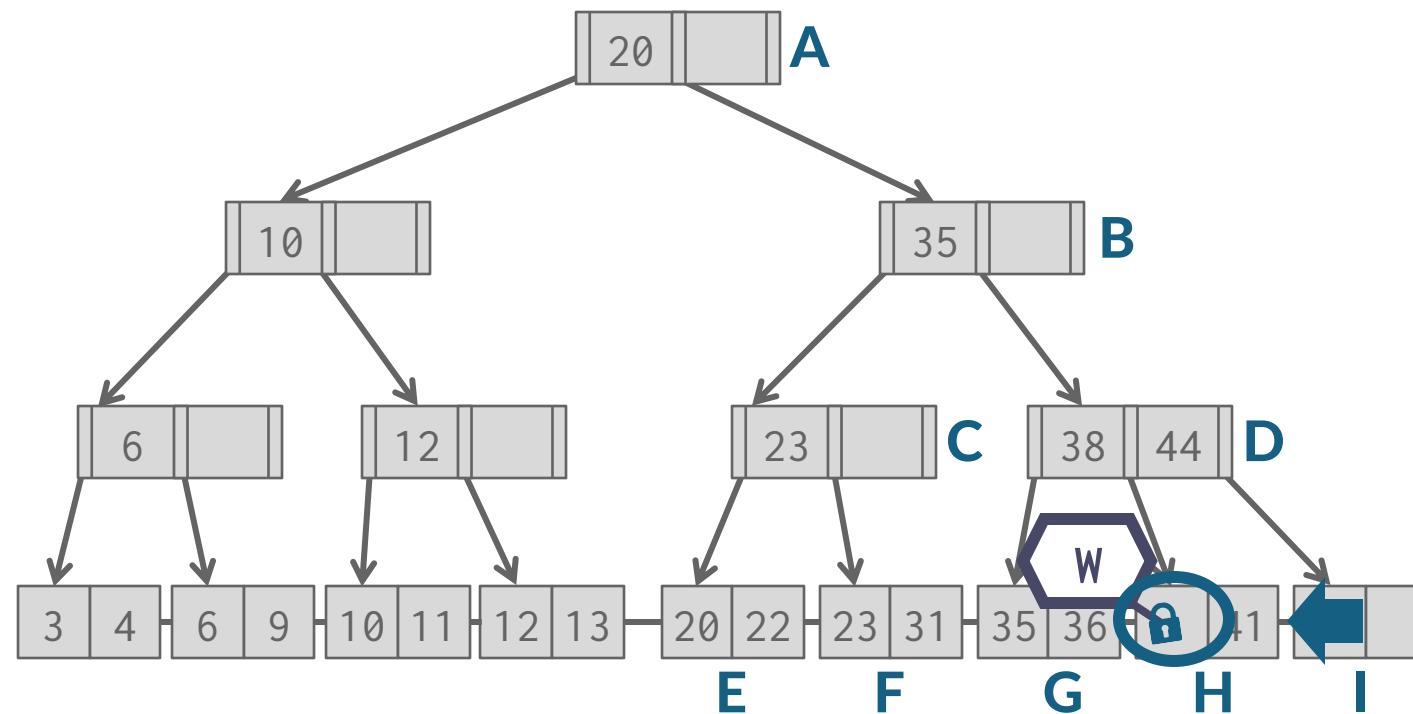
# Example #1 – Delete 38



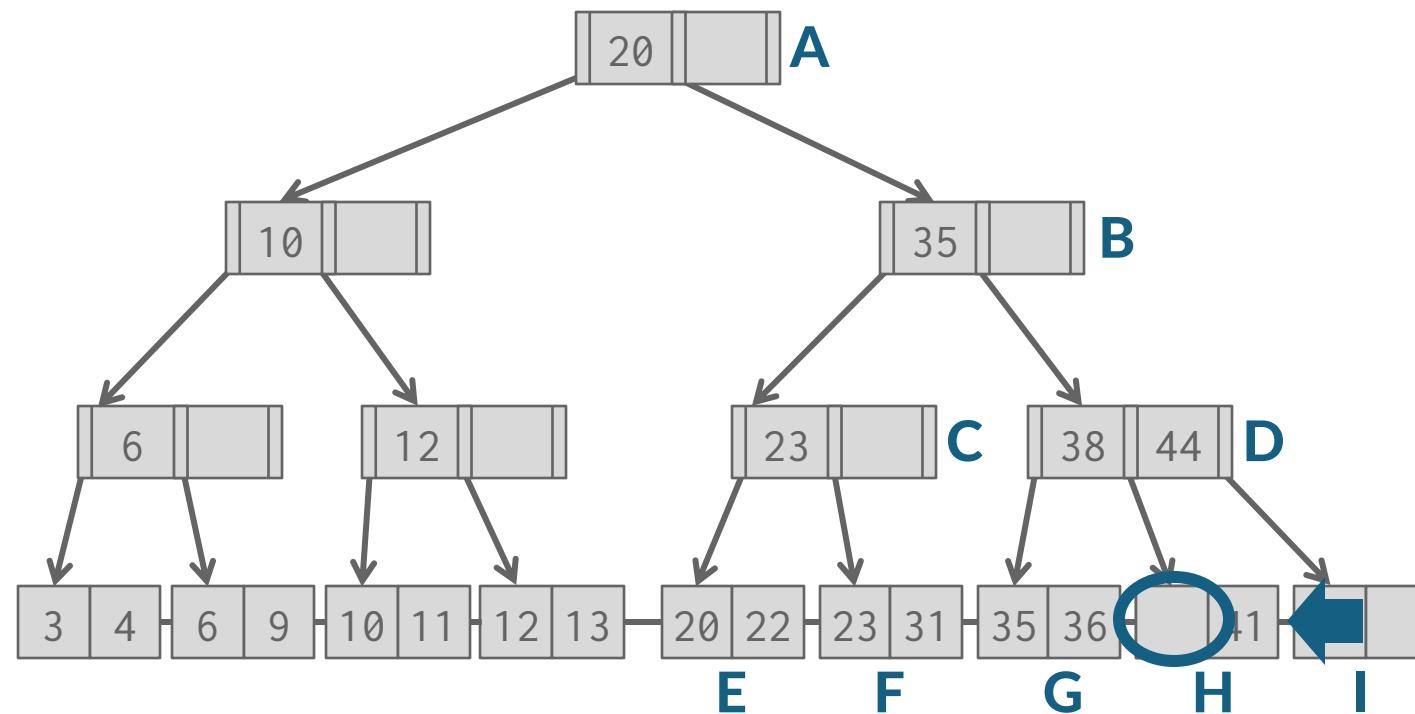
# Example #1 – Delete 38



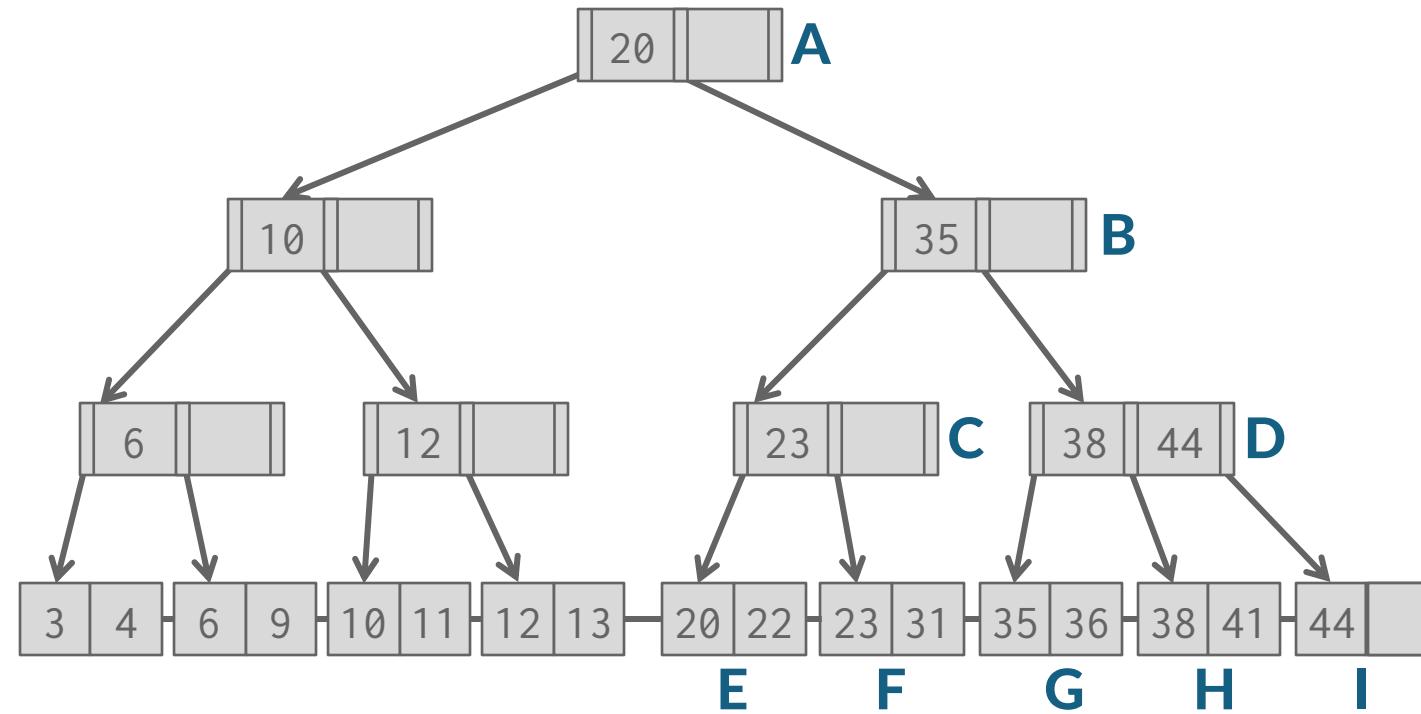
# Example #1 – Delete 38



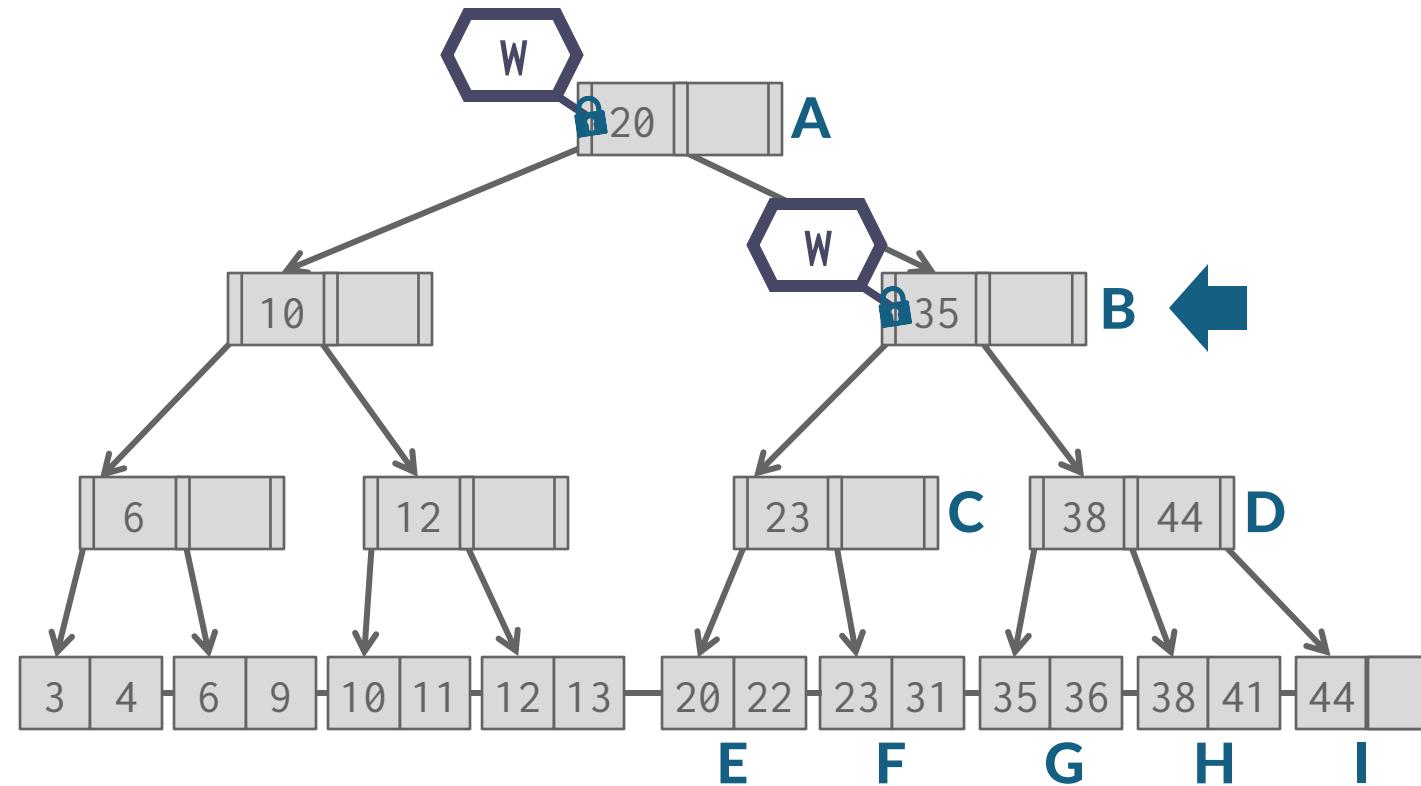
# Example #1 – Delete 38



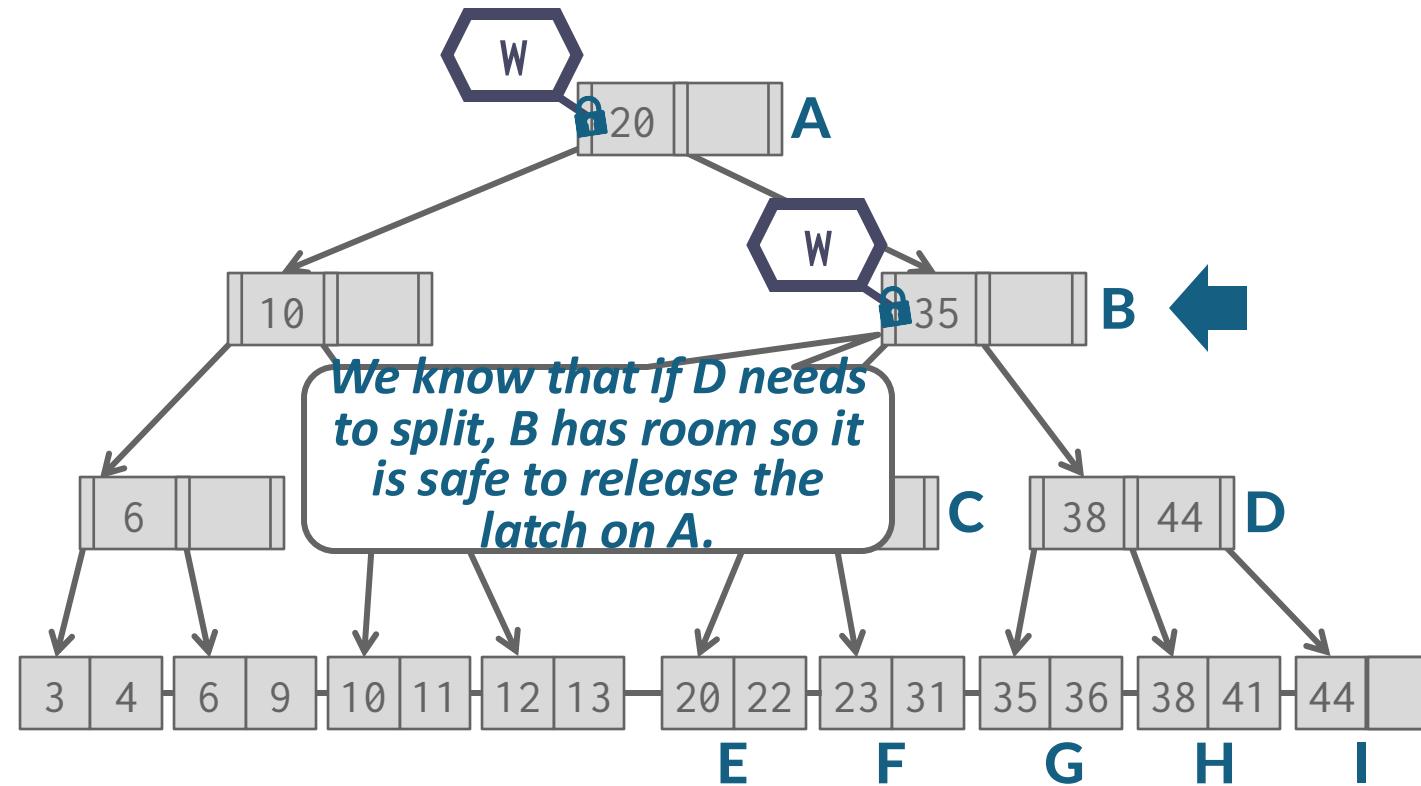
# Example #1 – Insert 45



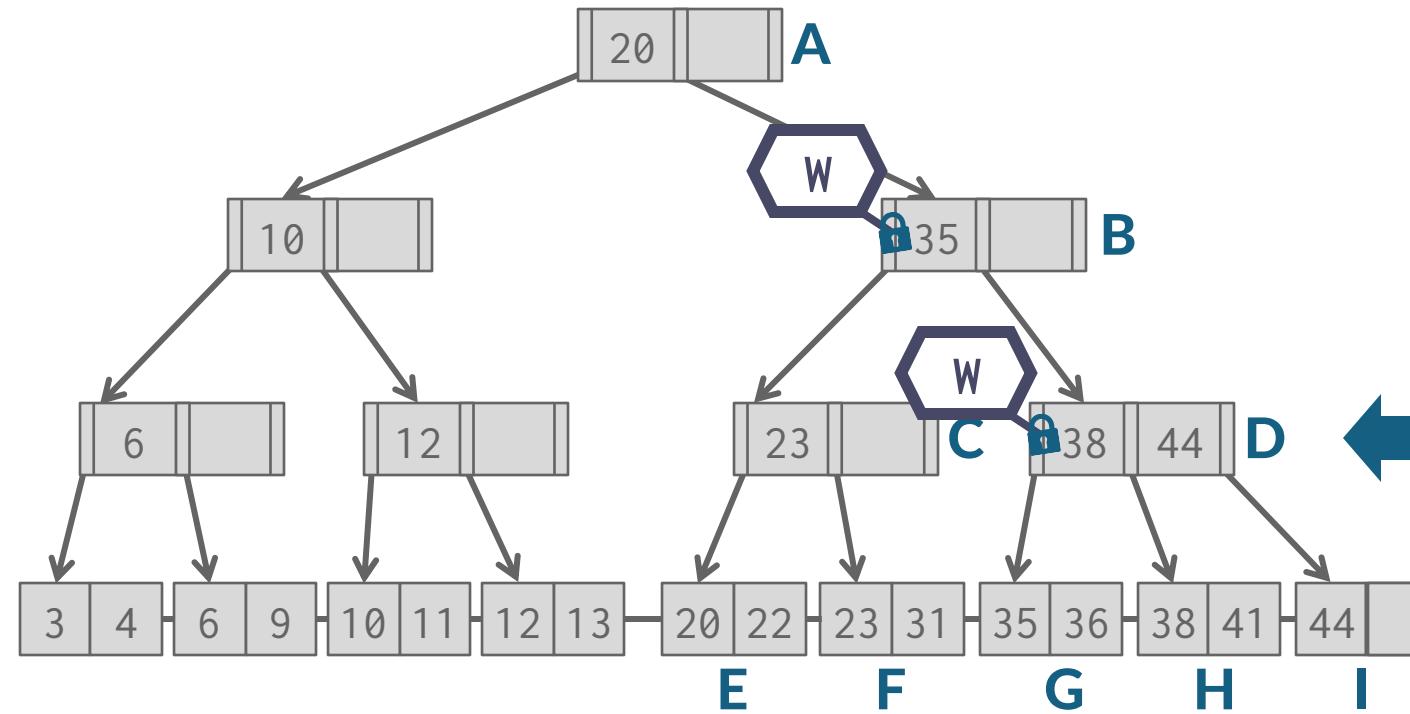
# Example #1 – Insert 45



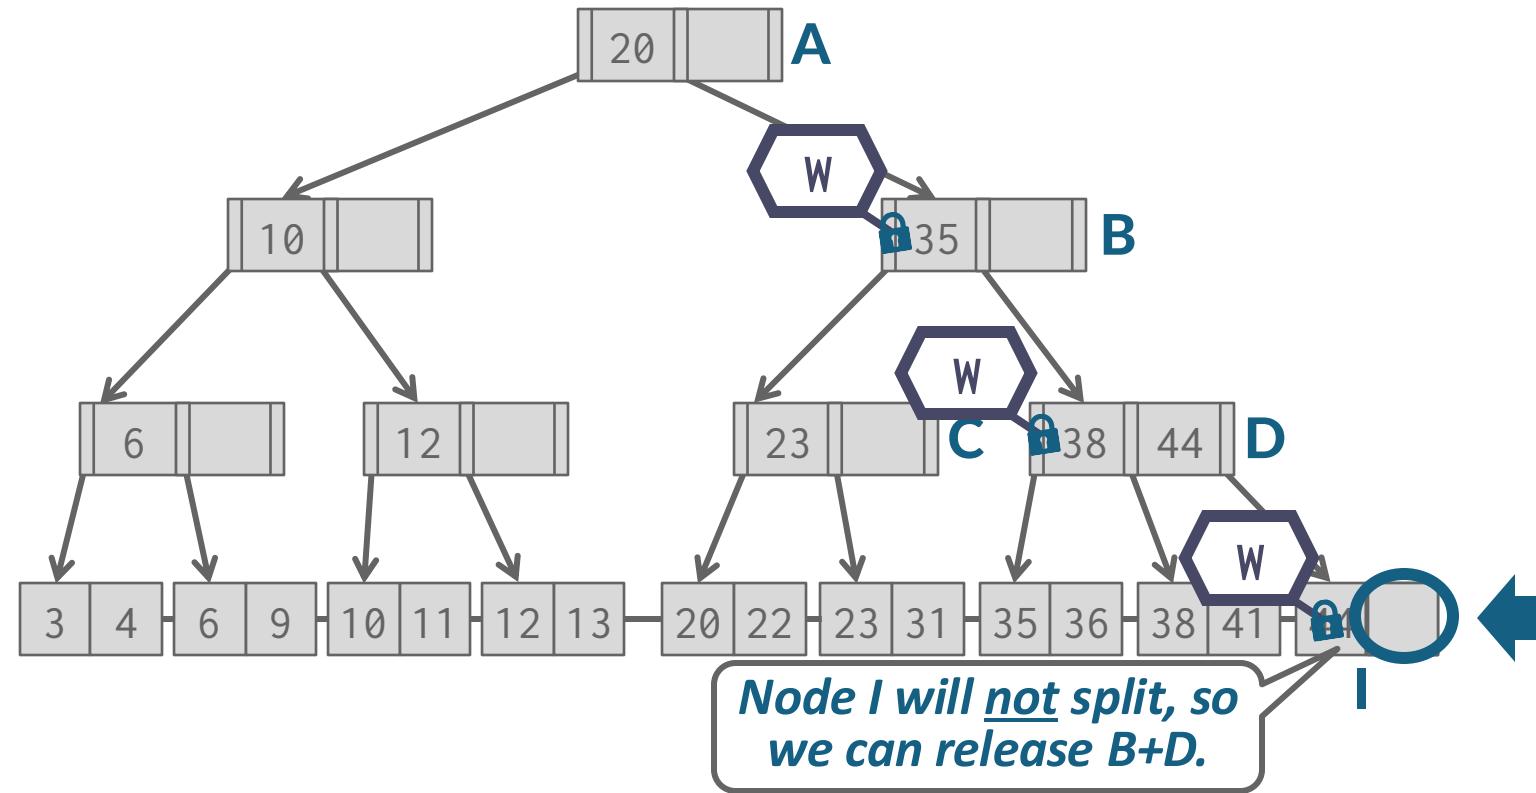
# Example #1 – Insert 45



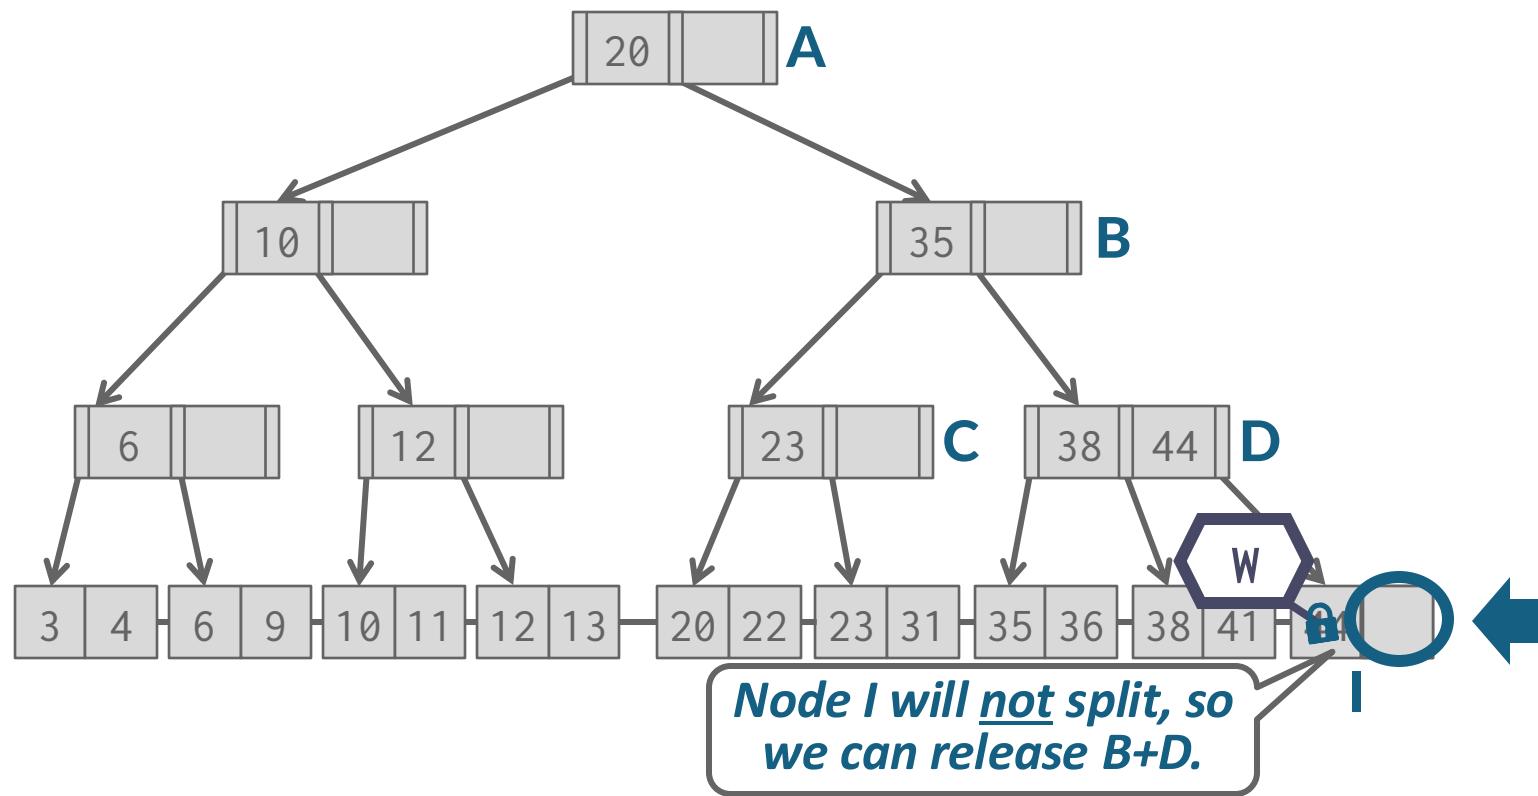
# Example #1 – Insert 45



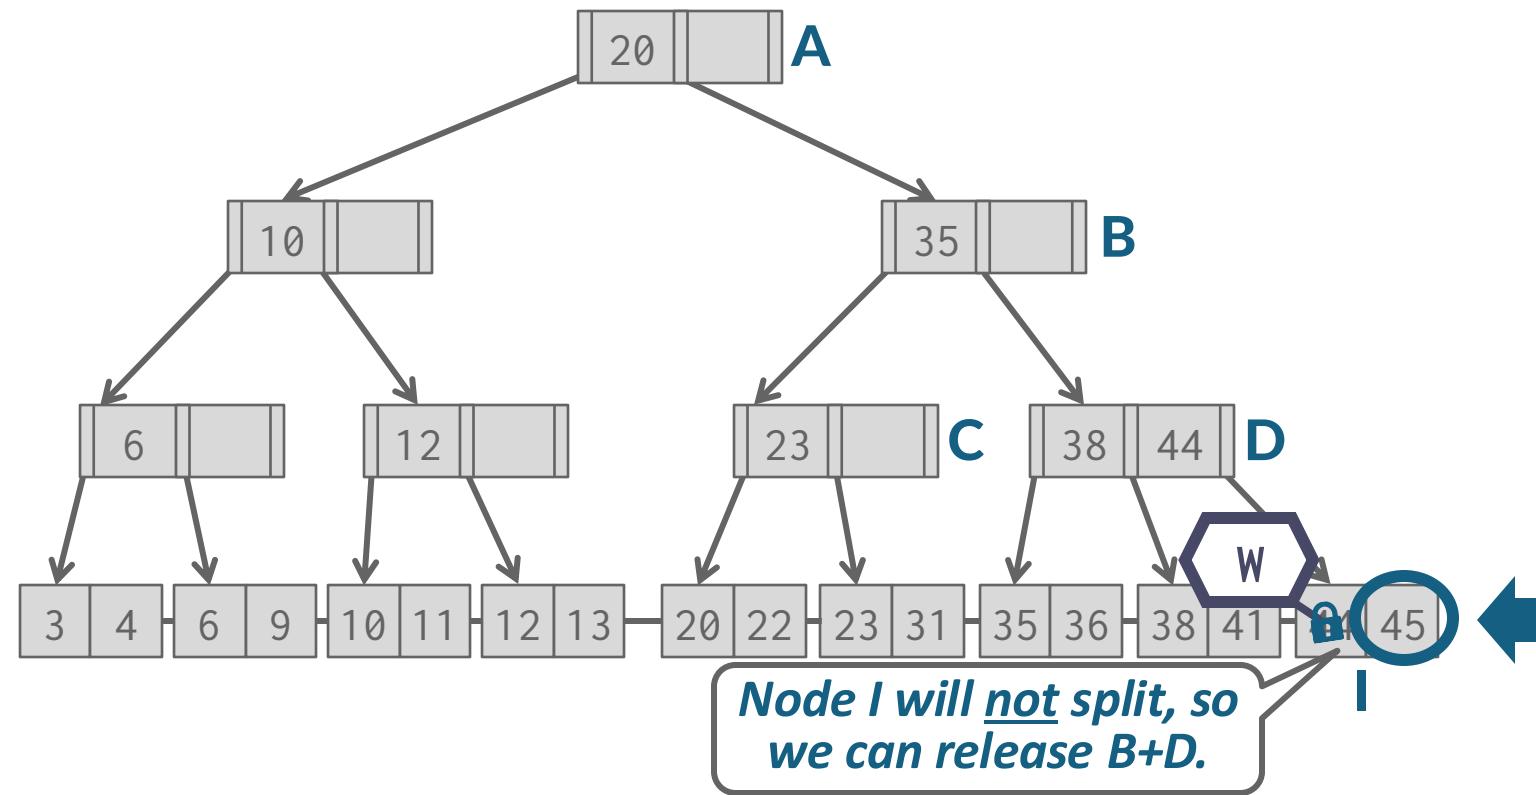
# Example #1 – Insert 45



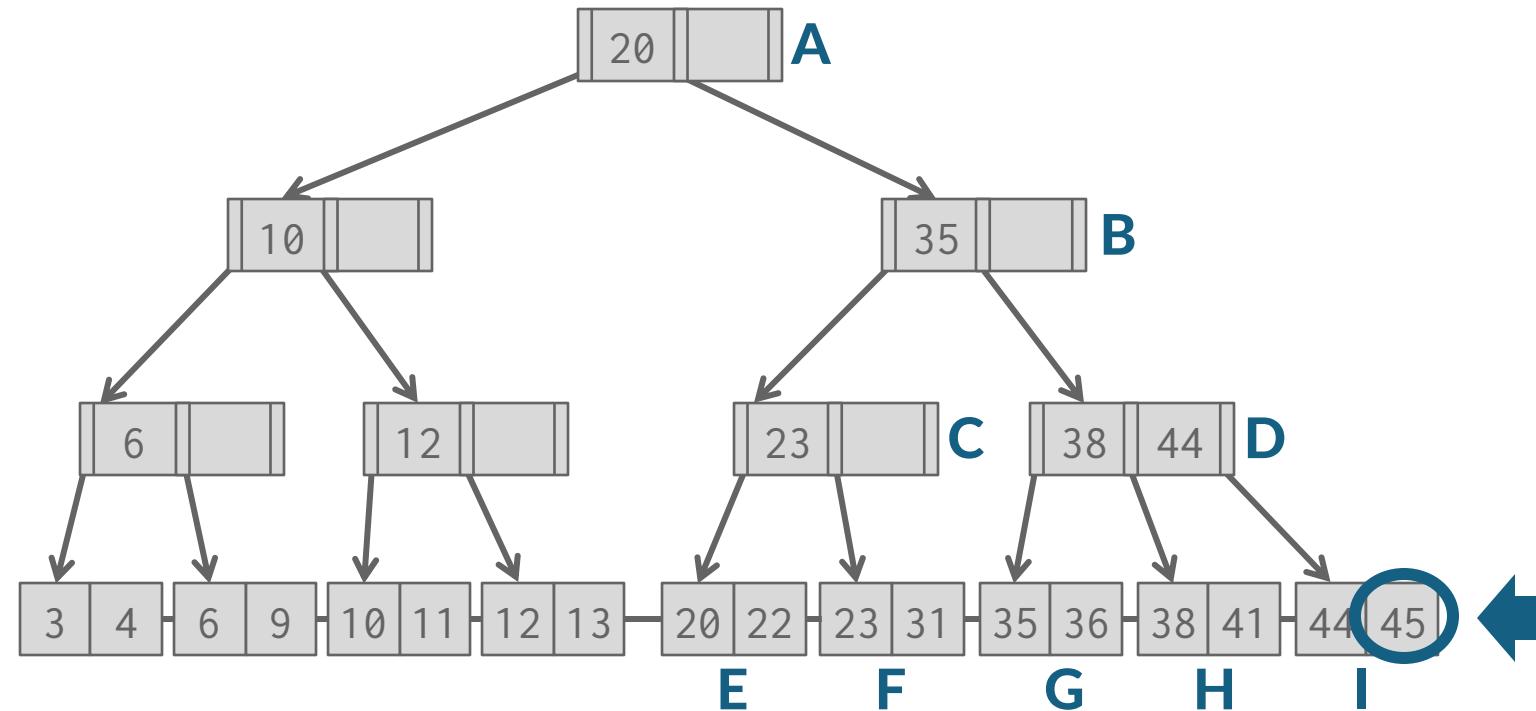
# Example #1 – Insert 45



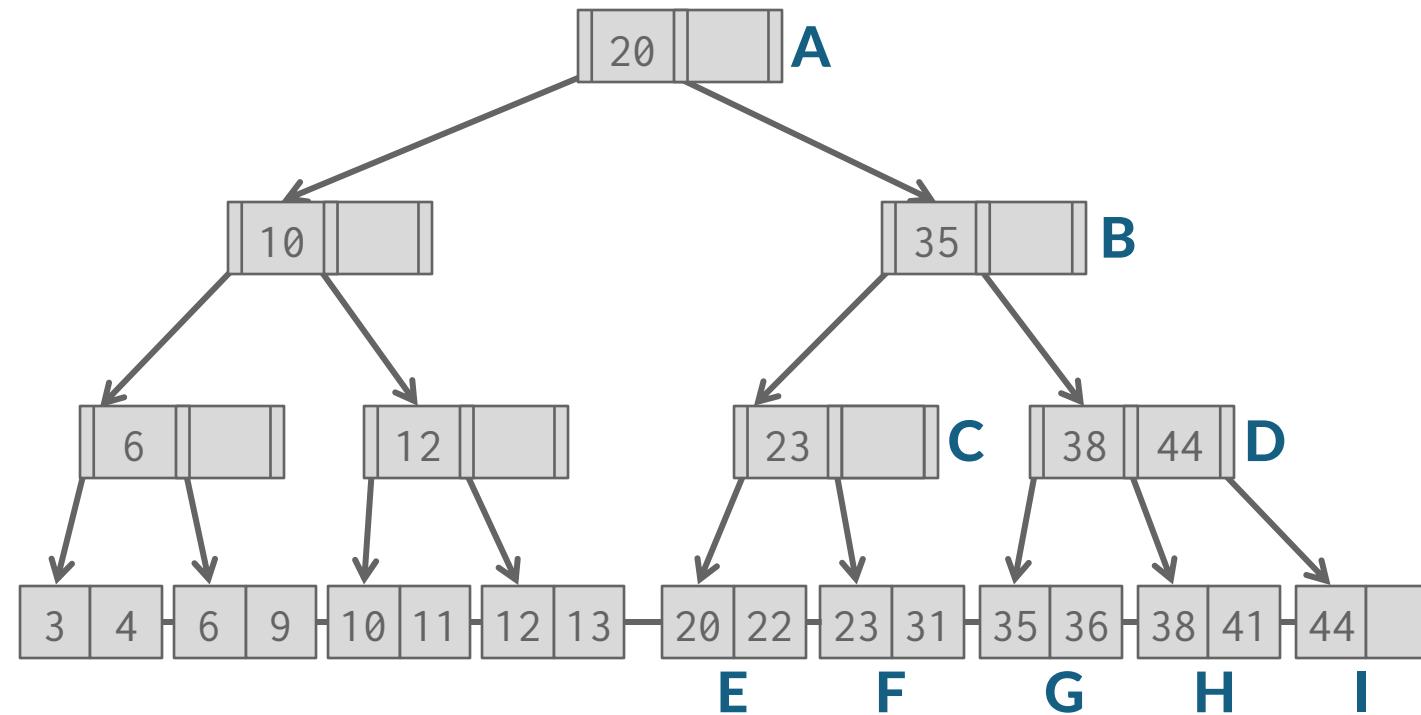
# Example #1 – Insert 45



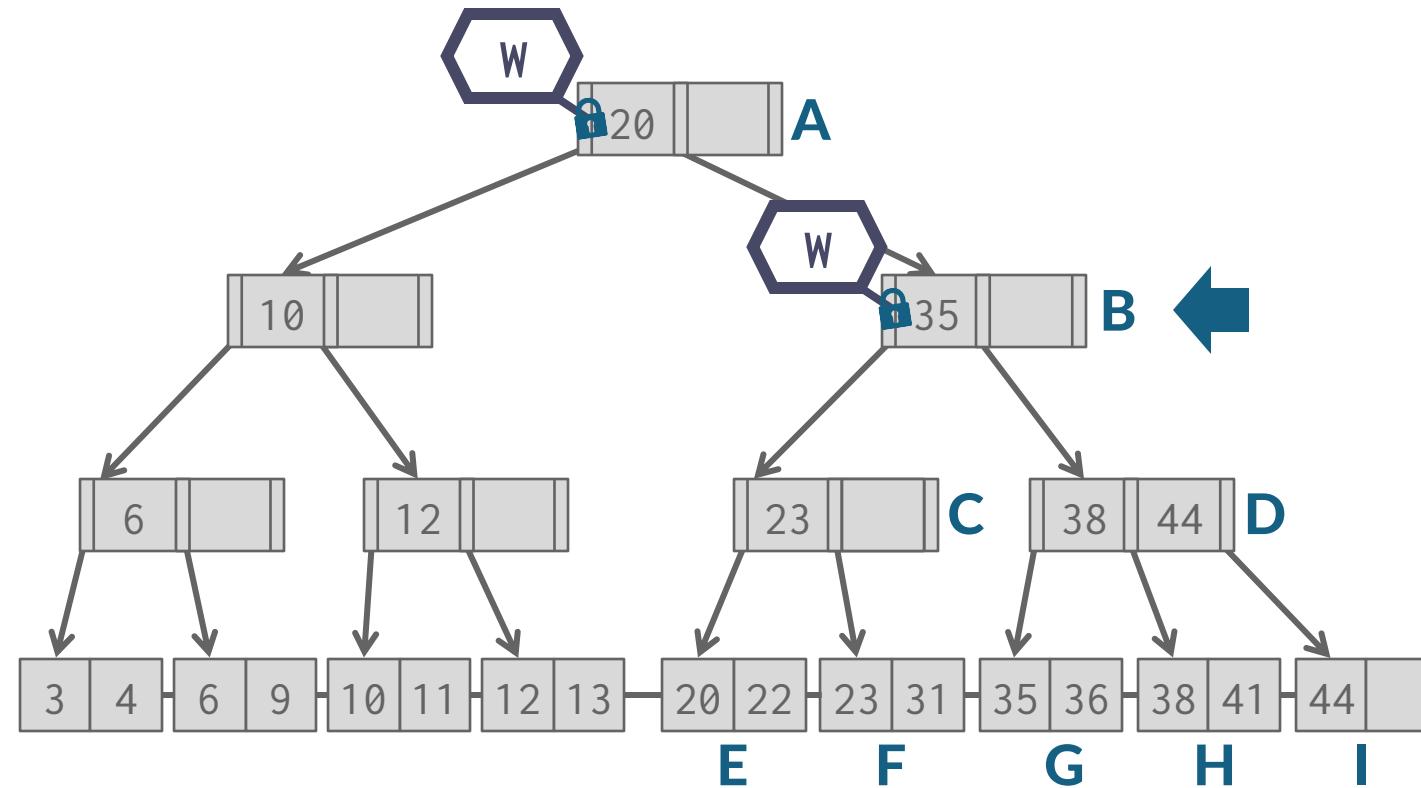
# Example #1 – Insert 45



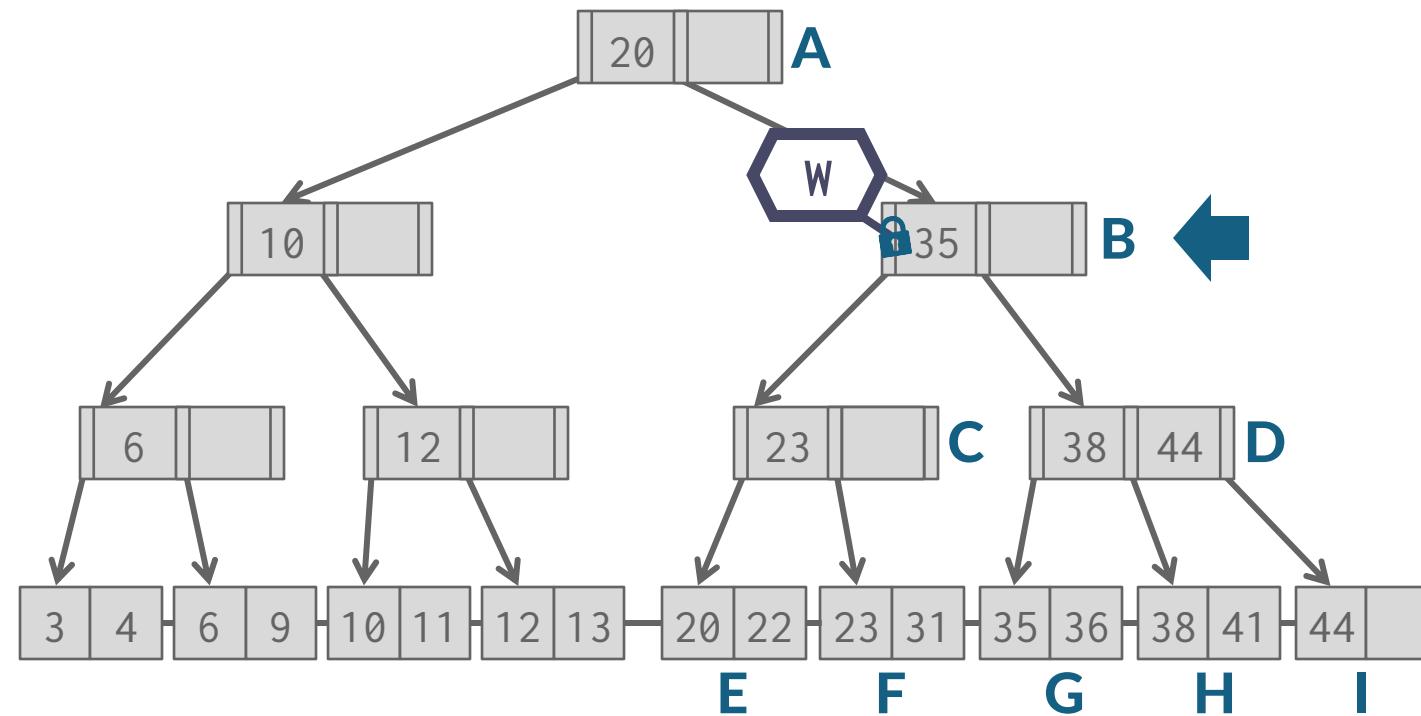
# Example #1 – Insert 25



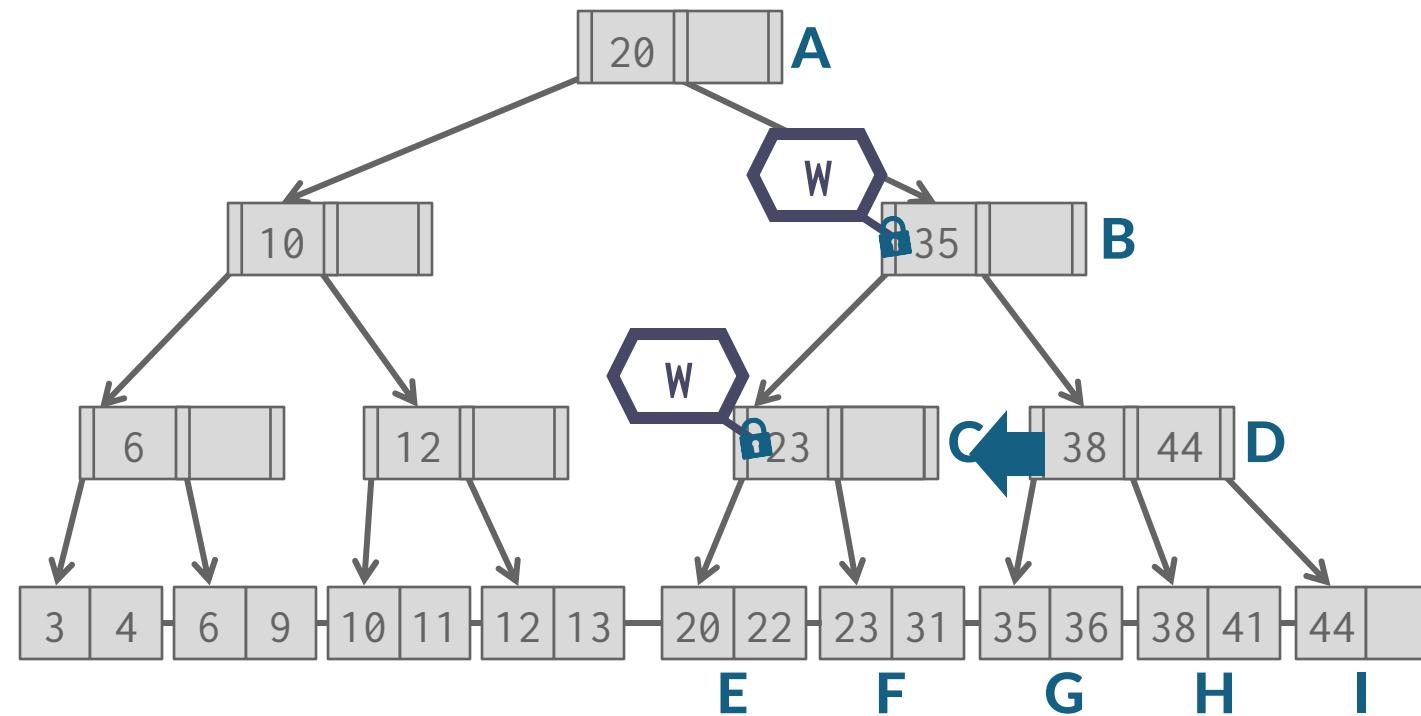
# Example #1 – Insert 25



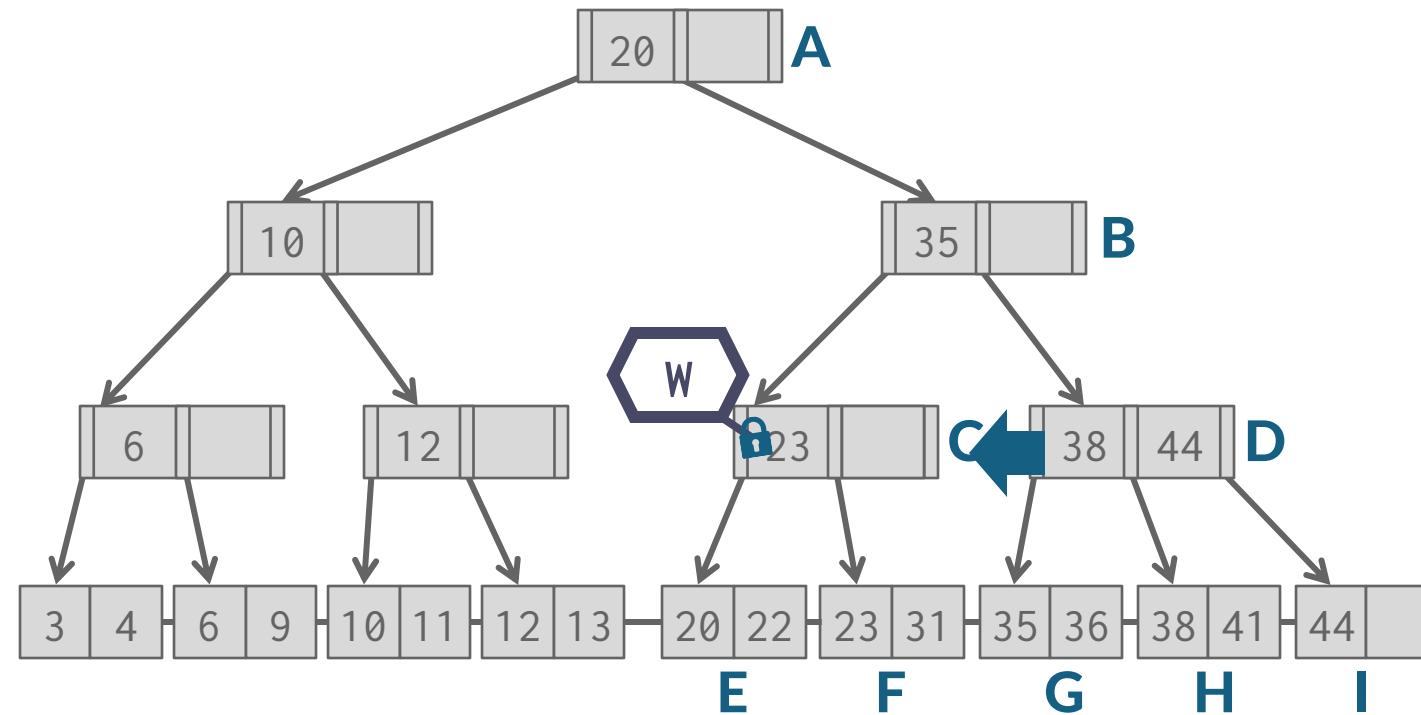
# Example #1 – Insert 25



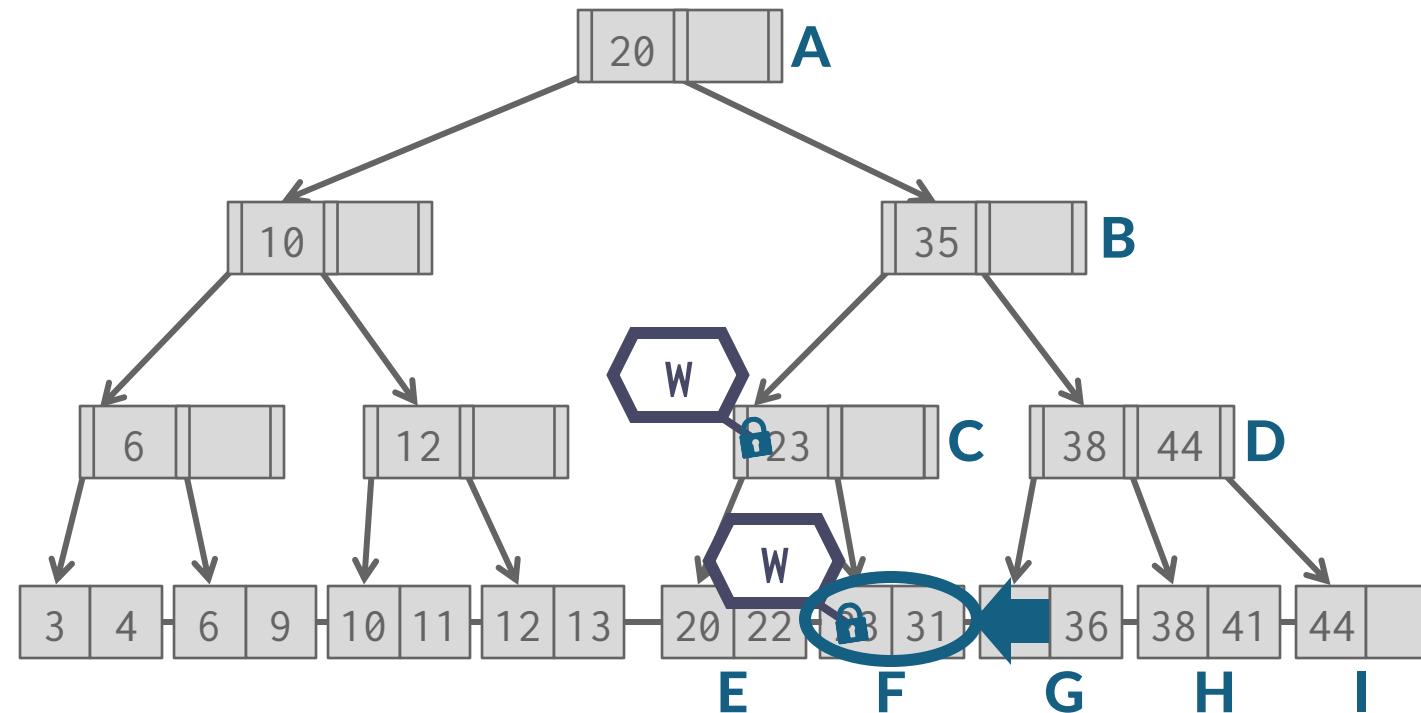
# Example #1 – Insert 25



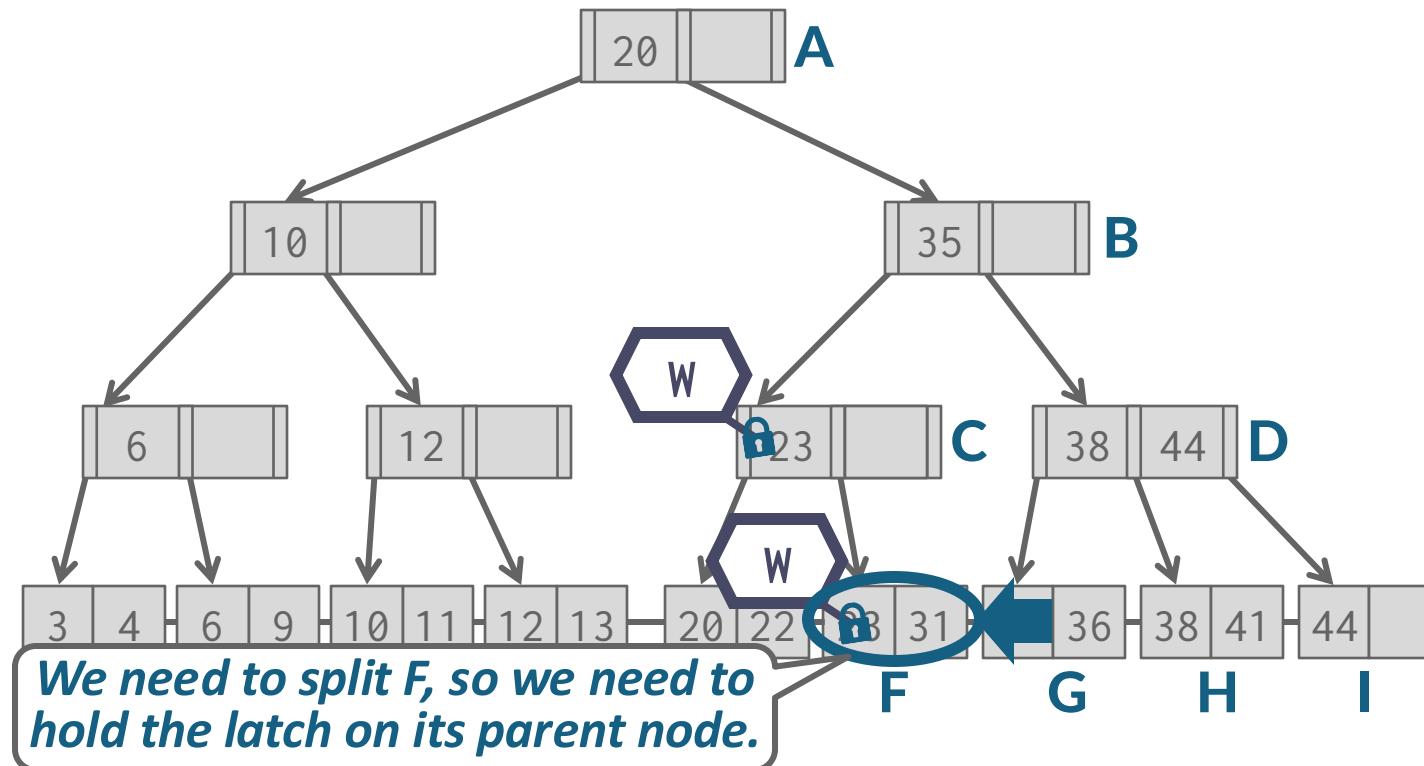
# Example #1 – Insert 25



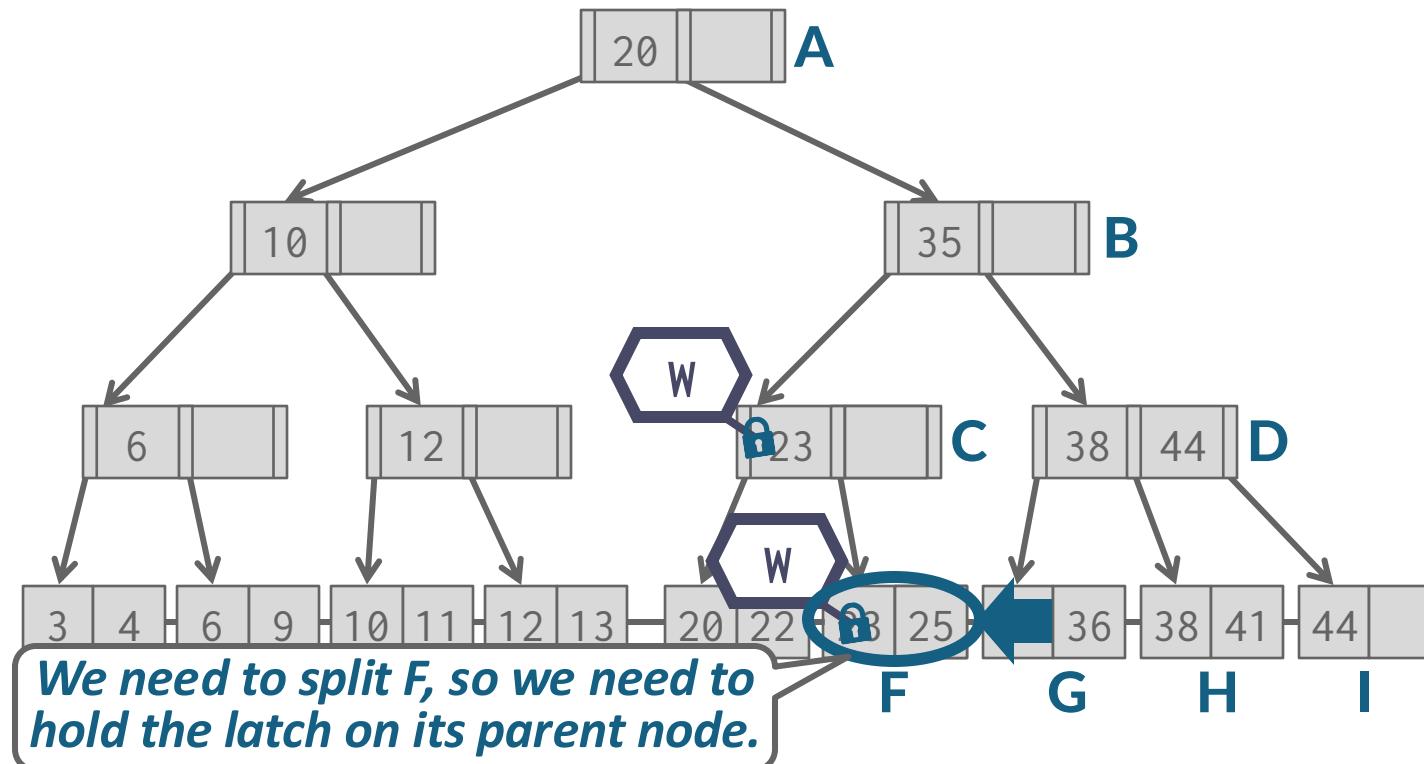
# Example #1 – Insert 25



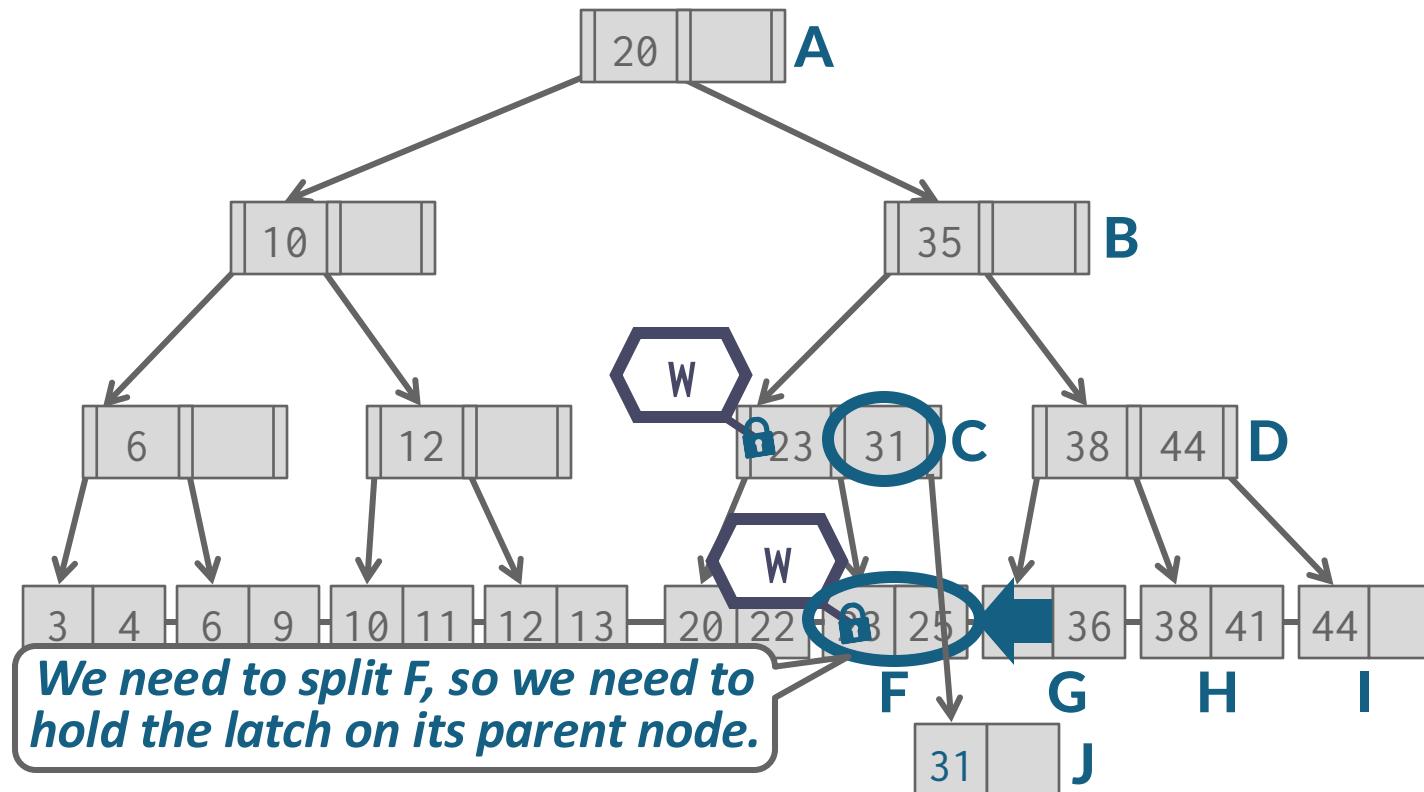
# Example #1 – Insert 25



# Example #1 – Insert 25



# Example #1 – Insert 25

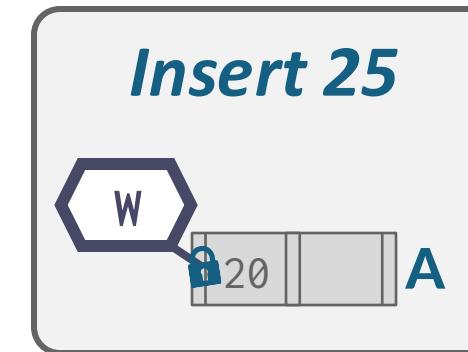
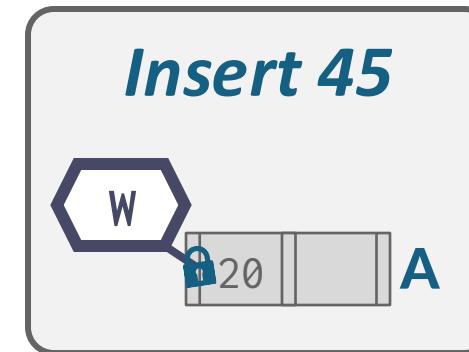
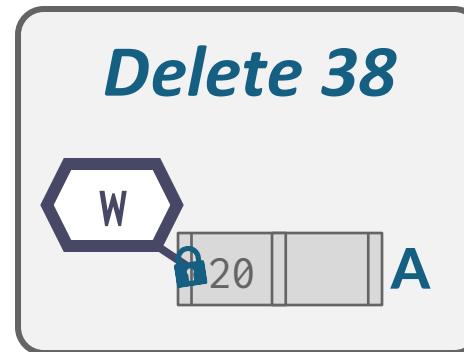


# Observation

- What was the first step that all the update examples did on the B+Tree?

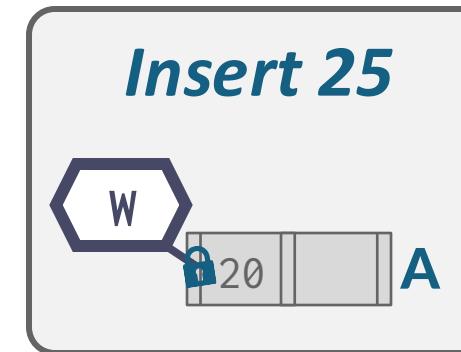
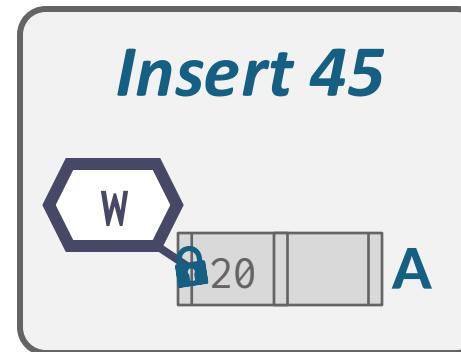
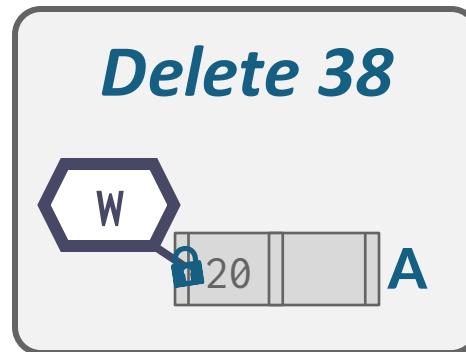
# Observation

- What was the first step that all the update examples did on the B+Tree?



# Observation

- What was the first step that all the update examples did on the B+Tree?



- Taking a write latch on the root every time becomes a bottleneck with higher concurrency.

# Better Latching Algorithm

- Most modifications to a B+Tree will not require a split or merge.
- Instead of assuming that there will be a split/merge, optimistically traverse the tree using read latches.
- If you guess wrong, repeat traversal with the pessimistic algorithm.

Acta Informatica 9, 1–21 (1977)

  
Springer-Verlag 1977

---

**Concurrency of Operations on B-Trees**

R. Bayer\* and M. Schkolnick  
IBM Research Laboratory, San Jose, CA 95193, USA

**Summary.** Concurrent operations on *B*-trees pose the problem of insuring that each operation can be carried out without interfering with other operations being performed simultaneously by other users. This problem can become critical if these structures are being used to support access paths, like indexes, to data base systems. In this case, serializing access to one of these indexes can create an unacceptable bottleneck for the entire system. Thus, there is a need for locking protocols that can assure integrity for each access while at the same time providing a maximum possible degree of concurrency. Another feature required from these protocols is that they be deadlock free, since the cost to resolve a deadlock may be high.

Recently, there has been some questioning on whether *B*-tree structures can support concurrent operations. In this paper, we examine the problem of concurrent access to *B*-trees. We present a deadlock free solution which can be tuned to specific requirements. An analysis is presented which allows the selection of parameters so as to satisfy these requirements.

The solution presented here uses simple locking protocols. Thus, we conclude that *B*-trees can be used advantageously in a multi-user environment.

**I. Introduction**

In this paper, we examine the problem of concurrent access to indexes which are maintained as *B*-trees. This type of organization was introduced by Bayer and McCreight [2] and some variants of it appear in Knuth [10] and Wedekind [13]. Performance studies of it were restricted to the single user environment. Recently, these structures have been examined for possible use in a multi-user (concurrent) environment. Some initial studies have been made about the feasibility of their use in this type of situation [1], [6], and [11].

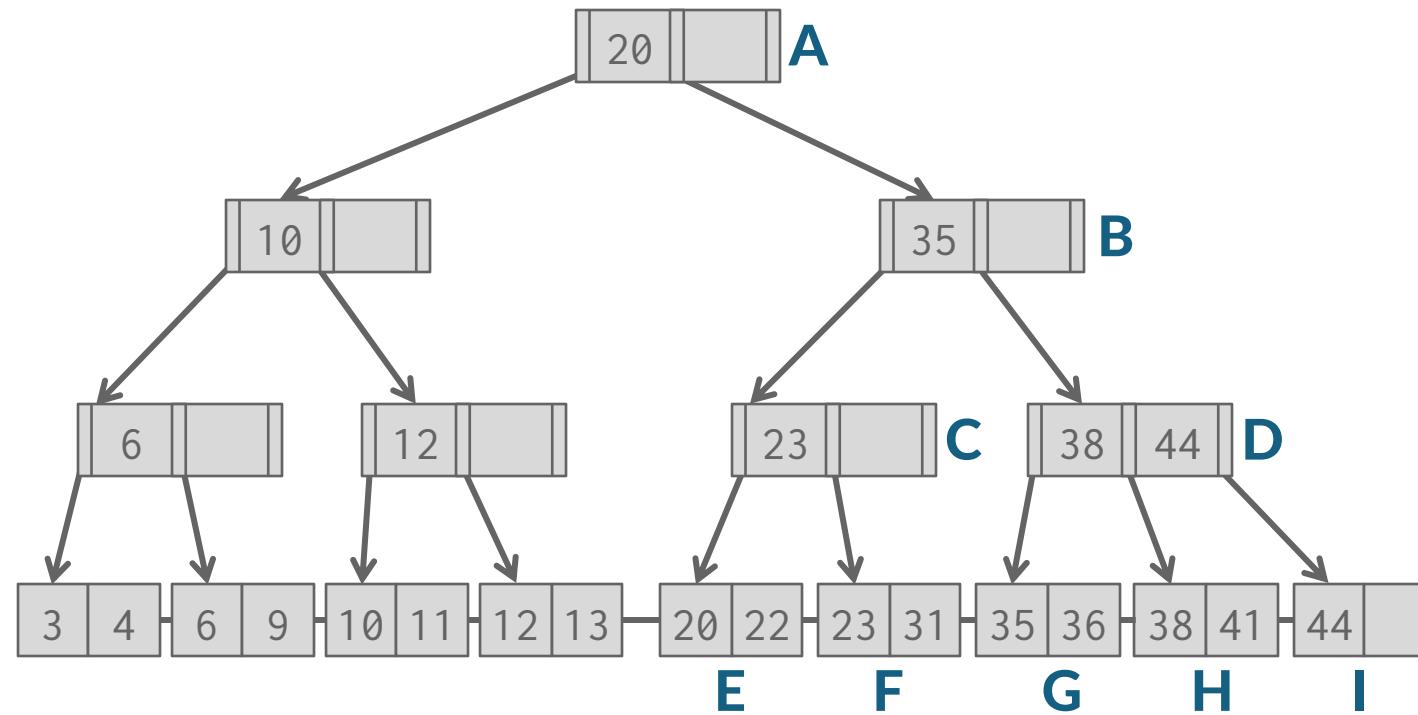
An accessing schema which achieves a high degree of concurrency in using the index will be presented. The schema allows dynamic tuning to adapt its performance to the profile of the current set of users. Another property of the

\* Permanent address: Institut für Informatik der Technischen Universität München, Arcistr. 21, D-8000 München 2, Germany (Fed. Rep.)

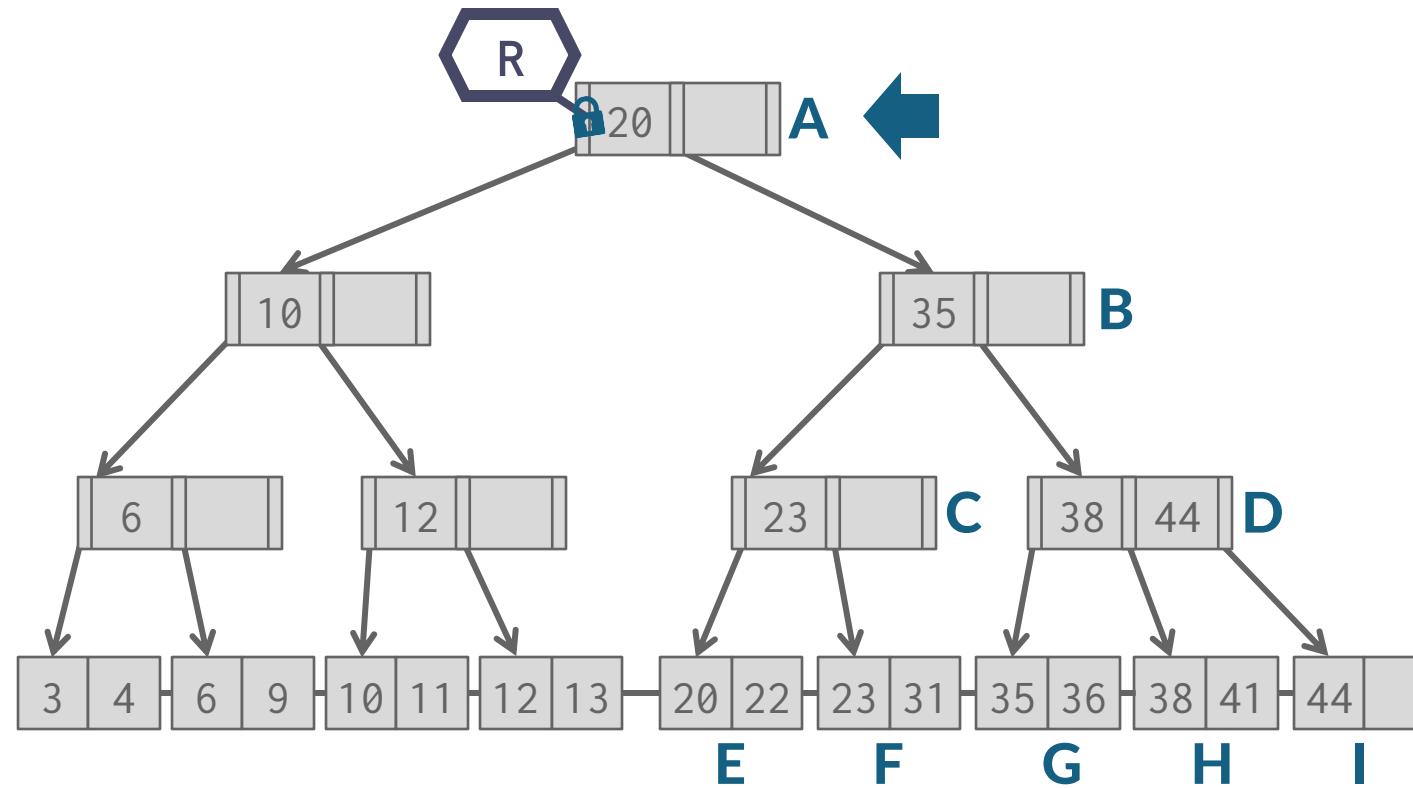
# Better Latching Algorithm

- **Search:** Same as before.
- **Insert/Delete:**
  - Set latches as if for search, get to leaf, and set **W** latch on leaf.
  - If leaf is not safe, release all latches, and restart thread using previous insert/delete protocol with write latches.
- This approach optimistically assumes that only leaf node will be modified; if not, **R** latches set on the first pass to leaf are wasteful.

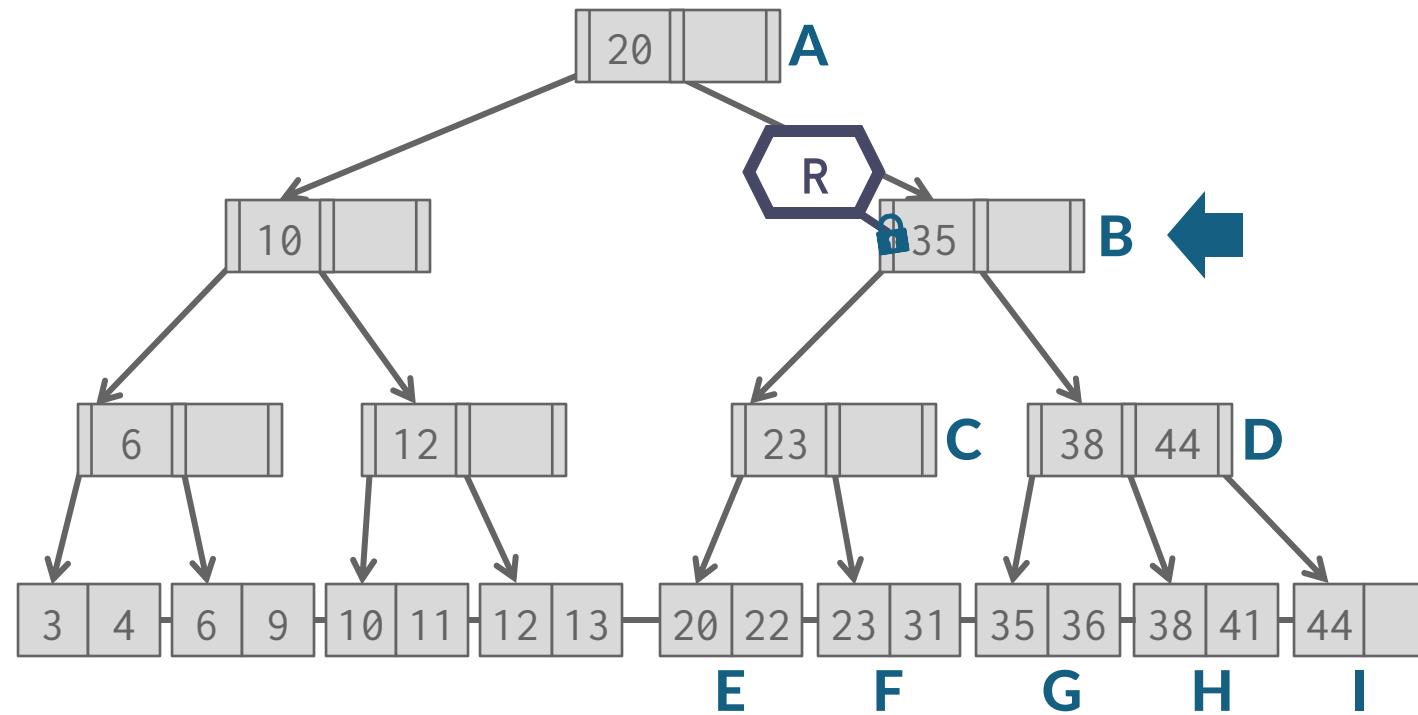
# Example #2 – Delete 38



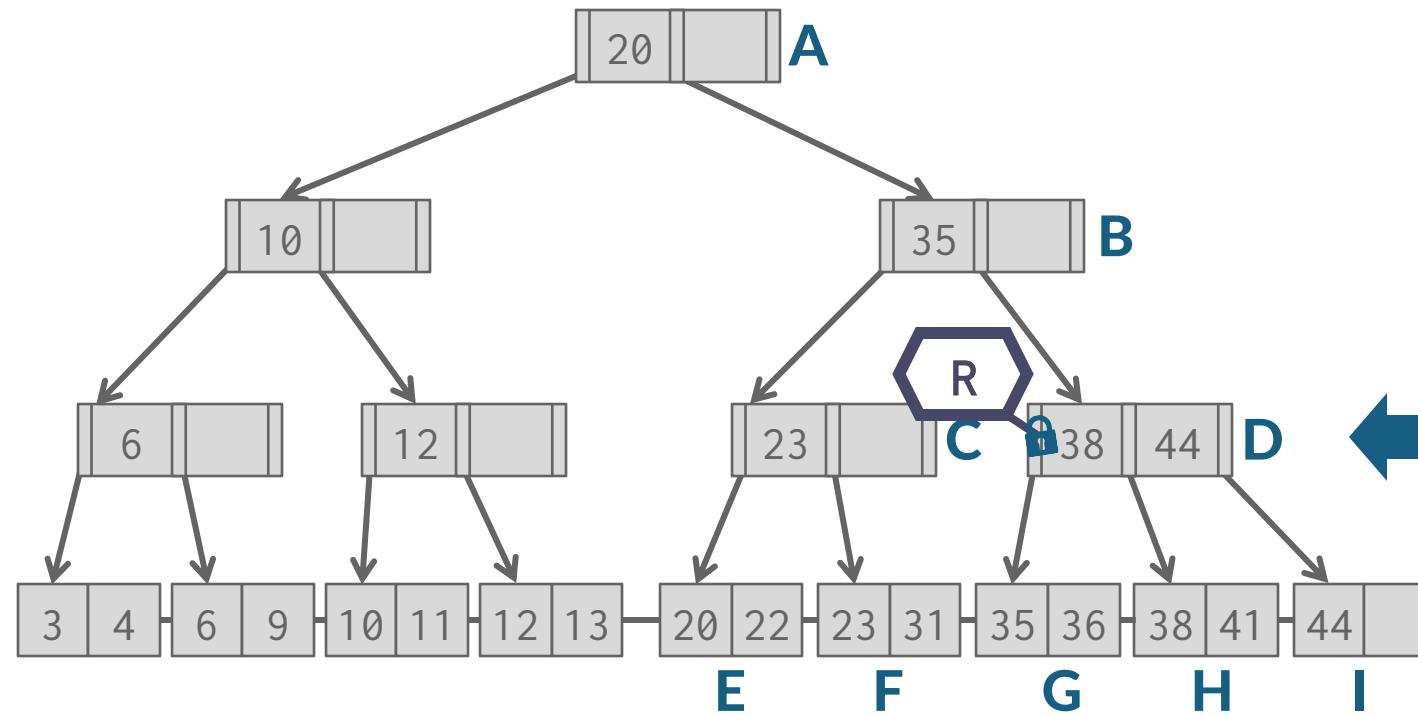
# Example #2 – Delete 38



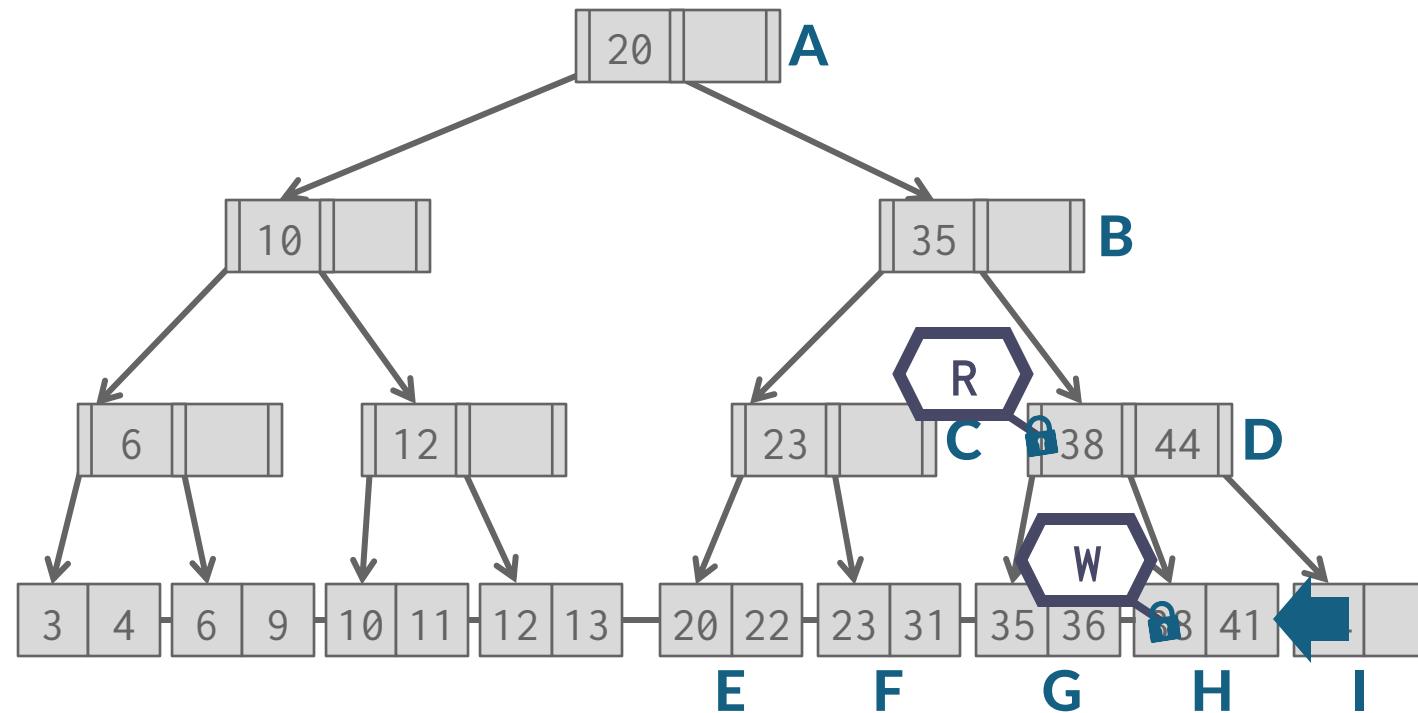
# Example #2 – Delete 38



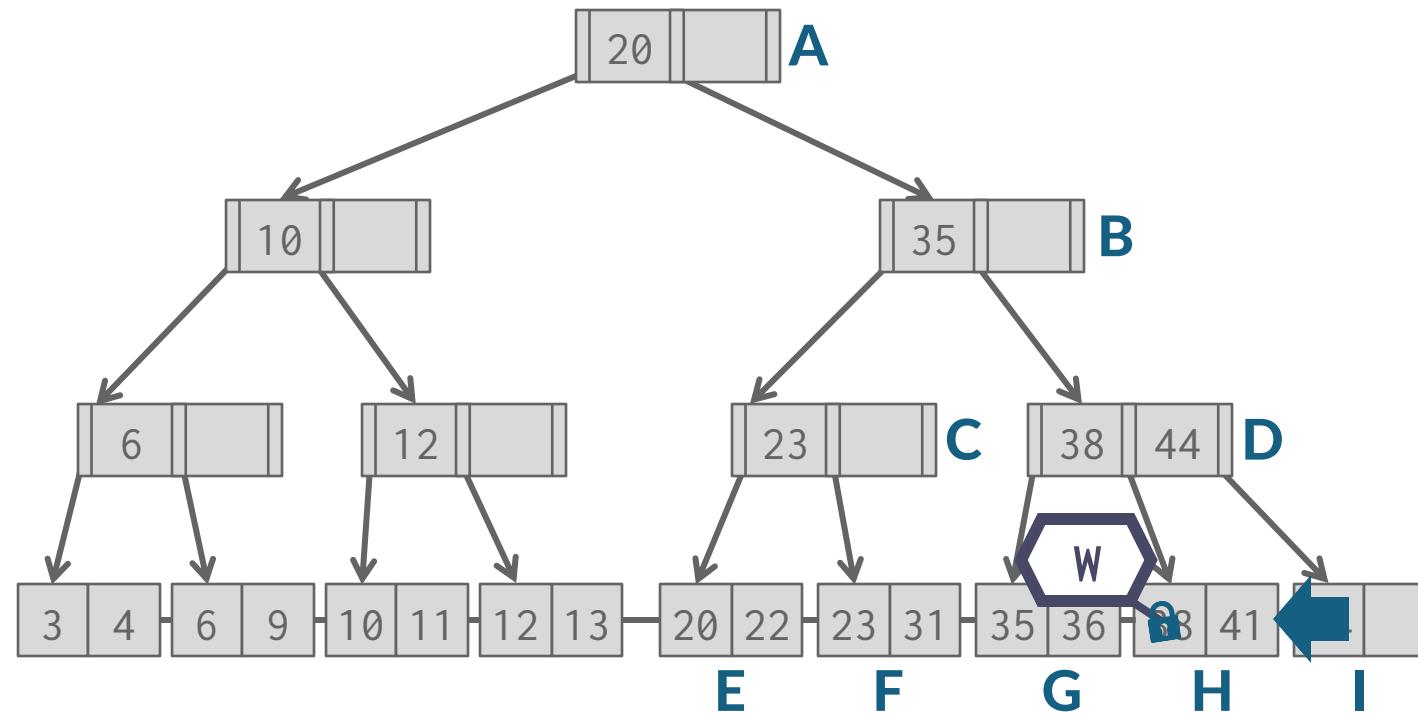
# Example #2 – Delete 38



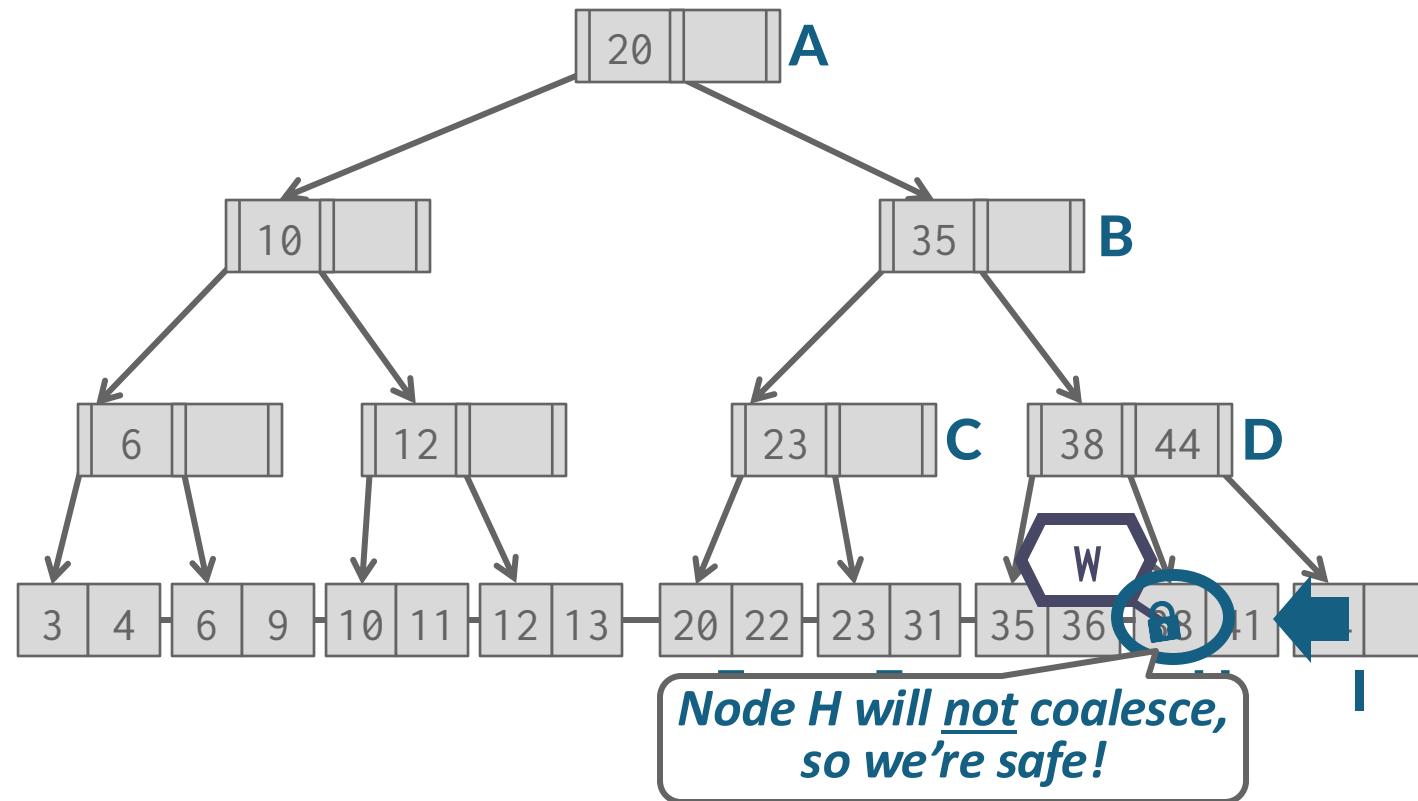
# Example #2 – Delete 38



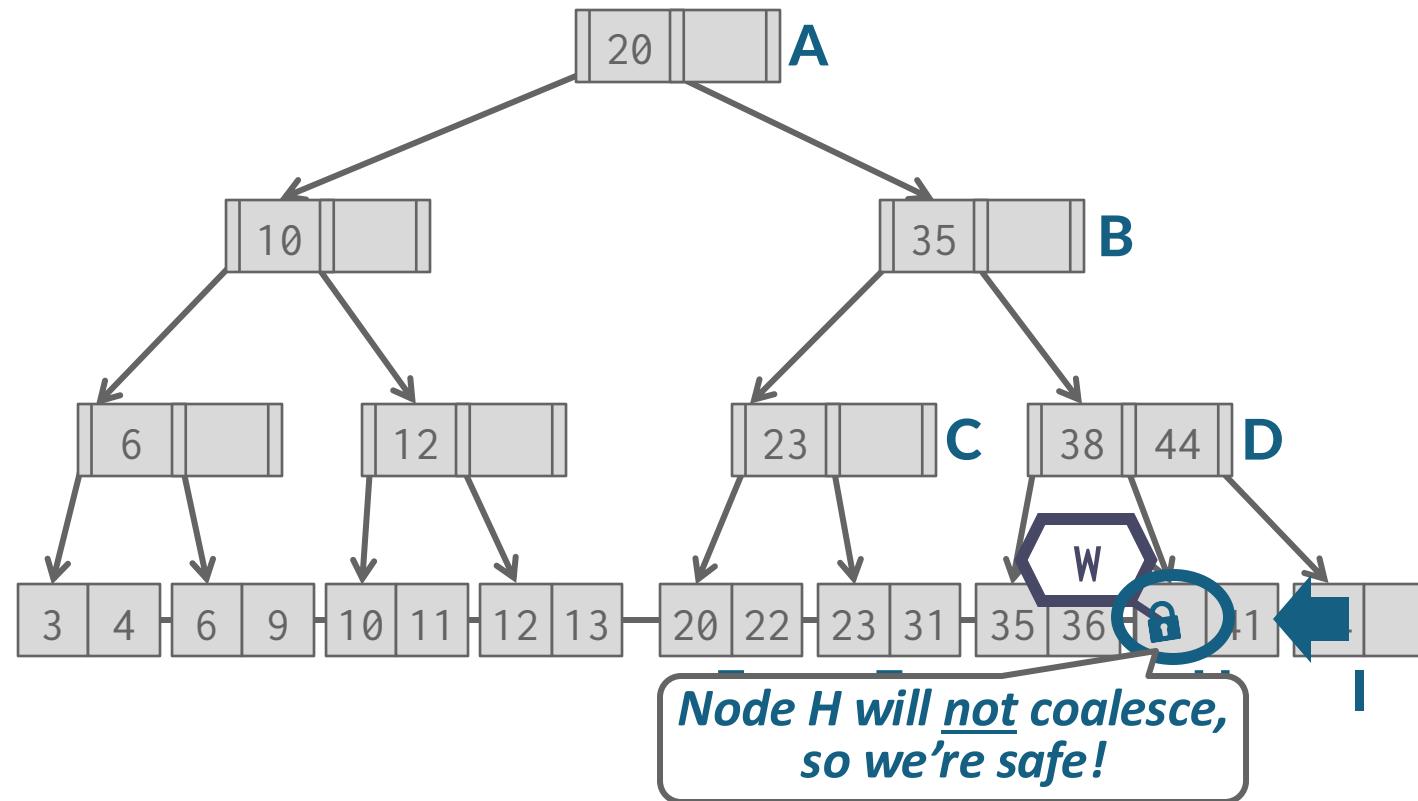
# Example #2 – Delete 38



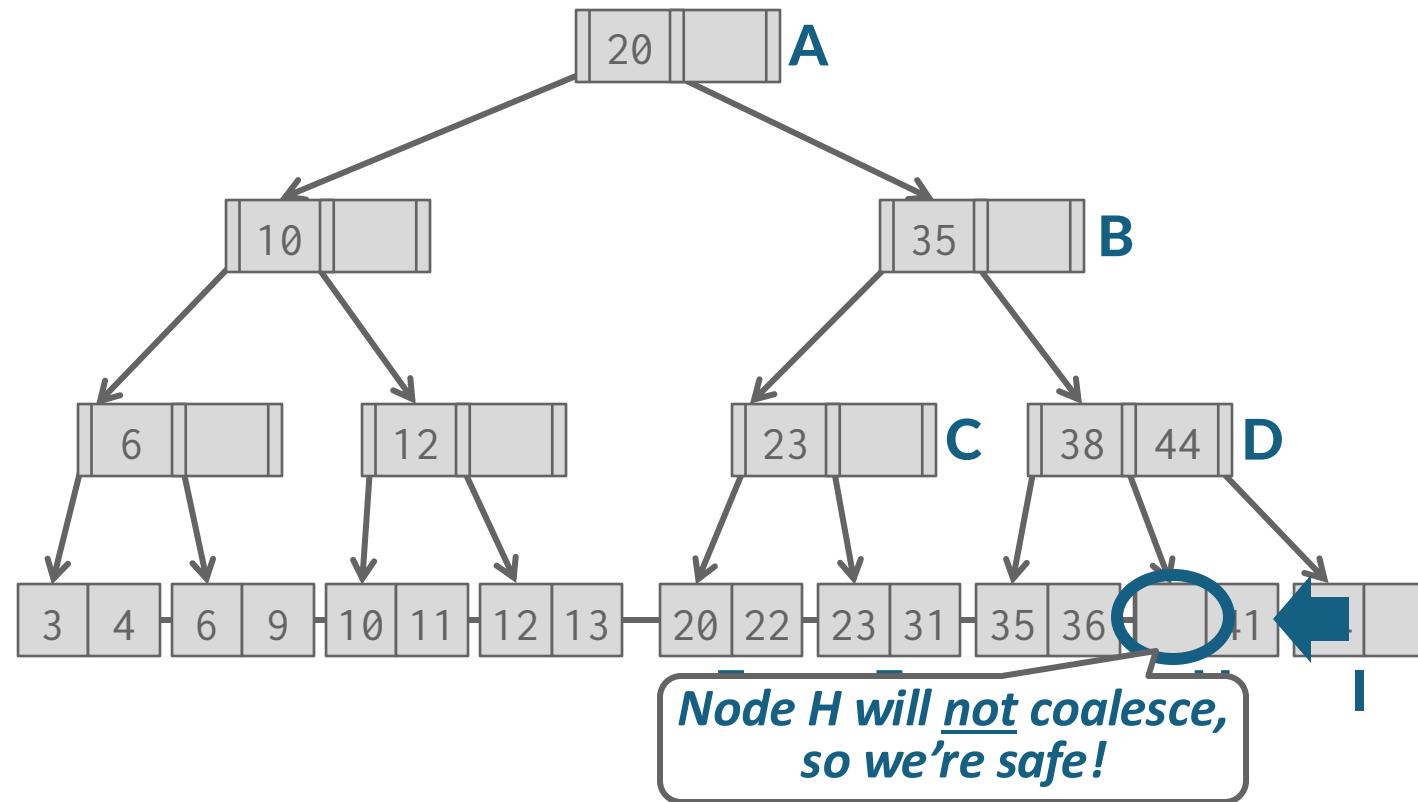
# Example #2 – Delete 38



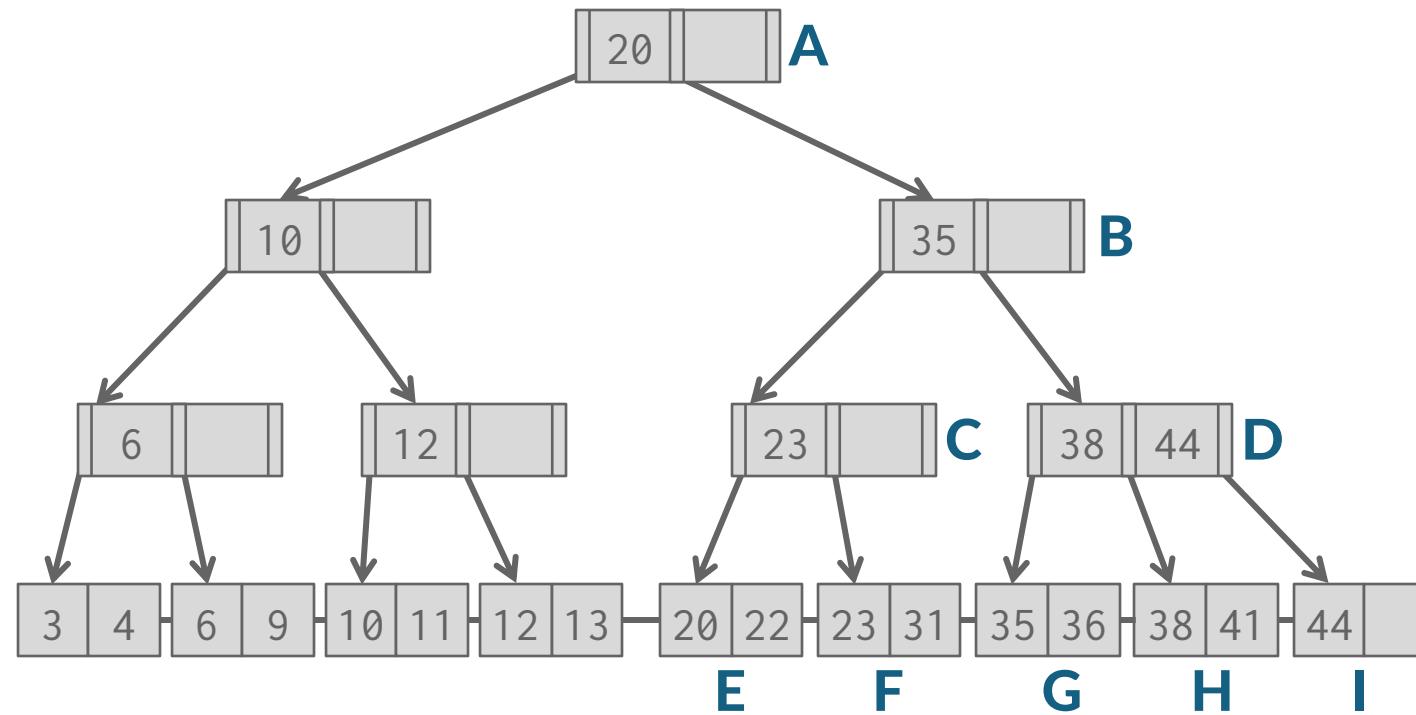
# Example #2 – Delete 38



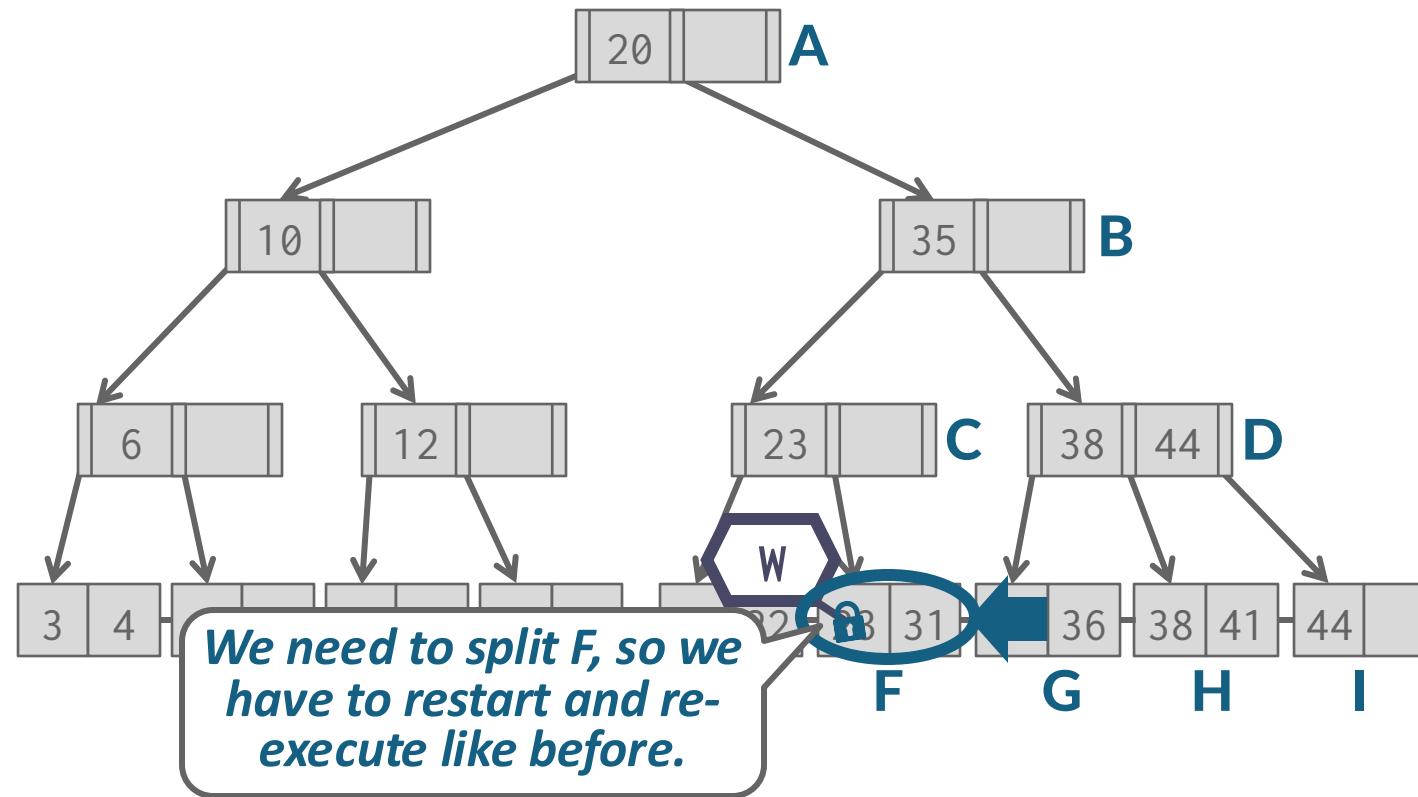
# Example #2 – Delete 38



# Example #2 – Insert 25



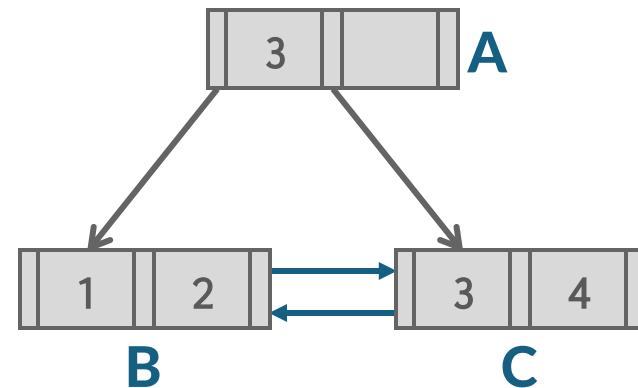
# Example #2 – Insert 25



# Observation

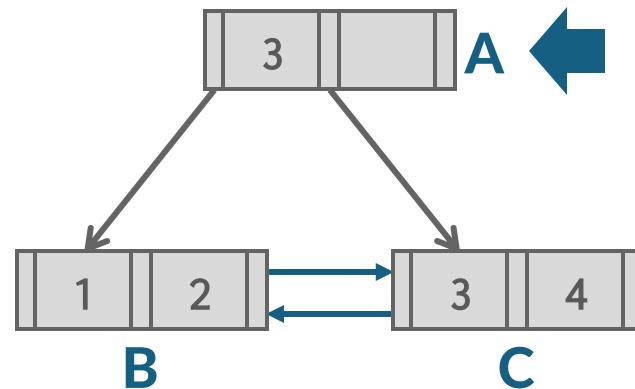
- The threads in all the examples so far have acquired latches in a “top-down” manner.
  - A thread can only acquire a latch from a node that is below its current node.
  - If the desired latch is unavailable, the thread must wait until it becomes available.
- But what if threads want to move from one leaf node to another leaf node?

# Leaf Node Scan Example #1

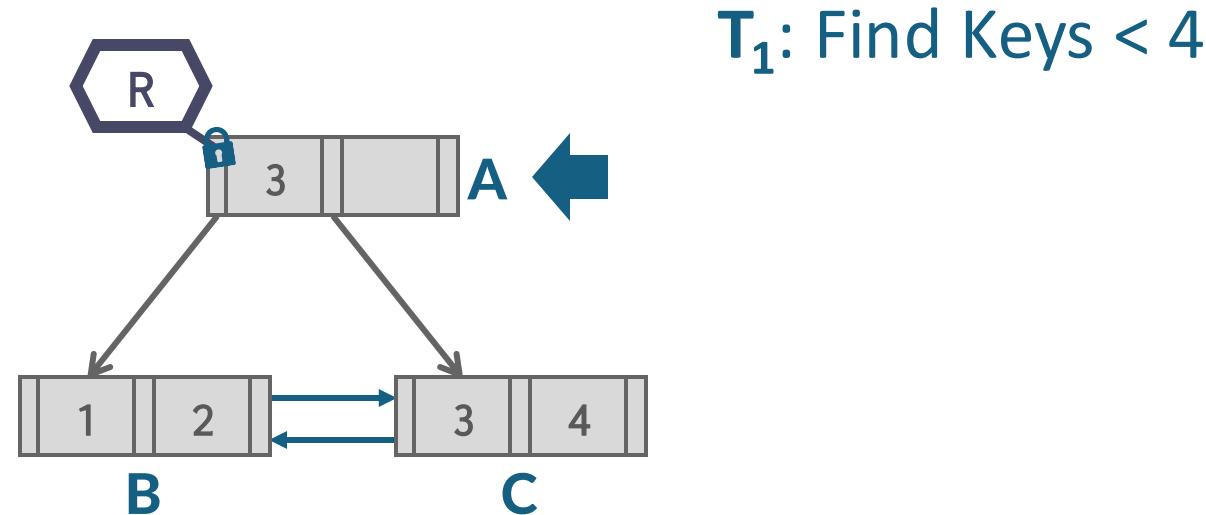


# Leaf Node Scan Example #1

$T_1$ : Find Keys < 4

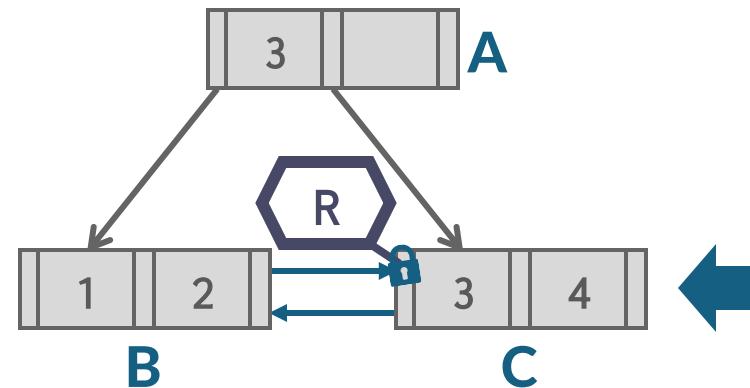


# Leaf Node Scan Example #1



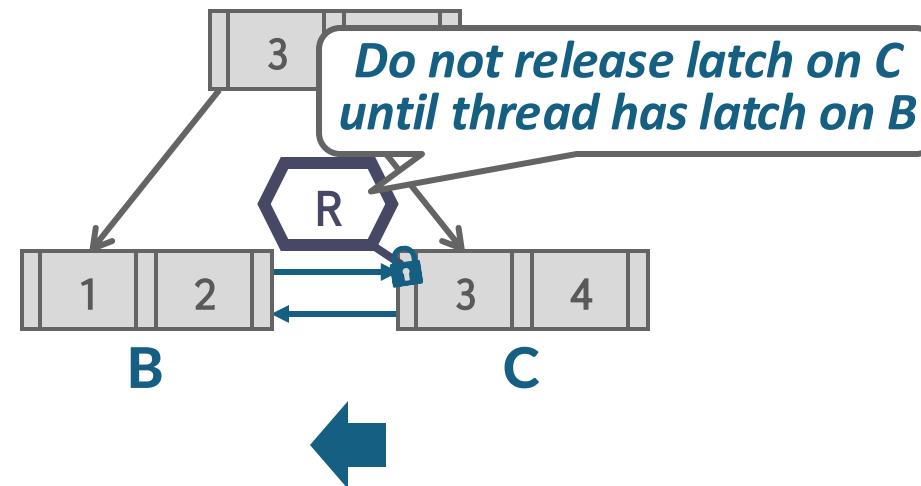
# Leaf Node Scan Example #1

$T_1$ : Find Keys < 4



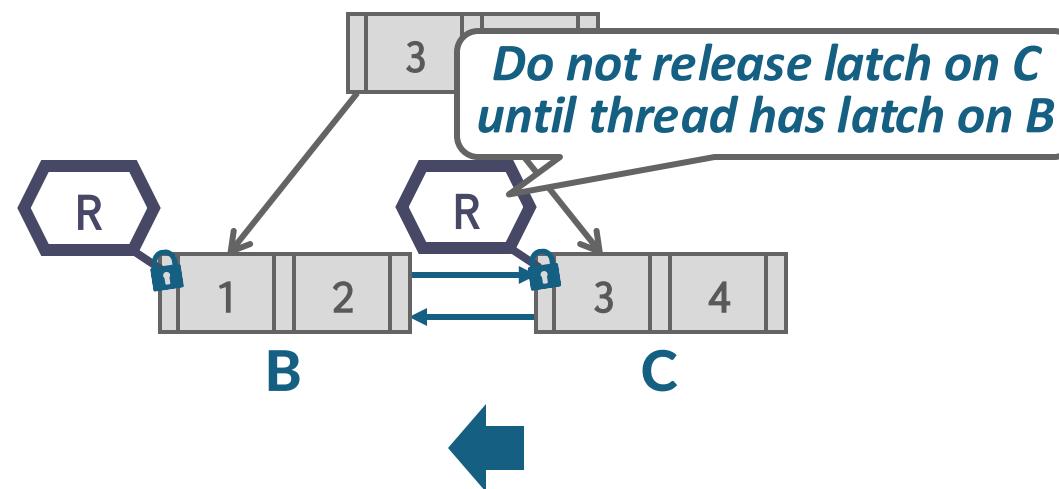
# Leaf Node Scan Example #1

$T_1$ : Find Keys < 4



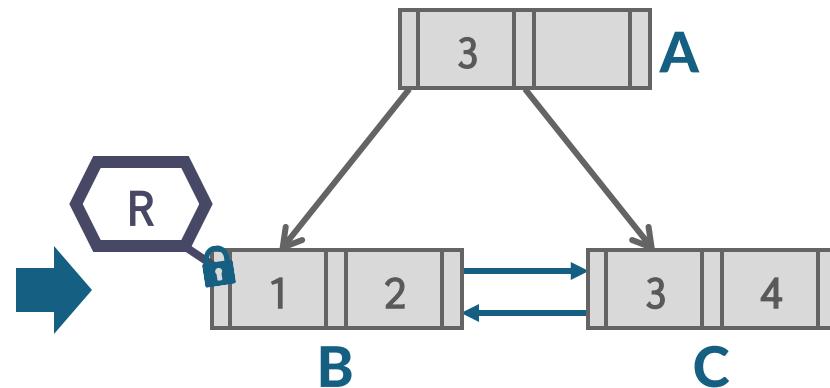
# Leaf Node Scan Example #1

$T_1$ : Find Keys < 4



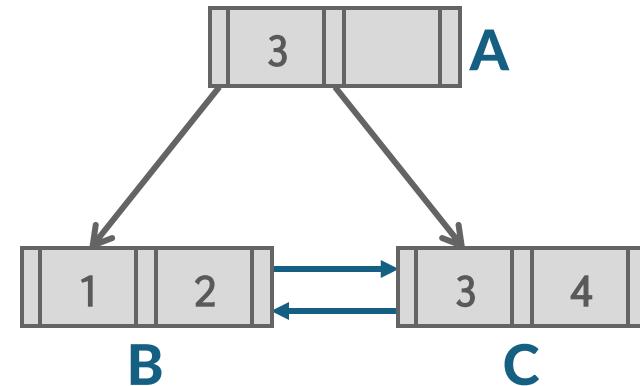
# Leaf Node Scan Example #1

$T_1$ : Find Keys < 4

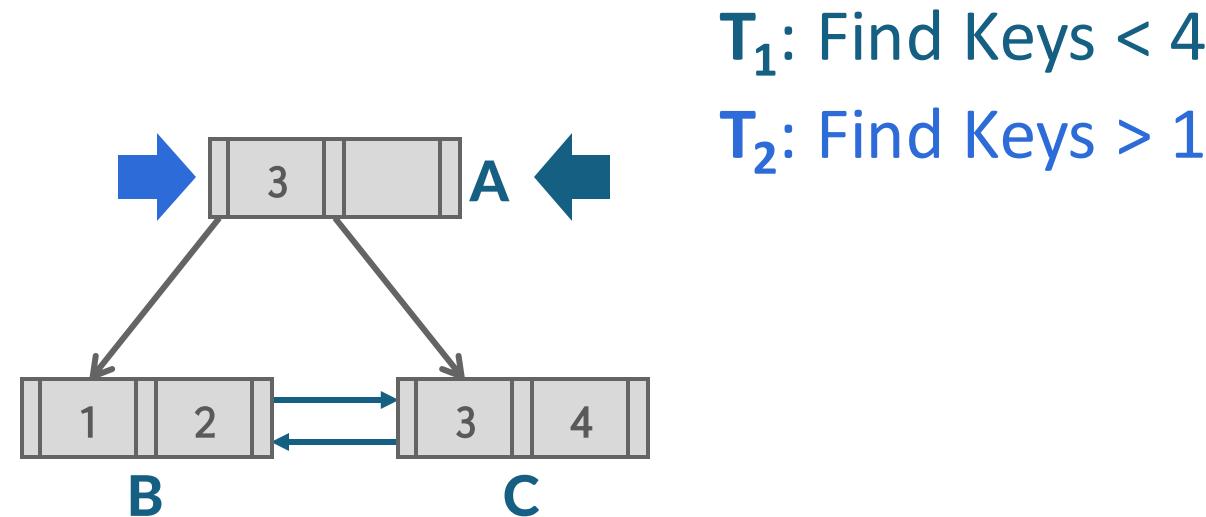


# Leaf Node Scan Example #2

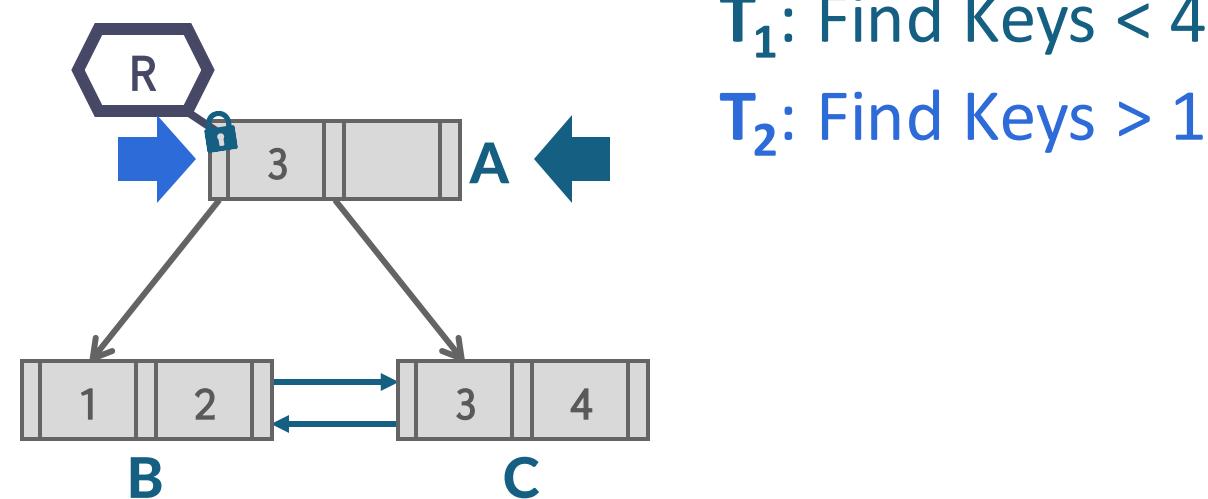
$T_1$ : Find Keys  $< 4$   
 $T_2$ : Find Keys  $> 1$



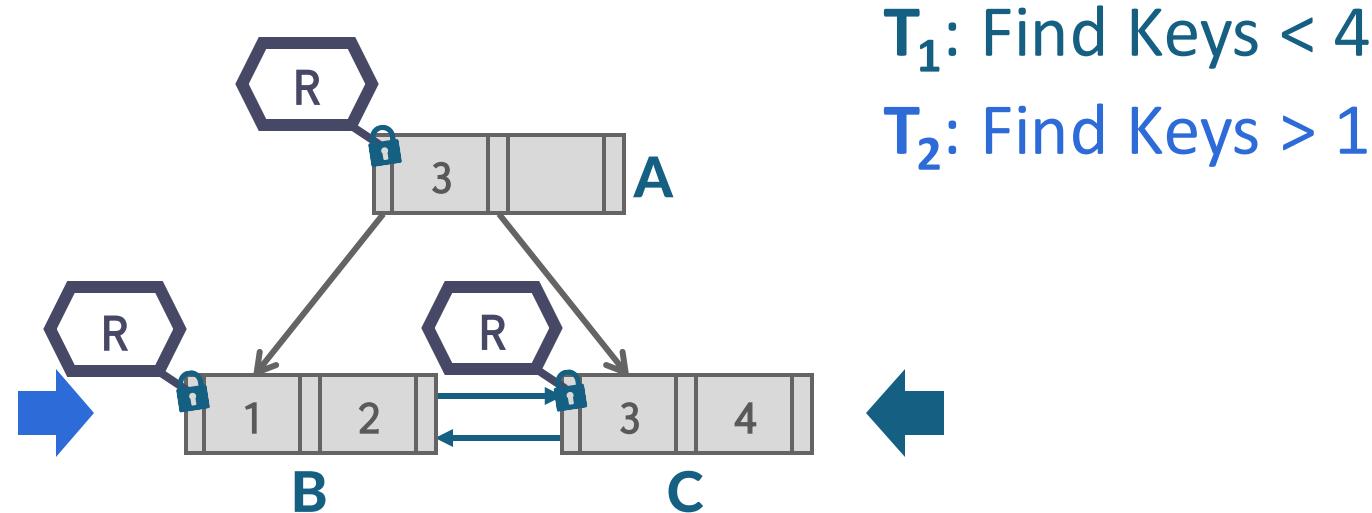
# Leaf Node Scan Example #2



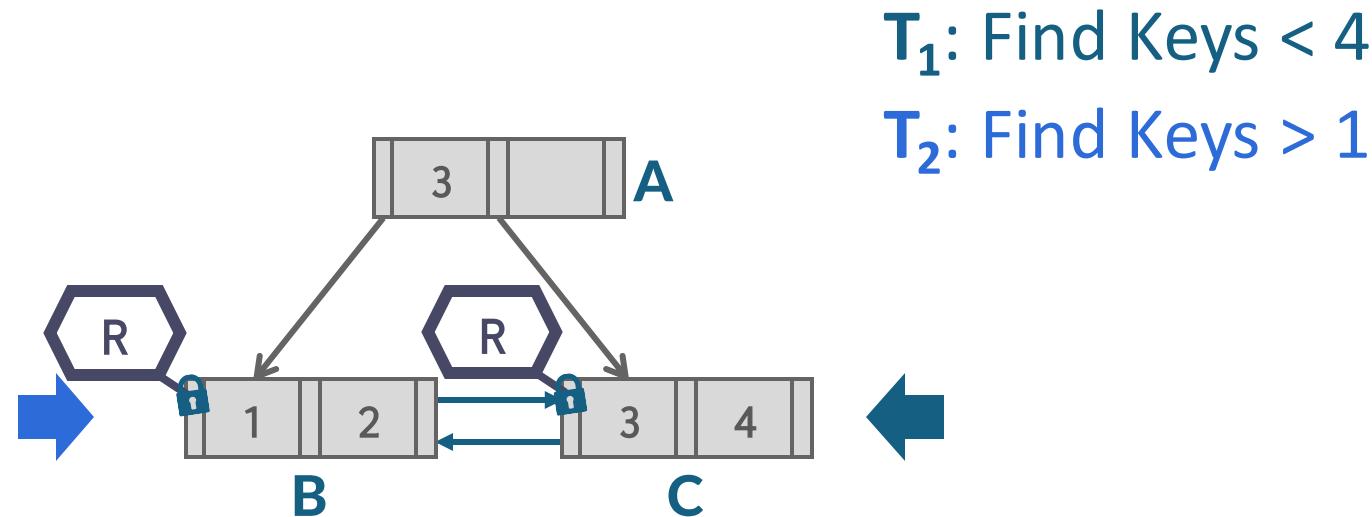
# Leaf Node Scan Example #2



# Leaf Node Scan Example #2

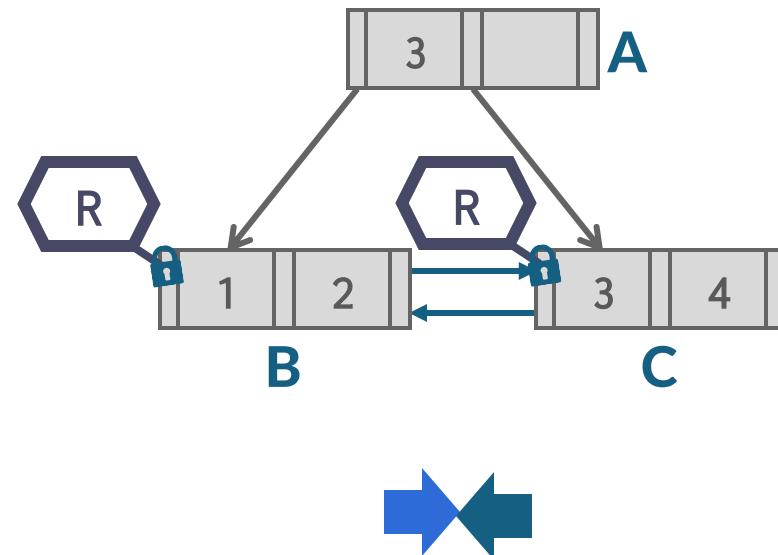


# Leaf Node Scan Example #2

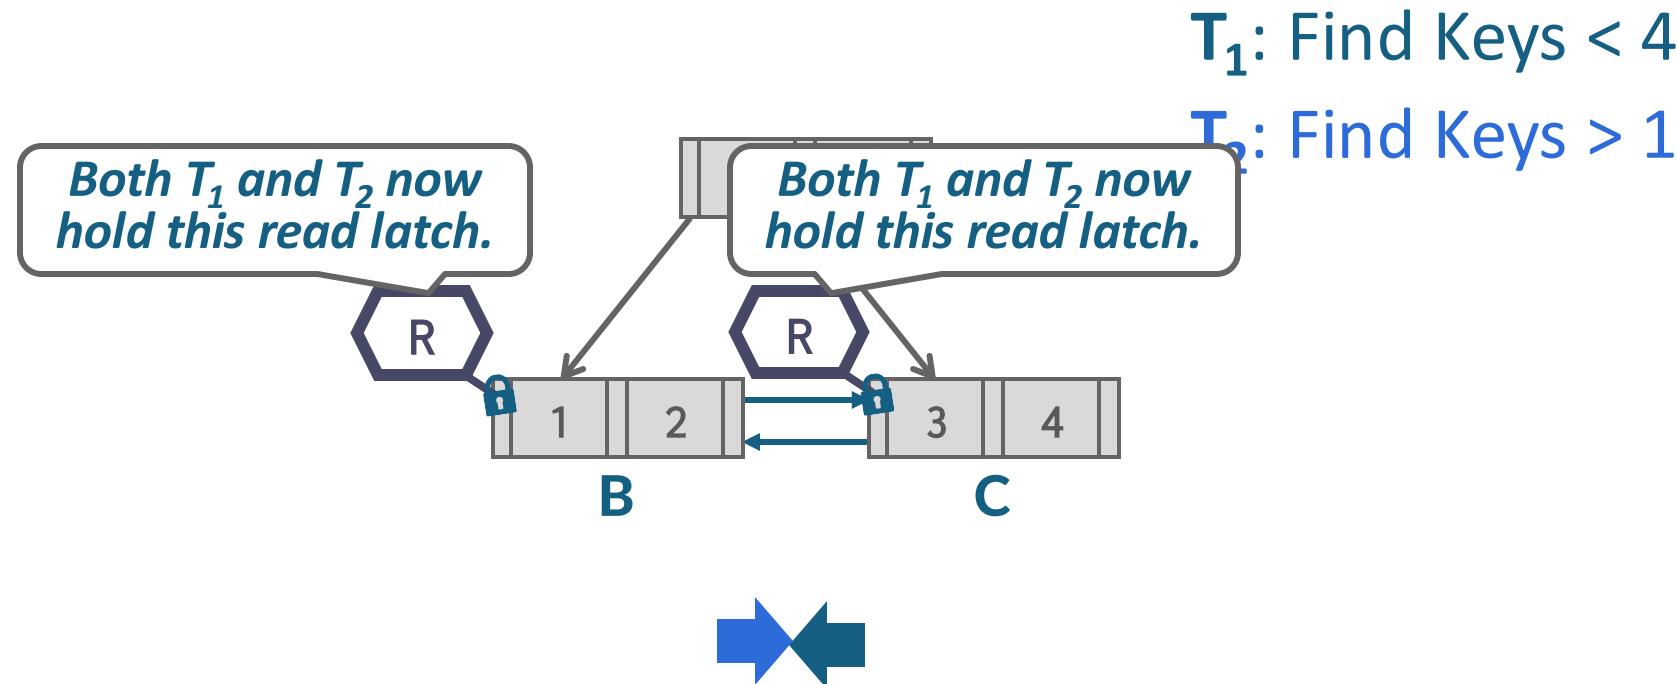


# Leaf Node Scan Example #2

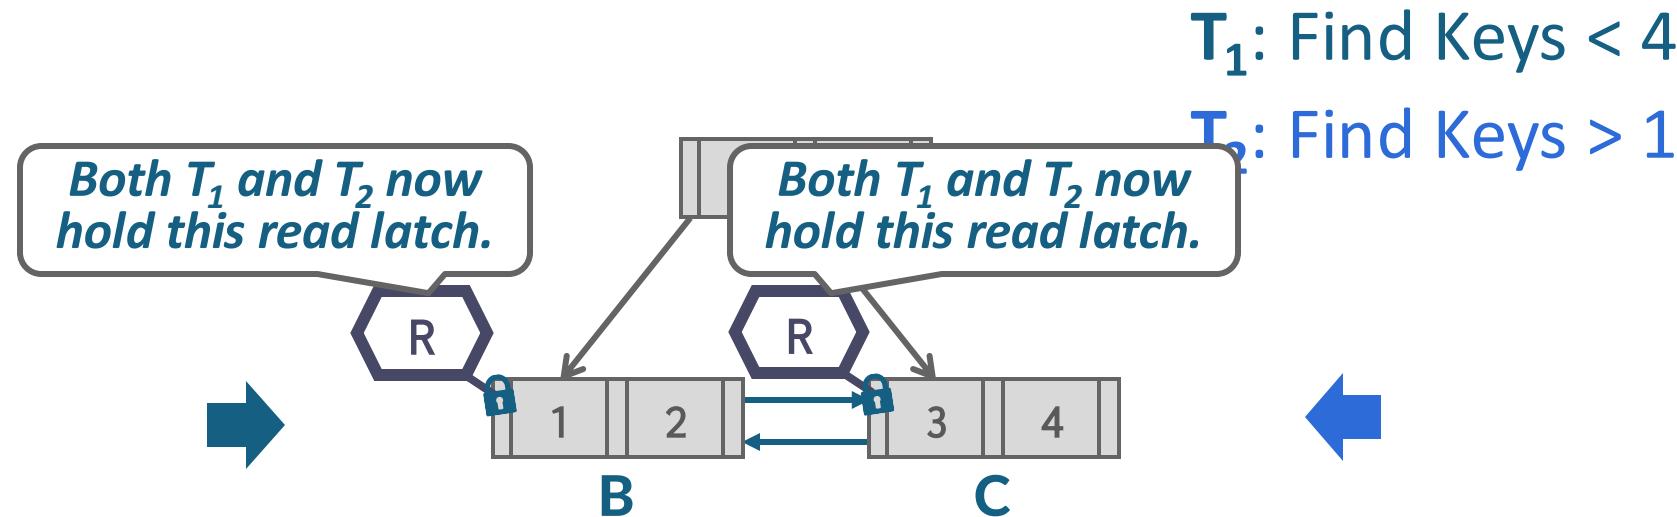
$T_1$ : Find Keys  $< 4$   
 $T_2$ : Find Keys  $> 1$



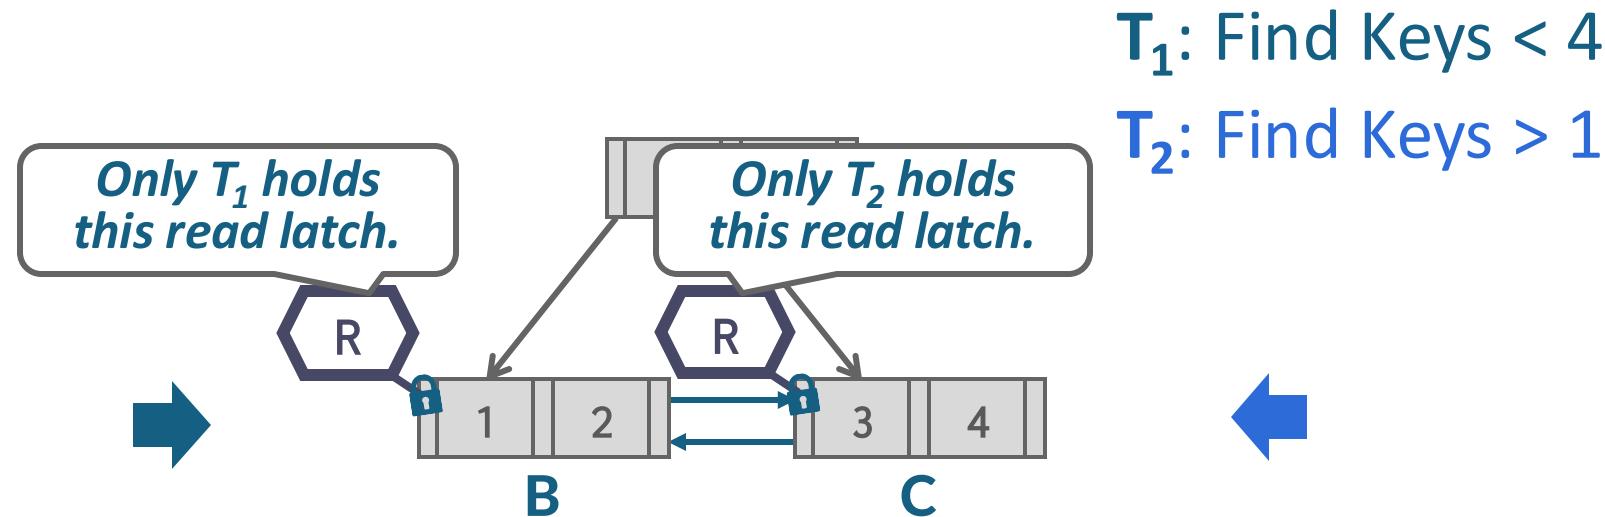
# Leaf Node Scan Example #2



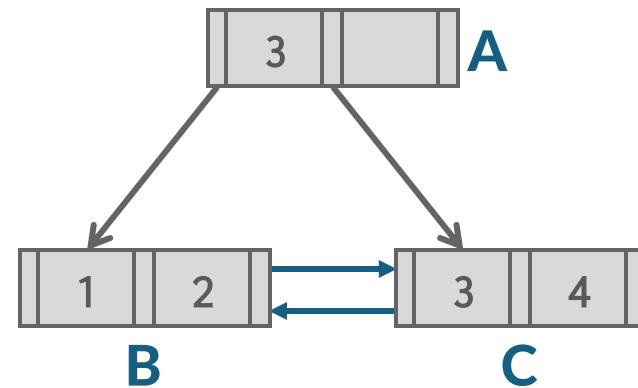
# Leaf Node Scan Example #2



# Leaf Node Scan Example #2

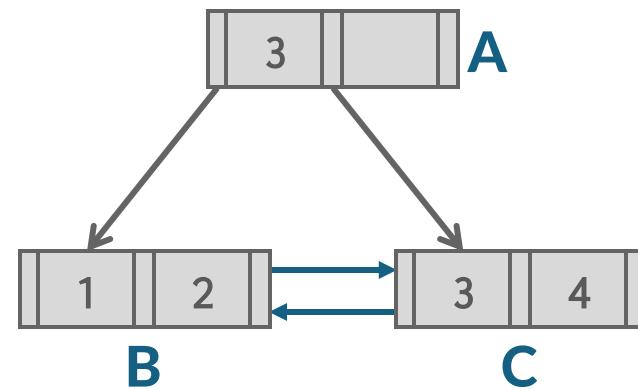


# Leaf Node Scan Example #3

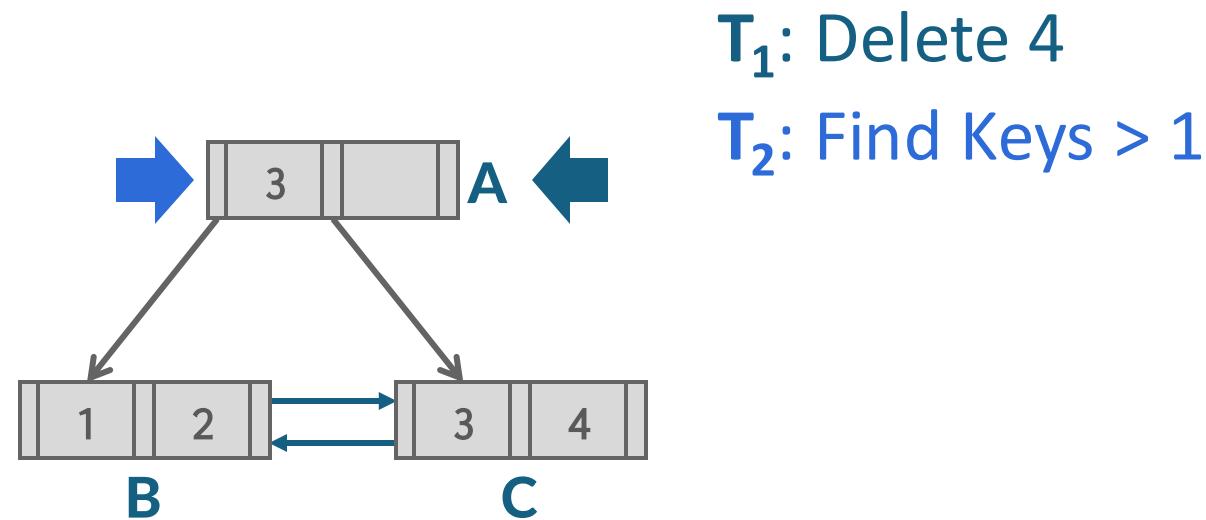


# Leaf Node Scan Example #3

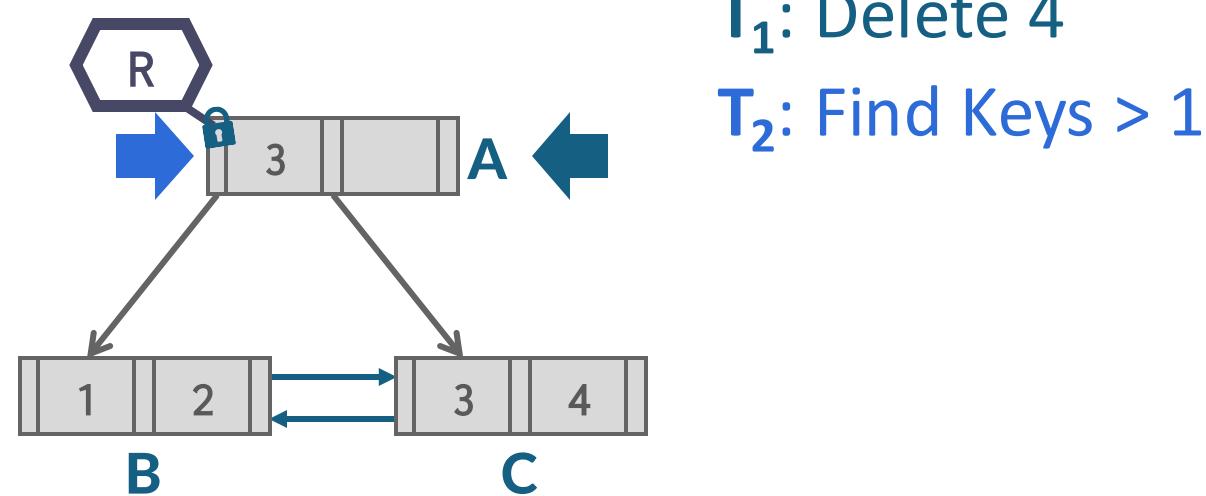
$T_1$ : Delete 4  
 $T_2$ : Find Keys > 1



# Leaf Node Scan Example #3

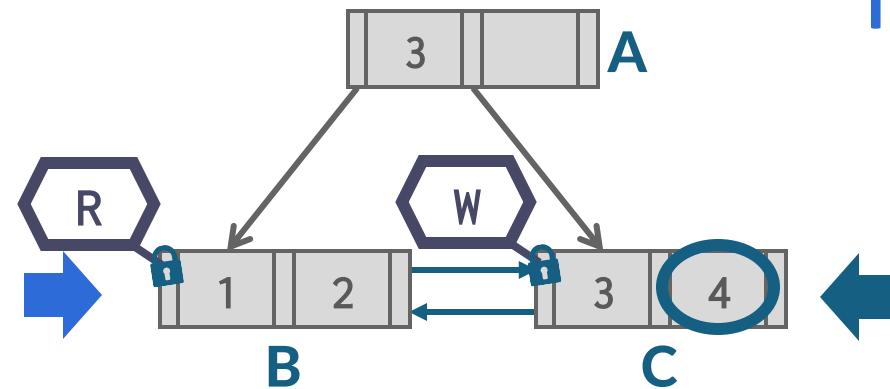


# Leaf Node Scan Example #3



# Leaf Node Scan Example #3

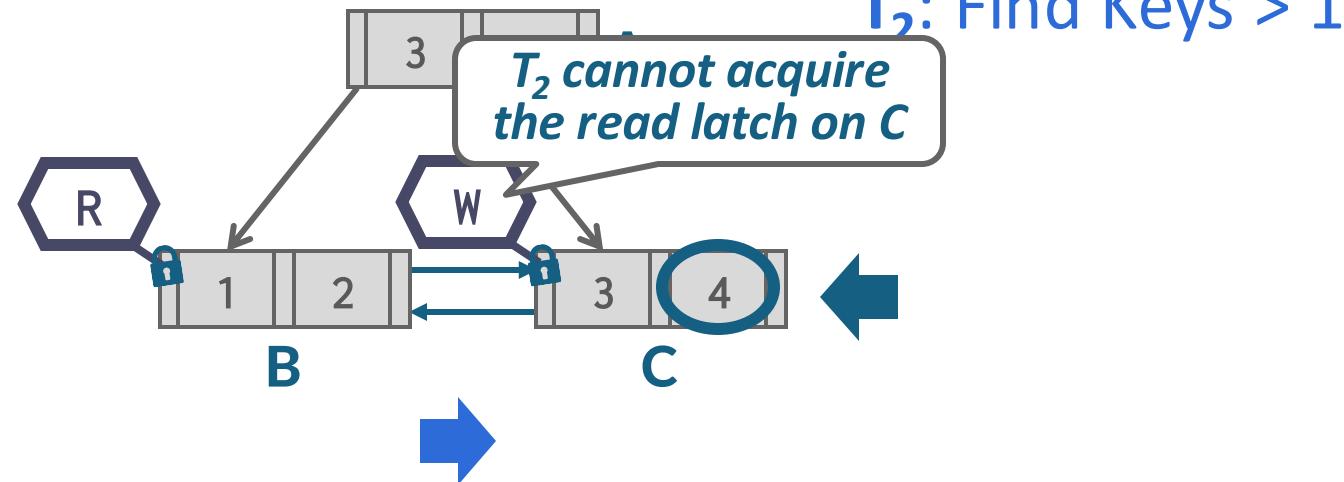
$T_1$ : Delete 4  
 $T_2$ : Find Keys > 1



# Leaf Node Scan Example #3

$T_1$ : Delete 4

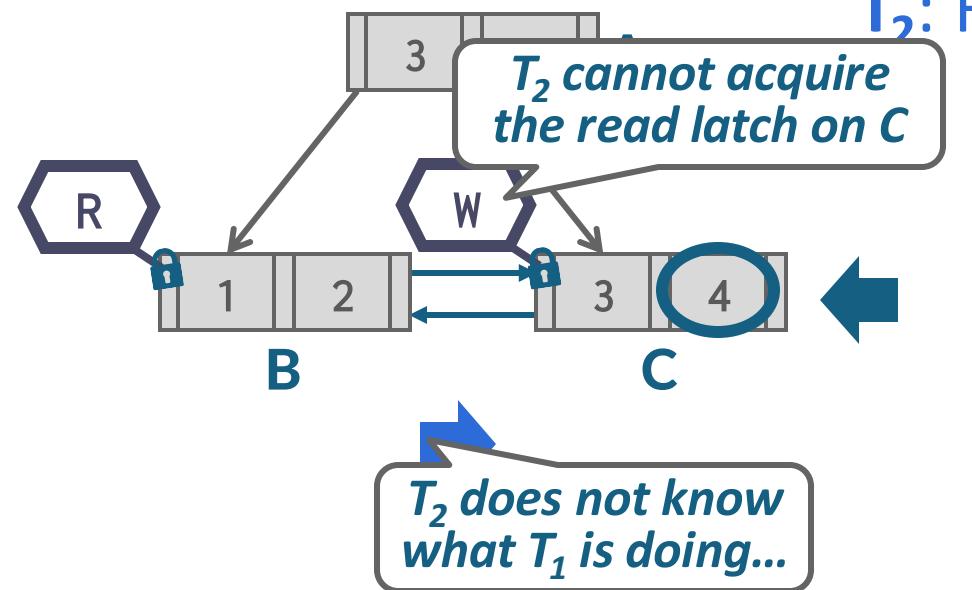
$T_2$ : Find Keys > 1



# Leaf Node Scan Example #3

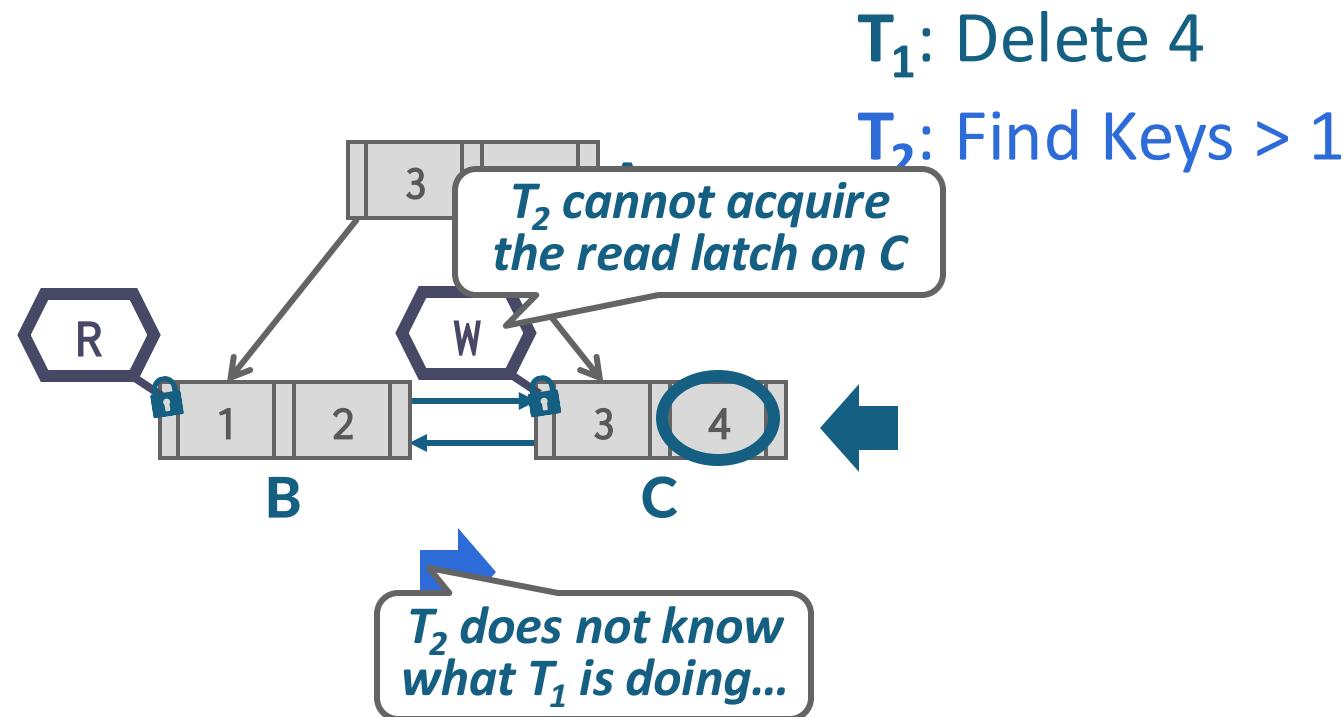
$T_1$ : Delete 4

$T_2$ : Find Keys > 1



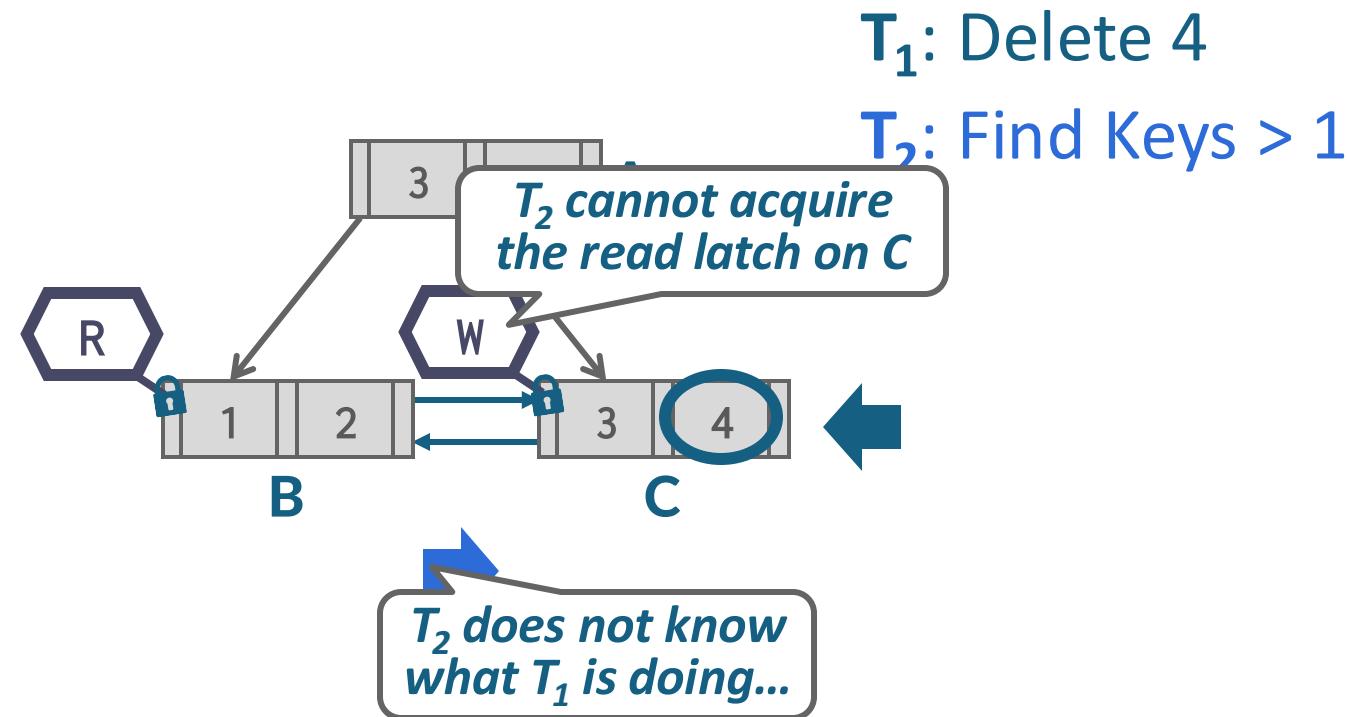
# Leaf Node Scan Example #3

T<sub>2</sub> Choices?



# Leaf Node Scan Example #3

**T<sub>2</sub> Choices?**  
⌚ Wait



# Leaf Node Scan Example #3

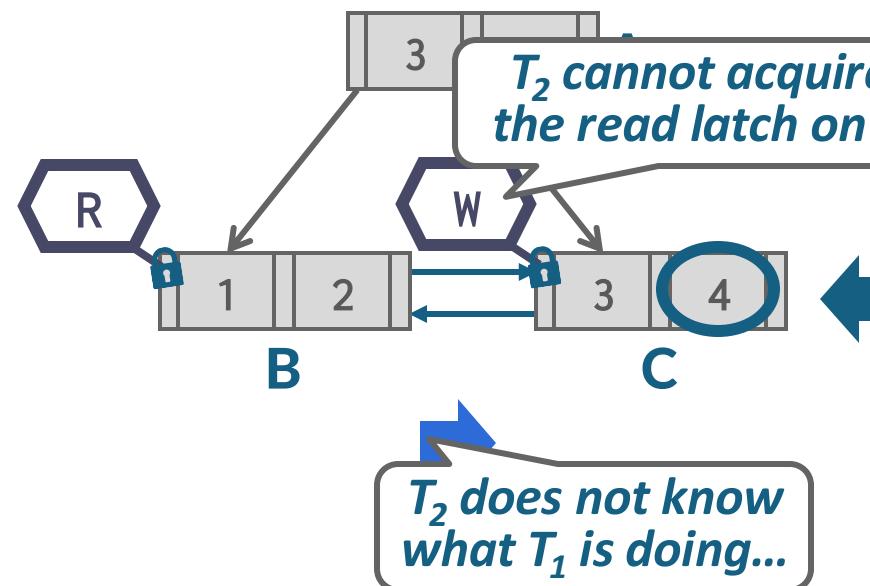
**T<sub>2</sub> Choices?**

 Wait

 Kill Ourselves

T<sub>1</sub>: Delete 4

T<sub>2</sub>: Find Keys > 1



# Leaf Node Scan Example #3

**T<sub>2</sub> Choices?**

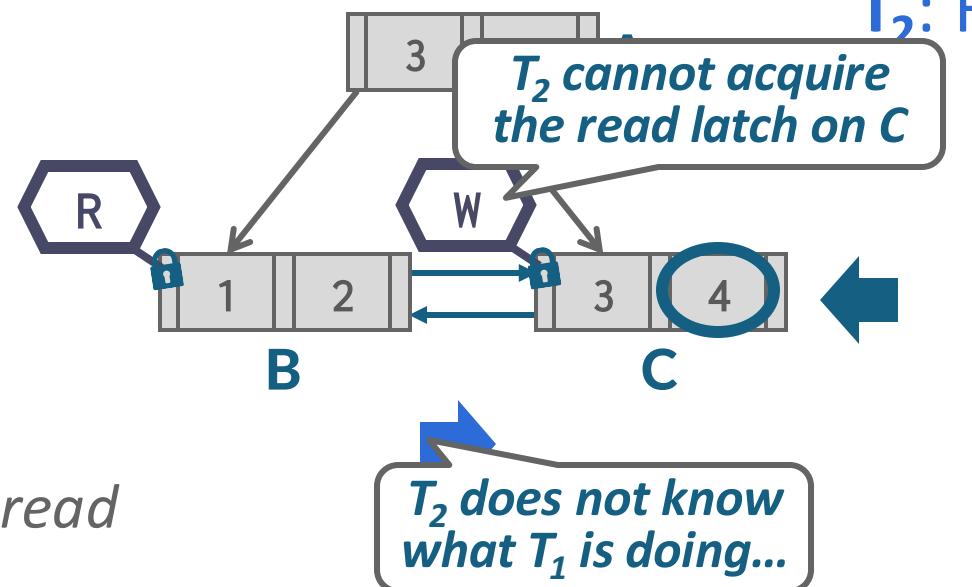
 Wait

 Kill Ourselves

 Kill Other Thread

T<sub>1</sub>: Delete 4

T<sub>2</sub>: Find Keys > 1



# Leaf Node Scan Example #3

**T<sub>2</sub> Choices?**

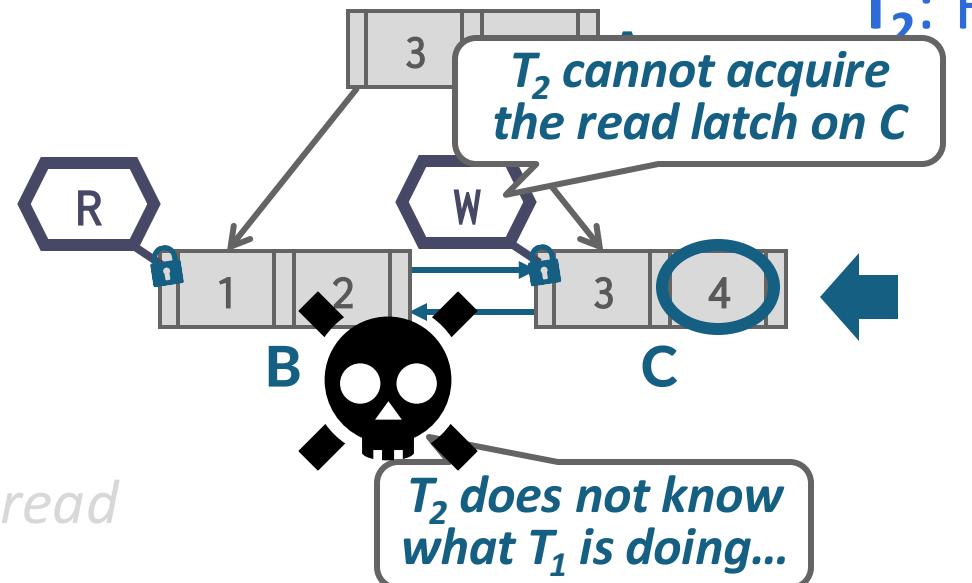
*Wait*

 *Kill Ourselves*

*Kill Other Thread*

**T<sub>1</sub>: Delete 4**

**T<sub>2</sub>: Find Keys > 1**



# Leaf Node Scans

- Latches do not support deadlock detection or avoidance. The only way we can deal with this problem is through coding discipline.
- The leaf node sibling latch acquisition protocol must support a “no-wait” mode.
- The DBMS’s data structures must cope with failed latch acquisitions.

# Conclusion

- Making a data structure thread-safe is notoriously difficult in practice.
- We focused on B+Trees, but the same high-level techniques are applicable to other data structures.

# Next Lecture

- We are finally going to discuss how to execute some queries...