



香港中文大學(深圳)
The Chinese University of Hong Kong, Shenzhen



Ack: Prof. Jignesh Patel @ CMU
Prof. Andy Pavlo @CMU

CSC3170

10: Sorting & Aggregations


Chenhao Ma

School of Data Science

The Chinese University of Hong Kong, Shenzhen


Course Status

- We are now going to talk about how to execute queries using the DBMS components we have discussed so far.
- Next several lectures:
 - Operator Algorithms
 - Query Processing Models
 - Runtime Architectures



Course Status


- We are now going to talk about how to execute queries using the DBMS components we have discussed so far.
- Next several lectures:
 - Operator Algorithms
 - Query Processing Models
 - Runtime Architectures



Query Plan

- The operators are arranged in a tree.
- Data flows from the leaves of the tree up towards the root.
 - We will discuss the granularity of the data movement later.
- The output of the root node is the result of the query.


```
SELECT R.id, S.cdate
  FROM R JOIN S
  ON R.id = S.id
 WHERE S.value > 100
```



Query Plan

- The operators are arranged in a tree.
- Data flows from the leaves of the tree up towards the root.
 - We will discuss the granularity of the data movement later.
- The output of the root node is the result of the query.

```
SELECT R.id, S.cdate
  FROM R JOIN S
  ON R.id = S.id
 WHERE S.value > 100
```



Disk-Oriented DBMS

- Just like it cannot assume that a table fits entirely in memory, a disk-oriented DBMS cannot assume that query results fit in memory.
- We will use the buffer pool to implement algorithms that need to spill to disk.
- We are also going to prefer algorithms that maximize the amount of sequential I/O.

Why Do We Need To Sort?

- Relational model/SQL is unsorted.
- Queries may request that tuples are sorted in a specific way (**ORDER BY**).
- But even if a query does not specify an order, we may still want to sort to do other things:
 - Trivial to support duplicate elimination (**DISTINCT**).
 - Bulk loading sorted tuples into a B+Tree index is faster.
 - Aggregations (**GROUP BY**).

In-Memory Sorting

- If data fits in memory, then we can use a standard sorting algorithm like Quicksort.

In-Memory Sorting

Most **database systems** use Quicksort for in-memory sorting.

In other **data platforms**, notably Python – the default sort algorithm is TimSort. It is a combination of insertion and binary merge sort. Often works well on real data.

- If data fits in memory, then we can use a standard sorting algorithm like Quicksort.

<https://www.toptal.com/developer/s/sorting-algorithms>
<https://visualgo.net/en/sorting>

In-Memory Sorting

Most **database systems** use Quicksort for in-memory sorting.

In other **data platforms**, notably Python – the default sort algorithm is TimSort. It is a combination of insertion and binary merge sort. Often works well on real data.

<https://www.toptal.com/developers/sorting-algorithms>
<https://visualgo.net/en/sorting>

- If data fits in memory, then we can use a standard sorting algorithm like Quicksort.
- If data does not fit in memory, then we need to use a technique that is aware of the cost of reading and writing disk pages ...

This Lecture

- Top-N Heap Sort
- External Merge Sort
- Aggregations

Top-N Heap Sort

- If a query contains an `ORDER BY` with a `LIMIT`, then the DBMS only needs to scan the data once to find the top-N elements.
- Ideal scenario for heapsort: if the top-N elements fit in memory.
 - Scan data once, maintain an in-memory sorted priority queue.

```
SELECT * FROM enrolled
ORDER BY sid
FETCH FIRST 4 ROWS
WITH TIES
```

Top-N Heap Sort

- If a query contains an `ORDER BY` with a `LIMIT`, then the DBMS only needs to scan the data once to find the top-N elements.
- Ideal scenario for heapsort: if the top-N elements fit in memory.
 - Scan data once, maintain an in-memory sorted priority queue.

```
SELECT * FROM enrolled
ORDER BY sid
FETCH FIRST 4 ROWS
WITH TIES
```

Top-N Heap Sort

- If a query contains an `ORDER BY` with a `LIMIT`, then the DBMS only needs to scan the data once to find the top-N elements.
- Ideal scenario for heapsort: if the top-N elements fit in memory.
 - Scan data once, maintain an in-memory sorted priority queue.

```
SELECT * FROM enrolled
ORDER BY sid
FETCH FIRST 4 ROWS
WITH TIES
```

Original Data

3	4	6	2	9	1	4	4	8
---	---	---	---	---	---	---	---	---

Top-N Heap Sort

- If a query contains an `ORDER BY` with a `LIMIT`, then the DBMS only needs to scan the data once to find the top-N elements.
- Ideal scenario for heapsort: if the top-N elements fit in memory.
 - Scan data once, maintain an in-memory sorted priority queue.

```
SELECT * FROM enrolled
ORDER BY sid
FETCH FIRST 4 ROWS
WITH TIES
```

Original Data

3	4	6	2	9	1	4	4	8
---	---	---	---	---	---	---	---	---

Sorted Heap

--	--	--	--

Top-N Heap Sort

- If a query contains an `ORDER BY` with a `LIMIT`, then the DBMS only needs to scan the data once to find the top-N elements.
- Ideal scenario for [heapsort](#): if the top-N elements fit in memory.
 - Scan data once, maintain an in-memory sorted priority queue.

```
SELECT * FROM enrolled
  ORDER BY sid
    FETCH FIRST 4 ROWS
      WITH TIES
```

Original Data

3	4	6	2	9	1	4	4	8
---	---	---	---	---	---	---	---	---



Sorted Heap


3			
---	--	--	--

Top-N Heap Sort

- If a query contains an `ORDER BY` with a `LIMIT`, then the DBMS only needs to scan the data once to find the top-N elements.
- Ideal scenario for [heapsort](#): if the top-N elements fit in memory.
 - Scan data once, maintain an in-memory sorted priority queue.

```
SELECT * FROM enrolled
  ORDER BY sid
    FETCH FIRST 4 ROWS
      WITH TIES
```

Original Data



Sorted Heap

4	3		
---	---	--	--

Top-N Heap Sort

- If a query contains an `ORDER BY` with a `LIMIT`, then the DBMS only needs to scan the data once to find the top-N elements.
- Ideal scenario for [heapsort](#): if the top-N elements fit in memory.
 - Scan data once, maintain an in-memory sorted priority queue.

```
SELECT * FROM enrolled
  ORDER BY sid
    FETCH FIRST 4 ROWS
      WITH TIES
```

Original Data

3	4	6	2	9	1	4	4	8
---	---	---	---	---	---	---	---	---



Sorted Heap

6	4	3	
---	---	---	--

Top-N Heap Sort

- If a query contains an `ORDER BY` with a `LIMIT`, then the DBMS only needs to scan the data once to find the top-N elements.
- Ideal scenario for heapsort: if the top-N elements fit in memory.
 - Scan data once, maintain an in-memory sorted priority queue.

```
SELECT * FROM enrolled
  ORDER BY sid
    FETCH FIRST 4 ROWS
      WITH TIES
```

Original Data

3	4	6	2	9	1	4	4	8
---	---	---	---	---	---	---	---	---



Sorted Heap

6	4	3	2
---	---	---	---

Top-N Heap Sort

- If a query contains an `ORDER BY` with a `LIMIT`, then the DBMS only needs to scan the data once to find the top-N elements.
- Ideal scenario for heapsort: if the top-N elements fit in memory.
 - Scan data once, maintain an in-memory sorted priority queue.

```
SELECT * FROM enrolled
  ORDER BY sid
    FETCH FIRST 4 ROWS
      WITH TIES
```

Original Data

3	4	6	2	9	1	4	4	8
---	---	---	---	---	---	---	---	---



Sorted Heap

6	4	3	2
---	---	---	---

Top-N Heap Sort

- If a query contains an `ORDER BY` with a `LIMIT`, then the DBMS only needs to scan the data once to find the top-N elements.
- Ideal scenario for heapsort: if the top-N elements fit in memory.
 - Scan data once, maintain an in-memory sorted priority queue.

```
SELECT * FROM enrolled
  ORDER BY sid
    FETCH FIRST 4 ROWS
      WITH TIES
```

Original Data

3	4	6	2	9	1	4	4	8
---	---	---	---	---	---	---	---	---



Skip!

Sorted Heap

6	4	3	2
---	---	---	---

Top-N Heap Sort

- If a query contains an `ORDER BY` with a `LIMIT`, then the DBMS only needs to scan the data once to find the top-N elements.
- Ideal scenario for heapsort: if the top-N elements fit in memory.
 - Scan data once, maintain an in-memory sorted priority queue.

```
SELECT * FROM enrolled
  ORDER BY sid
    FETCH FIRST 4 ROWS
      WITH TIES
```

Original Data

3	4	6	2	9	1	4	4	8
---	---	---	---	---	---	---	---	---



Sorted Heap

4	3	2	1
---	---	---	---

Top-N Heap Sort

- If a query contains an `ORDER BY` with a `LIMIT`, then the DBMS only needs to scan the data once to find the top-N elements.
- Ideal scenario for [heapsort](#): if the top-N elements fit in memory.
 - Scan data once, maintain an in-memory sorted priority queue.

```
SELECT * FROM enrolled
  ORDER BY sid
    FETCH FIRST 4 ROWS
      WITH TIES
```

Original Data

3	4	6	2	9	1	4	4	8
---	---	---	---	---	---	---	---	---



Sorted Heap

4	4	3	2	1				
---	---	---	---	---	--	--	--	--

Top-N Heap Sort

- If a query contains an `ORDER BY` with a `LIMIT`, then the DBMS only needs to scan the data once to find the top-N elements.
- Ideal scenario for heapsort: if the top-N elements fit in memory.
 - Scan data once, maintain an in-memory sorted priority queue.

```
SELECT * FROM enrolled
  ORDER BY sid
    FETCH FIRST 4 ROWS
      WITH TIES
```

Original Data

3	4	6	2	9	1	4	4	8
---	---	---	---	---	---	---	---	---



Sorted Heap

4	4	4	3	2	1		
---	---	---	---	---	---	--	--

Top-N Heap Sort

- If a query contains an `ORDER BY` with a `LIMIT`, then the DBMS only needs to scan the data once to find the top-N elements.
- Ideal scenario for heapsort: if the top-N elements fit in memory.
 - Scan data once, maintain an in-memory sorted priority queue.

```
SELECT * FROM enrolled
  ORDER BY sid
    FETCH FIRST 4 ROWS
      WITH TIES
```

Original Data

3	4	6	2	9	1	4	4	8
---	---	---	---	---	---	---	---	---



Sorted Heap

4	4	4	3	2	1		
---	---	---	---	---	---	--	--

Top-N Heap Sort

- If a query contains an `ORDER BY` with a `LIMIT`, then the DBMS only needs to scan the data once to find the top-N elements.
- Ideal scenario for [heapsort](#): if the top-N elements fit in memory.
 - Scan data once, maintain an in-memory sorted priority queue.

```
SELECT * FROM enrolled
  ORDER BY sid
    FETCH FIRST 4 ROWS
      WITH TIES
```

Original Data

3	4	6	2	9	1	4	4	8
---	---	---	---	---	---	---	---	---

Skip and done!



Sorted Heap

4	4	4	3	2	1		
---	---	---	---	---	---	--	--

Top-N Heap Sort


- If a query contains an `ORDER BY` with a `LIMIT`, then the DBMS only needs to scan the data once to find the top-N elements.
- Ideal scenario for [heapsort](#): if the top-N elements fit in memory.
 - Scan data once, maintain an in-memory sorted priority queue.

```
SELECT * FROM enrolled
  ORDER BY sid
    FETCH FIRST 4 ROWS
      WITH TIES
```

Original Data

3	4	6	2	9	1	4	4	8
---	---	---	---	---	---	---	---	---

Sorted Heap



External Merge Sort

External Merge Sort

- Divide-and-conquer algorithm that splits data into separate runs, sorts them individually, and then combines them into longer sorted runs.
- **Phase #1 – Sorting**
 - Sort chunks of data that fit in memory and then write back the sorted chunks to a file on disk.
- **Phase #2 – Merging**
 - Combine sorted runs into larger chunks.

Sorted Run

- A run is a list of key/value pairs.
- **Key:** The attribute(s) to compare to compute the sort order.
- **Value:** Two choices
 - Tuple (*early materialization*).
 - Record ID (*late materialization*).

Sorted Run

- A run is a list of key/value pairs.
- **Key:** The attribute(s) to compare to compute the sort order.
- **Value:** Two choices
 - Tuple (*early materialization*).
 - Record ID (*late materialization*).

Early Materialization

K1	< <i>Tuple Data</i> >
K2	< <i>Tuple Data</i> >
⋮	

Sorted Run

- A run is a list of key/value pairs.
- **Key:** The attribute(s) to compare to compute the sort order.
- **Value:** Two choices
 - Tuple (*early materialization*).
 - Record ID (*late materialization*).

Early Materialization

K1	<Tuple Data>
K2	<Tuple Data>
⋮	

Late Materialization




Sorted Run

- A run is a list of key/value pairs.
- **Key:** The attribute(s) to compare to compute the sort order.
- **Value:** Two choices
 - Tuple (*early materialization*).
 - Record ID (*late materialization*).

Early Materialization

K1	<Tuple Data>
K2	<Tuple Data>
⋮	

Late Materialization



2-Way External Merge Sort

- We will start with a simple example of a 2-way external merge sort.
 - “2” is the number of runs that we are going to merge into a new run for each pass.
- Data is broken up into N pages.
- The DBMS has a finite number of B buffer pool pages to hold input and output data.

Simplified 2-Way External Merge Sort

- **Pass #0**
 - Read one page of the table into memory
 - Sort page into a “run” and write it back to disk
 - Repeat until the whole table has been sorted into runs
- **Pass #1,2,3,...**
 - Recursively merge pairs of runs into runs twice as long
 - Need at least 3 buffer pages (2 for input, 1 for output)

Simplified 2-Way External Merge Sort

- In each pass, we read and write every page in the file.
- Number of passes
 $= 1 + \lceil \log_2 N \rceil$
- Total I/O cost
 $= 2N \cdot (\# \text{ of passes})$

Simplified 2-Way External Merge Sort

3,4	6,2	9,4	8,7	5,6	3,1	2	∅
-----	-----	-----	-----	-----	-----	---	---

- In each pass, we read and write every page in the file.
- Number of passes
 $= 1 + \lceil \log_2 N \rceil$
- Total I/O cost
 $= 2N \cdot (\# \text{ of passes})$


Simplified 2-Way External Merge Sort



- In each pass, we read and write every page in the file.
- Number of passes
 $= 1 + \lceil \log_2 N \rceil$
- Total I/O cost
 $= 2N \cdot (\# \text{ of passes})$


Simplified 2-Way External Merge Sort

- In each pass, we read and write every page in the file.
- Number of passes
 $= 1 + \lceil \log_2 N \rceil$
- Total I/O cost
 $= 2N \cdot (\# \text{ of passes})$




Simplified 2-Way External Merge Sort

- In each pass, we read and write every page in the file.
- Number of passes
 $= 1 + \lceil \log_2 N \rceil$
- Total I/O cost
 $= 2N \cdot (\# \text{ of passes})$




Simplified 2-Way External Merge Sort

- In each pass, we read and write every page in the file.
- Number of passes
 $= 1 + \lceil \log_2 N \rceil$
- Total I/O cost
 $= 2N \cdot (\# \text{ of passes})$




Simplified 2-Way External Merge Sort

- In each pass, we read and write every page in the file.
- Number of passes
 $= 1 + \lceil \log_2 N \rceil$
- Total I/O cost
 $= 2N \cdot (\# \text{ of passes})$




Simplified 2-Way External Merge Sort

- In each pass, we read and write every page in the file.
- Number of passes
 $= 1 + \lceil \log_2 N \rceil$
- Total I/O cost
 $= 2N \cdot (\# \text{ of passes})$




Simplified 2-Way External Merge Sort

- In each pass, we read and write every page in the file.
- Number of passes
 $= 1 + \lceil \log_2 N \rceil$
- Total I/O cost
 $= 2N \cdot (\# \text{ of passes})$




Simplified 2-Way External Merge Sort

- In each pass, we read and write every page in the file.
- Number of passes
 $= 1 + \lceil \log_2 N \rceil$
- Total I/O cost
 $= 2N \cdot (\# \text{ of passes})$




Simplified 2-Way External Merge Sort

- In each pass, we read and write every page in the file.
- Number of passes
 $= 1 + \lceil \log_2 N \rceil$
- Total I/O cost
 $= 2N \cdot (\# \text{ of passes})$




Simplified 2-Way External Merge Sort

- In each pass, we read and write every page in the file.
- Number of passes
 $= 1 + \lceil \log_2 N \rceil$
- Total I/O cost
 $= 2N \cdot (\# \text{ of passes})$




Simplified 2-Way External Merge Sort

- In each pass, we read and write every page in the file.
- Number of passes
 $= 1 + \lceil \log_2 N \rceil$
- Total I/O cost
 $= 2N \cdot (\# \text{ of passes})$




Simplified 2-Way External Merge Sort

- In each pass, we read and write every page in the file.
- Number of passes
 $= 1 + \lceil \log_2 N \rceil$
- Total I/O cost
 $= 2N \cdot (\# \text{ of passes})$




Simplified 2-Way External Merge Sort

- In each pass, we read and write every page in the file.
- Number of passes
 $= 1 + \lceil \log_2 N \rceil$
- Total I/O cost
 $= 2N \cdot (\# \text{ of passes})$




Simplified 2-Way External Merge Sort

- In each pass, we read and write every page in the file.
- Number of passes
 $= 1 + \lceil \log_2 N \rceil$
- Total I/O cost
 $= 2N \cdot (\# \text{ of passes})$




Simplified 2-Way External Merge Sort

- In each pass, we read and write every page in the file.
- Number of passes
 $= 1 + \lceil \log_2 N \rceil$
- Total I/O cost
 $= 2N \cdot (\# \text{ of passes})$




Simplified 2-Way External Merge Sort

- In each pass, we read and write every page in the file.
- Number of passes
 $= 1 + \lceil \log_2 N \rceil$
- Total I/O cost
 $= 2N \cdot (\# \text{ of passes})$




Simplified 2-Way External Merge Sort

- In each pass, we read and write every page in the file.
- Number of passes
 $= 1 + \lceil \log_2 N \rceil$
- Total I/O cost
 $= 2N \cdot (\# \text{ of passes})$




Simplified 2-Way External Merge Sort

- In each pass, we read and write every page in the file.
- Number of passes
 $= 1 + \lceil \log_2 N \rceil$
- Total I/O cost
 $= 2N \cdot (\# \text{ of passes})$




Simplified 2-Way External Merge Sort

- In each pass, we read and write every page in the file.
- Number of passes
 $= 1 + \lceil \log_2 N \rceil$
- Total I/O cost
 $= 2N \cdot (\# \text{ of passes})$



Simplified 2-Way External Merge Sort

- In each pass, we read and write every page in the file.
- Number of passes
 $= 1 + \lceil \log_2 N \rceil$
- Total I/O cost
 $= 2N \cdot (\# \text{ of passes})$



Simplified 2-Way External Merge Sort

- This algorithm only requires three buffer pool pages to perform the sorting ($B=3$).
 - Two input pages, one output page
- But even if we have more buffer space available ($B>3$), it does not effectively utilize them if the worker must block on disk I/O...

General External Merge Sort

- **Pass #0**
 - Use B buffer pages
 - Produce $\lceil N / B \rceil$ sorted runs of size B
- **Pass #1,2,3,...**
 - Merge $B-1$ runs (i.e., K-way merge)
- Number of passes = $1 + \lceil \log_{B-1} \lceil N / B \rceil \rceil$
- Total I/O Cost = $2N \cdot (\# \text{ of passes})$

Example

- Determine how many passes it takes to sort 108 pages with 5 buffer pool pages: $N=108$, $B=5$

Example

- Determine how many passes it takes to sort 108 pages with 5 buffer pool pages: $N=108$, $B=5$
 - Pass #0: $[N / B] = [108 / 5] = 22$ sorted runs of 5 pages each (last run is only 3 pages).

Example

- Determine how many passes it takes to sort 108 pages with 5 buffer pool pages: $N=108$, $B=5$
 - **Pass #0:** $[N / B] = [108 / 5] = 22$ sorted runs of 5 pages each (last run is only 3 pages).
 - **Pass #1:** $[N' / B-1] = [22 / 4] = 6$ sorted runs of 20 pages each (last run is only 8 pages).

Example

- Determine how many passes it takes to sort 108 pages with 5 buffer pool pages: $N=108$, $B=5$
 - **Pass #0:** $[N / B] = [108 / 5] = 22$ sorted runs of 5 pages each (last run is only 3 pages).
 - **Pass #1:** $[N' / B-1] = [22 / 4] = 6$ sorted runs of 20 pages each (last run is only 8 pages).
 - **Pass #2:** $[N'' / B-1] = [6 / 4] = 2$ sorted runs, first one has 80 pages and second one has 28 pages.

Example

- Determine how many passes it takes to sort 108 pages with 5 buffer pool pages: $N=108$, $B=5$
 - **Pass #0:** $[N / B] = [108 / 5] = 22$ sorted runs of 5 pages each (last run is only 3 pages).
 - **Pass #1:** $[N' / B-1] = [22 / 4] = 6$ sorted runs of 20 pages each (last run is only 8 pages).
 - **Pass #2:** $[N'' / B-1] = [6 / 4] = 2$ sorted runs, first one has 80 pages and second one has 28 pages.
 - **Pass #3:** Sorted file of 108 pages.

Example


- Determine how many passes it takes to sort 108 pages with 5 buffer pool pages: $N=108$, $B=5$
 - Pass #0: $[N / B] = [108 / 5] = 22$ sorted runs of 5 pages each (last run is only 3 pages).
 - Pass #1: $[N' / B-1] = [22 / 4] = 6$ sorted runs of 20 pages each (last run is only 8 pages).
 - Pass #2: $[N'' / B-1] = [6 / 4] = 2$ sorted runs, first one has 80 pages and second one has 28 pages.
 - Pass #3: Sorted file of 108 pages.
- $1 + \lceil \log_{B-1} [N / B] \rceil = 1 + \lceil \log_4 22 \rceil = 1 + \lceil 2.229... \rceil = 4 \text{ passes}$

Double Buffering Optimization

- Prefetch the next run in the background and store it in a second buffer while the system is processing the current run.
 - Reduces the wait time for I/O requests at each step by continuously utilizing the disk.


Double Buffering

- Prefetch the next run in the background and store it in a second buffer while the system is processing the current run.
 - Overlap CPU and I/O operations




Double Buffering

- Prefetch the next run in the background and store it in a second buffer while the system is processing the current run.
 - Overlap CPU and I/O operations




Double Buffering

- Prefetch the next run in the background and store it in a second buffer while the system is processing the current run.
 - Overlap CPU and I/O operations




Double Buffering

- Prefetch the next run in the background and store it in a second buffer while the system is processing the current run.
 - Overlap CPU and I/O operations




Double Buffering

- Prefetch the next run in the background and store it in a second buffer while the system is processing the current run.
 - Overlap CPU and I/O operations




Double Buffering

- Prefetch the next run in the background and store it in a second buffer while the system is processing the current run.
 - Overlap CPU and I/O operations




Double Buffering

- Prefetch the next run in the background and store it in a second buffer while the system is processing the current run.
 - Overlap CPU and I/O operations




Double Buffering

- Prefetch the next run in the background and store it in a second buffer while the system is processing the current run.
 - Overlap CPU and I/O operations




Double Buffering

- Prefetch the next run in the background and store it in a second buffer while the system is processing the current run.
 - Overlap CPU and I/O operations




Double Buffering

- Prefetch the next run in the background and store it in a second buffer while the system is processing the current run.
 - Overlap CPU and I/O operations



Double Buffering

- Prefetch the next run in the background and store it in a second buffer while the system is processing the current run.
 - Overlap CPU and I/O operations





Double Buffering

- Prefetch the next run in the background and store it in a second buffer while the system is processing the current run.
 - Overlap CPU and I/O operations

Impact: reduces the effective “B” by half.

Reduces response time.



Comparison Optimizations


- **Approach #1: Code Specialization**
 - Instead of providing a comparison function as a pointer to sorting algorithm, create a hardcoded version of sort that is specific to a key type.
- **Approach #2: Suffix Truncation**
 - First compare a binary prefix of long `VARCHAR` keys instead of slower string comparison. Fallback to slower version if prefixes are equal.

Using B+Trees For Sorting

- If the table that must be sorted already has a B+Tree index on the sort attribute(s), then we can use that to accelerate sorting.
- Retrieve tuples in desired sort order by simply traversing the leaf pages of the tree.
- Cases to consider:
 - Clustered B+Tree
 - Unclustered B+Tree


Case #1: Clustered B+Tree

- Traverse to the left-most leaf page, and then retrieve tuples from all leaf pages.
- This is always better than external sorting because there is no computational cost, and all disk access is sequential.




Case #1: Clustered B+Tree

- Traverse to the left-most leaf page, and then retrieve tuples from all leaf pages.
- This is always better than external sorting because there is no computational cost, and all disk access is sequential.




Case #2: Unclustered B+Tree

- Chase each pointer to the page that contains the data.
- This is almost always a bad idea.
In general, one I/O per data record.



Case #2: Unclustered B+Tree

- Chase each pointer to the page that contains the data.
- This is almost always a bad idea.
In general, one I/O per data record.



Aggregations

Aggregations

- Collapse values for a single attribute from multiple tuples into a single scalar value.
- The DBMS needs a way to quickly find tuples with the same distinguishing attributes for grouping.
- Two implementation choices:
 - Sorting
 - Hashing

Sorting Aggregation

```
SELECT DISTINCT cid
  FROM enrolled
 WHERE grade IN ('B', 'C')
 ORDER BY cid
```

enrolled (sid, cid, grade)

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53666	15-721	C
53655	15-445	C

Sorting Aggregation

```
SELECT DISTINCT cid
  FROM enrolled
 WHERE grade IN ('B', 'C')
 ORDER BY cid
```

enrolled (sid, cid, grade)

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53666	15-721	C
53655	15-445	C



Filter

sid	cid	grade
53666	15-445	C
53688	15-826	B
53666	15-721	C
53655	15-445	C

Sorting Aggregation

```
SELECT DISTINCT cid
  FROM enrolled
 WHERE grade IN ('B', 'C')
 ORDER BY cid
```

enrolled (sid, cid, grade)

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53666	15-721	C
53655	15-445	C

Filter

sid	cid	grade
53666	15-445	C
53688	15-826	B
53666	15-721	C
53655	15-445	C

Remove
Columns

cid
15-445
15-826
15-721
15-445

Sorting Aggregation

```
SELECT DISTINCT cid
  FROM enrolled
 WHERE grade IN ('B', 'C')
 ORDER BY cid
```

enrolled (sid, cid, grade)

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53666	15-721	C
53655	15-445	C

Filter

sid	cid	grade
53666	15-445	C
53688	15-826	B
53666	15-721	C
53655	15-445	C

Remove
Columns



cid
15-445
15-826
15-721
15-445

Sort



cid
15-445
15-445
15-721
15-826

Sorting Aggregation

```
SELECT DISTINCT cid
  FROM enrolled
 WHERE grade IN ('B', 'C')
 ORDER BY cid
```

enrolled (sid, cid, grade)

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53666	15-721	C
53655	15-445	C


Filter

sid	cid	grade
53666	15-445	C
53688	15-826	B
53666	15-721	C
53655	15-445	C


*Remove
Columns*

cid
15-445
15-826
15-721
15-445

 
Sort

*Eliminate
Duplicates*

Sorting Aggregation

```
SELECT DISTINCT cid
  FROM enrolled
 WHERE grade IN ('B', 'C')
 ORDER BY cid
```

enrolled (sid, cid, grade)

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53666	15-721	C
53655	15-445	C

Filter

sid	cid	grade
53666	15-445	C
53688	15-826	B
53666	15-721	C
53655	15-445	C

Remove
Columns



cid
15-445
15-826
15-721
15-445

Sort



cid
15-445
15-445
15-721
15-826

Eliminate
Duplicates

Alternatives To Sorting

- What if we do not need the data to be ordered?
 - Forming groups in **GROUP BY** (no ordering)
 - Removing duplicates in **DISTINCT** (no ordering)
- Hashing is a better alternative in this scenario.
 - Only need to remove duplicates, no need for ordering.
 - Can be computationally cheaper than sorting.

Hashing Aggregate

- Populate an ephemeral hash table as the DBMS scans the table. For each record, check whether there is already an entry in the hash table:
 - **DISTINCT**: Discard duplicate
 - **GROUP BY**: Perform aggregate computation
- If everything fits in memory, then this is easy.
- If the DBMS must spill data to disk, then we need to be smarter...

External Hashing Aggregate

- **Phase #1 - Partition**
 - Divide tuples into buckets based on hash key
 - Write them out to disk when they get full
- **Phase #2 - ReHash**
 - Build in-memory hash table for each partition and compute the aggregation

Phase #1: Partition

- Use a hash function h_1 to split tuples into partitions on disk.
 - A partition is one or more pages that contain the set of keys with the same hash value.
 - Partitions are “spilled” to disk via output buffers.
- Assume that we have B buffers.
- We will use $B-1$ buffers for the partitions and 1 buffer for the input data.

Phase #1: Partition

```
SELECT DISTINCT cid
  FROM enrolled
 WHERE grade IN ('B', 'C')
```

enrolled (sid, cid, grade)

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53666	15-721	C
53655	15-445	C

:

Phase #1: Partition

```
SELECT DISTINCT cid
  FROM enrolled
 WHERE grade IN ('B', 'C')
```

enrolled (sid, cid, grade)

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53666	15-721	C
53655	15-445	C

:


Filter

sid	cid	grade
53666	15-445	C
53688	15-826	B
53666	15-721	C
53655	15-445	C

:

Phase #1: Partition

```
SELECT DISTINCT cid
  FROM enrolled
 WHERE grade IN ('B', 'C')
```

enrolled (sid, cid, grade)

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53666	15-721	C
53655	15-445	C

:

sid	cid	grade
53666	15-445	C
53688	15-826	B
53666	15-721	C
53655	15-445	C

:

Filter

Remove
Columns

cid
15-445
15-826
15-721
15-445

:

Phase #1: Partition

```
SELECT DISTINCT cid
  FROM enrolled
 WHERE grade IN ('B', 'C')
```

Filter

sid	cid	grade
53666	15-445	C
53688	15-826	B
53666	15-721	C
53655	15-445	C

⋮


**Remove
Columns**

enrolled (sid, cid, grade)

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53666	15-721	C
53655	15-445	C

⋮

B-1 partitions



Phase #1: Partition

```
SELECT DISTINCT cid
  FROM enrolled
 WHERE grade IN ('B', 'C')
```

Filter

sid	cid	grade
53666	15-445	C
53688	15-826	B
53666	15-721	C
53655	15-445	C

⋮

**Remove
Columns**

enrolled (sid, cid, grade)


sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53666	15-721	C
53655	15-445	C

⋮

B-1 partitions

cid
15-445
15-826
15-721
15-445

⋮



Phase #1: Partition

```
SELECT DISTINCT cid
  FROM enrolled
 WHERE grade IN ('B', 'C')
```

Filter

sid	cid	grade
53666	15-445	C
53688	15-826	B
53666	15-721	C
53655	15-445	C

⋮


**Remove
Columns**

enrolled (sid, cid, grade)

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53666	15-721	C
53655	15-445	C

⋮

B-1 partitions



Phase #2: ReHash

- For each partition on disk:
 - Read it into memory and build an in-memory hash table based on a second hash function h_2 .
 - Then go through each bucket of this hash table to bring together matching tuples.
- This assumes that each partition fits in memory.

Phase #2: ReHash

```
SELECT DISTINCT cid
  FROM enrolled
 WHERE grade IN ('B', 'C')
```

enrolled (sid, cid, grade)

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53666	15-721	C
53655	15-445	C

Phase #2: ReHash

```
SELECT DISTINCT cid
  FROM enrolled
 WHERE grade IN ('B', 'C')
```

Phase #1 Buckets

15-445	15-445
15-445	15-445
15-445	15-445
15-445	15-445

15-826
15-826

:


enrolled (sid, cid, grade)

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53666	15-721	C
53655	15-445	C

Phase #2: ReHash

```
SELECT DISTINCT cid
  FROM enrolled
 WHERE grade IN ('B', 'C')
```

Phase #1 Buckets




enrolled (sid, cid, grade)

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53666	15-721	C
53655	15-445	C

Phase #2: ReHash

```
SELECT DISTINCT cid
  FROM enrolled
 WHERE grade IN ('B', 'C')
```

Phase #1 Buckets



15-826
15-826

:


enrolled (sid, cid, grade)

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53666	15-721	C
53655	15-445	C

Phase #2: ReHash

```
SELECT DISTINCT cid
  FROM enrolled
 WHERE grade IN ('B', 'C')
```

Phase #1 Buckets




`enrolled (sid, cid, grade)`

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53666	15-721	C
53655	15-445	C

Phase #2: ReHash

```
SELECT DISTINCT cid
  FROM enrolled
 WHERE grade IN ('B', 'C')
```

Phase #1 Buckets




enrolled (sid, cid, grade)

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53666	15-721	C
53655	15-445	C

Phase #2: ReHash

```
SELECT DISTINCT cid
  FROM enrolled
 WHERE grade IN ('B', 'C')
```

Phase #1 Buckets



enrolled (sid, cid, grade)

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53666	15-721	C
53655	15-445	C


Final Result

cid
15-445
15-826

Phase #2: ReHash

```
SELECT DISTINCT cid
  FROM enrolled
 WHERE grade IN ('B', 'C')
```

Phase #1 Buckets



enrolled (sid, cid, grade)

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53666	15-721	C
53655	15-445	C


Final Result

cid
15-445
15-826

Phase #2: ReHash

```
SELECT DISTINCT cid
  FROM enrolled
 WHERE grade IN ('B', 'C')
```

Phase #1 Buckets



enrolled (sid, cid, grade)

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53666	15-721	C
53655	15-445	C


Final Result

cid
15-445
15-826

Phase #2: ReHash

```
SELECT DISTINCT cid
  FROM enrolled
 WHERE grade IN ('B', 'C')
```

Phase #1 Buckets



enrolled (sid, cid, grade)

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53666	15-721	C
53655	15-445	C

Final Result

cid
15-445
15-826
15-721

Hashing Summarization

- During the rehash phase, store pairs of the form **(GroupKey→RunningVal)**
- When we want to insert a new tuple into the hash table:
 - If we find a matching **GroupKey**, just update the **RunningVal** appropriately
 - Else insert a new **GroupKey→RunningVal**

Hashing Summarization

```
SELECT cid, AVG(s.gpa)
  FROM student AS s, enrolled AS e
 WHERE s.sid = e.sid
 GROUP BY cid
```

*Phase #1
Buckets*

15-445
15-445


15-826

:

15-721


Hashing Summarization

```
SELECT cid, AVG(s.gpa)
  FROM student AS s, enrolled AS e
 WHERE s.sid = e.sid
 GROUP BY cid
```




Hashing Summarization

```
SELECT cid, AVG(s.gpa)
  FROM student AS s, enrolled AS e
 WHERE s.sid = e.sid
 GROUP BY cid
```



Hashing Summarization

```
SELECT cid, AVG(s.gpa)
  FROM student AS s, enrolled AS e
 WHERE s.sid = e.sid
 GROUP BY cid
```




Hashing Summarization

```
SELECT cid, AVG(s.gpa)
  FROM student AS s, enrolled AS e
 WHERE s.sid = e.sid
 GROUP BY cid
```

Running Totals

AVG(col)	\rightarrow	(COUNT, SUM)
MIN(col)	\rightarrow	(MIN)
MAX(col)	\rightarrow	(MAX)
SUM(col)	\rightarrow	(SUM)
COUNT(col)	\rightarrow	(COUNT)



Hash Table


key	value
15-445	(2, 7.32)
15-826	(1, 3.33)
15-721	(1, 2.89)

Hashing Summarization

```
SELECT cid, AVG(s.gpa)
  FROM student AS s, enrolled AS e
 WHERE s.sid = e.sid
 GROUP BY cid
```

Running Totals

AVG(col)	→ (COUNT, SUM)
MIN(col)	→ (MIN)
MAX(col)	→ (MAX)
SUM(col)	→ (SUM)
COUNT(col)	→ (COUNT)



Hash Table

key	value
15-445	(2, 7.32)
15-826	(1, 3.33)
15-721	(1, 2.89)

Final Result


cid	AVG(gpa)
15-445	3.66
15-826	3.33
15-721	2.89

Conclusion

- Choice of sorting vs. hashing is subtle and depends on optimizations done in each case.
- We already discussed the optimizations for sorting:
 - Chunk I/O into large blocks to amortize costs
 - Double-buffering to overlap CPU and I/O

Next Lecture

- Nested Loop Join
- Sort-Merge Join
- Hash Join



Next Lecture

- Nested Loop Join
- Sort-Merge Join
- Hash Join

