

CARNEGIE MELLON UNIVERSITY  
COMPUTER SCIENCE DEPARTMENT  
15-445/645 – DATABASE SYSTEMS (FALL 2025)  
PROF. ANDY PAVLO

Homework #3 (by Will) – Solutions  
Due: **Sunday October 3rd, 2025 @ 11:59pm**

**IMPORTANT:**

- Enter all of your answers into **Gradescope by 11:59pm on Sunday October 3rd, 2025.**
- **Plagiarism:** Homework may be discussed with other students, but all homework is to be completed **individually**.

For your information:

- Graded out of **100** points; **5** questions total
- Rough time estimate:  $\approx$ 4-6 hours (1-1.5 hours for each question)
- Each part is all or nothing. There is no partial credit.

*Revision : 2025/10/06 10:58*

Question	Points	Score
Linear Hashing and Cuckoo Hashing	18	
Extendible Hashing	20	
B+Tree	27	
Bloom Filter	20	
Alternate Index Structures	15	
Total:	100	

**Question 1: Linear Hashing and Cuckoo Hashing.....[18 points]****Graded by:**For warmup, consider the following *Linear Probe Hashing* schema:

1. The table has a size of 4 slots, each slot can only contain one key-value pair.
  2. The hashing function is  
 $h_1(x) = x \% 4$ .
  3. When there is a conflict, it finds the next free slot to insert key-value pairs.
  4. The original table is empty.
  5. Uses a tombstone when deleting a key.
- (a) [2 points] Insert key/value pairs (3, C) and (8, D). For (3, C), “3” is the key and “C” is the value. Select the value in each entry of the resulting table.
- i. Entry 0 (key % 4 = 0)  C  D  Empty
  - ii. Entry 1 (key % 4 = 1)  C  D  Empty
  - iii. Entry 2 (key % 4 = 2)  C  D  Empty
  - iv. Entry 3 (key % 4 = 3)  C  D  Empty

**Solution:** C is inserted into Entry 3, and D is inserted into Entry 0.

- (b) [2 points] After the changes from part (a), delete (8, D), insert key-value (0, E), insert (4, F), and lastly delete (3, C). Select the value in each entry of the resulting table.

- i. Entry 0 (key % 4 = 0)  Tombstone  C  D  E  F  Empty
- ii. Entry 1 (key % 4 = 1)  Tombstone  C  D  E  F  Empty
- iii. Entry 2 (key % 4 = 2)  Tombstone  C  D  E  F  Empty
- iv. Entry 3 (key % 4 = 3)  Tombstone  C  D  E  F  Empty

**Solution:** D is first deleted, which inserts a tombstone into Entry 0. E is then inserted into Entry 0 (since there is nothing there). Then, F is attempted to be inserted into Entry 0, but since it's occupied by E, F is inserted into Entry 1 instead. Finally C is deleted which leaves a tombstone in Entry 3.

Consider the following *Cuckoo Hashing* schema:

1. Both tables have a size of 4.
2. The hashing function of the first table returns the fourth and third least significant bits:  
 $h_1(x) = (x \gg 2) \& 0b11$ .
3. The hashing function of the second table returns the least significant two bits:  
 $h_2(x) = x \& 0b11$ .
4. When inserting, try table 1 first.
5. When replacement is necessary, first select an element in the *first* table.
6. The original entries in the table are shown below.

Table 1	Entry 0	Entry 1	Entry 2	Entry 3
Keys	-	-	-	445
Table 2	Entry 0	Entry 1	Entry 2	Entry 3
Keys	-	-	-	15

Figure 1: Initial contents of the hash tables.

- (a) [2 points] Select the sequence of insert operations that results in the initial state.

Insert 445, Insert 15    Insert 15, Insert 445    None of the above

**Solution:** 445 is inserted into Table 1 0b11 based on  $h_1$ , and 15 experiences a collision and is hashed to Table 2 0b11 based on  $h_2$ .

- (b) Starting from the initial contents, insert key 463 and then insert 789. Select the values in the resulting two tables.

i. Table 1

- $\alpha)$  [1 point] Entry 0 (0b00)  445  15  463  789  Empty
- $\beta)$  [1 point] Entry 1 (0b01)  445  15  463  789  Empty
- $\gamma)$  [1 point] Entry 2 (0b10)  445  15  463  789  Empty
- $\delta)$  [1 point] Entry 3 (0b11)  445  15  463  789  Empty

ii. Table 2

- $\alpha)$  [1 point] Entry 0 (0b00)  445  15  463  789  Empty
- $\beta)$  [1 point] Entry 1 (0b01)  445  15  463  789  Empty
- $\gamma)$  [1 point] Entry 2 (0b10)  445  15  463  789  Empty
- $\delta)$  [1 point] Entry 3 (0b11)  445  15  463  789  Empty

**Solution:** 463 tries to insert into both tables but conflicts with both, so 463 inserts into Entry 3 of Table 1, replacing 445. 445 is then rehashed into Table 2 Entry 1. 789 then inserts into Table 1 Entry 1.

- (c) [4 points] Consider completely empty tables using the same two hash functions. Select which sequence of insertions below will cause an infinite loop.

- [1, 5, 9, 13]
- [2, 7, 10, 14]
- [4, 10, 11, 15]
- [1, 9, 17, 25]
- None of the above

**Solution:** The sequence [1, 9, 17, 25] will cause an infinite loop because 4 keys map to 3 entries.

**Question 2: Extendible Hashing.....[20 points]****Graded by:**

Consider an extendible hashing structure such that:

- Each bucket can hold up to two records.
- The hashing function uses the lowest  $g$  bits, where  $g$  is the global depth.
- A new extendible hashing structure is initialized with  $g = 0$  and one empty bucket
- If multiple keys are provided in a question, assume they are inserted one after the other from left to right.

(a) Starting from an empty table, insert keys 10, 20.

- i. [1 point] What is the global depth of the resulting table?

0    1    2    3    4    None of the above

**Solution:** No split has occurred yet because the first bucket (on initialization) can hold 2 arbitrary values. Thus global depth is same as its initial value of 0.

- ii. [1 point] What is the local depth of the bucket containing 10?

0    1    2    3    4    None of the above

**Solution:** There is only one bucket (created on initialization), and it holds both 10 and 20. Since no split has occurred yet, the bucket has local depth  $d = 0$ .

(b) Starting from the result in (a), you insert keys 3, 4.

- i. [2 points] What is the global depth of the resulting table?

0    1    2    3    4    None of the above

**Solution:** After the inserts and splits, the table looks like the following:

Global depth = 2

b0 = 4,20 // at local depth 2

b2 = 10 // at local depth 2

b1,3 = 3 // at local depth 1

- ii. [2 points] What are the local depths of the buckets for each key?

3 (Depth 1), 4 (Depth 1), 10 (Depth 1), 20 (Depth 1)

3 (Depth 1), 4 (Depth 3), 10 (Depth 3), 20 (Depth 3)

3 (Depth 3), 4 (Depth 1), 10 (Depth 3), 20 (Depth 2)

3 (Depth 1), 4 (Depth 2), 10 (Depth 2), 20 (Depth 2)

3 (Depth 2), 4 (Depth 2), 10 (Depth 2), 20 (Depth 2)

None of the above

**Solution:** See the previous solution for an explanation.

(c) Starting from the result in (b), you insert keys 0, 12.

- i. [2 points] What is the global depth of the resulting table?

0    1    2    3    4    None of the above

**Solution:** 0 inserts into b0 and causes a split. 12 inserts into b4 and causes a split.

The updated table looks as follows: Global depth = 4

b0, 8 = 0 // at local depth 3

b2, 6, 10, 14 = 10 // at local depth 2

b1, 3, 5, 7, 9, 11, 13, 15 = 3 // at local depth 1

b4 = 4, 20 // at local depth 4

b12 = 12 // at local depth 4

- ii. [2 points] What are the local depths of the buckets for each new key?

- 0 (Depth (3), 12 (Depth 3))
- 0 (Depth (3), 12 (Depth 4))**
- 0 (Depth (4), 12 (Depth 3))
- 0 (Depth (4), 12 (Depth 4))
- 0 (Depth (2), 12 (Depth 4))
- None of the above

**Solution:** See the previous solution for an explanation.

- (d) [3 points] Starting from (c)'s result, which key(s), if inserted next, will **not** cause a split?

- 95**
- 100
- 38**
- 36
- None of the above

**Solution:** Only the selected values hash to buckets from c that are not full.

- (e) [3 points] Starting from the result in (c), which key(s), if inserted next, will cause a split and increase the table's global depth?

- 1
- 5
- 34
- 68**
- None of the above

**Solution:** The values must hash to b4, since it is the only full bucket whose local depth is equal to the global depth.


- (f) [4 points] Starting from an empty table, insert keys 64, 128, 256, 512. What is the global depth of the resulting table?

- 4
- 5
- 6
- 7
- 8**
- $\geq 9$

**Solution:** Since each bucket can hold at most two keys, three or more keys cannot hash to the same bucket without causing splits. When  $g = 8$ , 64 and 128 will each be mapped to their own bin. 256 and 512 will share the same bin.

**Question 3: B+Tree.....[27 points]****Graded by:**

Consider the following B+tree.

Figure 2: B+ Tree of order  $d = 4$  and height  $h = 2$ .


When answering the following questions, be sure to follow the procedures described in class and in your textbook. You can make the following assumptions:

- A left pointer in an internal node guides towards keys  $<$  than its corresponding key, while a right pointer guides towards keys  $\geq$ .
- A leaf node underflows when the number of **keys** goes below  $\lceil \frac{d-1}{2} \rceil$ .
- An internal node underflows when the number of **pointers** goes below  $\lceil \frac{d}{2} \rceil$ .


Note that B+ tree diagrams for this problem omit leaf pointers for convenience. The leaves of actual B+ trees are linked together via pointers, forming a singly linked list allowing for quick traversal through all keys.

(a) [4 points] Insert  $0^*$  into the B+tree. Select the resulting tree.


A)




B)



C)




D)




**Solution:** Inserting  $0^*$  adds one element in the left-most leaf. It should not cause any splits or merges.

(b) [5 points] Starting with the tree that results from (a), insert  $35^*$  and then  $45^*$ . Select the resulting tree.


A)




B)



■ C)




D)




**Solution:** Inserting  $35^*$  fills in the remaining space of the second leaf node (from the left). After inserting  $45^*$ , the second leaf node splits. As the root-level node is full, the root-level also splits.

(c) [8 points] Starting with the tree that results from (b), deletes  $80^*$  and then  $20^*$ . Select the resulting tree.


A)




B)



C)



D)



**Solution:** Deleting  $80^*$  causes the third leaf node to underflow and causes the right two leaf nodes to merge. After merging, the right internal node underflows, which triggers recursive merging.

Then deleting  $20^*$  causes the second leaf node to underflow. This leads the second and third leaf node to then merge into  $(30, 35, 45)$ .

- (d) i. [2 points] Threads must release their latches in the order they were acquired (i.e., FIFO).

True    False

**Solution:** Threads can release latches in any order.

- ii. [2 points] Under optimistic latch coupling, write threads only take the write latch on the root when they restart.

True    False

**Solution:** Using the optimistic latch coupling/crabbing scheme that we discussed in class, the thread will take a write latch on the root if it needs to restart.

- iii. [2 points] A DBMS primarily executes OLTP queries but periodically will execute OLAP style queries (e.g., analytics, book-keeping). The DBMS will benefit from using one buffer pool for inner node pages and a different buffer pool for leaf pages / table pages.

True    False

**Solution:** Because the DBMS is aware of whether a page is for a leaf or an inner node, and because B+Tree transformations never change a leaf into an inner node or vice-versa, it's straightforward for a DBMS to use a different buffer pool for inner node pages than for leaf node pages. Such a configuration would help prevent index leaf scans from sequentially flooding the buffer pool and harming the performance of OLTP-style queries on the same index.

- iv. [2 points] Under optimistic latch crabbing, read-only thread can drop its latch on the current page before acquiring the latch on the next page (e.g., child, sibling).

True    False

**Solution:** During traversal, a reader temporarily needs to hold a latch on both the parent and child (or two siblings in a leaf node scan) before releasing the latch on the parent page.

- v. [2 points] “No-Wait” mode for acquiring sibling latches prevents deadlock by allowing the acquirer to fail if another reader already owns the lock.

True    False

**Solution:** A “no-wait” mode prevents threads from getting stuck.

**Question 4: Bloom Filter.....[20 points]****Graded by:**

Assume that we have a bloom filter that is used to register words relating to Pokemon. The filter uses two hash functions  $h_1$  and  $h_2$  which hash the following words to the following values:

input	$h_1$	$h_2$
“Pikachu”	4721	8395
“Bulbasaur”	1568	4207
“Charmander”	9034	2876
“Squirtle”	6412	7559

- (a) [6 points] Suppose the filter has 8 bits initially set to 0:

bit 0	bit 1	bit 2	bit 3	bit 4	bit 5	bit 6	bit 7
0	0	0	0	0	0	0	0

Which bits will be set to 1 after “Pikachu” and “Bulbasaur” have been inserted?

- 0     1     2     3     4     5     6     7

**Solution:** Because the filter has 8 bits, we take the modulo of the hashed output and 8.

- (b) Suppose the filter has 8 bits set to the following values:

bit 0	bit 1	bit 2	bit 3	bit 4	bit 5	bit 6	bit 7
0	0	1	1	1	1	0	0

- i. [4 points] What will we learn using the above filter if we lookup “Charmander”?
- Charmander has been inserted
  - Charmander has not been inserted
  - Charmander may have been inserted
  - Not possible to know

**Solution:**  $9034 \bmod 8 = 2$ ;  $2876 \bmod 8 = 4$

Because both bits are 1, the filter will return true, meaning Charmander may have been inserted.

- ii. [4 points] What will we learn if we lookup “Squirtle”?
- Squirtle has been inserted
  - Squirtle has not been inserted
  - Squirtle may have been inserted
  - Not possible to know

**Solution:**  $6412 \bmod 8 = 4$ ;  $7559 \bmod 8 = 7$

Because bit 7 is 0, the filter will return false, so Squirtle has not been inserted.

- (c) [6 points] A colleague is interviewing a candidate and would like to first test your knowledge of bloom filters. The colleague has a list of prepared statements and would like you to identify which of them are true. Select all true statements.
- Bloom filters are more effective than hash indexes for exact-match (or lookup) queries.
  - Using more hash functions will *always* lower a bloom filter's false positive rate.
  - Bloom filters can reduce disk I/Os.**
  - Add and lookup operations on bloom filters are parallelizable.**
  - All of the above.

**Solution:** Using more hash functions can increase the likelihood of false positives due to overlapping bits.

For exact-match queries, query execution is better off using a hash index.

**Question 5: Alternate Index Structures ..... [15 points]****Graded by:**

- (a) **[5 points]** Your team is considering using a **Radix Tree** for indexing in a new database system. They consulted a large language model for some factual statements about Radix Trees but are unsure about the accuracy of the model's responses. They have asked you to identify all factually correct statements.

■ **Radix Trees are efficient for prefix queries.**

- Radix Trees support efficient substring searches (e.g., LIKE "%?%").
- Radix Trees require re-balancing after every insertion or deletion.

■ **The levels of a radix tree have no node requirements.**

■ **For datasets with lots of common prefixes, radix trees can use less space than B+Trees.**

- None of the above.

**Solution:** Radix Trees are indeed efficient for prefix queries.

Radix Trees do not support efficient substring searches; they are optimized for prefix matching.

Radix Trees do not require re-balancing like balanced trees (e.g., AVL or Red-Black Trees).

There is no requirement about the size of each level in a Radix Tree.

They can be more space-efficient than B+Trees for certain datasets, especially when keys share common prefixes.

- (b) **[5 points]** You are discussing index structures with a colleague. They want to compare **B+Trees**, **Skip Lists**, **Radix Trees**, and **Inverted Indexes**. Select all the true statements.

■ **It is *on average* cheaper to update skip lists than B+Trees.**

- B+Trees and Skip Lists both guarantee logarithmic complexity for lookups.
- B+Trees and Inverted Indexes are both efficient at handling substring (e.g., LIKE "%?%") queries.
- Skip Lists perform better than B+Trees for range queries.
- None of the above.

**Solution:** B+Trees are generally more expensive to update than Skip Lists due to the need for node splits and merges.

Skip Lists have **approximate** logarithmic complexity for lookups, but it is not guaranteed.

B+Trees are not efficient for substring queries.

B+Trees generally perform better than Skip Lists for range queries due to their sequential access patterns.

- (c) **[5 points]** Suppose you are trying to run the following query:

```
SELECT * FROM PRODUCTS WHERE description LIKE '%laptop%';
```

Assume that there is a non-clustering B+Tree index on description. Your query is running slowly. Which of the following choices (if any) would make this query go faster?

- Drop the index and build a **bloom filter** on description.
- Replace the index with a **hash index** on description.
- Replace the non-clustering B+Tree with a **clustering** B+Tree index on description.
- Replace the index with a **radix tree** on description.
- None of the above.**

**Solution:** All of these options would not substantially speed up the query. If such queries dominate the workload, the best approach would be to invest in an **inverted index**, which is specifically designed for text-based searches like the one in the query.