



香港中文大學(深圳)  
The Chinese University of Hong Kong, Shenzhen



# CSC3170

# Tutorial 9

School of Data Science  
The Chinese University of Hong Kong, Shenzhen

# Outline

- Join Algorithm
  - Block Nested Loop Join
  - Index Nested Loop Join
  - Sort-Merge Join
  - Hash Join
- Exercise

# Block Nested Loop Join

- If we have  $B$  buffers available:
  - Use  $B-2$  buffers for each block of the outer table.
  - Use one buffer for the inner table, one buffer for output.

```

foreach  $B - 2$  pages  $p_R \in R$ :
  foreach page  $p_S \in S$ :
    foreach tuple  $r \in B - 2$  pages:
      foreach tuple  $s \in p_s$ :
        if  $r$  and  $s$  match then emit
  
```

$M$  pages       $N$  pages

$m$  tuples       $n$  tuples

id	name
600	MethodMan
200	GZA
100	Andy
300	ODB
500	RZA
700	Ghostface
400	Raekwon

id	value	cdate
100	2222	11/4/2024
500	7777	11/4/2024
400	6666	11/4/2024
100	9999	11/4/2024
200	8888	11/4/2024

- Cost:  $M + (\lceil M / (B-2) \rceil \cdot N)$

# Index Nested Loop Join

```

foreach tuple  $r \in R$ :
    foreach tuple  $s \in \text{Index}(r_i = s_j)$ :
        if  $r$  and  $s$  match then emit
  
```

require an existing index of inner table

$M$  pages  
 $m$  tuples

R(id, name)	
id	name
600	MethodMan
200	GZA
100	Andy
300	ODB
500	RZA
700	Ghostface
400	Raekwon

S(id, value, cdate)		
id	value	cdate
100	2222	11/4/2024
500	7777	11/4/2024
400	6666	11/4/2024
100	9999	11/4/2024
200	8888	11/4/2024

Index(S.id)



$N$  pages  
 $n$  tuples

- Assume the cost of each index probe is some constant  $C$  per tuple.
- Cost:**  $M + (m \cdot C)$

# Sort–Merge Join

- **Phase #1: Sort**

- Sort both tables on the join key(s).
- You can use any appropriate sort algorithm
- These phases are distinct from the sort/merge phases of an external merge sort, from the previous class

- **Phase #2: Merge**

- Step through the two sorted tables with cursors and emit matching tuples.
- May need to backtrack depending on the join type.

# Sort–Merge Join

```

sort R,S on join keys
cursorR ← Rsorted, cursorS ← Ssorted
while cursorR and cursorS:
    if cursorR > cursorS:
        increment cursorS
    if cursorR < cursorS:
        increment cursorR
        backtrack cursorS (if necessary)
    elif cursorR and cursorS match:
        emit
        increment cursorS
    
```

```

do{
    if (!mark) {
        while(r<s) {advance r}
        while(s<r) {advance s}
        mark = s
    }
    if (r==s) {
        emit
        advance s
    }
    else{
        reset s to mark
        advance r
        mark = NULL
    }
}

```

# Sort–Merge Join

R(id, name)

id	name
600	MethodMan
200	GZA
100	Andy
300	ODB
500	RZA
700	Ghostface
200	GZA
400	Raekwon

S(id, value, cdate)

id	value	cdate
100	2222	11/4/2024
500	7777	11/4/2024
400	6666	11/4/2024
100	9999	11/4/2024
200	8888	11/4/2024

```
SELECT R.id, S.cdate
  FROM R JOIN S
  ON R.id = S.id
 WHERE S.value > 100
```

# Sort–Merge Join

- **Sort Cost (R):**  $2M \cdot (1 + \lceil \log_{B-1} [M / B] \rceil)$
- **Sort Cost (S):**  $2N \cdot (1 + \lceil \log_{B-1} [N / B] \rceil)$
- **Merge Cost:**  $(M + N)$
- **Total Cost: Sort + Merge**

Worst Case: join attribute of all the tuples in both relations contains the same value

Merge Cost:  $M \times N$

# Hash Join

- **Phase #1: Build**

- Scan the outer relation and populate a hash table using the hash function  $h_1$  on the join attributes.
- We can use any hash table that we discussed before but in practice linear probing works the best.

- **Phase #2: Probe**

- Scan the inner relation and use  $h_1$  on each tuple to jump to a location in the hash table and find a matching tuple.

# Hash Join–Grace Hash Join


Deal the problem that the tables can not fit in the memory

- **Partition Phase:** Hash both tables on the join attribute into partitions.
- **Probe Phase:** Compares tuples in corresponding partitions for each table.

# Hash Join–Grace Hash Join

## Partition Phase


- Hash  $R$  into  $k$  buckets.
- Hash  $S$  into  $k$  buckets with same hash function.
- Write buckets to disk when they get full.



# Hash Join–Grace Hash Join

## Probe Phase

- Read corresponding partitions into memory one pair at a time, hash join their contents.



# Hash Join–Grace Hash Join

- If we do not need recursive partitioning:
  - Cost:  $3(M + N)$
- **Partition phase:**
  - Read+write both tables
  - $2(M+N)$  I/Os
- **Probe phase:**
  - Read both tables (in total, one partition at a time)
  - $M+N$  I/Os

The algorithm is **recursively** applied when one or more of the partitions still does not fit into the available memory.

# Exercise

Consider relations X and Y to be joined on the common attribute. There are

- $B = 500$  pages in the buffer
  - Table X spans  $M = 2000$  pages with 40 tuples per page
  - Table Y spans  $N = 600$  pages with 450 tuples per page
1. What is the I/O cost of a block nested loop join with Y as the outer relation and X as the inner relation?

# Exercise

Consider relations X and Y to be joined on the common attribute. There are

- $B = 500$  pages in the buffer
  - Table X spans  $M = 2000$  pages with 40 tuples per page
  - Table Y spans  $N = 600$  pages with 450 tuples per page
1. What is the I/O cost of a block nested loop join with Y as the outer relation and X as the inner relation?

$$N + \lceil N/(B - 2) \rceil \cdot M = 600 + 2 \times 2000 = 4600$$

# Exercise

Consider relations X and Y to be joined on the common attribute. There are

- $B = 500$  pages in the buffer
  - Table X spans  $M = 2000$  pages with 40 tuples per page
  - Table Y spans  $N = 600$  pages with 450 tuples per page
2. What is the I/O cost of a Grace hash join with X as the outer relation and Y as the inner relation?

# Exercise

Consider relations X and Y to be joined on the common attribute. There are

- $B = 500$  pages in the buffer
  - Table X spans  $M = 2000$  pages with 40 tuples per page
  - Table Y spans  $N = 600$  pages with 450 tuples per page
2. What is the I/O cost of a Grace hash join with X as the outer relation and Y as the inner relation?

$$3(M + N) = 3 \times 2600 = 7800$$

# Q&A