



香港中文大學(深圳)
The Chinese University of Hong Kong, Shenzhen



CSC3170 Tutorial 5

School of Data Science
The Chinese University of Hong Kong, Shenzhen

Outline

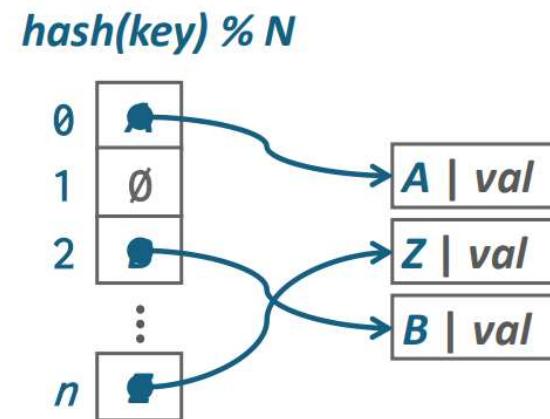
- Hash Tables
 - Concepts
 - Exercise

Hash Tables

- A hash table implements an unordered associative array that maps keys to values.
- It uses a hash function to compute an offset into this array for a given key, from which the desired value can be found.

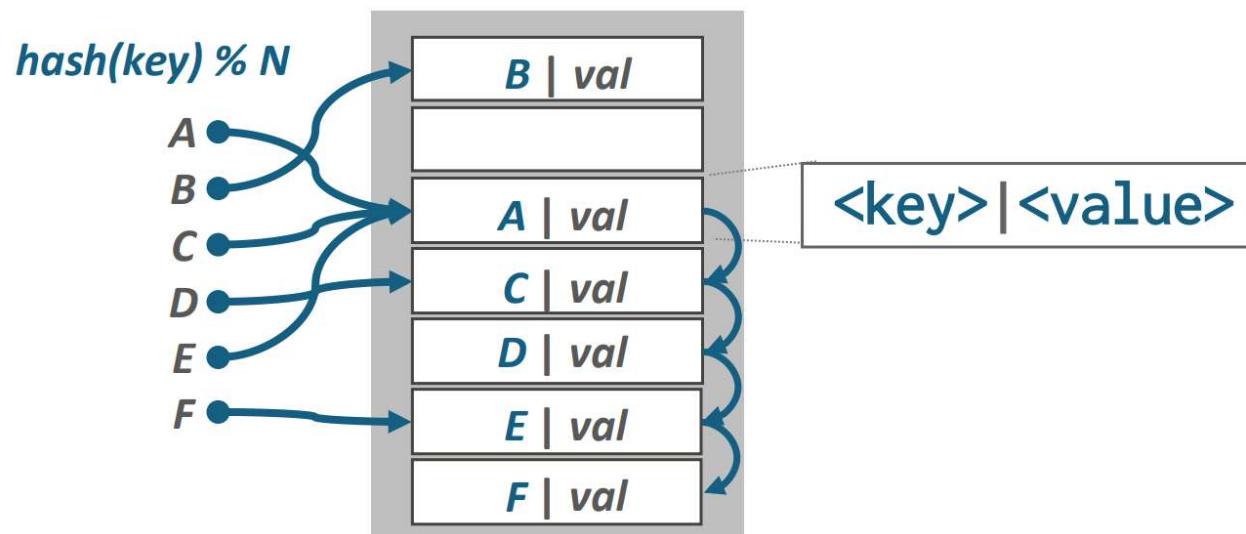
Static Hash Table

- Allocate a giant array that has one slot for every element you need to store.
- Hash function: $\text{key} \% N$
- Unrealistic assumptions



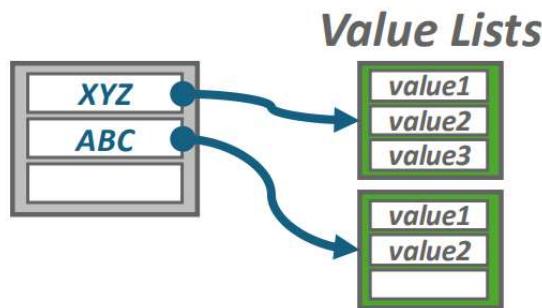
Linear Probe Hashing

- Similar to the static hash table.
- Resolve collisions by linearly searching for the next free slot in the table.



Non-Unique keys

- Separated Linked List
 - Store values in separate storage area of each key.
 - For instance, we have value lists for tuples with key = “XYZ” and



- Redundant keys
 - store duplicate key entries together in the hash table.

Cuckoo Hashing

- Use multiple hash functions to find multiple locations in the hash table to insert records.
 - On insert, check multiple locations and pick the one that is empty.
 - If no location is available, evict the element from one of them and then rehash it find a new location.
- Look-ups and deletions are always $O(1)$ because only one location per hash table is checked.

Exercise

Consider the following cuckoo hashing schema:

1. Both tables have a size of 4.
2. The hashing function of the first table returns the fourth and third least significant bits:
$$h_1(x) = (x \gg 2) \& \text{0b}11.$$
3. The hashing function of the second table returns the least significant two bits:
$$h_2(x) = x \& \text{0b}11.$$
4. When inserting, try table 1 first.
5. When replacement is necessary, first select an element in the second table.
6. The original entries in the table are shown in the figure below.

Exercise

- Which elements is first inserted, 11 or 9?
- Let's now do some insertion and deletion:
 - insert 16
 - delete 11
 - insert 17
 - insert 10
 - insert 33
 - delete 16
- After those operations, what is the smallest key that potentially cause an infinite loop?

Table 1



Table 2

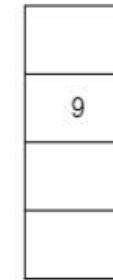


Figure 1: Initial contents of the hash tables.

Q&A